


Recipes for Object-Oriented and Functional Programming



Scala Cookbook

O'REILLY®

Alvin Alexander

Scala Cookbook

Save time and trouble when using Scala to build object-oriented, functional, and concurrent applications. With more than 250 ready-to-use recipes and 700 code examples, this comprehensive cookbook covers the most common problems you'll encounter when using the Scala language, libraries, and tools. It's ideal not only for experienced Scala developers, but also for programmers learning to use this JVM language.

Author Alvin Alexander (creator of DevDaily.com) provides solutions based on his experience using Scala for highly scalable, component-based applications that support concurrency and distribution. Packed with real-world scenarios, this book provides recipes for:

- Strings, numeric types, and control structures
- Classes, methods, objects, traits, and packaging
- Functional programming in a variety of situations
- Collections covering Scala's wealth of classes and methods
- Concurrency, using the Akka Actors library
- Using the Scala REPL and the Simple Build Tool (SBT)
- Web services on both the client and server sides
- Interacting with SQL and NoSQL databases
- Best practices in Scala development

Alvin Alexander holds a degree in Aerospace Engineering from Texas A&M University, and has been a Java/OOP programmer, instructor, and mentor since 1999. He's the founder of the Valley Programming software consulting business, and provides programming tutorials for Java, Ruby, Scala, and other topics at his website, alvinalexander.com.

US \$54.99

CAN \$57.99

ISBN: 978-1-449-33961-6



Twitter: @oreillymedia
facebook.com/oreilly

O'REILLY[®]
oreilly.com

Scala Cookbook

Alvin Alexander

O'REILLY®

Beijing • Cambridge • Farnham • Köln • Sebastopol • Tokyo

Scala Cookbook

by Alvin Alexander

Copyright © 2013 Alvin Alexander. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://my.safaribooksonline.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Editor: Courtney Nash

Production Editor: Rachel Steely

Copyeditor: Kim Cofer

Proofreader: Linley Dolby

Indexer: Ellen Troutman

Cover Designer: Karen Montgomery

Interior Designer: David Futato

Illustrator: Rebecca Demarest

August 2013: First Edition

Revision History for the First Edition:

2013-07-30: First release

See <http://oreilly.com/catalog/errata.csp?isbn=9781449339616> for release details.

Nutshell Handbook, the Nutshell Handbook logo, and the O'Reilly logo are registered trademarks of O'Reilly Media, Inc. *Scala Cookbook*, the image of a long-beaked echidna, and related trade dress are trademarks of O'Reilly Media, Inc.

Many of the designations used by manufacturers and sellers to distinguish their products are claimed as trademarks. Where those designations appear in this book, and O'Reilly Media, Inc., was aware of a trademark claim, the designations have been printed in caps or initial caps.

While every precaution has been taken in the preparation of this book, the publisher and author assume no responsibility for errors or omissions, or for damages resulting from the use of the information contained herein.

ISBN: 978-1-449-33961-6

[LSI]

For my mom, who loves cookbooks.

Table of Contents

Preface.....	xiii
1. Strings.....	1
1.1. Testing String Equality	4
1.2. Creating Multiline Strings	6
1.3. Splitting Strings	8
1.4. Substituting Variables into Strings	9
1.5. Processing a String One Character at a Time	13
1.6. Finding Patterns in Strings	18
1.7. Replacing Patterns in Strings	21
1.8. Extracting Parts of a String That Match Patterns	22
1.9. Accessing a Character in a String	24
1.10. Add Your Own Methods to the String Class	25
2. Numbers.....	31
2.1. Parsing a Number from a String	32
2.2. Converting Between Numeric Types (Casting)	36
2.3. Overriding the Default Numeric Type	37
2.4. Replacements for ++ and --	39
2.5. Comparing Floating-Point Numbers	41
2.6. Handling Very Large Numbers	43
2.7. Generating Random Numbers	45
2.8. Creating a Range, List, or Array of Numbers	47
2.9. Formatting Numbers and Currency	49
3. Control Structures.....	53
3.1. Looping with for and foreach	54
3.2. Using for Loops with Multiple Counters	60
3.3. Using a for Loop with Embedded if Statements (Guards)	62
3.4. Creating a for Comprehension (for/yield Combination)	63

3.5. Implementing break and continue	65
3.6. Using the if Construct Like a Ternary Operator	71
3.7. Using a Match Expression Like a switch Statement	72
3.8. Matching Multiple Conditions with One Case Statement	76
3.9. Assigning the Result of a Match Expression to a Variable	77
3.10. Accessing the Value of the Default Case in a Match Expression	78
3.11. Using Pattern Matching in Match Expressions	79
3.12. Using Case Classes in Match Expressions	86
3.13. Adding if Expressions (Guards) to Case Statements	87
3.14. Using a Match Expression Instead of isInstanceOf	88
3.15. Working with a List in a Match Expression	89
3.16. Matching One or More Exceptions with try/catch	91
3.17. Declaring a Variable Before Using It in a try/catch/finally Block	92
3.18. Creating Your Own Control Structures	95
4. Classes and Properties.....	99
4.1. Creating a Primary Constructor	100
4.2. Controlling the Visibility of Constructor Fields	104
4.3. Defining Auxiliary Constructors	108
4.4. Defining a Private Primary Constructor	112
4.5. Providing Default Values for Constructor Parameters	114
4.6. Overriding Default Accessors and Mutators	116
4.7. Preventing Getter and Setter Methods from Being Generated	119
4.8. Assigning a Field to a Block or Function	121
4.9. Setting Uninitialized var Field Types	122
4.10. Handling Constructor Parameters When Extending a Class	124
4.11. Calling a Superclass Constructor	127
4.12. When to Use an Abstract Class	129
4.13. Defining Properties in an Abstract Base Class (or Trait)	131
4.14. Generating Boilerplate Code with Case Classes	136
4.15. Defining an equals Method (Object Equality)	140
4.16. Creating Inner Classes	143
5. Methods.....	147
5.1. Controlling Method Scope	148
5.2. Calling a Method on a Superclass	152
5.3. Setting Default Values for Method Parameters	154
5.4. Using Parameter Names When Calling a Method	157
5.5. Defining a Method That Returns Multiple Items (Tuples)	159
5.6. Forcing Callers to Leave Parentheses off Accessor Methods	161
5.7. Creating Methods That Take Variable-Argument Fields	163
5.8. Declaring That a Method Can Throw an Exception	165

5.9. Supporting a Fluent Style of Programming	167
6. Objects.....	171
6.1. Object Casting	172
6.2. The Scala Equivalent of Java's .class	174
6.3. Determining the Class of an Object	174
6.4. Launching an Application with an Object	176
6.5. Creating Singletons with object	178
6.6. Creating Static Members with Companion Objects	180
6.7. Putting Common Code in Package Objects	182
6.8. Creating Object Instances Without Using the new Keyword	185
6.9. Implement the Factory Method in Scala with apply	189
7. Packaging and Imports.....	191
7.1. Packaging with the Curly Braces Style Notation	192
7.2. Importing One or More Members	193
7.3. Renaming Members on Import	195
7.4. Hiding a Class During the Import Process	196
7.5. Using Static Imports	197
7.6. Using Import Statements Anywhere	199
8. Traits.....	203
8.1. Using a Trait as an Interface	203
8.2. Using Abstract and Concrete Fields in Traits	206
8.3. Using a Trait Like an Abstract Class	207
8.4. Using Traits as Simple Mixins	208
8.5. Limiting Which Classes Can Use a Trait by Inheritance	209
8.6. Marking Traits So They Can Only Be Used by Subclasses of a Certain Type	211
8.7. Ensuring a Trait Can Only Be Added to a Type That Has a Specific Method	213
8.8. Adding a Trait to an Object Instance	215
8.9. Extending a Java Interface Like a Trait	216
9. Functional Programming.....	217
9.1. Using Function Literals (Anonymous Functions)	218
9.2. Using Functions as Variables	219
9.3. Defining a Method That Accepts a Simple Function Parameter	223
9.4. More Complex Functions	226
9.5. Using Closures	229
9.6. Using Partially Applied Functions	234
9.7. Creating a Function That Returns a Function	236

9.8. Creating Partial Functions	238
9.9. A Real-World Example	242
10. Collections.....	245
10.1. Understanding the Collections Hierarchy	246
10.2. Choosing a Collection Class	250
10.3. Choosing a Collection Method to Solve a Problem	255
10.4. Understanding the Performance of Collections	261
10.5. Declaring a Type When Creating a Collection	264
10.6. Understanding Mutable Variables with Immutable Collections	265
10.7. Make Vector Your “Go To” Immutable Sequence	266
10.8. Make ArrayBuffer Your “Go To” Mutable Sequence	268
10.9. Looping over a Collection with foreach	270
10.10. Looping over a Collection with a for Loop	272
10.11. Using zipWithIndex or zip to Create Loop Counters	276
10.12. Using Iterators	278
10.13. Transforming One Collection to Another with for/yield	279
10.14. Transforming One Collection to Another with map	282
10.15. Flattening a List of Lists with flatten	285
10.16. Combining map and flatten with flatMap	286
10.17. Using filter to Filter a Collection	289
10.18. Extracting a Sequence of Elements from a Collection	291
10.19. Splitting Sequences into Subsets (groupBy, partition, etc.)	293
10.20. Walking Through a Collection with the reduce and fold Methods	295
10.21. Extracting Unique Elements from a Sequence	300
10.22. Merging Sequential Collections	302
10.23. Merging Two Sequential Collections into Pairs with zip	304
10.24. Creating a Lazy View on a Collection	306
10.25. Populating a Collection with a Range	309
10.26. Creating and Using Enumerations	311
10.27. Tuples, for When You Just Need a Bag of Things	312
10.28. Sorting a Collection	315
10.29. Converting a Collection to a String with mkString	318
11. List, Array, Map, Set (and More).....	321
11.1. Different Ways to Create and Populate a List	322
11.2. Creating a Mutable List	324
11.3. Adding Elements to a List	325
11.4. Deleting Elements from a List (or ListBuffer)	328
11.5. Merging (Concatenating) Lists	330
11.6. Using Stream, a Lazy Version of a List	331
11.7. Different Ways to Create and Update an Array	333

11.8. Creating an Array Whose Size Can Change (ArrayBuffer)	335
11.9. Deleting Array and ArrayBuffer Elements	335
11.10. Sorting Arrays	337
11.11. Creating Multidimensional Arrays	338
11.12. Creating Maps	341
11.13. Choosing a Map Implementation	343
11.14. Adding, Updating, and Removing Elements with a Mutable Map	345
11.15. Adding, Updating, and Removing Elements with Immutable Maps	347
11.16. Accessing Map Values	349
11.17. Traversing a Map	350
11.18. Getting the Keys or Values from a Map	352
11.19. Reversing Keys and Values	352
11.20. Testing for the Existence of a Key or Value in a Map	353
11.21. Filtering a Map	354
11.22. Sorting an Existing Map by Key or Value	357
11.23. Finding the Largest Key or Value in a Map	360
11.24. Adding Elements to a Set	361
11.25. Deleting Elements from Sets	363
11.26. Using Sortable Sets	365
11.27. Using a Queue	367
11.28. Using a Stack	369
11.29. Using a Range	371
12. Files and Processes.....	375
12.1. How to Open and Read a Text File	375
12.2. Writing Text Files	381
12.3. Reading and Writing Binary Files	382
12.4. How to Process Every Character in a Text File	383
12.5. How to Process a CSV File	384
12.6. Pretending that a String Is a File	387
12.7. Using Serialization	389
12.8. Listing Files in a Directory	391
12.9. Listing Subdirectories Beneath a Directory	392
12.10. Executing External Commands	394
12.11. Executing External Commands and Using STDOUT	397
12.12. Handling STDOUT and STDERR for External Commands	399
12.13. Building a Pipeline of Commands	401
12.14. Redirecting the STDOUT and STDIN of External Commands	402
12.15. Using AND (&&) and OR () with Processes	404
12.16. Handling Wildcard Characters in External Commands	405
12.17. How to Run a Process in a Different Directory	406
12.18. Setting Environment Variables When Running Commands	407

12.19. An Index of Methods to Execute External Commands	408
13. Actors and Concurrency.....	411
13.1. Getting Started with a Simple Actor	414
13.2. Creating an Actor Whose Class Constructor Requires Arguments	418
13.3. How to Communicate Between Actors	419
13.4. Understanding the Methods in the Akka Actor Lifecycle	422
13.5. Starting an Actor	425
13.6. Stopping Actors	427
13.7. Shutting Down the Akka Actor System	432
13.8. Monitoring the Death of an Actor with watch	433
13.9. Simple Concurrency with Futures	436
13.10. Sending a Message to an Actor and Waiting for a Reply	445
13.11. Switching Between Different States with become	446
13.12. Using Parallel Collections	448
14. Command-Line Tasks.....	453
14.1. Getting Started with the Scala REPL	454
14.2. Pasting and Loading Blocks of Code into the REPL	459
14.3. Adding JAR Files and Classes to the REPL Classpath	461
14.4. Running a Shell Command from the REPL	462
14.5. Compiling with scalac and Running with scala	465
14.6. Disassembling and Decompiling Scala Code	466
14.7. Finding Scala Libraries	471
14.8. Generating Documentation with scaladoc	472
14.9. Faster Command-Line Compiling with fsc	479
14.10. Using Scala as a Scripting Language	480
14.11. Accessing Command-Line Arguments from a Script	483
14.12. Prompting for Input from a Scala Shell Script	485
14.13. Make Your Scala Scripts Run Faster	489
15. Web Services.....	491
15.1. Creating a JSON String from a Scala Object	491
15.2. Creating a JSON String from Classes That Have Collections	495
15.3. Creating a Simple Scala Object from a JSON String	500
15.4. Parsing JSON Data into an Array of Objects	501
15.5. Creating Web Services with Scalatra	503
15.6. Replacing XML Servlet Mappings with Scalatra Mounts	507
15.7. Accessing Scalatra Web Service GET Parameters	509
15.8. Accessing POST Request Data with Scalatra	510
15.9. Creating a Simple GET Request Client	514
15.10. Sending JSON Data to a POST URL	518

15.11. Getting URL Headers	519
15.12. Setting URL Headers When Sending a Request	520
15.13. Creating a GET Request Web Service with the Play Framework	521
15.14. POSTing JSON Data to a Play Framework Web Service	524
16. Databases and Persistence.....	527
16.1. Connecting to MySQL with JDBC	528
16.2. Connecting to a Database with the Spring Framework	530
16.3. Connecting to MongoDB and Inserting Data	533
16.4. Inserting Documents into MongoDB with insert, save, or +=	537
16.5. Searching a MongoDB Collection	539
16.6. Updating Documents in a MongoDB Collection	542
16.7. Accessing the MongoDB Document ID Field	544
16.8. Deleting Documents in a MongoDB Collection	545
16.9. A Quick Look at Slick	547
17. Interacting with Java.....	549
17.1. Going to and from Java Collections	549
17.2. Add Exception Annotations to Scala Methods to Work with Java	554
17.3. Using @SerialVersionUID and Other Annotations	556
17.4. Using the Spring Framework	557
17.5. Annotating varargs Methods	560
17.6. When Java Code Requires JavaBeans	562
17.7. Wrapping Traits with Implementations	565
18. The Simple Build Tool (SBT).....	569
18.1. Creating a Project Directory Structure for SBT	570
18.2. Compiling, Running, and Packaging a Scala Project with SBT	574
18.3. Running Tests with SBT and ScalaTest	579
18.4. Managing Dependencies with SBT	581
18.5. Controlling Which Version of a Managed Dependency Is Used	584
18.6. Creating a Project with Subprojects	586
18.7. Using SBT with Eclipse	588
18.8. Generating Project API Documentation	590
18.9. Specifying a Main Class to Run	591
18.10. Using GitHub Projects as Project Dependencies	593
18.11. Telling SBT How to Find a Repository (Working with Resolvers)	595
18.12. Resolving Problems by Getting an SBT Stack Trace	596
18.13. Setting the SBT Log Level	597
18.14. Deploying a Single, Executable JAR File	597
18.15. Publishing Your Library	601
18.16. Using Build.scala Instead of build.sbt	602

18.17. Using a Maven Repository Library with SBT	604
18.18. Building a Scala Project with Ant	606
19. Types.....	611
19.1. Creating Classes That Use Generic Types	614
19.2. Creating a Method That Takes a Simple Generic Type	617
19.3. Using Duck Typing (Structural Types)	618
19.4. Make Mutable Collections Invariant	620
19.5. Make Immutable Collections Covariant	622
19.6. Create a Collection Whose Elements Are All of Some Base Type	624
19.7. Selectively Adding New Behavior to a Closed Model	627
19.8. Building Functionality with Types	630
20. Idioms.....	635
20.1. Create Methods with No Side Effects (Pure Functions)	636
20.2. Prefer Immutable Objects	644
20.3. Think “Expression-Oriented Programming”	647
20.4. Use Match Expressions and Pattern Matching	650
20.5. Eliminate null Values from Your Code	654
20.6. Using the Option/Some/None Pattern	658
Index.....	667

Preface

This is a cookbook of problem-solving recipes about Scala, the most interesting programming language I've ever used. The book contains solutions to more than 250 common problems, shown with possibly more than 700 examples. (I haven't counted, but I suspect that's true.)

There are a few unique things about this book:

- As a cookbook, it's intended to save you time by providing solutions to the most common problems you'll encounter.
- Almost all of the examples are shown in the Scala interpreter. As a result, whether you're sitting by a computer, on a plane, or reading in your favorite recliner, you get the benefit of seeing their exact output. (Which often leads to, "Ah, so that's how that works.")
- The book covers not only the Scala language, but also has large chapters on Scala tools and libraries, including SBT, actors, the collections library (more than 100 pages), and JSON processing.

Just prior to its release, the book was updated to cover Scala 2.10.x and SBT 0.12.3.

The Scala Language

My (oversimplified) Scala elevator pitch is that it's a child of Ruby and Java: it's light, concise, and readable like Ruby, but it compiles to class files that you package as JAR files that run on the JVM; it uses traits and mixins, and feels dynamic, but it's statically typed. It uses the Actor model to simplify concurrent programming so you can keep those multicore processors humming. The name Scala comes from the word *scalable*, and true to that name, it's used to power the busiest websites in the world, including Twitter, Netflix, Tumblr, LinkedIn, Foursquare, and many more.

In my opinion, Scala is not a good language for teaching a Programming 101 class. Instead, it's a power language created for the professional programmer. Don't let that scare you, though. If you were my own brother and about to start a new project and could choose any programming language available, without hesitation I'd say, "Use Scala."

Here are a few more nuggets about Scala:

- It's a modern programming language created by Martin Odersky (the father of `javac`), influenced by Java, Ruby, Smalltalk, ML, Haskell, Erlang, and others.
- It's a *pure* object-oriented programming (OOP) language. Every variable is an object, and every "operator" is a method.
- It's also a functional programming (FP) language, so you can pass functions around as variables. You can write your code using OOP, FP, or both.
- Scala code runs on the JVM and lets you use the wealth of Java libraries that have been developed over the years.
- You can be productive on Day 1, but the language is deep, so as you go along you'll keep learning and finding newer, better ways to write code. Scala will change the way you think about programming—and that's a good thing.

Of all of Scala's benefits, what I like best is that it lets you write concise, readable code. The time a programmer spends *reading* code compared to the time spent *writing* code is said to be at least a 10:1 ratio, so writing code that's concise and readable is a big deal. Because Scala has these attributes, programmers say that it's *expressive*.

Solutions

I've always bought O'Reilly cookbooks for the solutions, and that's what this book is about: solving problems.

When using a cookbook, I usually think, "I have this problem, I need to iterate over the elements in an `Array`, what's the best way to do that?" I like to look at the table of contents, find a recipe, implement the solution, and move on. I tried to write each recipe with this use case in mind.

However, with a modern language like Scala, it may end up that I phrased my question wrong. Because of my prior programming experience I may have thought, "I need to iterate over the elements in an `Array`," but in reality my deeper *intent* was to loop over those elements for a reason, such as to transform them into a new collection. So it's nice when a recipe says, "Hey, I know you're here to read about how to loop over the elements in an `Array`, here's how you do that":

```
for (i <- Array(1,2,3)) println(i)
```

“But, if what you’re really trying to do is transform those elements into a new collection, what you want is a `for/yield` expression or `map` method”:

```
// for/yield
scala> for (i <- Array(1,2,3)) yield i * 2
res0: Array[Int] = Array(2, 4, 6)

// map
scala> Array(1,2,3).map(_ * 2)
res1: Array[Int] = Array(2, 4, 6)
```

(More on that `_` character shortly.)

To create the list of problems and solutions, I followed the “**Eat your own dog food**” philosophy. The recipes come from my own experience of creating Scala scripts, web applications, web services, Swing applications, and actor-based systems. As I developed the applications I needed, I encountered problems like these:

- Scala files tend to be very small; what’s the proper way to organize an application?
- It looks like SBT is the best build tool for Scala, but it’s different than Ant or Maven; how do I compile and package applications, and work with dependencies?
- Constructors are really different than Java; how do I create them? What code is generated when I declare constructor parameters and class fields?
- Actors are cool; how do I write a complete actor-based application?
- What, I shouldn’t use `null` values anymore? Why not? How do I code without them?
- I can pass a function around like any other variable? How do I do that, and what’s the benefit?
- Why are there so many collections classes, and why does each collection class have so many methods?
- I have all of this legacy Java code; can I still use it in Scala? If so, how?
- I’m starting to grok this. Now I need to know, what are the top five or ten “best practices” of writing Scala code?

Truthfully, I fell fast in love with everything about Scala except for one thing: the collections library seemed large and intimidating. I really enjoyed using Scala so I kept using the language, but whenever I needed a collection, I used a trusty old Java collection.

Then one day I got up the courage to dive into the collections library. I thought I’d hate it, but after struggling with it for a while, I suddenly “got” it. The light bulb went on over my head, and I suddenly understood not only the collections, but several other concepts I had been struggling with as well. I realized the collections library writers aren’t crazy; they’re brilliant.

Once I understood the collections library, I quit writing so many for loops, and started using collection methods like `filter`, `foreach`, and `map`. They made coding easier, and made my code more concise. These days I can't imagine a better way to write code like this:

```
// filter the items in a list
scala> val nums = List(1,2,3,4,5).filter(_ < 4)
nums: List[Int] = List(1, 2, 3)
```

The `_` wildcard character is discussed in several recipes, but as you can infer from that example, it's a placeholder for each element in the collection. The `filter` method loops through each element in the list, calling your `_ < 4` function on each iteration. That Scala one-liner is the equivalent of this Java code:

```
Integer[] intArray = {1,2,3,4,5};
List<Integer> nums = Arrays.asList(intArray);
List<Integer> filteredNums = new LinkedList<Integer>();
for (int n: nums) {
    if (n < 4) filteredNums.add(n);
}
```

The next example takes this a step further. It filters the elements as in the previous example, and then multiplies each element by the number 2 using the `map` method:

```
// filter the items, then double them
scala> val nums = List(1,2,3,4,5).filter(_ < 4).map(_ * 2)
nums: List[Int] = List(2, 4, 6)
```

If you think about how much code would be required to write this expression in another language, I think you'll agree that Scala is expressive.

(If you're new to Scala, examples like this are broken down into smaller chunks in the recipes.)

Audience

This book is intended for programmers who want to be able to quickly find solutions to problems they'll encounter when using Scala and its libraries and tools. I hope it will also be a good tool for developers who want to learn Scala. I'm a big believer in "learning by example," and this book is chock full of examples.

I generally assume that you have some experience with another programming language like C, C++, Java, Ruby, C#, PHP, Python, or similar. My own experience is with those languages, so I'm sure my writing is influenced by that background.

Another way to describe the audience for this book involves looking at different levels of software developers. In the article at scala-lang.org, Martin Odersky defines the following levels of computer programmers:

- Level A1: Beginning application programmer
- Level A2: Intermediate application programmer
- Level A3: Expert application programmer
- Level L1: Junior library designer
- Level L2: Senior library designer
- Level L3: Expert library designer

This book is primarily aimed at the application developers in the A1, A2, A3, and L1 categories. While helping those developers is my primary goal, I hope that L2 and L3 developers can also benefit from the many examples in this book—especially if they have no prior experience with functional programming, or they want to quickly get up to speed with Scala and its tools and libraries.

Contents of This Book

The first three chapters in this book cover some of the nuts and bolts of the Scala language.

Chapter 1, *Strings*, provides recipes for working with strings. Scala gets its basic `String` functionality from Java, but with the power of *implicit conversions*, Scala adds new functionality to strings through classes like `StringLike` and `StringOps`, which let Scala treat a `String` as a sequence of `Char`. The last recipe in the chapter shows how to add your own behavior to a `String` (or any other class) by creating an implicit conversion.

Chapter 2, *Numbers*, provides recipes for working with Scala’s numeric types. There are no `++` and `--` operators for working with numbers, and this chapter explains why, and demonstrates the other methods you can use. It also shows how to handle large numbers, currency, and how to compare floating-point numbers.

Chapter 3, *Control Structures*, demonstrates Scala’s built-in control structures, starting with `if/then` statements and `for` loops, and then provides solutions for working with `for/yield` loops (for comprehensions) and `for` expressions with embedded `if` statements (guards). Because `match` expressions are so important to Scala, several recipes show how to use them to solve a variety of problems.

The next five chapters continue to cover the Scala syntax, with an emphasis on organizing your projects with classes, methods, objects, traits, and packaging. Recipes on classes, methods, objects, and traits place an emphasis on object-oriented programming techniques.

Chapter 4, *Classes and Properties*, provides examples related to Scala classes and fields. Because Scala constructors are very different than Java constructors, several recipes show the ins and outs of writing both primary and auxiliary constructors. The chapter

also shows how to override the accessor and mutator methods that Scala automatically generates for your `val` and `var` variables. Several recipes show what *case classes* are and how to use them, and how to write `equals` methods.

Chapter 5, *Methods*, shows how to define methods to accept parameters, return values, use parameter names when calling methods, set default values for method parameters, create `varargs` fields, and write methods to support a fluent style of programming.

Chapter 6, *Objects*, covers “all things object.” Like Java, Scala uses the word *object* to refer to an instance of a class, but Scala also has an `object` keyword. This chapter covers topics like class casting, how to launch an application with an object, how to create the equivalent of Java’s static members, and how to write a class with a companion object so you can create new instances of a class without using the `new` keyword.

Chapter 7, *Packaging and Imports*, provides examples of Scala’s `package` and `import` statements, which provide more capabilities than the same Java keywords. This includes how to use the curly brace style for packaging, how to hide and rename members when you import them, and more.

Chapter 8, *Traits*, provides examples of the Scala trait. It begins by showing how to use a trait like a Java interface, and then gets into more advanced topics, such as how to use traits as “mixins,” and limit which members a trait can be mixed into using a variety of methods.

Although much of the book demonstrates functional programming (FP) techniques, **Chapter 9, *Functional Programming***, combines many FP recipes into one location. Solutions show how to define anonymous functions (function literals) and use them in a variety of situations. Recipes demonstrate how to define a method that accepts a function argument, how to return a function from a function, and how to use closures and partially applied functions.

The Scala collections library is rich and deep, so **Chapter 10, *Collections***, and **Chapter 11, *List, Array, Map, Set (and More)***, provide more than 100 pages of collection-related solutions.

Recipes in **Chapter 10, *Collections***, help you choose collection classes for specific needs, and then help you choose and use methods within a collection to solve specific problems, such as transforming one collection into a new collection, filtering a collection, and creating subgroups of a collection. More than 60 pages of recipes demonstrate solutions for writing for loops, `for/yield` expressions, using methods like `filter`, `foreach`, `groupBy`, `map`, and many more.

Chapter 11, *List, Array, Map, Set (and More)*, continues where **Chapter 10, *Collections***, leaves off, providing solutions for those specific collection types, as well as recipes for the `Queue`, `Stack`, and `Range` classes.

Chapter 12, *Files and Processes*, begins by providing solutions about reading and writing files with Scala, including CSV. After that, because the Scala library makes it much (much!) easier to work with external processes than Java, a collection of recipes demonstrates how to execute external commands and work with their I/O.

Chapter 13, *Actors and Concurrency*, provides solutions for the wonderful world of building concurrent applications (and engaging those multicore CPUs) with the Scala Actors library. Recipes in this chapter show solutions to common problems using the industrial-strength Akka Actors library that was integrated into the 2.10.x Scala release. Examples show how to build actor-based applications from the ground up, how to send messages to actors, how to receive and work with messages in actors, and how to kill actors and shut down the system. It also shows easy ways to run concurrent tasks with a Future, a terrific way to run simple computations in parallel.

Chapter 14, *Command-Line Tasks*, combines a collection of recipes centered around using Scala at the command line. It begins by showing tips on how to use the Scala REPL, and then shows how to use command-line tools like `scalac`, `scala`, `scaladoc`, and `fsc`. It also provides recipes showing how to use Scala as a scripting language, including how to precompile your Scala scripts to make them run faster.

Chapter 15, *Web Services*, shows how to use Scala on both the client and server sides of web services. On the server side, it shows how to use Scalatra and the Play Framework to develop RESTful web services, including how to use Scalatra with MongoDB. For both client and server code, it shows how to serialize and deserialize JSON and how to work with HTTP headers.

Chapter 16, *Databases and Persistence*, provides examples of how to interact with databases from Scala, including working with traditional SQL databases using JDBC and Spring JDBC, along with extensive coverage of how to work with MongoDB, a popular “NoSQL” database.

Chapter 17, *Interacting with Java*, shows how to solve the few problems you’ll encounter when working with Java code. While Scala code often *just works* when interacting with Java, there are a few gotchas. This chapter shows how to resolve problems related to the differences in the collections libraries, as well as problems you can run into when calling Scala code from Java.

Chapter 18, *The Simple Build Tool (SBT)*, is a comprehensive guide to the de-facto build tool for Scala applications. It starts by showing several ways to create an SBT project directory structure, and then shows how to include managed and unmanaged dependencies, build your projects, generate Scaladoc for your projects, deploy your projects, and more. Though I strongly recommend learning SBT, a recipe also shows how to use Ant to compile Scala projects.

Chapter 19, *Types*, provides recipes for working with Scala’s powerful type system. Starting right from the introduction, concepts such as type variance, bounds, and

constraints are demonstrated by example. Recipes demonstrate how to declare generics in class and method definitions, implement “duck typing,” and how to control which types your traits can be mixed into.

Chapter 20, *Idioms*, is unique for a cookbook, but because this is a book of solutions, I think it’s important to have a section dedicated to showing the best practices, i.e., how to write code “the Scala way.” Recipes show how to create methods with no side effects, how to work with immutable objects and collection types, how to think in terms of *expressions* (rather than statements), how to use pattern matching, and how to eliminate null values in your code.

Online Bonus Chapters

Because Scala is an incredibly rich and deep language, an additional three chapters consisting of more than 130 pages of *Scala Cookbook* content are available for readers who wish to explore Scala further. These bonus chapters are:

- *XML and XPath*
- *Testing and Debugging*
- *The Play Framework*

These chapters are available in PDF format, and can be downloaded at <http://examples.oreilly.com/9781449339616-files/>.

Installing the Software

Installing Scala is simple and should just take a few minutes.

On Unix systems (including Mac OS X), download the software from the **Scala download page** to a directory on your computer like `$HOME/scala`, and then add these lines to your `$HOME/.bash_profile` file (or its equivalent, depending on which login shell you’re using):

```
export SCALA_HOME=/Users/Al/scala
PATH=$PATH:/Users/Al/scala/bin
```

Once you’ve done this, when you open a new terminal window, you should have access to the `scala` and `scalac` commands at your command line.

You can follow a similar process if you’re using Microsoft Windows, or you can use an MSI installer. See the **Scala download page** for more information.

How the Code Listings Work

Most of the code listings in the book are shown in the Scala “Read-Eval-Print-Loop,” or *REPL*. If you’ve used `irb` with Ruby, the concept is the same: you type an expression, and the REPL evaluates the expression and prints the resulting output.

In the REPL examples, the code that’s shown in a bold font is what you type, and all the text that isn’t bold is output from the REPL.

You start the REPL from your operating system command line by executing the `scala` command:

```
$ scala
Welcome to Scala version 2.10.1
Type in expressions to have them evaluated.
Type :help for more information.

scala> _
```

Once the REPL has started, just type your expressions as input, and the REPL will evaluate them and show their output:

```
scala> val hello = "Hello, world"
hello: String = Hello, world

scala> Array(1,2,3).foreach(println)
1
2
3
```

The REPL is demonstrated more in the [Chapter 1](#) introduction and [Recipe 14.1](#), “Getting Started with the Scala REPL”. [Recipe 14.4](#) takes this a step further and shows how to customize the REPL environment.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

Constant width bold

Shows commands or other text that should be typed literally by the user.

Constant width italic

Shows text that should be replaced with user-supplied values or by values determined by context.



This icon signifies a tip, suggestion, or general note.



This icon indicates a warning or caution.

Using Code Examples

This book is here to help you get your job done. In general, if this book includes code examples, you may use the code in your programs and documentation. You do not need to contact us for permission unless you’re reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing a CD-ROM of examples from O’Reilly books does require permission. Answering a question by citing this book and quoting example code does not require permission. Incorporating a significant amount of example code from this book into your product’s documentation does require permission.

Supplemental material (code examples, exercises, etc.) is available for download at <https://github.com/alvinj>.

We appreciate, but do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: “*Scala Cookbook* by Alvin Alexander (O’Reilly). Copyright 2013 Alvin Alexander, 978-1-449-33961-6.”

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at permissions@oreilly.com.

Safari® Books Online



Safari Books Online (www.safaribooksonline.com) is an on-demand digital library that delivers expert **content** in both book and video form from the world’s leading authors in technology and business.

Technology professionals, software developers, web designers, and business and creative professionals use Safari Books Online as their primary resource for research, problem solving, learning, and certification training.

Safari Books Online offers a range of **product mixes** and pricing programs for **organizations**, **government agencies**, and **individuals**. Subscribers have access to thousands of books, training videos, and prepublication manuscripts in one fully searchable database from publishers like O'Reilly Media, Prentice Hall Professional, Addison-Wesley Professional, Microsoft Press, Sams, Que, Peachpit Press, Focal Press, Cisco Press, John Wiley & Sons, Syngress, Morgan Kaufmann, IBM Redbooks, Packt, Adobe Press, FT Press, Apress, Manning, New Riders, McGraw-Hill, Jones & Bartlett, Course Technology, and dozens **more**. For more information about Safari Books Online, please visit us **online**.

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at http://oreil.ly/Scala_CB.

To comment or ask technical questions about this book, send an email to bookquestions@oreilly.com.

For more information about our books, courses, conferences, and news, see our website at <http://www.oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>

Follow us on Twitter: <http://twitter.com/oreillymedia>

Watch us on YouTube: <http://www.youtube.com/oreillymedia>

Acknowledgments

Writing a book this large takes a lot of work, and I'd like to thank my editor, Courtney Nash, for keeping me sane during the speed bumps and generally being encouraging throughout the process.

Kim Cofer was the copy editor for this book, and I'd like to thank her for helping whip the book into shape, correcting my grammar issues regardless of how many times I repeated them, and for having good discussions about how to handle several issues in this book.

This book grew from about 540 pages during the first review to roughly 700 pages in its final release, and much of that was due to reviewers. All of the reviewers were helpful in different ways, but I'd especially like to thank **Eric Torreborre** and **Ryan LeCompte** for making it all the way through different versions of the book. Additional thanks go out to Rudi Farkas, Rahul Phulore, Jason Swartz, Hugo Sereno Ferreira, and Dean Wampler.


I'd also like to thank my friends and family members who encouraged me throughout the process. A special thanks goes to my sister Melissa, who helped bring my initial plain, wiki-style text into Microsoft Word, and styled everything correctly.

Finally, I'd like to thank Martin Odersky and his team for creating such an interesting programming language. I also owe his Programming Methods Laboratory at EFPL a special thank you for letting me use the Scala collections performance tables shown in **Recipe 10.4**.

Introduction

At first glance, a Scala String appears to be just a Java String. For instance, when you work in the Scala *Read-Evaluate-Print-Loop* (REPL) environment (see [Figure 1-1](#)) and print the name of a String literal, the REPL feedback tells you the type is `java.lang.String`:

```
scala> "Hello, world".getClass.getName  
res0: String = java.lang.String
```

A screenshot of a terminal window titled "alvinalexander.com". The terminal shows the Scala REPL environment. It starts with a prompt "\$ scala" followed by a welcome message: "Welcome to Scala version 2.10.0. Type in expressions to have them evaluated. Type :help for more information." Then, the user enters "scala> val s = \"Hello, world\"", and the REPL responds with "s: String = Hello, world". Next, the user enters "scala> s.length", and the REPL responds with "res0: Int = 12". Finally, the user enters "scala> \"hello\".foreach(println)", and the REPL prints "h", "e", "l", "l", "o" on separate lines. The prompt "scala> " is visible at the bottom of the terminal.

```
$ scala  
Welcome to Scala version 2.10.0.  
Type in expressions to have them evaluated.  
Type :help for more information.  
  
scala> val s = "Hello, world"  
s: String = Hello, world  
  
scala> s.length  
res0: Int = 12  
  
scala> "hello".foreach(println)  
h  
e  
l  
l  
o  
  
scala> 
```

Figure 1-1. The Scala REPL is an interactive environment where you can test Scala statements

Indeed, a Scala `String` is a Java `String`, so you can use all the normal Java string methods. You can create a string variable, albeit in the Scala way:

```
val s = "Hello, world"
```

You can get the length of a string:

```
s.length // 12
```

You can concatenate strings:

```
val s = "Hello" + " world"
```

These are all familiar operations. But because Scala offers the magic of *implicit conversions*, `String` instances also have access to all the methods of the `StringOps` class, so you can do many other things with them, such as treating a `String` instance as a sequence of characters. As a result, you can iterate over every character in the string using the `foreach` method:

```
scala> "hello".foreach(println)
h
e
l
l
o
```

You can treat a `String` as a sequence of characters in a `for` loop:

```
scala> for (c <- "hello") println(c)
h
e
l
l
o
```

You can also treat it as a sequence of bytes:

```
scala> s.getBytes.foreach(println)
104
101
108
108
111
```

Because there are many methods available on sequential collections, you can also use other functional methods like `filter`:

```
scala> val result = "hello world".filter(_ != 'l')
result: String = heo word
```


It's an oversimplification to say that this functionality comes from the `StringOps` class, but it's a useful illusion. The reality is that *some* of this functionality comes from `StringOps`, some comes from `StringLike`, some from `WrappedString`, and so on. If you dig into the Scala source code, you'll see that the rabbit hole goes deep, but it begins with the implicit conversion from `String` to `StringOps` in the `Predef` object.



When first learning Scala, take a look at the source code for the `Predef` object. It provides nice examples of many Scala programming features.

Figure 1-2, taken from the `StringOps` class Scaladoc page, shows the supertypes and type hierarchy for the `StringOps` class.




Figure 1-2. Supertypes and type hierarchy information for the `StringOps` class

Add Methods to Closed Classes

Even though the `String` class is declared as `final` in Java, you've seen that Scala somehow adds new functionality to it. This happens through the power of *implicit conversions*. **Recipe 1.9, "Accessing a Character in a String"**, demonstrates how to add your own methods to the `String` class using this technique.

As one more example of how this pattern helps a Scala `String` have both string and collection features, the following code uses the `drop` and `take` methods that are available on Scala sequences, along with the `capitalize` method from the `StringOps` class:

```
scala> "scala".drop(2).take(2).capitalize
res0: String = Al
```

In this chapter you'll see examples like this, and many more.

How Did the Preceding Example Work?

The `drop` and `take` methods are demonstrated in [Chapter 10](#), but in short, `drop` is a collection method that drops (discards) the number of elements that are specified from the beginning of the collection and keeps the remaining elements. When it's called on your string as `drop(2)`, it drops the first two characters from the string (`sc`), and returns the remaining elements:

```
scala> "scala".drop(2)
res0: String = ala
```

Next, the `take(2)` method retains the first two elements from the collection it's given, and discards the rest:

```
scala> "scala".drop(2).take(2)
res1: String = al
```

Finally, you treat the output from the `take(2)` method call like a `String` once again and call the `capitalize` method to get what you want:

```
scala> "scala".drop(2).take(2).capitalize
res2: String = Al
```

The `capitalize` method is in the `StringOps` class, but as a practical matter, you generally don't have to worry about that. When you're writing code in an IDE like Eclipse or IntelliJ and invoke the code assist keystroke, the `capitalize` method will appear in the list along with all the other methods that are available on a `String`.

If you're not familiar with chaining methods together like this, it's known as a *fluent* style of programming. See [Recipe 5.9, "Supporting a Fluent Style of Programming"](#), for more information.

1.1. Testing String Equality

Problem

You want to compare two strings to see if they're equal, i.e., whether they contain the same sequence of characters.

Solution

In Scala, you compare two `String` instances with the `==` operator. Given these strings:

```
scala> val s1 = "Hello"
s1: String = Hello
```

```
scala> val s2 = "Hello"
s2: String = Hello

scala> val s3 = "H" + "ello"
s3: String = Hello
```

You can test their equality like this:

```
scala> s1 == s2
res0: Boolean = true

scala> s1 == s3
res1: Boolean = true
```

A pleasant benefit of the `==` method is that it doesn't throw a `NullPointerException` on a basic test if a `String` is `null`:

```
scala> val s4: String = null
s4: String = null

scala> s3 == s4
res2: Boolean = false

scala> s4 == s3
res3: Boolean = false
```

If you want to compare two strings in a case-insensitive manner, you can convert both strings to uppercase or lowercase and compare them with the `==` method:

```
scala> val s1 = "Hello"
s1: String = Hello

scala> val s2 = "hello"
s2: String = hello

scala> s1.toUpperCase == s2.toUpperCase
res0: Boolean = true
```

However, be aware that calling a method on a `null` string can throw a `NullPointerException`:

```
scala> val s1: String = null
s1: String = null

scala> val s2: String = null
s2: String = null

scala> s1.toUpperCase == s2.toUpperCase
java.lang.NullPointerException // more output here ...
```

To compare two strings while ignoring their case, you can also fall back and use the `equalsIgnoreCase` of the Java `String` class:

```
scala> val a = "Marisa"
a: String = Marisa

scala> val b = "marisa"
b: String = marisa

scala> a.equalsIgnoreCase(b)
res0: Boolean = true
```

Discussion

In Scala, you test object equality with the `==` method. This is different than Java, where you use the `equals` method to compare two objects.

In Scala, the `==` method defined in the `AnyRef` class first checks for `null` values, and then calls the `equals` method on the first object (i.e., `this`) to see if the two objects are equal. As a result, you don't have to check for `null` values when comparing strings.



In idiomatic Scala, you *never* use `null` values. The discussion in this recipe is intended to help you understand how `==` works if you encounter a `null` value, presumably from working with a Java library, or some other library where `null` values were used.

If you're coming from a language like Java, any time you feel like using a `null`, use an `Option` instead. (I find it helpful to imagine that Scala doesn't even have a `null` keyword.) See [Recipe 20.6, "Using the Option/Some/None Pattern"](#), for more information and examples.

For more information on defining `equals` methods, see [Recipe 4.15, "Defining an equals Method \(Object Equality\)"](#).

1.2. Creating Multiline Strings

Problem

You want to create multiline strings within your Scala source code, like you can with the "heredoc" syntax of other languages.

Solution

In Scala, you create multiline strings by surrounding your text with three double quotes:

```
val foo = """This is
a multiline
String"""
```

Discussion

Although this works, the second and third lines in this example will end up with white-space at the beginning of their lines. If you print the string, it looks like this:

```
This is
  a multiline
    String
```

You can solve this problem in several different ways. First, you can left-justify every line after the first line of your string:

```
val foo = """This is
a multiline
String"""
```

A cleaner approach is to add the `stripMargin` method to the end of your multiline string and begin all lines after the first line with the pipe symbol (`|`):

```
val speech = """Four score and
|seven years ago""".stripMargin
```

If you don't like using the `|` symbol, you can use any character you like with the `stripMargin` method:

```
val speech = """Four score and
#seven years ago""".stripMargin('#')
```

All of these approaches yield the same result, a multiline string with each line of the string left justified:

```
Four score and
seven years ago
```

This results in a true multiline string, with a hidden `\n` character after the word “and” in the first line. To convert this multiline string into one continuous line you can add a `replaceAll` method after the `stripMargin` call, replacing all newline characters with blank spaces:

```
val speech = """Four score and
|seven years ago
|our fathers""".stripMargin.replaceAll("\n", " ")
```

This yields:

```
Four score and seven years ago our fathers
```

Another nice feature of Scala's multiline string syntax is that you can include single- and double-quotes without having to escape them:

```
val s = """This is known as a
|"multiline" string
|or 'heredoc' syntax.""".stripMargin.replaceAll("\n", " ")
```

This results in this string:

This is known as a "multiline" string or 'heredoc' syntax.

1.3. Splitting Strings

Problem

You want to split a string into parts based on a field separator, such as a string you get from a comma-separated value (CSV) or pipe-delimited file.

Solution

Use one of the `split` methods that are available on `String` objects:

```
scala> "hello world".split(" ")
res0: Array[java.lang.String] = Array(hello, world)
```

The `split` method returns an array of `String` elements, which you can then treat as a normal Scala `Array`:

```
scala> "hello world".split(" ").foreach(println)
hello
world
```

Discussion

The string that the `split` method takes can be a regular expression, so you can split a string on simple characters like a comma in a CSV file:

```
scala> val s = "eggs, milk, butter, Coco Puffs"
s: java.lang.String = eggs, milk, butter, Coco Puffs

// 1st attempt
scala> s.split(",")
res0: Array[java.lang.String] = Array(eggs, " milk", " butter", " Coco Puffs")
```

Using this approach, it's best to trim each string. Use the `map` method to call `trim` on each string before returning the array:

```
// 2nd attempt, cleaned up
scala> s.split(",").map(_.trim)
res1: Array[java.lang.String] = Array(eggs, milk, butter, Coco Puffs)
```

You can also split a string based on a regular expression. This example shows how to split a string on whitespace characters:

```
scala> "hello world, this is Al".split("\\s+")
res0: Array[java.lang.String] = Array(hello, world,, this, is, Al)
```

About that split method...

The `split` method is overloaded, with some versions of the method coming from the Java `String` class and some coming from the Scala `StringLike` class. For instance, if you call `split` with a `Char` argument instead of a `String` argument, you're using the `split` method from `StringLike`:

```
// split with a String argument
scala> "hello world".split(" ")
res0: Array[java.lang.String] = Array(hello, world)

// split with a Char argument
scala> "hello world".split(' ')
res1: Array[String] = Array(hello, world)
```

The subtle difference in that output—`Array[java.lang.String]` versus `Array[String]`—is a hint that something is different, but as a practical matter, this isn't important. Also, with the Scala IDE project integrated into Eclipse, you can see where each method comes from when the Eclipse “code assist” dialog is displayed. (IntelliJ IDEA and NetBeans may show similar information.)

1.4. Substituting Variables into Strings

Problem

You want to perform variable substitution into a string, like you can do with other languages, such as Perl, PHP, and Ruby.

Solution

Beginning with Scala 2.10 you can use *string interpolation* in a manner similar to other languages like Perl, PHP, and Ruby.

To use basic string interpolation in Scala, precede your string with the letter `s` and include your variables inside the string, with each variable name preceded by a `$` character. This is shown in the `println` statement in the following example:

```
scala> val name = "Fred"
name: String = Fred

scala> val age = 33
age: Int = 33

scala> val weight = 200.00
weight: Double = 200.0

scala> println(s"$name is $age years old, and weighs $weight pounds.")
Fred is 33 years old, and weighs 200.0 pounds.
```

According to [the official Scala string interpolation documentation](#), when you precede your string with the letter `s`, you're creating a *processed* string literal. This example uses the “`s` string interpolator,” which lets you embed variables inside a string, where they're replaced by their values. As stated in the documentation, “Prepending `s` to any string literal allows the usage of variables directly in the string.”

Using expressions in string literals

In addition to putting variables inside strings, you can include expressions inside a string by placing the expression inside curly braces. According to [the official string interpolation documentation](#), “Any arbitrary expression can be embedded in `{}`.”

In the following example, the value `1` is added to the variable `age` inside the string:

```
scala> println(s"Age next year: ${age + 1}")
Age next year: 34
```

This example shows that you can use an equality expression inside the curly braces:

```
scala> println(s"You are 33 years old: ${age == 33}")
You are 33 years old: true
```

You'll also need to use curly braces when printing object fields. The following example shows the correct approach:

```
scala> case class Student(name: String, score: Int)
defined class Student

scala> val hannah = Student("Hannah", 95)
hannah: Student = Student(Hannah,95)

scala> println(s"${hannah.name} has a score of ${hannah.score}")
Hannah has a score of 95
```

Attempting to print the values of the object fields without wrapping them in curly braces results in the wrong information being printed out:

```
// error: this is intentionally wrong
scala> println(s"$hannah.name has a score of $hannah.score")
Student(Hannah,95).name has a score of Student(Hannah,95).score
```

Because `$hannah.name` wasn't wrapped in curly braces, the wrong information was printed; in this case, the `toString` output of the `hannah` variable.

`s` is a method

The `s` that's placed before each string literal is actually a method. Though this seems slightly less convenient than just putting variables inside of strings, there are at least two benefits to this approach:

- Scala provides other off-the-shelf interpolation functions to give you more power.
- You can define your own string interpolation functions.

To see why this is a good thing, let's look at another string interpolation function.

The **f** string interpolator (**printf** style formatting)

In the example in the Solution, the `weight` was printed as `200.0`. This is okay, but what can you do if you want to add more decimal places to the weight, or remove them entirely?

This simple desire leads to the “**f** string interpolator,” which lets you use `printf` style formatting specifiers inside strings. The following examples show how to print the `weight`, first with two decimal places:

```
scala> println(f"$name is $age years old, and weighs $weight%.2f pounds.")
Fred is 33 years old, and weighs 200.00 pounds.
```

and then with no decimal places:

```
scala> println(f"$name is $age years old, and weighs $weight%.0f pounds.")
Fred is 33 years old, and weighs 200 pounds.
```

As demonstrated, to use this approach, just follow these steps:

1. Precede your string with the letter **f**.
2. Use `printf` style formatting specifiers immediately after your variables.



The most common `printf` format specifiers are shown in [Table 1-1](#) in the Discussion.

Though these examples used the `println` method, it's important to note that you can use string interpolation in other ways. For instance, you can assign the result of a variable substitution to a new variable, similar to calling `sprintf` in other languages:

```
scala> val out = f"$name, you weigh $weight%.0f pounds."
out: String = Fred, you weigh 200 pounds.
```

The **raw** interpolator

In addition to the **s** and **f** string interpolators, Scala 2.10 includes another interpolator named **raw**. The **raw** interpolator “performs no escaping of literals within the string.” The following example shows how **raw** compares to the **s** interpolator:

```
scala> s"foo\nbar"
res0: String =
foo
```

```
bar
```

```
scala> raw"foo\nbar"  
res1: String = foo\nbar
```

The `raw` interpolator is useful when you want to avoid having a sequence of characters like `\n` turn into a newline character.

Create your own interpolator

In addition to the `s`, `f`, and `raw` interpolators that are built into Scala 2.10, you can define your own interpolators. See the official [Scala String Interpolation documentation](#) for an example of how to create your own interpolator.



String interpolation does not work with pattern-matching statements in Scala 2.10. This feature is planned for inclusion in Scala 2.11.

Discussion

Prior to version 2.10, Scala didn't include the string interpolation functionality just described. If you need to use a release prior to Scala 2.10 for some reason, the solution is to call the `format` method on a string, as shown in the following examples:

```
scala> val name = "Fred"  
name: java.lang.String = Fred  
  
scala> val age = 33  
age: Int = 33  
  
scala> val s = "%s is %d years old".format(name, age)  
s: String = Fred is 33 years old  
  
scala> println("%s is %d years old".format(name, age))  
Fred is 33 years old
```

Just as with the string interpolation capability shown in the Solution, you can use this approach anywhere you want to format a string, such as a `toString` method:

```
override def toString: String =  
  "%s %s, age %d".format(firstName, lastName, age)
```

With either of these approaches, you can format your variables using all the usual `printf` specifiers. The most common format specifiers are shown in [Table 1-1](#).

Table 1-1. Common printf style format specifiers

Format specifier	Description
%c	Character
%d	Decimal number (integer, base 10)
%e	Exponential floating-point number
%f	Floating-point number
%i	Integer (base 10)
%o	Octal number (base 8)
%s	A string of characters
%u	Unsigned decimal (integer) number
%x	Hexadecimal number (base 16)
%%	Print a “percent” character
\%	Print a “percent” character

See Also

- This [printf cheat sheet](#) shows more [format specifiers and examples](#)
- This [Oracle Formatter page](#) shows examples and details
- [The official Scala String Interpolation documentation](#)

1.5. Processing a String One Character at a Time

Problem

You want to iterate through each character in a string, performing an operation on each character as you traverse the string.

Solution

Depending on your needs and preferences, you can use the `map` or `foreach` methods, a `for` loop, or other approaches. Here’s a simple example of how to create an uppercase string from an input string, using `map`:

```
scala> val upper = "hello, world".map(c => c.toUpperCase)
upper: String = HELLO, WORLD
```

As you’ll see in many examples throughout this book, you can shorten that code using the magic of Scala’s underscore character:

```
scala> val upper = "hello, world".map(_.toUpperCase)
upper: String = HELLO, WORLD
```

With any collection—such as a sequence of characters in a string—you can also chain collection methods together to achieve a desired result. In the following example, the `filter` method is called on the original `String` to create a new `String` with all occurrences of the lowercase letter “l” removed. That `String` is then used as input to the `map` method to convert the remaining characters to uppercase:

```
scala> val upper = "hello, world".filter(_ != 'l').map(_.toUpperCase)
upper: String = HEO, WORD
```

When you first start with Scala, you may not be comfortable with the `map` method, in which case you can use Scala’s `for` loop to achieve the same result. This example shows another way to print each character:

```
scala> for (c <- "hello") println(c)
h
e
l
l
o
```

To write a `for` loop to work like a `map` method, add a `yield` statement to the end of the loop. This `for/yield` loop is equivalent to the first two `map` examples:

```
scala> val upper = for (c <- "hello, world") yield c.toUpperCase
upper: String = HELLO, WORLD
```

Adding `yield` to a `for` loop essentially places the result from each loop iteration into a temporary holding area. When the loop completes, all of the elements in the holding area are returned as a single collection.

This `for/yield` loop achieves the same result as the third `map` example:

```
val result = for {
  c <- "hello, world"
  if c != 'l'
} yield c.toUpperCase
```

Whereas the `map` or `for/yield` approaches are used to transform one collection into another, the `foreach` method is typically used to operate on each element without returning a result. This is useful for situations like printing:

```
scala> "hello".foreach(println)
h
e
l
l
o
```

Discussion

Because Scala treats a string as a sequence of characters—and because of Scala’s background as both an object-oriented *and* functional programming language—you can

iterate over the characters in a string with the approaches shown. Compare those examples with a common Java approach:

```
String s = "Hello";
StringBuilder sb = new StringBuilder();
for (int i = 0; i < s.length(); i++) {
    char c = s.charAt(i);
    // do something with the character ...
    // sb.append ...
}
String result = sb.toString();
```

You'll see that the Scala approach is more concise, but still very readable. This combination of conciseness and readability lets you focus on solving the problem at hand. Once you get comfortable with Scala, it feels like the imperative code in the Java example obscures your business logic.



Wikipedia describes *imperative programming* like this:

Imperative programming is a programming paradigm that describes computation in terms of statements that change a program state ... imperative programs define sequences of commands for the computer to perform.

This is shown in the Java example, which defines a series of explicit statements that tell a computer how to achieve a desired result.

Understanding how map works

Depending on your coding preferences, you can pass large blocks of code to a `map` method. These two examples demonstrate the syntax for passing an algorithm to a `map` method:

```
// first example
"HELLO".map(c => (c.toByte+32).toChar)

// second example
"HELLO".map{ c =>
    (c.toByte+32).toChar
}
```

Notice that the algorithm operates on one `Char` at a time. This is because the `map` method in this example is called on a `String`, and `map` treats a `String` as a sequential collection of `Char` elements. The `map` method has an implicit loop, and in that loop, it passes one `Char` at a time to the algorithm it's given.

Although this algorithm is still short, imagine for a moment that it is longer. In this case, to keep your code clear, you might want to write it as a method (or function) that you can pass into the `map` method.

To write a method that you can pass into `map` to operate on the characters in a `String`, define it to take a single `Char` as input, then perform the logic on that `Char` inside the method. When the logic is complete, return whatever it is that your algorithm returns. Though the following algorithm is still short, it demonstrates how to create a custom method and pass that method into `map`:

```
// write your own method that operates on a character
scala> def toLower(c: Char): Char = (c.toByte+32).toChar
toLower: (c: Char)Char

// use that method with map
scala> "HELLO".map(toLower)
res0: String = hello
```

As an added benefit, the same method also works with the `for/yield` approach:

```
scala> val s = "HELLO"
s: java.lang.String = HELLO

scala> for (c <- s) yield toLower(c)
res1: String = hello
```



I've used the word "method" in this discussion, but you can also use functions here instead of methods. What's the difference between a method and a function?

Here's a quick look at a *function* equivalent to this `toLower` method:

```
val toLower = (c: Char) => (c.toByte+32).toChar
```

This function can be passed into `map` in the same way the previous `toLower` method was used:

```
scala> "HELLO".map(toLower)
res0: String = hello
```

For more information on functions and the differences between methods and functions, see [Chapter 9, Functional Programming](#).

A complete example

The following example demonstrates how to call the `getBytes` method on a `String`, and then pass a block of code into a `foreach` method to help calculate an Adler-32 checksum value on a `String`:

```
package tests

/**
 * Calculate the Adler-32 checksum using Scala.
 * @see http://en.wikipedia.org/wiki/Adler-32
 */
object Adler32Checksum {
```

```

val MOD_ADLER = 65521

def main(args: Array[String]) {
  val sum = adler32sum("Wikipedia")
  printf("checksum (int) = %d\n", sum)
  printf("checksum (hex) = %s\n", sum.toHexString)
}

def adler32sum(s: String): Int = {
  var a = 1
  var b = 0
  s.getBytes.foreach(char =>
    a = (char + a) % MOD_ADLER
    b = (b + a) % MOD_ADLER
  )
  // note: Int is 32 bits, which this requires
  b * 65536 + a    // or (b << 16) + a
}

```

The `getBytes` method returns a sequential collection of bytes from a `String` as follows:

```

scala> "hello".getBytes
res0: Array[Byte] = Array(104, 101, 108, 108, 111)

```

Adding the `foreach` method call after `getBytes` lets you operate on each `Byte` value:

```

scala> "hello".getBytes.foreach(println)
104
101
108
108
111

```

You use `foreach` in this example instead of `map`, because the goal is to loop over each `Byte` in the `String`, and do something with each `Byte`, but you don't want to return anything from the loop.

See Also

- Under the covers, the Scala compiler translates a `for` loop into a `foreach` method call. This gets more complicated if the loop has one or more `if` statements (guards) or a `yield` expression. This is discussed in detail in [Recipe 3.1, “Looping with `for` and `foreach`”](#) and I also provide examples on my website at alvinalexander.com. The full details are presented in [Section 6.19](#) of the current Scala Language Specification.
- [The Adler-32 checksum algorithm](#)

1.6. Finding Patterns in Strings

Problem

You need to determine whether a `String` contains a regular expression pattern.

Solution

Create a `Regex` object by invoking the `.r` method on a `String`, and then use that pattern with `findFirstIn` when you're looking for one match, and `findAllIn` when looking for all matches.

To demonstrate this, first create a `Regex` for the pattern you want to search for, in this case, a sequence of one or more numeric characters:

```
scala> val numPattern = "[0-9]+".r
numPattern: scala.util.matching.Regex = [0-9]+
```

Next, create a sample `String` you can search:

```
scala> val address = "123 Main Street Suite 101"
address: java.lang.String = 123 Main Street Suite 101
```

The `findFirstIn` method finds the first match:

```
scala> val match1 = numPattern.findFirstIn(address)
match1: Option[String] = Some(123)
```

(Notice that this method returns an `Option[String]`. I'll dig into that in the Discussion.)

When looking for multiple matches, use the `findAllIn` method:

```
scala> val matches = numPattern.findAllIn(address)
matches: scala.util.matching.Regex.MatchIterator = non-empty iterator
```

As you can see, `findAllIn` returns an iterator, which lets you loop over the results:

```
scala> matches.foreach(println)
123
101
```

If `findAllIn` doesn't find any results, an empty iterator is returned, so you can still write your code just like that—you don't need to check to see if the result is `null`. If you'd rather have the results as an `Array`, add the `toArray` method after the `findAllIn` call:

```
scala> val matches = numPattern.findAllIn(address).toArray
matches: Array[String] = Array(123, 101)
```

If there are no matches, this approach yields an empty `Array`. Other methods like `toList`, `toSeq`, and `toVector` are also available.

Discussion

Using the `.r` method on a `String` is the easiest way to create a `Regex` object. Another approach is to import the `Regex` class, create a `Regex` instance, and then use the instance in the same way:

```
scala> import scala.util.matching.Regex
import scala.util.matching.Regex

scala> val numPattern = new Regex("[0-9]+")
numPattern: scala.util.matching.Regex = [0-9]+

scala> val address = "123 Main Street Suite 101"
address: java.lang.String = 123 Main Street Suite 101

scala> val match1 = numPattern.findFirstIn(address)
match1: Option[String] = Some(123)
```

Although this is a bit more work, it's also more obvious. I've found that it can be easy to overlook the `.r` at the end of a `String` (and then spend a few minutes wondering how the code I saw could possibly work).

Handling the Option returned by `findFirstIn`

As mentioned in the Solution, the `findFirstIn` method finds the first match in the `String` and returns an `Option[String]`:

```
scala> val match1 = numPattern.findFirstIn(address)
match1: Option[String] = Some(123)
```

The `Option/Some/None` pattern is discussed in detail in [Recipe 20.6](#), but the simple way to think about an `Option` is that it's a container that holds either zero or one values. In the case of `findFirstIn`, if it succeeds, it returns the string "123" as a `Some(123)`, as shown in this example. However, if it fails to find the pattern in the string it's searching, it will return a `None`, as shown here:

```
scala> val address = "No address given"
address: String = No address given

scala> val match1 = numPattern.findFirstIn(address)
match1: Option[String] = None
```

To summarize, a method defined to return an `Option[String]` will either return a `Some(String)`, or a `None`.

The normal way to work with an `Option` is to use one of these approaches:

- Call `getOrElse` on the value.
- Use the `Option` in a match expression.
- Use the `Option` in a `foreach` loop.

[Recipe 20.6](#) describes those approaches in detail, but they're demonstrated here for your convenience.

With the `getOrElse` approach, you attempt to “get” the result, while also specifying a default value that should be used if the method failed:

```
scala> val result = numPattern.findFirstIn(address).getOrElse("no match")
result: String = 123
```

Because an `Option` is a collection of zero or one elements, an experienced Scala developer will also use a `foreach` loop in this situation:

```
numPattern.findFirstIn(address).foreach { e =>
  // perform the next step in your algorithm,
  // operating on the value 'e'
}
```

A `match` expression also provides a very readable solution to the problem:

```
match1 match {
  case Some(s) => println(s"Found: $s")
  case None =>
}
```

See [Recipe 20.6](#) for more information.

To summarize this approach, the following REPL example shows the complete process of creating a `Regex`, searching a `String` with `findFirstIn`, and then using a `foreach` loop on the resulting match:

```
scala> val numPattern = "[0-9]+".r
numPattern: scala.util.matching.Regex = [0-9]+

scala> val address = "123 Main Street Suite 101"
address: String = 123 Main Street Suite 101

scala> val match1 = numPattern.findFirstIn(address)
match1: Option[String] = Some(123)

scala> match1.foreach { e =>
  |   println(s"Found a match: $e")
  | }
Found a match: 123
```

See Also

- [The `StringOps` class](#)
- [The `Regex` class](#)

- [Recipe 20.6, “Using the Option/Some/None Pattern”](#) provides more information on Option

1.7. Replacing Patterns in Strings

Problem

You want to search for regular-expression patterns in a string, and replace them.

Solution

Because a `String` is immutable, you can't perform find-and-replace operations directly on it, but you can create a new `String` that contains the replaced contents. There are several ways to do this.

You can call `replaceAll` on a `String`, remembering to assign the result to a new variable:

```
scala> val address = "123 Main Street".replaceAll("[0-9]", "x")
address: java.lang.String = xxx Main Street
```

You can create a regular expression and then call `replaceAllIn` on that expression, again remembering to assign the result to a new string:

```
scala> val regex = "[0-9]".r
regex: scala.util.matching.Regex = [0-9]

scala> val newAddress = regex.replaceAllIn("123 Main Street", "x")
newAddress: String = xxx Main Street
```

To replace only the first occurrence of a pattern, use the `replaceFirst` method:

```
scala> val result = "123".replaceFirst("[0-9]", "x")
result: java.lang.String = x23
```

You can also use `replaceFirstIn` with a `Regex`:

```
scala> val regex = "H".r
regex: scala.util.matching.Regex = H

scala> val result = regex.replaceFirstIn("Hello world", "J")
result: String = Jello world
```

See Also

[Recipe 1.6, “Finding Patterns in Strings”](#) for examples of how to find patterns in strings

1.8. Extracting Parts of a String That Match Patterns

Problem

You want to extract one or more parts of a string that match the regular-expression patterns you specify.

Solution

Define the regular-expression patterns you want to extract, placing parentheses around them so you can extract them as “regular-expression groups.” First, define the desired pattern:

```
val pattern = "([0-9]+) ([A-Za-z]+)".r
```

Next, extract the regex groups from the target string:

```
val pattern(count, fruit) = "100 Bananas"
```

This code extracts the numeric field and the alphabetic field from the given string as two separate variables, `count` and `fruit`, as shown in the Scala REPL:

```
scala> val pattern = "([0-9]+) ([A-Za-z]+)".r
pattern: scala.util.matching.Regex = ([0-9]+) ([A-Za-z]+)

scala> val pattern(count, fruit) = "100 Bananas"
count: String = 100
fruit: String = Bananas
```

Discussion

The syntax shown here may feel a little unusual because it seems like you’re defining `pattern` as a `val` field twice, but this syntax is more convenient and readable in a real-world example.

Imagine you’re writing the code for a search engine like Google, and you want to let people search for movies using a wide variety of phrases. To be really convenient, you’ll let them type any of these phrases to get a listing of movies near Boulder, Colorado:

```
"movies near 80301"
"movies 80301"
"80301 movies"
"movie: 80301"
"movies: 80301"
"movies near boulder, co"
"movies near boulder, colorado"
```

One way you can allow all these phrases to be used is to define a series of regular-expression patterns to match against them. Just define your expressions, and then attempt to match whatever the user types against all the possible expressions you're willing to allow.

For example purposes, you'll just allow these two simplified patterns:

```
// match "movies 80301"
val MoviesZipRE = "movies (\\d{5})".r

// match "movies near boulder, co"
val MoviesNearCityStateRE = "movies near ([a-z]+), ([a-z]{2})".r
```

Once you've defined the patterns you want to allow, you can match them against whatever text the user enters, using a match expression. In this example, you'll call a fictional method named `getSearchResults` when a match occurs:

```
textUserTyped match {
  case MoviesZipRE(zip) => getSearchResults(zip)
  case MoviesNearCityStateRE(city, state) => getSearchResults(city, state)
  case _ => println("did not match a regex")
}
```

As you can see, this syntax makes your match expressions very readable. For both patterns you're matching, you call an overloaded version of the `getSearchResults` method, passing it the `zip` field in the first case, and the `city` and `state` fields in the second case.

The two regular expressions shown in this example will match strings like this:

```
"movies 80301"
"movies 99676"
"movies near boulder, co"
"movies near talkeetna, ak"
```

It's important to note that with this technique, the regular expressions must match the *entire* user input. With the regex patterns shown, the following strings will fail because they have a blank space at the end of the line:

```
"movies 80301 "
"movies near boulder, co "
```

You can solve this particular problem by trimming the input string or using a more complicated regular expression, which you'll want to do anyway in the “real world.”

As you can imagine, you can use this same pattern-matching technique in many different circumstances, including matching date and time formats, street addresses, people's names, and many other situations.

See Also

- [Recipe 3.7, “Using a Match Expression Like a switch Statement”](#) for more match expression examples
- [Recipe 14.12, “Prompting for Input from a Scala Shell Script”](#) shows another example of this technique

1.9. Accessing a Character in a String

Problem

You want to get a character at a specific position in a string.

Solution

You *could* use the Java `charAt` method:

```
scala> "hello".charAt(0)
res0: Char = h
```

However, the preferred approach is to use Scala’s Array notation:

```
scala> "hello"(0)
res1: Char = h
```

```
scala> "hello"(1)
res2: Char = e
```

Discussion

When looping over the characters in a string, you’ll normally use the `map` or `foreach` methods, but if for some reason those approaches won’t work for your situation, you can treat a `String` as an `Array`, and access each character with the array notation shown.

The Scala array notation is different than Java because in Scala it’s really a method call, with some nice syntactic sugar added. You write your code like this, which is convenient and easy to read:

```
scala> "hello"(1)
res0: Char = e
```

But behind the scenes, Scala converts your code into this:

```
scala> "hello".apply(1)
res1: Char = e
```

This little bit of syntactic sugar is explained in detail in [Recipe 6.8, “Creating Object Instances Without Using the new Keyword”](#).

1.10. Add Your Own Methods to the String Class

Problem

Rather than create a separate library of `String` utility methods, like a `StringUtilities` class, you want to add your own behavior(s) to the `String` class, so you can write code like this:

```
"HAL".increment
```

Instead of this:

```
StringUtilities.increment("HAL")
```

Solution

In Scala 2.10, you define an implicit class, and then define methods within that class to implement the behavior you want.

You can see this in the REPL. First, define your implicit class and method(s):

```
scala> implicit class StringImprovements(s: String) {  
  | def increment = s.map(c => (c + 1).toChar)  
  | }  
defined class StringImprovements
```

Then invoke your method on any `String`:

```
scala> val result = "HAL".increment  
result: String = IBM
```

In real-world code, this is just slightly more complicated. According to [SIP-13, Implicit Classes](#), “An implicit class must be defined in a scope where method definitions are allowed (not at the top level).” This means that your implicit class must be defined inside a class, object, or package object.

Put the implicit class in an object

One way to satisfy this condition is to put the implicit class inside an object. For instance, you can place the `StringImprovements` implicit class in an object such as a `StringUtils` object, as shown here:

```
package com.alvinalexander.utils  
  
object StringUtils {  
  implicit class StringImprovements(val s: String) {  
    def increment = s.map(c => (c + 1).toChar)  
  }  
}
```

```
}
}
```

You can then use the `increment` method somewhere else in your code, after adding the proper `import` statement:

```
package foo.bar

import com.alvinalexander.utils.StringUtils._

object Main extends App {
  println("HAL".increment)
}
```

Put the implicit class in a package object

Another way to satisfy the requirement is to put the implicit class in a *package object*. With this approach, place the following code in a file named *package.scala*, in the appropriate directory. If you're using SBT, you should place the file in the *src/main/scala/com/alvinalexander* directory of your project, containing the following code:

```
package com.alvinalexander

package object utils {

  implicit class StringImprovements(val s: String) {
    def increment = s.map(c => (c + 1).toChar)
  }

}
```

When you need to use the `increment` method in some other code, use a slightly different `import` statement from the previous example:

```
package foo.bar

import com.alvinalexander.utils._

object MainDriver extends App {
  println("HAL".increment)
}
```



See [Recipe 6.7, “Putting Common Code in Package Objects”](#) for more information about package objects.

Using versions of Scala prior to version 2.10

If for some reason you need to use a version of Scala prior to version 2.10, you'll need to take a slightly different approach. In this case, define a method named `increment` in a normal Scala class:

```
class StringImprovements(val s: String) {  
  def increment = s.map(c => (c + 1).toChar)  
}
```

Next, define another method to handle the implicit conversion:

```
implicit def stringToString(s: String) = new StringImprovements(s)
```

The `String` parameter in the `stringToString` method essentially links the `String` class to the `StringImprovements` class.

Now you can use `increment` as in the earlier examples:

```
"HAL".increment
```

Here's what this looks like in the REPL:

```
scala> class StringImprovements(val s: String) {  
      |   def increment = s.map(c => (c + 1).toChar)  
      | }  
defined class StringImprovements  
  
scala> implicit def stringToString(s: String) = new StringImprovements(s)  
stringToString: (s: String)StringImprovements  
  
scala> "HAL".increment  
res0: String = IBM
```

Discussion

As you just saw, in Scala, you can add new functionality to closed classes by writing implicit conversions and bringing them into scope when you need them. A major benefit of this approach is that you don't have to extend existing classes to add the new functionality. For instance, there's no need to create a new class named `MyString` that extends `String`, and then use `MyString` throughout your code instead of `String`; instead, you define the behavior you want, and then add that behavior to *all* `String` objects in the current scope when you add the `import` statement.

Note that you can define as many methods as you need in your implicit class. The following code shows both `increment` and `decrement` methods, along with a method named `hideAll` that returns a `String` with all characters replaced by the `*` character:

```
implicit class StringImprovements(val s: String) {  
  def increment = s.map(c => (c + 1).toChar)  
  def decrement = s.map(c => (c - 1).toChar)
```

```
def hideAll = s.replaceAll(".", "*")
}
```

Notice that except for the `implicit` keyword before the class name, the `StringImprovements` class and its methods are written as usual.

By simply bringing the code into scope with an `import` statement, you can use these methods, as shown here in the REPL:

```
scala> "HAL".increment
res0: String = IBM
```

Here's a simplified description of how this works:

1. The compiler sees a string literal "HAL."
2. The compiler sees that you're attempting to invoke a method named `increment` on the `String`.
3. Because the compiler can't find that method on the `String` class, it begins looking around for implicit conversion methods that are in scope and accepts a `String` argument.
4. This leads the compiler to the `StringImprovements` class, where it finds the `increment` method.

That's an oversimplification of what happens, but it gives you the general idea of how implicit conversions work.



For more details on what's happening here, see [SIP-13, Implicit Classes](#).

Annotate your method return type

It's recommended that the return type of implicit method definitions should be annotated. If you run into a situation where the compiler can't find your implicit methods, or you just want to be explicit when declaring your methods, add the return type to your method definitions.

In the `increment`, `decrement`, and `hideAll` methods shown here, the return type of `String` is made explicit:

```
implicit class StringImprovements(val s: String) {
  // being explicit that each method returns a String
  def increment: String = s.map(c => (c + 1).toChar)
  def decrement: String = s.map(c => (c - 1).toChar)
  def hideAll: String = s.replaceAll(".", "*")
}
```

Returning other types

Although all of the methods shown so far have returned a `String`, you can return any type from your methods that you need. The following class demonstrates several different types of string conversion methods:

```
implicit class StringImprovements(val s: String) {  
  def increment = s.map(c => (c + 1).toChar)  
  def decrement = s.map(c => (c - 1).toChar)  
  def hideAll: String = s.replaceAll(".", "*")  
  def plusOne = s.toInt + 1  
  def asBoolean = s match {  
    case "0" | "zero" | "" | " " => false  
    case _ => true  
  }  
}
```

With these new methods you can now perform `Int` and `Boolean` conversions, in addition to the `String` conversions shown earlier:

```
scala> "4".plusOne  
res0: Int = 5  
  
scala> "0".asBoolean  
res1: Boolean = false  
  
scala> "1".asBoolean  
res2: Boolean = true
```

Note that all of these methods have been simplified to keep them short and readable. In the real world, you'll want to add some error-checking.

Introduction

In Scala, all the numeric types are objects, including `Byte`, `Char`, `Double`, `Float`, `Int`, `Long`, and `Short`. These seven numeric types extend the `AnyVal` trait, as do the `Unit` and `Boolean` classes, which are considered to be “nonnumeric value types.”

As shown in [Table 2-1](#), the seven built-in numeric types have the same data ranges as their Java primitive equivalents.

Table 2-1. Data ranges of Scala’s built-in numeric types

Data type	Range
<code>Char</code>	16-bit unsigned Unicode character
<code>Byte</code>	8-bit signed value
<code>Short</code>	16-bit signed value
<code>Int</code>	32-bit signed value
<code>Long</code>	64-bit signed value
<code>Float</code>	32-bit IEEE 754 single precision float
<code>Double</code>	64-bit IEEE 754 single precision float

In addition to those types, `Boolean` can have the values `true` or `false`.

If you ever need to know the exact values of the data ranges, you can find them in the Scala REPL:

```
scala> Short.MinValue
res0: Short = -32768

scala> Short.MaxValue
res1: Short = 32767
```

```
scala> Int.MinValue
res2: Int = -2147483648

scala> Float.MinValue
res3: Float = -3.4028235E38
```

In addition to these basic numeric types, it's helpful to understand the `BigInt` and `BigDecimal` classes, as well as the methods in the `scala.math` package. These are all covered in this chapter.

Complex Numbers and Dates

If you need more powerful math classes than those that are included with the standard Scala distribution, check out the [Spire project](#), which includes classes like `Rational`, `Complex`, `Real`, and more; and [ScalaLab](#), which offers Matlab-like scientific computing in Scala.

For processing dates, the Java [Joda Time](#) project is popular and well documented. A project named `nscala-time` implements a Scala wrapper around Joda Time, and lets you write date expressions in a more Scala-like way, including these examples:

```
DateTime.now // returns org.joda.time.DateTime
DateTime.now + 2.months
DateTime.nextMonth < DateTime.now + 2.months
(2.hours + 45.minutes + 10.seconds).millis
```

2.1. Parsing a Number from a String

Problem

You want to convert a `String` to one of Scala's numeric types.

Solution

Use the `to*` methods that are available on a `String` (courtesy of the `StringLike` trait):

```
scala> "100".toInt
res0: Int = 100

scala> "100".toDouble
res1: Double = 100.0

scala> "100".toFloat
res2: Float = 100.0

scala> "1".toLong
res3: Long = 1

scala> "1".toShort
```

```
res4: Short = 1
```

```
scala> "1".toByte  
res5: Byte = 1
```

Be careful, because these methods can throw the usual Java `NumberFormatException`:

```
scala> "foo".toInt  
java.lang.NumberFormatException: For input string: "foo"  
    at java.lang.NumberFormatException.forInputString(NumberFormatException.java)  
    at java.lang.Integer.parseInt(Integer.java:449)  
    ... more output here ...
```

`BigInt` and `BigDecimal` instances can also be created directly from strings (and can also throw a `NumberFormatException`):

```
scala> val b = BigInt("1")  
b: scala.math.BigInt = 1  
  
scala> val b = BigDecimal("3.14159")  
b: scala.math.BigDecimal = 3.14159
```

Handling a base and radix

If you need to perform calculations using bases other than 10, you'll find the `toInt` method in the Scala `Int` class doesn't have a method that lets you pass in a base and radix. To solve this problem, use the `parseInt` method in the `java.lang.Integer` class, as shown in these examples:

```
scala> Integer.parseInt("1", 2)  
res0: Int = 1  
  
scala> Integer.parseInt("10", 2)  
res1: Int = 2  
  
scala> Integer.parseInt("100", 2)  
res2: Int = 4  
  
scala> Integer.parseInt("1", 8)  
res3: Int = 1  
  
scala> Integer.parseInt("10", 8)  
res4: Int = 8
```

If you're a fan of implicit conversions, you can create an implicit class and method to help solve the problem. As described in [Recipe 1.10, "Add Your Own Methods to the String Class"](#) create the implicit conversion as follows:

```
implicit class StringToInt(s: String) {  
  def toInt(radix: Int) = Integer.parseInt(s, radix)  
}
```

Defining this implicit class (and bringing it into scope) adds a `toInt` method that takes a radix argument to the `String` class, which you can now call instead of calling `Integer.parseInt`:

```
scala> implicit class StringToInt(s: String) {  
  | def toInt(radix: Int) = Integer.parseInt(s, radix)  
  | }  
defined class StringToInt  
  
scala> "1".toInt(2)  
res0: Int = 1  
  
scala> "10".toInt(2)  
res1: Int = 2  
  
scala> "100".toInt(2)  
res2: Int = 4  
  
scala> "100".toInt(8)  
res3: Int = 64  
  
scala> "100".toInt(16)  
res4: Int = 256
```

See [Recipe 1.10](#) for more details on how to implement this solution outside of the REPL.

Discussion

If you've used Java to convert a `String` to a numeric data type, then the `NumberFormatException` is familiar. However, Scala doesn't have checked exceptions, so you'll probably want to handle this situation differently.

First, you don't have to declare that Scala methods can throw an exception, so it's perfectly legal to declare a Scala method like this:

```
// not required to declare "throws NumberFormatException"  
def toInt(s: String) = s.toInt
```

If you're going to allow an exception to be thrown like this, callers of your method might appreciate knowing that this can happen. Consider adding a Scaladoc comment to your method in this case.

If you prefer to declare that your method can throw an exception, mark it with the `@throws` annotation, as shown here:

```
@throws(classOf[NumberFormatException])  
def toInt(s: String) = s.toInt
```

This approach is required if the method will be called from Java code, as described in [Recipe 17.2, "Add Exception Annotations to Scala Methods to Work with Java"](#).

However, in Scala, situations like this are often handled with the `Option/Some/None` pattern, as described in [Recipe 20.6, “Using the Option/Some/None Pattern”](#). With this approach, define the `toInt` method like this:

```
def toInt(s: String):Option[Int] = {  
  try {  
    Some(s.toInt)  
  } catch {  
    case e: NumberFormatException => None  
  }  
}
```

Now you can call the `toInt` method in several different ways, depending on your needs. One way is with `getOrElse`:

```
println(toInt("1").getOrElse(0)) // 1  
println(toInt("a").getOrElse(0)) // 0  
  
// assign the result to x  
val x = toInt(aString).getOrElse(0)
```

Another approach is to use a match expression. You can write a match expression to print the `toInt` result like this:

```
toInt(aString) match {  
  case Some(n) => println(n)  
  case None => println("Boom! That wasn't a number.")  
}
```

You can also write a match expression as follows to assign the result to a variable:

```
val result = toInt(aString) match {  
  case Some(x) => x  
  case None => 0 // however you want to handle this  
}
```

If these examples haven't yet sold you on the `Option/Some/None` approach, you'll see in [Chapter 10](#) and [Chapter 11](#) that this pattern is incredibly helpful and convenient when working with collections.

Alternatives to Option

If you like the `Option/Some/None` concept, but need access to the exception information, there are several additional possibilities:

- Try, Success, and Failure (introduced in Scala 2.10)
- Either, Left, and Right

These alternate approaches are discussed in [Recipe 20.6, “Using the Option/Some/None Pattern”](#). (The new Try/Success/Failure approach is especially appealing.)

See Also

- [Recipe 20.6, “Using the Option/Some/None Pattern”](#)
- [The StringLike trait](#)

2.2. Converting Between Numeric Types (Casting)

Problem

You want to convert from one numeric type to another, such as from an `Int` to a `Double`.

Solution

Instead of using the “cast” approach in Java, use the `to*` methods that are available on all numeric types. These methods can be demonstrated in the REPL (note that you need to hit `Tab` at the end of the first example):

```
scala> val b = a.to[Tab]
toByte      toChar    toDouble    toFloat    toInt      toLong
toShort     toString

scala> 19.45.toInt
res0: Int = 19

scala> 19.toFloat
res1: Float = 19.0

scala> 19.toDouble
res2: Double = 19.0

scala> 19.toLong
res3: Long = 19

scala> val b = a.toFloat
b: Float = 1945.0
```

Discussion

In Java, you convert from one numeric type to another by casting the types, like this:

```
int a = (int) 100.00;
```

But in Scala, you use the `to*` methods, as shown in this recipe.

If you want to avoid potential conversion errors when casting from one numeric type to another, you can use the related `isValid` methods to test whether the type can be

converted before attempting the conversion. For instance, a `Double` object (via `RichDouble`) has methods like `isValidInt` and `isValidShort`:

```
scala> val a = 1000L
a: Long = 1000

scala> a.isValidByte
res0: Boolean = false

scala> a.isValidShort
res1: Boolean = true
```

See Also

[The `RichDouble` class](#)

2.3. Overriding the Default Numeric Type

Problem

Scala automatically assigns types to numeric values when you assign them, and you need to override the default type it assigns as you create a numeric field.

Solution

If you assign `1` to a variable, Scala assigns it the type `Int`:

```
scala> val a = 1
a: Int = 1
```

The following examples show one way to override simple numeric types:

```
scala> val a = 1d
a: Double = 1.0

scala> val a = 1f
a: Float = 1.0

scala> val a = 1000L
a: Long = 1000
```

Another approach is to annotate the variable with a type, like this:

```
scala> val a = 0: Byte
a: Byte = 0

scala> val a = 0: Int
a: Int = 0

scala> val a = 0: Short
a: Short = 0
```

```
scala> val a = 0: Double
a: Double = 0.0
```

```
scala> val a = 0: Float
a: Float = 0.0
```

Spacing after the colon isn't important, so you can use this format, if preferred:

```
val a = 0:Byte
```

According to the [Scala Style Guide](#), those examples show the preferred style for annotating types, but personally I prefer the following syntax when assigning types to variables, specifying the type after the variable name:

```
scala> val a:Byte = 0
a: Byte = 0
```

```
scala> val a:Int = 0
a: Int = 0
```

You can create hex values by preceding the number with a leading 0x or 0X, and you can store them as an Int or Long:

```
scala> val a = 0x20
a: Int = 32
```

```
// if you want to store the value as a Long
scala> val a = 0x20L
a: Long = 32
```



In some rare instances, you may need to take advantage of *type ascription*. [Stack Overflow](#) shows a case where it's advantageous to upcast a String to an Object. The technique is shown here:

```
scala> val s = "Dave"
s: String = Dave
```

```
scala> val p = s: Object
p: Object = Dave
```

As you can see, the technique is similar to this recipe. This upcasting is known as *type ascription*. The official Scala documentation describes type ascription as follows:

Ascription is basically just an up-cast performed at compile time for the sake of the type checker. Its use is not common, but it does happen on occasion. The most often seen case of ascription is invoking a `varargs` method with a single Seq parameter.

Discussion

It's helpful to know about this approach when creating object instances. The general syntax looks like this:

```
// general case
var [name]:[Type] = [initial value]

// example
var a:Short = 0
```

This form can be helpful when you need to initialize numeric var fields in a class:

```
class Foo {
  var a: Short = 0      // specify a default value
  var b: Short = _      // defaults to 0
}
```

As shown, you can use the underscore character as a placeholder when assigning an initial value. This works when creating class variables, but doesn't work in other places, such as inside a method. For numeric types this isn't an issue—you can just assign the type the value zero—but with most other types, you can use this approach inside a method:

```
var name = null.asInstanceOf[String]
```

Better yet, use the Option/Some/None pattern. It helps eliminate null values from your code, which is a very good thing. You'll see this pattern used in the best Scala libraries and frameworks, such as the Play Framework. An excellent example of this approach is shown in [Recipe 12.4, “How to Process Every Character in a Text File”](#).

See [Recipe 20.5, “Eliminate null Values from Your Code”](#) and [Recipe 20.6, “Using the Option/Some/None Pattern”](#) for more discussion of this important topic.

See Also

- [The Scala Style Guide](#)
- The [Stack Overflow URL](#) mentioned in the note in the Solution

2.4. Replacements for ++ and --

Problem

You want to increment or decrement numbers using operators like ++ and -- that are available in other languages, but Scala doesn't have these operators.

Solution

Because `val` fields are immutable, they can't be incremented or decremented, but `var` `Int` fields can be mutated with the `+=` and `-=` methods:

```
scala> var a = 1
a: Int = 1

scala> a += 1

scala> println(a)
2

scala> a -= 1

scala> println(a)
1
```

As an added benefit, you use similar methods for multiplication and division:

```
scala> var i = 1
i: Int = 1

scala> i *= 2

scala> println(i)
2

scala> i *= 2

scala> println(i)
4

scala> i /= 2

scala> println(i)
2
```

Note that these symbols aren't operators; they're implemented as methods that are available on `Int` fields declared as a `var`. Attempting to use them on `val` fields results in a compile-time error:

```
scala> val x = 1
x: Int = 1

scala> x += 1
<console>:9: error: value += is not a member of Int
      x += 1
        ^
```



As mentioned, the symbols `+=`, `-=`, `*=`, and `/=` aren't operators, they're *methods*. This approach of building functionality with libraries instead of operators is a consistent pattern in Scala. Actors, for instance, are not built into the language, but are instead implemented as a library. See the Dr. Dobbs link in the See Also for Martin Odersky's discussion of this philosophy.

Discussion

Another benefit of this approach is that you can call methods of the same name on other types besides `Int`. For instance, the `Double` and `Float` classes have methods of the same name:

```
scala> var x = 1d
x: Double = 1.0
```

```
scala> x += 1
```

```
scala> println(x)
2.0
```

```
scala> var x = 1f
x: Float = 1.0
```

```
scala> x += 1
```

```
scala> println(x)
2.0
```

See Also

Martin Odersky discusses how Actors are added into Scala as a library on drdobbs.com.

2.5. Comparing Floating-Point Numbers

Problem

You need to compare two floating-point numbers, but as in some other programming languages, two floating-point numbers that *should* be equivalent may not be.

Solution

As in Java and many other languages, you solve this problem by creating a method that lets you specify the precision for your comparison. The following “approximately equals” method demonstrates the approach:

```
def ==(x: Double, y: Double, precision: Double) = {
  if ((x - y).abs < precision) true else false
}
```

You can use this method like this:

```
scala> val a = 0.3
a: Double = 0.3

scala> val b = 0.1 + 0.2
b: Double = 0.30000000000000004

scala> ==(a, b, 0.0001)
res0: Boolean = true

scala> ==(b, a, 0.0001)
res1: Boolean = true
```

Discussion

When you begin working with floating-point numbers, you quickly learn that 0.1 plus 0.1 is 0.2:

```
scala> 0.1 + 0.1
res38: Double = 0.2
```

But 0.1 plus 0.2 isn't exactly 0.3:

```
scala> 0.1 + 0.2
res37: Double = 0.30000000000000004
```

This subtle inaccuracy makes comparing two floating-point numbers a real problem:

```
scala> val a = 0.3
a: Double = 0.3

scala> val b = 0.1 + 0.2
b: Double = 0.30000000000000004

scala> a == b
res0: Boolean = false
```

As a result, you end up writing your own functions to compare floating-point numbers with a precision (or tolerance).

As you saw in Recipe 1.11, you can define an implicit conversion to add a method like this to the Double class. This makes the following code very readable:

```
if (a ~= b) ...
```

Or, you can add the same method to a utilities object, if you prefer:

```
object MathUtils {
  def ==(x: Double, y: Double, precision: Double) = {
    if ((x - y).abs < precision) true else false
  }
}
```



```
}  
}
```

which you can then invoke like a static method:

```
println(MathUtils.~=(a, b, 0.000001))
```

With an implicit conversion, the name `~=` is very readable, but in a utilities object like this, it doesn't look quite right, so it might be better named `approximatelyEqual`, `equalWithinTolerance`, or some other name.

See Also

- [Floating-point accuracy problems](#)
- [Arbitrary-precision arithmetic](#)
- What every computer scientist should know about [floating-point arithmetic](#)

2.6. Handling Very Large Numbers

Problem

You're writing an application and need to use very large integer or decimal numbers.

Solution

Use the Scala `BigInt` and `BigDecimal` classes. You can create a `BigInt`:

```
scala> var b = BigInt(1234567890)  
b: scala.math.BigInt = 1234567890
```

or a `BigDecimal`:

```
scala> var b = BigDecimal(123456.789)  
b: scala.math.BigDecimal = 123456.789
```

Unlike their Java equivalents, these classes support all the operators you're used to using with numeric types:

```
scala> b + b  
res0: scala.math.BigInt = 2469135780  
  
scala> b * b  
res1: scala.math.BigInt = 1524157875019052100  
  
scala> b += 1  
  
scala> println(b)  
1234567891
```

You can convert them to other numeric types:

```
scala> b.toInt
res2: Int = 1234567891

scala> b.toLong
res3: Long = 1234567891

scala> b.toFloat
res4: Float = 1.23456794E9

scala> b.toDouble
res5: Double = 1.234567891E9
```

To help avoid errors, you can also test them first to see if they can be converted to other numeric types:

```
scala> b.isValidByte
res6: Boolean = false

scala> b.isValidChar
res7: Boolean = false

scala> b.isValidShort
res8: Boolean = false

scala> if (b.isValidInt) b.toInt
res9: AnyVal = 1234567890
```

Discussion

Although the Scala `BigInt` and `BigDecimal` classes are backed by the Java `BigInteger` and `BigDecimal` classes, they are simpler to use than their Java counterparts. As you can see in the examples, they work just like other numeric types, and they're also mutable (as you saw in the `+=` example). These are nice improvements over the Java classes.

Before using `BigInt` or `BigDecimal`, you can check the maximum values that the other Scala numeric types can handle in [Table 1-1](#), or by checking their `MaxValue` in the REPL:

```
scala> Byte.MaxValue
res0: Byte = 127

scala> Short.MaxValue
res1: Short = 32767

scala> Int.MaxValue
res2: Int = 2147483647

scala> Long.MaxValue
res3: Long = 9223372036854775807
```

```
scala> Double.MaxValue  
res4: Double = 1.7976931348623157E308
```

Depending on your needs, you may also be able to use the `PositiveInfinity` and `NegativeInfinity` of the standard numeric types:

```
scala> Double.PositiveInfinity  
res0: Double = Infinity
```

```
scala> Double.NegativeInfinity  
res1: Double = -Infinity
```

```
scala> 1.7976931348623157E308 > Double.PositiveInfinity  
res45: Boolean = false
```

See Also

- [The Java `BigInteger` class](#)
- [The Scala `BigInt` class](#)
- [The Scala `BigDecimal` class](#)

2.7. Generating Random Numbers

Problem

You need to create random numbers, such as when testing an application, performing a simulation, and many other situations.

Solution

Create random numbers with the Scala `scala.util.Random` class. You can create random integers:

```
scala> val r = scala.util.Random  
r: scala.util.Random = scala.util.Random@13eb41e5
```

```
scala> r.nextInt  
res0: Int = -1323477914
```

You can limit the random numbers to a maximum value:

```
scala> r.nextInt(100)  
res1: Int = 58
```

In this use, the `Int` returned is between 0 (inclusive) and the value you specify (exclusive), so specifying 100 returns an `Int` from 0 to 99.

You can also create random `Float` values:

```
// returns a value between 0.0 and 1.0
scala> r.nextFloat
res2: Float = 0.50317204
```

You can create random Double values:

```
// returns a value between 0.0 and 1.0
scala> r.nextDouble
res3: Double = 0.6946000981900997
```

You can set the seed value using an Int or Long when creating the Random object:

```
scala> val r = new scala.util.Random(100)
r: scala.util.Random = scala.util.Random@bbf4061
```

You can also set the seed value after a Random object has been created:

```
r.setSeed(1000L)
```

Discussion

The `Random` class handles all the usual use cases, including creating numbers, setting the maximum value of a random number range, and setting a seed value. You can also generate random characters:

```
// random characters
scala> r.nextPrintableChar
res0: Char = H
```

```
scala> r.nextPrintableChar
res1: Char = r
```

Scala makes it easy to create a random-length range of numbers, which is especially useful for testing:

```
// create a random length range
scala> var range = 0 to r.nextInt(10)
range: scala.collection.immutable.Range.Inclusive = Range(0, 1, 2, 3)
```

```
scala> range = 0 to r.nextInt(10)
range: scala.collection.immutable.Range.Inclusive = Range(0, 1)
```

You can add a `for`/`yield` loop to modify the numbers:

```
scala> for (i <- 0 to r.nextInt(10)) yield i * 2
res0: scala.collection.immutable.IndexedSeq[Int] = Vector(0, 2, 4)
```

You can easily create random-length ranges of other types. Here's a random-length collection of up to 10 `Float` values:

```
scala> for (i <- 0 to r.nextInt(10)) yield (i * r.nextFloat)
res1: scala.collection.immutable.IndexedSeq[Float] =
  Vector(0.0, 0.71370363, 1.0783684)
```

Here's a random-length collection of "printable characters":

```
scala> for (i <- 0 to r.nextInt(10)) yield r.nextPrintableChar
res2: scala.collection.immutable.IndexedSeq[Char] = Vector(x, K, ^, z, w)
```

Be careful with the `nextPrintableChar` method. A better approach may be to control the characters you use, as shown in my “How to create a list of alpha or alphanumeric characters” article, shown in the See Also.

Conversely, you can create a sequence of known length, filled with random numbers:

```
scala> for (i <- 1 to 5) yield r.nextInt(100)
res3: scala.collection.immutable.IndexedSeq[Int] = Vector(88, 94, 58, 96, 82)
```

See Also

- [The Scala Random class](#)
- [Recipe 11.29, “Using a Range”](#), provides examples of how to create and use ranges
- [My article](#) on how to create a list of alpha or alphanumeric characters
- An additional recipe for [generating random strings](#)

2.8. Creating a Range, List, or Array of Numbers

Problem

You need to create a range, list, or array of numbers, such as in a for loop, or for testing purposes.

Solution

Use the `to` method of the `Int` class to create a `Range` with the desired elements:

```
scala> val r = 1 to 10
r: scala.collection.immutable.Range.Inclusive = Range(1, 2, 3, 4, 5,
6, 7, 8, 9, 10)
```

You can set the step with the `by` method:

```
scala> val r = 1 to 10 by 2
r: scala.collection.immutable.Range = Range(1, 3, 5, 7, 9)
```

```
scala> val r = 1 to 10 by 3
r: scala.collection.immutable.Range = Range(1, 4, 7, 10)
```

Ranges are commonly used in for loops:

```
scala> for (i <- 1 to 5) println(i)
1
2
3
```

```
4  
5
```

When creating a Range, you can also use `until` instead of `to`:

```
scala> for (i <- 1 until 5) println(i)  
1  
2  
3  
4
```

Discussion

Scala makes it easy to create a range of numbers. The first three examples shown in the Solution create a Range. You can easily convert a Range to other sequences, such as an Array or List, like this:

```
scala> val x = 1 to 10 toArray  
x: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)  
  
scala> val x = 1 to 10 toList  
x: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

Although this *infix notation* syntax is clear in many situations (such as for loops), it's generally preferable to use this syntax:

```
scala> val x = (1 to 10).toList  
x: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)  
  
scala> val x = (1 to 10).toArray  
x: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

The magic that makes this process work is the `to` and `until` methods, which you'll find in the `RichInt` class. When you type the following portion of the code, you're actually invoking the `to` method of the `RichInt` class:

```
1 to
```

You can demonstrate that `to` is a method on an `Int` by using this syntax in the REPL:

```
1.to(10)
```



Although the infix notation `(1 to 10)` shown in most of these examples can make your code more readable, Rahul Phulore has a post on Stack Overflow where he advises against using it for anything other than internal DSLs. The link to that post is shown in the See Also.

Combine this with [Recipe 2.7, “Generating Random Numbers”](#) and you can create a random-length range, which can be useful for testing:

```
scala> var range = 0 to scala.util.Random.nextInt(10)
range: scala.collection.immutable.Range.Inclusive = Range(0, 1, 2, 3)
```

By using a range with the `for/yield` construct, you don't have to limit your ranges to sequential numbers:

```
scala> for (i <- 1 to 5) yield i * 2
res0: scala.collection.immutable.IndexedSeq[Int] = Vector(2, 4, 6, 8, 10)
```

You also don't have to limit your ranges to just integers:

```
scala> for (i <- 1 to 5) yield i.toDouble
res1: scala.collection.immutable.IndexedSeq[Double] =
  Vector(1.0, 2.0, 3.0, 4.0, 5.0)
```

See Also

- [The Scala `RichInt` class](#)
- Rahul Phulore's [post](#), where he advises not using the infix notation

2.9. Formatting Numbers and Currency

Problem

You want to format numbers or currency to control decimal places and commas, typically for printed output.

Solution

For basic number formatting, use the `f` string interpolator shown in [Recipe 1.4, “Substituting Variables into Strings”](#):

```
scala> val pi = scala.math.Pi
pi: Double = 3.141592653589793
```

```
scala> println(f"$pi%1.5f")
3.14159
```

A few more examples demonstrate the technique:

```
scala> f"$pi%1.5f"
res0: String = 3.14159
```

```
scala> f"$pi%1.2f"
res1: String = 3.14
```

```
scala> f"$pi%06.2f"
res2: String = 003.14
```

If you're using a version of Scala prior to 2.10, or prefer the explicit use of the `format` method, you can write the code like this instead:

```
scala> "%06.2f".format(pi)
res3: String = 003.14
```

A simple way to add commas is to use the `getIntegerInstance` method of the `java.text.NumberFormat` class:

```
scala> val formatter = java.text.NumberFormat.getIntegerInstance
formatter: java.text.NumberFormat = java.text.DecimalFormat@674dc

scala> formatter.format(10000)
res0: String = 10,000

scala> formatter.format(1000000)
res1: String = 1,000,000
```

You can also set a locale with the `getIntegerInstance` method:

```
scala> val locale = new java.util.Locale("de", "DE")
locale: java.util.Locale = de_DE

scala> val formatter = java.text.NumberFormat.getIntegerInstance(locale)
formatter: java.text.NumberFormat = java.text.DecimalFormat@674dc

scala> formatter.format(1000000)
res2: String = 1.000.000
```

You can handle floating-point values with a formatter returned by `getInstance`:

```
scala> val formatter = java.text.NumberFormat.getInstance
formatter: java.text.NumberFormat = java.text.DecimalFormat@674dc

scala> formatter.format(10000.33)
res0: String = 10,000.33
```

For currency output, use the `getCurrencyInstance` formatter:

```
scala> val formatter = java.text.NumberFormat.getCurrencyInstance
formatter: java.text.NumberFormat = java.text.DecimalFormat@67500

scala> println(formatter.format(123.456789))
$123.46

scala> println(formatter.format(1234.56789))
$1,234.57

scala> println(formatter.format(12345.6789))
$12,345.68

scala> println(formatter.format(123456.789))
$123,456.79
```

This approach handles international currency:


```
scala> import java.util.{Currency, Locale}
import java.util.{Currency, Locale}

scala> val de = Currency.getInstance(new Locale("de", "DE"))
de: java.util.Currency = EUR

scala> formatter.setCurrency(de)

scala> println(formatter.format(123456.789))
EUR123,456.79
```

Discussion

This recipe falls back to the Java approach for printing currency and other formatted numeric fields, though of course the currency solution depends on how you handle currency in your applications. In my work as a consultant, I've seen most companies handle currency using the Java `BigDecimal` class, and others create their own custom currency classes, which are typically wrappers around `BigDecimal`.

See Also

- [My printf cheat sheet](#).
- The [Joda Money library](#) is a Java library for handling currency, and is currently at version 0.8.
- JSR 354: Money and Currency API, is also being developed in the Java Community Process. See jcp.org for more information.

Control Structures

Introduction

The control structures in Scala start off similar to their Java counterparts, and then diverge in some wonderful ways. For instance, Scala's `if/then/else` structure is similar to Java, but can also be used to return a value. As a result, though Java has a special syntax for a ternary operator, in Scala you just use a normal `if` statement to achieve the ternary effect:

```
val x = if (a) y else z
```

The `try/catch/finally` structure is similar to Java, though Scala uses pattern matching in the `catch` clause. This differs from Java, but because it's consistent with other uses of pattern matching in Scala, it's easy to remember.

When you get to the `for` loop, things really start to get interesting. Its basic use is similar to Java, but with the addition of *guards* and other conveniences, the Scala `for` loop rapidly departs from its Java counterpart. For instance, in Scala you *could* write two `for` loops as follows to read every line in a file and then operate on each character in each line:

```
for (line <- source.getLines) {  
  for {  
    char <- line  
    if char.isLetter  
  } // char algorithm here ...  
}
```

But with Scala's `for` loop mojo, you can write this code even more concisely:

```
for {  
  line <- source.getLines  
  char <- line  
  if char.isLetter  
} // char algorithm here ...
```

The rabbit hole goes even deeper, because a Scala *for comprehension* lets you easily apply an algorithm to one collection to generate a new collection:

```
scala> val nieces = List("emily", "hannah", "mercedes", "porsche")
nieces: List[String] = List(emily, hannah, mercedes, porsche)

scala> for (n <- nieces) yield n.capitalize
res0: List[String] = List(Emily, Hannah, Mercedes, Porsche)
```

Similarly, in its most basic use, a Scala *match expression* can look like a Java switch statement, but because you can match any object, extract information from matched objects, add guards to case statements, return values, and more, match expressions are a major feature of the Scala language.

3.1. Looping with for and foreach

Problem

You want to iterate over the elements in a collection, either to operate on each element in the collection, or to create a new collection from the existing collection.

Solution

There are many ways to loop over Scala collections, including for loops, while loops, and collection methods like `foreach`, `map`, `flatMap`, and more. This solution focuses primarily on the for loop and `foreach` method.

Given a simple array:

```
val a = Array("apple", "banana", "orange")
```

I prefer to iterate over the array with the following for loop syntax, because it's clean and easy to remember:

```
scala> for (e <- a) println(e)
apple
banana
orange
```

When your algorithm requires multiple lines, use the same for loop syntax, and perform your work in a block:

```
scala> for (e <- a) {
  |   // imagine this requires multiple lines
  |   val s = e.toUpperCase
  |   println(s)
  | }
APPLE
BANANA
ORANGE
```

Returning values from a for loop

Those examples perform an operation using the elements in an array, but they don't return a value you can use, such as a new array. In cases where you want to build a new collection from the input collection, use the `for/yield` combination:

```
scala> val newArray = for (e <- a) yield e.toUpperCase
newArray: Array[java.lang.String] = Array(APPLE, BANANA, ORANGE)
```

The `for/yield` construct returns a value, so in this case, the array `newArray` contains uppercase versions of the three strings in the initial array. Notice that an input `Array` yields an `Array` (and not something else, like a `Vector`).

When your algorithm requires multiple lines of code, perform the work in a block after the `yield` keyword:

```
scala> val newArray = for (e <- a) yield {
  |   // imagine this requires multiple lines
  |   val s = e.toUpperCase
  |   s
  | }
newArray: Array[java.lang.String] = Array(APPLE, BANANA, ORANGE)
```

for loop counters

If you need access to a counter inside a `for` loop, use one of the following approaches. First, you can access array elements with a counter like this:

```
for (i <- 0 until a.length) {
  println(s"$i is ${a(i)}")
}
```

That loop yields this output:

```
0 is apple
1 is banana
2 is orange
```

Scala collections also offer a `zipWithIndex` method that you can use to create a loop counter:

```
scala> for ((e, count) <- a.zipWithIndex) {
  |   println(s"$count is $e")
  | }
0 is apple
1 is banana
2 is orange
```

See [Recipe 10.11, “Using `zipWithIndex` or `zip` to Create Loop Counters](#)”, for more examples of how to use `zipWithIndex`.

Generators and guards

On a related note, the following example shows how to use a `Range` to execute a loop three times:

```
scala> for (i <- 1 to 3) println(i)
1
2
3
```

The `1 to 3` portion of the loop creates a `Range`, as shown in the REPL:

```
scala> 1 to 3
res0: scala.collection.immutable.Range.Inclusive = Range(1, 2, 3)
```

Using a `Range` like this is known as using a *generator*. The next recipe demonstrates how to use this technique to create multiple loop counters.

Recipe 3.3 demonstrates how to use guards (if statements in for loops), but here's a quick preview:

```
scala> for (i <- 1 to 10 if i < 4) println(i)
1
2
3
```

Looping over a Map

When iterating over keys and values in a `Map`, I find this to be the most concise and readable for loop:

```
val names = Map("fname" -> "Robert",
                "lname" -> "Goren")
for ((k,v) <- names) println(s"key: $k, value: $v")
```

See **Recipe 11.17, “Traversing a Map”** for more examples of how to iterate over the elements in a `Map`.

Discussion

An important lesson from the for loop examples is that when you use the `for/yield` combination with a collection, you're building and returning a new collection, but when you use a for loop *without* `yield`, you're just operating on each element in the collection—you're not creating a new collection. The `for/yield` combination is referred to as a *for comprehension*, and in its basic use, it works just like the `map` method. It's discussed in more detail in **Recipe 3.4, “Creating a for Comprehension (for/yield Combination)”**.

In some ways Scala reminds me of the Perl slogan, “There's more than one way to do it,” and iterating over a collection provides some great examples of this. With the wealth of methods that are available on collections, it's important to note that a for loop may not even be the best approach to a particular problem; the methods `foreach`, `map`,

`flatMap`, `collect`, `reduce`, etc., can often be used to solve your problem without requiring an explicit `for` loop.

For example, when you're working with a collection, you can also iterate over each element by calling the `foreach` method on the collection:

```
scala> a.foreach(println)
apple
banana
orange
```

When you have an algorithm you want to run on each element in the collection, just use the anonymous function syntax:

```
scala> a.foreach(e => println(e.toUpperCase))
APPLE
BANANA
ORANGE
```

As before, if your algorithm requires multiple lines, perform your work in a block:

```
scala> a.foreach { e =>
  |   val s = e.toUpperCase
  |   println(s)
  | }
APPLE
BANANA
ORANGE
```

How `for` loops are translated

As you work with Scala, it's helpful to understand how `for` loops are translated by the compiler. The [Scala Language Specification](#) provides details on precisely how a `for` loop is translated under various conditions. I encourage you to read the Specification for details on the rules, but a simplification of those rules can be stated as follows:

1. A simple `for` loop that iterates over a collection is translated to a `foreach` method call on the collection.
2. A `for` loop with a *guard* (see [Recipe 3.3](#)) is translated to a sequence of a `withFilter` method call on the collection followed by a `foreach` call.
3. A `for` loop with a `yield` expression is translated to a `map` method call on the collection.
4. A `for` loop with a `yield` expression and a guard is translated to a `withFilter` method call on the collection, followed by a `map` method call.

Again, the Specification is more detailed than this, but those statements will help get you started in the right direction.

These statements can be demonstrated with a series of examples. Each of the following examples starts with a `for` loop, and the code in each example will be compiled with the following `scalac` command:

```
$ scalac -Xprint:parse Main.scala
```

This command provides some initial output about how the Scala compiler translates the `for` loops into other code.

As a first example, start with the following code in a file named *Main.scala*:

```
class Main {  
  for (i <- 1 to 10) println(i)  
}
```

This code is intentionally small and trivial so you can see how the `for` loop is translated by the compiler.

When you compile this code with the `scalac -Xprint:parse` command, the full output looks like this:

```
$ scalac -Xprint:parse Main.scala  
  
[[syntax trees at end of parser]] // Main.scala  
package <empty> {  
  class Main extends scala.AnyRef {  
    def <init>() = {  
      super.<init>();  
      ()  
    };  
    1.to(10).foreach(((i) => println(i)))  
  }  
}
```

For this example, the important part of the output is the area that shows the `for` loop was translated by the compiler into the following code:

```
1.to(10).foreach(((i) => println(i)))
```

As you can see, the Scala compiler translates a simple `for` loop over a collection into a `foreach` method call on the collection.



If you compile the file with the `-Xprint:all` option instead of `-Xprint:parse`, you'll see that the code is further translated into the following code:

```
scala.this.Predef.intWrapper(1).to(10).foreach[Unit]  
(((i: Int) => scala.this.Predef.println(i)))
```

The code continues to get more and more detailed as the compiler phases continue, but for this demonstration, only the first step in the translation process is necessary.

Note that although I use a `Range` in these examples, the compiler behaves similarly for other collections. For example, if I replace the `Range` in the previous example with a `List`, like this:

```
// original List code
val nums = List(1,2,3)
for (i <- nums) println(i)
```

the `for` loop is still converted by the compiler into a `foreach` method call:

```
// translation performed by the compiler
nums.foreach(((i) => println(i)))
```

Given this introduction, the following series of examples demonstrates how various `for` loops are translated by the Scala 2.10 compiler. Here's the first example again, showing both the *input* code I wrote and the *output* code from the compiler:

```
// #1 - input (my code)
for (i <- 1 to 10) println(i)

// #1 - compiler output
1.to(10).foreach(((i) => println(i)))
```

Next, I'll use the same `for` loop but add a guard condition (an `if` statement) to it:

```
// #2 - input code
for {
  i <- 1 to 10
  if i % 2 == 0
} println(i)

// #2 - translated output
1.to(10).withFilter(((i) => i.$percent(2).$eqSeq(0))).foreach(((i) =>
  println(i)))
```

As shown, a simple, single guard is translated into a `withFilter` method call on the collection, followed by a `foreach` call.

The same `for` loop with two guards is translated into two `withFilter` calls:

```
// #3 - input code
for {
  i <- 1 to 10
  if i != 1
  if i % 2 == 0
} println(i)

// #3 - translated output
1.to(10).withFilter(((i) => i.$bangSeq(1)))
    .withFilter(((i)
    => i.$percent(2).$eqSeq(0))).foreach(((i) => println(i)))
```

Next, I'll add a `yield` statement to the initial `for` loop:

```
// #4 - input code
for { i <- 1 to 10 } yield i
```

```
// #4 - output
1.to(10).map((i) => i)
```

As shown, when a `yield` statement is used, the compiler translates the `for/yield` code into a `map` method call on the collection.

Here's the same `for/yield` combination with a guard added in:

```
// #5 - input code (for loop, guard, and yield)
for {
  i <- 1 to 10
  if i % 2 == 0
} yield i
```

```
// #5 - translated code
1.to(10).withFilter((i) => i.$percent(2).$eq$eq(0))).map((i) => i)
```

As in the previous examples, the guard is translated into a `withFilter` method call, and the `for/yield` code is translated into a `map` method call.

These examples demonstrate how the translations are made by the Scala compiler, and I encourage you to create your own examples to see how they're translated by the compiler into other code. The `-Xprint:parse` option shows a small amount of compiler output, while the `-Xprint:all` option produces hundreds of lines of output for some of these examples, showing all the steps in the compilation process.

For more details, see the Scala Language Specification for exact rules on the `for` loop translation process. The details are currently in Section 6.19, “For Comprehensions and For Loops,” of the Specification.

See Also

The [Scala Language Specification](#) in PDF format

3.2. Using for Loops with Multiple Counters

Problem

You want to create a loop with multiple counters, such as when iterating over a multi-dimensional array.

Solution

You can create a `for` loop with two counters like this:

```
scala> for (i <- 1 to 2; j <- 1 to 2) println(s"i = $i, j = $j")
i = 1, j = 1
i = 1, j = 2
i = 2, j = 1
i = 2, j = 2
```

When doing this, the preferred style for multiline for loops is to use curly brackets:

```
for {
  i <- 1 to 2
  j <- 1 to 2
} println(s"i = $i, j = $j")
```

Similarly, you can use three counters like this:

```
for {
  i <- 1 to 3
  j <- 1 to 5
  k <- 1 to 10
} println(s"i = $i, j = $j, k = $k")
```

This is useful when looping over a multidimensional array. Assuming you create a small two-dimensional array like this:

```
val array = Array.ofDim[Int](2,2)
array(0)(0) = 0
array(0)(1) = 1
array(1)(0) = 2
array(1)(1) = 3
```

you can print each element of the array like this:

```
scala> for {
  |   i <- 0 to 1
  |   j <- 0 to 1
  | } println(s"($i)($j) = ${array(i)(j)}")
(0)(0) = 0
(0)(1) = 1
(1)(0) = 2
(1)(1) = 3
```

Discussion

Ranges created with the <- symbol in for loops are referred to as *generators*, and you can easily use multiple generators in one loop.

As shown in the examples, the recommended style for writing longer for loops is to use curly braces:

```
for {
  i <- 1 to 2
  j <- 2 to 3
} println(s"i = $i, j = $j")
```

This style is more scalable than other styles; in this case, “scalable” means that it continues to be readable as you add more generators and guards to the expression.

See Also

The [Scala Style Guide page on formatting control structures](#)

3.3. Using a for Loop with Embedded if Statements (Guards)

Problem

You want to add one or more conditional clauses to a for loop, typically to filter out some elements in a collection while working on the others.

Solution

Add an if statement after your generator, like this:

```
// print all even numbers
scala> for (i <- 1 to 10 if i % 2 == 0) println(i)
2
4
6
8
10
```

or using the preferred curly brackets style, like this:

```
for {
  i <- 1 to 10
  if i % 2 == 0
} println(i)
```

These if statements are referred to as filters, filter expressions, or *guards*, and you can use as many guards as are needed for the problem at hand. This loop shows a hard way to print the number 4:

```
for {
  i <- 1 to 10
  if i > 3
  if i < 6
  if i % 2 == 0
} println(i)
```

Discussion

Using guards with for loops can make for concise and readable code, but you can also use the traditional approach:

```
for (file <- files) {  
  if (hasSoundFileExtension(file) && !soundFileIsLong(file)) {  
    soundFiles += file  
  }  
}
```

However, once you become comfortable with Scala's for loop syntax, I think you'll find it makes the code more readable, because it separates the looping and filtering concerns from the business logic:

```
for {  
  file <- files  
  if passesFilter1(file)  
  if passesFilter2(file)  
} doSomething(file)
```

As a final note, because guards are generally intended to filter collections, you may want to use one of the many filtering methods that are available to collections (`filter`, `take`, `drop`, etc.) instead of a for loop, depending on your needs.

3.4. Creating a for Comprehension (for/yield Combination)

Problem

You want to create a new collection from an existing collection by applying an algorithm (and potentially one or more guards) to each element in the original collection.

Solution

Use a `yield` statement with a for loop and your algorithm to create a new collection from an existing collection.

For instance, given an array of lowercase strings:

```
scala> val names = Array("chris", "ed", "maurice")  
names: Array[String] = Array(chris, ed, maurice)
```

you can create a new array of capitalized strings by combining `yield` with a for loop and a simple algorithm:

```
scala> val capNames = for (e <- names) yield e.capitalize  
capNames: Array[String] = Array(Chris, Ed, Maurice)
```

Using a for loop with a `yield` statement is known as a *for comprehension*.

If your algorithm requires multiple lines of code, perform the work in a block after the `yield` keyword:

```
scala> val lengths = for (e <- names) yield {  
  |   // imagine that this required multiple lines of code  
  |   e.length  
  | }  
lengths: Array[Int] = Array(5, 2, 7)
```

Except for rare occasions, the collection type returned by a `for` comprehension is the same type that you begin with. For instance, if the collection you're looping over is an `ArrayBuffer`:

```
var fruits = scala.collection.mutable.ArrayBuffer[String]()  
fruits += "apple"  
fruits += "banana"  
fruits += "orange"
```

the collection your loop returns will also be an `ArrayBuffer`:

```
scala> val out = for (e <- fruits) yield e.toUpperCase  
out: scala.collection.mutable.ArrayBuffer[java.lang.String] =  
  ArrayBuffer(APPLE, BANANA, ORANGE)
```

If your input collection is a `List`, the `for/yield` loop will return a `List`:

```
scala> val fruits = "apple" :: "banana" :: "orange" :: Nil  
fruits: List[java.lang.String] = List(apple, banana, orange)  
  
scala> val out = for (e <- fruits) yield e.toUpperCase  
out: List[java.lang.String] = List(APPLE, BANANA, ORANGE)
```

Discussion

If you're new to using `yield` with a `for` loop, it can help to think of the loop like this:

- When it begins running, the `for/yield` loop immediately creates a new, empty collection that is of the same type as the input collection. For example, if the input type is a `Vector`, the output type will also be a `Vector`. You can think of this new collection as being like a bucket.
- On each iteration of the `for` loop, a new output element is created from the current element of the input collection. When the output element is created, it's placed in the bucket.
- When the loop finishes running, the entire contents of the bucket are returned.

That's a simplification of the process, but I find it helpful when explaining the process.

Writing a basic `for/yield` expression without a guard is just like calling the `map` method on a collection. For instance, the following `for` comprehension converts all the strings in the `fruits` collection to uppercase:

```
scala> val out = for (e <- fruits) yield e.toUpperCase
out: List[String] = List(APPLE, BANANA, ORANGE)
```

Calling the `map` method on the collection does the same thing:

```
scala> val out = fruits.map(_.toUpperCase)
out: List[String] = List(APPLE, BANANA, ORANGE)
```

When I first started learning Scala, I wrote all of my code using `for/yield` expressions until the `map` light bulb went on one day.

See Also

- Comparisons between `for` comprehensions and `map` are shown in more detail in [Recipe 10.13, “Transforming One Collection to Another with `for/yield`”](#) and [Recipe 10.14, “Transforming One Collection to Another with `map`”](#).
- The official [Scala website](#) offers an introduction to sequence comprehensions

3.5. Implementing break and continue

Problem

You have a situation where you need to use a `break` or `continue` construct, but Scala doesn't have `break` or `continue` keywords.

Solution

It's true that Scala doesn't have `break` and `continue` keywords, but it does offer similar functionality through `scala.util.control.Breaks`.

The following code demonstrates the Scala “`break`” and “`continue`” approach:

```
package com.alvinalexander.breakandcontinue

import util.control.Breaks._

object BreakAndContinueDemo extends App {

  println("\n=== BREAK EXAMPLE ===")
  breakable {
    for (i <- 1 to 10) {
      println(i)
      if (i > 4) break // break out of the for loop
    }
  }

  println("\n=== CONTINUE EXAMPLE ===")
  val searchMe = "peter piper picked a peck of pickled peppers"
```

```

var numPs = 0
for (i <- 0 until searchMe.length) {
  breakable {
    if (searchMe.charAt(i) != 'p') {
      break // break out of the 'breakable', continue the outside loop
    } else {
      numPs += 1
    }
  }
}
println("Found " + numPs + " p's in the string.")
}

```

Here's the output from the code:

```

=== BREAK EXAMPLE ===
1
2
3
4
5

=== CONTINUE EXAMPLE ===
Found 9 p's in the string.

```

(The “pickled peppers” example comes from a `continue` example in the Java documentation. More on this at the end of the recipe.)

The following discussions describe how this code works.

The break example

The break example is pretty easy to reason about. Again, here's the code:

```

breakable {
  for (i <- 1 to 10) {
    println(i)
    if (i > 4) break // break out of the for loop
  }
}

```

In this case, when `i` becomes greater than 4, the `break` “keyword” is reached. At this point an exception is thrown, and the `for` loop is exited. The `breakable` “keyword” essentially catches the exception, and the flow of control continues with any other code that might be after the `breakable` block.

Note that `break` and `breakable` aren't actually keywords; they're methods in `scala.util.control.Breaks`. In Scala 2.10, the `break` method is declared as follows to throw an instance of a `BreakControl` exception when it's called:

```

private val breakException = new BreakControl

def break(): Nothing = { throw breakException }

```


The breakable method is defined to catch a BreakControl exception, like this:

```
def breakable(op: => Unit) {  
  try {  
    op  
  } catch {  
    case ex: BreakControl =>  
      if (ex ne breakException) throw ex  
  }  
}
```



See [Recipe 3.18](#) for examples of how to implement your own control structures in a manner similar to the Breaks library.

The continue example

Given the explanation for the break example, you can now reason about how the “continue” example works. Here’s the code again:

```
val searchMe = "peter piper picked a peck of pickled peppers"  
var numPs = 0  
  
for (i <- 0 until searchMe.length) {  
  breakable {  
    if (searchMe.charAt(i) != 'p') {  
      break // break out of the 'breakable', continue the outside loop  
    } else {  
      numPs += 1  
    }  
  }  
}  
  
println("Found " + numPs + " p's in the string.")
```

Following the earlier explanation, as the code walks through the characters in the String variable named searchMe, if the current character is not the letter p, the code breaks out of the if/then statement, and the loop continues executing.

As before, what really happens is that the break method is reached, an exception is thrown, and that exception is caught by breakable. The exception serves to break out of the if/then statement, and catching it allows the for loop to continue executing with the next element.

General syntax

The general syntax for implementing break and continue functionality is shown in the following examples, which are partially written in pseudocode, and compared to their Java equivalents.

To implement a *break*, this Scala:

```
breakable {  
  for (x <- xs) {  
    if (cond)  
      break  
  }  
}
```

corresponds to this Java:

```
for (X x : xs) {  
  if (cond) break;  
}
```

To implement *continue* functionality, this Scala:

```
for (x <- xs) {  
  breakable {  
    if (cond)  
      break  
  }  
}
```

corresponds to this Java:

```
for (X x : xs) {  
  if (cond) continue;  
}
```

About that continue example...

The continue example shown is a variation of the Java continue example shown on the [Oracle website](#). If you know Scala, you know that there are better ways to solve this particular problem. For instance, a direct approach is to use the count method with a simple anonymous function:

```
val count = searchMe.count(_ == 'p')
```

When this code is run, count is again 9.

Nested loops and labeled breaks

In some situations, you may need nested break statements. Or, you may prefer labeled break statements. In either case, you can create labeled breaks as shown in the following example:

```
package com.alvinalexander.labeledbreaks  
  
object LabeledBreakDemo extends App {  
  
  import scala.util.control._  
  
  val Inner = new Breaks  
  val Outer = new Breaks
```

```

Outer.breakable {
  for (i <- 1 to 5) {
    Inner.breakable {
      for (j <- 'a' to 'e') {
        if (i == 1 && j == 'c') Inner.break else println(s"i: $i, j: $j")
        if (i == 2 && j == 'b') Outer.break
      }
    }
  }
}

```

In this example, if the first `if` condition is met, an exception is thrown and caught by `Inner.breakable`, and the outer `for` loop continues. But if the second `if` condition is triggered, control of flow is sent to `Outer.breakable`, and both loops are exited. Running this object results in the following output:

```

i: 1, j: a
i: 1, j: b
i: 2, j: a

```

Use the same approach if you prefer labeled breaks. This example shows how you can use the same technique with just one `break` method call:

```

import scala.util.control._

val Exit = new Breaks
Exit.breakable {
  for (j <- 'a' to 'e') {
    if (j == 'c') Exit.break else println(s"j: $j")
  }
}

```

Discussion

If you don't like using `break` and `continue`, there are several other ways to attack these problems.

For instance, if you want to add monkeys to a barrel, but only until the barrel is full, you can use a simple boolean test to break out of a `for` loop:

```

var barrelIsFull = false
for (monkey <- monkeyCollection if !barrelIsFull) {
  addMonkeyToBarrel(monkey)
  barrelIsFull = checkIfBarrelIsFull
}

```

Another approach is to place your algorithm inside a function, and then return from the function when the desired condition is reached. In the following example, the `sumToMax` function returns early if `sum` becomes greater than `limit`:

```
// calculate a sum of numbers, but limit it to a 'max' value
def sumToMax(arr: Array[Int], limit: Int): Int = {
  var sum = 0
  for (i <- arr) {
    sum += i
    if (sum > limit) return limit
  }
  sum
}
val a = Array.range(0,10)
println(sumToMax(a, 10))
```

A common approach in functional programming is to use recursive algorithms. This is demonstrated in a recursive approach to a *factorial* function, where the condition `n == 1` results in a break from the recursion:

```
def factorial(n: Int): Int = {
  if (n == 1) 1
  else n * factorial(n - 1)
}
```

Note that this example does not use *tail recursion* and is therefore not an optimal approach, especially if the starting value `n` is very large. A more optimal solution takes advantage of tail recursion:

```
import scala.annotation.tailrec

def factorial(n: Int): Int = {
  @tailrec def factorialAcc(acc: Int, n: Int): Int = {
    if (n <= 1) acc
    else factorialAcc(n * acc, n - 1)
  }
  factorialAcc(1, n)
}
```

Note that you can use the `@tailrec` annotation in situations like this to confirm that your algorithm is tail recursive. If you use this annotation and your algorithm isn't tail recursive, the compiler will complain. For instance, if you attempt to use this annotation on the first version of the `factorial` method, you'll get the following compile-time error:

```
Could not optimize @tailrec annotated method factorial: it contains a recursive
call not in tail position
```

See Also

The Java `continue` example mentioned can be found on the [Oracle website](#).

There are many Scala recursive factorial examples on the Internet; here are two of the best discussions:

- A nice discussion about [tail recursion and trampolines](#)
- [Tail-call optimization in Scala](#)

3.6. Using the if Construct Like a Ternary Operator

Problem

You'd like to use a Scala `if` expression like a ternary operator to solve a problem in a concise, expressive way.

Solution

This is a bit of a trick problem, because unlike Java, in Scala there is no special ternary operator; just use an `if/else` expression:

```
val absValue = if (a < 0) -a else a
```

Because an `if` expression returns a value, you can embed it into a print statement:

```
println(if (i == 0) "a" else "b")
```

You can use it in another expression, such as this portion of a `hashCode` method:

```
hash = hash * prime + (if (name == null) 0 else name.hashCode)
```

Discussion

The Java documentation page shown in the See Also states that the Java conditional operator `?:` “is known as the *ternary operator* because it uses three operands.” Unlike some other languages, Scala doesn't have a special operator for this use case.

In addition to the examples shown, the combination of (a) `if` statements returning a result, and (b) Scala's syntax for defining methods makes for concise code:

```
def abs(x: Int) = if (x >= 0) x else -x
def max(a: Int, b: Int) = if (a > b) a else b
val c = if (a > b) a else b
```

See Also

“Equality, Relational, and Conditional Operators” on the [Java Tutorials page](#)

3.7. Using a Match Expression Like a switch Statement

Problem

You have a situation where you want to create something like a simple Java integer-based switch statement, such as matching the days in a week, months in a year, and other situations where an integer maps to a result.

Solution

To use a Scala match expression like a Java switch statement, use this approach:

```
// i is an integer
i match {
  case 1 => println("January")
  case 2 => println("February")
  case 3 => println("March")
  case 4 => println("April")
  case 5 => println("May")
  case 6 => println("June")
  case 7 => println("July")
  case 8 => println("August")
  case 9 => println("September")
  case 10 => println("October")
  case 11 => println("November")
  case 12 => println("December")
  // catch the default with a variable so you can print it
  case whoa => println("Unexpected case: " + whoa.toString)
}
```

That example shows how to take an action based on a match. A more functional approach returns a value from a match expression:

```
val month = i match {
  case 1 => "January"
  case 2 => "February"
  case 3 => "March"
  case 4 => "April"
  case 5 => "May"
  case 6 => "June"
  case 7 => "July"
  case 8 => "August"
  case 9 => "September"
  case 10 => "October"
  case 11 => "November"
  case 12 => "December"
  case _ => "Invalid month" // the default, catch-all
}
```

The @switch annotation

When writing simple match expressions like this, it's recommend to use the @switch annotation. This annotation provides a warning at compile time if the switch can't be compiled to a `tableswitch` or `lookupswitch`.

Compiling your match expression to a `tableswitch` or `lookupswitch` is better for performance, because it results in a branch table rather than a decision tree. When a value is given to the expression, it can jump directly to the result rather than working through the decision tree.

Here's the official description from the [@switch annotation documentation](#):

"An annotation to be applied to a match expression. If present, the compiler will verify that the match has been compiled to a `tableswitch` or `lookupswitch`, and issue an error if it instead compiles into a series of conditional expressions."

The effect of the @switch annotation is demonstrated with a simple example. First, place the following code in a file named *SwitchDemo.scala*:

```
// Version 1 - compiles to a tableswitch
import scala.annotation.switch

class SwitchDemo {

    val i = 1
    val x = (i: @switch) match {
        case 1 => "One"
        case 2 => "Two"
        case _ => "Other"
    }

}
```

Then compile the code as usual:

```
$ scalac SwitchDemo.scala
```

Compiling this class produces no warnings and creates the *SwitchDemo.class* output file. Next, disassemble that file with this `javap` command:

```
$ javap -c SwitchDemo
```

The output from this command shows a `tableswitch`, like this:

```
16: tableswitch{ //1 to 2
      1: 50;
      2: 45;
      default: 40 }
```

This shows that Scala was able to optimize your match expression to a `tableswitch`. (This is a good thing.)

Next, make a minor change to the code, replacing the integer literal 2 with a value:

```
import scala.annotation.switch

// Version 2 - leads to a compiler warning
class SwitchDemo {

  val i = 1
  val Two = 2 // added
  val x = (i: @switch) match {
    case 1 => "One"
    case Two => "Two" // replaced the '2'
    case _ => "Other"
  }
}
```

Again, compile the code with `scalac`, but right away you'll see a warning message:

```
$ scalac SwitchDemo.scala
SwitchDemo.scala:7: warning: could not emit switch for @switch annotated match
  val x = (i: @switch) match {
            ^
one warning found
```

This warning message is saying that neither a `tableswitch` nor `lookupswitch` could be generated for the match expression. You can confirm this by running the `javap` command on the `SwitchDemo.class` file that was generated. When you look at that output, you'll see that the `tableswitch` shown in the previous example is now gone.

In his book, *Scala In Depth* (Manning), Joshua Suereth states that the following conditions must be true for Scala to apply the `tableswitch` optimization:

1. The matched value must be a known integer.
2. The matched expression must be “simple.” It can't contain any type checks, if statements, or extractors.
3. The expression must also have its value available at compile time.
4. There should be more than two case statements.

For more information on how JVM switches work, see the Oracle document, [Compiling Switches](#).

Discussion

As demonstrated in other recipes, you aren't limited to matching only integers; the match expression is incredibly flexible:

```
def getClassAsString(x: Any): String = x match {
  case s: String => s + " is a String"
  case i: Int => "Int"
  case f: Float => "Float"
```



```

case l: List[_] => "List"
case p: Person => "Person"
case _ => "Unknown"
}

```

Handling the default case

The examples in the Solution showed the two ways you can handle the default, “catch all” case. First, if you’re not concerned about the value of the default match, you can catch it with the `_` wildcard:

```
case _ => println("Got a default match")
```

Conversely, if you are interested in what fell down to the default match, assign a variable name to it. You can then use that variable on the right side of the expression:

```
case default => println(default)
```

Using the name `default` often makes the most sense and leads to readable code, but you can use any legal name for the variable:

```
case oops => println(oops)
```

You can generate a `MatchError` if you don’t handle the default case. Given this match expression:

```

i match {
  case 0 => println("0 received")
  case 1 => println("1 is good, too")
}

```

if `i` is a value other than `0` or `1`, the expression throws a `MatchError`:

```

scala.MatchError: 42 (of class java.lang.Integer)
  at .<init>(<console>:9)
  at .<clinit>(<console>)
  much more error output here ...

```

So unless you’re intentionally writing a *partial function*, you’ll want to handle the default case. (See [Recipe 9.8, “Creating Partial Functions”](#), for more information on partial functions.)

Do you really need a switch statement?

Of course you don’t really need a switch statement if you have a data structure that maps month numbers to month names. In that case, just use a `Map`:

```

val monthNumberToName = Map(
  1 -> "January",
  2 -> "February",
  3 -> "March",
  4 -> "April",
  5 -> "May",
  6 -> "June",
  7 -> "July",

```

```

    8 -> "August",
    9 -> "September",
   10 -> "October",
   11 -> "November",
   12 -> "December"
)

val monthName = monthNumberToName(4)
println(monthName) // prints "April"

```

See Also

- The [@switch annotation documentation](#).
- The Oracle document, [Compiling Switches](#), discusses the `tableswitch` and `lookupswitch`.
- A [tableswitch and lookupswitch differences](#) discussion.

3.8. Matching Multiple Conditions with One Case Statement

Problem

You have a situation where several match conditions require that the same business logic be executed, and rather than repeating your business logic for each case, you'd like to use one copy of the business logic for the matching conditions.

Solution

Place the match conditions that invoke the same business logic on one line, separated by the `|` (pipe) character:

```

val i = 5
i match {
  case 1 | 3 | 5 | 7 | 9 => println("odd")
  case 2 | 4 | 6 | 8 | 10 => println("even")
}

```

This same syntax works with strings and other types. Here's an example based on a `String` match:

```

val cmd = "stop"
cmd match {
  case "start" | "go" => println("starting")
  case "stop" | "quit" | "exit" => println("stopping")
  case _ => println("doing nothing")
}

```

This example shows how to match multiple case objects:

```
trait Command
case object Start extends Command
case object Go extends Command
case object Stop extends Command
case object Whoa extends Command

def executeCommand(cmd: Command) = cmd match {
  case Start | Go => start()
  case Stop | Whoa => stop()
}
```

As demonstrated, the ability to define multiple possible matches for each case statement can simplify your code.

See Also

See [Recipe 3.13, “Adding if Expressions \(Guards\) to Case Statements”](#), for a related approach.

3.9. Assigning the Result of a Match Expression to a Variable

Problem

You want to return a value from a match expression and assign it to a variable, or use a match expression as the body of a method.

Solution

To assign a variable to the result of a match expression, insert the variable assignment before the expression, as with the variable `evenOrOdd` in this example:

```
val evenOrOdd = someNumber match {
  case 1 | 3 | 5 | 7 | 9 => println("odd")
  case 2 | 4 | 6 | 8 | 10 => println("even")
}
```

This approach is commonly used to create short methods or functions. For example, the following method implements the Perl definitions of `true` and `false`:

```
def isTrue(a: Any) = a match {
  case 0 | "" => false
  case _ => true
}
```

You'll hear that Scala is an “expression-oriented programming (EOP) language,” which Wikipedia defines as, “a programming language where every (or nearly every) construction is an expression and thus yields a value.” The ability to return values from `if` statements and `match` expressions helps Scala meet this definition.

See Also

- [Recipe 20.3, “Think “Expression-Oriented Programming””](#)
- The Expression-Oriented Programming page on [Wikipedia](#)

3.10. Accessing the Value of the Default Case in a Match Expression

Problem

You want to access the value of the default, “catch all” case when using a match expression, but you can't access the value when you match it with the `_` wildcard syntax.

Solution

Instead of using the `_` wildcard character, assign a variable name to the default case:

```
i match {  
  case 0 => println("1")  
  case 1 => println("2")  
  case default => println("You gave me: " + default)  
}
```

By giving the default match a variable name, you can access the variable on the right side of the statement.

Discussion

The key to this recipe is in using a variable name for the default match instead of the usual `_` wildcard character.

The name you assign can be any legal variable name, so instead of naming it `default`, you can name it something else, such as `whoa`:

```
i match {  
  case 0 => println("1")  
  case 1 => println("2")  
  case whoa => println("You gave me: " + whoa)  
}
```

It's important to provide a default match. Failure to do so can cause a `MatchError`:

```
scala> 3 match {
  |   case 1 => println("one")
  |   case 2 => println("two")
  |   // no default match
  | }
scala.MatchError: 3 (of class java.lang.Integer)
many more lines of output ...
```

3.11. Using Pattern Matching in Match Expressions

Problem

You need to match one or more patterns in a match expression, and the pattern may be a constant pattern, variable pattern, constructor pattern, sequence pattern, tuple pattern, or type pattern.

Solution

Define a case statement for each pattern you want to match. The following method shows examples of many different types of patterns you can use in match expressions:

```
def echoWhatYouGaveMe(x: Any): String = x match {

  // constant patterns
  case 0 => "zero"
  case true => "true"
  case "hello" => "you said 'hello'"
  case Nil => "an empty List"

  // sequence patterns
  case List(0, _, _) => "a three-element list with 0 as the first element"
  case List(1, _) => "a list beginning with 1, having any number of elements"
  case Vector(1, _) => "a vector starting with 1, having any number of elements"

  // tuples
  case (a, b) => s"got $a and $b"
  case (a, b, c) => s"got $a, $b, and $c"

  // constructor patterns
  case Person(first, "Alexander") => s"found an Alexander, first name = $first"
  case Dog("Suka") => "found a dog named Suka"

  // typed patterns
  case s: String => s"you gave me this string: $s"
  case i: Int => s"thanks for the int: $i"
  case f: Float => s"thanks for the float: $f"
  case a: Array[Int] => s"an array of int: ${a.mkString(",")}"
  case as: Array[String] => s"an array of strings: ${as.mkString(",")}"
```

```

case d: Dog => s"dog: ${d.name}"
case list: List[_] => s"thanks for the List: $list"
case m: Map[_ , _] => m.toString

// the default wildcard pattern
case _ => "Unknown"
}

```

The large match expression in this method shows the different categories of patterns described in the book, *Programming in Scala* (Artima), by Odersky, et al, including constant patterns, sequence patterns, tuple patterns, constructor patterns, and typed patterns.

You can test this match expression in a variety of ways. For the purposes of this example, I created the following object to test the `echoWhatYouGaveMe` method:

```

object LargeMatchTest extends App {

  case class Person(firstName: String, lastName: String)
  case class Dog(name: String)

  // trigger the constant patterns
  println(echoWhatYouGaveMe(0))
  println(echoWhatYouGaveMe(true))
  println(echoWhatYouGaveMe("hello"))
  println(echoWhatYouGaveMe(nil))

  // trigger the sequence patterns
  println(echoWhatYouGaveMe(List(0,1,2)))
  println(echoWhatYouGaveMe(List(1,2)))
  println(echoWhatYouGaveMe(List(1,2,3)))
  println(echoWhatYouGaveMe(Vector(1,2,3)))

  // trigger the tuple patterns
  println(echoWhatYouGaveMe((1,2)))           // two element tuple
  println(echoWhatYouGaveMe((1,2,3)))         // three element tuple

  // trigger the constructor patterns
  println(echoWhatYouGaveMe(Person("Melissa", "Alexander")))
  println(echoWhatYouGaveMe(Dog("Suka")))

  // trigger the typed patterns
  println(echoWhatYouGaveMe("Hello, world"))
  println(echoWhatYouGaveMe(42))
  println(echoWhatYouGaveMe(42F))
  println(echoWhatYouGaveMe(Array(1,2,3)))
  println(echoWhatYouGaveMe(Array("coffee", "apple pie")))
  println(echoWhatYouGaveMe(Dog("Fido")))
  println(echoWhatYouGaveMe(List("apple", "banana")))
  println(echoWhatYouGaveMe(Map(1->"Al", 2->"Alexander")))
}

```

```
// trigger the wildcard pattern
println(echoWhatYouGaveMe("33d"))

}
```

Running this object results in the following output:

```
zero
true
you said 'hello'
an empty List

a three-element list with 0 as the first element
a list beginning with 1 and having any number of elements
a list beginning with 1 and having any number of elements
a vector beginning with 1 and having any number of elements
a list beginning with 1 and having any number of elements

got 1 and 2
got 1, 2, and 3

found an Alexander, first name = Melissa
found a dog named Suka

you gave me this string: Hello, world
thanks for the int: 42
thanks for the float: 42.0
an array of int: 1,2,3
an array of strings: coffee,apple pie
dog: Fido
thanks for the List: List(apple, banana)
Map(1 -> Al, 2 -> Alexander)

you gave me this string: 33d
```

Note that in the match expression, the List and Map statements that were written like this:

```
case list: List[_] => s"thanks for the List: $list"
case m: Map[_,_] => m.toString
```

could have been written as this instead:

```
case m: Map[a,b] => m.toString
case list: List[x] => s"thanks for the List: $list"
```

I prefer the underscore syntax because it makes it clear that I'm not concerned about what's stored in the List or Map. Actually, there are times that I might be interested in what's stored in the List or Map, but because of type erasure in the JVM, that becomes a difficult problem.



When I first wrote this example, I wrote the `List` expression as follows:

```
case l: List[Int] => "List"
```

If you're familiar with *type erasure* on the Java platform, you may know that this won't work. The Scala compiler kindly lets you know about this problem with this warning message:

```
Test1.scala:7: warning: non-variable type argument Int in type pattern
List[Int] is unchecked since it is eliminated by erasure
case l: List[Int] => "List[Int]"
      ^
```

If you're not familiar with type erasure, I've included a link in the See Also section of this recipe that describes how it works on the JVM.

Discussion

Typically when using this technique, your method will expect an instance that inherits from a base class or trait, and then your case statements will reference subtypes of that base type. This was inferred in the `echoWhatYouGaveMe` method, where every Scala type is a subtype of `Any`. The following code shows a more obvious example of this technique.

In my **Blue Parrot application**, which either plays a sound file or “speaks” the text it's given at random intervals, I have a method that looks like this:

```
import java.io.File

sealed trait RandomThing

case class RandomFile(f: File) extends RandomThing
case class RandomString(s: String) extends RandomThing

class RandomNoiseMaker {

  def makeRandomNoise(t: RandomThing) = t match {
    case RandomFile(f) => playSoundFile(f)
    case RandomString(s) => speak(s)
  }

}
```

The `makeRandomNoise` method is declared to take a `RandomThing` type, and then the match expression handles its two subtypes, `RandomFile` and `RandomString`.

Patterns

The large match expression in the Solution shows a variety of patterns that are defined in the book *Programming in Scala*. These patterns are briefly described in the following paragraphs.

Constant patterns

A constant pattern can only match itself. Any literal may be used as a constant. If you specify a 0 as the literal, only an Int value of 0 will be matched. Examples:

```
case 0 => "zero"
case true => "true"
```

Variable patterns

This was not shown in the large match example in the Solution—it’s discussed in detail in [Recipe 3.10, “Accessing the Value of the Default Case in a Match Expression”](#)—but a *variable pattern* matches any object just like the `_` wildcard character. Scala binds the variable to whatever the object is, which lets you use the variable on the right side of the case statement. For example, at the end of a match expression you can use the `_` wildcard character like this to catch “anything else”:

```
case _ => s"Hmm, you gave me something ..."
```

But with a variable pattern you can write this instead:

```
case foo => s"Hmm, you gave me a $foo"
```

See [Recipe 3.10](#) for more information.

Constructor patterns

The *constructor pattern* lets you match a constructor in a case statement. As shown in the examples, you can specify constants or variable patterns as needed in the constructor pattern:

```
case Person(first, "Alexander") => s"found an Alexander, first name = $first"
case Dog("Suka") => "found a dog named Suka"
```

Sequence patterns

You can match against sequences like List, Array, Vector, etc. Use the `_` character to stand for one element in the sequence, and use `_*` to stand for “zero or more elements,” as shown in the examples:

```
case List(0, _, _) => "a three-element list with 0 as the first element"
case List(1, _*) => "a list beginning with 1, having any number of elements"
case Vector(1, _*) => "a vector beginning with 1 and having any number ..."
```

Tuple patterns

As shown in the examples, you can match *tuple patterns* and access the value of each element in the tuple. You can also use the `_` wildcard if you’re not interested in the value of an element:

```
case (a, b, c) => s"3-elem tuple, with values $a, $b, and $c"
case (a, b, c, _) => s"4-elem tuple: got $a, $b, and $c"
```

Type patterns

In the following example, `str: String` is a *typed pattern*, and `str` is a *pattern variable*:

```
case str: String => s"you gave me this string: $str"
```

As shown in the examples, you can access the pattern variable on the right side of the expression after declaring it.

Adding variables to patterns

At times you may want to add a variable to a pattern. You can do this with the following general syntax:

```
variableName @ pattern
```

As the book, *Programming in Scala*, states, “This gives you a variable-binding pattern. The meaning of such a pattern is to perform the pattern match as normal, and if the pattern succeeds, set the variable to the matched object just as with a simple variable pattern.”

The usefulness of this is best shown by demonstrating the problem it solves. Suppose you had the `List` pattern that was shown earlier:

```
case List(1, _) => "a list beginning with 1, having any number of elements"
```

As demonstrated, this lets you match a `List` whose first element is 1, but so far, the `List` hasn't been accessed on the right side of the expression. When accessing a `List`, you know that you can do this:

```
case list: List[_] => s"thanks for the List: $list"
```

so it seems like you should try this with a sequence pattern:

```
case list: List(1, _) => s"thanks for the List: $list"
```

Unfortunately, this fails with the following compiler error:

```
Test2.scala:22: error: '=>' expected but '(' found.
    case list: List(1, _) => s"thanks for the List: $list"
                        ^
one error found
```

The solution to this problem is to add a variable-binding pattern to the sequence pattern:

```
case list @ List(1, _) => s"$list"
```

This code compiles, and works as expected, giving you access to the `List` on the right side of the statement.

The following code demonstrates this example and the usefulness of this approach:

```
case class Person(firstName: String, lastName: String)

object Test2 extends App {

  def matchType(x: Any): String = x match {

    //case x: List(1, _) => s"$x" // doesn't compile
  }
}
```

```

    case x @ List(1, _) => s"$x"           // works; prints the list

    //case Some(_) => "got a Some" // works, but can't access the Some
    //case Some(x) => s"$x" // works, returns "foo"
    case x @ Some(_) => s"$x"           // works, returns "Some(foo)"

    case p @ Person(first, "Doe") => s"$p" // works, returns "Person(John,Doe)"

}

println(matchType(List(1,2,3)))           // prints "List(1, 2, 3)"
println(matchType(Some("foo")))           // prints "Some(foo)"
println(matchType(Person("John", "Doe"))) // prints "Person(John,Doe)"

}

```

In the two `List` examples inside the `match` expression, the commented-out line of code won't compile, but the second example shows how to assign the variable `x` to the `List` object it matches. When this line of code is matched with the `println(matchType(List(1,2,3)))` call, it results in the output `List(1, 2, 3)`.

The first `Some` example shows that you can match a `Some` with the approach shown, but you can't access its information on the righthand side of the expression. The second example shows how you can access the value inside the `Some`, and the third example takes this a step further, giving you access to the `Some` object itself. When it's matched by the second `println` call, it prints `Some(foo)`, demonstrating that you now have access to the `Some` object.

Finally, this approach is used to match a `Person` whose last name is `Doe`. This syntax lets you assign the result of the pattern match to the variable `p`, and then access that variable on the right side of the expression.

Using `Some` and `None` in match expressions

To round out these examples, you'll often use `Some` and `None` with match expressions. For instance, assume you have a `toInt` method defined like this:

```

def toInt(s: String): Option[Int] = {
  try {
    Some(Integer.parseInt(s.trim))
  } catch {
    case e: Exception => None
  }
}

```

In some situations, you may want to use this method with a match expression, like this:

```

toInt("42") match {
  case Some(i) => println(i)
  case None => println("That wasn't an Int.")
}

```

Inside the match expression you just specify the `Some` and `None` cases as shown to handle the success and failure conditions. See [Recipe 20.6](#) for more examples of using `Option`, `Some`, and `None`.

See Also

- A discussion of getting around type erasure when using match expressions on [Stack Overflow](#)
- [My Blue Parrot application](#)
- [The “Type Erasure” documentation](#)

3.12. Using Case Classes in Match Expressions

Problem

You want to match different case classes (or case objects) in a match expression, such as when receiving messages in an actor.

Solution

Use the different patterns shown in the previous recipe to match case classes and objects, depending on your needs.

The following example demonstrates how to use patterns to match case classes and case objects in different ways, depending primarily on what information you need on the right side of each case statement. In this example, the `Dog` and `Cat` case classes and the `Woodpecker` case object are different subtypes of the `Animal` trait:

```
trait Animal
case class Dog(name: String) extends Animal
case class Cat(name: String) extends Animal
case object Woodpecker extends Animal

object CaseClassTest extends App {

  def determineType(x: Animal): String = x match {
    case Dog(moniker) => "Got a Dog, name = " + moniker
    case _:Cat => "Got a Cat (ignoring the name)"
    case Woodpecker => "That was a Woodpecker"
    case _ => "That was something else"
  }
```

```
println(determineType(new Dog("Rocky")))
println(determineType(new Cat("Rusty the Cat")))
println(determineType(Woodpecker))

}
```

When the code is compiled and run, the output is:

```
Got a Dog, name = Rocky
Got a Cat (ignoring the name)
That was a Woodpecker
```

In this example, if the `Dog` class is matched, its name is extracted and used in the print statement on the right side of the expression. To show that the variable name used when extracting the name can be any legal variable name, I use the name `moniker`.

When matching a `Cat`, I want to ignore the name, so I use the syntax shown to match any `Cat` instance. Because `Woodpecker` is defined as a case object and has no name, it is also matched as shown.

3.13. Adding if Expressions (Guards) to Case Statements

Problem

You want to add qualifying logic to a case statement in a match expression, such as allowing a range of numbers, or matching a pattern, but only if that pattern matches some additional criteria.

Solution

Add an *if guard* to your case statement. Use it to match a range of numbers:

```
i match {
  case a if 0 to 9 contains a => println("0-9 range: " + a)
  case b if 10 to 19 contains b => println("10-19 range: " + b)
  case c if 20 to 29 contains c => println("20-29 range: " + c)
  case _ => println("Hmmm...")
}
```

Use it to match different values of an object:

```
num match {
  case x if x == 1 => println("one, a lonely number")
  case x if (x == 2 || x == 3) => println(x)
  case _ => println("some other value")
}
```

You can reference class fields in your *if guards*. Imagine here that `x` is an instance of a `Stock` class that has `symbol` and `price` fields:

```
stock match {
  case x if (x.symbol == "XYZ" && x.price < 20) => buy(x)
  case x if (x.symbol == "XYZ" && x.price > 50) => sell(x)
  case _ => // do nothing
}
```

You can also extract fields from case classes and use those in your guards:

```
def speak(p: Person) = p match {
  case Person(name) if name == "Fred" => println("Yubba dubba doo")
  case Person(name) if name == "Bam Bam" => println("Bam bam!")
  case _ => println("Watch the Flintstones!")
}
```

Discussion

You can use this syntax whenever you want to add simple matches to your case statements on the left side of the expression.

Note that all of these examples could be written by putting the `if` tests on the right side of the expressions, like this:

```
case Person(name) =>
  if (name == "Fred") println("Yubba dubba doo")
  else if (name == "Bam Bam") println("Bam bam!")
```

However, for many situations, your code will be simpler and easier to read by joining the `if` guard directly with the case statement.

3.14. Using a Match Expression Instead of `isInstanceOf`

Problem

You want to write a block of code to match one type, or multiple different types.

Solution

You *can* use the `isInstanceOf` method to test the type of an object:

```
if (x.isInstanceOf[Foo]) { do something ... }
```

However, some programmers discourage this approach, and in other cases, it may not be convenient. In these instances, you can handle the different expected types in a match expression.

For example, you may be given an object of unknown type, and want to determine if the object is an instance of a `Person`:

```
def isPerson(x: Any): Boolean = x match {
  case p: Person => true
}
```

```

    case _ => false
  }

```

Or you may be given an object that extends a known supertype, and then want to take different actions based on the exact subtype. In the following example, the `printInfo` method is given a `SentientBeing`, and then handles the subtypes differently:

```

trait SentientBeing
trait Animal extends SentientBeing
case class Dog(name: String) extends Animal
case class Person(name: String, age: Int) extends SentientBeing

// later in the code ...
def printInfo(x: SentientBeing) = x match {
  case Person(name, age) => // handle the Person
  case Dog(name) => // handle the Dog
}

```

Discussion

As shown, a match expression lets you match multiple types, so using it to replace the `isInstanceOf` method is just a natural use of the case syntax and the general pattern-matching approach used in Scala applications.

In simple examples, the `isInstanceOf` method can be a simpler approach to determining whether an object matches a type:

```

if (o.isInstanceOf[Person]) { // handle this ...

```

However, with more complex needs, a match expression is more readable than an `if/else` statement.

3.15. Working with a List in a Match Expression

Problem

You know that a `List` data structure is a little different than other collection data structures. It's built from *cons* cells and ends in a `Nil` element. You want to use this to your advantage when working with a match expression, such as when writing a recursive function.

Solution

You can create a `List` like this:

```

val x = List(1, 2, 3)

```

or like this, using *cons* cells and a `Nil` element:

```

val y = 1 :: 2 :: 3 :: Nil

```

When writing a recursive algorithm, you can take advantage of the fact that the last element in a `List` is a `Nil` object. For instance, in the following `listToString` method, if the current element is not `Nil`, the method is called recursively with the remainder of the `List`, but if the current element is `Nil`, the recursive calls are stopped and an empty `String` is returned, at which point the recursive calls unwind:

```
def listToString(list: List[String]): String = list match {  
  case s :: rest => s + " " + listToString(rest)  
  case Nil => ""  
}
```

Running this example in the REPL yields the following result:

```
scala> val fruits = "Apples" :: "Bananas" :: "Oranges" :: Nil  
fruits: List[java.lang.String] = List(Apples, Bananas, Oranges)  
  
scala> listToString(fruits)  
res0: String = "Apples Bananas Oranges "
```

The same approach of (a) handling the `Nil` condition and (b) handling the remainder of the `List` can be used when dealing with a `List` of other types:

```
def sum(list: List[Int]): Int = list match {  
  case Nil => 1  
  case n :: rest => n + sum(rest)  
}  
  
def multiply(list: List[Int]): Int = list match {  
  case Nil => 1  
  case n :: rest => n * multiply(rest)  
}
```

These methods are demonstrated in the REPL:

```
scala> val nums = List(1,2,3,4,5)  
nums: List[Int] = List(1, 2, 3, 4, 5)  
  
scala> sum(nums)  
res0: Int = 16  
  
scala> multiply(nums)  
res1: Int = 120
```

Discussion

When using this recipe, be sure to handle the `Nil` case, or you'll get the following error in the REPL:

```
warning: match is not exhaustive!
```

In the real world (outside the REPL), you'll get a `MatchError`:


```
Exception in thread "main" scala.MatchError: List()  
(of class scala.collection.immutable.Nil$)
```

See Also

[Recipe 3.11, “Using Pattern Matching in Match Expressions”](#), for more examples of using a match expression with multiple types

3.16. Matching One or More Exceptions with try/catch

Problem

You want to catch one or more exceptions in a try/catch block.

Solution

The Scala try/catch/finally syntax is similar to Java, but it uses the match expression approach in the catch block:

```
val s = "Foo"  
try {  
  val i = s.toInt  
} catch {  
  case e: Exception => e.printStackTrace  
}
```

When you need to catch and handle multiple exceptions, just add the exception types as different case statements:

```
try {  
  openAndReadAFile(filename)  
} catch {  
  case e: FileNotFoundException => println("Couldn't find that file.")  
  case e: IOException => println("Had an IOException trying to read that file")  
}
```

Discussion

As shown, the Scala match expression syntax is used to match different possible exceptions. If you're not concerned about which specific exceptions might be thrown, and want to catch them all and do something with them (such as log them), use this syntax:

```
try {  
  openAndReadAFile("foo")  
} catch {  
  case t: Throwable => t.printStackTrace()  
}
```

You can also catch them all and ignore them like this:

```

try {
  val i = s.toInt
} catch {
  case _: Throwable => println("exception ignored")
}

```

As with Java, you can throw an exception from a `catch` clause, but because Scala doesn't have checked exceptions, you don't need to specify that a method throws the exception. This is demonstrated in the following example, where the method isn't annotated in any way:

```

// nothing required here
def toInt(s: String): Option[Int] =
  try {
    Some(s.toInt)
  } catch {
    case e: Exception => throw e
  }

```

If you prefer to declare the exceptions that your method throws, or you need to interact with Java, add the `@throws` annotation to your method definition:

```

@throws(classOf[NumberFormatException])
def toInt(s: String): Option[Int] =
  try {
    Some(s.toInt)
  } catch {
    case e: NumberFormatException => throw e
  }

```

See Also

- [Recipe 5.8, “Declaring That a Method Can Throw an Exception”](#) for more examples of declaring that a method can throw an exception
- [Recipe 2.1, “Parsing a Number from a String”](#) for more examples of a `toInt` method

3.17. Declaring a Variable Before Using It in a try/catch/finally Block

Problem

You want to use an object in a `try` block, and need to access it in the `finally` portion of the block, such as when you need to call a `close` method on an object.

Solution

In general, declare your field as an `Option` before the `try/catch` block, then create a `Some` inside the `try` clause. This is shown in the following example, where the fields `in` and `out` are declared before the `try/catch` block, and assigned inside the `try` clause:

```
import java.io._

object CopyBytes extends App {

  var in = None: Option[FileInputStream]
  var out = None: Option[FileOutputStream]

  try {
    in = Some(new FileInputStream("/tmp/Test.class"))
    out = Some(new FileOutputStream("/tmp/Test.class.copy"))
    var c = 0
    while ({c = in.get.read; c != -1}) {
      out.get.write(c)
    }
  } catch {
    case e: IOException => e.printStackTrace
  } finally {
    println("entered finally ...")
    if (in.isDefined) in.get.close
    if (out.isDefined) out.get.close
  }
}
```

In this code, `in` and `out` are assigned to `None` before the `try` clause, and then reassigned to `Some` values inside the `try` clause if everything succeeds. Therefore, it's safe to call `in.get` and `out.get` in the `while` loop, because if an exception had occurred, flow control would have switched to the `catch` clause, and then the `finally` clause before leaving the method.

Normally I tell people that I wish the `get` and `isDefined` methods on `Option` would be deprecated, but this is one of the few times where I think their use is acceptable, and they lead to more readable code.

Another approach you can employ inside the `try` clause is to use the `foreach` approach with a `Some`:

```
try {
  in = Some(new FileInputStream("/tmp/Test.class"))
  out = Some(new FileOutputStream("/tmp/Test.class.copy"))
  in.foreach { inputStream =>
    out.foreach { outputStream =>
      var c = 0
      while ({c = inputStream.read; c != -1}) {
        outputStream.write(c)
      }
    }
  }
}
```

```

    }
  }
} // ...

```

This is still readable with two variables, and eliminates the `get` method calls, but wouldn't be practical with more variables.

Discussion

One key to this recipe is knowing the syntax for declaring `Option` fields that aren't initially populated:

```

var in = None: Option[FileInputStream]
var out = None: Option[FileOutputStream]

```

I had a hard time remembering this until I came up with a little mnemonic, “Var x has No Option yeT,” where I capitalize the “T” there to stand for “type.” In my brain it looks like this:

```

var x has No Option[yeT]

```

From there it's a simple matter to get to this:

```

var x = None: Option[Type]

```

When I first started working with Scala, the only way I could think to write this code was using `null` values. The following code demonstrates the approach I used in an application that checks my email accounts. The `store` and `inbox` fields in this code are declared as `null` fields that have the `Store` and `Folder` types (from the *javax.mail* package):

```

// (1) declare the null variables
var store: Store = null
var inbox: Folder = null

try {
  // (2) use the variables/fields in the try block
  store = session.getStore("imaps")
  inbox = getFolder(store, "INBOX")
  // rest of the code here ...
  catch {
    case e: NoSuchProviderException => e.printStackTrace
    case me: MessagingException => me.printStackTrace
  } finally {
    // (3) call close() on the objects in the finally clause
    if (inbox != null) inbox.close
    if (store != null) store.close
  }
}

```

However, working in Scala gives you a chance to forget that `null` values even exist, so this is *not* a recommended approach. See [Recipe 20.5, “Eliminate null Values from Your Code”](#), for examples of how to rid your code of `null` values.

See Also

The code shown in this recipe is a Scala version of [this Oracle “Byte Streams” example](#).

3.18. Creating Your Own Control Structures

Problem

You want to define your own control structures to improve the Scala language, simplify your own code, or create a DSL for others to use.

Solution

The creators of the Scala language made a conscious decision not to implement some keywords in Scala, and instead implemented functionality through Scala libraries. This was demonstrated in [Recipe 3.5, “Implementing `break` and `continue`”](#), which showed that although the Scala language doesn’t have `break` and `continue` keywords, you can achieve the same functionality through library methods.

As a simple example of creating what appears to be a control structure, imagine for a moment that for some reason you don’t like the `while` loop and want to create your own `whilst` loop, which you can use like this:

```
package foo

import com.alvinalexander.controls.Whilst._

object WhilstDemo extends App {

    var i = 0
    whilst (i < 5) {
        println(i)
        i += 1
    }

}
```

To create your own `whilst` control structure, define a function named `whilst` that takes two parameter lists. The first parameter list handles the test condition—in this case, `i < 5`—and the second parameter list is the block of code the user wants to run.

You could implement this as a method that's just a wrapper around the `while` operator:

```
// 1st attempt
def whilst(testCondition: => Boolean)(codeBlock: => Unit) {
  while (testCondition) {
    codeBlock
  }
}
```

But a more interesting approach is to implement the `whilst` method without calling `while`. This is shown in a complete object here:

```
package com.alvinalexander.controls

import scala.annotation.tailrec

object Whilst {

  // 2nd attempt
  @tailrec
  def whilst(testCondition: => Boolean)(codeBlock: => Unit) {
    if (testCondition) {
      codeBlock
      whilst(testCondition)(codeBlock)
    }
  }
}
```

In this code, the `testCondition` is evaluated once, and if the condition is true, the `codeBlock` is executed, and then `whilst` is called recursively. This approach lets you keep checking the condition without needing a `while` or `for` loop.

Discussion

In the second `whilst` example, I used a recursive call to keep the loop running, but in a simpler example, you don't need recursion. For example, assume you want a control structure that takes two test conditions, and if both evaluate to true, you'll run a block of code that's supplied. An expression using that control structure might look like this:

```
doubleif(age > 18)(numAccidents == 0) { println("Discount!") }
```

In this case, define a function that takes three parameter lists:

```
// two 'if' condition tests
def doubleif(test1: => Boolean)(test2: => Boolean)(codeBlock: => Unit) {
  if (test1 && test2) {
    codeBlock
  }
}
```

Because `doubleIf` only needs to perform one test and doesn't need to loop indefinitely, there's no need for a recursive call in its method body. It simply checks the two test conditions, and if they evaluate to true, the `codeBlock` is executed.

See Also

- One of my favorite uses of this technique is shown in the book, *Beginning Scala* (Apress), by David Pollak. I describe how it works on [my website](#).
- The `Scala Breaks` class is demonstrated in [Recipe 3.5](#). Its source code is simple, and provides another example of how to implement a [control structure](#).

Classes and Properties

Introduction

Although Scala and Java share many similarities, the declaration of classes, class constructors, and the control of field visibility are some of the biggest differences between the two languages. Whereas Java tends to be more verbose (yet obvious), Scala is more concise, and the code you write ends up generating other code.

Recipes in this chapter will help you get through the initial learning curve related to Scala classes and fields by demonstrating how class constructors work, and the code the Scala compiler generates on your behalf when you declare constructor parameters and class fields using the `val`, `var`, and `private` keywords.

Because the Scala compiler generates accessors and mutators based on your field declarations, you may wonder how to override those methods, and this chapter provides recipes showing how to override that generated code.

Additionally, because Scala automatically sets the field type based on the value you assign, you may wonder, “What happens when a field has no initial value?” For instance, you may want to create an uninitialized field as an instance of an `Address` class. As you think about this you start typing the following code, and then wonder how to complete it:

```
var address = ?    // how to create an uninitialized Address?
```

This chapter shows the solution to that problem, demonstrates how declaring a class as a *case class* results in more than 20 additional methods being generated, shows how to write `equals` methods that work with class inheritance, and much more.



In Java, it seems correct to refer to *accessor* and *mutator* methods as “getter” and “setter” methods, primarily because of the JavaBeans standard. In this chapter, I use the terms interchangeably, but to be clear, Scala does not follow the JavaBeans naming convention for accessor and mutator methods.

4.1. Creating a Primary Constructor

Problem

You want to create a primary constructor for a class, and you quickly find that the approach is different than Java.

Solution

The primary constructor of a Scala class is a combination of:

- The constructor parameters
- Methods that are called in the body of the class
- Statements and expressions that are executed in the body of the class

Fields declared in the body of a Scala class are handled in a manner similar to Java; they are assigned when the class is first instantiated.

The following class demonstrates constructor parameters, class fields, and statements in the body of a class:

```
class Person(var firstName: String, var lastName: String) {  
  
    println("the constructor begins")  
  
    // some class fields  
    private val HOME = System.getProperty("user.home")  
    var age = 0  
  
    // some methods  
    override def toString = s"$firstName $lastName is $age years old"  
    def printHome { println(s"HOME = $HOME") }  
    def printFullName { println(this) } // uses toString  
  
    printHome  
    printFullName  
    println("still in the constructor")  
  
}
```

Because the methods in the body of the class are part of the constructor, when an instance of a `Person` class is created, you'll see the output from the `println` statements at the beginning and end of the class declaration, along with the call to the `printHome` and `printFullName` methods near the bottom of the class:

```
scala> val p = new Person("Adam", "Meyer")
the constructor begins
HOME = /Users/Al
Adam Meyer is 0 years old
still in the constructor
```

Discussion

If you're coming to Scala from Java, you'll find that the process of declaring a primary constructor in Scala is quite different. In Java it's fairly obvious when you're in the main constructor and when you're not, but Scala blurs this distinction. However, once you understand the approach, it also makes your class declarations more concise than Java class declarations.

In the example shown, the two constructor arguments `firstName` and `lastName` are defined as `var` fields, which means that they're variable, or mutable; they can be changed after they're initially set. Because the fields are mutable, Scala generates both accessor and mutator methods for them. As a result, given an instance `p` of type `Person`, you can change the values like this:

```
p.firstName = "Scott"
p.lastName = "Jones"
```

and you can access them like this:

```
println(p.firstName)
println(p.lastName)
```

Because the `age` field is declared as a `var`, it's also visible, and can be mutated and accessed:

```
p.age = 30
println(p.age)
```

The field `HOME` is declared as a `private val`, which is like making it `private` and `final` in a Java class. As a result, it can't be accessed directly by other objects, and its value can't be changed.

When you call a method in the body of the class—such as the call near the bottom of the class to the `printFullName` method—that method call is also part of the constructor. You can verify this by compiling the code to a `Person.class` file with `scalac`, and then decompiling it back into Java source code with a tool like the [JAD decompiler](#). After doing so, this is what the `Person` class constructor looks like:

```

public Person(String firstName, String lastName)
{
    super();
    this.firstName = firstName;
    this.lastName = lastName;
    Predef$.MODULE$.println("the constructor begins");
    age = 0;
    printHome();
    printFullName();
    Predef$.MODULE$.println("still in the constructor");
}

```

This clearly shows the `printHome` and `printFullName` methods call in the `Person` constructor, as well as the initial `age` being set.

When the code is decompiled, the constructor parameters and class fields appear like this:

```

private String firstName;
private String lastName;
private final String HOME = System.getProperty("user.home");
private int age;

```



Anything defined within the body of the class other than method declarations is a part of the primary class constructor. Because auxiliary constructors must always call a previously defined constructor in the same class, auxiliary constructors will also execute the same code.

A comparison with Java

The following code shows the equivalent Java version of the `Person` class:

```

// java
public class Person {

    private String firstName;
    private String lastName;
    private final String HOME = System.getProperty("user.home");
    private int age;

    public Person(String firstName, String lastName) {
        super();
        this.firstName = firstName;
        this.lastName = lastName;
        System.out.println("the constructor begins");
        age = 0;
        printHome();
        printFullName();
        System.out.println("still in the constructor");
    }
}

```

```

public String firstName() { return firstName; }
public String lastName() { return lastName; }
public int age() { return age; }

public void firstName_$eq(String firstName) {
    this.firstName = firstName;
}

public void lastName_$eq(String lastName) {
    this.lastName = lastName;
}

public void age_$eq(int age) {
    this.age = age;
}

public String toString() {
    return firstName + " " + lastName + " is " + age + " years old";
}

public void printHome() {
    System.out.println(HOME);
}

public void printFullName() {
    System.out.println(this);
}
}

```

As you can see, this is quite a bit lengthier than the equivalent Scala code. With constructors, I find that Java code is more verbose, but obvious; you don't have to reason much about what the compiler is doing for you.

Those _\$eq methods

The names of the mutator methods that are generated may look a little unusual:

```

public void firstName_$eq(String firstName) { ...
public void age_$eq(int age) { ...

```

These names are part of the Scala syntactic sugar for mutating `var` fields, and not anything you normally have to think about. For instance, the following `Person` class has a `var` field named `name`:

```

class Person {
    var name = ""
    override def toString = s"name = $name"
}

```

Because `name` is a `var` field, Scala generates accessor and mutator methods for it. What you don't normally see is that when the code is compiled, the mutator method is named

`name_$eq`. You don't see that because with Scala's syntactic sugar, you mutate the field like this:

```
p.name = "Ron Artest"
```

However, behind the scenes, Scala converts that line of code into this code:

```
p.name_$eq("Ron Artest")
```

To demonstrate this, you can run the following object that calls the mutator method in both ways (not something that's normally done):

```
object Test extends App {  
  
    val p = new Person  
  
    // the 'normal' mutator approach  
    p.name = "Ron Artest"  
    println(p)  
  
    // the 'hidden' mutator method  
    p.name_$eq("Metta World Peace")  
    println(p)  
  
}
```

When this code is run, it prints this output:

```
name = Ron Artest  
name = Metta World Peace
```

Again, there's no reason to call the `name_$eq` method in the real world, but when you get into overriding mutator methods, it's helpful to understand how this translation process works.

Summary

As shown with the equivalent Scala and Java classes, the Java code is verbose, but it's also straightforward. The Scala code is more concise, but you have to look at the constructor parameters to understand whether getters and setters are being generated for you, and you have to know that any method that's called in the body of the class is really being called from the primary constructor. This was a little confusing when I first started working with Scala, but it quickly became second nature.

4.2. Controlling the Visibility of Constructor Fields

Problem

You want to control the visibility of fields that are used as constructor parameters in a Scala class.

Solution

As shown in the following examples, the visibility of constructor fields in a Scala class is controlled by whether the fields are declared as `val`, `var`, without either `val` or `var`, and whether `private` is also added to the fields.

Here's the short version of the solution:

- If a field is declared as a `var`, Scala generates both getter and setter methods for that field.
- If the field is a `val`, Scala generates only a getter method for it.
- If a field doesn't have a `var` or `val` modifier, Scala gets conservative, and doesn't generate a getter or setter method for the field.
- Additionally, `var` and `val` fields can be modified with the `private` keyword, which prevents getters and setters from being generated.

See the examples that follow for more details.

var fields

If a constructor parameter is declared as a `var`, the value of the field *can* be changed, so Scala generates both getter and setter methods for that field. In the following examples, the constructor parameter name is declared as a `var`, so the field can be accessed and mutated:

```
scala> class Person(var name: String)
defined class Person

scala> val p = new Person("Alvin Alexander")
p: Person = Person@369e58be

// getter
scala> p.name
res0: String = Alvin Alexander

// setter
scala> p.name = "Fred Flintstone"
p.name: String = Fred Flintstone

scala> p.name
res1: String = Fred Flintstone
```

As shown, Scala does not follow the JavaBean naming convention when generating accessor and mutator methods.

val fields

If a constructor field is defined as a `val`, the value of the field *can't* be changed once it's been set; it's immutable (like `final` in Java). Therefore it makes sense that it should have an accessor method, and should *not* have a mutator method:

```
scala> class Person(val name: String)
defined class Person

scala> val p = new Person("Alvin Alexander")
p: Person = Person@3f9f332b

scala> p.name
res0: String = Alvin Alexander

scala> p.name = "Fred Flintstone"
<console>:11: error: reassignment to val
    p.name = "Fred Flintstone"
        ^
```

The last example fails because a mutator method is not generated for a `val` field.

Fields without val or var

When neither `val` nor `var` are specified on constructor parameters, the visibility of the field becomes very restricted, and Scala doesn't generate accessor or mutator methods:

```
scala> class Person(name: String)
defined class Person

scala> val p = new Person("Alvin Alexander")
p: Person = Person@144b6a6c

scala> p.name
<console>:12: error: value name is not a member of Person
    p.name
        ^
```

Adding private to val or var

In addition to these three basic configurations, you can add the `private` keyword to a `val` or `var` field. This keyword prevents getter and setter methods from being generated, so the field can only be accessed from within members of the class:

```
scala> class Person(private var name: String) { def getName {println(name)} }
defined class Person

scala> val p = new Person("Alvin Alexander")
p: Person = Person@3cb7cee4

scala> p.name
<console>:10: error: variable name in class Person cannot be accessed in Person
    p.name
        ^
```



```
scala> p.getName
Alvin Alexander
```

Attempting to access `p.name` fails because a getter method is not generated for the `name` field, so callers can't access it directly, but `p.getName` works because it can access the `name` field.

Discussion

If this is a little confusing, it helps to think about the choices the compiler has when generating code for you. When a field is defined as a `val`, by definition its value can't be changed, so it makes sense to generate a getter, but no setter. By definition, the value of a `var` field *can* be changed, so generating both a getter and setter make sense for it.

The `private` setting on a constructor parameter gives you additional flexibility. When it's added to a `val` or `var` field, the getter and setter methods are generated as before, but they're marked `private`. (I rarely use this feature, but it's there if you need it.)

The accessors and mutators that are generated for you based on these settings are summarized in [Table 4-1](#).

Table 4-1. The effect of constructor parameter settings

Visibility	Accessor?	Mutator?
<code>var</code>	Yes	Yes
<code>val</code>	Yes	No
Default visibility (no <code>var</code> or <code>val</code>)	No	No
Adding the <code>private</code> keyword to <code>var</code> or <code>val</code>	No	No

You can also manually add your own accessor and mutator methods. See [Recipe 4.6, “Overriding Default Accessors and Mutators”](#), for more information.

Case classes

Parameters in the constructor of a *case class* differ from these rules in one way. Case class constructor parameters are `val` by default. So if you define a case class field without adding `val` or `var`, like this:

```
case class Person(name: String)
```

you can still access the field, just as if it were defined as a `val`:

```
scala> val p = Person("Dale Cooper")
p: Person = Person(Dale Cooper)

scala> p.name
res0: String = Dale Cooper
```

Although this is slightly different than a “regular” class, it’s a nice convenience and has to do with the way case classes are intended to be used in functional programming, i.e., as immutable records. See [Recipe 4.14, “Generating Boilerplate Code with Case Classes”](#), for more information about how case classes work.

4.3. Defining Auxiliary Constructors

Problem

You want to define one or more auxiliary constructors for a class to give consumers of the class different ways to create object instances.

Solution

Define the auxiliary constructors as methods in the class with the name `this`. You can define multiple auxiliary constructors, but they must have different signatures (parameter lists). Also, each constructor must call one of the previously defined constructors.

The following example demonstrates a primary constructor and three auxiliary constructors:

```
// primary constructor
class Pizza (var crustSize: Int, var crustType: String) {

    // one-arg auxiliary constructor
    def this(crustSize: Int) {
        this(crustSize, Pizza.DEFAULT_CRUST_TYPE)
    }

    // one-arg auxiliary constructor
    def this(crustType: String) {
        this(Pizza.DEFAULT_CRUST_SIZE, crustType)
    }

    // zero-arg auxiliary constructor
    def this() {
        this(Pizza.DEFAULT_CRUST_SIZE, Pizza.DEFAULT_CRUST_TYPE)
    }

    override def toString = s"A $crustSize inch pizza with a $crustType crust"
}

object Pizza {
    val DEFAULT_CRUST_SIZE = 12
    val DEFAULT_CRUST_TYPE = "THIN"
}
```

Given these constructors, the same pizza can be created in the following ways:

```

val p1 = new Pizza(Pizza.DEFAULT_CRUST_SIZE, Pizza.DEFAULT_CRUST_TYPE)
val p2 = new Pizza(Pizza.DEFAULT_CRUST_SIZE)
val p3 = new Pizza(Pizza.DEFAULT_CRUST_TYPE)
val p4 = new Pizza

```

Discussion

There are several important points to this recipe:

- Auxiliary constructors are defined by creating methods named `this`.
- Each auxiliary constructor must begin with a call to a previously defined constructor.
- Each constructor must have a different signature.
- One constructor calls another constructor with the name `this`.

In the example shown, all of the auxiliary constructors call the primary constructor, but this isn't necessary; an auxiliary constructor just needs to call one of the previously defined constructors. For instance, the auxiliary constructor that takes the `crustType` parameter could have been written like this:

```

def this(crustType: String) {
    this(Pizza.DEFAULT_CRUST_SIZE)
    this.crustType = Pizza.DEFAULT_CRUST_TYPE
}

```

Another important part of this example is that the `crustSize` and `crustType` parameters are declared in the primary constructor. This isn't necessary, but doing this lets Scala generate the accessor and mutator methods for those parameters for you. You could start to write a similar class as follows, but this approach requires more code:

```

class Pizza () {

    var crustSize = 0
    var crustType = ""

    def this(crustSize: Int) {
        this()
        this.crustSize = crustSize
    }

    def this(crustType: String) {
        this()
        this.crustType = crustType
    }
}

```

```
// more constructors here ...

override def toString = s"A $crustSize inch pizza with a $crustType crust"

}
```

To summarize, if you want the accessors and mutators to be generated for you, put them in the primary constructor.



Although the approach shown in the Solution is perfectly valid, before creating multiple class constructors like this, take a few moments to read [Recipe 4.5, “Providing Default Values for Constructor Parameters”](#). Using that recipe can often eliminate the need for multiple constructors.

Generating auxiliary constructors for case classes

A *case class* is a special type of class that generates a *lot* of boilerplate code for you. Because of the way they work, adding what appears to be an auxiliary constructor to a case class is different than adding an auxiliary constructor to a “regular” class. This is because they’re not really constructors: they’re *apply* methods in the companion object of the class.

To demonstrate this, assume that you start with this case class in a file named *Person.scala*:

```
// initial case class
case class Person (var name: String, var age: Int)
```

This lets you create a new *Person* instance without using the *new* keyword, like this:

```
val p = Person("John Smith", 30)
```

This appears to be a different form of a constructor, but in fact, it’s a little syntactic sugar—a factory method, to be precise. When you write this line of code:

```
val p = Person("John Smith", 30)
```

behind the scenes, the Scala compiler converts it into this:

```
val p = Person.apply("John Smith", 30)
```

This is a call to an *apply* method in the companion object of the *Person* class. You don’t see this, you just see the line that you wrote, but this is how the compiler translates your code. As a result, if you want to add new “constructors” to your case class, you write new *apply* methods. (To be clear, the word “constructor” is used loosely here.)

For instance, if you decide that you want to add auxiliary constructors to let you create new *Person* instances (a) without specifying any parameters, and (b) by only specifying

their name, the solution is to add `apply` methods to the companion object of the `Person` case class in the *Person.scala* file:

```
// the case class
case class Person (var name: String, var age: Int)

// the companion object
object Person {

  def apply() = new Person("<no name>", 0)
  def apply(name: String) = new Person(name, 0)

}
```

The following test code demonstrates that this works as desired:

```
object CaseClassTest extends App {

  val a = Person()           // corresponds to apply()
  val b = Person("Pam")      // corresponds to apply(name: String)
  val c = Person("William Shatner", 82)

  println(a)
  println(b)
  println(c)

  // verify the setter methods work
  a.name = "Leonard Nimoy"
  a.age = 82
  println(a)

}
```

This code results in the following output:

```
Person(<no name>,0)
Person(Pam,0)
Person(William Shatner,82)
Person(Leonard Nimoy,82)
```

See Also

- [Recipe 6.8, “Creating Object Instances Without Using the new Keyword”](#), demonstrates how to implement the `apply` method in a companion object so you can create instances of a class without having to use the `new` keyword (or declare your class as a case class).
- [Recipe 4.5, “Providing Default Values for Constructor Parameters”](#), demonstrates an approach that can often eliminate the need for auxiliary constructors.
- [Recipe 4.14, “Generating Boilerplate Code with Case Classes”](#), details the nuts and bolts of how case classes work.

4.4. Defining a Private Primary Constructor

Problem

You want to make the primary constructor of a class private, such as to enforce the Singleton pattern.

Solution

To make the primary constructor private, insert the `private` keyword in between the class name and any parameters the constructor accepts:

```
// a private no-args primary constructor
class Order private { ...

// a private one-arg primary constructor
class Person private (name: String) { ...
```

As shown in the REPL, this keeps you from being able to create an instance of the class:

```
scala> class Person private (name: String)
defined class Person

scala> val p = new Person("Mercedes")
<console>:9: error: constructor Person in class Person cannot be accessed
in object $iw
    val p = new Person("Mercedes")
           ^
```

Discussion

A simple way to enforce the Singleton pattern in Scala is to make the primary constructor private, then put a `getInstance` method in the *companion object* of the class:

```
class Brain private {
  override def toString = "This is the brain."
}

object Brain {
  val brain = new Brain
  def getInstance = brain
}

object SingletonTest extends App {

  // this won't compile
  // val brain = new Brain
```

```
// this works
val brain = Brain.getInstance
println(brain)
}
```

You don't have to name the accessor method `getInstance`; it's only used here because of the Java convention.



A *companion object* is simply an object that's defined in the same file as a class, where the object and class have the same name. If you declare a class named `Foo` in a file named `Foo.scala`, and then declare an object named `Foo` in that same file, the `Foo` object is the companion object for the `Foo` class.

A companion object has several purposes, and one purpose is that any method declared in a companion object will appear to be a static method on the object. See [Recipe 6.6](#) for more information on creating the equivalent of Java's static methods, and [Recipe 6.8](#) for examples of how (and why) to define apply methods in a companion object.

Utility classes

Depending on your needs, creating a private class constructor may not be necessary at all. For instance, in Java you'd create a file utilities class by defining static methods in a Java class, but in Scala you do the same thing by putting all the methods in a Scala *object*:

```
object FileUtils {

  def readFile(filename: String) = {
    // code here ...
  }

  def writeToFile(filename: String, contents: String) {
    // code here ...
  }
}
```

This lets consumers of your code call these methods like this:

```
val contents = FileUtils.readFile("input.txt")
FileUtils.writeToFile("output.txt", content)
```

Because only an object is defined, code like this won't compile:

```
val utils = new FileUtils // won't compile
```

So in this case, there's no need for a private class constructor; just don't define a class.

4.5. Providing Default Values for Constructor Parameters

Problem

You want to provide a default value for a constructor parameter, which gives other classes the option of specifying that parameter when calling the constructor, or not.

Solution

Give the parameter a default value in the constructor declaration. Here's a simple declaration of a `Socket` class with one constructor parameter named `timeout` that has a default value of 10000:

```
class Socket (val timeout: Int = 10000)
```

Because the parameter is defined with a default value, you can call the constructor without specifying a timeout value, in which case you get the default value:

```
scala> val s = new Socket
s: Socket = Socket@7862af46
```

```
scala> s.timeout
res0: Int = 10000
```

You can also specify the desired timeout value when creating a new `Socket`:

```
scala> val s = new Socket(5000)
s: Socket = Socket@6df5205c
```

```
scala> s.timeout
res1: Int = 5000
```

If you prefer the approach of using named parameters when calling a constructor (or method), you can also use this approach to construct a new `Socket`:

```
scala> val s = new Socket(timeout=5000)
s: Socket = Socket@52aaf3d2
```

```
scala> s.timeout
res0: Int = 5000
```

Discussion

This recipe demonstrates a powerful feature that can eliminate the need for auxiliary constructors. As shown in the Solution, the following single constructor is the equivalent of two constructors:

```
class Socket (val timeout: Int = 10000)
```

If this feature didn't exist, two constructors would be required to get the same functionality; a primary one-arg constructor and an auxiliary zero-args constructor:


```
class Socket(val timeout: Int) {

    def this() = this(10000)
    override def toString = s"timeout: $timeout"

}
```

Multiple parameters

Taking this approach a step further, you can provide default values for multiple constructor parameters:

```
class Socket(val timeout: Int = 1000, val linger: Int = 2000) {
    override def toString = s"timeout: $timeout, linger: $linger"
}
```

Though you’ve defined only one constructor, your class now appears to have three constructors:

```
scala> println(new Socket)
timeout: 1000, linger: 2000

scala> println(new Socket(3000))
timeout: 3000, linger: 2000

scala> println(new Socket(3000, 4000))
timeout: 3000, linger: 4000
```

Using named parameters

As shown in the Solution, you can also provide the names of constructor parameters when creating objects, in a manner similar to Objective-C and other languages. This means you can also create new `Socket` instances like this:

```
println(new Socket(timeout=3000, linger=4000))
println(new Socket(linger=4000, timeout=3000))
println(new Socket(timeout=3000))
println(new Socket(linger=4000))
```

See [Recipe 5.4, “Using Parameter Names When Calling a Method”](#), for more examples of how to use parameter names in method calls.

See Also

[Recipe 4.3, “Defining Auxiliary Constructors”](#), for more information on creating auxiliary class constructors

4.6. Overriding Default Accessors and Mutators

Problem

You want to override the getter or setter methods that Scala generates for you.

Solution

This is a bit of a trick problem, because you can't override the getter and setter methods Scala generates for you, at least not if you want to stick with the Scala naming conventions. For instance, if you have a class named `Person` with a constructor parameter named `name`, and attempt to create getter and setter methods according to the Scala conventions, your code won't compile:

```
// error: this won't work
class Person(private var name: String) {
  // this line essentially creates a circular reference
  def name = name
  def name_=(aName: String) { name = aName }
}
```

Attempting to compile this code generates three errors:

```
Person.scala:3: error: overloaded method name needs result type
  def name = name
      ^
Person.scala:4: error: ambiguous reference to overloaded definition,
both method name_ in class Person of type (aName: String)Unit
and method name_ in class Person of type (x$1: String)Unit
match argument types (String)
  def name_=(aName: String) { name = aName }
      ^
Person.scala:4: error: method name_= is defined twice
  def name_=(aName: String) { name = aName }
      ^
three errors found
```

I'll examine these problems more in the Discussion, but the short answer is that both the constructor parameter and the getter method are named `name`, and Scala won't allow that.

To solve this problem, change the name of the field you use in the class constructor so it won't collide with the name of the getter method you want to use. A common approach is to add a leading underscore to the parameter name, so if you want to manually create a getter method called `name`, use the parameter name `_name` in the constructor, then declare your getter and setter methods according to the Scala conventions:

```
class Person(private var _name: String) {
  def name = _name // accessor
```

```
def name_=(aName: String) { _name = aName } // mutator
}
```

Notice the constructor parameter is declared `private` and `var`. The `private` keyword keeps Scala from exposing that field to other classes, and the `var` lets the value of the field be changed.

Creating a getter method named `name` and a setter method named `name_ =` conforms to the Scala convention and lets a consumer of your class write code like this:

```
val p = new Person("Jonathan")
p.name = "Jony" // setter
println(p.name) // getter
```

If you don't want to follow this Scala naming convention for getters and setters, you can use any other approach you want. For instance, you can name your methods `getName` and `setName`, following the JavaBean style. (However, if JavaBeans are what you really want, you may be better off using the `@BeanProperty` annotation, as described in [Recipe 17.6, “When Java Code Requires JavaBeans”](#).)

Discussion

When you define a constructor parameter to be a `var` field, Scala makes the field `private` to the class and automatically generates getter and setter methods that other classes can use to access the field. For instance, given a simple class like this:

```
class Stock (var symbol: String)
```

after the class is compiled with `scalac`, you'll see this signature when you disassemble it with `javap`:

```
$ javap Stock

public class Stock extends java.lang.Object{
  public java.lang.String symbol();
  public void symbol_$eq(java.lang.String);
  public Stock(java.lang.String);
}
```

You can see that the Scala compiler generated two methods: a getter named `symbol` and a setter named `symbol_$eq`. This second method is the same as a method you'd name `symbol_ =`, but Scala needs to translate the `=` symbol to `$eq` to work with the JVM.

That second method name is a little unusual, but it follows a Scala convention, and when it's mixed with some syntactic sugar, it lets you set the `symbol` field on a `Stock` instance like this:

```
stock.symbol = "GOOG"
```

The way this works is that behind the scenes, Scala converts that line of code into this line of code:

```
stock.symbol_$eq("GOOG")
```

You generally never have to think about this, unless you want to override the mutator method.

Summary

As shown in the Solution, the recipe for overriding default getter and setter methods is:

1. Create a `private var` constructor parameter with a name you want to reference from within your class. In the example in the Solution, the field is named `_name`.
2. Define getter and setter names that you want other classes to use. In the Solution the getter name is `name`, and the setter name is `name_ =` (which, combined with Scala's syntactic sugar, lets users write `p.name = "Jony"`).
3. Modify the body of the getter and setter methods as desired.

It's important to remember the `private` setting on your field. If you forget to control the access with `private` (or `private[this]`), you'll end up with getter/setter methods for the field you meant to hide. For example, in the following code, I intentionally left the `private` modifier off of the `_symbol` constructor parameter:

```
// intentionally left the 'private' modifier off _symbol
class Stock (var _symbol: String) {

  // getter
  def symbol = _symbol

  // setter
  def symbol_ = (s: String) {
    this.symbol = s
    println(s"symbol was updated, new value is $symbol")
  }

}
```

Compiling and disassembling this code shows the following class signature, including two methods I “accidentally” made visible:

```
public class Stock extends java.lang.Object{
  public java.lang.String _symbol();           // error
  public void _symbol_$eq(java.lang.String);  // error
  public java.lang.String symbol();
  public void symbol_$eq(java.lang.String);
  public Stock(java.lang.String);
}
```

Correctly adding `private` to the `_symbol` field results in the correct signature in the disassembled code:

```
public class Stock extends java.lang.Object{
  public java.lang.String symbol();           // println(stock.symbol)
```

```

    public void symbol_$eq(java.lang.String); // stock.symbol = "AAPL"
    public Stock(java.lang.String);
}

```

Note that while these examples used fields in a class constructor, the same principles hold true for fields defined inside a class.

4.7. Preventing Getter and Setter Methods from Being Generated

Problem

When you define a class field as a `var`, Scala automatically generates getter and setter methods for the field, and defining a field as a `val` automatically generates a getter method, but you don't want either a getter or setter.

Solution

Define the field with the `private` or `private[this]` access modifiers, as shown with the `currentPrice` field in this example:

```

class Stock {

    // getter and setter methods are generated
    var delayedPrice: Double = _

    // keep this field hidden from other classes
    private var currentPrice: Double = _
}

```

When you compile this class with `scalac`, and then disassemble it with `javap`, you'll see this interface:

```

// Compiled from "Stock.scala"
public class Stock extends java.lang.Object implements scala.ScalaObject{
    public double delayedPrice();
    public void delayedPrice_$eq(double);
    public Stock();
}

```

This shows that getter and setter methods are defined for the `delayedPrice` field, and there are no getter or setter methods for the `currentPrice` field, as desired.

Discussion

Defining a field as `private` limits the field so it's only available to instances of the same class, in this case instances of the `Stock` class. To be clear, any instance of a `Stock` class can access a private field of any other `Stock` instance.

As an example, the following code yields true when the Driver object is run, because the `isHigher` method in the `Stock` class can access the `price` field both (a) in its object, and (b) in the other `Stock` object it's being compared to:

```
class Stock {  
  // a private field can be seen by any Stock instance  
  private var price: Double = _  
  def setPrice(p: Double) { price = p }  
  def isHigher(that: Stock): Boolean = this.price > that.price  
}  
  
object Driver extends App {  
  
  val s1 = new Stock  
  s1.setPrice(20)  
  
  val s2 = new Stock  
  s2.setPrice(100)  
  
  println(s2.isHigher(s1))  
}
```

Object-private fields

Defining a field as `private[this]` takes this privacy a step further, and makes the field *object-private*, which means that it can only be accessed from the object that contains it. Unlike `private`, the field can't also be accessed by other instances of the same type, making it more private than the plain `private` setting.

This is demonstrated in the following example, where changing `private` to `private[this]` in the `Stock` class no longer lets the `isHigher` method compile:

```
class Stock {  
  // a private[this] var is object-private, and can only be seen  
  // by the current instance  
  private[this] var price: Double = _  
  
  def setPrice(p: Double) { price = p }  
  
  // error: this method won't compile because price is now object-private  
  def isHigher(that: Stock): Boolean = this.price > that.price  
}
```

Attempting to compile this class generates the following error:

```
Stock.scala:5: error: value price is not a member of Stock  
  def isHigher(that: Stock): Boolean = this.price > that.price  
                                         ^  
one error found
```

4.8. Assigning a Field to a Block or Function

Problem

You want to initialize a field in a class using a block of code, or by calling a function.

Solution

Set the field equal to the desired block of code or function. Optionally, define the field as lazy if the algorithm requires a long time to run.

In the following example, the field `text` is set equal to a block of code, which either returns (a) the text contained in a file, or (b) an error message, depending on whether the file exists and can be read:

```
class Foo {  
  
    // set 'text' equal to the result of the block of code  
    val text = {  
        var lines = ""  
        try {  
            lines = io.Source.fromFile("/etc/passwd").getLines.mkString  
        } catch {  
            case e: Exception => lines = "Error happened"  
        }  
        lines  
    }  
  
    println(text)  
}  
  
object Test extends App {  
    val f = new Foo  
}
```

Because the assignment of the code block to the `text` field and the `println` statement are both in the body of the `Foo` class, they are in the class's constructor, and will be executed when a new instance of the class is created. Therefore, compiling and running this example will either print the contents of the file, or the “Error happened” message from the catch block.

In a similar way, you can assign a class field to the results of a method or function:

```
class Foo {  
    import scala.xml.XML  
  
    // assign the xml field to the result of the load method  
    val xml = XML.load("http://example.com/foo.xml")  
}
```

```
// more code here ...  
}
```

Discussion

When it makes sense, define a field like this to be *lazy*, meaning it won't be evaluated until it is accessed. To demonstrate this, ignore the potential for errors and shorten the class to this:

```
class Foo {  
    val text =  
        io.Source.fromFile("/etc/passwd").getLines().foreach(println)  
}  
  
object Test extends App {  
    val f = new Foo  
}
```

When this code is compiled and run on a Unix system, the contents of the */etc/passwd* file are printed. That's interesting, but notice what happens when you change the block to define the *text* field as *lazy*:

```
class Foo {  
    lazy val text =  
        io.Source.fromFile("/etc/passwd").getLines().foreach(println)  
}  
  
object Test extends App {  
    val f = new Foo  
}
```

When this code is compiled and run, there is no output, because the *text* field isn't initialized until it's accessed. That's how a *lazy* field works.

Defining a field as *lazy* is a useful approach when the field might not be accessed in the normal processing of your algorithms, or if running the algorithm will take a long time, and you want to defer that to a later time.

4.9. Setting Uninitialized var Field Types

Problem

You want to set the type for an uninitialized *var* field in a class, so you begin to write code like this:

```
var x =
```

and then wonder how to finish writing the expression.

Solution

In general, define the field as an `Option`. For certain types, such as `String` and numeric fields, you can specify default initial values.

For instance, imagine that you're starting a social network, and to encourage people to sign up, you only ask for a username and password during the registration process. Therefore, you define username and password as fields in your class constructor:

```
case class Person(var username: String, var password: String) ...
```

However, later on, you'll also want to get other information from users, including their age, first name, last name, and address. Declaring those first three `var` fields is simple:

```
var age = 0
var firstName = ""
var lastName = ""
```

But what do you do when you get to the address?

The solution is to define the address field as an `Option`, as shown here:

```
case class Person(var username: String, var password: String) {

  var age = 0
  var firstName = ""
  var lastName = ""
  var address = None: Option[Address]

}

case class Address(city: String, state: String, zip: String)
```

Later, when a user provides an address, you can assign it using a `Some[Address]`, like this:

```
val p = Person("alvinalexander", "secret")
p.address = Some(Address("Talkeetna", "AK", "99676"))
```

When you need to access the address field, there are a variety of approaches you can use, and these are discussed in detail in [Recipe 20.6](#). As one example, if you want to print the fields of an `Address`, calling `foreach` on the address field works well:

```
p.address.foreach { a =>
  println(a.city)
  println(a.state)
  println(a.zip)
}
```

If the field hasn't been assigned, address is a `None`, and calling `foreach` on it does no harm, the loop is just skipped over. If the address field is assigned, it will be a `Some[Address]`, so the `foreach` loop will be entered and the data printed.

Discussion

In a related situation, setting the type on numeric `var` fields can occasionally be interesting. For instance, it's easy to create an `Int` or `Double` field:

```
var i = 0    // Int
var d = 0.0  // Double
```

In those cases, the compiler automatically defaults to the desired types, but what if you want a different numeric type? This approach lets you give each field the proper type, and a default value:

```
var b: Byte   = 0
var c: Char   = 0
var f: Float  = 0
var l: Long   = 0
var s: Short  = 0
```

See Also

- The [Option class](#)
- Don't set fields like this to `null`; Scala provides a terrific opportunity for you to get away from ever using `null` values again. See [Recipe 20.5, “Eliminate null Values from Your Code”](#), for ways to eliminate common uses of `null` values.
- In many Scala frameworks, such as the Play Framework, fields like this are commonly declared as `Option` values. See [Recipe 20.6, “Using the Option/Some/None Pattern”](#), for a detailed discussion of this approach.

4.10. Handling Constructor Parameters When Extending a Class

Problem

You want to extend a base class, and need to work with the constructor parameters declared in the base class, as well as new parameters in the subclass.

Solution

Declare your base class as usual with `val` or `var` constructor parameters. When defining a subclass constructor, leave the `val` or `var` declaration off of the fields that are common to both classes. Then define new constructor parameters in the subclass as `val` or `var` fields, as usual.

For example, first define a `Person` base class:

```
class Person (var name: String, var address: Address) {
  override def toString = if (address == null) name else s"$name @ $address"
}
```

Next define Employee as a subclass of Person, so that it takes the constructor parameters name, address, and age. The name and address parameters are common to the parent Person class, so leave the var declaration off of those fields, but age is new, so declare it as a var:

```
class Employee (name: String, address: Address, var age: Int)
  extends Person (name, address) {
  // rest of the class
}
```

With this Employee class and an Address case class:

```
case class Address (city: String, state: String)
```

you can create a new Employee as follows:

```
val teresa = new Employee("Teresa", Address("Louisville", "KY"), 25)
```

By placing all that code in the REPL, you can see that all of the fields work as expected:

```
scala> teresa.name
res0: String = Teresa

scala> teresa.address
res1: Address = Address(Louisville,KY)

scala> teresa.age
res2: Int = 25
```

Discussion

To understand how constructor parameters in a subclass work, it helps to understand how the Scala compiler translates your code. Because the following Person class defines its constructor parameters as var fields:

```
class Person (var name: String, var address: Address) {
  override def toString = if (address == null) name else s"$name @ $address"
}
```

the Scala compiler generates both accessor and mutator methods for the class. You can demonstrate this by compiling and then disassembling the Person class.

First, put this code in a file named *Person.scala*:

```
case class Address (city: String, state: String)

class Person (var name: String, var address: Address) {
  override def toString = if (address == null) name else s"$name @ $address"
}
```

Then compile the code with `scalac`, and disassemble the *Person.class* file with `javap`:

```
$ javap Person
Compiled from "Person.scala"
public class Person extends java.lang.Object implements scala.ScalaObject{
    public java.lang.String name();
    public void name_$eq(java.lang.String);
    public Address address();
    public void address_$eq(Address);
    public java.lang.String toString();
    public Person(java.lang.String, Address);
}
```

As shown, the `Person` class contains the `name`, `name_$eq`, `address`, and `address_$eq` methods, which are the accessor and mutator methods for the `name` and `address` fields. (See [Recipe 6.8](#) for an explanation of how those mutator methods work.)

This raises the question, if you define an `Employee` class that extends `Person`, how should you handle the `name` and `address` fields in the `Employee` constructor? Assuming `Employee` adds no new parameters, there are at least two main choices:

```
// Option 1: define name and address as 'var'
class Employee (var name: String, var address: Address)
extends Person (name, address) { ... }

// Option 2: define name and address without var or val
class Employee (name: String, address: Address)
extends Person (name, address) { ... }
```

Because Scala has already generated the getter and setter methods for the `name` and `address` fields in the `Person` class, the solution is to declare the `Employee` constructor without `var` declarations:

```
// this is correct
class Employee (name: String, address: Address)
extends Person (name, address) { ... }
```

Because you don't declare the parameters in `Employee` as `var`, Scala won't attempt to generate methods for those fields. You can demonstrate this by adding the `Employee` class definition to the code in *Person.scala*:

```
case class Address (city: String, state: String)
class Person (var name: String, var address: Address) {
    override def toString = if (address == null) name else s"$name @ $address"
}
class Employee (name: String, address: Address)
extends Person (name, address) {
    // code here ...
}
```

Compiling the code with `scalac` and then disassembling the *Employee.class* file with `javap`, you see the following, expected result:

```
$ javap Employee
Compiled from "Person.scala"
public class Employee extends Person implements scala.ScalaObject{
    public Employee(java.lang.String, Address);
}
```

The `Employee` class extends `Person`, and Scala did not generate any methods for the `name` and `address` fields. Therefore, the `Employee` class inherits that behavior from `Person`.

While this example shows how Scala works with `var` fields, you can follow the same line of reasoning with `val` fields as well.

4.11. Calling a Superclass Constructor

Problem

You want to control the superclass constructor that's called when you create constructors in a subclass.

Solution

This is a bit of a trick question, because you *can* control the superclass constructor that's called by the primary constructor in a subclass, but you *can't* control the superclass constructor that's called by an auxiliary constructor in the subclass.

When you define a subclass in Scala, you control the superclass constructor that's called by its primary constructor when you define the `extends` portion of the subclass declaration. For instance, in the following code, the `Dog` class is defined to call the primary constructor of the `Animal` class, which is a one-arg constructor that takes `name` as its parameter:

```
class Animal (var name: String) {
    // ...
}
class Dog (name: String) extends Animal (name) {
    // ...
}
```

However, if the `Animal` class has multiple constructors, the primary constructor of the `Dog` class can call any of those constructors.

For example, the primary constructor of the `Dog` class in the following code calls the one-arg auxiliary constructor of the `Animal` class by specifying that constructor in its `extends` clause:

```
// (1) primary constructor
class Animal (var name: String, var age: Int) {
```

```

// (2) auxiliary constructor
def this (name: String) {
    this(name, 0)
}

override def toString = s"$name is $age years old"
}

// calls the Animal one-arg constructor
class Dog (name: String) extends Animal (name) {
    println("Dog constructor called")
}

```

Alternatively, it could call the two-arg primary constructor of the `Animal` class:

```

// call the two-arg constructor
class Dog (name: String) extends Animal (name, 0) {
    println("Dog constructor called")
}

```

Auxiliary constructors

Regarding auxiliary constructors, because the first line of an auxiliary constructor must be a call to another constructor of the current class, there is no way for auxiliary constructors to call a superclass constructor.

As you can see in the following code, the primary constructor of the `Employee` class can call any constructor in the `Person` class, but the auxiliary constructors of the `Employee` class must call a previously defined constructor of its own class with the `this` method as its first line:

```

case class Address (city: String, state: String)
case class Role (role: String)

class Person (var name: String, var address: Address) {

    // no way for Employee auxiliary constructors to call this constructor
    def this (name: String) {
        this(name, null)
        address = null
    }

    override def toString = if (address == null) name else s"$name @ $address"
}

class Employee (name: String, role: Role, address: Address)
extends Person (name, address) {

    def this (name: String) {
        this(name, null, null)
    }
}

```

```

def this (name: String, role: Role) {
  this(name, role, null)
}

def this (name: String, address: Address) {
  this(name, null, address)
}
}

```

Therefore, there's no direct way to control which superclass constructor is called from an auxiliary constructor in a subclass. In fact, because each auxiliary constructor must call a previously defined constructor in the same class, all auxiliary constructors will eventually call the same superclass constructor that's called from the subclass's primary constructor.

4.12. When to Use an Abstract Class

Problem

Scala has traits, and a trait is more flexible than an abstract class, so you wonder, “When should I use an abstract class?”

Solution

There are two main reasons to use an abstract class in Scala:

- You want to create a base class that requires constructor arguments.
- The code will be called from Java code.

Regarding the first reason, traits don't allow constructor parameters:

```

// this won't compile
trait Animal(name: String)

```

So, use an abstract class whenever a base behavior must have constructor parameters:

```

abstract class Animal(name: String)

```

Regarding the second reason, if you're writing code that needs to be accessed from Java, you'll find that Scala traits with implemented methods can't be called from Java code. If you run into this situation, see [Recipe 17.7, “Wrapping Traits with Implementations”](#), for solutions to that problem.

Discussion

Use an abstract class instead of a trait when the base functionality must take constructor parameters. However, be aware that a class can extend only one abstract class.

Abstract classes work just like Java in that you can define some methods that have complete implementations, and other methods that have no implementation and are therefore abstract. To declare that a method is abstract, just leave the body of the method undefined:

```
def speak // no body makes the method abstract
```

There is no need for an `abstract` keyword; simply leaving the body of the method undefined makes it abstract. This is consistent with how abstract methods in traits are defined.

In the following example, the methods `save`, `update`, and `delete` are defined in the abstract class `BaseController`, but the `connect`, `getStatus`, and `setServerName` methods have no method body, and are therefore abstract:

```
abstract class BaseController(db: Database) {  
  
  def save { db.save }  
  def update { db.update }  
  def delete { db.delete }  
  
  // abstract  
  def connect  
  
  // an abstract method that returns a String  
  def getStatus: String  
  
  // an abstract method that takes a parameter  
  def setServerName(serverName: String)  
}
```

When a class extends the `BaseController` class, it must implement the `connect`, `getStatus`, and `setServerName` methods, or be declared abstract. Attempting to extend `BaseController` without implementing those methods yields a “class needs to be abstract” error, as shown in the REPL:

```
scala> class WidgetController(db: Database) extends BaseController(db)  
<console>:9: error: class WidgetController needs to be abstract, since:  
method setServerName in class BaseController of type (serverName: String)Unit  
is not defined  
method getStatus in class BaseController of type => String is not defined  
method connect in class BaseController of type => Unit is not defined  
class WidgetController(db: Database) extends BaseController(db)  
  ^
```

Because a class can extend only one abstract class, when you’re trying to decide whether to use a trait or abstract class, always use a trait, unless you have this specific need to have constructor arguments in your base implementation.

4.13. Defining Properties in an Abstract Base Class (or Trait)

Problem

You want to define abstract or concrete properties in an abstract base class (or trait) that can be referenced in all child classes.

Solution

You can declare both `val` and `var` fields in an abstract class (or trait), and those fields can be abstract or have concrete implementations. All of these variations are shown in this recipe.

Abstract `val` and `var` fields

The following example demonstrates an `Animal` trait with abstract `val` and `var` fields, along with a simple concrete method named `sayHello`, and an override of the `toString` method:

```
abstract class Pet (name: String) {  
    val greeting: String  
    var age: Int  
    def sayHello { println(greeting) }  
    override def toString = s"I say $greeting, and I'm $age"  
}
```

The following `Dog` and `Cat` classes extend the `Animal` class and provide values for the `greeting` and `age` fields. Notice that the fields are again specified as `val` or `var`:

```
class Dog (name: String) extends Pet (name) {  
    val greeting = "Woof"  
    var age = 2  
}  
  
class Cat (name: String) extends Pet (name) {  
    val greeting = "Meow"  
    var age = 5  
}
```

The functionality can be demonstrated with a simple driver object:

```
object AbstractFieldsDemo extends App {  
    val dog = new Dog("Fido")  
    val cat = new Cat("Morris")  
  
    dog.sayHello  
    cat.sayHello  
  
    println(dog)  
    println(cat)
```

```

    // verify that the age can be changed
    cat.age = 10
    println(cat)
}

```

The resulting output looks like this:

```

Woof
Meow
I say Woof, and I'm 2
I say Meow, and I'm 5
I say Meow, and I'm 10

```

Concrete field implementations are presented in the Discussion, because it helps to understand how the Scala compiler translates your code in the preceding examples.

Discussion

As shown, you can declare abstract fields in an abstract class as either `val` or `var`, depending on your needs. The way abstract fields work in abstract classes (or traits) is interesting:

- An abstract `var` field results in getter and setter methods being generated for the field.
- An abstract `val` field results in a getter method being generated for the field.
- When you define an abstract field in an abstract class or trait, the Scala compiler does *not* create a field in the resulting code; it only generates the methods that correspond to the `val` or `var` field.

In the example shown in the Solution, if you look at the code that's created by `scalac` using the `-Xprint:all` option, or by decompiling the resulting `Pet.class` file, you won't find greeting or age fields. For instance, if you decompile the class, the output shows only methods in the class, no fields:

```

import scala.*;
import scala.runtime.BoxesRunTime;

public abstract class Pet
{
    public abstract String greeting();
    public abstract int age();
    public abstract void age_$eq(int i);

    public void sayHello() {
        Predef$.MODULE$.println(greeting());
    }

    public String toString(){

```

```

    // code omitted
}

public Pet(String name){}
}

```

Because of this, when you provide concrete values for these fields in your concrete classes, you must again define your fields to be `val` or `var`. Because the fields don't actually exist in the abstract base class (or trait), the `override` keyword is not necessary.

As another result of this, you may see developers define a `def` that takes no parameters in the abstract base class rather than defining a `val`. They can then define a `val` in the concrete class, if desired. This technique is demonstrated in the following code:

```

abstract class Pet (name: String) {
    def greeting: String
}

class Dog (name: String) extends Pet (name) {
    val greeting = "Woof"
}

object Test extends App {
    val dog = new Dog("Fido")
    println(dog.greeting)
}

```

Given this background, it's time to examine the use of concrete `val` and `var` fields in abstract classes.

Concrete `val` fields in abstract classes

When defining a concrete `val` field in an abstract class, you can provide an initial value, and then override that value in concrete subclasses:

```

abstract class Animal {
    val greeting = "Hello" // provide an initial value
    def sayHello { println(greeting) }
    def run
}

class Dog extends Animal {
    override val greeting = "Woof" // override the value
    def run { println("Dog is running") }
}

```

In this example, the `greeting` variable is created in both classes. To demonstrate this, running the following code:

```

abstract class Animal {
    val greeting = { println("Animal"); "Hello" }
}

```

```

class Dog extends Animal {
  override val greeting = { println("Dog"); "Woof" }
}

object Test extends App {
  new Dog
}

```

results in this output, showing that both values are created:

```

Animal
Dog

```

To prove this, you can also decompile both the `Animal` and `Dog` classes, where you'll find the `greeting` declared like this:

```
private final String greeting = "Hello";
```

To prevent a concrete `val` field in an abstract base class from being overridden in a subclass, declare the field as a `final val`:

```

abstract class Animal {
  final val greeting = "Hello" // made the field 'final'
}

class Dog extends Animal {
  val greeting = "Woof" // this line won't compile
}

```

Concrete var fields in abstract classes

You can also give `var` fields an initial value in your trait or abstract class, and then refer to them in your concrete subclasses, like this:

```

abstract class Animal {
  var greeting = "Hello"
  var age = 0
  override def toString = s"I say $greeting, and I'm $age years old."
}

class Dog extends Animal {
  greeting = "Woof"
  age = 2
}

```

In this case, these fields are declared and assigned in the abstract base class, as shown in the decompiled code for the `Animal` class:

```

private String greeting;
private int age;

```

```
public Animal(){
    greeting = "Hello";
    age = 0;
}

// more code ...
```

Because the fields are declared and initialized in the abstract `Animal` base class, there's no need to redeclare the fields as `val` or `var` in the concrete `Dog` subclass.

You can verify this by looking at the code the Scala compiler generates for the `Dog` class. When you compile the code with `scalac -Xprint:all`, and look at the last lines of output, you'll see how the compiler has converted the `Dog` class:

```
class Dog extends Animal {
    def <init>(): Dog = {
        Dog.super.<init>();
        Dog.this.greeting_=("Woof");
        Dog.this.age_=(2);
        ()
    }
}
```

Because the fields are concrete fields in the abstract base class, they just need to be reassigned in the concrete `Dog` class.

Don't use null

As discussed in many recipes in this book, including [Recipe 20.5, “Eliminate null Values from Your Code”](#), you shouldn't use null values in these situations. If you're tempted to use a null, instead initialize the fields using the `Option/Some/None` pattern. The following example demonstrates how to initialize `val` and `var` fields with this approach:

```
trait Animal {
    val greeting: Option[String]
    var age: Option[Int] = None
    override def toString = s"I say $greeting, and I'm $age years old."
}

class Dog extends Animal {
    val greeting = Some("Woof")
    age = Some(2)
}

object Test extends App {
    val d = new Dog
    println(d)
}
```

Running this `Test` object yields the following output:

```
I say Some(Woof), and I'm Some(2) years old.
```

See Also

See [Recipe 5.2, “Calling a Method on a Superclass”](#), for more examples of how to call methods on superclasses.

4.14. Generating Boilerplate Code with Case Classes

Problem

You’re working with match expressions, actors, or other situations where you want to use the *case class* syntax to generate boilerplate code, including accessor and mutator methods, along with `apply`, `unapply`, `toString`, `equals`, and `hashCode` methods, and more.

Solution

Define your class as a *case class*, defining any parameters it needs in its constructor:

```
// name and relation are 'val' by default
case class Person(name: String, relation: String)
```

Defining a class as a case class results in a lot of boilerplate code being generated, with the following benefits:

- An `apply` method is generated, so you don’t need to use the `new` keyword to create a new instance of the class.
- Accessor methods are generated for the constructor parameters because case class constructor parameters are `val` by default. Mutator methods are also generated for parameters declared as `var`.
- A good, default `toString` method is generated.
- An `unapply` method is generated, making it easy to use case classes in match expressions.
- `equals` and `hashCode` methods are generated.
- A `copy` method is generated.

When you define a class as a case class, you don’t have to use the `new` keyword to create a new instance:

```
scala> case class Person(name: String, relation: String)
defined class Person

// "new" not needed before Person
scala> val emily = Person("Emily", "niece")
emily: Person = Person(Emily,niece)
```

Case class constructor parameters are `val` by default, so accessor methods are generated for the parameters, but mutator methods are not generated:

```
scala> emily.name
res0: String = Emily

scala> emily.name = "Fred"
<console>:10: error: reassignment to val
    emily.name = "Fred"
           ^
```

By defining a case class constructor parameter as a `var`, both accessor and mutator methods are generated:

```
scala> case class Company (var name: String)
defined class Company

scala> val c = Company("Mat-Su Valley Programming")
c: Company = Company(Mat-Su Valley Programming)

scala> c.name
res0: String = Mat-Su Valley Programming

scala> c.name = "Valley Programming"
c.name: String = Valley Programming
```

Case classes also have a good default `toString` method implementation:

```
scala> emily
res0: Person = Person(Emily,niece)
```

Because an `unapply` method is automatically created for a case class, it works well when you need to extract information in match expressions, as shown here:

```
scala> emily match { case Person(n, r) => println(n, r) }
(Emily,niece)
```

Case classes also have generated `equals` and `hashCode` methods, so instances can be compared:

```
scala> val hannah = Person("Hannah", "niece")
hannah: Person = Person(Hannah,niece)

scala> emily == hannah
res1: Boolean = false
```

A case class even creates a `copy` method that is helpful when you need to clone an object, and change some of the fields during the process:

```
scala> case class Employee(name: String, loc: String, role: String)
defined class Employee

scala> val fred = Employee("Fred", "Anchorage", "Salesman")
fred: Employee = Employee(Fred,Anchorage,Salesman)
```

```
scala> val joe = fred.copy(name="Joe", role="Mechanic")
joe: Employee = Employee(Joe,Anchorage,Mechanic)
```

Discussion

Case classes are primarily intended to create “immutable records” that you can easily use in pattern-matching expressions. Indeed, pure FP developers look at case classes as being similar to immutable records found in ML, Haskell, and other languages.

Perhaps as a result of this, case class constructor parameters are `val` by default. As a reviewer of this book with an FP background wrote, “Case classes allow `var` fields, but then you are subverting their very purpose.”

Generated code

As shown in the Solution, when you create a case class, Scala generates a wealth of code for your class. To see the code that’s generated for you, first compile a simple case class, then disassemble it with `javap`. For example, put this code in a file named *Person.scala*:

```
case class Person(var name: String, var age: Int)
```

Then compile the file:

```
$ scalac Person.scala
```

This creates two class files, *Person.class* and *Person\$.class*. Disassemble *Person.class* with this command:

```
$ javap Person
```

This results in the following output, which is the public signature of the class:

```
Compiled from "Person.scala"
public class Person extends java.lang.Object ✎
implements scala.ScalaObject,scala.Product,scala.Serializable{
    public static final scala.Function1 tupled();
    public static final scala.Function1 curry();
    public static final scala.Function1 curried();
    public scala.collection.Iterator productIterator();
    public scala.collection.Iterator productElements();
    public java.lang.String name();
    public void name_$eq(java.lang.String);
    public int age();
    public void age_$eq(int);
    public Person copy(java.lang.String, int);
    public int copy$default$2();
    public java.lang.String copy$default$1();
    public int hashCode();
    public java.lang.String toString();
    public boolean equals(java.lang.Object);
    public java.lang.String productPrefix();
```



```

    public int productArity();
    public java.lang.Object productElement(int);
    public boolean canEqual(java.lang.Object);
    public Person(java.lang.String, int);
}

```

Then disassemble *Person\$.class*:

```

$ javap Person$

Compiled from "Person.scala"
public final class Person$ extends scala.runtime.AbstractFunction2
implements scala.ScalaObject,scala.Serializable{
    public static final Person$ MODULE$;
    public static {};
    public final java.lang.String toString();
    public scala.Option unapply(Person);
    public Person apply(java.lang.String, int);
    public java.lang.Object readResolve();
    public java.lang.Object apply(java.lang.Object, java.lang.Object);
}

```

As you can see, Scala generates a *lot* of source code when you declare a class as a case class.

As a point of comparison, if you remove the keyword `case` from that code (making it a “regular” class), compile it, and then disassemble it, Scala only generates the following code:

```

public class Person extends java.lang.Object{
    public java.lang.String name();
    public void name_$eq(java.lang.String);
    public int age();
    public void age_$eq(int);
    public Person(java.lang.String, int);
}

```

That’s a big difference. The case class results in 22 more methods than the “regular” class. If you need the functionality, this is a good thing. However, if you don’t need all this additional functionality, consider using a “regular” class declaration instead. For instance, if you just want to be able to create new instances of a class without the `new` keyword, like this:

```
val p = Person("Alex")
```

create an `apply` method in the companion object of a “regular” class, as described in [Recipe 6.8, “Creating Object Instances Without Using the new Keyword”](#). Remember, there isn’t anything in a case class you can’t code for yourself.

See Also

- [Recipe 4.3, “Defining Auxiliary Constructors”](#), shows how to write additional apply methods so a case class can appear to have multiple constructors.
- A discussion of extractors on the [official Scala website](#).

4.15. Defining an equals Method (Object Equality)

Problem

You want to define an equals method for your class so you can compare object instances to each other.

Solution

Like Java, you define an equals method (and hashCode method) in your class to compare two instances, but unlike Java, you then use the == method to compare the equality of two instances.

There are many ways to write equals methods. The following example shows one possible way to define an equals method and its corresponding hashCode method:

```
class Person (name: String, age: Int) {  
  
    def canEqual(a: Any) = a.isInstanceOf[Person]  
  
    override def equals(that: Any): Boolean =  
        that match {  
            case that: Person => that.canEqual(this) && this.hashCode == that.hashCode  
            case _ => false  
        }  
  
    override def hashCode: Int = {  
        val prime = 31  
        var result = 1  
        result = prime * result + age;  
        result = prime * result + (if (name == null) 0 else name.hashCode)  
        return result  
    }  
}
```

This example shows a modified version of a hashCode method that Eclipse generated for a similar Java class. It also uses a canEqual method, which will be explained shortly.

With the equals method defined, you can compare instances of a Person with ==, as demonstrated in the following tests:

```
import org.scalatest.FunSuite

class PersonTests extends FunSuite {

  // these first two instances should be equal
  val nimoy = new Person("Leonard Nimoy", 82)
  val nimoy2 = new Person("Leonard Nimoy", 82)
  val shatner = new Person("William Shatner", 82)
  val ed = new Person("Ed Chigliak", 20)

  // all tests pass
  test("nimoy == nimoy") { assert(nimoy == nimoy) }
  test("nimoy == nimoy2") { assert(nimoy == nimoy2) }
  test("nimoy2 == nimoy") { assert(nimoy2 == nimoy) }
  test("nimoy != shatner") { assert(nimoy != shatner) }
  test("shatner != nimoy") { assert(shatner != nimoy) }
  test("nimoy != null") { assert(nimoy != null) }
  test("nimoy != String") { assert(nimoy != "Leonard Nimoy") }
  test("nimoy != ed") { assert(nimoy != ed) }

}
```

As noted in the code comments, all of these tests pass.



These tests were created with the ScalaTest FunSuite, which is similar to writing unit tests with JUnit.

Discussion

The first thing to know about Scala and the `equals` method is that, unlike Java, you compare the equality of two objects with `==`. In Java, the `==` operator compares “reference equality,” but in Scala, `==` is a method you use on each class to compare the equality of two instances, calling your `equals` method under the covers.

As mentioned, there are many ways to implement `equals` methods, and the code in the Solution shows just one possible approach. The book *Programming in Scala* contains one chapter of more than 25 pages on “object equality,” so this is a big topic.

An important benefit of the approach shown in the Solution is that you can continue to use it when you use inheritance in classes. For instance, in the following code, the `Employee` class extends the `Person` class that’s shown in the Solution:

```
class Employee(name: String, age: Int, var role: String)
  extends Person(name, age)
{

  override def canEqual(a: Any) = a.isInstanceOf[Employee]
```

```

override def equals(that: Any): Boolean =
  that match {
    case that: Employee =>
      that.canEqual(this) && this.hashCode == that.hashCode
    case _ => false
  }

override def hashCode: Int = {
  val ourHash = if (role == null) 0 else role.hashCode
  super.hashCode + ourHash
}
}

```

This code uses the same approach to the `canEqual`, `equals`, and `hashCode` methods, and I like that consistency. Just as important as the consistency is the accuracy of the approach, especially when you get into the business of comparing instances of a child class to instances of any of its parent classes. In the case of the `Person` and `Employee` code shown, these classes pass all of the following tests:

```

class EmployeeTests extends FunSuite with BeforeAndAfter {

  // these first two instance should be equal
  val eNimoy1 = new Employee("Leonard Nimoy", 82, "Actor")
  val eNimoy2 = new Employee("Leonard Nimoy", 82, "Actor")
  val pNimoy = new Person("Leonard Nimoy", 82)
  val eShatner = new Employee("William Shatner", 82, "Actor")

  test("eNimoy1 == eNimoy1") { assert(eNimoy1 == eNimoy1) }
  test("eNimoy1 == eNimoy2") { assert(eNimoy1 == eNimoy2) }
  test("eNimoy2 == eNimoy1") { assert(eNimoy2 == eNimoy1) }
  test("eNimoy != pNimoy") { assert(eNimoy1 != pNimoy) }
  test("pNimoy != eNimoy") { assert(pNimoy != eNimoy1) }

}

```

All the tests pass, including the comparison of the `eNimoy` and `pNimoy` objects, which are instances of the `Employee` and `Person` classes, respectively.

Theory

The Scaladoc for the `equals` method of the `Any` class states, “any implementation of this method should be an *equivalence relation*.” The documentation states that an equivalence relation should have these three properties:

- It is *reflexive*: for any instance `x` of type `Any`, `x.equals(x)` should return true.
- It is *symmetric*: for any instances `x` and `y` of type `Any`, `x.equals(y)` should return true if and only if `y.equals(x)` returns true.

- It is *transitive*: for any instances *x*, *y*, and *z* of type `AnyRef`, if `x.equals(y)` returns true and `y.equals(z)` returns true, then `x.equals(z)` should return true.

Therefore, if you override the `equals` method, you should verify that your implementation remains an equivalence relation.

See Also

- The Artima website has an excellent related article titled [How to Write an Equality Method in Java](#).
- Eric Torreborre shares an excellent `canEqual` example on [GitHub](#).
- “Equivalence relation” defined on [Wikipedia](#).
- The [Scala Any class](#).

4.16. Creating Inner Classes

Problem

You want to create a class as an inner class to help keep the class out of your public API, or to otherwise encapsulate your code.

Solution

Declare one class inside another class. In the following example, a case class named `Thing` is declared inside of a class named `PandorasBox`:

```
class PandorasBox {

  case class Thing (name: String)

  var things = new collection.mutable.ArrayBuffer[Thing]()
  things += Thing("Evil Thing #1")
  things += Thing("Evil Thing #2")

  def addThing(name: String) { things += new Thing(name) }

}
```

This lets users of `PandorasBox` access the collection of `things` inside the box, while code outside of `PandorasBox` generally doesn't have to worry about the concept of a `Thing`:

```
object ClassInAClassExample extends App {

    val p = new PandorasBox
    p.things.foreach(println)

}
```

As shown, you can access the things in PandorasBox with the things method. You can also add new things to PandorasBox by calling the addThing method:

```
p.addThing("Evil Thing #3")
p.addThing("Evil Thing #4")
```

Discussion

The concept of a “class within a class” is different in Scala than in Java. As described [on the official Scala website](#), “Opposed to Java-like languages where such inner classes are members of the enclosing class, in Scala, such inner classes are bound to the outer object.” The following code demonstrates this:

```
object ClassInObject extends App {

    // inner classes are bound to the object
    val oc1 = new OuterClass
    val oc2 = new OuterClass
    val ic1 = new oc1.InnerClass
    val ic2 = new oc2.InnerClass
    ic1.x = 10
    ic2.x = 20
    println(s"ic1.x = ${ic1.x}")
    println(s"ic2.x = ${ic2.x}")

}

class OuterClass {
    class InnerClass {
        var x = 1
    }
}
```

Because inner classes are bound to their object instances, when that code is run, it prints the following output:

```
ic1.x = 10
ic2.x = 20
```

There are many other things you can do with inner classes, such as include a class inside an object or an object inside a class:

```
object InnerClassDemo2 extends App {

    // class inside object
    println(new OuterObject.InnerClass().x)
```

```
// object inside class
println(new OuterClass().InnerObject.y)

}

object OuterObject {
  class InnerClass {
    var x = 1
  }
}

class OuterClass {
  object InnerObject {
    val y = 2
  }
}
```

See Also

The Scala website has a page on [Inner Classes](#).

Introduction

Conceptually, Scala methods are similar to Java methods in that they are behaviors you add to a class. However, they differ significantly in their implementation details. The following example shows some of the differences between Java and Scala when defining a simple method that takes an integer argument and returns a string:

```
// java
public String doSomething(int x) {
    // code here
}

// scala
def doSomething(x: Int): String = {
    // code here
}
```

This is just a start, though. Scala methods can be written even more concisely. This method takes an `Int`, adds 1 to it, and returns the resulting `Int` value:

```
def plusOne(i: Int) = i + 1
```

Notice that the return type didn't have to be specified, and parentheses around the short method body aren't required.

In addition to the differences shown in these simple examples, there are other differences between Java and Scala methods, including:

- Specifying method access control (visibility)
- The ability to set default values for method parameters

- The ability to specify the names of method parameters when calling a method
- How you declare the exceptions a method can throw
- Using varargs fields in methods

This chapter demonstrates all of these method-related features.

5.1. Controlling Method Scope

Problem

Scala methods are public by default, and you want to control their scope in ways similar to Java.

Solution

Scala lets you control method visibility in a more granular and powerful way than Java. In order from “most restrictive” to “most open,” Scala provides these scope options:

- Object-private scope
- Private
- Package
- Package-specific
- Public

These scopes are demonstrated in the examples that follow.

Object-private scope

The most restrictive access is to mark a method as *object-private*. When you do this, the method is available only to the current instance of the current object. Other instances of the same class cannot access the method.

You mark a method as object-private by placing the access modifier `private[this]` before the method declaration:

```
private[this] def isFoo = true
```

In the following example, the method `doFoo` takes an instance of a `Foo` object, but because the `isFoo` method is declared as an object-private method, the code won't compile:

```
class Foo {

  private[this] def isFoo = true

  def doFoo(other: Foo) {
    if (other.isFoo) { // this line won't compile

```

```

        // ...
    }
}

```

The code won't compile because the current `Foo` instance can't access the `isFoo` method of the other instance, because `isFoo` is declared as `private[this]`. As you can see, the object-private scope is extremely restrictive.

Private scope

A slightly less restrictive access is to mark a method `private`, which makes the method available to (a) the current class and (b) other instances of the current class. This is the same as marking a method `private` in Java. By changing the access modifier from `private[this]` to `private`, the code will now compile:

```

class Foo {

  private def isFoo = true

  def doFoo(other: Foo) {
    if (other.isFoo) { // this now compiles
      // ...
    }
  }
}

```

By making a method `private`, it is not available to subclasses. The following code won't compile because the `heartBeat` method is private to the `Animal` class:

```

class Animal {
  private def heartBeat {}
}

class Dog extends Animal {
  heartBeat // won't compile
}

```

Protected scope

Marking a method `protected` makes the method available to subclasses, so the following code will compile:

```

class Animal {
  protected def breathe {}
}

class Dog extends Animal {
  breathe
}

```

The meaning of `protected` is slightly different in Scala than in Java. In Java, `protected` methods can be accessed by other classes in the same package, but this isn't true in Scala. The following code won't compile because the `Jungle` class can't access the `breathe` method of the `Animal` class, even though they're in the same package:

```
package world {

  class Animal {
    protected def breathe {}
  }

  class Jungle {
    val a = new Animal
    a.breathe // error: this line won't compile
  }

}
```

Package scope

To make a method available to all members of the current package—what would be called “package scope” in Java—mark the method as being `private` to the current package with the `private[packageName]` syntax.

In the following example, the method `doX` can be accessed by other classes in the same package (the `model` package), but the method `doY` is available only to the `Foo` class:

```
package com.acme.coolapp.model {

  class Foo {
    private[model] def doX {}
    private def doY {}
  }

  class Bar {
    val f = new Foo
    f.doX // compiles
    f.doY // won't compile
  }

}
```

More package-level control

Beyond making a method available to classes in the current package, Scala gives you more control and lets you make a method available at different levels in a class hierarchy. The following example demonstrates how you can make the methods `doX`, `doY`, and `doZ` available to different package levels:

```
package com.acme.coolapp.model {
  class Foo {
    private[model] def doX {}
```

```

        private[coolapp] def doY {}
        private[acme] def doZ {}
    }
}

import com.acme.coolapp.model._

package com.acme.coolapp.view {
    class Bar {
        val f = new Foo
        f.doX // won't compile
        f.doY
        f.doZ
    }
}

package com.acme.common {
    class Bar {
        val f = new Foo
        f.doX // won't compile
        f.doY // won't compile
        f.doZ
    }
}

```

In this example, the methods can be seen as follows:

- The method `doX` can be seen by other classes in the `model` package (`com.acme.coolapp.model`).
- The method `doY` can be seen by all classes under the `coolapp` package level.
- The method `doZ` can be seen by all classes under the `acme` level.

As you can see, this approach allows a fine-grained level of access control.

Public scope

If no access modifier is added to the method declaration, the method is public. In the following example, any class in any package can access the `doX` method:

```

package com.acme.coolapp.model {
    class Foo {
        def doX {}
    }
}

package org.xyz.bar {
    class Bar {
        val f = new com.acme.coolapp.model.Foo
        f.doX
    }
}

```

Discussion

The Scala approach to access modifiers is different than Java. Though it offers more power than Java, it's also a little more complicated.

Table 5-1 describes the levels of access control that were demonstrated in the examples in the Solution.

Table 5-1. Descriptions of Scala's access control modifiers

Access modifier	Description
<code>private[this]</code>	The method is available only to the current instance of the class it's declared in.
<code>private</code>	The method is available to the current instance and other instances of the class it's declared in.
<code>protected</code>	The method is available only to instances of the current class and subclasses of the current class.
<code>private[model]</code>	The method is available to all classes beneath the <code>com.acme.coolapp.model</code> package.
<code>private[coolapp]</code>	The method is available to all classes beneath the <code>com.acme.coolapp</code> package.
<code>private[acme]</code>	The method is available to all classes beneath the <code>com.acme</code> package.
(no modifier)	The method is public.

5.2. Calling a Method on a Superclass

Problem

To keep your code DRY (“Don’t Repeat Yourself”), you want to invoke a method that’s already defined in a parent class or trait.

Solution

In the basic use case, the syntax to invoke a method in an immediate parent class is the same as Java: Use `super` to refer to the parent class, and then provide the method name. The following Android method (written in Scala) demonstrates how to call a method named `onCreate` that’s defined in the `Activity` parent class:

```
class WelcomeActivity extends Activity {  
  override def onCreate(bundle: Bundle) {  
    super.onCreate(bundle)  
    // more code here ...  
  }  
}
```

As with Java, you can call multiple superclass methods if necessary:

```
class FourLeggedAnimal {  
  def walk { println("I'm walking") }  
  def run { println("I'm running") }  
}
```

```
class Dog extends FourLeggedAnimal {
  def walkThenRun {
    super.walk
    super.run
  }
}
```

Running this code in the Scala REPL yields:

```
scala> val suka = new Dog
suka: Dog = Dog@239bf795

scala> suka.walkThenRun
I'm walking
I'm running
```

Controlling which trait you call a method from

If your class inherits from multiple traits, and those traits implement the same method, you can select not only a method name, but also a trait name when invoking a method using `super`. For instance, given this class hierarchy:

```
trait Human {
  def hello = "the Human trait"
}

trait Mother extends Human {
  override def hello = "Mother"
}

trait Father extends Human {
  override def hello = "Father"
}
```

The following code shows different ways to invoke the `hello` method from the traits the `Child` class inherits from. This example shows that by mixing in the `Human`, `Mother`, and `Father` traits, you can call `super.hello`, or be more specific by calling `super[Mother].hello`, `super[Father].hello`, or `super[Human].hello`:

```
class Child extends Human with Mother with Father {
  def printSuper = super.hello
  def printMother = super[Mother].hello
  def printFather = super[Father].hello
  def printHuman = super[Human].hello
}
```

If you construct a test object to run this code:

```
object Test extends App {
  val c = new Child
  println(s"c.printSuper = ${c.printSuper}")
  println(s"c.printMother = ${c.printMother}")
}
```

```

    println(s"c.printFather = ${c.printFather}")
    println(s"c.printHuman = ${c.printHuman}")
  }

```

you can see the output:

```

c.printSuper = Father
c.printMother = Mother
c.printFather = Father
c.printHuman = the Human trait

```

As shown, when a class inherits from multiple traits, and those traits have a common method name, you can choose which trait to run the method from with the `super[traitName].methodName` syntax.

Note that when using this technique, you can't continue to reach up through the parent class hierarchy unless you directly extend the target class or trait using the `extends` or `with` keywords. For instance, the following code won't compile because `Dog` doesn't directly extend the `Animal` trait:

```

trait Animal {
  def walk { println("Animal is walking") }
}

class FourLeggedAnimal extends Animal {
  override def walk { println("I'm walking on all fours") }
}

class Dog extends FourLeggedAnimal {
  def walkThenRun {
    super.walk           // works
    super[FourLeggedAnimal].walk // works
    super[Animal].walk   // error: won't compile
  }
}

```

If you attempt to compile the code, you'll get the error, "Animal does not name a parent class of class Dog." You can get around that error by adding `with Animal` to your class declaration (but whether or not that's really a good idea is another story):

```

class Dog extends FourLeggedAnimal with Animal {

```

5.3. Setting Default Values for Method Parameters

Problem

You want to set default values for method parameters so the method can optionally be called without those parameters having to be assigned.

Solution

Specify the default value for parameters in the method signature. In the following code, the `timeout` field is assigned a default value of 5000, and the `protocol` field is given a default value of "http":

```
class Connection {  
  def makeConnection(timeout: Int = 5000, protocol: = "http") {  
    println("timeout = %d, protocol = %s".format(timeout, protocol))  
    // more code here  
  }  
}
```

This method can now be called in the following ways:

```
c.makeConnection()  
c.makeConnection(2000)  
c.makeConnection(3000, "https")
```

The results are demonstrated in the REPL:

```
scala> val c = new Connection  
c: Connection = Connection@385db088  
  
scala> c.makeConnection()  
timeout = 5000, protocol = http  
  
scala> c.makeConnection(2000)  
timeout = 2000, protocol = http  
  
scala> c.makeConnection(3000, "https")  
timeout = 3000, protocol = https
```

Note that empty parentheses are used in the first example. Attempting to call this method without parentheses results in an error:

```
scala> c.makeConnection  
<console>:10: error: missing arguments for method makeConnection in Connection;  
follow this method with `_' to treat it as a partially applied function  
      c.makeConnection  
        ^
```

The reason for this error is discussed in [Recipe 9.6, “Using Partially Applied Functions”](#).

If you like to call methods with the names of the method parameters, the method `makeConnection` can also be called in these ways:

```
c.makeConnection(timeout=10000)  
c.makeConnection(protocol="https")  
c.makeConnection(timeout=10000, protocol="https")
```

Discussion

Just as with constructor parameters, you can provide default values for method arguments. Because you have provided defaults, the consumer of your method can either supply an argument to override the default or skip the argument, letting it use its default value.

Arguments are assigned from left to right, so the following call assigns no arguments and uses the default values for both `timeout` and `protocol`:

```
c.makeConnection()
```

This call sets `timeout` to 2000 and leaves `protocol` to its default:

```
c.makeConnection(2000)
```

This call sets both the `timeout` and `protocol`:

```
c.makeConnection(3000, "https")
```

Note that you can't set the `protocol` only with this approach, but as shown in the Solution, you can use a named parameter:

```
c.makeConnection(protocol="https")
```

If your method provides a mix of some fields that offer default values and others that don't, list the fields that have default values last. To demonstrate the problem, the following example assigns a default value to the first argument and does not assign a default to the second argument:

```
class Connection {  
  // intentional error  
  def makeConnection(timeout: Int = 5000, protocol: String) {  
    println("timeout = %d, protocol = %s".format(timeout, protocol))  
    // more code here  
  }  
}
```

This code compiles, but you won't be able to take advantage of the default, as shown in the REPL errors:

```
scala> c.makeConnection(1000)  
<console>:10: error: not enough arguments for method makeConnection:  
(timeout: Int, protocol: String)Unit.  
Unspecified value parameter protocol.  
      c.makeConnection(1000)  
                        ^
```

```
scala> c.makeConnection("https")  
<console>:10: error: not enough arguments for method makeConnection:  
(timeout: Int, protocol: String)Unit.  
Unspecified value parameter protocol.
```

```
c.makeConnection("https")
      ^
```

By changing the method so the first field doesn't have a default and the last field does, the default method call can now be used:

```
class Connection {
  // corrected implementation
  def makeConnection(timeout: Int, protocol: String = "http") {
    println("timeout = %d, protocol = %s".format(timeout, protocol))
    // more code here
  }
}

scala> c.makeConnection(1000)
timeout = 1000, protocol = http

scala> c.makeConnection(1000, "https")
timeout = 1000, protocol = https
```

5.4. Using Parameter Names When Calling a Method

Problem

You prefer a coding style where you specify the method parameter names when calling a method.

Solution

The general syntax for calling a method with named parameters is this:

```
methodName(param1=value1, param2=value2, ...)
```

This is demonstrated in the following example.

Given this definition of a Pizza class:

```
class Pizza {
  var crustSize = 12
  var crustType = "Thin"
  def update(crustSize: Int, crustType: String) {
    this.crustSize = crustSize
    this.crustType = crustType
  }
  override def toString = {
    "A %d inch %s crust pizza.".format(crustSize, crustType)
  }
}
```

you can create a Pizza:

```
val p = new Pizza
```

You can then update the `Pizza`, specifying the field names and corresponding values when you call the `update` method:

```
p.update(crustSize = 16, crustType = "Thick")
```

This approach has the added benefit that you can place the fields in any order:

```
p.update(crustType = "Pan", crustSize = 14)
```

Discussion

You can confirm that this example works by running it in the Scala REPL:

```
scala> val p = new Pizza
p: Pizza = A 12 inch Thin crust pizza.

scala> p.updatePizza(crustSize = 16, crustType = "Thick")

scala> println(p)
A 16 inch Thick crust pizza.

scala> p.updatePizza(crustType = "Pan", crustSize = 14)

scala> println(p)
A 14 inch Pan crust pizza.
```

The ability to use named parameters when calling a method is available in other languages, including Objective-C. Although this approach is more verbose, it can also be more readable.

This technique is especially useful when several parameters have the same type, such as having several `Boolean` or `String` parameters in a method. For instance, compare this method call:

```
engage(true, true, true, false)
```

to this one:

```
engage(speedIsSet = true,
       directionIsSet = true,
       picardSaidMakeItSo = true,
       turnedOffParkingBrake = false)
```

When a method specifies default values for its parameters, as demonstrated in [Recipe 5.3](#), you can use this approach to specify only the parameters you want to override.

For instance, the `scala.xml.Utility` object has a method named `serialize` that takes seven parameters. However, default values are defined for each parameter in the method declaration, so if you need to change only one parameter, such as whether you want comments stripped from the output, you need to specify only that one parameter, in addition to your XML node:

```
Utility.serialize(myNode, stripComments = true)
```

The combination of these two recipes makes for a powerful approach.

5.5. Defining a Method That Returns Multiple Items (Tuples)

Problem

You want to return multiple values from a method, but don't want to wrap those values in a makeshift class.

Solution

Although you can return objects from methods just as in other OOP languages, Scala also lets you return multiple values from a method using *tuples*. First, define a method that returns a tuple:

```
def getStockInfo = {  
  // other code here ...  
  ("NFLX", 100.00, 101.00) // this is a Tuple3  
}
```

Then call that method, assigning variable names to the expected return values:

```
val (symbol, currentPrice, bidPrice) = getStockInfo
```

Running this example in the REPL demonstrates how this works:

```
scala> val (symbol, currentPrice, bidPrice) = getStockInfo  
symbol: java.lang.String = NFLX  
currentPrice: Double = 100.0  
bidPrice: Double = 101.0
```

Discussion

In Java, when it would be convenient to be able to return multiple values from a method, the typical workaround is to return those values in a one-off “wrapper” class. For instance, you might create a temporary wrapper class like this:

```
// java  
public class StockInfo {  
  String symbol;  
  double currentPrice;  
  double bidPrice;  
  
  public StockInfo(String symbol, double currentPrice, double bidPrice) {  
    this.symbol = symbol;  
    this.currentPrice = currentPrice;  
    this.bidPrice = bidPrice;  
  }  
}
```

```
}  
}
```

Then you could return an instance of this class from a method, like this:

```
return new StockInfo("NFLX", 100.00, 101.00);
```

In Scala you don't need to create a wrapper like this; you can just return the data as a tuple.

Working with tuples

In the example shown in the Solution, the `getStockInfo` method returned a tuple with three elements, so it is a `Tuple3`. Tuples can contain up to 22 variables and are implemented as `Tuple1` through `Tuple22` classes. As a practical matter, you don't have to think about those specific classes; just create a new tuple by enclosing elements inside parentheses, as shown.

To demonstrate a `Tuple2`, if you wanted to return only two elements from a method, just put two elements in the parentheses:

```
def getStockInfo = ("NFLX", 100.00)  
  
val (symbol, currentPrice) = getStockInfo
```

If you don't want to assign variable names when calling the method, you can set a variable equal to the tuple the method returns, and then access the tuple values using the following tuple underscore syntax:

```
scala> val result = getStockInfo  
x: (java.lang.String, Double, Double) = (NFLX,100.0)  
  
scala> result._1  
res0: java.lang.String = NFLX  
  
scala> result._2  
res1: Double = 100.0
```

As shown, tuple values can be accessed by position as `result._1`, `result._2`, and so on. Though this approach can be useful in some situations, your code will generally be clearer if you assign variable names to the values:

```
val (symbol, currentPrice) = getStockInfo
```

See Also

- [The Tuple3 class](#)
- [Recipe 10.27, “Tuples, for When You Just Need a Bag of Things”](#) for more tuple examples

5.6. Forcing Callers to Leave Parentheses off Accessor Methods

Problem

You want to enforce a coding style where getter/accessor methods can't have parentheses when they are invoked.

Solution

Define your getter/accessor method without parentheses after the method name:

```
class Pizza {  
  // no parentheses after crustSize  
  def crustSize = 12  
}
```

This forces consumers of your class to call `crustSize` without parentheses:

```
scala> val p = new Pizza  
p: Pizza = Pizza@3a3e8692  
  
// this fails because of the parentheses  
scala> p.crustSize()  
<console>:10: error: Int does not take parameters  
      p.crustSize()  
                ^  
  
// this works  
scala> p.crustSize  
res0: Int = 12
```



Coming from a Java background, I originally named this method `getCrustSize`, but the Scala convention is to drop “get” from methods like this, hence the method name `crustSize`.

Discussion

The recommended strategy for calling getter methods that have no side effects is to leave the parentheses off when calling the method. As stated in the [Scala Style Guide](#):

Methods which act as accessors of any sort ... should be declared *without* parentheses, except if they have side effects.

According to the style guide, because a simple accessor method like `crustSize` does not have side effects, it should not be called with parentheses, and this recipe demonstrates how to enforce this convention.

Although this recipe shows how to force callers to leave parentheses off methods when calling simple getters, there is no way to force them to use parentheses for side-effecting methods. This is only a convention, albeit a convention that I like and use these days. Although it's usually obvious that a method named `printStuff` is probably going to print some output, a little warning light goes off in my head when I see it called as `printStuff()` instead.

Side Effects

It's said that a purely functional program has no side effects. So what is a side effect?

According to Wikipedia, a function is said to have a side effect “if, in addition to returning a value, it also modifies some state or has an observable interaction with calling functions or the outside world.”

Side effects include things like:

- Writing or printing output.
- Reading input.
- Mutating the state of a variable that was given as input, changing data in a data structure, or modifying the value of a field in an object.
- Throwing an exception, or stopping the application when an error occurs.
- Calling other functions that have side effects.

In theory, pure functions are much easier to test. Imagine writing an addition function, such as `+`. Given the two numbers 1 and 2, the result will always be 3. A pure function like this is a simple matter of (a) immutable data coming in, and (b) a result coming out; nothing else happens. Because a function like this has no side effects, it's simple to test.

See [Recipe 20.1, “Create Methods with No Side Effects \(Pure Functions\)”](#), for more details on writing pure functions. Also, see the Wikipedia discussion on side effects in functional programming (FP) applications for [more details and examples](#).

See Also

The [Scala Style Guide on naming conventions and parentheses](#)

5.7. Creating Methods That Take Variable-Argument Fields

Problem

To make a method more flexible, you want to define a method parameter that can take a variable number of arguments, i.e., a `varargs` field.

Solution

Define a *varargs* field in your method declaration by adding a `*` character after the field type:

```
def printAll(strings: String*) {  
    strings.foreach(println)  
}
```

Given that method declaration, the `printAll` method can be called with zero or more parameters:

```
// these all work  
printAll()  
printAll("foo")  
printAll("foo", "bar")  
printAll("foo", "bar", "baz")
```

Use `_*` to adapt a sequence

As shown in the following example, you can use Scala's `_*` operator to adapt a sequence (`Array`, `List`, `Seq`, `Vector`, etc.) so it can be used as an argument for a `varargs` field:

```
// a sequence of strings  
val fruits = List("apple", "banana", "cherry")  
  
// pass the sequence to the varargs field  
printAll(fruits: _*)
```

If you come from a Unix background, it may be helpful to think of `_*` as a “splat” operator. This operator tells the compiler to pass each element of the sequence to `printAll` as a separate argument, instead of passing `fruits` as a single argument.

Discussion

When declaring that a method has a field that can contain a variable number of arguments, the `varargs` field must be the last field in the method signature. Attempting to define a field in a method signature *after* a `varargs` field is an error:

```
scala> def printAll(strings: String*, i: Int) {  
    |   strings.foreach(println)  
    | }
```

```
<console>:7: error: *-parameter must come last
    def printAll(strings: String*, i: Int) {
          ^
```

As an implication of that rule, a method can have only one varargs field.

As demonstrated in the Solution, if a field is a varargs field, you don't have to supply any arguments for it. For instance, in a method that has only one varargs field, you can call it with no arguments:

```
scala> def printAll(numbers: Int*) {
      |   numbers.foreach(println)
      | }
printAll: (numbers: Int*)Unit

scala> printAll()
```

This case reveals some of the inner workings of how Scala handles varargs fields. By defining a varargs method that can take multiple integers, and then calling that method (a) with arguments, and (b) without arguments, you can see how Scala handles the two situations:

```
def printAll(numbers: Int*) {
  println(numbers.getClass)
}

scala> printAll(1, 2, 3)
class scala.collection.mutable.WrappedArray$ofInt

scala> printAll()
class scala.collection.immutable.Nil$
```

While the first situation reveals how Scala handles the normal “one or more arguments” situation, treating the “no args” situation as a `Nil$` in the second situation keeps your code from throwing a `NullPointerException`.

Although the resulting types are different, as a practical matter, this isn't too important. You'll typically use a loop inside a method to handle a varargs field, and either of the following examples work fine whether the method is called with zero or multiple parameters:

```
// version 1
def printAll(numbers: Int*) {
  numbers.foreach(println)
}

// version 2
def printAll(numbers: Int*) {
  for (i <- numbers) println
}
```

5.8. Declaring That a Method Can Throw an Exception

Problem

You want to declare that a method can throw an exception, either to alert callers to this fact or because your method will be called from Java code.

Solution

Use the `@throws` annotation to declare the exception(s) that can be thrown. To declare that one exception can be thrown, place the annotation just before the method signature:

```
@throws(classOf[Exception])
override def play {
  // exception throwing code here ...
}
```

To indicate that a method can throw multiple exceptions, list them all before the method signature:

```
@throws(classOf[IOException])
@throws(classOf[LineUnavailableException])
@throws(classOf[UnsupportedAudioFileException])
def playSoundFileWithJavaAudio {
  // exception throwing code here ...
}
```

Discussion

The two examples shown are from an open source project I created that lets developers play WAV, AIFF, MP3, and other types of sound files. I declared that these two methods can throw exceptions for two reasons. First, whether the consumers are using Scala or Java, if they're writing robust code, they'll want to know that something failed.

Second, if they're using Java, the `@throws` annotation is the Scala way of providing the throws method signature to Java consumers. It's equivalent to declaring that a method throws an exception with this Java syntax:

```
public void play() throws FooException {
  // code here ...
}
```

It's important to note that Scala's philosophy regarding checked exceptions is different than Java's. Scala doesn't require that methods declare that exceptions can be thrown, and it also doesn't require calling methods to catch them. This is easily demonstrated in the REPL:

```
// 1) it's not necessary to state that a method throws an exception
scala> def boom {
    |   throw new Exception
    | }
boom: Unit

// 2) it's not necessary to wrap 'boom' in a try/catch block, but ...
scala> boom
java.lang.Exception
    at .boom(<console>:8)
    // much more exception output here ...
```

Although Scala doesn't require that exceptions are checked, if you fail to test for them, they'll blow up your code just like they do in Java. In the following example, the second `println` statement is never reached because the `boom` method throws its exception:

```
object BoomTest extends App {

    def boom { throw new Exception }

    println("Before boom")

    boom

    // this line is never reached
    println("After boom")

}
```

Java Exception Types

As a quick review, Java has (a) checked exceptions, (b) descendants of `Error`, and (c) descendants of `RuntimeException`. Like checked exceptions, `Error` and `RuntimeException` have many subclasses, such as `RuntimeException`'s famous offspring, `NullPointerException`.

According to the Java documentation for [the `Exception` class](#), “The class `Exception` and any subclasses that are not also subclasses of `RuntimeException` are checked exceptions. Checked exceptions need to be declared in a method or constructor's throws clause if they can be thrown by the execution of the method or constructor and propagate outside the method or constructor boundary.”

The following links provide more information on Java exceptions and exception handling:

- [The Three Kinds of \(Java\) Exceptions](#)
- [Unchecked Exceptions—The Controversy](#)
- [Wikipedia discussion of checked exceptions](#)

- [Java tutorial on exception handling](#)
- [Java Exception class](#)

See Also

[Recipe 17.2, “Add Exception Annotations to Scala Methods to Work with Java”](#), for other examples of adding exception annotations to methods

5.9. Supporting a Fluent Style of Programming

Problem

You want to create an API so developers can write code in a **fluent** programming style, also known as **method chaining**.

Solution

A fluent style of programming lets users of your API write code by chaining method calls together, as in this example:

```
person.setFirstName("Leonard")
      .setLastName("Nimoy")
      .setAge(82)
      .setCity("Los Angeles")
      .setState("California")
```

To support this style of programming:

- If your class can be extended, specify `this.type` as the return type of fluent style methods.
- If you're sure that your class won't be extended, you can optionally return `this` from your fluent style methods.

The following code demonstrates how to specify `this.type` as the return type of the `set*` methods:

```
class Person {

  protected var fname = ""
  protected var lname = ""

  def setFirstName(firstName: String): this.type = {
    fname = firstName
    this
  }
}
```

```

    def setLastName(lastName: String): this.type = {
        lname = lastName
        this
    }
}

class Employee extends Person {

    protected var role = ""

    def setRole(role: String): this.type = {
        this.role = role
        this
    }

    override def toString = {
        "%s, %s, %s".format(fname, lname, role)
    }

}

```

The following test object demonstrates how these methods can be chained together:

```

object Main extends App {

    val employee = new Employee

    // use the fluent methods
    employee.setFirstName("Al")
              .setLastName("Alexander")
              .setRole("Developer")
    println(employee)

}

```

Discussion

If you're sure your class won't be extended, specifying `this.type` as the return type of your `set*` methods isn't necessary; you can just return the `this` reference at the end of each fluent style method. This is shown in the `addTopping`, `setCrustSize`, and `setCrustType` methods of the following `Pizza` class, which is declared to be `final`:

```

final class Pizza {

    import scala.collection.mutable.ArrayBuffer

    private val toppings = ArrayBuffer[String]()
    private var crustSize = 0
    private var crustType = ""

    def addTopping(topping: String) = {
        toppings += topping
    }
}

```

```

    this
  }

  def setCrustSize(crustSize: Int) = {
    this.crustSize = crustSize
    this
  }

  def setCrustType(crustType: String) = {
    this.crustType = crustType
    this
  }

  def print() {
    println(s"crust size: $crustSize")
    println(s"crust type: $crustType")
    println(s"toppings: $toppings")
  }
}

```

This class is demonstrated with the following driver program:

```

object FluentPizzaTest extends App {

  val p = new Pizza
  p.setCrustSize(14)
  .setCrustType("thin")
  .addTopping("cheese")
  .addTopping("green olives")
  .print()

}

```

This results in the following output:

```

crust size: 14
crust type: thin
toppings:  ArrayBuffer(cheese, green olives)

```

Returning `this` in your methods works fine if you're sure your class won't be extended, but if your class can be extended—as in the first example where the `Employee` class extended the `Person` class—explicitly setting `this.type` as the return type of your `set*` methods ensures that the fluent style will continue to work in your subclasses. In this example, this makes sure that methods like `setFirstName` on an `Employee` object return an `Employee` reference and not a `Person` reference.

See Also

- [Definition of a fluent interface](#)
- [Method chaining](#)
- Martin Fowler's [discussion of a fluent interface](#)

Introduction

The word “object” has a dual meaning in Scala. As with Java, you use it to refer to an instance of a class, but in Scala, `object` is also a keyword.

The first three recipes in this chapter look at an object as an instance of a class, show how to cast objects from one type to another, demonstrate the Scala equivalent of Java’s `.class` approach, and show how to determine the class of an object.

The remaining recipes demonstrate how the `object` keyword is used for other purposes. You’ll see how to use it to launch Scala applications and to create Singletons. There’s also a special type of object known as a *package object*. Using a package object is entirely optional, but it provides a nice little out-of-the-way place where you can put code that’s common to all classes and objects in a particular package level in your application. For instance, Scala’s root-level package object contains many lines of code like this:

```
type Throwable = java.lang.Throwable
type Exception = java.lang.Exception
type Error = java.lang.Error

type Seq[+A] = scala.collection.Seq[A]
val Seq = scala.collection.Seq
```

Declaring those type definitions in Scala’s root package object helps to make the rest of the code a little bit cleaner, and also keeps these definitions from cluttering up other files.

You’ll also see how to create a *companion object* to solve several problems. For instance, one use of a companion object is to create the equivalent of Java’s static members. You can also use a companion object so consumers of its corresponding class won’t need to use the `new` keyword to create an instance of the class. For example, notice how the `new` keyword isn’t required before each `Person` instance in this code:

```
val siblings = List(Person("Kim"), Person("Julia"), Person("Kenny"))
```

These solutions, and a few more, are presented in this chapter.

6.1. Object Casting

Problem

You need to cast an instance of a class from one type to another, such as when creating objects dynamically.

Solution

Use the `asInstanceOf` method to cast an instance to the desired type. In the following example, the object returned by the `lookup` method is cast to an instance of a class named `Recognizer`:

```
val recognizer = cm.lookup("recognizer").asInstanceOf[Recognizer]
```

This Scala code is equivalent to the following Java code:

```
Recognizer recognizer = (Recognizer)cm.lookup("recognizer");
```

The `asInstanceOf` method is defined in the Scala `Any` class and is therefore available on all objects.

Discussion

In dynamic programming, it's often necessary to cast from one type to another. This approach is needed when using the Spring Framework and instantiating beans from an application context file:

```
// open/read the application context file
val ctx = new ClassPathXmlApplicationContext("applicationContext.xml")

// instantiate our dog and cat objects from the application context
val dog = ctx.getBean("dog").asInstanceOf[Animal]
val cat = ctx.getBean("cat").asInstanceOf[Animal]
```

It's used when reading a YAML configuration file:

```
val yaml = new Yaml(new Constructor(classOf[EmailAccount]))
val emailAccount = yaml.load(text).asInstanceOf[EmailAccount]
```

The example shown in the Solution comes from code I wrote to work with an open source Java speech recognition library named Sphinx-4. With this library, many properties are defined in an XML file, and then you create recognizer and microphone objects dynamically. In a manner similar to Spring, this requires reading an XML configuration file, then casting instances to the specific types you want:

```

val cm = new ConfigurationManager("config.xml")

// instance of Recognizer
val recognizer = cm.lookup("recognizer").asInstanceOf[Recognizer]

// instance of Microphone
val microphone = cm.lookup("microphone").asInstanceOf[Microphone]

```

The `asInstanceOf` method isn't limited to only these situations. You can use it to cast numeric types:

```

scala> val a = 10
a: Int = 10

scala> val b = a.asInstanceOf[Long]
b: Long = 10

scala> val c = a.asInstanceOf[Byte]
c: Byte = 10

```

It can be used in more complicated code, such as when you need to interact with Java and send it an array of Object instances:

```

val objects = Array("a", 1)
val arrayOfObject = objects.asInstanceOf[Array[Object]]
AJavaClass.sendObjects(arrayOfObject)

```

It's demonstrated in [Chapter 15](#) like this:

```

import java.net.{URL, HttpURLConnection}
val connection = (new URL(url)).openConnection.asInstanceOf[HttpURLConnection]

```

Be aware that as with Java, this type of coding can lead to a `ClassCastException`, as demonstrated in this REPL example:

```

scala> val i = 1
i: Int = 1

scala> i.asInstanceOf[String]
ClassCastException: java.lang.Integer cannot be cast to java.lang.String

```

As usual, use a `try/catch` expression to handle this situation.

See Also

- [Recipe 2.2, “Converting Between Numeric Types \(Casting\)”](#), for more numeric type casting recipes
- [The Any class](#)
- [The Sphinx-4 project](#)

6.2. The Scala Equivalent of Java's .class

Problem

When an API requires that you pass in a `Class`, you'd call `.class` on an object in Java, but that doesn't work in Scala.

Solution

Use the Scala `classOf` method instead of Java's `.class`. The following example shows how to pass a class of type `TargetDataLine` to a method named `DataLine.Info`:

```
val info = new DataLine.Info(classOf[TargetDataLine], null)
```

By contrast, the same method call would be made like this in Java:

```
// java
info = new DataLine.Info(TargetDataLine.class, null);
```

The `classOf` method is defined in the Scala `Predef` object and is therefore available in all classes without requiring an import.

Discussion

This approach also lets you begin with simple reflection techniques. The following REPL example demonstrates how to access the methods of the `String` class:

```
scala> val stringClass = classOf[String]
stringClass: Class[String] = class java.lang.String

scala> stringClass.getMethods
res0: Array[java.lang.reflect.Method] = Array(public boolean
java.lang.String.equals(java.lang.Object), public java.lang.String
(output goes on for a while ...))
```

See Also

- [Oracle's "Retrieving Class Objects" document](#)
- The [Scala Predef](#) object

6.3. Determining the Class of an Object

Problem

Because you don't have to explicitly declare types with Scala, you may occasionally want to print the class/type of an object to understand how Scala works, or to debug code.

Solution

When you want to learn about the types Scala is automatically assigning on your behalf, call the `getClass` method on the object.

For instance, when I was first trying to understand how varargs fields work, I called `getClass` on a method argument, and found that the class my method was receiving varied depending on the situation. Here's the method declaration:

```
def printAll(numbers: Int*) {  
  println("class: " + numbers.getClass)  
}
```

Calling the `printAll` method with and without arguments demonstrates the two classes Scala assigns to the `numbers` field under the different conditions:

```
scala> printAll(1, 2, 3)  
class scala.collection.mutable.WrappedArray$ofInt  
  
scala> printAll()  
class scala.collection.immutable.Nil$
```

This technique can be very useful when working with something like Scala's XML library, so you can understand which classes you're working with in different situations. For instance, the following example shows that the `<p>` tag contains one child element, which is of class `scala.xml.Text`:

```
scala> val hello = <p>Hello, world</p>  
hello: scala.xml.Elem = <p>Hello, world</p>  
  
scala> hello.child.foreach(e => println(e.getClass))  
class scala.xml.Text
```

However, by adding a `<br />` tag inside the `<p>` tags, there are now three child elements of two different types:

```
scala> val hello = <p>Hello, <br/>world</p>  
hello: scala.xml.Elem = <p>Hello, <br/>world</p>  
  
scala> hello.child.foreach(e => println(e.getClass))  
class scala.xml.Text  
class scala.xml.Elem  
class scala.xml.Text
```

When you can't see information like this in your IDE, using this `getClass` approach is very helpful.

Discussion

When I can't see object types in an IDE, I write little tests like this in the REPL. The usual pattern is to call `getClass` on the object of interest, passing in different parameters to see how things work:

```
scala> def printClass(c: Any) { println(c.getClass) }
printClass: (c: Any)Unit

scala> printClass(1)
class java.lang.Integer

scala> printClass("yo")
class java.lang.String
```

In the first example shown in the Solution, the types Scala assigns to the number parameter don't matter too much; it was more a matter of curiosity about how things work. The actual method looks like the following code, and for my purposes, the only important thing is that each class Scala uses supports a `foreach` method:

```
def printAll(numbers: Int*) {
  numbers.foreach(println)
}
```

As desired, this method can be called with and without parameters:

```
scala> printAll(1,2,3)
1
2
3

scala> printAll()
(no output)
```

6.4. Launching an Application with an Object

Problem

You want to start an application with a `main` method, or provide the entry point for a script.

Solution

There are two ways to create a launching point for your application: define an object that extends the `App` trait, or define an object with a properly defined `main` method.

For the first solution, define an object that extends the `App` trait. Using this approach, the following code creates a simple but complete Scala application:

```
object Hello extends App {
  println("Hello, world")
}
```

The code in the body of the object is automatically run, just as if it were inside a `main` method.

Just save that code to a file named *Hello.scala*, compile it with `scalac`, and then run it with `scala`, like this:

```
$ scalac Hello.scala

$ scala Hello
Hello, world
```

When using this approach, any command-line arguments to your application are implicitly available through an `args` object, which is inherited from the `App` trait. The `args` object is an instance of `Array[String]`, just as if you had declared a `main` method yourself. The following code demonstrates how to use the `args` object:

```
object Hello extends App {
  if (args.length == 1)
    println(s"Hello, ${args(0)}")
  else
    println("I didn't get your name.")
}
```

After it's been compiled, this program yields the following results:

```
$ scala Hello
I didn't get your name.

$ scala Hello Joe
Hello, Joe
```

The second approach to launching an application is to manually implement a `main` method with the correct signature in an object, in a manner similar to Java:

```
object Hello2 {
  def main(args: Array[String]) {
    println("Hello, world")
  }
}
```

This is also a simple but complete application.

Discussion

Note that in both cases, Scala applications are launched from an *object*, not a class.

I tend to use the `App` trait for both scripts and larger applications, but you can use either approach. I recommend reviewing the source code for the `App` trait to better understand what it performs. The source code is available from the URL in the See Also section.

The Scaladoc for the `App` trait currently includes two caveats:

1. It should be noted that this trait is implemented using the `DelayedInit` functionality, which means that fields of the object will not have been initialized before the `main` method has been executed.
2. It should also be noted that the `main` method will not normally need to be overridden: the purpose is to turn the whole class body into the “main method.” You should only choose to override it if you know what you are doing.

See the Scaladoc for the `App` and `DelayedInit` traits for more information.

See Also

- [The `App` trait](#).
- [The `DelayedInit` trait](#).
- The shell script examples in [Chapter 14](#) demonstrate more examples of the `App` trait.

6.5. Creating Singletons with object

Problem

You want to create a Singleton object to ensure that only one instance of a class exists.

Solution

Create Singleton objects in Scala with the `object` keyword. For instance, you might create a Singleton object to represent something like a keyboard, mouse, or perhaps a cash register in a pizza restaurant:

```
object CashRegister {  
  def open { println("opened") }  
  def close { println("closed") }  
}
```

With `CashRegister` defined as an object, there can be only one instance of it, and its methods are called just like static methods on a Java class:

```
object Main extends App {  
  CashRegister.open  
  CashRegister.close  
}
```


This pattern is also common when creating utility methods, such as this `DateUtils` object:

```
import java.util.Calendar
import java.text.SimpleDateFormat

object DateUtils {

  // as "Thursday, November 29"
  def getCurrentDate: String = getCurrentDateTime("EEEE, MMMM d")

  // as "6:20 p.m."
  def getCurrentTime: String = getCurrentDateTime("K:m aa")

  // a common function used by other date/time functions
  private def getCurrentDateTime(dateTimeFormat: String): String = {
    val dateFormat = new SimpleDateFormat(dateTimeFormat)
    val cal = Calendar.getInstance()
    dateFormat.format(cal.getTime())
  }
}
```

Because these methods are defined in an object instead of a class, they can be called in the same way as a static method in Java:

```
scala> DateUtils.getCurrentTime
res0: String = 10:13 AM

scala> DateUtils.getCurrentDate
res1: String = Friday, July 6
```

Singleton objects also make great reusable messages when using actors. If you have a number of actors that can all receive start and stop messages, you can create Singletons like this:

```
case object StartMessage
case object StopMessage
```

You can then use those objects as messages that can be sent to actors:

```
inputValve ! StopMessage
outputValve ! StopMessage
```

See [Chapter 13, *Actors and Concurrency*](#), for more examples of this approach.

Discussion

In addition to creating objects in this manner, you can give the appearance that a class has both static and nonstatic methods using an approach known as a “companion object.” See the following recipe for examples of that approach.

6.6. Creating Static Members with Companion Objects

Problem

You want to create a class that has instance methods and static methods, but unlike Java, Scala does not have a `static` keyword.

Solution

Define nonstatic (instance) members in your *class*, and define members that you want to appear as “static” members in an *object* that has the same name as the class, and is in the same file as the class. This object is known as a *companion object*.

Using this approach lets you create what appear to be static members on a class (both fields and methods), as shown in this example:

```
// Pizza class
class Pizza (var crustType: String) {
  override def toString = "Crust type is " + crustType
}

// companion object
object Pizza {
  val CRUST_TYPE_THIN = "thin"
  val CRUST_TYPE_THICK = "thick"
  def getFoo = "Foo"
}
```

With the `Pizza` class and `Pizza` object defined in the same file (presumably named `Pizza.scala`), members of the `Pizza` object can be accessed just as static members of a Java class:

```
println(Pizza.CRUST_TYPE_THIN)
println(Pizza.getFoo)
```

You can also create a new `Pizza` instance and use it as usual:

```
var p = new Pizza(Pizza.CRUST_TYPE_THICK)
println(p)
```



If you're coming to Scala from a language other than Java, "static" methods in Java are methods that can be called directly on a class, without requiring an instance of the class. For instance, here's an example of a method named `increment` in a Scala object named `StringUtils`:

```
object StringUtils {  
  def increment(s: String) = s.map(c => (c + 1).toChar)  
}
```

Because it's defined inside an object (not a class), the `increment` method can be called directly on the `StringUtils` object, without requiring an instance of `StringUtils` to be created:

```
scala> StringUtils.increment("HAL")  
res0: String = IBM
```

In fact, when an object is defined like this without a corresponding class, you *can't* create an instance of it. This line of code won't compile:

```
val utils = new StringUtils
```

Discussion

Although this approach is different than Java, the recipe is straightforward:

- Define your class and object in the same file, giving them the same name.
- Define members that should appear to be "static" in the object.
- Define nonstatic (instance) members in the class.

Accessing private members

It's also important to know that a class and its companion object can access each other's private members. In the following code, the "static" method `double` in the object can access the private variable `secret` of the class `Foo`:

```
class Foo {  
  private val secret = 2  
}  
  
object Foo {  
  // access the private class field 'secret'  
  def double(foo: Foo) = foo.secret * 2  
}  
  
object Driver extends App {  
  val f = new Foo  
  println(Foo.double(f)) // prints 4  
}
```

Similarly, in the following code, the instance member `printObj` can access the private field `obj` of the object `Foo`:

```
class Foo {  
  // access the private object field 'obj'  
  def printObj { println(s"I can see ${Foo.obj}") }  
}  
  
object Foo {  
  private val obj = "Foo's object"  
}  
  
object Driver extends App {  
  val f = new Foo  
  f.printObj  
}
```

6.7. Putting Common Code in Package Objects

Problem

You want to make functions, fields, and other code available at a package level, without requiring a class or object.

Solution

Put the code you want to make available to all classes within a package in a *package object*.

By convention, put your code in a file named *package.scala* in the directory where you want your code to be available. For instance, if you want your code to be available to all classes in the `com.alvinalexander.myapp.model` package, create a file named *package.scala* in the `com/alvinalexander/myapp/model` directory of your project.

In the *package.scala* source code, remove the word `model` from the end of the package statement, and use that name to declare the name of the package object. Including a blank line, the first three lines of your file will look like this:

```
package com.alvinalexander.myapp  
  
package object model {
```

Now write the rest of your code as you normally would. The following example shows how to create a field, method, enumeration, and type definition in your package object:

```
package com.alvinalexander.myapp  
  
package object model {  
  
  // field  
  val MAGIC_NUM = 42
```

```

// method
def echo(a: Any) { println(a) }

// enumeration
object Margin extends Enumeration {
  type Margin = Value
  val TOP, BOTTOM, LEFT, RIGHT = Value
}

// type definition
type MutableMap[K, V] = scala.collection.mutable.Map[K, V]
val MutableMap = scala.collection.mutable.Map
}

```

You can now access this code directly from within other classes, traits, and objects in the package `com.alvinalexander.myapp.model` as shown here:

```

package com.alvinalexander.myapp.model

object MainDriver extends App {

  // access our method, constant, and enumeration
  echo("Hello, world")
  echo(MAGIC_NUM)
  echo(Margin.LEFT)

  // use our MutableMap type (scala.collection.mutable.Map)
  val mm = MutableMap("name" -> "Al")
  mm += ("password" -> "123")
  for ((k,v) <- mm) printf("key: %s, value: %s\n", k, v)
}

```

Discussion

The most confusing part about package objects is where to put them, along with what their package and object names should be.

Where to put them isn't too hard; by convention, create a file named *package.scala* in the directory where you want your code to be available. In the example shown, I want the package code to be available in the `com.alvinalexander.myapp.model` package, so I put the file *package.scala* in the *com/alvinalexander/myapp/model* source code directory:

```

+-- com
  +-- alvinalexander
    +-- myapp
      +-- model
        +-- package.scala

```

In regards to the first few lines of the *package.scala* source code, simply start with the usual name of the package:

```
package com.alvinalexander.myapp.model
```

Then take the name of the last package level (*model*) off that statement, leaving you with this:

```
package com.alvinalexander.myapp
```

Then use that name (*model*) as the name of your package object:

```
package object model {
```

As shown earlier, the first several lines of your *package.scala* file will look like this:

```
package com.alvinalexander.myapp
```

```
package object model {
```

The [Scala package object documentation](#) states, “Any kind of definition that you can put inside a class, you can also put at the top level of a package.” In my experience, package objects are a great place to put methods and functions that are common to the package, as well as constants, enumerations, and implicit conversions.

As described in [the second page of the Scala package object documentation](#), “The standard Scala package also has its package object. Because `scala._` is automatically imported into every Scala file, the definitions of this object are available without prefix.” If you create something like a `StringBuilder` or `Range`, you’re using this code.

See Also

Scala’s root package object is full of type aliases, like these:

```
type Throwable = java.lang.Throwable
type Exception = java.lang.Exception
type Error     = java.lang.Error

type RuntimeException = java.lang.RuntimeException
type NullPointerException = java.lang.NullPointerException
type ClassCastException = java.lang.ClassCastException
```

Like the `Predef` object, its source code is worth looking at if you want to know more about how Scala works. You can find its source by following the “source” link on its [Scaladoc page](#).

- [An introduction to package objects](#)
- [The Scala package object](#)

6.8. Creating Object Instances Without Using the new Keyword

Problem

You've seen that Scala code looks cleaner when you don't always have to use the `new` keyword to create a new instance of a class, like this:

```
val a = Array(Person("John"), Person("Paul"))
```

So you want to know how to write your code to make your classes work like this.

Solution

There are two ways to do this:

- Create a companion object for your class, and define an `apply` method in the companion object with the desired constructor signature.
- Define your class as a *case class*.

You'll look at both approaches next.

Creating a companion object with an apply method

To demonstrate the first approach, define a `Person` class and `Person` object in the same file. Define an `apply` method in the object that takes the desired parameters. This method is essentially the constructor of your class:

```
class Person {  
  var name: String = _  
}  
  
object Person {  
  def apply(name: String): Person = {  
    var p = new Person  
    p.name = name  
    p  
  }  
}
```

Given this definition, you can create new `Person` instances without using the `new` keyword, as shown in these examples:

```
val dawn = Person("Dawn")  
val a = Array(Person("Dan"), Person("Elijah"))
```

The `apply` method in a companion object is treated specially by the Scala compiler and lets you create new instances of your class without requiring the `new` keyword. (More on this in the Discussion.)

Declare your class as a case class

The second solution to the problem is to declare your class as a *case class*, defining it with the desired constructor:

```
case class Person (var name: String)
```

This approach also lets you create new class instances without requiring the `new` keyword:

```
val p = Person("Fred Flinstone")
```

With case classes, this works because the case class generates an `apply` method in a companion object for you. However, it's important to know that a case class creates *much* more code for you than just the `apply` method. This is discussed in depth in the Discussion.

Discussion

An `apply` method defined in the companion object of a class is treated specially by the Scala compiler. There is essentially a little syntactic sugar baked into Scala that converts this code:

```
val p = Person("Fred Flinstone")
```

into this code:

```
val p = Person.apply("Fred Flinstone")
```

The `apply` method is basically a factory method, and Scala's little bit of syntactic sugar lets you use the syntax shown, creating new class instances without using the `new` keyword.

Providing multiple constructors with additional apply methods

To create multiple constructors when manually defining your own `apply` method, just define multiple `apply` methods in the companion object that provide the constructor signatures you want:

```
class Person {  
  var name = ""  
  var age = 0  
}  
  
object Person {  
  // a one-arg constructor  
  def apply(name: String): Person = {  
    var p = new Person  
    p.name = name  
    p  
  }  
}
```



```
// a two-arg constructor
def apply(name: String, age: Int): Person = {
  var p = new Person
  p.name = name
  p.age = age
  p
}
}
```

You can now create a new `Person` instance in these ways:

```
val fred = Person("Fred")
val john = Person("John", 42)
```

I’m using the term “constructor” loosely here, but each `apply` method does define a different way to construct an instance.

Providing multiple constructors for case classes

To provide multiple constructors for a case class, it’s important to know what the case class declaration actually does.

If you look at the code the Scala compiler generates for the case class example, you’ll see that see it creates two output files, *Person\$.class* and *Person.class*. If you disassemble *Person\$.class* with the `javap` command, you’ll see that it contains an `apply` method, along with many others:

```
$ javap Person$
Compiled from "Person.scala"
public final class Person$ extends scala.runtime.AbstractFunction1
implements scala.ScalaObject,scala.Serializable{
    public static final Person$ MODULE$;
    public static {};
    public final java.lang.String toString();
    public scala.Option unapply(Person);
    public Person apply(java.lang.String); // the apply method (returns a Person)
    public java.lang.Object readResolve();
    public java.lang.Object apply(java.lang.Object);
}
```

You can also disassemble *Person.class* to see what it contains. For a simple class like this, it contains an additional 20 methods; this hidden bloat is one reason some developers don’t like case classes.



See [Recipe 4.14, “Generating Boilerplate Code with Case Classes”](#), for a thorough discussion of what code is generated for case classes, and why.

Note that the `apply` method in the disassembled code accepts one `String` argument:

```
public Person apply(java.lang.String);
```

That `String` corresponds to the `name` field in your case class constructor:

```
case class Person (var name: String)
```

So, it's important to know that when a case class is created, it writes the accessor and (optional) mutator methods *only* for the default constructor. As a result, (a) it's best to define all class parameters in the default constructor, and (b) write `apply` methods for the auxiliary constructors you want.

This is demonstrated in the following code, which I place in a file named *Person.scala*:

```
// want accessor and mutator methods for the name and age fields
case class Person (var name: String, var age: Int)

// define two auxiliary constructors
object Person {

  def apply() = new Person("<no name>", 0)
  def apply(name: String) = new Person(name, 0)

}
```

Because `name` and `age` are declared as `var` fields, accessor and mutator methods will both be generated. Also, two `apply` methods are declared in the object: a no-args constructor, and a one-arg constructor.

As a result, you can create instances of your class in three different ways, as demonstrated in the following code:

```
object Test extends App {

  val a = Person()
  val b = Person("Al")
  val c = Person("William Shatner", 82)

  println(a)
  println(b)
  println(c)

  // test the mutator methods
  a.name = "Leonard Nimoy"
  a.age = 82
  println(a)

}
```

Running this test object results in the following output:

```
Person(<no name>,0)
Person(Al,0)
Person(William Shatner,82)
Person(Leonard Nimoy,82)
```

For more information on case classes, see [Recipe 4.14, “Generating Boilerplate Code with Case Classes”](#).

6.9. Implement the Factory Method in Scala with apply

Problem

To let subclasses declare which type of object should be created, and to keep the object creation point in one location, you want to implement the factory method in Scala.

Solution

One approach to this problem is to take advantage of how a Scala companion object’s `apply` method works. Rather than creating a “get” method for your factory, you can place the factory’s decision-making algorithm in the `apply` method.

For instance, suppose you want to create an `Animal` factory that returns instances of `Cat` and `Dog` classes, based on what you ask for. By writing an `apply` method in the companion object of an `Animal` class, users of your factory can create new `Cat` and `Dog` instances like this:

```
val cat = Animal("cat") // creates a Cat
val dog = Animal("dog") // creates a Dog
```

To implement this behavior, create a parent `Animal` trait:

```
trait Animal {
  def speak
}
```

In the same file, create (a) a companion object, (b) the classes that extend the base trait, and (c) a suitable `apply` method:

```
object Animal {

  private class Dog extends Animal {
    override def speak { println("woof") }
  }

  private class Cat extends Animal {
    override def speak { println("meow") }
  }

  // the factory method
  def apply(s: String): Animal = {
    if (s == "dog") new Dog
    else new Cat
  }
}
```

This lets you run the desired code:

```
val cat = Animal("cat") // returns a Cat
val dog = Animal("dog") // returns a Dog
```

You can test this by pasting the `Animal` trait and object into the REPL, and then issuing these statements:

```
scala> val cat = Animal("cat")
cat: Animal = Animal$Cat@486f8860
```

```
scala> cat.speak
meow
```

```
scala> val dog = Animal("dog")
dog: Animal = Animal$Dog@412798c1
```

```
scala> dog.speak
woof
```

As you can see, this approach works as desired.

Discussion

You have a variety of ways to implement this solution, so experiment with different approaches, in particular how you want to make the `Cat` and `Dog` classes accessible. The idea of the factory method is to make sure that concrete instances can only be created through the factory; therefore, the class constructors should be hidden from all other classes. The code here shows one possible solution to this problem.

If you don't like using the `apply` method as the factory interface, you can create the usual “get” method in the companion object, as shown in the `getAnimal` method here:

```
// an alternative factory method (use one or the other)
def getAnimal(s: String): Animal = {
  if (s == "dog") return new Dog
  else return new Cat
}
```

Using this method instead of the `apply` method, you now create new `Animal` instances like this:

```
val cat = Animal.getAnimal("cat") // returns a Cat
val dog = Animal.getAnimal("dog") // returns a Dog
```

Either approach is fine; consider this recipe as a springboard for your own solution.

See Also

[Recipe 6.8](#) for more examples of implementing the `apply` method

Packaging and Imports

Introduction

Scala's packaging approach is similar to Java, but it's more flexible. In addition to using the `package` statement at the top of a class file, you can use a curly brace packaging style, similar to C++ and C# namespaces.

The Scala approach to importing members is also similar to Java, and more flexible. With Scala you can:

- Place import statements anywhere
- Import classes, packages, or objects
- Hide and rename members when you import them

All of these approaches are demonstrated in this chapter.

It's helpful to know that in Scala, two packages are implicitly imported for you:

- `java.lang._`
- `scala._`

In Scala, the `_` character is similar to the `*` character in Java, so these statements refer to every member in those packages.

In addition to those packages, all members from the `scala.Predef` object are imported into your applications implicitly.

A great suggestion from the book *Beginning Scala* by David Pollak (Apress), is to dig into the source code of the **Predef object**. The code isn't too long, and it demonstrates many of the features of the Scala language. Many implicit conversions are brought into

scope by the `Predef` object, as well as methods like `println`, `readLine`, `assert`, and `require`.

7.1. Packaging with the Curly Braces Style Notation

Problem

You want to use a nested style package notation, similar to the namespace notation in C++ and C#.

Solution

Wrap one or more classes in a set of curly braces with a package name, as shown in this example:

```
package com.acme.store {  
  class Foo { override def toString = "I am com.acme.store.Foo" }  
}
```

The canonical name of the class is `com.acme.store.Foo`. It's just as though you declared the code like this:

```
package com.acme.store  
  
class Foo { override def toString = "I am com.acme.store.Foo" }
```

With this approach, you can place multiple packages in one file. You can also nest packages using this “curly braces” style.

The following example creates three `Foo` classes, all of which are in different packages, to demonstrate how to include one package inside another:

```
// a package containing a class named Foo  
package orderentry {  
  class Foo { override def toString = "I am orderentry.Foo" }  
}  
  
// one package nested inside the other  
package customers {  
  class Foo { override def toString = "I am customers.Foo" }  
  
  package database {  
    // this Foo is different than customers.Foo or orderentry.Foo  
    class Foo { override def toString = "I am customers.database.Foo" }  
  }  
}  
  
// a simple object to test the packages and classes  
object PackageTests extends App {  
  println(new orderentry.Foo)
```

```
println(new customers.Foo)
println(new customers.database.Foo)
}
```

If you place this code in a file, and then compile and run it, you'll get the following output:

```
I am orderentry.Foo
I am customers.Foo
I am customers.database.Foo
```

This demonstrates that each Foo class is indeed in a different package.

As shown in the first example, package names don't have to be limited to just one level. You can define multiple levels of depth at one time:

```
package com.alvinalexander.foo {
  class Foo { override def toString = "I am com.alvinalexander.foo.Foo" }
}
```

Discussion

You can create Scala packages with the usual Java practice of declaring a package name at the top of the file:

```
package foo.bar.baz

class Foo {
  override def toString = "I'm foo.bar.baz.Foo"
}
```

In most cases, I use this packaging approach, but because Scala code can be much more concise than Java, the alternative curly brace packaging syntax can be very convenient when you want to declare multiple classes and packages in one file.

7.2. Importing One or More Members

Problem

You want to import one or more members into the scope of your current program.

Solution

This is the syntax for importing one class:

```
import java.io.File
```

You can import multiple classes the Java way:

```
import java.io.File
import java.io.IOException
import java.io.FileNotFoundException
```

Or you can import several classes the Scala way:

```
import java.io.{File, IOException, FileNotFoundException}
```

Use the following syntax to import everything from the `java.io` package:

```
import java.io._
```

The `_` character in this example is similar to the `*` wildcard character in Java. If the `_` character feels unusual, it helps to know that it's used consistently throughout the Scala language as a wildcard character, and that consistency is very nice.

Discussion

The concept of importing code into the current scope is similar between Java and Scala, but Scala is more flexible. Scala lets you:

- Place import statements anywhere, including the top of a class, within a class or object, within a method, or within a block of code
- Import classes, packages, or objects
- Hide and rename members when you import them

Syntactically, the two big differences are the curly brace syntax, known as the *import selector clause*, and the use of the `_` wildcard character instead of Java's `*` wildcard. The advantages of the import selector clause are demonstrated further in Recipes [7.3](#) and [7.4](#).

Placing import statements anywhere

In Scala you can place an import statement anywhere. For instance, because Scala makes it easy to include multiple classes in the same file, you may want to separate your import statements so the common imports are declared at the top of the file, and the imports specific to each class are within each class specification:

```
package foo

import java.io.File
import java.io.PrintWriter

class Foo {
  import javax.swing.JFrame // only visible in this class
  // ...
}

class Bar {
  import scala.util.Random // only visible in this class
```



```
// ...  
}
```

You can also place import statements inside methods, functions, or blocks:

```
class Bar {  
  def doBar = {  
    import scala.util.Random  
    println("")  
  }  
}
```

See [Recipe 7.6, “Using Import Statements Anywhere”](#), for more examples and details about the use of import statements.

7.3. Renaming Members on Import

Problem

You want to rename members when you import them to help avoid namespace collisions or confusion.

Solution

Give the class you’re importing a new name when you import it with this syntax:

```
import java.util.{ArrayList => JavaList}
```

Then, within your code, refer to the class by the alias you’ve given it:

```
val list = new JavaList[String]
```

You can also rename multiple classes at one time during the import process:

```
import java.util.{Date => JDate, HashMap => JHashMap}
```

Because you’ve created these aliases during the import process, the original (real) name of the class can’t be used in your code. For instance, in the last example, the following code will fail because the compiler can’t find the `java.util.HashMap` class:

```
// error: this won't compile because HashMap was renamed  
// during the import process  
val map = new HashMap[String, String]
```

Discussion

As shown, you can create a new name for a class when you import it, and can then refer to it by the new name, or alias. The book *Programming in Scala*, by Odersky, et al (Artima). The book refers to this as a *renaming clause*.

This can be very helpful when trying to avoid namespace collisions and confusion. Class names like `Listener`, `Message`, `Handler`, `Client`, `Server`, and many more are all very common, and it can be helpful to give them an alias when you import them.

From a strategy perspective, you can either rename all classes that might be conflicting or confusing:

```
import java.util.{HashMap => JavaHashMap}
import scala.collection.mutable.{Map => ScalaMutableMap}
```

or you can just rename one class to clarify the situation:

```
import java.util.{HashMap => JavaHashMap}
import scala.collection.mutable.Map
```

As an interesting combination of several recipes, not only can you rename classes on import, but you can even rename class members. As an example of this, in shell scripts I tend to rename the `println` method to a shorter name, as shown here in the REPL:

```
scala> import System.out.{println => p}
import System.out.{println=>p}

scala> p("hello")
hello
```

7.4. Hiding a Class During the Import Process

Problem

You want to hide one or more classes while importing other members from the same package.

Solution

To hide a class during the import process, use the renaming syntax shown in [Recipe 7.3, “Renaming Members on Import”](#), but point the class name to the `_` wildcard character. The following example hides the `Random` class, while importing everything else from the `java.util` package:

```
import java.util.{Random => _, _}
```

This can be confirmed in the REPL:

```
scala> import java.util.{Random => _, _}
import java.util.{Random=>_, _}

// can't access Random
scala> val r = new Random
<console>:10: error: not found: type Random
      val r = new Random
                  ^
```

```
// can access other members
scala> new ArrayList
res0: java.util.ArrayList[Nothing] = []
```

In that example, the following portion of the code is what “hides” the Random class:

```
import java.util.{Random => _}
```

The second `_` character inside the curly braces is the same as stating that you want to import everything else in the package, like this:

```
import java.util._
```

Note that the `_` import wildcard must be in the last position. It yields an error if you attempt to use it in other positions:

```
scala> import java.util.{_, Random => _}
<console>:1: error: Wildcard import must be in last position
import java.util.{_, Random => _}
                  ^
```

This is because you may want to hide multiple members during the import process, and to do, so you need to list them first.

To hide multiple members, list them before using the final wildcard import:

```
scala> import java.util.{List => _, Map => _, Set => _, _}
import java.util.{List=>_, Map=>_, Set=>_, _}

scala> new ArrayList
res0: java.util.ArrayList[Nothing] = []
```

This ability to hide members on import is useful when you need many members from one package, and therefore want to use the `_` wildcard syntax, but you also want to hide one or more members during the import process, typically due to naming conflicts.

7.5. Using Static Imports

Problem

You want to import members in a way similar to the Java static import approach, so you can refer to the member names directly, without having to prefix them with their class name.

Solution

Use this syntax to import all members of the Java Math class:

```
import java.lang.Math._
```

You can now access these members without having to precede them with the class name:

```
scala> import java.lang.Math._
import java.lang.Math._

scala> val a = sin(0)
a: Double = 0.0

scala> val a = cos(PI)
a: Double = -1.0
```

The Java Color class also demonstrates the usefulness of this technique:

```
scala> import java.awt.Color._
import java.awt.Color._

scala> println(RED)
java.awt.Color[r=255,g=0,b=0]

scala> val currentColor = BLUE
currentColor: java.awt.Color = java.awt.Color[r=0,g=0,b=255]
```

Enumerations are another great candidate for this technique. Given a Java enum like this:

```
package com.alvinalexander.dates;

public enum Day {
    SUNDAY, MONDAY, TUESDAY, WEDNESDAY, THURSDAY, FRIDAY, SATURDAY
}
```

you can import and use this enumeration in a Scala program like this:

```
import com.alvinalexander.dates.Day._

// somewhere after the import statement
if (date == SUNDAY || date == SATURDAY) println("It's the weekend.")
```

Discussion

Although some developers don't like static imports, I find that this approach makes enums more readable. Just specifying the name of a class or enum before the constant makes the code less readable:

```
if (date == Day.SUNDAY || date == Day.SATURDAY) {
    println("It's the weekend.")
}
```

With the static import approach there's no need for the leading "Day." in the code, and it's easier to read.

7.6. Using Import Statements Anywhere

Problem

You want to use an import statement anywhere, generally to limit the scope of the import, to make the code more clear, or to organize your code.

Solution

You can place an import statement almost anywhere inside a program. As with Java, you can import members at the top of a class definition, and then use the imported resource later in your code:

```
package foo

import scala.util.Random

class ImportTests {
  def printRandom {
    val r = new Random
  }
}
```

You can import members inside a class:

```
package foo

class ImportTests {
  import scala.util.Random
  def printRandom {
    val r = new Random
  }
}
```

This limits the scope of the import to the code in the class that comes after the import statement.

You can limit the scope of an import to a method:

```
def getRandomWaitTimeInMinutes: Int = {
  import com.alvinalexander.pandorasbox._
  val p = new Pandora
  p.release
}
```

You can even place an import statement inside a block, limiting the scope of the import to only the code that follows the statement, inside that block. In the following example, the field `r1` is declared correctly, because it's within the block and after the import statement, but the declaration for field `r2` won't compile, because the `Random` class is not in scope at that point:

```
def printRandom {
  {
    import scala.util.Random
    val r1 = new Random // this is fine
  }
  val r2 = new Random // error: not found: type Random
}
```

Discussion

Import statements are read in the order of the file, so where you place them in a file also limits their scope. The following code won't compile because I attempt to reference the `Random` class before the import statement is declared:

```
// this doesn't work because the import is after the attempted reference
class ImportTests {
  def printRandom {
    val r = new Random // fails
  }
}
import scala.util.Random
```

When you want to include multiple classes and packages in one file, you can combine import statements and the curly brace packaging approach to limit the scope of the import statements, as shown in these examples:

```
package orderentry {
  import foo._
  // more code here ...
}

package customers {
  import bar._
  // more code here ...

  package database {
    import baz._
    // more code here ...
  }
}
```

In this example, members can be accessed as follows:

- Code in the `orderentry` package can access members of `foo`, but can't access members of `bar` or `baz`.
- Code in `customers` and `customers.database` can't access members of `foo`.
- Code in `customers` can access members of `bar`.
- Code in `customers.database` can access members in `bar` and `baz`.

The same concept applies when defining multiple classes in one file:

```
package foo

// available to all classes defined below
import java.io.File
import java.io.PrintWriter

class Foo {
    // only available inside this class
    import javax.swing.JFrame
    // ...
}

class Bar {
    // only available inside this class
    import scala.util.Random
    // ...
}
```

Although placing import statements at the top of a file or just before they're used can be a matter of style, I find this flexibility to be useful when placing multiple classes or packages in one file. In these cases, it's nice to keep the imports in a small scope to limit namespace issues, and also to make the code easier to refactor as it grows.

Introduction

In its most basic use, a Scala trait is just like a Java interface. When you're faced with situations where you would have used an interface in Java, just think "trait" in Scala.

Just as Java classes can implement multiple interfaces, Scala classes can extend multiple traits. As you'll see in the recipes in this chapter, this is done with the `extends` and `with` keywords, so when a class (or object) extends multiple traits, you'll see code like this:

```
class Woodpecker extends Bird with TreeScaling with Pecking
```

However, using traits as interfaces only scratches the surface of what they can do. Traits have much more power than Java interfaces because, just like abstract methods in Java, they can also have implemented methods. However, unlike Java's abstract classes, you can mix more than one trait into a class, and a trait can also control what classes it can be mixed into.

This chapter provides examples of the many uses of Scala traits.

8.1. Using a Trait as an Interface

Problem

You're used to creating interfaces in other languages like Java and want to create something like that in Scala.

Solution

You can use a trait just like a Java interface. As with interfaces, just declare the methods in your trait that you want extending classes to implement:

```

trait BaseSoundPlayer {
  def play
  def close
  def pause
  def stop
  def resume
}

```

If the methods don't take any argument, you only need to declare the names of the methods after the `def` keyword, as shown. If a method should require parameters, list them as usual:

```

trait Dog {
  def speak(whatToSay: String)
  def wagTail(enabled: Boolean)
}

```

When a class extends a trait, it uses the `extends` and `with` keywords. When extending one trait, use `extends`:

```

class Mp3SoundPlayer extends BaseSoundPlayer { ...

```

When extending a class and one or more traits, use `extends` for the class, and `with` for subsequent traits:

```

class Foo extends BaseClass with Trait1 with Trait2 { ...

```

When a class extends multiple traits, use `extends` for the first trait, and `with` for subsequent traits:

```

class Foo extends Trait1 with Trait2 with Trait3 with Trait4 { ...

```

Unless the class implementing a trait is abstract, it must implement all of the abstract trait methods:

```

class Mp3SoundPlayer extends BaseSoundPlayer {
  def play { // code here ... }
  def close { // code here ... }
  def pause { // code here ... }
  def stop { // code here ... }
  def resume { // code here ... }
}

```

If a class extends a trait but does not implement the abstract methods defined in that trait, it must be declared abstract:

```

// must be declared abstract because it does not implement
// all of the BaseSoundPlayer methods
abstract class SimpleSoundPlayer extends BaseSoundPlayer {
  def play { ... }
  def close { ... }
}

```

In other uses, one trait can extend another trait:

```

trait Mp3BaseSoundFilePlayer extends BaseSoundFilePlayer {
  def getBasicPlayer: BasicPlayer
  def getBasicController: BasicController
  def setGain(volume: Double)
}

```

Discussion

As demonstrated, at their most basic level, traits can be used just like Java interfaces. In your trait, just declare the methods that need to be implemented by classes that want to extend your trait.

Classes extend your trait using either the `extends` or `with` keywords, according to these simple rules:

- If a class extends one trait, use the `extends` keyword.
- If a class extends multiple traits, use `extends` for the first trait and `with` to extend (mix in) the other traits.
- If a class extends a class (or abstract class) and a trait, always use `extends` before the class name, and use `with` before the trait name(s).

You can also use fields in your traits. See the next recipe for examples.

As shown in the `WaggingTail` trait in the following example, not only can a trait be used like a Java interface, but it can also provide method implementations, like an abstract class in Java:

```

abstract class Animal {
  def speak
}

trait WaggingTail {
  def startTail { println("tail started") }
  def stopTail { println("tail stopped") }
}

trait FourLeggedAnimal {
  def walk
  def run
}

class Dog extends Animal with WaggingTail with FourLeggedAnimal {
  // implementation code here ...
  def speak { println("Dog says 'woof'") }
  def walk { println("Dog is walking") }
  def run { println("Dog is running") }
}

```

This ability is discussed in detail in [Recipe 8.3, “Using a Trait Like an Abstract Class”](#).

When a class has multiple traits, such as the `WaggingTail` and `FourLeggedAnimal` traits in this example, those traits are said to be *mixed in* to the class. The term “mixed in” is also used when extending a single object instance with a trait, like this:

```
val f = new Foo with Trait1
```

This feature is discussed more in [Recipe 8.8, “Adding a Trait to an Object Instance”](#).

8.2. Using Abstract and Concrete Fields in Traits

Problem

You want to put abstract or concrete fields in your traits so they are declared in one place and available to all types that implement the trait.

Solution

Define a field with an initial value to make it *concrete*; otherwise, don’t assign it an initial value to make it *abstract*. This trait shows several examples of abstract and concrete fields with `var` and `val` types:

```
trait PizzaTrait {  
  var numToppings: Int    // abstract  
  var size = 14           // concrete  
  val maxNumToppings = 10 // concrete  
}
```

In the class that extends the trait, you’ll need to define the values for the abstract fields, or make the class abstract. The following `Pizza` class demonstrates how to set the values for the `numToppings` and `size` fields in a concrete class:

```
class Pizza extends PizzaTrait {  
  var numToppings = 0    // 'override' not needed  
  size = 16              // 'var' and 'override' not needed  
}
```

Discussion

As shown in the example, fields of a trait can be declared as either `var` or `val`. You don’t need to use the `override` keyword to override a `var` field in a subclass (or trait), but you do need to use it to override a `val` field:

```
trait PizzaTrait {  
  val maxNumToppings: Int  
}  
  
class Pizza extends PizzaTrait {  
  override val maxNumToppings = 10 // 'override' is required  
}
```

Overriding `var` and `val` fields is discussed more in [Recipe 4.13, “Defining Properties in an Abstract Base Class \(or Trait\)”](#).

8.3. Using a Trait Like an Abstract Class

Problem

You want to use a trait as something like an abstract class in Java.

Solution

Define methods in your trait just like regular Scala methods. In the class that extends the trait, you can override those methods or use them as they are defined in the trait.

In the following example, an implementation is provided for the `speak` method in the `Pet` trait, so implementing classes don't have to override it. The `Dog` class chooses not to override it, whereas the `Cat` class does:

```
trait Pet {  
  def speak { println("Yo") } // concrete implementation  
  def comeToMaster           // abstract method  
}  
  
class Dog extends Pet {  
  // don't need to implement 'speak' if you don't need to  
  def comeToMaster { ("I'm coming!") }  
}  
  
class Cat extends Pet {  
  // override the speak method  
  override def speak { ("meow") }  
  def comeToMaster { ("That's not gonna happen.") }  
}
```

If a class extends a trait without implementing its abstract methods, it must be defined as abstract. Because `FlyingPet` does not implement `comeToMaster`, it must be declared as abstract:

```
abstract class FlyingPet extends Pet {  
  def fly { ("I'm flying!") }  
}
```

Discussion

Although Scala has abstract classes, it's much more common to use traits than abstract classes to implement base behavior. A class can extend only one abstract class, but it can implement multiple traits, so using traits is more flexible.

See Also

- Like Java, you use `super.foo` to call a method named `foo` in an immediate superclass. When a class mixes in multiple traits—and those traits implement a method declared by a common ancestor—you can be more specific, and specify which trait you'd like to invoke a method on. See [Recipe 5.2, “Calling a Method on a Super-class”](#), for more information.
- See [Recipe 4.12, “When to Use an Abstract Class”](#), for information on when to use an abstract class instead of a trait. (Spoiler: Use an abstract class (a) when you want to define a base behavior, and that behavior requires a constructor with parameters, and (b) in some situations when you need to interact with Java.)

8.4. Using Traits as Simple Mixins

Problem

You want to design a solution where multiple traits can be mixed into a class to provide a robust design.

Solution

To implement a simple *mixin*, define the methods you want in your trait, then add the trait to your class using `extends` or `with`. For instance, the following code defines a `Tail` trait:

```
trait Tail {  
  def wagTail { println("tail is wagging") }  
  def stopTail { println("tail is stopped") }  
}
```

You can use this trait with an abstract `Pet` class to create a `Dog`:

```
abstract class Pet (var name: String) {  
  def speak // abstract  
  def ownerIsHome { println("excited") }  
  def jumpForJoy { println("jumping for joy") }  
}  
  
class Dog (name: String) extends Pet (name) with Tail {  
  def speak { println("woof") }  
  override def ownerIsHome {  
    wagTail  
    speak  
  }  
}
```

The `Dog` class extends the abstract class `Pet` and mixes in the `Tail` trait, and can use the methods defined by both `Pet` and `Tail`:

```
object Test extends App {  
  val zeus = new Dog("Zeus")  
  zeus.ownerIsHome  
  zeus.jumpForJoy  
}
```

In summary, the `Dog` class gets behavior from both the abstract `Pet` class and the `Tail` trait; this is something you can't do in Java.



To see a great demonstration of the power of mixins, read Artima's short "[Stackable Trait Pattern](#)" article. By defining traits and classes as *base*, *core*, and *stackable* components, they demonstrate how sixteen different classes can be derived from three traits by "stacking" the traits together.

See Also

When you develop traits, you may want to limit the classes they can be mixed into. The classes a trait can be mixed into can be limited using the following techniques:

- [Recipe 8.5](#) shows how to limit which classes can use a trait by declaring inheritance.
- [Recipe 8.6](#) shows how to mark traits so they can only be used by subclasses of a certain type.
- [Recipe 8.7](#) demonstrates the technique to use to make sure a trait can only be mixed into classes that have a specific method.
- Also, see Artima's "[Stackable Trait Pattern](#)" article.

8.5. Limiting Which Classes Can Use a Trait by Inheritance

Problem

You want to limit a trait so it can only be added to classes that extend a superclass or another trait.

Solution

Use the following syntax to declare a trait named `TraitName`, where `TraitName` can only be mixed into classes that extend a type named `SuperThing`, where `SuperThing` may be a trait, class, or abstract class:

```
trait [TraitName] extends [SuperThing]
```

For instance, in the following example, `Starship` and `StarfleetWarpCore` both extend the common superclass `StarfleetComponent`, so the `StarfleetWarpCore` trait *can* be mixed into the `Starship` class:

```
class StarfleetComponent
trait StarfleetWarpCore extends StarfleetComponent
class Starship extends StarfleetComponent with StarfleetWarpCore
```

However, in the following example, the `Warbird` class *can't* extend the `StarfleetWarpCore` trait, because `Warbird` and `StarfleetWarpCore` don't share the same superclass:

```
class StarfleetComponent
trait StarfleetWarpCore extends StarfleetComponent
class RomulanStuff

// won't compile
class Warbird extends RomulanStuff with StarfleetWarpCore
```

Attempting to compile this second example yields this error:

```
error: illegal inheritance; superclass RomulanStuff
  is not a subclass of the superclass StarfleetComponent
  of the mixin trait StarfleetWarpCore
class Warbird extends RomulanStuff with StarfleetWarpCore
                                   ^
```

Discussion

A trait inheriting from a class is not a common occurrence, and in general, [Recipes 8.6](#) and [Recipe 8.7](#) are more commonly used to limit the classes a trait can be mixed into.

However, when this situation occurs, you can see how inheritance can be used. As long as a class and a trait share the same superclass (`Starship` and `StarfleetWarpCore` extend `StarfleetComponent`) the code will compile, but if the superclasses are different (`Warbird` and `StarfleetWarpCore` have different superclasses), the code will not compile.

As a second example, in modeling a large pizza store chain that has a corporate office and many small retail stores, the legal department creates a rule that people who deliver pizzas to customers must be a subclass of `StoreEmployee` and cannot be a subclass of `CorporateEmployee`. To enforce this, begin by defining your base classes:

```
abstract class Employee
class CorporateEmployee extends Employee
class StoreEmployee extends Employee
```

Someone who delivers food can only be a `StoreEmployee`, so you enforce this requirement in the `DeliversFood` trait using inheritance like this:


```
trait DeliversFood extends StoreEmployee
```

Now you can define a `DeliveryPerson` class like this:

```
// this is allowed
class DeliveryPerson extends StoreEmployee with DeliversFood
```

Because the `DeliversFood` trait can only be mixed into classes that extend `StoreEmployee`, the following line of code won't compile:

```
// won't compile
class Receptionist extends CorporateEmployee with DeliversFood
```

Discussion

It seems rare that a trait and a class the trait will be mixed into should both have the same superclass, so I suspect the need for this recipe is also rare. When you want to limit the classes a trait can be mixed into, don't create an artificial inheritance tree to use this recipe; use one of the following recipes instead.

See Also

- [Recipe 8.6](#) to see how to mark traits so they can only be used by subclasses of a certain type
- [Recipe 8.7](#) to make sure a trait can only be mixed into a class that has a specific method

8.6. Marking Traits So They Can Only Be Used by Subclasses of a Certain Type

Problem

You want to mark your trait so it can only be used by types that extend a given base type.

Solution

To make sure a trait named `MyTrait` can only be mixed into a class that is a subclass of a type named `BaseType`, begin your trait with a `this: BaseType =>` declaration, as shown here:

```
trait MyTrait {
  this: BaseType =>
```

For instance, to make sure a `StarFleetWarpCore` can only be used in a `Starship`, mark the `StarFleetWarpCore` trait like this:

```
trait StarfleetWarpCore {
  this: Starship =>
  // more code here ...
}
```

Given that declaration, this code will work:

```
class Starship
class Enterprise extends Starship with StarfleetWarpCore
```

But other attempts like this will fail:

```
class RomulanShip
// this won't compile
class Warbird extends RomulanShip with StarfleetWarpCore
```

This second example fails with an error message similar to this:

```
error: illegal inheritance;
self-type Warbird does not conform to StarfleetWarpCore's selftype
StarfleetWarpCore with Starship
class Warbird extends RomulanShip with StarfleetWarpCore
      ^
```

Discussion

As shown in the error message, this approach is referred to as a *self type*. The [Scala Glossary](#) includes this statement as part of its description of a self type:

“Any concrete class that mixes in the trait must ensure that its type conforms to the trait’s self type.”

A trait can also require that any type that wishes to extend it must extend multiple other types. The following `WarpCore` definition requires that any type that wishes to mix it in must extend `WarpCoreEjector` and `FireExtinguisher`, in addition to extending `Starship`:

```
trait WarpCore {
  this: Starship with WarpCoreEjector with FireExtinguisher =>
}
```

Because the following `Enterprise` definition matches that signature, this code compiles:

```
class Starship
trait WarpCoreEjector
trait FireExtinguisher

// this works
class Enterprise extends Starship
  with WarpCore
  with WarpCoreEjector
  with FireExtinguisher
```

However, if the `Enterprise` doesn't extend `Starship`, `WarpCoreEjector`, and `FireExtinguisher`, the code won't compile. Once again, the compiler shows that the self-type signature is not correct:

```
// won't compile
class Enterprise extends Starship with WarpCore with WarpCoreEjector

error: illegal inheritance;
self-type Enterprise does not conform to WarpCore's selftype WarpCore
with Starship with WarpCoreEjector with FireExtinguisher

class Enterprise extends Starship with WarpCore with WarpCoreEjector
                                     ^
```

See Also

- [Recipe 8.5](#) shows how to limit which classes can use a trait by declaring inheritance
- [Recipe 8.7](#) demonstrates the technique to use to make sure a trait can only be mixed into classes that have a specific method
- [The Scala Glossary](#)

8.7. Ensuring a Trait Can Only Be Added to a Type That Has a Specific Method

Problem

You only want to allow a trait to be mixed into a type (class, abstract class, or trait) that has a method with a given signature.

Solution

Use a variation of the self-type syntax that lets you declare that any class that attempts to mix in the trait must implement the method you specify.

In the following example, the `WarpCore` trait requires that any classes that attempt to mix it in must have an `ejectWarpCore` method:

```
trait WarpCore {
  this: { def ejectWarpCore(password: String): Boolean } =>
}
```

It further states that the `ejectWarpCore` method must accept a `String` argument and return a `Boolean` value.

The following definition of the Enterprise class meets these requirements, and will therefore compile:

```
class Starship {
  // code here ...
}

class Enterprise extends Starship with WarpCore {
  def ejectWarpCore(password: String): Boolean = {
    if (password == "password") {
      println("ejecting core")
      true
    } else {
      false
    }
  }
}
```

A trait can also require that a class have multiple methods. To require more than one method, just add the additional method signatures inside the block:

```
trait WarpCore {
  this: {
    def ejectWarpCore(password: String): Boolean
    def startWarpCore: Unit
  } =>
}

class Starship

class Enterprise extends Starship with WarpCore {
  def ejectWarpCore(password: String): Boolean = {
    if (password == "password") { println("core ejected"); true } else false
  }
  def startWarpCore { println("core started") }
}
```

Discussion

This approach is known as a *structural type*, because you're limiting what classes the trait can be mixed into by stating that the class must have a certain structure, i.e., the methods you've defined. In the examples shown, limits were placed on what classes the WarpCore trait can be mixed into.

See Also

- [Recipe 8.5](#) shows how to limit which classes can use a trait by declaring inheritance.
- [Recipe 8.6](#) shows how to mark traits so they can only be used by subclasses of a certain type.

8.8. Adding a Trait to an Object Instance

Problem

Rather than add a trait to an entire class, you just want to add a trait to an object instance when the object is created.

Solution

Add the trait to the object when you construct it. This is demonstrated in a simple example:

```
class DavidBanner

trait Angry {
  println("You won't like me ...")
}

object Test extends App {
  val hulk = new DavidBanner with Angry
}
```

When you compile and run this code, it will print, “You won’t like me ...”, because the hulk object is created when the DavidBanner class is instantiated with the Angry trait, which has the print statement shown in its constructor.

Discussion

As a more practical matter, you might mix in something like a debugger or logging trait when constructing an object to help debug that object:

```
trait Debugger {
  def log(message: String) {
    // do something with message
  }
}

// no debugger
val child = new Child

// debugger added as the object is created
val problemChild = new ProblemChild with Debugger
```

This makes the log method available to the problemChild instance.

8.9. Extending a Java Interface Like a Trait

Problem

You want to implement a Java interface in a Scala application.

Solution

In your Scala application, use the `extends` and `with` keywords to implement your Java interfaces, just as though they were Scala traits.

Given these three Java interfaces:

```
// java
public interface Animal {
    public void speak();
}

public interface Wagging {
    public void wag();
}

public interface Running {
    public void run();
}
```

you can create a `Dog` class in Scala with the usual `extends` and `with` keywords, just as though you were using traits:

```
// scala
class Dog extends Animal with Wagging with Running {
    def speak { println("Woof") }
    def wag { println("Tail is wagging!") }
    def run { println("I'm running!") }
}
```

The difference is that Java interfaces don't implement behavior, so if you're defining a class that extends a Java interface, you'll need to implement the methods, or declare the class abstract.

Functional Programming

Introduction

Scala is both an object-oriented programming (OOP) and a functional programming (FP) language. This chapter demonstrates functional programming techniques, including the ability to define functions and pass them around as instances. Just like you create a `String` instance in Java and pass it around, you can define a function as a variable and pass it around. I'll demonstrate many examples and advantages of this capability in this chapter.

As a language that supports functional programming, Scala encourages an **expression-oriented programming (EOP) model**. Simply put, in EOP, every statement (expression) yields a value. This paradigm can be as obvious as an `if/else` statement returning a value:

```
val greater = if (a > b) a else b
```

It can also be as surprising as a `try/catch` statement returning a value:

```
val result = try {  
  aString.toInt  
} catch {  
  case _ => 0  
}
```

Although EOP is casually demonstrated in many examples in this book, it's helpful to be consciously aware of this way of thinking in the recipes that follow.

9.1. Using Function Literals (Anonymous Functions)

Problem

You want to use an anonymous function—also known as a *function literal*—so you can pass it into a method that takes a function, or to assign it to a variable.

Solution

Given this List:

```
val x = List.range(1, 10)
```

you can pass an anonymous function to the List's `filter` method to create a new List that contains only even numbers:

```
val evens = x.filter((i: Int) => i % 2 == 0)
```

The REPL demonstrates that this expression indeed yields a new List of even numbers:

```
scala> val evens = x.filter((i: Int) => i % 2 == 0)
evens: List[Int] = List(2, 4, 6, 8)
```

In this solution, the following code is a function literal (also known as an *anonymous function*):

```
(i: Int) => i % 2 == 0
```

Although that code works, it shows the most explicit form for defining a function literal. Thanks to several Scala shortcuts, the expression can be simplified to this:

```
val evens = x.filter(_ % 2 == 0)
```

In the REPL, you see that this returns the same result:

```
scala> val evens = x.filter(_ % 2 == 0)
evens: List[Int] = List(2, 4, 6, 8)
```

Discussion

In this example, the original function literal consists of the following code:

```
(i: Int) => i % 2 == 0
```

When examining this code, it helps to think of the `=>` symbol as a *transformer*, because the expression transforms the parameter list on the left side of the symbol (an `Int` named `i`) into a new result using the algorithm on the right side of the symbol (in this case, an expression that results in a `Boolean`).

As mentioned, this example shows the long form for defining an anonymous function, which can be simplified in several different ways. The first example shows the most explicit form:


```
val evens = x.filter((i: Int) => i % 2 == 0)
```

Because the Scala compiler can infer from the expression that `i` is an `Int`, the `Int` declaration can be dropped off:

```
val evens = x.filter(i => i % 2 == 0)
```

Because Scala lets you use the `_` wildcard instead of a variable name when the parameter appears only once in your function, this code can be simplified even more:

```
val evens = x.filter(_ % 2 == 0)
```

In other examples, you can simplify your anonymous functions further. For instance, beginning with the most explicit form, you can print each element in the list using this anonymous function with the `foreach` method:

```
x.foreach((i: Int) => println(i))
```

As before, the `Int` declaration isn't required:

```
x.foreach((i) => println(i))
```

Because there is only one argument, the parentheses around the `i` parameter aren't needed:

```
x.foreach(i => println(i))
```

Because `i` is used only once in the body of the function, the expression can be further simplified with the `_` wildcard:

```
x.foreach(println(_))
```

Finally, if a function literal consists of one statement that takes a single argument, you need not explicitly name and specify the argument, so the statement can finally be reduced to this:

```
x.foreach(println)
```

9.2. Using Functions as Variables

Problem

You want to pass a function around like a variable, just like you pass `String`, `Int`, and other variables around in an object-oriented programming language.

Solution

Use the syntax shown in [Recipe 9.1](#) to define a function literal, and then assign that literal to a variable.

The following code defines a function literal that takes an `Int` parameter and returns a value that is twice the amount of the `Int` that is passed in:

```
(i: Int) => { i * 2 }
```

As mentioned in [Recipe 9.1](#), you can think of the `=>` symbol as a *transformer*. In this case, the function transforms the `Int` value `i` to an `Int` value that is twice the value of `i`.

You can now assign that function literal to a variable:

```
val double = (i: Int) => { i * 2 }
```

The variable `double` is an instance, just like an instance of a `String`, `Int`, or other type, but in this case, it's an instance of a function, known as a *function value*. You can now invoke `double` just like you'd call a method:

```
double(2) // 4  
double(3) // 6
```

Beyond just invoking `double` like this, you can also pass it to any method (or function) that takes a function parameter with its signature. For instance, because the `map` method of a sequence is a generic method that takes an input parameter of type `A` and returns a type `B`, you can pass the `double` method into the `map` method of an `Int` sequence:

```
scala> val list = List.range(1, 5)  
list: List[Int] = List(1, 2, 3, 4)
```

```
scala> list.map(double)  
res0: List[Int] = List(2, 4, 6, 8)
```

Welcome to the world of functional programming.

Discussion

You can declare a function literal in at least two different ways. I generally prefer the following approach, which implicitly infers that the following function's return type is `Boolean`:

```
val f = (i: Int) => { i % 2 == 0 }
```

In this case, the Scala compiler is smart enough to look at the body of the function and determine that it returns a `Boolean` value. As a human, it's also easy to look at the code on the right side of the expression and see that it returns a `Boolean`, so I usually leave the explicit `Boolean` return type off the function declaration.

However, if you prefer to explicitly declare the return type of a function literal, or want to do so because your function is more complex, the following examples show different forms you can use to explicitly declare that your function returns a `Boolean`:

```
val f: (Int) => Boolean = i => { i % 2 == 0 }  
val f: Int => Boolean = i => { i % 2 == 0 }  
val f: Int => Boolean = i => i % 2 == 0  
val f: Int => Boolean = _ % 2 == 0
```

A second example helps demonstrate the difference of these approaches. These functions all take two `Int` parameters and return a single `Int` value, which is the sum of the two input values:

```
// implicit approach
val add = (x: Int, y: Int) => { x + y }
val add = (x: Int, y: Int) => x + y

// explicit approach
val add: (Int, Int) => Int = (x,y) => { x + y }
val add: (Int, Int) => Int = (x,y) => x + y
```

As shown, the curly braces around the body of the function in these simple examples are optional, but they are required when the function body grows to more than one expression:

```
val addThenDouble: (Int, Int) => Int = (x,y) => {
  val a = x + y
  2 * a
}
```

Using a method like an anonymous function

Scala is very flexible, and just like you can define an anonymous function and assign it to a variable, you can also define a method and then pass it around like an instance variable. Again using a modulus example, you can define a method in any of these ways:

```
def modMethod(i: Int) = i % 2 == 0
def modMethod(i: Int) = { i % 2 == 0 }
def modMethod(i: Int): Boolean = i % 2 == 0
def modMethod(i: Int): Boolean = { i % 2 == 0 }
```

Any of these methods can be passed into collection methods that expect a function that has one `Int` parameter and returns a `Boolean`, such as the `filter` method of a `List[Int]`:

```
val list = List.range(1, 10)
list.filter(modMethod)
```

Here's what that looks like in the REPL:

```
scala> def modMethod(i: Int) = i % 2 == 0
modMethod: (i: Int)Boolean

scala> val list = List.range(1, 10)
list: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> list.filter(modMethod)
res0: List[Int] = List(2, 4, 6, 8)
```

As noted, this is similar to the process of defining a function literal and assigning it to a variable. The following function works just like the previous method:

```
val modFunction = (i: Int) => i % 2 == 0
list.filter(modFunction)
```

At a coding level, the obvious difference is that `modMethod` is a *method* defined in a class, whereas `modFunction` is a *function* that's assigned to a variable. Under the covers, `modFunction` is an instance of the **Function1 trait**, which defines a function that takes one argument. (The `scala` package defines other similar traits, including `Function0`, `Function2`, and so on, up to `Function22`.)

Assigning an existing function/method to a function variable

Continuing our exploration, you can assign an existing method or function to a function variable. For instance, you can create a new function named `c` from the `scala.math.cos` method using either of these approaches:

```
scala> val c = scala.math.cos _
c: Double => Double = <function1>
```

```
scala> val c = scala.math.cos(_)
c: Double => Double = <function1>
```

This is called a *partially applied function*. It's partially applied because the `cos` method requires one argument, which you have not yet supplied (more on this in [Recipe 9.6](#)).

Now that you have `c`, you can use it just like you would have used `cos`:

```
scala> c(0)
res0: Double = 1.0
```

If you're not familiar with this syntax, this is a place where the REPL can be invaluable. If you attempt to assign the `cos` function/method to a variable, the REPL tells you what's wrong:

```
scala> val c = scala.math.cos
<console>:11: error: missing arguments for method cos in class MathCommon;
follow this method with `_' to treat it as a partially applied function
    val c = scala.math.cos
                        ^
```

The following example shows how to use this same technique on the `scala.math.pow` method, which takes two parameters:

```
scala> val p = scala.math.pow(_, _)
pow: (Double, Double) => Double = <function2>
```

```
scala> p(scala.math.E, 2)
res0: Double = 7.3890560989306495
```

If this seems like an interesting language feature, but you're wondering where it would be useful, see [Recipe 9.6, "Using Partially Applied Functions"](#), for more information.

Summary notes:

- Think of the `=>` symbol as a transformer. It transforms the input data on its left side to some new output data, using the algorithm on its right side.
- Use `def` to define a method, `val`, to create a function.
- When assigning a function to a variable, a *function literal* is the code on the right side of the expression.
- A *function value* is an object, and extends the `FunctionN` traits in the main `scala` package, such as `Function0` for a function that takes no parameters.

See Also

The `Function1` trait

9.3. Defining a Method That Accepts a Simple Function Parameter

Problem

You want to create a method that takes a simple function as a method parameter.

Solution

This solution follows a three-step process:

1. Define your method, including the signature for the function you want to take as a method parameter.
2. Define one or more functions that match this signature.
3. Sometime later, pass the function(s) as a parameter to your method.

To demonstrate this, define a method named `executeFunction`, which takes a function as a parameter. The method will take one parameter named `callback`, which is a function. That function must have no input parameters and must return nothing:

```
def executeFunction(callback:() => Unit) {  
  callback()  
}
```

Two quick notes:

- The `callback:()` syntax defines a function that has no parameters. If the function had parameters, the types would be listed inside the parentheses.
- The `=> Unit` portion of the code indicates that this method returns nothing.

I'll discuss this syntax more shortly.

Next, define a function that matches this signature. The following function named `sayHello` takes no input parameters and returns nothing:

```
val sayHello = () => { println("Hello") }
```

In the last step of the recipe, pass the `sayHello` function to the `executeFunction` method:

```
executeFunction(sayHello)
```

The REPL demonstrates how this works:

```
scala> def executeFunction(callback:() => Unit) { callback() }  
executeFunction: (callback: () => Unit)Unit  
  
scala> val sayHello = () => { println("Hello") }  
sayHello: () => Unit = <function0>  
  
scala> executeFunction(sayHello)  
Hello
```

Discussion

There isn't anything special about the `callback` name used in this example. When I first learned how to pass functions to methods, I preferred the name `callback` because it made the meaning clear, but it's just the name of a method parameter. These days, just as I often name an `Int` parameter `i`, I name a function parameter `f`:

```
def executeFunction(f:() => Unit) {  
  f()  
}
```

The part that is special is that the function that's passed in must match the function signature you define. In this case, you've declared that the function that's passed in must take no arguments and must return nothing:

```
f:() => Unit
```

The general syntax for defining a function as a method parameter is:

```
parameterName: (parameterType(s)) => returnType
```

In the example, the `parameterName` is `f`, the `parameterType` is empty because you don't want the function to take any parameters, and the return type is `Unit` because you don't want the function to return anything:

```
executeFunction(f:() => Unit)
```

To define a function that takes a `String` and returns an `Int`, use one of these two signatures:

```
executeFunction(f:String => Int)
executeFunction(f:(String) => Int)
```

See the next recipe for more function signature examples.

Scala's Unit

The Scala `Unit` shown in these examples is similar to Java's `Void` class. It's used in situations like this to indicate that the function returns nothing ... or perhaps nothing of interest.

As a quick look into its effect, first define a method named `plusOne`, which does what its name implies:

```
scala> def plusOne(i: Int) = i + 1
plusOne: (i: Int)Int
```

```
scala> plusOne(1)
res0: Int = 2
```

When it's called, `plusOne` adds 1 to its input parameter, and returns that result as an `Int`.

Now, modify `plusOne` to declare that it returns `Unit`:

```
scala> def plusOne(i: Int): Unit = i + 1
plusOne: (i: Int)Unit
```

```
scala> plusOne(1)
(returns nothing)
```

Because you explicitly stated that `plusOne` returns `Unit`, there's no result in the REPL when `plusOne(1)` is called.

This isn't a common use of `Unit`, but it helps to demonstrate its effect.

See Also

Scala's *call-by-name* functionality provides a very simple way to pass a block of code into a function or method. See [Recipe 19.8, "Building Functionality with Types"](#), for several call-by-name examples.

9.4. More Complex Functions

Problem

You want to define a method that takes a function as a parameter, and that function may have one or more input parameters, and may also return a value.

Solution

Following the approach described in the previous recipe, define a method that takes a function as a parameter. Specify the function signature you expect to receive, and then execute that function inside the body of the method.

The following example defines a method named `exec` that takes a function as an input parameter. That function must take one `Int` as an input parameter and return nothing:

```
def exec(callback: Int => Unit) {  
    // invoke the function we were given, giving it an Int parameter  
    callback(1)  
}
```

Next, define a function that matches the expected signature. The following `plusOne` function matches that signature, because it takes an `Int` argument and returns nothing:

```
val plusOne = (i: Int) => { println(i+1) }
```

Now you can pass `plusOne` into the `exec` function:

```
exec(plusOne)
```

Because the function is called inside the method, this prints the number 2.

Any function that matches this signature can be passed into the `exec` method. To demonstrate this, define a new function named `plusTen` that also takes an `Int` and returns nothing:

```
val plusTen = (i: Int) => { println(i+10) }
```

Now you can pass it into your `exec` function, and see that it also works:

```
exec(plusTen)    // prints 11
```

Although these examples are simple, you can see the power of the technique: you can easily swap in interchangeable algorithms. As long as your function signature matches what your method expects, your algorithms can do anything you want. This is comparable to swapping out algorithms in the OOP *Strategy* design pattern.

Discussion

The general syntax for describing a function as a method parameter is this:


```
parameterName: (parameterType(s)) => returnType
```

Therefore, to define a function that takes a `String` and returns an `Int`, use one of these two signatures:

```
executeFunction(f:(String) => Int)
```

```
// parentheses are optional when the function has only one parameter  
executeFunction(f:String => Int)
```

To define a function that takes two `Ints` and returns a `Boolean`, use this signature:

```
executeFunction(f:(Int, Int) => Boolean)
```

The following `exec` method expects a function that takes `String`, `Int`, and `Double` parameters and returns a `Seq[String]`:

```
exec(f:(String, Int, Double) => Seq[String])
```

As shown in the Solution, if a function doesn't return anything, declare its return type as `Unit`:

```
exec(f:(Int) => Unit)  
exec(f:Int => Unit)
```

Passing in a function with other parameters

A function parameter is just like any other method parameter, so a method can accept other parameters in addition to a function.

The following code demonstrates this in a simple example. First, define a simple function:

```
val sayHello = () => println("Hello")
```

Next, define a method that takes this function as a parameter and also takes a second `Int` parameter:

```
def executeXTimes(callback:() => Unit, numTimes: Int) {  
  for (i <- 1 to numTimes) callback()  
}
```

Next, pass the function value and an `Int` into the method:

```
scala> executeXTimes(sayHello, 3)  
Hello  
Hello  
Hello
```

Though that was a simple example, this technique can be used to pass variables into the method that can then be used by the function, inside the method body. To see how this works, create a method named `executeAndPrint` that takes a function and two `Int` parameters:

```
def executeAndPrint(f:(Int, Int) => Int, x: Int, y: Int) {
    val result = f(x, y)
    println(result)
}
```

This method is more interesting than the previous method, because it takes the `Int` parameters it's given and passes those parameters to the function it's given in this line of code:

```
val result = f(x, y)
```

To show how this works, create two functions that match the signature of the function that `executeAndPrint` expects, a `sum` function and a `multiply` function:

```
val sum = (x: Int, y: Int) => x + y
val multiply = (x: Int, y: Int) => x * y
```

Now you can call `executeAndPrint` like this, passing in the different functions, along with two `Int` parameters:

```
executeAndPrint(sum, 2, 9)    // prints 11
executeAndPrint(multiply, 3, 9) // prints 27
```

This is cool, because the `executeAndPrint` method doesn't know what algorithm is actually run. All it knows is that it passes the parameters `x` and `y` to the function it is given and then prints the result from that function. This is similar to defining an interface in Java and then providing concrete implementations of the interface in multiple classes.

Here's one more example of this three-step process:

```
// 1 - define the method
def exec(callback: (Any, Any) => Unit, x: Any, y: Any) {
    callback(x, y)
}

// 2 - define a function to pass in
val printTwoThings =(a: Any, b: Any) => {
    println(a)
    println(b)
}

// 3 - pass the function and some parameters to the method
case class Person(name: String)
exec(printTwoThings, "Hello", Person("Dave"))
```

Note that in all of the previous examples where you created functions with the `val` keyword, you could have created methods, and the examples would still work. For instance, you can define `printTwoThings` as a method, and `exec` still works:

```
// 2a - define a method to pass in
def printTwoThings (a: Any, b: Any) {
    println(a)
}
```

```

    println(b)
  }

  // 3a - pass the printTwoThings method to the exec method
  case class Person(name: String)
  exec(printTwoThings, "Hello", Person("Dave"))

```

Behind the scenes, there are differences between these two approaches—for instance, a function implements one of the `Function0` to `Function22` traits—but Scala is forgiving, and lets you pass in either a method or function, as long as the signature is correct.

9.5. Using Closures

Problem

You want to pass a function around like a variable, and while doing so, you want that function to be able to refer to one or more fields that were in the same scope as the function when it was declared.

Solution

To demonstrate a closure in Scala, use the following simple (but complete) example:

```

package otherscope {

  class Foo {
    // a method that takes a function and a string, and passes the string into
    // the function, and then executes the function
    def exec(f:(String) => Unit, name: String) {
      f(name)
    }
  }

}

object ClosureExample extends App {

  var hello = "Hello"
  def sayHello(name: String) { println(s"$hello, $name") }

  // execute sayHello from the exec method foo
  val foo = new otherscope.Foo
  foo.exec(sayHello, "Al")

  // change the local variable 'hello', then execute sayHello from
  // the exec method of foo, and see what happens
  hello = "Hola"
  foo.exec(sayHello, "Lorenzo")

}

```

To test this code, save it as a file named *ClosureExample.scala*, then compile and run it. When it's run, the output will be:

```
Hello, Al  
Hola, Lorenzo
```

If you're coming to Scala from Java or another OOP language, you might be asking, "How could this possibly work?" Not only did the `sayHello` method reference the variable `hello` from within the `exec` method of the `Foo` class on the first run (where `hello` was no longer in scope), but on the second run, it also picked up the change to the `hello` variable (from `Hello` to `Hola`). The simple answer is that Scala supports closure functionality, and this is how closures work.

As Dean Wampler and Alex Payne describe in their book *Programming Scala* (O'Reilly), there are two *free variables* in the `sayHello` method: `name` and `hello`. The `name` variable is a formal parameter to the function; this is something you're used to.

However, `hello` is not a formal parameter; it's a reference to a variable in the enclosing scope (similar to the way a method in a Java class can refer to a field in the same class). Therefore, the Scala compiler creates a closure that encompasses (or "closes over") `hello`.



You could continue to pass the `sayHello` method around so it gets farther and farther away from the scope of the `hello` variable, but in an effort to keep this example simple, it's only passed to one method in a class in a different package. You can verify that `hello` is not in scope in the `Foo` class by attempting to print its value in that class or in its `exec` method, such as with `println(hello)`. You'll find that the code won't compile because `hello` is not in scope there.

Discussion

In my research, I've found many descriptions of closures, each with slightly different terminology. Wikipedia defines a closure like this:

"In computer science, a closure (also lexical closure or function closure) is a function together with a referencing environment for the non-local variables of that function. A closure allows a function to access variables outside its immediate lexical scope."

In his excellent article, *Closures in Ruby*, Paul Cantrell states, "a closure is a block of code which meets three criteria." He defines the criteria as follows:

1. The block of code can be passed around as a value, and
2. It can be executed on demand by anyone who has that value, at which time
3. It can refer to variables from the context in which it was created (i.e., it is closed with respect to variable access, in the mathematical sense of the word “closed”).

Personally, I like to think of a closure as being like quantum entanglement, which Einstein referred to as “a spooky action at a distance.” Just as quantum entanglement begins with two elements that are together and then separated—but somehow remain aware of each other—a closure begins with a function and a variable defined in the same scope, which are then separated from each other. When the function is executed at some other point in space (scope) and time, it is magically still aware of the variable it referenced in their earlier time together, and even picks up any changes to that variable.

As shown in the Solution, to create a closure in Scala, just define a function that refers to a variable that’s in the same scope as its declaration. That function can be used later, even when the variable is no longer in the function’s current scope, such as when the function is passed to another class, method, or function.

Any time you run into a situation where you’re passing around a function, and wish that function could refer to a variable like this, a closure can be a solution. The variable can be a collection, an `Int` you use as a counter or limit, or anything else that helps to solve a problem. The value you refer to can be a `val`, or as shown in the example, a `var`.

A second example

If you’re new to closures, another example may help demonstrate them. First, start with a simple function named `isOfVotingAge`. This function tests to see if the age given to the function is greater than or equal to 18:

```
val isOfVotingAge = (age: Int) => age >= 18
isOfVotingAge(16) // false
isOfVotingAge(20) // true
```

Next, to make your function more flexible, instead of hardcoding the value 18 into the function, you can take advantage of this closure technique, and let the function refer to the variable `votingAge` that’s in scope when you define the function:

```
var votingAge = 18
val isOfVotingAge = (age: Int) => age >= votingAge
```

When called, `isOfVotingAge` works as before:

```
isOfVotingAge(16) // false
isOfVotingAge(20) // true
```

You can now pass `isOfVotingAge` around to other methods and functions:

```
def printResult(f: Int => Boolean, x: Int) {
  println(f(x))
}
printResult(isOfVotingAge, 20) // true
```

Because you defined `votingAge` as a `var`, you can reassign it. How does this affect `printResult`? Let's see:

```
// change votingAge in one scope
votingAge = 21

// the change to votingAge affects the result
printResult(isOfVotingAge, 20) // now false
```

Cool. The field and function are still entangled.

Using closures with other data types

In the two examples shown so far, you've worked with simple `String` and `Int` fields, but closures can work with any data type, including collections. For instance, in the following example, the function named `addToBasket` is defined in the same scope as an `ArrayBuffer` named `fruits`:

```
import scala.collection.mutable.ArrayBuffer
val fruits = ArrayBuffer("apple")

// the function addToBasket has a reference to fruits
val addToBasket = (s: String) => {
  fruits += s
  println(fruits.mkString(", "))
}
```

As with the previous example, the `addToBasket` function can now be passed around as desired, and will always have a reference to the `fruits` field. To demonstrate this, define a method that accepts a function with `addToBasket`'s signature:

```
def buyStuff(f: String => Unit, s: String) {
  f(s)
}
```

Then pass `addToBasket` and a `String` parameter to the method:

```
scala> buyStuff(addToBasket, "cherries")
cherries

scala> buyStuff(addToBasket, "grapes")
cherries, grapes
```

As desired, the elements are added to your `ArrayBuffer`.

Note that the `buyStuff` method would typically be in another class, but this example demonstrates the basic idea.

A comparison to Java

If you're coming to Scala from Java, or an OOP background in general, it may help to see a comparison between this closure technique and what you can currently do in Java. (In Java, there are some closure-like things you can do with inner classes, and closures are intended for addition to Java 8 in [Project Lambda](#). But this example attempts to show a simple OOP example.)

The following example shows how a `sayHello` method and the `helloPhrase` string are encapsulated in the class `Greeter`. In the `main` method, the first two examples with `Al` and `Lorenzo` show how the `sayHello` method can be called directly.

At the end of the `main` method, the `greeter` instance is passed to an instance of the `Bar` class, and `greeter`'s `sayHello` method is executed from there:

```
public class SimulatedClosure {

    public static void main (String[] args) {
        Greeter greeter = new Greeter();
        greeter.setHelloPhrase("Hello");
        greeter.sayHello("Al");           // "Hello, Al"

        greeter.setHelloPhrase("Hola");
        greeter.sayHello("Lorenzo");     // "Hola, Lorenzo"

        greeter.setHelloPhrase("Yo");
        Bar bar = new Bar(greeter);      // pass the greeter instance to a new Bar
        bar.sayHello("Adrian");          // invoke greeter.sayHello via Bar
    }
}

class Greeter {

    private String helloPhrase;

    public void setHelloPhrase(String helloPhrase) {
        this.helloPhrase = helloPhrase;
    }

    public void sayHello(String name) {
        System.out.println(helloPhrase + ", " + name);
    }
}

class Bar {

    private Greeter greeter;

    public Bar (Greeter greeter) {
        this.greeter = greeter;
    }
}
```

```

    }

    public void sayHello(String name) {
        greeter.sayHello(name);
    }
}

```

Running this code prints the following output:

```

Hello, Al
Hola, Lorenzo
Yo, Adrian

```

The end result is similar to the Scala closure approach, but the big differences in this example are that you’re passing around a Greeter instance (instead of a function), and sayHello and the helloPhrase are encapsulated in the Greeter class. In the Scala closure solution, you passed around a function that was coupled with a field from another scope.

See Also

- The voting age example in this recipe was inspired by Mario Gleichmann’s example in [Functional Scala: Closures](#).
- Paul Cantrell’s article, [Closures in Ruby](#).
- [Recipe 3.18, “Creating Your Own Control Structures”](#), demonstrates the use of multiple parameter lists.
- [Java 8’s Project Lambda](#).

9.6. Using Partially Applied Functions

Problem

You want to eliminate repetitively passing variables into a function by (a) passing common variables into the function to (b) create a new function that is preloaded with those values, and then (c) use the new function, passing it only the unique variables it needs.

Solution

The classic example of a partially applied function begins with a simple sum function:

```
val sum = (a: Int, b: Int, c: Int) => a + b + c

```


There's nothing special about this `sum` function, it's just a normal function. But things get interesting when you supply two of the parameters when calling the function, but don't provide the third parameter:

```
val f = sum(1, 2, _: Int)
```

Because you haven't provided a value for the third parameter, the resulting variable `f` is a *partially applied function*. You can see this in the REPL:

```
scala> val sum = (a: Int, b: Int, c: Int) => a + b + c
sum: (Int, Int, Int) => Int = <function3>
```

```
scala> val f = sum(1, 2, _: Int)
f: Int => Int = <function1>
```

The result in the REPL shows that `f` is a function that implements the `function1` trait, meaning that it takes one argument. Looking at the rest of the signature, you see that it takes an `Int` argument, and returns an `Int` value.

When you give `f` an `Int`, such as the number 3, you magically get the sum of the three numbers that have been passed into the two functions:

```
scala> f(3)
res0: Int = 6
```

The first two numbers (1 and 2) were passed into the original `sum` function; that process created the new function named `f`, which is a partially applied function; then, some time later in the code, the third number (3) was passed into `f`.

Discussion

In functional programming languages, when you call a function that has parameters, you are said to be applying the function to the parameters. When all the parameters are passed to the function—something you always do in Java—you have fully applied the function to all of the parameters. But when you give only a subset of the parameters to the function, the result of the expression is a partially applied function.

As demonstrated in the example, this partially applied function is a variable that you can pass around. This variable is called a *function value*, and when you later provide all the parameters needed to complete the function value, the original function is executed and a result is yielded.

This technique has many advantages, including the ability to make life easier for the consumers of a library you create. For instance, when working with HTML, you may want a function that adds a prefix and a suffix to an HTML snippet:

```
def wrap(prefix: String, html: String, suffix: String) = {
  prefix + html + suffix
}
```

If at a certain point in your code, you know that you always want to add the same prefix and suffix to different HTML strings, you can apply those two parameters to the function, without applying the `html` parameter:

```
val wrapWithDiv = wrap("<div>", _: String, "</div>")
```

Now you can call the new `wrapWithDiv` function, just passing it the HTML you want to wrap:

```
scala> wrapWithDiv("<p>Hello, world</p>")
res0: String = <div><p>Hello, world</p></div>

scala> wrapWithDiv("<img src=\"/images/foo.png\" />")
res1: String = <div></div>
```

The `wrapWithDiv` function is preloaded with the `<div>` tags you applied, so it can be called with just one argument: the HTML you want to wrap.

As a nice benefit, you can still call the original `wrap` function if you want:

```
wrap("<pre>", "val x = 1", "</pre>")
```

You can use partially applied functions to make programming easier by binding some arguments—typically some form of local arguments—and leaving the others to be filled in.

9.7. Creating a Function That Returns a Function

Problem

You want to return a function (algorithm) from a function or method.

Solution

Define a function that returns an algorithm (an anonymous function), assign that to a new function, and then call that new function.

The following code declares an anonymous function that takes a `String` argument and returns a `String`:

```
(s: String) => { prefix + " " + s }
```

You can return that anonymous function from the body of another function as follows:

```
def saySomething(prefix: String) = (s: String) => {
  prefix + " " + s
}
```

Because `saySomething` returns a function, you can assign that resulting function to a variable. The `saySomething` function requires a `String` argument, so give it one as you create the resulting function `sayHello`:

```
val sayHello = saySomething("Hello")
```

The `sayHello` function is now equivalent to your anonymous function, with the `prefix` set to `hello`. Looking back at the anonymous function, you see that it takes a `String` parameter and returns a `String`, so you pass it a `String`:

```
sayHello("Al")
```

Here's what these steps look like in the REPL:

```
scala> def saySomething(prefix: String) = (s: String) => {
      |   prefix + " " + s
      | }
saySomething: (prefix: String)String => java.lang.String

scala> val sayHello = saySomething("Hello")
sayHello: String => java.lang.String = <function1>

scala> sayHello("Al")
res0: java.lang.String = Hello Al
```

Discussion

If you're new to functional programming, it can help to break this down a little. You can break the expression down into its two components. On the left side of the `=` symbol you have a normal method declaration:

```
def saySomething(prefix: String)
```

On the right side of the `=` is a function literal (also known as an anonymous function):

```
(s: String) => { prefix + " " + s }
```

Another example

As you can imagine, you can use this approach any time you want to encapsulate an algorithm inside a function. A bit like a `Factory` or `Strategy` pattern, the function your method returns can be based on the input parameter it receives. For example, create a greeting method that returns an appropriate greeting based on the language specified:

```
def greeting(language: String) = (name: String) => {
  language match {
    case "english" => "Hello, " + name
    case "spanish" => "Buenos dias, " + name
  }
}
```

If it's not clear that `greeting` is returning a function, you can make the code a little more explicit:

```
def greeting(language: String) = (name: String) => {
  val english = () => "Hello, " + name
  val spanish = () => "Buenos dias, " + name
  language match {
```

```

    case "english" => println("returning 'english' function")
                        english()
    case "spanish" => println("returning 'spanish' function")
                        spanish()
  }
}

```

Here's what this second method looks like when it's invoked in the REPL:

```

scala> val hello = greeting("english")
hello: String => java.lang.String = <function1>

scala> val buenosDias = greeting("spanish")
buenosDias: String => java.lang.String = <function1>

scala> hello("Al")
returning 'english' function
res0: java.lang.String = Hello, Al

scala> buenosDias("Lorenzo")
returning 'spanish' function
res1: java.lang.String = Buenos dias, Lorenzo

```

You can use this recipe any time you want to encapsulate one or more functions behind a method, and is similar in that effect to the Factory and Strategy patterns.

See Also

[My Java Factory Pattern example](#)

9.8. Creating Partial Functions

Problem

You want to define a function that will only work for a subset of possible input values, or you want to define a series of functions that only work for a subset of input values, and combine those functions to completely solve a problem.

Solution

A *partial function* is a function that does not provide an answer for every possible input value it can be given. It provides an answer only for a subset of possible data, and defines the data it can handle. In Scala, a partial function can also be queried to determine if it can handle a particular value.

As a simple example, imagine a normal function that divides one number by another:

```
val divide = (x: Int) => 42 / x
```

As defined, this function blows up when the input parameter is zero:

```
scala> divide(0)
java.lang.ArithmeticException: / by zero
```

Although you can handle this particular situation by catching and throwing an exception, Scala lets you define the `divide` function as a `PartialFunction`. When doing so, you also explicitly state that the function is defined when the input parameter is not zero:

```
val divide = new PartialFunction[Int, Int] {
  def apply(x: Int) = 42 / x
  def isDefinedAt(x: Int) = x != 0
}
```

With this approach, you can do several nice things. One thing you can do is test the function before you attempt to use it:

```
scala> divide.isDefinedAt(1)
res0: Boolean = true

scala> if (divide.isDefinedAt(1)) divide(1)
res1: AnyVal = 42

scala> divide.isDefinedAt(0)
res2: Boolean = false
```

This isn't all you can do with partial functions. You'll see shortly that other code can take advantage of partial functions to provide elegant and concise solutions.

Whereas that `divide` function is explicit about what data it handles, partial functions are often written using case statements:

```
val divide2: PartialFunction[Int, Int] = {
  case d: Int if d != 0 => 42 / d
}
```

Although this code doesn't explicitly implement the `isDefinedAt` method, it works exactly the same as the previous `divide` function definition:

```
scala> divide2.isDefinedAt(0)
res0: Boolean = false

scala> divide2.isDefinedAt(1)
res1: Boolean = true
```

The PartialFunction explained

The `PartialFunction` Scaladoc describes a partial function in this way:

A partial function of type `PartialFunction[A, B]` is a unary function where the domain does not necessarily include all values of type `A`. The function `isDefinedAt` allows [you] to test dynamically if a value is in the domain of the function.

This helps to explain why the last example with the match expression (case statement) works: the `isDefinedAt` method dynamically tests to see if the given value is in the domain of the function (i.e., it is handled, or accounted for).

The signature of the `PartialFunction` trait looks like this:

```
trait PartialFunction[-A, +B] extends (A) => B
```

As discussed in other recipes, the `=>` symbol can be thought of as a transformer, and in this case, the `(A) => B` can be interpreted as a function that transforms a type `A` into a resulting type `B`.

The example method transformed an input `Int` into an output `Int`, but if it returned a `String` instead, it would be declared like this:

```
PartialFunction[Int, String]
```

For example, the following method uses this signature:

```
// converts 1 to "one", etc., up to 5
val convertLowNumToString = new PartialFunction[Int, String] {
  val nums = Array("one", "two", "three", "four", "five")
  def apply(i: Int) = nums(i-1)
  def isDefinedAt(i: Int) = i > 0 && i < 6
}
```

orElse and andThen

A terrific feature of partial functions is that you can chain them together. For instance, one method may only work with even numbers, and another method may only work with odd numbers. Together they can solve all integer problems.

In the following example, two functions are defined that can each handle a small number of `Int` inputs, and convert them to `String` results:

```
// converts 1 to "one", etc., up to 5
val convert1to5 = new PartialFunction[Int, String] {
  val nums = Array("one", "two", "three", "four", "five")
  def apply(i: Int) = nums(i-1)
  def isDefinedAt(i: Int) = i > 0 && i < 6
}

// converts 6 to "six", etc., up to 10
val convert6to10 = new PartialFunction[Int, String] {
  val nums = Array("six", "seven", "eight", "nine", "ten")
  def apply(i: Int) = nums(i-6)
  def isDefinedAt(i: Int) = i > 5 && i < 11
}
```

Taken separately, they can each handle only five numbers. But combined with `orElse`, they can handle ten:

```
scala> val handle1to10 = convert1to5 orElse convert6to10
handle1to10: PartialFunction[Int,String] = <function1>
```

```
scala> handle1to10(3)
res0: String = three

scala> handle1to10(8)
res1: String = eight
```

The `orElse` method comes from the `Scala PartialFunction` trait, which also includes the `andThen` method to further help chain partial functions together.

Discussion

It's important to know about partial functions, not just to have another tool in your toolbox, but because they are used in the APIs of some libraries, including the Scala collections library.

One example of where you'll run into partial functions is with the `collect` method on collections' classes. The `collect` method takes a partial function as input, and as its Scaladoc describes, `collect` "Builds a new collection by applying a partial function to all elements of this list on which the function is defined."

For instance, the `divide` function shown earlier is a partial function that is not defined at the `Int` value zero. Here's that function again:

```
val divide: PartialFunction[Int, Int] = {
  case d: Int if d != 0 => 42 / d
}
```

If you attempt to use this function with the `map` method, it will explode with a `MatchError`:

```
scala> List(0,1,2) map { divide }
scala.MatchError: 0 (of class java.lang.Integer)
stack trace continues ...
```

However, if you use the same function with the `collect` method, it works fine:

```
scala> List(0,1,2) collect { divide }
res0: List[Int] = List(42, 21)
```

This is because the `collect` method is written to test the `isDefinedAt` method for each element it's given. As a result, it doesn't run the `divide` algorithm when the input value is 0 (but does run it for every other element).

You can see the `collect` method work in other situations, such as passing it a `List` that contains a mix of data types, with a function that works only with `Int` values:

```
scala> List(42, "cat") collect { case i: Int => i + 1 }
res0: List[Int] = List(43)
```

Because it checks the `isDefinedAt` method under the covers, `collect` can handle the fact that your anonymous function can't work with a `String` as input.

The `PartialFunction` Scaladoc demonstrates this same technique in a slightly different way. In the first example, it shows how to create a list of even numbers by defining a `PartialFunction` named `isEven`, and using that function with the `collect` method:

```
scala> val sample = 1 to 5
sample: scala.collection.immutable.Range.Inclusive = Range(1, 2, 3, 4, 5)

scala> val isEven: PartialFunction[Int, String] = {
  |   case x if x % 2 == 0 => x + " is even"
  | }
isEven: PartialFunction[Int,String] = <function1>

scala> val evenNumbers = sample collect isEven
evenNumbers: scala.collection.immutable.IndexedSeq[String] =
  Vector(2 is even, 4 is even)
```

Similarly, an `isOdd` function can be defined, and the two functions can be joined by `orElse` to work with the `map` method:

```
scala> val isOdd: PartialFunction[Int, String] = {
  |   case x if x % 2 == 1 => x + " is odd"
  | }
isOdd: PartialFunction[Int,String] = <function1>

scala> val numbers = sample map (isEven orElse isOdd)
numbers: scala.collection.immutable.IndexedSeq[String] =
  Vector(1 is odd, 2 is even, 3 is odd, 4 is even, 5 is odd)
```



Portions of this recipe were inspired by Erik Bruchez's blog post, titled, "Scala partial functions (without a PhD)."

See Also

- [Erik Bruchez's blog post](#)
- [PartialFunction trait](#)
- [Wikipedia definition of a partial function](#)

9.9. A Real-World Example

Problem

Understanding functional programming concepts is one thing; putting them into practice in a real project is another. You'd like to see a real example of them in action.

Solution

To demonstrate some of the techniques introduced in this chapter, the following example shows one way to implement Newton's Method, a mathematical method that can be used to solve the roots of equations.

As you can see from the code, the method named `newtonsMethod` takes functions as its first two parameters. It also takes two other `Double` parameters, and returns a `Double`. The two functions that are passed in should be the original equation (`fx`) and the derivative of that equation (`fxPrime`).

The method `newtonsMethodHelper` also takes two functions as parameters, so you can see how the functions are passed from `newtonsMethod` to `newtonsMethodHelper`.

Here is the complete source code for this example:

```
object NewtonsMethod {

    def main(args: Array[String]) {
        driver
    }

    /**
     * A "driver" function to test Newton's method.
     * Start with (a) the desired  $f(x)$  and  $f'(x)$  equations,
     * (b) an initial guess and (c) tolerance values.
     */
    def driver {
        // the  $f(x)$  and  $f'(x)$  functions
        val fx = (x: Double) => 3*x + math.sin(x) - math.pow(math.E, x)
        val fxPrime = (x: Double) => 3 + math.cos(x) - math.pow(Math.E, x)

        val initialGuess = 0.0
        val tolerance = 0.00005

        // pass  $f(x)$  and  $f'(x)$  to the Newton's Method function, along with
        // the initial guess and tolerance
        val answer = newtonsMethod(fx, fxPrime, initialGuess, tolerance)

        println(answer)
    }

    /**
     * Newton's Method for solving equations.
     * @todo check that  $|f(x_{\text{Next}})|$  is greater than a second tolerance value
     * @todo check that  $f'(x) \neq 0$ 
     */
    def newtonsMethod(fx: Double => Double,
                      fxPrime: Double => Double,
                      x: Double,
                      tolerance: Double): Double = {

        var x1 = x
```

```

var xNext = newtonsMethodHelper(fx, fxPrime, x1)
while (math.abs(xNext - x1) > tolerance) {
  x1 = xNext
  println(xNext) // debugging (intermediate values)
  xNext = newtonsMethodHelper(fx, fxPrime, x1)
}

xNext
}

/**
 * This is the "x2 = x1 - f(x1)/f'(x1)" calculation
 */
def newtonsMethodHelper(fx: Double => Double,
                        fxPrime: Double => Double,
                        x: Double): Double = {
  x - fx(x) / fxPrime(x)
}
}

```

Discussion

As you can see, a majority of this code involves defining functions, passing those functions to methods, and then invoking the functions from within a method.

The method name `newtonsMethod` will work for any two functions `fx` and `fxPrime`, where `fxPrime` is the derivative of `fx` (within the limits of the “to do” items that are not implemented).

To experiment with this example, try changing the functions `fx` and `fxPrime`, or implement the @todo items in `newtonsMethod`.



The algorithm shown comes from an old textbook titled *Applied Numerical Analysis*, by Gerald and Wheatley, where the approach was demonstrated in pseudocode.

See Also

- [More details on this example](#)
- [Newton's Method](#)

Introduction

Scala's collection classes are rich, deep, and differ significantly from the Java collections, all of which makes learning them a bit of a speed bump for developers coming to Scala from Java.

When a Java developer first comes to Scala, she might think, “Okay, I’ll use lists and arrays, right?” Well, not really. The Scala `List` class is very different from the Java `List` classes—including the part where it’s immutable—and although the Scala `Array` is an improvement on the Java array in most ways, it’s not even recommended as the “go to” sequential collection class.

Because there are many collections classes to choose from, and each of those classes offers many methods, a goal of this chapter (and the next) is to help guide you through this plethora of options to find the solutions you need. Recipes will help you decide which collections to use in different situations, and also choose a method to solve a problem. To help with this, the methods that are common to all collections are shown in this chapter, and methods specific to collections like `List`, `Array`, `Map`, and `Set` are shown in [Chapter 11](#).

A Few Important Concepts

There are a few important concepts to know when working with the methods of the Scala collection classes:

- What a predicate is
- What an anonymous function is
- Implied loops

A *predicate* is simply a method, function, or anonymous function that takes one or more parameters and returns a Boolean value. For instance, the following method returns true or false, so it's a predicate:

```
def isEven (i: Int) = if (i % 2 == 0) true else false
```

That's a simple concept, but you'll hear the term so often when working with collection methods that it's important to mention it.

The concept of an *anonymous function* is also important. They're described in depth in [Recipe 9.1](#), but here's an example of the long form for an anonymous function:

```
(i: Int) => i % 2 == 0
```

Here's the short form of the same function:

```
_ % 2 == 0
```

That doesn't look like much by itself, but when it's combined with the `filter` method on a collection, it makes for a lot of power in just a little bit of code:

```
scala> val list = List.range(1, 10)
list: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val events = list.filter(_ % 2 == 0)
events: List[Int] = List(2, 4, 6, 8)
```

This is a nice lead-in into the third topic: *implied loops*. As you can see from that example, the `filter` method contains a loop that applies your function to every element in the collection and returns a new collection. You could live without the `filter` method and write equivalent code like this:

```
for {
  e <- list
  if e % 2 == 0
} yield e
```

But I think you'll agree that the `filter` approach is both more concise and easier to read.

Collection methods like `filter`, `foreach`, `map`, `reduceLeft`, and many more have loops built into their algorithms. As a result, you'll write far fewer loops when writing Scala code than with another language like Java.

10.1. Understanding the Collections Hierarchy

Problem

The Scala collections hierarchy is very rich (deep and wide), and understanding how it's organized can be helpful when choosing a collection to solve a problem.

Solution

Figure 10-1, which shows the traits from which the `Vector` class inherits, demonstrates some of the complexity of the Scala collections hierarchy.




Figure 10-1. The traits inherited by the `Vector` class

Because Scala classes can inherit from traits, and well-designed traits are granular, a class hierarchy can look like this. However, don't let Figure 10-1 throw you for a loop: you don't need to know all those traits to use a `Vector`. In fact, using a `Vector` is straightforward:

```
val v = Vector(1, 2, 3)
v.sum           // 6
v.filter(_ > 1) // Vector(2, 3)
v.map(_ * 2)    // Vector(2, 4, 6)
```

At a high level, Scala's collection classes begin with the `Traversable` and `Iterable` traits, and extend into the three main categories of sequences (`Seq`), sets (`Set`), and maps (`Map`). Sequences further branch off into *indexed* and *linear* sequences, as shown in Figure 10-2.




Figure 10-2. A high-level view of the Scala collections

The `Traversable` trait lets you traverse an entire collection, and its Scaladoc states that it “implements the behavior common to all collections in terms of a `foreach` method,” which lets you traverse the collection repeatedly.

The `Iterable` trait defines an *iterator*, which lets you loop through a collection’s elements one at a time, but when using an iterator, the collection can be traversed only once, because each element is consumed during the iteration process.

Sequences

Digging a little deeper into the *sequence* hierarchy, Scala contains a large number of sequences, many of which are shown in [Figure 10-3](#).




Figure 10-3. A portion of the Scala sequence hierarchy

These traits and classes are described in [Tables 10-1](#) through [10-4](#).

As shown in [Figure 10-3](#), sequences branch off into two main categories: *indexed sequences* and *linear sequences* (linked lists). An `IndexedSeq` indicates that random access of elements is efficient, such as accessing an `Array` element as `arr(5000)`. By default, specifying that you want an `IndexedSeq` with Scala 2.10.x creates a `Vector`:

```
scala> val x = IndexedSeq(1,2,3)
x: IndexedSeq[Int] = Vector(1, 2, 3)
```

A `LinearSeq` implies that the collection can be efficiently split into head and tail components, and it’s common to work with them using the `head`, `tail`, and `isEmpty` methods. Note that creating a `LinearSeq` creates a `List`, which is a singly linked list:

```
scala> val seq = scala.collection.immutable.LinearSeq(1,2,3)
seq: scala.collection.immutable.LinearSeq[Int] = List(1, 2, 3)
```

Maps

Like a Java Map, Ruby Hash, or Python dictionary, a Scala Map is a collection of key/value pairs, where all the keys must be unique. The most common map classes are shown in Figure 10-4.




Figure 10-4. Common map classes

Map traits and classes are discussed in Table 10-5. When you just need a simple, *immutable* map, you can create one without requiring an import:

```
scala> val m = Map(1 -> "a", 2 -> "b")
m: scala.collection.immutable.Map[Int,java.lang.String] = Map(1 -> a, 2 -> b)
```

The *mutable* map is not in scope by default, so you must import it (or specify its full path) to use it:

```
scala> val m = collection.mutable.Map(1 -> "a", 2 -> "b")
m: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 1 -> a)
```

Sets

Like a Java Set, a Scala Set is a collection of unique elements. The common set classes are shown in Figure 10-5.




Figure 10-5. Common set classes

Set traits and classes are discussed in Table 10-6, but as a quick preview, if you just need an immutable set, you can create it like this, without needing an import statement:

```
scala> val set = Set(1, 2, 3)
set: scala.collection.immutable.Set[Int] = Set(1, 2, 3)
```

Just like a map, if you want to use a mutable set, you must import it, or specify its complete path:

```
scala> val s = collection.mutable.Set(1, 2, 3)
s: scala.collection.mutable.Set[Int] = Set(1, 2, 3)
```

More collection classes

There are many additional collection traits and classes, including `Stream`, `Queue`, `Stack`, and `Range`. You can also create *views* on collections (like a database view); use iterators; and work with the `Option`, `Some`, and `None` types as collections. All of these classes (and objects) are demonstrated in this and the next chapter.

Strict and lazy collections

Collections can also be thought of in terms of being *strict* or *lazy*. See the next recipe for a discussion of these terms.

10.2. Choosing a Collection Class

Problem

You want to choose a Scala collection class to solve a particular problem.

Solution

There are three main categories of collection classes to choose from:

- `Sequence`
- `Map`
- `Set`

A *sequence* is a linear collection of elements and may be indexed or linear (a linked list). A *map* contains a collection of key/value pairs, like a Java `Map`, Ruby `Hash`, or Python dictionary. A *set* is a collection that contains no duplicate elements.

In addition to these three main categories, there are other useful collection types, including `Stack`, `Queue`, and `Range`. There are a few other classes that act like collections, including tuples, enumerations, and the `Option/Some/None` and `Try/Success/Failure` classes.

Choosing a sequence

When choosing a *sequence* (a sequential collection of elements), you have two main decisions:

- Should the sequence be indexed (like an array), allowing rapid access to any elements, or should it be implemented as a linked list?
- Do you want a mutable or immutable collection?

As of Scala 2.10, the recommended, general-purpose, “go to” sequential collections for the combinations of mutable/immutable and indexed/linear are shown in [Table 10-1](#).

Table 10-1. Scala’s general-purpose sequential collections

	Immutable	Mutable
Indexed	Vector	ArrayBuffer
Linear (Linked lists)	List	ListBuffer

As an example of reading that table, if you want an immutable, indexed collection, in general you should use a `Vector`; if you want a mutable, indexed collection, use an `ArrayBuffer` (and so on).

While those are the general-purpose recommendations, there are many more sequence alternatives. The most common *immutable* sequence choices are shown in [Table 10-2](#).

Table 10-2. Main immutable sequence choices

	IndexedSeq	LinearSeq	Description
List		✓	A singly linked list. Suited for recursive algorithms that work by splitting the head from the remainder of the list.
Queue		✓	A first-in, first-out data structure.
Range	✓		A range of integer values.
Stack		✓	A last-in, first-out data structure.
Stream		✓	Similar to <code>List</code> , but it’s lazy and persistent. Good for a large or infinite sequence, similar to a Haskell <code>List</code> .
String	✓		Can be treated as an immutable, indexed sequence of characters.
Vector	✓		The “go to” immutable, indexed sequence. The Scaladoc describes it as, “Implemented as a set of nested arrays that’s efficient at splitting and joining.”

The most common *mutable* sequence choices are shown in [Table 10-3](#). `Queue` and `Stack` are also in this table because there are mutable and immutable versions of these classes.

Table 10-3. Main mutable sequence choices

	IndexedSeq	LinearSeq	Description
Array	✓		Backed by a Java array, its elements are mutable, but it can’t change in size.
ArrayBuffer		✓	The “go to” class for a mutable, sequential collection. The amortized cost for appending elements is constant.

	IndexedSeq	LinearSeq	Description
ArrayStack	✓		A last-in, first-out data structure. Prefer over Stack when performance is important.
DoubleLinkedList		✓	Like a singly linked list, but with a prev method as well. The documentation states, “The additional links make element removal very fast.”
LinkedList		✓	A mutable, singly linked list.
ListBuffer		✓	Like an ArrayBuffer, but backed by a list. The documentation states, “If you plan to convert the buffer to a list, use ListBuffer instead of ArrayBuffer.” Offers constant-time prepend and append; most other operations are linear.
MutableList		✓	A mutable, singly linked list with constant-time append.
Queue		✓	A first-in, first-out data structure.
Stack		✓	A last-in, first-out data structure. (The documentation suggests that an ArrayStack is slightly more efficient.)
StringBuilder	✓		Used to build strings, as in a loop. Like the Java StringBuilder.

In addition to the information shown in these tables, performance can be a consideration. See [Recipe 10.4, “Understanding the Performance of Collections”](#), if performance is important to your selection process.

When creating an API for a library, you may want to refer to your sequences in terms of their superclasses. [Table 10-4](#) shows the traits that are often used when referring generically to a collection in an API.

Table 10-4. Traits commonly used in library APIs

Trait	Description
IndexedSeq	Implies that random access of elements is efficient.
LinearSeq	Implies that linear access to elements is efficient.
Seq	Used when it isn’t important to indicate that the sequence is indexed or linear in nature.

Of course if the collection you’re returning can be *very* generic, you can also refer to the collections as `Iterable` or `Traversable`. This is the rough equivalent of declaring that a Java method returns `Collection`.

You can also learn more about declaring the type a method returns by looking at the “code assist” tool in your IDE. For instance, when I create a new `Vector` in Eclipse and then look at the methods available on a `Vector` instance, I see that the methods return types such as `GenSeqLike`, `IndexedSeqLike`, `IterableLike`, `TraversableLike`, and `TraversableOnce`. You don’t have to be this specific with the types your methods return—certainly not initially—but it’s usually a good practice to identify the *intent* of what you’re really returning, so you can declare these more specific types once you get used to them.

Choosing a map

Choosing a map class is easier than choosing a sequence. There are the base mutable and immutable map classes, a `SortedMap` trait to keep elements in sorted order by key, a `LinkedHashMap` to store elements in insertion order, and a few other maps for special purposes. These options are shown in [Table 10-5](#). (Quotes in the descriptions come from the Scaladoc for each class.)

Table 10-5. Common map choices, including whether immutable or mutable versions are available

	Immutable	Mutable	Description
HashMap	✓	✓	The immutable version “implements maps using a hash trie”; the mutable version “implements maps using a hashtable.”
LinkedHashMap		✓	“Implements mutable maps using a hashtable.” Returns elements by the order in which they were inserted.
ListMap	✓	✓	A map implemented using a list data structure. Returns elements in the opposite order by which they were inserted, as though each element is inserted at the head of the map.
Map	✓	✓	The base map, with both mutable and immutable implementations.
SortedMap	✓		A base trait that stores its keys in sorted order. (Creating a variable as a <code>SortedMap</code> currently returns a <code>TreeMap</code> .)
TreeMap	✓		An immutable, sorted map, implemented as a red-black tree.
WeakHashMap		✓	A hash map with weak references, it’s a wrapper around <code>java.util.WeakHashMap</code> .

You can also create a thread-safe mutable map by mixing the `SynchronizedMap` trait into the map implementation you want. See the [map discussion in the Scala Collections Overview](#) for more information.

Choosing a set

Choosing a set is similar to choosing a map. There are base mutable and immutable set classes, a `SortedSet` to return elements in sorted order by key, a `LinkedHashSet` to store elements in insertion order, and a few other sets for special purposes. The common classes are shown in [Table 10-6](#). (Quotes in the descriptions come from the Scaladoc for each class.)

Table 10-6. Common set choices, including whether immutable or mutable versions are available

	Immutable	Mutable	
BitSet	✓	✓	A set of “non-negative integers represented as variable-size arrays of bits packed into 64-bit words.” Used to save memory when you have a set of integers.
HashSet	✓	✓	The immutable version “implements sets using a hash trie”; the mutable version “implements sets using a hashtable.”

	Immutable	Mutable	
LinkedHashSet		✓	A mutable set implemented using a hashtable. Returns elements in the order in which they were inserted.
ListSet	✓		A set implemented using a list structure.
TreeSet	✓	✓	The immutable version “implements immutable sets using a tree.” The mutable version is a mutable SortedSet with “an immutable AVL Tree as underlying data structure.”
Set	✓	✓	Generic base traits, with both mutable and immutable implementations.
SortedSet	✓	✓	A base trait. (Creating a variable as a SortedSet returns a TreeSet.)

You can also create a thread-safe mutable set by mixing the `SynchronizedSet` trait into the set implementation you want. See the [Scala Collections Overview discussion of maps and sets](#) for more information.

Types that act like collections

Scala offers many other collection types, and some types that act like collections. [Table 10-7](#) provides descriptions of several types that act somewhat like collections, even though they aren’t.

Table 10-7. Other collections classes (and types that act like collections)

	Description
Enumeration	A finite collection of constant values (i.e., the days in a week or months in a year).
Iterator	An iterator isn’t a collection; instead, it gives you a way to access the elements in a collection. It does, however, define many of the methods you’ll see in a normal collection class, including <code>foreach</code> , <code>map</code> , <code>flatMap</code> , etc. You can also convert an iterator to a collection when needed.
Option	Acts as a collection that contains zero or one elements. The <code>Some</code> class and <code>None</code> object extend <code>Option</code> . <code>Some</code> is a container for one element, and <code>None</code> holds zero elements.
Tuple	Supports a heterogeneous collection of elements. There is no one “ <code>Tuple</code> ” class; tuples are implemented as case classes ranging from <code>Tuple1</code> to <code>Tuple22</code> , which support 1 to 22 elements.

Strict and lazy collections

To understand strict and lazy collections, it helps to first understand the concept of a transformer method. A *transformer method* is a method that constructs a new collection from an existing collection. This includes methods like `map`, `filter`, `reverse`, etc.—any method that transforms the input collection to a new output collection.

Given that definition, collections can also be thought of in terms of being strict or lazy. In a *strict* collection, memory for the elements is allocated immediately, and all of its elements are immediately evaluated when a transformer method is invoked. In a *lazy* collection, memory for the elements is not allocated immediately, and transformer methods do not construct new elements until they are demanded.

All of the collection classes except `Stream` are strict, but the other collection classes can be converted to a lazy collection by creating a *view* on the collection. See [Recipe 10.24, “Creating a Lazy View on a Collection”](#), for more information on this approach.

See Also

- In addition to my own experience using the collections, most of the information used to create these tables comes from the Scaladoc of each type, and the [Scala Collections Overview documentation](#).
- [Recipe 10.1, “Understanding the Collections Hierarchy”](#).
- [Recipe 10.4, “Understanding the Performance of Collections”](#).

10.3. Choosing a Collection Method to Solve a Problem

Problem

There is a large number of methods available to Scala collections, and you need to choose a method to solve a problem.

Solution

The Scala collection classes provide a wealth of methods that can be used to manipulate data. Most methods take either a function or a predicate as an argument. (A *predicate* is just a function that returns a Boolean.)

The methods that are available are listed in two ways in this recipe. In the next few paragraphs, the methods are grouped into categories to help you easily find what you need. In the tables that follow, a brief description and method signature is provided.

Methods organized by category

Filtering methods

Methods that can be used to filter a collection include `collect`, `diff`, `distinct`, `drop`, `dropWhile`, `filter`, `filterNot`, `find`, `foldLeft`, `foldRight`, `head`, `headOption`, `init`, `intersect`, `last`, `lastOption`, `reduceLeft`, `reduceRight`, `remove`, `slice`, `tail`, `take`, `takeWhile`, and `union`.

Transformer methods

Transformer methods take at least one input collection to create a new output collection, typically using an algorithm you provide. They include `+`, `++`, `-`, `--`, `diff`, `distinct`, `collect`, `flatMap`, `map`, `reverse`, `sortWith`, `takeWhile`, `zip`, and `zipWithIndex`.

Grouping methods

These methods let you take an existing collection and create multiple groups from that one collection. These methods include `groupBy`, `partition`, `sliding`, `span`, `splitAt`, and `unzip`.

Informational and mathematical methods

These methods provide information about a collection, and include `canEqual`, `contains`, `containsSlice`, `count`, `endsWith`, `exists`, `find`, `forAll`, `hasDefiniteSize`, `indexOf`, `indexOfSlice`, `indexWhere`, `isDefinedAt`, `isEmpty`, `lastIndexOf`, `lastIndexOfSlice`, `lastIndexWhere`, `max`, `min`, `nonEmpty`, `product`, `segmentLength`, `size`, `startsWith`, `sum`. The methods `foldLeft`, `foldRight`, `reduceLeft`, and `reduceRight` can also be used with a function you supply to obtain information about a collection.

Others

A few other methods are hard to categorize, including `par`, `view`, `flatten`, `foreach`, and `mkString`. `par` creates a parallel collection from an existing collection; `view` creates a lazy view on a collection (see [Recipe 10.24](#)); `flatten` converts a list of lists down to one list; `foreach` is like a `for` loop, letting you iterate over the elements in a collection; `mkString` lets you build a `String` from a collection.

There are even more methods than those listed here. For instance, there's a collection of `to*` methods that let you convert the current collection (a `List`, for example) to other collection types (`Array`, `Buffer`, `Vector`, etc.). Check the Scaladoc for your collection class to find more built-in methods.

Common collection methods

The following tables list the most common collection methods.

Table 10-8 lists methods that are common to all collections via `Traversable`. The following symbols are used in the first column of the table:

- `c` refers to a collection
- `f` refers to a function
- `p` refers to a predicate
- `n` refers to a number
- `op` refers to a simple operation (usually a simple function)

Additional methods for mutable and immutable collections are listed in [Tables 10-9 and 10-10](#), respectively.

Table 10-8. Common methods on Traversable collections

Method	Description
<code>c collect f</code>	Builds a new collection by applying a partial function to all elements of the collection on which the function is defined.
<code>c count p</code>	Counts the number of elements in the collection for which the predicate is satisfied.
<code>c1 diff c2</code>	Returns the difference of the elements in <code>c1</code> and <code>c2</code> .
<code>c drop n</code>	Returns all elements in the collection except the first <code>n</code> elements.
<code>c dropWhile p</code>	Returns a collection that contains the “longest prefix of elements that satisfy the predicate.”
<code>c exists p</code>	Returns <code>true</code> if the predicate is true for any element in the collection.
<code>c filter p</code>	Returns all elements from the collection for which the predicate is <code>true</code> .
<code>c filterNot p</code>	Returns all elements from the collection for which the predicate is <code>false</code> .
<code>c find p</code>	Returns the first element that matches the predicate as <code>Some[A]</code> . Returns <code>None</code> if no match is found.
<code>c flatten</code>	Converts a collection of collections (such as a list of lists) to a single collection (single list).
<code>c flatMap f</code>	Returns a new collection by applying a function to all elements of the collection <code>c</code> (like <code>map</code>), and then flattening the elements of the resulting collections.
<code>c foldLeft(z)(op)</code>	Applies the operation to successive elements, going from left to right, starting at element <code>z</code> .
<code>c foldRight(z)(op)</code>	Applies the operation to successive elements, going from right to left, starting at element <code>z</code> .
<code>c forAll p</code>	Returns <code>true</code> if the predicate is true for all elements, <code>false</code> otherwise.
<code>c foreach f</code>	Applies the function <code>f</code> to all elements of the collection.
<code>c groupBy f</code>	Partitions the collection into a <code>Map</code> of collections according to the function.
<code>c hasDefiniteSize</code>	Tests whether the collection has a finite size. (Returns <code>false</code> for a <code>Stream</code> or <code>Iterator</code> , for example.)
<code>c head</code>	Returns the first element of the collection. Throws a <code>NoSuchElementException</code> if the collection is empty.
<code>c headOption</code>	Returns the first element of the collection as <code>Some[A]</code> if the element exists, or <code>None</code> if the collection is empty.
<code>c init</code>	Selects all elements from the collection except the last one. Throws an <code>UnsupportedOperationException</code> if the collection is empty.
<code>c1 intersect c2</code>	On collections that support it, it returns the intersection of the two collections (the elements common to both collections).
<code>c isEmpty</code>	Returns <code>true</code> if the collection is empty, <code>false</code> otherwise.
<code>c last</code>	Returns the last element from the collection. Throws a <code>NoSuchElementException</code> if the collection is empty.
<code>c lastOption</code>	Returns the last element of the collection as <code>Some[A]</code> if the element exists, or <code>None</code> if the collection is empty.
<code>c map f</code>	Creates a new collection by applying the function to all the elements of the collection.
<code>c max</code>	Returns the largest element from the collection.
<code>c min</code>	Returns the smallest element from the collection.

Method	Description
<code>c.nonEmpty</code>	Returns <code>true</code> if the collection is not empty.
<code>c.par</code>	Returns a parallel implementation of the collection, e.g., <code>Array</code> returns <code>ParArray</code> .
<code>c.partition p</code>	Returns two collections according to the predicate algorithm.
<code>c.product</code>	Returns the multiple of all elements in the collection.
<code>c.reduceLeft op</code>	The same as <code>foldLeft</code> , but begins at the first element of the collection.
<code>c.reduceRight op</code>	The same as <code>foldRight</code> , but begins at the last element of the collection.
<code>c.reverse</code>	Returns a collection with the elements in reverse order. (Not available on <code>Traversable</code> , but common to most collections, from <code>GenSeqLike</code> .)
<code>c.size</code>	Returns the size of the collection.
<code>c.slice(from, to)</code>	Returns the interval of elements beginning at element <code>from</code> and ending at element <code>to</code> .
<code>c.sortWith f</code>	Returns a version of the collection sorted by the comparison function <code>f</code> .
<code>c.span p</code>	Returns a collection of two collections; the first created by <code>c.takeWhile(p)</code> , and the second created by <code>c.dropWhile(p)</code> .
<code>c.splitAt n</code>	Returns a collection of two collections by splitting the collection <code>c</code> at element <code>n</code> .
<code>c.sum</code>	Returns the sum of all elements in the collection.
<code>c.tail</code>	Returns all elements from the collection except the first element.
<code>c.take n</code>	Returns the first <code>n</code> elements of the collection.
<code>c.takeWhile p</code>	Returns elements from the collection while the predicate is <code>true</code> . Stops when the predicate becomes <code>false</code> .
<code>c1.union c2</code>	Returns the union (all elements) of two collections.
<code>c.unzip</code>	The opposite of <code>zip</code> , breaks a collection into two collections by dividing each element into two pieces, as in breaking up a collection of <code>Tuple2</code> elements.
<code>c.view</code>	Returns a nonstrict (lazy) view of the collection.
<code>c1.zip c2</code>	Creates a collection of pairs by matching the element 0 of <code>c1</code> with element 0 of <code>c2</code> , element 1 of <code>c1</code> with element 1 of <code>c2</code> , etc.
<code>c.zipWithIndex</code>	Zips the collection with its indices.

Mutable collection methods

Table 10-9 shows the common methods for mutable collections. (Although these are all methods, they're often referred to as operators, because that's what they look like.)

Table 10-9. Common operators (methods) on mutable collections

Operator (method)	Description
<code>c += x</code>	Adds the element <code>x</code> to the collection <code>c</code> .
<code>c += (x, y, z)</code>	Adds the elements <code>x</code> , <code>y</code> , and <code>z</code> to the collection <code>c</code> .
<code>c1 += c2</code>	Adds the elements in the collection <code>c2</code> to the collection <code>c1</code> .
<code>c -= x</code>	Removes the element <code>x</code> from the collection <code>c</code> .
<code>c -= (x, y, z)</code>	Removes the elements <code>x</code> , <code>y</code> , and <code>z</code> from the collection <code>c</code> .

Operator (method)	Description
<code>c1 --= c2</code>	Removes the elements in the collection <code>c2</code> from the collection <code>c1</code> .
<code>c(n) = x</code>	Assigns the value <code>x</code> to the element <code>c(n)</code> .
<code>c clear</code>	Removes all elements from the collection.
<code>c remove n</code> <code>c.remove(n, len)</code>	Removes the element at position <code>n</code> , or the elements beginning at position <code>n</code> and continuing for length <code>len</code> .

There are additional methods, but these are the most common. See the Scaladoc for the mutable collection you're working with for more methods.

Immutable collection operators

Table 10-10 shows the common methods for working with immutable collections. Note that immutable collections can't be modified, so the result of each expression in the first column must be assigned to a new variable. (Also, see **Recipe 10.6** for details on using a mutable variable with an immutable collection.)

Table 10-10. Common operators (methods) on immutable collections

Operator (method)	Description
<code>c1 ++ c2</code>	Creates a new collection by appending the elements in the collection <code>c2</code> to the collection <code>c1</code> .
<code>c :+ e</code>	Returns a new collection with the element <code>e</code> appended to the collection <code>c</code> .
<code>e +: c</code>	Returns a new collection with the element <code>e</code> prepended to the collection <code>c</code> .
<code>e :: list</code>	Returns a <code>List</code> with the element <code>e</code> prepended to the <code>List</code> named <code>list</code> . (<code>::</code> works only on <code>List</code> .)
<code>c drop n</code>	The two methods <code>-</code> and <code>--</code> have been deprecated, so use the filtering methods listed in Table 10-8 to return a new collection with the desired elements removed. Examples of some of these filtering methods are shown here.
<code>c dropWhile p</code>	
<code>c filter p</code>	
<code>c filterNot p</code>	
<code>c head</code>	
<code>c tail</code>	
<code>c take n</code>	
<code>c takeWhile p</code>	

Again, this table lists only the most common methods available on immutable collections. There are other methods available, such as the `--` method on a `Set`. See the Scaladoc for your current collection for even more methods.

Maps

Maps have additional methods, as shown in **Table 10-11**. In this table, the following symbols are used in the first column:

- `m` refers to a map
- `mm` refers to a mutable map
- `k` refers to a key

- `p` refers to a predicate (a function that returns `true` or `false`)
- `v` refers to a map value
- `c` refers to a collection

Table 10-11. Common methods for immutable and mutable maps

Map method	Description
Methods for immutable maps	
<code>m - k</code>	Returns a map with the key <code>k</code> (and its corresponding value) removed.
<code>m - (k1, k2, k3)</code>	Returns a map with the keys <code>k1</code> , <code>k2</code> , and <code>k3</code> removed.
<code>m -- c</code>	Returns a map with the keys in the collection removed. (Although <code>List</code> is shown, this can be any sequential collection.)
<code>m -- List(k1, k2)</code>	
Methods for mutable maps	
<code>mm += (k -> v)</code>	Add the key/value pair(s) to the mutable map <code>mm</code> .
<code>mm += (k1 -> v1, k2 -> v2)</code>	
<code>mm ++= c</code>	Add the elements in the collection <code>c</code> to the mutable map <code>mm</code> .
<code>mm ++= List(3 -> "c")</code>	
<code>mm -= k</code>	Remove map entries from the mutable map <code>mm</code> based on the given key(s).
<code>mm -= (k1, k2, k3)</code>	
<code>mm -= c</code>	Remove the map entries from the mutable map <code>mm</code> based on the keys in the collection <code>c</code> .
Methods for both mutable and immutable maps	
<code>m(k)</code>	Returns the value associated with the key <code>k</code> .
<code>m contains k</code>	Returns <code>true</code> if the map <code>m</code> contains the key <code>k</code> .
<code>m filter p</code>	Returns a map whose keys and values match the condition of the predicate <code>p</code> .
<code>m filterKeys p</code>	Returns a map whose keys match the condition of the predicate <code>p</code> .
<code>m get k</code>	Returns the value for the key <code>k</code> as <code>Some[A]</code> if the key is found, <code>None</code> otherwise.
<code>m getOrElse(k, d)</code>	Returns the value for the key <code>k</code> if the key is found, otherwise returns the default value <code>d</code> .
<code>m isDefinedAt k</code>	Returns <code>true</code> if the map contains the key <code>k</code> .
<code>m keys</code>	Returns the keys from the map as an <code>Iterable</code> .
<code>m keyIterator</code>	Returns the keys from the map as an <code>Iterator</code> .
<code>m keySet</code>	Returns the keys from the map as a <code>Set</code> .
<code>m mapValues f</code>	Returns a new map by applying the function <code>f</code> to every value in the initial map.
<code>m values</code>	Returns the values from the map as an <code>Iterable</code> .
<code>m valuesIterator</code>	Returns the values from the map as an <code>Iterator</code> .

For additional methods, see the Scaladoc for the [mutable](#) and [immutable](#) map classes.

Discussion

As you can see, Scala collection classes contain a wealth of methods (and methods that appear to be operators). Understanding these methods will help you become more productive, because as you understand them, you’ll write less code and fewer loops, and instead write short functions and predicates to work with these methods.

10.4. Understanding the Performance of Collections

Problem

When choosing a collection for an application where performance is extremely important, you want to choose the right collection for the algorithm.

Solution

In many cases, you can reason about the performance of a collection by understanding its basic structure. For instance, a `List` is a singly linked list. It’s not indexed, so if you need to access the one-millionth element of a `List` as `list(1000000)`, that will be slower than accessing the one-millionth element of an `Array`, because the `Array` is indexed, whereas accessing the element in the `List` requires traversing the length of the `List`.

In other cases, it may help to look at the tables. For instance, [Table 10-13](#) shows that the append operation on a `Vector` is `eC`, “effectively constant time.” As a result, I know I can create a large `Vector` in the REPL very quickly like this:

```
var v = Vector[Int]()
for (i <- 1 to 50000) v = v :+ i
```

However, as the table shows, the append operation on a `List` requires linear time, so attempting to create a `List` of the same size takes a much (much!) longer time.

With permission from [EFPL](#), the tables in this recipe have been reproduced from [scala-lang.org](#).

Before looking at the performance tables, [Table 10-12](#) shows the performance characteristic keys that are used in the other tables that follow.

Table 10-12. Performance characteristic keys for the subsequent tables

Key	Description
C	The operation takes (fast) constant time.
eC	The operation takes effectively constant time, but this might depend on some assumptions, such as maximum length of a vector, or distribution of hash keys.
aC	The operation takes amortized constant time. Some invocations of the operation might take longer, but if many operations are performed, on average only constant time per operation is taken.
Log	The operation takes time proportional to the logarithm of the collection size.

Key	Description
L	The operation is linear, so the time is proportional to the collection size.
-	The operation is not supported.

Table 10-13 shows the performance characteristics for operations on immutable and mutable sequential collections.

Table 10-13. Performance characteristics for sequential collections

	head	tail	apply	update	prepend	append	insert
Immutable							
List	C	C	L	L	C	L	-
Stream	C	C	L	L	C	L	-
Vector	eC	eC	eC	eC	eC	eC	-
Stack	C	C	L	L	C	C	L
Queue	aC	aC	L	L	L	C	-
Range	C	C	C	-	-	-	-
String	C	L	C	L	L	L	-
Mutable							
ArrayBuffer	C	L	C	C	L	aC	L
ListBuffer	C	L	L	L	C	C	L
StringBuilder	C	L	C	C	L	aC	L
MutableList	C	L	L	L	C	C	L
Queue	C	L	L	L	C	C	L
ArraySeq	C	L	C	C	-	-	-
Stack	C	L	L	L	C	L	L
ArrayStack	C	L	C	C	aC	L	L
Array	C	L	C	C	-	-	-

Table 10-14 describes the column headings used in **Table 10-13**.

*Table 10-14. Descriptions of the column headings for **Table 10-13***

Operation	Description
head	Selecting the first element of the sequence.
tail	Producing a new sequence that consists of all elements of the sequence except the first one.
apply	Indexing.
update	Functional update for immutable sequences, side-effecting update (with update) for mutable sequences.

Operation	Description
prepend	Adding an element to the front of the sequence. For immutable sequences, this produces a new sequence. For mutable sequences, it modifies the existing sequence.
append	Adding an element at the end of the sequence. For immutable sequences, this produces a new sequence. For mutable sequences, it modifies the existing sequence.
insert	Inserting an element at an arbitrary position in the sequence. This is supported directly only for mutable sequences.

Map and set performance characteristics

Table 10-15 shows the performance characteristics for maps and sets.

Table 10-15. The performance characteristics for maps and sets

	lookup	add	remove	min
Immutable				
HashSet/HashMap	eC	eC	eC	L
TreeSet/TreeMap	Log	Log	Log	Log
BitSet	C	L	L	eC
ListMap	L	L	L	L
Mutable				
HashSet/HashMap	eC	eC	eC	L
WeakHashMap	eC	eC	eC	L
BitSet	C	aC	C	eC
TreeSet	Log	Log	Log	Log

Table 10-16 provides descriptions for the column headings used in **Table 10-15**.

Table 10-16. Descriptions of the column headings used in Table 10-15

Operation	Description
lookup	Testing whether an element is contained in a set, or selecting a value associated with a map key.
add	Adding a new element to a set or key/value pair to a map.
remove	Removing an element from a set or a key from a map.
min	The smallest element of the set, or the smallest key of a map.

See Also

- The tables in this recipe have been reproduced from the following URL, with permission from the [Programming Methods Laboratory of EFPL](#).
- [The Programming Methods Laboratory of EFPL](#).

10.5. Declaring a Type When Creating a Collection

Problem

You want to create a collection of mixed types, and Scala isn't automatically assigning the type you want.

Solution

In the following example, if you don't specify a type, Scala automatically assigns a type of `Double` to the list:

```
scala> val x = List(1, 2.0, 33D, 400L)
x: List[Double] = List(1.0, 2.0, 33.0, 400.0)
```

If you'd rather have the collection be of type `AnyVal` or `Number`, specify the type in brackets before your collection declaration:

```
scala> val x = List[Number](1, 2.0, 33D, 400L)
x: List[java.lang.Number] = List(1, 2.0, 33.0, 400)
```

```
scala> val x = List[AnyVal](1, 2.0, 33D, 400L)
x: List[AnyVal] = List(1, 2.0, 33.0, 400)
```

Discussion

By manually specifying a type, in this case `Number`, you control the collection type. This is useful any time a list contains mixed types or multiple levels of inheritance. For instance, given this type hierarchy:

```
trait Animal
trait FurryAnimal extends Animal
case class Dog(name: String) extends Animal
case class Cat(name: String) extends Animal
```

create a sequence with a `Dog` and a `Cat`:

```
scala> val x = Array(Dog("Fido"), Cat("Felix"))
x: Array[Product with Serializable with Animal] = Array(Dog(Fido), Cat(Felix))
```

As shown, Scala assigns a type of `Product with Serializable with Animal`. If you just want an `Array[Animal]`, manually specify the desired type:

```
scala> val x = Array[Animal](Dog("Fido"), Cat("Felix"))
x: Array[Animal] = Array(Dog(Fido), Cat(Felix))
```

This may not seem like a big deal, but imagine declaring a class with a method that returns this array:

```
class AnimalKingdom {
  def animals = Array(Dog("Fido"), Cat("Felix"))
}
```

When you generate the Scaladoc for this class, the `animals` method will show the “Product with Serializable” in its Scaladoc:

```
def animals: Array[Product with Serializable with Animal]
```

If you’d rather have it appear like this in your Scaladoc:

```
def animals: Array[Animal]
```

manually assign the type, as shown in the Solution:

```
def animals = Array[Animal](Dog("Fido"), Cat("Felix"))
```

10.6. Understanding Mutable Variables with Immutable Collections

Problem

You may have seen that mixing a mutable variable (`var`) with an immutable collection causes surprising behavior. For instance, when you create an immutable `Vector` as a `var`, it appears you can somehow add new elements to it:

```
scala> var sisters = Vector("Melinda")
sisters: collection.immutable.Vector[String] = Vector(Melinda)

scala> sisters = sisters :+ "Melissa"
sisters: collection.immutable.Vector[String] = Vector(Melinda, Melissa)

scala> sisters = sisters :+ "Marisa"
sisters: collection.immutable.Vector[String] = Vector(Melinda, Melissa, Marisa)

scala> sisters.foreach(println)
Melinda
Melissa
Marisa
```

How can this be?

Solution

Though it looks like you’re mutating an immutable collection, what’s really happening is that the `sisters` variable points to a new collection each time you use the `:+` method. The `sisters` variable is mutable—like a non-`final` field in Java—so it’s actually being reassigned to a new collection during each step. The end result is similar to these lines of code:

```
var sisters = Vector("Melinda")
sisters = Vector("Melinda", "Melissa")
sisters = Vector("Melinda", "Melissa", "Marisa")
```

In the second and third lines of code, the `sisters` reference has been changed to point to a new collection.

You can demonstrate that the vector itself is immutable. Attempting to mutate one of its elements—which doesn’t involve reassigning the variable—results in an error:

```
scala> sisters(0) = "Molly"
<console>:12: error: value update is not a member of
scala.collection.immutable.Vector[String]
      sisters(0) = "Molly"
              ^
```

Summary

When you first start working with Scala, the behavior of a mutable variable with an immutable collection can be surprising. To be clear about *variables*:

- A mutable variable (`var`) can be reassigned to point at new data.
- An immutable variable (`val`) is like a `final` variable in Java; it can never be reassigned.

To be clear about *collections*:

- The elements in a mutable collection (like `ArrayBuffer`) can be changed.
- The elements in an immutable collection (like `Vector`) cannot be changed.

See Also

[Recipe 20.2, “Prefer Immutable Objects”](#), discusses the use of mutable variables with immutable collections, and its opposite, using immutable variables with mutable collections as a “best practice.”

10.7. Make Vector Your “Go To” Immutable Sequence

Problem

You want a fast, general-purpose, immutable, sequential collection type for your Scala applications.

Solution

The `Vector` class was introduced in Scala 2.8 and is now considered to be the “go to,” general-purpose immutable data structure. (`Vector` is an indexed, immutable sequential collection. Use a `List` if you prefer working with a linear, immutable sequential collection. See [Recipe 10.2, “Choosing a Collection Class”](#), for more details.)

Create and use a `Vector` just like other immutable, indexed sequences. You can create them and access elements efficiently by index:

```
scala> val v = Vector("a", "b", "c")
v: scala.collection.immutable.Vector[java.lang.String] = Vector(a, b, c)

scala> v(0)
res0: java.lang.String = a
```

You can't modify a vector, so you “add” elements to an existing vector as you assign the result to a new variable:

```
scala> val a = Vector(1, 2, 3)
a: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3)

scala> val b = a ++ Vector(4, 5)
b: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3, 4, 5)
```

Use the updated method to replace one element in a `Vector` while assigning the result to a new variable:

```
scala> val c = b.updated(0, "x")
c: scala.collection.immutable.Vector[java.lang.String] = Vector(x, b, c)
```

You can also use all the usual filtering methods to get just the elements you want out of a vector:

```
scala> val a = Vector(1, 2, 3, 4, 5)
a: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3, 4, 5)

scala> val b = a.take(2)
b: scala.collection.immutable.Vector[Int] = Vector(1, 2)

scala> val c = a.filter(_ > 2)
c: scala.collection.immutable.Vector[Int] = Vector(3, 4, 5)
```

In those examples, I created each variable as a `val` and assigned the output to a new variable just to be clear, but you can also declare your variable as a `var` and reassign the result back to the same variable:

```
scala> var a = Vector(1, 2, 3)
a: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3)

scala> a = a ++ Vector(4, 5)
a: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3, 4, 5)
```

Discussion

The “concrete, immutable collections classes” page from the scala-lang.org website states the following:

`Vector` is a collection type (introduced in Scala 2.8) that addresses the inefficiency for random access on lists. Vectors allow accessing any element of the list in ‘effectively’ constant time ... Because vectors strike a good balance between fast random selections and fast random functional updates, they are currently the default implementation of immutable indexed sequences...

In his book, *Scala In Depth* (Manning Publications), Joshua Suereth offers the rule, “When in Doubt, Use Vector.” He writes, “Vector is the most flexible, efficient collection in the Scala collections library.”

As noted in [Recipe 10.1](#), if you create an instance of an `IndexedSeq`, Scala returns a `Vector`:

```
scala> val x = IndexedSeq(1,2,3)
x: IndexedSeq[Int] = Vector(1, 2, 3)
```

As a result, I’ve seen some developers create an `IndexedSeq` in their code, rather than a `Vector`, to be more generic and to allow for potential future changes.

See Also

- [The `Vector` class](#)
- The “concrete, immutable collections classes” discussion of the [Vector class](#)

10.8. Make `ArrayBuffer` Your “Go To” Mutable Sequence

Problem

You want to use a general-purpose, mutable sequence in your Scala applications.

Solution

Just as the `Vector` is the recommended “go to” class for immutable, sequential collections, the `ArrayBuffer` class is recommended as the general-purpose class for *mutable* sequential collections. (`ArrayBuffer` is an indexed sequential collection. Use `ListBuffer` if you prefer a linear sequential collection that is mutable. See [Recipe 10.2](#), “Choosing a Collection Class”, for more information.)

To use an `ArrayBuffer`, first import it:

```
import scala.collection.mutable.ArrayBuffer
```

You can then create an empty `ArrayBuffer`:

```
var fruits = ArrayBuffer[String]()
var ints = ArrayBuffer[Int]()
```

Or you can create an `ArrayBuffer` with initial elements:

```
var nums = ArrayBuffer(1, 2, 3)
```

Like other mutable collection classes, you add elements using the `+=` and `++=` methods:

```
scala> var nums = ArrayBuffer(1, 2, 3)
nums: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3)

// add one element
scala> nums += 4
res0: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3, 4)

// add two or more elements (method has a varargs parameter)
scala> nums += (5, 6)
res1: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3, 4, 5, 6)

// add elements from another collection
scala> nums ++= List(7, 8)
res2: scala.collection.mutable.ArrayBuffer[Int] =
  ArrayBuffer(1, 2, 3, 4, 5, 6, 7, 8)
```

You remove elements with the `-=` and `--=` methods:

```
// remove one element
scala> nums -= 9
res3: scala.collection.mutable.ArrayBuffer[Int] =
  ArrayBuffer(1, 2, 3, 4, 5, 6, 7, 8)

// remove two or more elements
scala> nums -= (7, 8)
res4: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3, 4, 5, 6)

// remove elements specified by another sequence
scala> nums --= Array(5, 6)
res5: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3, 4)
```

Discussion

Those are the methods I generally use to add and remove elements from an `ArrayBuffer`. However, there are many more:

```
val a = ArrayBuffer(1, 2, 3) // ArrayBuffer(1, 2, 3)
a.append(4) // ArrayBuffer(1, 2, 3, 4)
a.append(5, 6) // ArrayBuffer(1, 2, 3, 4, 5, 6)
a.appendAll(Seq(7,8)) // ArrayBuffer(1, 2, 3, 4, 5, 6, 7, 8)
a.clear // ArrayBuffer()

val a = ArrayBuffer(9, 10) // ArrayBuffer(9, 10)
a.insert(0, 8) // ArrayBuffer(8, 9, 10)
a.insert(0, 6, 7) // ArrayBuffer(6, 7, 8, 9, 10)
a.insertAll(0, Vector(4, 5)) // ArrayBuffer(4, 5, 6, 7, 8, 9, 10)
a.prepend(3) // ArrayBuffer(3, 4, 5, 6, 7, 8, 9, 10)
a.prepend(1, 2) // ArrayBuffer(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
a.prependAll(Array(0)) // ArrayBuffer(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

```

val a = ArrayBuffer.range('a', 'h') // ArrayBuffer(a, b, c, d, e, f, g)
a.remove(0)                          // ArrayBuffer(b, c, d, e, f, g)
a.remove(2, 3)                       // ArrayBuffer(b, c, g)

val a = ArrayBuffer.range('a', 'h') // ArrayBuffer(a, b, c, d, e, f, g)
a.trimStart(2)                       // ArrayBuffer(c, d, e, f, g)
a.trimEnd(2)                         // ArrayBuffer(c, d, e)

```

See [the Scaladoc](#) for more methods that you can use to modify an `ArrayBuffer`.

The `ArrayBuffer` Scaladoc provides these details about `ArrayBuffer` performance: “Append, update, and random access take constant time (amortized time). Prepends and removes are linear in the buffer size.” The `ArrayBuffer` documentation also states, “array buffers are useful for efficiently building up a large collection whenever the new items are always added to the end.”

If you need a mutable sequential collection that works more like a `List` (i.e., a linear sequence rather than an indexed sequence), use `ListBuffer` instead of `ArrayBuffer`. The Scala documentation on the `ListBuffer` states, “A `ListBuffer` is like an array buffer except that it uses a linked list internally instead of an array. If you plan to convert the buffer to a list once it is built up, use a list buffer instead of an array buffer.”

See Also

- [ArrayBuffer discussion](#)
- [ArrayBuffer Scaladoc](#)
- [ListBuffer discussion](#)

10.9. Looping over a Collection with `foreach`

Problem

You want to iterate over the elements in a collection with the `foreach` method.

Solution

The `foreach` method takes a function as an argument. The function you define should take an element as an input parameter, and should not return anything. The input parameter type should match the type stored in the collection. As `foreach` executes, it passes one element at a time from the collection to your function until it reaches the last element in the collection.

The `foreach` method applies your function to each element of the collection, but it doesn't return a value. Because it doesn't return anything, it's said that it's used for its "side effect."

As an example, a common use of `foreach` is to output information:

```
scala> val x = Vector(1, 2, 3)
x: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3)

scala> x.foreach((i: Int) => println(i))
1
2
3
```

That's the longhand way of writing that code. For most expressions, Scala can infer the type, so specifying `i: Int` isn't necessary:

```
args.foreach(i => println(i))
```

You can further shorten this expression by using the ubiquitous underscore wildcard character instead of using a temporary variable:

```
args.foreach(println(_))
```

In a situation like this, where a function literal consists of one statement that takes a single argument, it can be condensed to this form:

```
args.foreach(println)
```

For a simple case like this, the syntax in the last example is typically used.

Discussion

As long as your function (or method) takes one parameter of the same type as the elements in the collection and returns nothing (`Unit`), it can be called from a `foreach` method. In the following example, the `printIt` method takes a `Char`, does something with it, and returns nothing:

```
def printIt(c: Char) { println(c) }
```

Because a `String` is a sequence of type `Char`, `printIt` can be called in a `foreach` method on a `String` as follows:

```
"HAL".foreach(c => printIt(c))
"HAL".foreach(printIt)
```

If your algorithm is used only once, you don't have to declare it as a method or function; just pass it to `foreach` as a function literal:

```
"HAL".foreach((c: Char) => println(c))
```

To declare a multiline function, use this format:

```
val longWords = new StringBuilder

"Hello world it's Al".split(" ").foreach{ e =>
  if (e.length > 4) longWords.append(s" $e")
  else println("Not added: " + e)
}
```

To understand this example, it may be helpful to know the `split` method used in that function creates an `Array[String]`, as shown here:

```
scala> "Hello world it's Al".split(" ")
res0: Array[java.lang.String] = Array(Hello, world, it's, Al)
```

In addition to using the `foreach` method on sequential collections, it's also available on the `Map` class. The `Map` implementation of `foreach` passes two parameters to your function. You can handle those parameters as a tuple:

```
val m = Map("fname" -> "Tyler", "lname" -> "LeDude")
m foreach (x => println(s"${x._1} -> ${x._2}"))
```

However, I generally prefer the following approach:

```
movieRatings.foreach {
  case(movie, rating) => println(s"key: $movie, value: $rating")
}
```

See [Recipe 11.17, “Traversing a Map”](#), for other ways to iterate over a map.

Scala's `for` loop provides another powerful way to iterate over the elements in a collection. See [Recipe 10.10, “Looping over a Collection with a for Loop”](#), for more information.

10.10. Looping over a Collection with a for Loop

Problem

You want to loop over the elements in a collection using a `for` loop, possibly creating a new collection from the existing collection using the `for/yield` combination.

Solution

You can loop over any `Traversable` type (basically any sequence) using a `for` loop:

```
scala> val fruits = Traversable("apple", "banana", "orange")
fruits: Traversable[String] = List(apple, banana, orange)

scala> for (f <- fruits) println(f)
apple
banana
orange
```

```
scala> for (f <- fruits) println(f.toUpperCase)
APPLE
BANANA
ORANGE
```

If your algorithm is long, perform the work in a block following a for loop:

```
scala> val fruits = Array("apple", "banana", "orange")
fruits: Array[String] = Array(apple, banana, orange)
```

```
scala> for (f <- fruits) {
  |   // imagine this required multiple lines
  |   val s = f.toUpperCase
  |   println(s)
  | }
APPLE
BANANA
ORANGE
```

This example shows one approach to using a counter inside a for loop:

```
scala> for (i <- 0 until fruits.size) println(s"element $i is ${fruits(i)}")
element 0 is apple
element 1 is banana
element 2 is orange
```

You can also use the `zipWithIndex` method when you need a loop counter:

```
scala> for ((elem, count) <- fruits.zipWithIndex) {
  |   println(s"element $count is $elem")
  | }
element 0 is apple
element 1 is banana
element 2 is orange
```



When using `zipWithIndex`, consider calling `view` before `zipWithIndex`:

```
// added a call to 'view'
for ((elem, count) <- fruits.view.zipWithIndex) {
  println(s"element $count is $elem")
}
```

See the next recipe for details.

Using `zip` with a `Stream` is another way to generate a counter:

```
scala> for ((elem, count) <- fruits.zip(Stream from 1)) {
  |   println(s"element $count is $elem")
  | }
element 1 is apple
element 2 is banana
element 3 is orange
```

See the next recipe for details on using `zipWithIndex` and `zip` to create loop counters.

If you just need to do something *N* times, using a `Range` works well:

```
scala> for (i <- 1 to 3) println(i)
1
2
3
```

In that example, the expression `1 to 3` creates a `Range`, which you can demonstrate in the REPL:

```
scala> 1 to 3
res0: scala.collection.immutable.Range.Inclusive = Range(1, 2, 3)
```

Again you can use a block inside curly braces when your algorithm gets long:

```
scala> for (i <- 1 to 3) {
  |   // do whatever you want in this block
  |   println(i)
  | }
1
2
3
```

The `for/yield` construct

The previous examples show how to operate on each element in a sequence, but they don't return a value. As with the `foreach` examples in the previous recipe, they're used for their side effect.

To build a new collection from an input collection, use the `for/yield` construct. The following example shows how to build a new array of uppercase strings from an input array of lowercase strings:

```
scala> val fruits = Array("apple", "banana", "orange")
fruits: Array[java.lang.String] = Array(apple, banana, orange)

scala> val newArray = for (e <- fruits) yield e.toUpperCase
newArray: Array[java.lang.String] = Array(APPLE, BANANA, ORANGE)
```

The `for/yield` construct returns (yields) a new collection from the input collection by applying your algorithm to the elements of the input collection, so the array `newArray` contains uppercase versions of the three strings in the initial array. Using `for/yield` like this is known as a *for comprehension*.

If your `for/yield` processing requires multiple lines of code, perform the work in a block after the `yield` keyword:


```
scala> val newArray = for (fruit <- fruits) yield {
  |   // imagine this required multiple lines
  |   val upper = fruit.toUpperCase
  |   upper
  | }
newArray: Array[java.lang.String] = Array(APPLE, BANANA, ORANGE)
```

If your algorithm is long, or you want to reuse it, first define it in a method (or function):

```
def upperReverse(s: String) = {
  // imagine this is a long algorithm
  s.toUpperCase.reverse
}
```

then use the method with the for/yield loop:

```
scala> val newArray = for (fruit <- fruits) yield upperReverse(fruit)
newArray: Array[String] = Array(ELPPA, ANANAB, EGNARO)
```

Maps

You can also iterate over a Map nicely using a for loop:

```
scala> val names = Map("fname" -> "Ed", "lname" -> "Chigliak")
names: scala.collection.immutable.Map[String,String] =
  Map(fname -> Ed, lname -> Chigliak)

scala> for ((k,v) <- names) println(s"key: $k, value: $v")
key: fname, value: Ed
key: lname, value: Chigliak
```

See [Recipe 11.17, “Traversing a Map”](#), for more examples of iterating over a map.

Discussion

When using a for loop, the <- symbol can be read as “in,” so the following statement can be read as “for i in 1 to 3, do ...”:

```
for (i <- 1 to 3) { // more code here ...
```

As demonstrated in [Recipe 3.3, “Using a for Loop with Embedded if Statements \(Guards\)”](#), you can also combine a for loop with if statements, which are known as *guards*:

```
for {
  file <- files
  if file.isFile
  if file.getName.endsWith(".txt")
} doSomething(file)
```

See that recipe for more examples of using guards with for loops.

See Also

- [Recipe 3.3, “Using a for Loop with Embedded if Statements \(Guards\)”](#)
- [Recipe 10.9, “Looping over a Collection with foreach”](#)
- [Recipe 10.13, “Transforming One Collection to Another with for/yield”](#)

10.11. Using zipWithIndex or zip to Create Loop Counters

Problem

You want to loop over a sequential collection, and you’d like to have access to a counter in the loop, without having to manually create a counter.

Solution

Use the `zipWithIndex` or `zip` methods to create a counter automatically. Assuming you have a sequential collection of days:

```
val days = Array("Sunday", "Monday", "Tuesday", "Wednesday",  
                 "Thursday", "Friday", "Saturday")
```

you can print the elements in the collection with a counter using the `zipWithIndex` and `foreach` methods:

```
days.zipWithIndex.foreach {  
  case(day, count) => println(s"$count is $day")  
}
```

As you’ll see in the Discussion, this works because `zipWithIndex` returns a series of `Tuple2` elements in an `Array`, like this:

```
Array((Sunday,0), (Monday,1), ...
```

and the `case` statement in the `foreach` loop matches a `Tuple2`.

You can also use `zipWithIndex` with a `for` loop:

```
for ((day, count) <- days.zipWithIndex) {  
  println(s"$count is $day")  
}
```

Both loops result in the following output:

```
0 is Sunday  
1 is Monday  
2 is Tuesday  
3 is Wednesday  
4 is Thursday
```

```
5 is Friday
6 is Saturday
```

When using `zipWithIndex`, the counter always starts at 0. You can also use the `zip` method with a `Stream` to create a counter. This gives you a way to control the starting value:

```
scala> for ((day,count) <- days.zip(Stream from 1)) {
  |   println(s"day $count is $day")
  | }
```

Discussion

When `zipWithIndex` is used on a sequence, it returns a sequence of `Tuple2` elements, as shown in this example:

```
scala> val list = List("a", "b", "c")
list: List[String] = List(a, b, c)

scala> val zwi = list.zipWithIndex
zwi: List[(String, Int)] = List((a,0), (b,1), (c,2))
```

Because `zipWithIndex` creates a new sequence from the existing sequence, you may want to call `view` before invoking `zipWithIndex`, like this:

```
scala> val zwi2 = list.view.zipWithIndex
zwi2: scala.collection.SeqView[(String, Int),Seq[_]] = SeqViewZ(...)
```

As shown, this creates a lazy view on the original list, so the tuple elements won't be created until they're needed. Because of this behavior, calling `view` before calling `zipWithIndex` is recommended at the first two links in the See Also section. However, my own experience concurs with the performance shown in the third link in the See Also section, where *not* using a view performs better. If performance is a concern, try your loop both ways, and also try manually incrementing a counter.

As mentioned, the `zip` and `zipWithIndex` methods both return a sequence of `Tuple2` elements. Therefore, your `foreach` method can also look like this:

```
days.zipWithIndex.foreach { d =>
  println(s"${d._2} is ${d._1}")
}
```

However, I think the approaches shown in the Solution are more readable.

As shown in the previous recipe, you can also use a range with a `for` loop to create a counter:

```
val fruits = Array("apple", "banana", "orange")
for (i <- 0 until fruits.size) println(s"element $i is ${fruits(i)}")
```

See [Recipe 10.24, “Creating a Lazy View on a Collection”](#), for more information on using views.

See Also

- A blog post on using `zipWithIndex` in [several use cases](#)
- A discussion of using `zipWithIndex` in a [for loop](#)
- A discussion of performance related to using a view with `zipWithIndex`
- [SeqView trait](#)

10.12. Using Iterators

Problem

You want (or need) to work with an iterator in a Scala application.

Solution

Although using an iterator with `hasNext()` and `next()` is a common way to loop over a collection in Java, they aren't commonly used in Scala, because Scala collections have methods like `map` and `foreach` that let you implement algorithms more concisely. To be clear, in Scala, I've never directly written code like this:

```
// don't do this
val it = collection.iterator
while (it.hasNext) ...
```

That being said, sometimes you'll run into an iterator, with one of the best examples being the `io.Source.fromFile` method. This method returns an iterator, which makes sense, because when you're working with very large files, it's not practical to read the entire file into memory.

An important part of using an iterator is knowing that it's exhausted after you use it. As you access each element, you mutate the iterator, and the previous element is discarded. For instance, if you use `foreach` to iterate over an iterator's elements, the call works the first time:

```
scala> val it = Iterator(1,2,3)
it: Iterator[Int] = non-empty iterator

scala> it.foreach(println)
1
2
3
```

But when you attempt the same call a second time, you won't get any output, because the iterator has been exhausted:

```
scala> it.foreach(println)
(no output here)
```

An iterator isn't a collection; instead, it gives you a way to access the elements in a collection, one by one. But an iterator does define many of the methods you'll see in a normal collection class, including `foreach`, `map`, `flatMap`, `collect`, etc. You can also convert an iterator to a collection when needed:

```
val it = Iterator(1,2,3)
it.toArray
```

The REPL output shows the collections you can create from an iterator:

```
scala> it.to[Tab]
toArray      toBuffer      toIndexedSeq  toIterable    toIterator
toList      toMap         toSeq         toSet         toStream
toString    toTraversable
```

See Also

- An [introduction to Scala iterators](#)
- The [Iterator trait](#)

10.13. Transforming One Collection to Another with `for`/`yield`

Problem

You want to create a new collection from an existing collection by transforming the elements with an algorithm.

Solution

Use the `for`/`yield` construct and your algorithm to create the new collection. For instance, starting with a basic collection:

```
scala> val a = Array(1, 2, 3, 4, 5)
a: Array[Int] = Array(1, 2, 3, 4, 5)
```

You can create a copy of that collection by just “yielding” each element (with no algorithm):

```
scala> for (e <- a) yield e
res0: Array[Int] = Array(1, 2, 3, 4, 5)
```

You can create a new collection where each element is twice the value of the original:

```
scala> for (e <- a) yield e * 2
res1: Array[Int] = Array(2, 4, 6, 8, 10)
```

You can determine the modulus of each element:

```
scala> for (e <- a) yield e % 2
res2: Array[Int] = Array(1, 0, 1, 0, 1)
```

This example converts a list of strings to uppercase:

```
scala> val fruits = Vector("apple", "banana", "lime", "orange")
fruits: Vector[String] = Vector(apple, banana, lime, orange)

scala> val ucFruits = for (e <- fruits) yield e.toUpperCase
ucFruits: Vector[String] = Vector(APPLE, BANANA, LIME, ORANGE)
```

Your algorithm can return whatever collection is needed. This approach converts the original collection into a sequence of `Tuple2` elements:

```
scala> for (i <- 0 until fruits.length) yield (i, fruits(i))
res0: scala.collection.immutable.IndexedSeq[(Int, String)] =
  Vector((0,apple), (1,banana), (2,lime), (3,orange))
```

This algorithm yields a sequence of `Tuple2` elements that contains each original string along with its length:

```
scala> for (f <- fruits) yield (f, f.length)
res1: Vector[(String, Int)] = Vector((apple,5), (banana,6), (lime,4), (orange,6))
```

If your algorithm takes multiple lines, include it in a block after the `yield`:

```
scala> val x = for (e <- fruits) yield {
  |   // imagine this required multiple lines
  |   val s = e.toUpperCase
  |   s
  | }
x: Vector[String] = List(APPLE, BANANA, LIME, ORANGE)
```

Given a `Person` class and a list of friend's names like this:

```
case class Person (name: String)
val friends = Vector("Mark", "Regina", "Matt")
```

a `for/yield` loop can yield a collection of `Person` instances:

```
scala> for (f <- friends) yield Person(f)
res0: Vector[Person] = Vector(Person(Mark), Person(Regina), Person(Matt))
```

You can include `if` statements (guards) in a `for` comprehension to filter elements:

```
scala> val x = for (e <- fruits if e.length < 6) yield e.toUpperCase
x: List[java.lang.String] = List(APPLE, LIME)
```

Discussion

This combination of a `for` loop and `yield` statement is known as a *for comprehension* or *sequence comprehension*. It yields a new collection from an existing collection.

If you're new to using the `for/yield` construct, it can help to think that it has a bucket or temporary holding area on the side. As each element from the original collection is operated on with `yield` and your algorithm, it's added to that bucket. Then, when the `for` loop is finished iterating over the entire collection, all of the elements in the bucket are returned (yielded) by the expression.

In general, the collection type that's returned by a `for` comprehension will be the same type that you begin with. If you begin with an `ArrayBuffer`, you'll end up with an `ArrayBuffer`:

```
scala> val fruits = scala.collection.mutable.ArrayBuffer("apple", "banana")
fruits: scala.collection.mutable.ArrayBuffer[java.lang.String] =
  ArrayBuffer(apple, banana)

scala> val x = for (e <- fruits) yield e.toUpperCase
x: scala.collection.mutable.ArrayBuffer[java.lang.String] =
  ArrayBuffer(APPLE, BANANA)
```

A `List` returns a `List`:

```
scala> val fruits = "apple" :: "banana" :: "orange" :: Nil
fruits: List[java.lang.String] = List(apple, banana, orange)

scala> val x = for (e <- fruits) yield e.toUpperCase
x: List[java.lang.String] = List(APPLE, BANANA, ORANGE)
```

However, as shown in the Solution, this isn't always the case.

Using guards

When you add guards to a `for` comprehension and want to write it as a multiline expression, the recommended coding style is to use curly braces rather than parentheses:

```
for {
  file <- files
  if hasSoundFileExtension(file)
  if !soundFileIsLong(file)
} yield file
```

This makes the code more readable, especially when the list of guards becomes long. See [Recipe 3.3, “Using a `for` Loop with Embedded `if` Statements \(Guards\)”](#), more information on using guards.

When using guards, the resulting collection can end up being a different size than the input collection:

```
scala> val cars = Vector("Mercedes", "Porsche", "Tesla")
cars: Vector[String] = Vector(Mercedes, Porsche, Tesla)
```

```
scala> for {
  |   c <- cars
  |   if c.startsWith("M")
  | } yield c
res0: Vector[String] = Vector(Mercedes)
```

In fact, if none of the car names had matched the `startsWith` test, that code would return an empty `Vector`.

When I first started working with Scala I always used a `for/yield` expression to do this kind of work, but one day I realized that I could achieve the same result more concisely using the `map` method. The next recipe demonstrates how to use `map` to create a new collection from an existing collection.

See Also

- [Recipe 3.1, “Looping with `for` and `foreach`”](#), provides detailed examples of how `for` loops are translated by the Scala compiler into `foreach` and `map` method calls.
- [Recipe 3.3, “Using a `for` Loop with Embedded `if` Statements \(Guards\)”](#), provides more examples of using guards.

10.14. Transforming One Collection to Another with `map`

Problem

Like the previous recipe, you want to transform one collection into another by applying an algorithm to every element in the original collection.

Solution

Rather than using the `for/yield` combination shown in the previous recipe, call the `map` method on your collection, passing it a function, an anonymous function, or method to transform each element. This is shown in the following examples, where each `String` in a `List` is converted to begin with a capital letter:

```
scala> val helpers = Vector("adam", "kim", "melissa")
helpers: scala.collection.immutable.Vector[java.lang.String] =
  Vector(adam, kim, melissa)

// the long form
scala> val caps = helpers.map(e => e.capitalize)
caps: scala.collection.immutable.Vector[String] = Vector(Adam, Kim, Melissa)

// the short form
```



```
scala> val caps = helpers.map(_.capitalize)
caps: scala.collection.immutable.Vector[String] = Vector(Adam, Kim, Melissa)
```

The next example shows that an array of `String` can be converted to an array of `Int`:

```
scala> val names = Array("Fred", "Joe", "Jonathan")
names: Array[java.lang.String] = Array(Fred, Joe, Jonathan)

scala> val lengths = names.map(_.length)
lengths: Array[Int] = Array(4, 3, 8)
```

The `map` method comes in handy if you want to convert a collection to a list of XML elements:

```
scala> val nieces = List("Aleka", "Christina", "Molly")
nieces: List[String] = List(Aleka, Christina, Molly)

scala> val elems = nieces.map(niece => <li>{niece}</li>)
elems: List[scala.xml.Elem] =
  List(<li>Aleka</li>, <li>Christina</li>, <li>Molly</li>)
```

Using a similar technique, you can convert the collection directly to an XML literal:

```
scala> val ul = <ul>{nieces.map(i => <li>{i}</li>)}</ul>
ul: scala.xml.Elem = <ul><li>Aleka</li><li>Christina</li><li>Molly</li></ul>
```

A function that's passed into `map` can be as complicated as necessary. An example in the Discussion shows how to use a multiline anonymous function with `map`. When your algorithm gets longer, rather than using an anonymous function, define the function (or method) first, and then pass it into `map`:

```
// imagine this is a long method
scala> def plusOne(c: Char): Char = (c.toByte+1).toChar
plusOne: (c: Char)Char

scala> "HAL".map(plusOne)
res0: String = IBM
```

When writing a method to work with `map`, define the method to take a single parameter that's the same type as the collection. In this case, `plusOne` is defined to take a `Char`, because a `String` is a collection of `Char` elements. The return type of the method can be whatever you need for your algorithm. For instance, the previous `names.map(_.length)` example showed that a function applied to a `String` can return an `Int`.

Unlike the `for/yield` approach shown in the previous recipe, the `map` method also works well when writing a chain of method calls. For instance, you can split a `String` into an array of strings, then trim the blank spaces from those strings:

```
scala> val s = " eggs, milk, butter, Coco Puffs "
s: String = " eggs, milk, butter, Coco Puffs "
```

```
scala> val items = s.split(",").map(_.trim)
items: Array[String] = Array(eggs, milk, butter, Coco Puffs)
```

This works because `split` creates an `Array[String]`, and `map` applies the `trim` method to each element in that array before returning the final array.

Discussion

For simple cases, using `map` is the same as using a basic `for/yield` loop:

```
scala> val people = List("adam", "kim", "melissa")
people: List[java.lang.String] = List(adam, kim, melissa)

// map
scala> val caps1 = people.map(_.capitalize)
caps1: List[String] = List(Adam, Kim, Melissa)

// for/yield
scala> val caps2 = for (f <- people) yield f.capitalize
caps2: List[String] = List(Adam, Kim, Melissa)
```

But once you add a guard, a `for/yield` loop is no longer directly equivalent to just a `map` method call. If you attempt to use an `if` statement in the algorithm you pass to a `map` method, you'll get a very different result:

```
scala> val fruits = List("apple", "banana", "lime", "orange", "raspberry")
fruits: List[java.lang.String] = List(apple, banana, lime, orange, raspberry)

scala> val newFruits = fruits.map( f =>
  |   if (f.length < 6) f.toUpperCase
  | )
newFruits: List[Any] = List(APPLE, (), LIME, (), ())
```

You *could* filter the result *after* calling `map` to clean up the result:

```
scala> newFruits.filter(_ != ())
res0: List[Any] = List(APPLE, LIME)
```

But in this situation, it helps to think of an `if` statement as being a filter, so the correct solution is to first filter the collection, and then call `map`:

```
scala> val fruits = List("apple", "banana", "lime", "orange", "raspberry")
fruits: List[String] = List(apple, banana, lime, orange, raspberry)

scala> fruits.filter(_.length < 6).map(_.toUpperCase)
res1: List[String] = List(APPLE, LIME)
```

See Also

[Recipe 3.1, “Looping with `for` and `foreach`”](#), provides detailed examples of how `for` loops are translated by the Scala compiler into `foreach` and `map` method calls.

10.15. Flattening a List of Lists with flatten

Problem

You have a list of lists (a sequence of sequences) and want to create one list (sequence) from them.

Solution

Use the `flatten` method to convert a list of lists into a single list. To demonstrate this, first create a list of lists:

```
scala> val lol = List(List(1,2), List(3,4))
lol: List[List[Int]] = List(List(1, 2), List(3, 4))
```

Calling the `flatten` method on this list of lists creates one new list:

```
scala> val result = lol.flatten
result: List[Int] = List(1, 2, 3, 4)
```

As shown, `flatten` does what its name implies, flattening the lists held inside the outer list into one resulting list.

Though I use the term “list” here, the `flatten` method isn’t limited to a `List`; it works with other sequences (`Array`, `ArrayBuffer`, `Vector`, etc.) as well:

```
scala> val a = Array(Array(1,2), Array(3,4))
a: Array[Array[Int]] = Array(Array(1, 2), Array(3, 4))

scala> a.flatten
res0: Array[Int] = Array(1, 2, 3, 4)
```

In the real world, you might use `flatten` to convert a list of couples attending a wedding into a single list of all people attending the wedding. Calling `flatten` on a `List[List[String]]` does the job:

```
scala> val couples = List(List("kim", "al"), List("julia", "terry"))
couples: List[List[String]] = List(List(kim, al), List(julia, terry))

scala> val people = couples.flatten
people: List[String] = List(kim, al, julia, terry)
```

If you really want to have fun, capitalize each element in the resulting list and then sort the list:

```
scala> val people = couples.flatten.map(_.capitalize).sorted
people: List[String] = List(Al, Julia, Kim, Terry)
```

This helps to demonstrate the power of the Scala collections methods. (Imagine trying to write that code with only a `for` loop.)

In a social-networking application, you might do the same thing with a list of friends, and their friends:

```
val myFriends = List("Adam", "David", "Frank")
val adamsFriends = List("Nick K", "Bill M")
val davidsFriends = List("Becca G", "Kenny D", "Bill M")
val friendsOfFriends = List(adamsFriends, davidsFriends)
```

Because `friendsOfFriends` is a list of lists, you can use `flatten` to accomplish many tasks with it, such as creating a unique list of the friends of your friends:

```
scala> val uniqueFriendsOfFriends = friendsOfFriends.flatten.distinct
uniqueFriendsOfFriends: List[String] = List(Nick K, Bill M, Becca G, Kenny D)
```

The `flatten` method is useful in at least two other situations. First, because a `String` is a sequence of `Char`, you can flatten a list of strings into a list of characters:

```
scala> val list = List("Hello", "world")
list: List[java.lang.String] = List(Hello, world)

scala> list.flatten
res0: List[Char] = List(H, e, l, l, o, w, o, r, l, d)
```

Second, because an `Option` can be thought of as a container that holds zero or one elements, `flatten` has a very useful effect on a sequence of `Some` and `None` elements. It pulls the values out of the `Some` elements to create the new list, and drops the `None` elements:

```
scala> val x = Vector(Some(1), None, Some(3), None)
x: Vector[Option[Int]] = Vector(Some(1), None, Some(3), None)

scala> x.flatten
res1: Vector[Int] = Vector(1, 3)
```

10.16. Combining map and flatten with flatMap

Problem

When you first come to Scala, the `flatMap` method can seem very foreign, so you'd like to understand how to use it and see where it can be applied.

Solution

Use `flatMap` in situations where you run `map` followed by `flatten`. The specific situation is this:

- You're using `map` (or a `for/yield` expression) to create a new collection from an existing collection.

- The resulting collection is a list of lists.
- You call `flatten` immediately after `map` (or a `for/yield` expression).

When you're in this situation, you can use `flatMap` instead.

The next example shows how to use `flatMap` with an `Option`. In this example, you're told that you should calculate the sum of the numbers in a list, with one catch: the numbers are all strings, and some of them won't convert properly to integers. Here's the list:

```
val bag = List("1", "2", "three", "4", "one hundred seventy five")
```

To solve the problem, you begin by creating a “string to integer” conversion method that returns either `Some[Int]` or `None`, based on the `String` it's given:

```
def toInt(in: String): Option[Int] = {
  try {
    Some(Integer.parseInt(in.trim))
  } catch {
    case e: Exception => None
  }
}
```

With this method in hand, the resulting solution is surprisingly simple:

```
scala> bag.flatMap(toInt).sum
res0: Int = 7
```

Discussion

To see how this works, break the problem down into smaller steps. First, here's what happens when you use `map` on the initial collection of strings:

```
scala> bag.map(toInt)
res0: List[Option[Int]] = List(Some(1), Some(2), None, Some(4), None)
```

The `map` method applies the `toInt` function to each element in the collection, and returns a list of `Some[Int]` and `None` values. But the `sum` method needs a `List[Int]`; how do you get there from here?

As shown in the previous recipe, `flatten` works very well with a list of `Some` and `None` elements. It extracts the values from the `Some` elements while discarding the `None` elements:

```
scala> bag.map(toInt).flatten
res1: List[Int] = List(1, 2, 4)
```

This makes finding the sum easy:

```
scala> bag.map(toInt).flatten.sum
res2: Int = 7
```

Now, whenever I see `map` followed by `flatten`, I think “flat map,” so I get back to the earlier solution:

```
scala> bag.flatMap(toInt).sum
res3: Int = 7
```

(Actually, I think, “map flat,” but the method is named `flatMap`.)

As you can imagine, once you get the original list down to a `List[Int]`, you can call any of the powerful collections methods to get what you want:

```
scala> bag.flatMap(toInt).filter(_ > 1)
res4: List[Int] = List(2, 4)
```

```
scala> bag.flatMap(toInt).takeWhile(_ < 4)
res5: List[Int] = List(1, 2)
```

```
scala> bag.flatMap(toInt).partition(_ > 3)
res6: (List[Int], List[Int]) = (List(4),List(1, 2))
```

As a second example of using `flatMap`, imagine you have a method that finds all the subwords from a word you give it. Skipping the implementation for a moment, if you call the method with the string then, it should work as follows:

```
scala> subWords("then")
res0: List[String] = List(then, hen, the)
```

(`subWords` should also return the string `he`, but it’s in beta.)

With that method (mostly) working, it can be called on a list of words with `map`:

```
scala> val words = List("band", "start", "then")
words: List[java.lang.String] = List(band, start, then)
```

```
scala> words.map(subWords)
res0: List[List[String]] =
  List(List(band, and, ban), List(start, tart, star), List(then, hen, the))
```

Very cool, you have a list of subwords for all the given words. One problem, though: `map` gave you a list of lists. What to do? Call `flatten`:

```
scala> words.map(subWords).flatten
res1: List[String] = List(band, and, ban, start, tart, star, then, hen, the)
```

Success! You have a list of all the subwords from the original list of words. But notice what you did: You called `map`, then `flatten`. Enter “map flat,” er, `flatMap`:

```
scala> words.flatMap(subWords)
res2: List[String] = List(band, and, ban, start, tart, star, then, hen, the)
```

General rule: Whenever you think `map` followed by `flatten`, use `flatMap`. Eventually your brain will skip over the intermediate steps.

As for the implementation of `subWords` ... well, it’s a work in progress:

```
def subWords(word: String) = List(word, word.tail, word.take(word.length-1))
```

See Also

Recipe 20.6, “Using the Option/Some/None Pattern”, shows another `flatMap` example.

10.17. Using filter to Filter a Collection

Problem

You want to filter the items in a collection to create a new collection that contains only the elements that match your filtering criteria.

Solution

As listed in [Recipe 10.3, “Choosing a Collection Method to Solve a Problem”](#), a variety of methods can be used to filter the elements of an input collection to produce a new output collection. This recipe demonstrates the `filter` method.

To use `filter` on your collection, give it a predicate to filter the collection elements as desired. Your predicate should accept a parameter of the same type that the collection holds, evaluate that element, and return `true` to keep the element in the new collection, or `false` to filter it out. Remember to assign the results of the filtering operation to a new variable.

For instance, the following example shows how to create a list of even numbers from an input list using a modulus algorithm:

```
scala> val x = List.range(1, 10)
x: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9)

// create a list of all the even numbers in the list
scala> val evens = x.filter(_ % 2 == 0)
evens: List[Int] = List(2, 4, 6, 8)
```

As shown, `filter` returns all elements from a sequence that return `true` when your function/predicate is called. There’s also a `filterNot` method that returns all elements from a list for which your function returns `false`.

Discussion

The main methods you can use to filter a collection are listed in [Recipe 10.3](#), and are repeated here for your convenience: `collect`, `diff`, `distinct`, `drop`, `dropWhile`, `filter`, `filterNot`, `find`, `foldLeft`, `foldRight`, `head`, `headOption`, `init`, `intersect`, `last`,

lastOption, reduceLeft, reduceRight, remove, slice, tail, take, takeWhile, and union.

Unique characteristics of filter compared to these other methods include:

- filter walks through all of the elements in the collection; some of the other methods stop before reaching the end of the collection.
- filter lets you supply a predicate (a function that returns true or false) to filter the elements.

How you filter the elements in your collection is entirely up to your algorithm. The following examples show a few ways to filter a list of strings:

```
scala> val fruits = Set("orange", "peach", "apple", "banana")
fruits: scala.collection.immutable.Set[java.lang.String] =
  Set(orange, peach, apple, banana)

scala> val x = fruits.filter(_.startsWith("a"))
x: scala.collection.immutable.Set[String] = Set(apple)

scala> val y = fruits.filter(_.length > 5)
y: scala.collection.immutable.Set[String] = Set(orange, banana)
```

Your filtering function can be as complicated as needed. When your algorithm gets long, you can pass a multiline block of code into filter:

```
scala> val list = "apple" :: "banana" :: 1 :: 2 :: Nil
list: List[Any] = List(apple, banana, 1, 2)

scala> val strings = list.filter {
  |   case s: String => true
  |   case _ => false
  | }
strings: List[Any] = List(apple, banana)
```

You can also put your algorithm in a separate method (or function) and then pass it into filter:

```
def onlyStrings(a: Any) = a match {
  case s: String => true
  case _ => false
}

val strings = list.filter(onlyStrings)
```

The following example demonstrates that you can filter a list as many times as needed:

```
def getFileContentsWithoutBlanksComments(canonicalFilename: String):
List[String] = {
  io.Source.fromFile(canonicalFilename)
    .getLines
    .toList
}
```



```

        .filter(_._trim != "")
        .filter(_._charAt(0) != '#')
    }

```

The two keys to using `filter` are:

- Your algorithm should return `true` for the elements you want to keep and `false` for the other elements
- Remember to assign the results of the `filter` method to a new variable; `filter` doesn't modify the collection it's invoked on

See Also

The `collect` method can also be used as a filtering method. Because it uses partial functions, it's described in detail in [Recipe 9.8, “Creating Partial Functions”](#).

10.18. Extracting a Sequence of Elements from a Collection

Problem

You want to extract a sequence of contiguous elements from a collection, either by specifying a starting position and length, or a function.

Solution

There are quite a few collection methods you can use to extract a contiguous list of elements from a sequence, including `drop`, `dropWhile`, `head`, `headOption`, `init`, `last`, `lastOption`, `slice`, `tail`, `take`, `takeWhile`.

Given the following Array:

```

scala> val x = (1 to 10).toArray
x: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)

```

The `drop` method drops the number of elements you specify from the beginning of the sequence:

```

scala> val y = x.drop(3)
y: Array[Int] = Array(4, 5, 6, 7, 8, 9, 10)

```

The `dropWhile` method drops elements as long as the predicate you supply returns `true`:

```
scala> val y = x.dropWhile(_ < 6)
y: List[Int] = List(6, 7, 8, 9, 10)
```

The `dropRight` method works like `drop`, but starts at the end of the collection and works forward, dropping elements from the end of the sequence:

```
scala> val y = x.dropRight(4)
y: Array[Int] = Array(1, 2, 3, 4, 5, 6)
```

`take` extracts the first N elements from the sequence:

```
scala> val y = x.take(3)
y: Array[Int] = Array(1, 2, 3)
```

`takeWhile` returns elements as long as the predicate you supply returns `true`:

```
scala> val y = x.takeWhile(_ < 5)
y: Array[Int] = Array(1, 2, 3, 4)
```

`takeRight` works the same way `take` works, but starts at the end of the sequence and moves forward, taking the specified number of elements from the end of the sequence:

```
scala> val y = x.takeRight(3)
y: Array[Int] = Array(8, 9, 10)
```

`slice(from, until)` returns a sequence beginning at the index `from` until the index `until`, not including `until`, and assuming a zero-based index:

```
scala> val peeps = List("John", "Mary", "Jane", "Fred")
peeps: List[String] = List(John, Mary, Jane, Fred)

scala> peeps.slice(1,3)
res0: List[String] = List(Mary, Jane)
```

All of these methods provide another way of filtering a collection, with their distinguishing feature being that they return a contiguous sequence of elements.

Even more methods

There are even more methods you can use. Given this list:

```
scala> val nums = (1 to 5).toArray
nums: Array[Int] = Array(1, 2, 3, 4, 5)
```

the comments after the following expressions show the values that are returned by each expression:

```
nums.head           // 1
nums.headOption     // Some(1)
nums.init           // Array(1, 2, 3, 4)
nums.last           // 5
nums.lastOption     // Some(5)
nums.tail           // Array(2, 3, 4, 5)
```

Hopefully the use of most of those methods is obvious. Two that might need a little explanation are `init` and `tail`. The `init` method returns all elements from the sequence

except for the last element. The `tail` method returns all of the elements except the first one.

See the Scaladoc for any sequence (`List`, `Array`, etc.) for more methods.

10.19. Splitting Sequences into Subsets (`groupBy`, `partition`, etc.)

Problem

You want to partition a sequence into two or more different sequences (subsets) based on an algorithm or location you define.

Solution

Use the `groupBy`, `partition`, `span`, or `splitAt` methods to partition a sequence into subsequences. The `sliding` and `unzip` methods can also be used to split sequences into subsequences, though `sliding` can generate many subsequences, and `unzip` primarily works on a sequence of `Tuple2` elements.

The `groupBy`, `partition`, and `span` methods let you split a sequence into subsets according to a function, whereas `splitAt` lets you split a collection into two sequences by providing an index number, as shown in these examples:

```
scala> val x = List(15, 10, 5, 8, 20, 12)
x: List[Int] = List(15, 10, 5, 8, 20, 12)

scala> val y = x.groupBy(_ > 10)
y: Map[Boolean,List[Int]] =
  Map(false -> List(10, 5, 8), true -> List(15, 20, 12))

scala> val y = x.partition(_ > 10)
y: (List[Int], List[Int]) = (List(15, 20, 12), List(10, 5, 8))

scala> val y = x.span(_ < 20)
y: (List[Int], List[Int]) = (List(15, 10, 5, 8), List(20, 12))

scala> val y = x.splitAt(2)
y: (List[Int], List[Int]) = (List(15, 10), List(5, 8, 20, 12))
```

The `groupBy` method partitions the collection into a `Map` of subcollections based on your function. The `true` map contains the elements for which your predicate returned `true`, and the `false` map contains the elements that returned `false`.

The `partition`, `span`, and `splitAt` methods create a `Tuple2` of sequences that are of the same type as the original collection. The `partition` method creates two lists, one containing values for which your predicate returned `true`, and the other containing the

elements that returned false. The `span` method returns a `Tuple2` based on your predicate `p`, consisting of “the longest prefix of this list whose elements all satisfy `p`, and the rest of this list.” The `splitAt` method splits the original list according to the element index value you supplied.

When a `Tuple2` of sequences is returned, its two sequences can be accessed like this:

```
scala> val (a,b) = x.partition(_ > 10)
a: List[Int] = List(15, 20, 12)
b: List[Int] = List(10, 5, 8)
```

The sequences in the `Map` that `groupBy` creates can be accessed like this:

```
scala> val groups = x.groupBy(_ > 10)
groups: scala.collection.immutable.Map[Boolean,List[Int]] =
  Map(false -> List(10, 5, 8), true -> List(15, 20, 12))

scala> val trues = groups(true)
trues: List[Int] = List(15, 20, 12)

scala> val falses = groups(false)
falses: List[Int] = List(10, 5, 8)
```

The `sliding(size, step)` method is an interesting creature that can be used to break a sequence into many groups. It can be called with just a `size`, or both a `size` and `step`:

```
scala> val nums = (1 to 5).toArray
nums: Array[Int] = Array(1, 2, 3, 4, 5)

// size = 2
scala> nums.sliding(2).toList
res0: List[Array[Int]] = List(Array(1, 2), Array(2, 3), Array(3, 4), Array(4, 5))

// size = 2, step = 2
scala> nums.sliding(2,2).toList
res1: List[Array[Int]] = List(Array(1, 2), Array(3, 4), Array(5))

// size = 2, step = 3
scala> nums.sliding(2,3).toList
res2: List[Array[Int]] = List(Array(1, 2), Array(4, 5))
```

As shown, `sliding` works by passing a “sliding window” over the original sequence, returning sequences of a length given by `size`. The `step` parameter lets you skip over elements, as shown in the last two examples. In my experience, the first two examples are the most useful, first with a default step size of 1, and then when step matches size.

The `unzip` method is also interesting. It can be used to take a sequence of `Tuple2` values and create two resulting lists: one that contains the first element of each tuple, and another that contains the second element from each tuple:

```
scala> val listOfTuple2s = List((1,2), ('a', 'b'))
listOfTuple2s: List[(AnyVal, AnyVal)] = List((1,2), (a,b))
```

```
scala> val x = listOfTuple2s.unzip
x: (List[AnyVal], List[AnyVal]) = (List(1, a),List(2, b))
```

For instance, given a list of couples, you can unzip the list to create a list of women and a list of men:

```
scala> val couples = List(("Kim", "Al"), ("Julia", "Terry"))
couples: List[(String, String)] = List((Kim,Al), (Julia,Terry))
```

```
scala> val (women, men) = couples.unzip
women: List[String] = List(Kim, Julia)
men: List[String] = List(Al, Terry)
```

As you might guess from its name, the unzip method is the opposite of zip:

```
scala> val women = List("Kim", "Julia")
women: List[String] = List(Kim, Julia)
```

```
scala> val men = List("Al", "Terry")
men: List[String] = List(Al, Terry)
```

```
scala> val couples = women zip men
couples: List[(String, String)] = List((Kim,Al), (Julia,Terry))
```

See the Scaladoc for any sequence (List, Array, etc.) for more methods.

10.20. Walking Through a Collection with the reduce and fold Methods

Problem

You want to walk through all of the elements in a sequence, comparing two neighboring elements as you walk through the collection.

Solution

Use the `reduceLeft`, `foldLeft`, `reduceRight`, and `foldRight` methods to walk through the elements in a sequence, applying your function to neighboring elements to yield a new result, which is then compared to the next element in the sequence to yield a new result. (Related methods, such as `scanLeft` and `scanRight`, are also shown in the Discussion.)

For example, use `reduceLeft` to walk through a sequence from left to right (from the first element to the last). `reduceLeft` starts by comparing the first two elements in the collection with your algorithm, and returns a result. That result is compared with

the third element, and that comparison yields a new result. That result is compared to the fourth element to yield a new result, and so on.

If you've never used these methods before, you'll see that they give you a surprising amount of power. The best way to show this is with some examples. First, create a sample collection to experiment with:

```
scala> val a = Array(12, 6, 15, 2, 20, 9)
a: Array[Int] = Array(12, 6, 15, 2, 20, 9)
```

Given that sequence, use `reduceLeft` to determine different properties about the collection. The following example shows how to get the sum of all the elements in the sequence:

```
scala> a.reduceLeft(_ + _)
res0: Int = 64
```

Don't let the underscores throw you for a loop; they just stand for the two parameters that are passed into your function. You can write that code like this, if you prefer:

```
a.reduceLeft((x,y) => x + y)
```

The following examples show how to use `reduceLeft` to get the product of all elements in the sequence, the smallest value in the sequence, and the largest value:

```
scala> a.reduceLeft(_ * _)
res1: Int = 388800
```

```
scala> a.reduceLeft(_ min _)
res2: Int = 2
```

```
scala> a.reduceLeft(_ max _)
res3: Int = 20
```

Show each step in the process

You can demonstrate how `reduceLeft` works by creating a larger function. The following function does a “max” comparison like the last example, but has some extra debugging code so you can see how `reduceLeft` works as it marches through the sequence. Here's the function:

```
// returns the max of the two elements
val findMax = (x: Int, y: Int) => {
  val winner = x max y
  println(s"compared $x to $y, $winner was larger")
  winner
}
```

Now call `reduceLeft` again on the array, this time giving it the `findMax` function:

```
scala> a.reduceLeft(findMax)
compared 12 to 6, 12 was larger
compared 12 to 15, 15 was larger
compared 15 to 2, 15 was larger
```

```
compared 15 to 20, 20 was larger
compared 20 to 9, 20 was larger
res0: Int = 20
```

The output shows how `reduceLeft` marches through the elements in the sequence, and how it called the function at each step. Here's how the process works:

- `reduceLeft` starts by calling `findMax` to test the first two elements in the array, 12 and 6. `findMax` returned 12, because 12 is larger than 6.
- `reduceLeft` takes that result (12), and calls `findMax(12, 15)`. 12 is the result of the first comparison, and 15 is the next element in the collection. 15 is larger, so it becomes the new result.
- `reduceLeft` keeps taking the result from the function and comparing it to the next element in the collection, until it marches through all the elements in the collection, ending up with the result, 20.

The code that `reduceLeft` uses under the hood looks like this:

```
// you provide the sequence 'seq' and the function 'f'
var result = seq(0)
for (i <- 1 until seq.length) {
  val next = seq(i)
  result = f(result, next)
}
```

Feeding different algorithms into this loop lets you extract different types of information from your sequence. Wrapping the algorithm in a method also makes for very concise code.

One subtle but important note about `reduceLeft`: the function (or method) you supply must return the same data type that's stored in the collection. This is necessary so `reduceLeft` can compare the result of your function to the next element in the collection.

Working with other sequences and types

As you can imagine, the type contained in the sequence can be anything you need. For instance, determining the longest or shortest string in a sequence of strings is a matter of walking through the elements in the sequence with a function to compare the lengths of two strings:

```
scala> val peeps = Vector("al", "hannah", "emily", "christina", "aleka")
peeps: scala.collection.immutable.Vector[java.lang.String] =
  Vector(al, hannah, emily, christina, aleka)

// longest
scala> peeps.reduceLeft((x,y) => if (x.length > y.length) x else y)
res0: String = christina
```

```
// shortest
scala> peeps.reduceLeft((x,y) => if (x.length < y.length) x else y)
res1: String = al
```

If this had been a collection of `Person` instances, you could run a similar algorithm on each person's name to get the longest and shortest names.

foldLeft, reduceRight, and foldRight

The `foldLeft` method works just like `reduceLeft`, but it lets you set a seed value to be used for the first element. The following examples demonstrate a “sum” algorithm, first with `reduceLeft` and then with `foldLeft`, to demonstrate the difference:

```
scala> val a = Array(1, 2, 3)
a: Array[Int] = Array(1, 2, 3)

scala> a.reduceLeft(_ + _)
res0: Int = 6

scala> a.foldLeft(20)(_ + _)
res1: Int = 26

scala> a.foldLeft(100)(_ + _)
res2: Int = 106
```

In the last two examples, `foldLeft` uses 20 and then 100 for its first element, which affects the resulting sum as shown.

If you haven't seen syntax like that before, `foldLeft` takes two parameter lists. The first parameter list takes one field, the seed value. The second parameter list is the block of code you want to run (your algorithm). [Recipe 3.18, “Creating Your Own Control Structures”](#), demonstrates the use of multiple parameter lists.

The `reduceRight` and `foldRight` methods work the same as `reduceLeft` and `foldLeft`, respectively, but they begin at the end of the collection and work from right to left, i.e., from the end of the collection back to the beginning.

The difference between reduceLeft and reduceRight

In many algorithms, it may not matter if you call `reduceLeft` or `reduceRight`. In that case, you can call `reduce` instead. The `reduce` Scaladoc states, “The order in which operations are performed on elements is unspecified and may be nondeterministic.”

But some algorithms will yield a big difference. For example, given this `divide` function:

```
val divide = (x: Double, y: Double) => {
  val result = x / y
  println(s"divided $x by $y to yield $result")
  result
}
```

and this array:


```
val a = Array(1.0, 2.0, 3.0)
```

`reduceLeft` and `reduceRight` yield a significantly different result:

```
scala> a.reduceLeft(divide)
divided 1.0 by 2.0 to yield 0.5
divided 0.5 by 3.0 to yield 0.16666666666666666
res0: Double = 0.16666666666666666
```

```
scala> a.reduceRight(divide)
divided 2.0 by 3.0 to yield 0.6666666666666666
divided 1.0 by 0.6666666666666666 to yield 1.5
res1: Double = 1.5
```

scanLeft and scanRight

Two methods named `scanLeft` and `scanRight` walk through a sequence in a manner similar to `reduceLeft` and `reduceRight`, but they return a sequence instead of a single value.

For instance, `scanLeft` “Produces a collection containing cumulative results of applying the operator going left to right.” To understand how it works, create another function with a little debug code in it:

```
val product = (x: Int, y: Int) => {
  val result = x * y
  println(s"multiplied $x by $y to yield $result")
  result
}
```

Here’s what `scanLeft` looks like when it’s used with that function and a seed value:

```
scala> val a = Array(1, 2, 3)
a: Array[Int] = Array(1, 2, 3)

scala> a.scanLeft(10)(product)
multiplied 10 by 1 to yield 10
multiplied 10 by 2 to yield 20
multiplied 20 by 3 to yield 60
res0: Array[Int] = Array(10, 10, 20, 60)
```

As you can see, `scanLeft` returns a new sequence, rather than a single value. The `scanRight` method works the same way, but marches through the collection from right to left.

There are a few more related methods, including `reduce` (which was mentioned earlier), `reduceLeftOption`, and `reduceRightOption`.

If you’re curious about the statement in the `reduce` method Scaladoc that, “The order in which operations are performed on elements is unspecified and may be nondeterministic,” run this code in the REPL:

```

val findMax = (x: Int, y: Int) => {
  Thread.sleep(10)
  val winner = x max y
  println(s"compared $x to $y, $winner was larger")
  winner
}

val a = Array.range(0,50)
a.par.reduce(findMax)

```

You'll see that the elements in the sequence are indeed compared in a nondeterministic order.

10.21. Extracting Unique Elements from a Sequence

Problem

You have a collection that contains duplicate elements, and you want to remove the duplicates.

Solution

Call the `distinct` method on the collection:

```

scala> val x = Vector(1, 1, 2, 3, 3, 4)
x: scala.collection.immutable.Vector[Int] = Vector(1, 1, 2, 3, 3, 4)

scala> val y = x.distinct
y: scala.collection.immutable.Vector[Int] = Vector(1, 2, 3, 4)

```

The `distinct` method returns a new collection with the duplicate values removed. Remember to assign the result to a new variable. This is required for both immutable and mutable collections.

If you happen to need a `Set`, converting the collection to a `Set` is another way to remove the duplicate elements:

```

scala> val s = x.toSet
s: scala.collection.immutable.Set[Int] = Set(1, 2, 3, 4)

```

By definition a `Set` can only contain unique elements, so converting an `Array`, `List`, `Vector`, or other sequence to a `Set` removes the duplicates. In fact, this is how `distinct` works. The source code for the `distinct` method in `GenSeqLike` shows that it uses an instance of `mutable.HashSet`.

Using `distinct` with your own classes

To use `distinct` with your own class, you'll need to implement the `equals` and `hashCode` methods. For example, the following class *will* work with `distinct` because it implements those methods:

```

class Person(firstName: String, lastName: String) {

  override def toString = s"$firstName $lastName"

  def canEqual(a: Any) = a.isInstanceOf[Person]

  override def equals(that: Any): Boolean =
    that match {
      case that: Person => that.canEqual(this) && this.hashCode == that.hashCode
      case _ => false
    }

  override def hashCode: Int = {
    val prime = 31
    var result = 1
    result = prime * result + lastName.hashCode;
    result = prime * result + (if (firstName == null) 0 else firstName.hashCode)
    return result
  }
}

object Person {
  def apply(firstName: String, lastName: String) =
    new Person(firstName, lastName)
}

```

You can demonstrate that this class works with `distinct` by placing the following code in the REPL:

```

val dale1 = new Person("Dale", "Cooper")
val dale2 = new Person("Dale", "Cooper")
val ed = new Person("Ed", "Hurley")
val list = List(dale1, dale2, ed)
val uniques = list.distinct

```

The last two lines look like this in the REPL:

```

scala> val list = List(dale1, dale2, ed)
list: List[Person] = List(Dale Cooper, Dale Cooper, Ed Hurley)

scala> val uniquePeople = list.distinct
uniquePeople: List[Person] = List(Dale Cooper, Ed Hurley)

```

If you remove either the `equals` method or `hashCode` method, you'll see that `distinct` won't work as desired.

See Also

You can find the source code for the `SeqLike` trait (and its `distinct` method) by following the *Source* link on its [Scaladoc page](#).

10.22. Merging Sequential Collections

Problem

You want to join two sequences into one sequence, either keeping all of the original elements, finding the elements that are common to both collections, or finding the difference between the two sequences.

Solution

There are a variety of solutions to this problem, depending on your needs:

- Use the `++=` method to merge a sequence into a mutable sequence.
- Use the `++` method to merge two mutable or immutable sequences.
- Use collection methods like `union`, `diff`, and `intersect`.

Use the `++=` method to merge a sequence (any `TraversableOnce`) into a mutable collection like an `ArrayBuffer`:

```
scala> val a = collection.mutable.ArrayBuffer(1,2,3)
a: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3)

scala> a += Seq(4,5,6)
res0: a.type = ArrayBuffer(1, 2, 3, 4, 5, 6)
```

Use the `++` method to merge two mutable or immutable collections while assigning the result to a new variable:

```
scala> val a = Array(1,2,3)
a: Array[Int] = Array(1, 2, 3)

scala> val b = Array(4,5,6)
b: Array[Int] = Array(4, 5, 6)

scala> val c = a ++ b
c: Array[Int] = Array(1, 2, 3, 4, 5, 6)
```

You can also use methods like `union` and `intersect` to combine sequences to create a resulting sequence:

```
scala> val a = Array(1,2,3,4,5)
a: Array[Int] = Array(1, 2, 3, 4, 5)

scala> val b = Array(4,5,6,7,8)
b: Array[Int] = Array(4, 5, 6, 7, 8)

// elements that are in both collections
scala> val c = a.intersect(b)
c: Array[Int] = Array(4, 5)
```

```
// all elements from both collections
scala> val c = a.union(b)
c: Array[Int] = Array(1, 2, 3, 4, 5, 4, 5, 6, 7, 8)

// distinct elements from both collections
scala> val c = a.union(b).distinct
c: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8)
```

The `diff` method results depend on which sequence it's called on:

```
scala> val c = a.diff(b)
c: Array[Int] = Array(1, 2, 3)

scala> val c = b.diff(a)
c: Array[Int] = Array(6, 7, 8)
```

The Scaladoc for the `diff` method states that it returns, “a new list which contains all elements of this list except some of occurrences of elements that also appear in that. If an element value `x` appears `n` times in that, then the first `n` occurrences of `x` will not form part of the result, but any following occurrences will.”

The objects that correspond to most collections also have a `concat` method:

```
scala> Array.concat(a, b)
res0: Array[Int] = Array(1, 2, 3, 4, 4, 5, 6, 7)
```

If you happen to be working with a `List`, the `:::` method prepends the elements of one list to another list:

```
scala> val a = List(1,2,3,4)
a: List[Int] = List(1, 2, 3, 4)

scala> val b = List(4,5,6,7)
b: List[Int] = List(4, 5, 6, 7)

scala> val c = a ::: b
c: List[Int] = List(1, 2, 3, 4, 4, 5, 6, 7)
```

Discussion

You can also use the `diff` method to get the *relative complement* of two sets.



The relative complement of a set `A` with respect to a set `B` is the set of elements in `B` that are not in `A`.

On a recent project, I needed to find the elements in one list that weren't in another list. I did this by converting the lists to sets, and then using the `diff` method to compare the two sets. For instance, given these two arrays:

```
val a = Array(1,2,3,11,4,12,4,5)
val b = Array(6,7,4,5)
```

you can find the relative complement of each array by first converting them to sets, and then comparing them with the `diff` method:

```
// the elements in a that are not in b
scala> val c = a.toSet diff b.toSet
c: scala.collection.immutable.Set[Int] = Set(1, 2, 12, 3, 11)

// the elements in b that are not in a
scala> val d = b.toSet diff a.toSet
d: scala.collection.immutable.Set[Int] = Set(6, 7)
```

If desired, you can then sum those results to get the list of elements that are either in the first set or the second set, but not both sets:

```
scala> val complement = c ++ d
complement: scala.collection.immutable.Set[Int] = Set(1, 6, 2, 12, 7, 3, 11)
```

This works because `diff` returns a set that contains the elements in the current set (`this`) that are not in the other set (`that`).

You can also use the `--` method to get the same result:

```
scala> val c = a.toSet -- b.toSet
c: scala.collection.immutable.Set[Int] = Set(1, 2, 12, 3, 11)

scala> val d = b.toSet -- a.toSet
d: scala.collection.immutable.Set[Int] = Set(6, 7)
```

Subtracting the intersection of the two sets also yields the same result:

```
scala> val i = a.intersect(b)
i: Array[Int] = Array(4, 5)

scala> val c = a.toSet -- i.toSet
c: scala.collection.immutable.Set[Int] = Set(1, 2, 12, 3, 11)

scala> val d = b.toSet -- i.toSet
d: scala.collection.immutable.Set[Int] = Set(6, 7)
```

10.23. Merging Two Sequential Collections into Pairs with `zip`

Problem

You want to merge data from two sequential collections into a collection of key/value pairs.

Solution

Use the `zip` method to join two sequences into one:

```
scala> val women = List("Wilma", "Betty")
women: List[String] = List(Wilma, Betty)

scala> val men = List("Fred", "Barney")
men: List[String] = List(Fred, Barney)

scala> val couples = women.zip(men)
couples: List[(String, String)] = List((Wilma,Fred), (Betty,Barney))
```

This creates an `Array` of `Tuple2` elements, which is a merger of the two original sequences.

This code shows one way to loop over the resulting collection:

```
scala> for ((wife, husband) <- couples) {
  |   println(s"$wife is married to $husband")
  | }
Wilma is married to Fred
Betty is married to Barney
```

Once you have a sequence of tuples like `couples`, you can convert it to a `Map`, which may be more convenient:

```
scala> val couplesMap = couples.toMap
couplesMap: scala.collection.immutable.Map[String,String] =
  Map(Wilma -> Fred, Betty -> Barney)
```

Discussion

If one collection contains more items than the other collection, the items at the end of the longer collection will be dropped. In the previous example, if the `prices` collection contained only one element, the resulting collection will contain only one `Tuple2`:

```
// three elements
scala> val products = Array("breadsticks", "pizza", "soft drink")
products: Array[String] = Array(breadsticks, pizza, soft drink)

// one element
scala> val prices = Array(4)
prices: Array[Int] = Array(4)

// one resulting element
scala> val productsWithPrice = products.zip(prices)
productsWithPrice: Array[(String, Int)] = Array((breadsticks,4))
```

Note that the `unzip` method is the reverse of `zip`:

```
scala> val (a,b) = productsWithPrice.unzip
```

```
a: collection.mutable.IndexedSeq[String] =  
  ArrayBuffer(breadsticks, pizza, soft drink)  
  
b: collection.mutable.IndexedSeq[Double] =  
  ArrayBuffer(4.0, 10.0, 1.5)
```

See Also

Recipes [10.10](#), [10.11](#), and [10.19](#) demonstrate other uses of the `zip` method (and `zipWithIndex`).

10.24. Creating a Lazy View on a Collection

Problem

You're working with a large collection and want to create a “lazy” version of it so it will only compute and return results as they are actually needed.

Solution

Except for the `Stream` class, whenever you create an instance of a Scala collection class, you're creating a *strict* version of the collection. This means that if you create a collection that contains one million elements, memory is allocated for all of those elements immediately. This is the way things normally work in a language like Java.

In Scala you can optionally create a *view* on a collection. A view makes the result non-strict, or *lazy*. This changes the resulting collection, so when it's used with a transformer method, the elements will only be calculated as they are accessed, and not “eagerly,” as they normally would be. (A *transformer method* is a method that transforms an input collection into a new output collection, as described in the Discussion.)

You can see the effect of creating a view on a collection by creating one `Range` without a view, and a second one with a view:

```
scala> 1 to 100  
res0: scala.collection.immutable.Range.Inclusive =  
  Range(1, 2, 3, 4, ... 98, 99, 100)  
  
scala> (1 to 100).view  
res0: java.lang.Object with  
  scala.collection.SeqView[Int,scala.collection.immutable.IndexedSeq[Int]] =  
  SeqView(...)
```

Creating the `Range` without a view shows what you expect, a `Range` with 100 elements. However, the `Range` with the view shows different output in the REPL, showing something called a `SeqView`.

The signature of the `SeqView` shows:

- `Int` is the type of the view's elements.
- The `scala.collection.immutable.IndexedSeq[Int]` portion of the output indicates the type you'll get if you force the collection back to a "normal," strict collection.

You can see this when you force the view back to a normal collection:

```
scala> val view = (1 to 100).view
view: java.lang.Object with
  scala.collection.SeqView[Int,scala.collection.immutable.IndexedSeq[Int]] =
  SeqView(...)

scala> val x = view.force
x: scala.collection.immutable.IndexedSeq[Int] =
  Vector(1, 2, 3, ... 98, 99, 100)
```

There are several ways to see the effect of adding a view to a collection. First, you'll see that using a method like `foreach` doesn't seem to change when using a view:

```
(1 to 100).foreach(println)
(1 to 100).view.foreach(println)
```

Both of those expressions will print 100 elements to the console. Because `foreach` isn't a transformer method, the result is unaffected.

However, calling a `map` method with and without a view has dramatically different results:

```
scala> (1 to 100).map { _ * 2 }
res1: scala.collection.immutable.IndexedSeq[Int] =
  Vector(2, 4, 6, ... 196, 198, 200)

scala> (1 to 100).view.map { _ * 2 }
res0: scala.collection.SeqView[Int,Seq[_]] = SeqViewM(...)
```

These results are different because `map` is a transformer method. A fun way to further demonstrate this difference is with the following code:

```
val x = (1 to 1000).view.map { e =>
  Thread.sleep(10)
  e * 2
}
```

If you run that code as shown, it will return immediately, returning a `SeqView` as before. But if you remove the `view` method call, the code block will take about 10 seconds to run.

Discussion

The Scala documentation states that a view “constructs only a proxy for the result collection, and its elements get constructed only as one demands them ... A view is a special

kind of collection that represents some base collection, but implements all transformers lazily.”

A *transformer* is a method that constructs a new collection from an existing collection. This includes methods like `map`, `filter`, `reverse`, and many more. When you use these methods, you’re transforming the input collection to a new output collection.

This helps to explain why the `foreach` method prints the same result for a strict collection and its view: it’s not a transformer method. But the `map` method, and other transformer methods like `reverse`, treat the view in a lazy manner:

```
scala> l.reverse
res0: List[Int] = List(3, 2, 1)

scala> l.view.reverse
res1: scala.collection.SeqView[Int,List[Int]] = SeqViewR(...)
```

At the end of the Solution you saw this block of code:

```
val x = (1 to 1000).view.map { e =>
  Thread.sleep(10)
  e * 2
}
```

As mentioned, that code returns a `SeqView` immediately. But when you go to print the elements in `x`, like this:

```
x.foreach(print)
```

there will be a 10 ms pause before each element is printed. The elements are being “demanded” in this line of code, so the penalty of the `Thread.sleep` method call is paid as each element is yielded.

Use cases

There are two primary use cases for using a view:

- Performance
- To treat a collection like a database view

Regarding performance, assume that you get into a situation where you may (or may not) have to operate on a collection of a billion elements. You certainly want to avoid running an algorithm on a billion elements if you don’t have to, so using a view makes sense here.

The second use case lets you use a Scala view on a collection just like a database view. The following examples show how a collection view works like a database view:

```
// create a normal array
scala> val arr = (1 to 10).toArray
arr: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

```
// create a view on the array
scala> val view = arr.view.slice(2, 5)
view: scala.collection.mutable.IndexedSeqView[Int,Array[Int]] = SeqViewS(...)

// modify the array
scala> arr(2) = 42

// the view is affected:
scala> view.foreach(println)
42
4
5

// change the elements in the view
scala> view(0) = 10
scala> view(1) = 20
scala> view(2) = 30

// the array is affected:
scala> arr
res0: Array[Int] = Array(1, 2, 10, 20, 30, 6, 7, 8, 9, 10)
```

Changing the elements in the array updates the view, and changing the elements referenced by the view changes the elements in the array. When you need to modify a subset of elements in a collection, creating a view on the original collection and modifying the elements in the view can be a powerful way to achieve this goal.

As a final note, don't confuse using a view with saving memory when creating a collection. Both of the following approaches will generate a "java.lang.OutOfMemoryError: Java heap space" error in the REPL:

```
val a = Array.range(0,123456789)
val a = Array.range(0,123456789).view
```

The benefit of using a view in regards to performance comes with how the view works with transformer methods.

See Also

[An introduction to Scala views](#)

10.25. Populating a Collection with a Range

Problem

You want to populate a List, Array, Vector, or other sequence with a Range.

Solution

Call the `range` method on sequences that support it, or create a `Range` and convert it to the desired sequence.

In the first approach, the `range` method is available on the companion object of supported types like `Array`, `List`, `Vector`, `ArrayBuffer`, and others:

```
scala> Array.range(1, 5)
res0: Array[Int] = Array(1, 2, 3, 4)

scala> List.range(0, 10)
res1: List[Int] = List(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> Vector.range(0, 10, 2)
res2: collection.immutable.Vector[Int] = Vector(0, 2, 4, 6, 8)
```

For some of the collections, such as `List` and `Array`, you can also create a `Range` and convert it to the desired sequence:

```
scala> val a = (0 until 10).toArray
a: Array[Int] = Array(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val list = 1 to 10 by 2 toList
list: List[Int] = List(1, 3, 5, 7, 9)

scala> val list = (1 to 10).by(2).toList
list: List[Int] = List(1, 3, 5, 7, 9)
```

The REPL shows the collections that can be created directly from a `Range`:

```
toArray    toBuffer      toIndexedSeq  toIterable  toIterator
toList     toMap          toSeq         toSet       toStream
toString   toTraversable
```

Using this approach is useful for some collections, like `Set`, which don't offer a `range` method:

```
// intentional error
scala> val set = Set.range(0, 5)
<console>:7: error: value range is not a member of object
scala.collection.immutable.Set
    val set = Set.range(0,5)
                      ^

scala> val set = (0 until 10 by 2).toSet
set: scala.collection.immutable.Set[Int] = Set(0, 6, 2, 8, 4)
```

You can also use a `Range` to create a sequence of characters:

```
scala> val letters = ('a' to 'f').toList
letters: List[Char] = List(a, b, c, d, e, f)
```

```
scala> val letters = ('a' to 'f').by(2).toList
letters: List[Char] = List(a, c, e)
```

As shown in many recipes, ranges are also very useful in for loops:

```
scala> for (i <- 1 until 10 by 2) println(i)
1
3
5
7
9
```

Discussion

By using the map method with a Range, you can create a sequence with elements other than type Int or Char:

```
scala> val map = (1 to 5).map(_ * 2.0)
map: collection.immutable.IndexedSeq[Double] = Vector(2.0, 4.0, 6.0, 8.0, 10.0)
```

Using a similar approach, you can also return a sequence of Tuple2 elements:

```
scala> val map = (1 to 5).map(e => (e,e))
map: scala.collection.immutable.IndexedSeq[(Int, Int)] =
  Vector((1,1), (2,2), (3,3), (4,4), (5,5))
```

That sequence easily converts to a Map:

```
scala> val map = (1 to 5).map(e => (e,e)).toMap
map: scala.collection.immutable.Map[Int,Int] =
  Map(5 -> 5, 1 -> 1, 2 -> 2, 3 -> 3, 4 -> 4)
```

10.26. Creating and Using Enumerations

Problem

You want to use an enumeration (a set of named values that act as constants) in your application.

Solution

Extend the scala Enumeration class to create your enumeration:

```
package com.acme.app {
  object Margin extends Enumeration {
    type Margin = Value
    val TOP, BOTTOM, LEFT, RIGHT = Value
  }
}
```

Then import the enumeration to use it in your application:

```

object Main extends App {

  import com.acme.app.Margin._

  // use an enumeration value in a test
  var currentMargin = TOP

  // later in the code ...
  if (currentMargin == TOP) println("working on Top")

  // print all the enumeration values
  import com.acme.app.Margin
  Margin.values foreach println

}

```

Enumerations are useful tool for creating groups of constants, such as days of the week, weeks of the year, and many other situations where you have a group of related, constant values.

You can also use the following approach, but it generates about four times as much code as an Enumeration, most of which you won't need if your sole purpose is to use it like an enumeration:

```

// a much "heavier" approach
package com.acme.app {
  trait Margin
  case object TOP extends Margin
  case object RIGHT extends Margin
  case object BOTTOM extends Margin
  case object LEFT extends Margin
}

```

See Also

[Scala Enumeration class](#)

10.27. Tuples, for When You Just Need a Bag of Things

Problem

You want to create a small collection of heterogeneous elements.

Solution

A tuple gives you a way to store a group of heterogeneous items in a container, which is useful in many situations.

Create a tuple by enclosing the desired elements between parentheses. This is a two-element tuple:

```
scala> val d = ("Debi", 95)
d: (String, Int) = (Debi,95)
```

Notice that it contains two different types. The following example shows a three-element tuple:

```
scala> case class Person(name: String)
defined class Person

scala> val t = (3, "Three", new Person("Al"))
t: (Int, java.lang.String, Person) = (3,Three,Person(Al))
```

You can access tuple elements using an underscore construct:

```
scala> t._1
res1: Int = 3

scala> t._2
res2: java.lang.String = Three

scala> t._3
res3: Person = Person(Al)
```

I usually prefer to assign them to variables using pattern matching:

```
scala> val(x, y, z) = (3, "Three", new Person("Al"))
x: Int = 3
y: String = Three
z: Person = Person(Al)
```

A nice feature of this approach is that if you don't want all of the elements from the tuple, just use the `_` wildcard character in place of the elements you don't want:

```
scala> val (x, y, _) = t
x: Int = 3
y: java.lang.String = Three

scala> val (x, _, _) = t
x: Int = 3

scala> val (x, _, z) = t
x: Int = 3
z: Person = Person(Al)
```

A two-element tuple is an instance of the `Tuple2` class, and a tuple with three elements is an instance of the `Tuple3` class. (More on this in the Discussion.) As shown earlier, you can create a `Tuple2` like this:

```
scala> val a = ("AL", "Alabama")
a: (java.lang.String, java.lang.String) = (AL,Alabama)
```

You can also create it using these approaches:

```
scala> val b = "AL" -> "Alabama"
b: (java.lang.String, java.lang.String) = (AL,Alabama)

scala> val c = ("AL" -> "Alabama")
c: (java.lang.String, java.lang.String) = (AL,Alabama)
```

When you check the class created by these examples, you'll find they're all of type `Tuple2`:

```
scala> c.getClass
res0: java.lang.Class[_ <: (java.lang.String, java.lang.String)] =
      class scala.Tuple2
```

This syntax is very convenient for other uses, including the creation of maps:

```
val map = Map("AL" -> "Alabama")
```

Discussion

The tuple is an interesting construct. There is no single “`Tuple`” class; instead, the API defines tuple case classes from `Tuple2` through `Tuple22`, meaning that you can have from 2 to 22 elements in a tuple.

A common use case for a tuple is returning multiple items from a method. See [Recipe 5.5, “Defining a Method That Returns Multiple Items \(Tuples\)”](#), for an example of this.

Though a tuple isn't a collection, you can treat a tuple as a collection when needed by creating an iterator:

```
scala> val x = ("AL" -> "Alabama")
x: (java.lang.String, java.lang.String) = (AL,Alabama)

scala> val it = x.productIterator
it: Iterator[Any] = non-empty iterator

scala> for (e <- it) println(e)
AL
Alabama
```

Be aware that like any other iterator, after it's used once, it will be exhausted. Attempting to print the elements a second time yields no output:

```
scala> for (e <- it) println(e)
// no output here
```

Create a new iterator if you need to loop over the elements a second time.

You can also convert a tuple to a collection:

```
scala> val t = ("AL", "Alabama")
t: (String, String) = (AL,Alabama)
```



```
scala> t.productIterator.toArray
res0: Array[Any] = Array(AL, Alabama)
```

See Also

- The [Tuple2 class](#)
- [Recipe 5.5, “Defining a Method That Returns Multiple Items \(Tuples\)”](#)

10.28. Sorting a Collection

Problem

You want to sort a sequential collection. Or, you want to implement the `Ordered` trait in a custom class so you can use the `sorted` method, or operators like `<`, `<=`, `>`, and `>=` to compare instances of your class.

Solution

See [Recipe 11.10, “Sorting Arrays”](#), for information on how to sort an `Array`. Otherwise, use the `sorted` or `sortWith` methods to sort a collection.

The `sorted` method can sort collections with type `Double`, `Float`, `Int`, and any other type that has an implicit `scala.math.Ordering`:

```
scala> val a = List(10, 5, 8, 1, 7).sorted
a: List[Int] = List(1, 5, 7, 8, 10)
```

```
scala> val b = List("banana", "pear", "apple", "orange").sorted
b: List[String] = List(apple, banana, orange, pear)
```

The “rich” versions of the numeric classes (like `RichInt`) and the `StringOps` class all extend the `Ordered` trait, so they can be used with the `sorted` method. (More on the `Ordered` trait in the Discussion.)

The `sortWith` method lets you provide your own sorting function. The following examples demonstrate how to sort a collection of `Int` or `String` in both directions:

```
scala> List(10, 5, 8, 1, 7).sortWith(_ < _)
res1: List[Int] = List(1, 5, 7, 8, 10)
```

```
scala> List(10, 5, 8, 1, 7).sortWith(_ > _)
res2: List[Int] = List(10, 8, 7, 5, 1)
```

```
scala> List("banana", "pear", "apple", "orange").sortWith(_ < _)
res3: List[java.lang.String] = List(apple, banana, orange, pear)
```

```
scala> List("banana", "pear", "apple", "orange").sortWith(_ > _)
res4: List[java.lang.String] = List(pear, orange, banana, apple)
```

Your sorting function can be as complicated as it needs to be. For example, you can access methods on the elements during the sort, such as the following example, which sorts a list of strings by the string length:

```
scala> List("banana", "pear", "apple", "orange").sortWith(_.length < _.length)
res5: List[java.lang.String] = List(pear, apple, banana, orange)

scala> List("banana", "pear", "apple", "orange").sortWith(_.length > _.length)
res6: List[java.lang.String] = List(banana, orange, apple, pear)
```

In the same way the length method is called on a String, you can call a method on any class you want to sort. If your sorting method gets longer, first declare it as a method:

```
def sortByLength(s1: String, s2: String) = {
  println("comparing %s and %s".format(s1, s2))
  s1.length > s2.length
}
```

Then use it by passing it into the sortWith method:

```
scala> List("banana", "pear", "apple").sortWith(sortByLength)
comparing banana and pear
comparing pear and apple
comparing apple and pear
comparing banana and apple
res0: List[String] = List(banana, apple, pear)
```

Discussion

If the type a sequence is holding doesn't have an implicit Ordering, you won't be able to sort it with sorted. For instance, given this basic class:

```
class Person (var name: String) {
  override def toString = name
}
```

create a List[Person]:

```
val ty = new Person("Tyler")
val al = new Person("Al")
val paul = new Person("Paul")
val dudes = List(ty, al, paul)
```

If you try to sort this list in the REPL, you'll see an error stating that the Person class doesn't have an implicit Ordering:

```
scala> dudes.sorted
<console>:13: error: No implicit Ordering defined for Person.
      dudes.sorted
      ^
```

You can't use `sorted` with the `Person` class as it's written, but you can write a simple anonymous function to sort the `Person` elements by the `name` field using `sortWith`:

```
scala> val sortedDudes = dudes.sortWith(_.name < _.name)
sortedDudes: Array[Person] = Array(Al, Paul, Tyler)
```

```
scala> val sortedDudes = dudes.sortWith(_.name > _.name)
sortedDudes: Array[Person] = Array(Tyler, Paul, Al)
```

Mix in the `Ordered` trait

If you'd rather use the `Person` class with the `sorted` method, just mix the `Ordered` trait into the `Person` class, and implement a `compare` method. This technique is shown in the following code:

```
class Person (var name: String) extends Ordered [Person]
{
  override def toString = name

  // return 0 if the same, negative if this < that, positive if this > that
  def compare (that: Person) = {
    if (this.name == that.name)
      0
    else if (this.name > that.name)
      1
    else
      -1
  }
}
```

This new `Person` class can be used with `sorted`.

The `compare` method is what provides the sorting capability. As shown in the comment, `compare` should work like this:

- Return 0 if the two objects are the same (equal, typically using the `equals` method of your class)
- Return a negative value if `this` is less than `that`
- Return a positive value if `this` is greater than `that`

How you determine whether one instance is greater than another instance is entirely up to your `compare` algorithm.

Note that because this `compare` algorithm only compares two `String` values, it could have been written like this:

```
def compare (that: Person) = this.name.compare(that.name)
```

However, I wrote it as shown in the first example to be clear about the approach.

An added benefit of mixing the `Ordered` trait into your class is that it also lets you compare object instances directly in your code:

```
if (al > ty) println("Al") else println("Tyler")
```

This works because the `Ordered` trait implements the `<=`, `<`, `>`, and `>=` methods, and calls your `compare` method to make those comparisons.

See Also

For more information, the `Ordered` and `Ordering` Scaladoc is excellent, with good examples of this approach, and other approaches.

- [The `Ordering` trait](#)
- [The `Ordered` trait](#)

10.29. Converting a Collection to a String with `mkString`

Problem

You want to convert elements of a collection to a `String`, possibly adding a field separator, prefix, and suffix.

Solution

Use the `mkString` method to print a collection as a `String`. Given a simple collection:

```
val a = Array("apple", "banana", "cherry")
```

you can print the collection elements using `mkString`:

```
scala> a.mkString
res1: String = applebananacherry
```

That doesn't look too good, so add a separator:

```
scala> a.mkString(" ")
res2: String = apple banana cherry
```

That's better. Use a comma and a space to create a CSV string:

```
scala> a.mkString(", ")
res3: String = apple, banana, cherry
```

The `mkString` method is overloaded, so you can also add a prefix and suffix:

```
scala> a.mkString("[", ", ", " ]")
res4: String = [apple, banana, cherry]
```

If you happen to have a list of lists that you want to convert to a `String`, such as the following array of arrays, first `flatten` the collection, and then call `mkString`:

```
scala> val a = Array(Array("a", "b"), Array("c", "d"))
a: Array[Array[java.lang.String]] = Array(Array(a, b), Array(c, d))

scala> a.flatten.mkString(", ")
res5: String = a, b, c, d
```

Discussion

You can also use the `toString` method on a collection, but it returns the name of the collection with the elements in the collection listed inside parentheses:

```
scala> val v = Vector("apple", "banana", "cherry")
v: scala.collection.immutable.Vector[String] = Vector(apple, banana, cherry)

scala> v.toString
res0: String = Vector(apple, banana, cherry)
```

List, Array, Map, Set (and More)

Introduction

Whereas [Chapter 10](#) covers collections in general, this chapter provides recipes that are specific to the following collection types:

- `List`
- `Array` (and `ArrayBuffer`)
- `Map`
- `Set`

It also provides a few recipes for special-purpose collections like `Queue`, `Stack`, `Range`, and `Stream`. The following paragraphs provide a brief introduction to the `List`, `Array`, `Map`, and `Set` classes.

List

If you're coming to Scala from Java, you'll quickly see that despite their names, the Scala `List` class is nothing like the Java `List` classes, such as the popular Java `ArrayList`. The Scala `List` class is immutable, so its size as well as the elements it refers to can't change. It's implemented as a linked list, and is generally thought of in terms of its `head`, `tail`, and `isEmpty` methods. Therefore, most operations on a `List` involve recursive algorithms, where the algorithm splits the list into its head and tail components.

Array (and ArrayBuffer)

A Scala `Array` is an interesting collection type. The Scaladoc for the `Array` class states, "Arrays are mutable, indexed collections of values." The class is mutable in that its elements can be changed, but once the size of an `Array` is set, it can never grow or shrink.

Although the `Array` is often demonstrated in Scala examples, and often shows up in the Scala API and third-party APIs, the recommendation with Scala 2.10.x is to use the `Vector` class as your “go to” *immutable*, indexed sequence class, and `ArrayBuffer` as your *mutable*, indexed sequence of choice. In keeping with this suggestion, in my real-world code, I use `Vector` and `ArrayBuffer` for those use cases, and then convert them to an `Array` when needed.

Maps

A Scala `Map` is a collection of key/value pairs, like a Java `Map`, Ruby `Hash`, or Python dictionary. One big difference between a Scala `Map` and the Java `map` classes is that the default `Map` in Scala is immutable, so if you’re not used to working with immutable collections, this can be a big surprise when you attempt to add, delete, or change elements in the map. The techniques of using both immutable and mutable map traits are demonstrated in this chapter.

Sets

A Scala `Set` is also like a Java `Set`. It’s a collection that contains only unique elements, where “uniqueness” is determined by the `==` method of the type the set contains. If you attempt to add duplicate elements to a set, the set silently ignores the request. Scala has both mutable and immutable versions of its base `Set` implementation and offers additional set classes for other needs, such as having a sorted set.

11.1. Different Ways to Create and Populate a List

Problem

You want to create and populate a `List`.

Solution

There are many ways to create and initially populate a `List`:

```
// 1
scala> val list = 1 :: 2 :: 3 :: Nil
list: List[Int] = List(1, 2, 3)

// 2
scala> val list = List(1, 2, 3)
x: List[Int] = List(1, 2, 3)

// 3a
scala> val x = List(1, 2.0, 33D, 4000L)
x: List[Double] = List(1.0, 2.0, 33.0, 4000.0)
```



```
// 3b
scala> val x = List[Number](1, 2.0, 33D, 4000L)
x: List[java.lang.Number] = List(1, 2.0, 33.0, 4000)

// 4
scala> val x = List.range(1, 10)
x: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val x = List.range(0, 10, 2)
x: List[Int] = List(0, 2, 4, 6, 8)

// 5
scala> val x = List.fill(3)("foo")
x: List[String] = List(foo, foo, foo)

// 6
scala> val x = List.tabulate(5)(n => n * n)
x: List[Int] = List(0, 1, 4, 9, 16)

// 7
scala> val x = collection.mutable.ListBuffer(1, 2, 3).toList
x: List[Int] = List(1, 2, 3)

// 8
scala> "foo".toList
res0: List[Char] = List(f, o, o)
```

The first two approaches shown are the most common and straightforward ways to create a `List`. Examples 3a and 3b show how you can manually control the `List` type when your collection has mixed types. When the type isn't manually set in Example 3a, it ends up as a `List[Double]`, and in 3b it's manually set to be a `List[Number]`.

Examples 4 through 6 show different ways to create and populate a `List` with data. Examples 7 and 8 show that many collection types also have a `toList` method that converts their data to a `List`.

Going back to the first example, it shows the `::` method for creating a `List`, which will be new to Java developers. As shown, the `::` method (called *cons*) takes two arguments: a *head* element, which is a single element, and a *tail*, which is another `List`. When a `List` is constructed like this, it must end with a `Nil` element.

It's important to know that the Scala `List` class is not like Java `List` classes, such as the Java `ArrayList`. For example, [Recipe 17.1, “Going to and from Java Collections”](#) shows that a `java.util.List` converts to a Scala `Buffer` or `Seq`, not a Scala `List`.

The following quote from the Scala `List` Scaladoc discusses the important properties of the `List` class:

This class is optimal for last-in-first-out (LIFO), stack-like access patterns. If you need another access pattern, for example, random access or FIFO, consider using a collection more suited to this than `List`. `List` has $O(1)$ prepend and head/tail access. Most other operations are $O(n)$ on the number of elements in the list.

See [Recipe 10.4, “Understanding the Performance of Collections”](#) for more information on the `List` performance characteristics.

See Also

- [The `List` class](#).
- [Recipe 3.15, “Working with a List in a Match Expression”](#), shows how to handle a `List` in a match expression, especially the `Nil` element.
- [Recipe 10.4, “Understanding the Performance of Collections”](#), discusses `List` class performance.
- [Recipe 17.1, “Going to and from Java Collections”](#), demonstrates how to convert back and forth between Scala and Java collections.

11.2. Creating a Mutable List

Problem

You want to use a mutable list (a `LinearSeq`, as opposed to an `IndexedSeq`), but a `List` isn't mutable.

Solution

Use a `ListBuffer`, and convert the `ListBuffer` to a `List` when needed.

The following examples demonstrate how to create a `ListBuffer`, and then add and remove elements, and then convert it to a `List` when finished:

```
import scala.collection.mutable.ListBuffer

var fruits = new ListBuffer[String]()

// add one element at a time to the ListBuffer
fruits += "Apple"
fruits += "Banana"
fruits += "Orange"
```

```

// add multiple elements
fruits += ("Strawberry", "Kiwi", "Pineapple")

// remove one element
fruits -= "Apple"

// remove multiple elements
fruits -= ("Banana", "Orange")

// remove multiple elements specified by another sequence
fruits -= Seq("Kiwi", "Pineapple")

// convert the ListBuffer to a List when you need to
val fruitsList = fruits.toList

```

Discussion

Because a `List` is immutable, if you need to create a list that is constantly changing, the preferred approach is to use a `ListBuffer` while the list is being modified, then convert it to a `List` when a `List` is needed.

The `ListBuffer` Scaladoc states that a `ListBuffer` is “a `Buffer` implementation backed by a list. It provides constant time prepend and append. Most other operations are linear.” So, don’t use `ListBuffer` if you want to access elements arbitrarily, such as accessing items by index (like `list(10000)`); use `ArrayBuffer` instead. See [Recipe 10.4, “Understanding the Performance of Collections”](#) for more information.

Although you can’t modify the elements in a `List`, you can create a new `List` from an existing one, typically prepending items to the original list with the `::` method:

```

scala> val x = List(2)
x: List[Int] = List(2)

scala> val y = 1 :: x
y: List[Int] = List(1, 2)

scala> val z = 0 :: y
z: List[Int] = List(0, 1, 2)

```

This is discussed more in [Recipe 11.3, “Adding Elements to a List”](#).

11.3. Adding Elements to a List

Problem

You want to add elements to a `List` that you’re working with.

Solution

“How do I add elements to a `List`?” is a bit of a trick question, because a `List` is immutable, so you can’t actually add elements to it. If you want a `List` that is constantly changing, use a `ListBuffer` (as described in [Recipe 11.2](#)), and then convert it to a `List` when necessary.

To work with a `List`, the general approach is to prepend items to the list with the `::` method while assigning the results to a new `List`:

```
scala> val x = List(2)
x: List[Int] = List(2)

scala> val y = 1 :: x
y: List[Int] = List(1, 2)

scala> val z = 0 :: y
z: List[Int] = List(0, 1, 2)
```

Rather than continually reassigning the result of this operation to a new variable, you can declare your variable as a `var`, and reassign the result to it:

```
scala> var x = List(2)
x: List[Int] = List(2)

scala> x = 1 :: x
x: List[Int] = List(1, 2)

scala> x = 0 :: x
x: List[Int] = List(0, 1, 2)
```

As these examples illustrate, the `::` method is right-associative; lists are constructed from right to left, which you can see in this example:

```
scala> val list1 = 3 :: Nil
list1: List[Int] = List(3)

scala> val list2 = 2 :: list1
list2: List[Int] = List(2, 3)

scala> val list3 = 1 :: list2
list3: List[Int] = List(1, 2, 3)
```



Any Scala method that ends with a `:` character is evaluated from right to left. This means that the method is invoked on the right operand. You can see how this works by analyzing the following code, where both methods print the number 42:

```
object RightAssociativeExample extends App {  
  val f1 = new Printer  
  f1 >> 42  
  42 >>: f1  
}  
  
class Printer {  
  def >>(i: Int) { println(s"$i") }  
  def >>:(i: Int) { println(s"$i") }  
}
```

The two methods can also be invoked like this:

```
f1.>>(42)  
f1.>>:(42)
```

but by defining the second method to end in a colon, it can be used as a right-associative operator.

Though using `::` is very common, there are additional methods that let you prepend or append single elements to a `List`:

```
scala> val x = List(1)  
x: List[Int] = List(1)  
  
scala> val y = 0 +: x  
y: List[Int] = List(0, 1)  
  
scala> val y = x :+ 2  
y: List[Int] = List(1, 2)
```

You can also merge lists to create a new list. See [Recipe 11.5](#) for examples.

Discussion

If you're not comfortable using a `List`, but want to use a mutable, linear list, see [Recipe 11.2, "Creating a Mutable List"](#) for examples of how to use the `ListBuffer` class. The `ListBuffer` is a mutable, *linear* sequence (as opposed to an *indexed* sequence, like an `Array` or `ArrayBuffer`), and is similar to working with a `StringBuffer` or `StringBuilder` in Java. Just as you'd convert those classes to a `String` when needed, you convert a `ListBuffer` to a `List` when needed. Programmers from other backgrounds may be more comfortable with the `::` approach. A nice benefit of Scala is that it offers both options.

See Also

- [Recipe 11.2, “Creating a Mutable List”](#)
- [Recipe 10.4, “Understanding the Performance of Collections”](#)

11.4. Deleting Elements from a List (or ListBuffer)

Problem

You want to delete elements from a `List` or `ListBuffer`.

Solution

A `List` is immutable, so you can't delete elements from it, but you can filter out the elements you don't want while you assign the result to a new variable:

```
scala> val originalList = List(5, 1, 4, 3, 2)
originalList: List[Int] = List(5, 1, 4, 3, 2)

scala> val newList = originalList.filter(_ > 2)
newList: List[Int] = List(5, 4, 3)
```

Rather than continually assigning the result of operations like this to a new variable, you can declare your variable as a `var` and reassign the result of the operation back to itself:

```
scala> var x = List(5, 1, 4, 3, 2)
x: List[Int] = List(5, 1, 4, 3, 2)

scala> x = x.filter(_ > 2)
x: List[Int] = List(5, 4, 3)
```

See [Chapter 10](#) for other ways to get subsets of a collection using methods like `filter`, `partition`, `splitAt`, `take`, and more.

ListBuffer

If you're going to be modifying a list frequently, it may be better to use a `ListBuffer` instead of a `List`. A `ListBuffer` is mutable, so you can remove items from it using all the methods for mutable sequences shown in [Chapter 10](#). For example, assuming you've created a `ListBuffer` like this:

```
import scala.collection.mutable.ListBuffer
val x = ListBuffer(1, 2, 3, 4, 5, 6, 7, 8, 9)
```

You can delete one element at a time, by value:

```
scala> x -= 5
res0: x.type = ListBuffer(1, 2, 3, 4, 6, 7, 8, 9)
```

You can delete two or more elements at once:

```
scala> x -= (2, 3)
res1: x.type = ListBuffer(1, 4, 6, 7, 8, 9)
```

(That method looks like it takes a tuple, but it's actually defined to take two parameters and a third varargs field.)

You can delete elements by position:

```
scala> x.remove(0)
res2: Int = 1

scala> x
res3: scala.collection.mutable.ListBuffer[Int] = ListBuffer(4, 6, 7, 8, 9)
```

You can use `remove` to delete from a given starting position and provide the number of elements to delete:

```
scala> x.remove(1, 3)

scala> x
res4: scala.collection.mutable.ListBuffer[Int] = ListBuffer(4, 9)
```

You can also use `--=` to delete multiple elements that are specified in another collection:

```
scala> val x = ListBuffer(1, 2, 3, 4, 5, 6, 7, 8, 9)
x: scala.collection.mutable.ListBuffer[Int] = ListBuffer(
  1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> x --= Seq(1,2,3)
res0: x.type = ListBuffer(4, 5, 6, 7, 8, 9)
```

Discussion

When you first start using Scala, the wealth of methods whose names are only symbols (`+=`, `/=`, `::=`, etc.) can seem daunting, but the `-=` and `--=` methods are used consistently across mutable collections, so it quickly becomes second nature to use them.

See Also

- Recipes [10.17](#) through [10.19](#) show many ways to filter collections (filtering is a way of deleting).
- [Recipe 10.3, “Choosing a Collection Method to Solve a Problem”](#).

11.5. Merging (Concatenating) Lists

Problem

You want to merge/concatenate the contents of two lists.

Solution

Merge two lists using the `++`, `concat`, or `:::` methods. Given these two lists:

```
scala> val a = List(1,2,3)
a: List[Int] = List(1, 2, 3)

scala> val b = List(4,5,6)
b: List[Int] = List(4, 5, 6)
```

you can use the `++` method as shown in the following example. It's used consistently across immutable collections, so it's easy to remember:

```
scala> val c = a ++ b
c: List[Int] = List(1, 2, 3, 4, 5, 6)
```

If you work with the `List` class frequently, you may prefer using `:::` as a way to create a new list from two existing lists:

```
scala> val c = a ::: b
c: List[Int] = List(1, 2, 3, 4, 5, 6)
```

The `concat` method on the `List` object also works:

```
scala> val c = List.concat(a, b)
c: List[Int] = List(1, 2, 3, 4, 5, 6)
```

Discussion

Perhaps because I come from a Java background, I don't work with the `List` class too often, so I can't remember some of its custom methods without looking at its Scaladoc. As a result, I prefer the `++` method, because it's consistently used across immutable collections.

However, keep in mind what the `List` class is good at. As its Scaladoc states, "This class is optimal for last-in-first-out (LIFO), stack-like access patterns. If you need another access pattern, for example, random access or FIFO, consider using a collection more suited to this than `List`." See [Recipe 10.4, "Understanding the Performance of Collections"](#) for a discussion of `List` class performance.

See Also

[The `List` class](#)

11.6. Using Stream, a Lazy Version of a List

Problem

You want to use a collection that works like a `List` but invokes its transformer methods (`map`, `filter`, etc.) lazily.

Solution

A `Stream` is like a `List`, except that its elements are computed lazily, in a manner similar to how a *view* creates a lazy version of a collection. Because `Stream` elements are computed lazily, a `Stream` can be long ... infinitely long. Like a view, only the elements that are accessed are computed. Other than this behavior, a `Stream` behaves similar to a `List`.

Just like a `List` can be constructed with `::`, a `Stream` can be constructed with the `#::` method, using `Stream.empty` at the end of the expression instead of `Nil`:

```
scala> val stream = 1 #:: 2 #:: 3 #:: Stream.empty
stream: scala.collection.immutable.Stream[Int] = Stream(1, ?)
```

The REPL output shows that the stream begins with the number 1 but uses a `?` to denote the end of the stream. This is because the end of the stream hasn't been evaluated yet.

For example, given a `Stream`:

```
scala> val stream = (1 to 1000000000).toStream
stream: scala.collection.immutable.Stream[Int] = Stream(1, ?)
```

you can attempt to access the head and tail of the stream. The head is returned immediately:

```
scala> stream.head
res0: Int = 1
```

but the tail isn't evaluated yet:

```
scala> stream.tail
res1: scala.collection.immutable.Stream[Int] = Stream(2, ?)
```

The `?` symbol is the way a lazy collection shows that the end of the collection hasn't been evaluated yet.

As discussed in [Recipe 10.24, "Creating a Lazy View on a Collection"](#), *transformer methods* are computed lazily, so when transformers are called, you see the familiar `?` character that indicates the end of the stream hasn't been evaluated yet:

```
scala> stream.take(3)
res0: scala.collection.immutable.Stream[Int] = Stream(1, ?)

scala> stream.filter(_ < 200)
```

```
res1: scala.collection.immutable.Stream[Int] = Stream(1, ?)

scala> stream.filter(_ > 200)
res2: scala.collection.immutable.Stream[Int] = Stream(201, ?)

scala> stream.map { _ * 2 }
res3: scala.collection.immutable.Stream[Int] = Stream(2, ?)
```

However, be careful with methods that aren't transformers. Calls to the following *strict* methods are evaluated immediately and can easily cause `java.lang.OutOfMemoryError` errors:

```
stream.max
stream.size
stream.sum
```



Transformer methods are collection methods that convert a given input collection to a new output collection, based on an algorithm you provide to transform the data. This includes methods like `map`, `filter`, and `reverse`. When using these methods, you're transforming the input collection to a new output collection. Methods like `max`, `size`, and `sum` don't fit that definition, so they attempt to operate on the `Stream`, and if the `Stream` requires more memory than you can allocate, you'll get the `java.lang.OutOfMemoryError`.

As a point of comparison, if I had attempted to use a `List` in these examples, I would have encountered a `java.lang.OutOfMemoryError` as soon as I attempted to create the `List`:

```
val list = (1 to 100000000).toList
```

Using a `Stream` gives you a chance to specify a huge list, and begin working with its elements:

```
stream(0) // returns 1
stream(1) // returns 2
// ...
stream(10) // returns 11
```

See Also

- [A discussion of Scala's concrete, immutable collections classes, including `Stream`](#)
- [Recipe 10.24, "Creating a Lazy View on a Collection"](#)

11.7. Different Ways to Create and Update an Array

Problem

You want to create and optionally populate an Array.

Solution

There are many different ways to define and populate an Array. You can create an array with initial values, in which case Scala can determine the array type implicitly:

```
scala> val a = Array(1,2,3)
a: Array[Int] = Array(1, 2, 3)

scala> val fruits = Array("Apple", "Banana", "Orange")
fruits: Array[String] = Array(Apple, Banana, Orange)
```

If you don't like the type Scala determines, you can assign it manually:

```
// scala makes this Array[Double]
scala> val x = Array(1, 2.0, 33D, 400L)
x: Array[Double] = Array(1.0, 2.0, 33.0, 400.0)

// manually override the type
scala> val x = Array[Number](1, 2.0, 33D, 400L)
x: Array[java.lang.Number] = Array(1, 2.0, 33.0, 400)
```

You can define an array with an initial size and type, and then populate it later:

```
// create an array with an initial size
val fruits = new Array[String](3)

// somewhere later in the code ...
fruits(0) = "Apple"
fruits(1) = "Banana"
fruits(2) = "Orange"
```

You can create a var reference to an array in a class, and then assign it later:

```
// this uses a null. don't do this in the real world
var fruits: Array[String] = _

// later ...
fruits = Array("apple", "banana")
```

The following examples show a handful of other ways to create and populate an Array:

```
scala> val x = Array.range(1, 10)
x: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val x = Array.range(0, 10, 2)
x: Array[Int] = Array(0, 2, 4, 6, 8)
```

```
scala> val x = Array.fill(3)("foo")
x: Array[String] = Array(foo, foo, foo)

scala> val x = Array.tabulate(5)(n => n * n)
x: Array[Int] = Array(0, 1, 4, 9, 16)

scala> val x = List(1, 2, 3).toArray
x: Array[Int] = Array(1, 2, 3)

scala> "Hello".toArray
res0: Array[Char] = Array(H, e, l, l, o)
```

Discussion

The `Array` is an interesting creature: It's *mutable* in that its elements can be changed, but it's *immutable* in that its size cannot be changed. The first link in the See Also section provides this information about the `Array`:

Scala arrays correspond one-to-one to Java arrays. That is, a Scala array `Array[Int]` is represented as a Java `int[]`, an `Array[Double]` is represented as a Java `double[]` and a `Array[String]` is represented as a Java `String[]`.

The `Array` is an *indexed* sequential collection, so accessing and changing values by their index position is straightforward and fast. Once you've created an `Array`, access its elements by enclosing the desired element number in parentheses:

```
scala> val a = Array(1, 2, 3)
a: Array[Int] = Array(1, 2, 3)

scala> a(0)
res0: Int = 1
```

Just as you access an array element by index, you update elements in a similar way:

```
scala> a(0) = 10

scala> a(1) = 20

scala> a(2) = 30

scala> a
res1: Array[Int] = Array(10, 20, 30)
```

See Also

- [A thorough discussion of `Array`, including background on its implementation.](#)
- [Recipe 10.4, “Understanding the Performance of Collections”](#) discusses `Array` class performance.

11.8. Creating an Array Whose Size Can Change (ArrayBuffer)

Problem

You want to create an array whose size can change, i.e., a completely mutable array.

Solution

An Array is mutable in that its elements can change, but its size can't change. To create a mutable, indexed sequence whose size can change, use the ArrayBuffer class.

To use an ArrayBuffer, import it into scope and then create an instance. You can declare an ArrayBuffer without initial elements, and then add them later:

```
import scala.collection.mutable.ArrayBuffer
var characters = ArrayBuffer[String]()
characters += "Ben"
characters += "Jerry"
characters += "Dale"
```

You can add elements when you create the ArrayBuffer, and continue to add elements later:

```
// initialize with elements
val characters = collection.mutable.ArrayBuffer("Ben", "Jerry")

// add one element
characters += "Dale"

// add two or more elements (method has a varargs parameter)
characters += ("Gordon", "Harry")

// add multiple elements with any TraversableOnce type
characters ++= Seq("Andy", "Big Ed")

// append one or more elements (uses a varargs parameter)
characters.append("Laura", "Lucy")
```

Those are the most common ways to add elements to an ArrayBuffer (and other mutable sequences). The next recipe demonstrates methods to delete ArrayBuffer elements.

11.9. Deleting Array and ArrayBuffer Elements

Problem

You want to delete elements from an Array or ArrayBuffer.

Solution

An `ArrayBuffer` is a mutable sequence, so you can delete elements with the usual `-=`, `remove`, and `clear` methods.

You can remove one or more elements with `-=`:

```
import scala.collection.mutable.ArrayBuffer
val x = ArrayBuffer('a', 'b', 'c', 'd', 'e')

// remove one element
x -= 'a'

// remove multiple elements (methods defines a varargs param)
x -= ('b', 'c')
```

Use `-=` to remove multiple elements that are declared in another collection (any collection that extends `TraversableOnce`):

```
val x = ArrayBuffer('a', 'b', 'c', 'd', 'e')
x -= Seq('a', 'b')
x -= Array('c')
x -= Set('d')
```

Use the `remove` method to delete one element by its position in the `ArrayBuffer`, or a series of elements beginning at a starting position:

```
scala> val x = ArrayBuffer('a', 'b', 'c', 'd', 'e', 'f')
x: scala.collection.mutable.ArrayBuffer[Char] = ArrayBuffer(a, b, c, d, e, f)

scala> x.remove(0)
res0: Char = a

scala> x
res1: scala.collection.mutable.ArrayBuffer[Char] = ArrayBuffer(b, c, d, e, f)

scala> x.remove(1, 3)

scala> x
res2: scala.collection.mutable.ArrayBuffer[Char] = ArrayBuffer(b, f)
```

In these examples, the collection that contains the elements to be removed can be any collection that extends `TraversableOnce`, so `removeThese` can be a `Seq`, `Array`, `Vector`, and many other types that extend `TraversableOnce`.

The `clear` method removes all the elements from an `ArrayBuffer`:

```
scala> var a = ArrayBuffer(1,2,3,4,5)
a: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer(1, 2, 3, 4, 5)

scala> a.clear
```

```
scala> a
res0: scala.collection.mutable.ArrayBuffer[Int] = ArrayBuffer()
```

You can also use the usual Scala filtering methods (drop, filter, take, etc.) to filter elements out of a collection; just remember to assign the result to a new variable.

Array

The size of an Array can't be changed, so you can't directly delete elements. You can reassign the elements in an Array, which has the effect of replacing them:

```
scala> val a = Array("apple", "banana", "cherry")
a: Array[String] = Array(apple, banana, cherry)

scala> a(0) = ""

scala> a(1) = null

scala> a
res0: Array[String] = Array("", null, cherry)
```

You can also filter elements out of one array while you assign the result to a new array:

```
scala> val a = Array("apple", "banana", "cherry")
a: Array[String] = Array(apple, banana, cherry)

scala> val b = a.filter(!_._contains("apple"))
b: Array[String] = Array(banana, cherry)
```

Use other filtering methods (drop, slice, take, etc.) in the same way.

If you define the array variable as a `var`, you can assign the result back to itself, which gives the appearance of deleting elements using filtering:

```
scala> var a = Array("apple", "banana", "cherry")
a: Array[String] = Array(apple, banana, cherry)

scala> a = a.take(2)
a: Array[String] = [Ljava.lang.String;@e41a882

scala> a
res0: Array[String] = Array(apple, banana)
```

11.10. Sorting Arrays

Problem

You want to sort the elements in an Array (or ArrayBuffer).

Solution

If you're working with an `Array` that holds elements that have an `implicit Ordering`, you can sort the `Array` in place using the `scala.util.Sorting.quickSort` method. For example, because the `String` class has an `implicit Ordering`, it can be used with `quickSort`:

```
scala> val fruits = Array("cherry", "apple", "banana")
fruits: Array[String] = Array(cherry, apple, banana)

scala> scala.util.Sorting.quickSort(fruits)

scala> fruits
res0: Array[String] = Array(apple, banana, cherry)
```

Notice that `quickSort` sorts the `Array` in place; there's no need to assign the result to a new variable.

This example works because the `String` class (via `StringOps`) has an `implicit Ordering`. `Sorting.quickSort` can also sort arrays with the base numeric types like `Double`, `Float`, and `Int`, because they also have an `implicit Ordering`.

Other solutions

If the type an `Array` is holding doesn't have an `implicit Ordering`, you can either modify it to mix in the `Ordered` trait (which gives it an `implicit Ordering`), or sort it using the `sorted`, `sortWith`, or `sortBy` methods. These approaches are shown in [Recipe 10.29](#).

Also, there are no unique sorting approaches for an `ArrayBuffer`, so see [Recipe 10.29](#) for an example of how to sort it as well.

See Also

The Scaladoc for the `Ordered` and `Ordering` traits is very good. The header information in both documents shows good examples of the approaches shown in this recipe and [Recipe 10.29](#).

- [The Sorting object](#)
- [The Ordering trait](#)
- [The Ordered trait](#)

11.11. Creating Multidimensional Arrays

Problem

You need to create a multidimensional array, i.e., an array with two or more dimensions.

Solution

There are two main solutions:

- Use `Array.ofDim` to create a multidimensional array. You can use this approach to create arrays of up to five dimensions. With this approach you need to know the number of rows and columns at creation time.
- Create arrays of arrays as needed.

Both approaches are shown in this solution.

Using `Array.ofDim`

Use the `Array.ofDim` method to create the array you need:

```
scala> val rows = 2
rows: Int = 2

scala> val cols = 3
cols: Int = 3

scala> val a = Array.ofDim[String](rows, cols)
a: Array[Array[String]] = Array(Array(null, null, null), Array(null, null, null))
```

After declaring the array, add elements to it:

```
a(0)(0) = "a"
a(0)(1) = "b"
a(0)(2) = "c"
a(1)(0) = "d"
a(1)(1) = "e"
a(1)(2) = "f"
```

Access the elements using parentheses, similar to a one-dimensional array:

```
scala> val x = a(0)(0)
x: String = a
```

Iterate over the array with a for loop:

```
scala> for {
  |   i <- 0 until rows
  |   j <- 0 until cols
  | } println(s"($i)($j) = ${a(i)(j)}")
(0)(0) = a
(0)(1) = b
(0)(2) = c
(1)(0) = d
(1)(1) = e
(1)(2) = f
```

To create an array with more dimensions, just follow that same pattern. Here's the code for a three-dimensional array:

```

val x, y, z = 10
val a = Array.ofDim[Int](x,y,z)
for {
  i <- 0 until x
  j <- 0 until y
  k <- 0 until z
} println(s"($i)($j)($k) = ${a(i)(j)(k)}")

```

Using an array of arrays

Another approach is to create an array whose elements are arrays:

```

scala> val a = Array( Array("a", "b", "c"), Array("d", "e", "f") )
a: Array[Array[String]] = Array(Array(a, b, c), Array(d, e, f))

scala> a(0)
res0: Array[String] = Array(a, b, c)

scala> a(0)(0)
res1: String = a

```

This gives you more control of the process, and lets you create “ragged” arrays (where each contained array may be a different size):

```

scala> val a = Array(Array("a", "b", "c"), Array("d", "e"))
a: Array[Array[String]] = Array(Array(a, b, c), Array(d, e))

```

You can declare your variable as a `var` and create the same array in multiple steps:

```

scala> var arr = Array(Array("a", "b", "c"))
arr: Array[Array[String]] = Array(Array(a, b, c))

scala> arr += Array(Array("d", "e"))

scala> arr
res0: Array[Array[String]] = Array(Array(a, b, c), Array(d, e))

```

Note in this example that the variable `arr` was created as a `var`, which lets you assign the output from the `+=` operator back to it. This gives the illusion that you’ve modified the *contents* of `arr`, but in reality, you’ve modified `arr`’s *reference* so it points at a new collection. (See [Recipe 10.6, “Understanding Mutable Variables with Immutable Collections”](#) for more information.)

Discussion

Decompiling the `Array.ofDim` solution helps to understand how this works behind the scenes. Create the following Scala class in a file named `Test.scala`:

```

class Test {
  val arr = Array.ofDim[String](2, 3)
}

```

If you compile that class with `scalac`, and then decompile it with a tool like JAD, you can see the Java code that's created:

```
private final String arr[][];
```

Similarly, creating a Scala three-dimensional Array like this:

```
val arr = Array.ofDim[String](2, 2, 2)
```

results in a Java array like this:

```
private final String arr[][][];
```

As you might expect, the code generated by using the “array of arrays” approach is more complicated. This is a case where using a decompiler can help you understand how Scala works, i.e., what code it generates for you.

Finally, the `Array.ofDim` approach is unique to the `Array` class; there is no `ofDim` method on a `List`, `Vector`, `ArrayBuffer`, etc. But the “array of arrays” solution is not unique to the `Array` class. You can have a “list of lists,” “vector of vectors,” and so on.

11.12. Creating Maps

Problem

You want to use a mutable or immutable `Map` in a Scala application.

Solution

To use an immutable map, you don't need an import statement, just create a `Map`:

```
scala> val states = Map("AL" -> "Alabama", "AK" -> "Alaska")
states: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska)
```

This expression creates an immutable `Map` with type `[String, String]`. For the first element, the string `AL` is the key, and `Alabama` is the value.

As noted, you don't need an import statement to use a basic, immutable `Map`. The Scala `Predef` object brings the immutable `Map` trait into scope by defining a type alias:

```
type Map[A, +B] = immutable.Map[A, B]
val Map         = immutable.Map
```

To create a *mutable* map, either use an import statement to bring it into scope, or specify the full path to the `scala.collection.mutable.Map` class when you create an instance. You can define a mutable `Map` that has initial elements:

```
scala> var states = collection.mutable.Map("AL" -> "Alabama")
states: scala.collection.mutable.Map[String,String] = Map(AL -> Alabama)
```

You can also create an empty, mutable `Map` initially, and add elements to it later:

```
scala> var states = collection.mutable.Map[String, String]()
states: scala.collection.mutable.Map[String,String] = Map()

scala> states += ("AL" -> "Alabama")
res0: scala.collection.mutable.Map[String,String] = Map(AL -> Alabama)
```

Discussion

Like maps in other programming languages, maps in Scala are a collection of key/value pairs. If you've used maps in Java, dictionaries in Python, or a hash in Ruby, Scala maps are straightforward. You only need to know a couple of new things, including the methods available on map classes, and the specialty maps that can be useful in certain situations, such as having a sorted map.

Note that the syntax that's used inside parentheses in a map creates a `Tuple2`:

```
"AL" -> "Alabama"
```

Because you can also declare a `Tuple2` as `("AL", "Alabama")`, you may also see maps created like this:

```
scala> val states = Map( ("AL", "Alabama"), ("AK", "Alaska") )
states: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska)
```

Use whichever style you prefer.

When I want to be clear that I'm using a mutable map, I normally specify the full path to the mutable `Map` class when I create the instance, as shown in the Solution. Another technique you can use it to give the mutable `Map` an alias when you import it, and then refer to it using that alias, as shown here:

```
import scala.collection.mutable.{Map => MMap}

object Test extends App {

  // MMap is really scala.collection.mutable.Map
  val m = MMap(1 -> 'a')
  for((k,v) <- m) println(s"$k, $v")

}
```

This technique is described more in [Recipe 7.3, “Renaming Members on Import”](#).

See Also

- [The Map trait](#)
- [The Predef object](#)

11.13. Choosing a Map Implementation

Problem

You need to choose a map class for a particular problem.

Solution

Scala has a wealth of map types to choose from, and you can even use Java map classes.

If you're looking for a basic map class, where sorting or insertion order doesn't matter, you can either choose the default, immutable `Map`, or import the mutable `Map`, as shown in the previous recipe.

If you want a map that returns its elements in sorted order by keys, use a `SortedMap`:

```
scala> import scala.collection.SortedMap
import scala.collection.SortedMap

scala> val grades = SortedMap("Kim" -> 90,
  | "Al" -> 85,
  | "Melissa" -> 95,
  | "Emily" -> 91,
  | "Hannah" -> 92
  | )
grades: scala.collection.SortedMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)
```

If you want a map that remembers the insertion order of its elements, use a `LinkedHashMap` or `ListMap`. Scala only has a *mutable* `LinkedHashMap`, and it returns its elements in the order you inserted them:

```
scala> import scala.collection.mutable.LinkedHashMap
import scala.collection.mutable.LinkedHashMap

scala> var states = LinkedHashMap("IL" -> "Illinois")
states: scala.collection.mutable.LinkedHashMap[String,String] =
  Map(IL -> Illinois)

scala> states += ("KY" -> "Kentucky")
res0: scala.collection.mutable.LinkedHashMap[String,String] =
  Map(IL -> Illinois, KY -> Kentucky)

scala> states += ("TX" -> "Texas")
res1: scala.collection.mutable.LinkedHashMap[String,String] =
  Map(IL -> Illinois, KY -> Kentucky, TX -> Texas)
```

Scala has both mutable and immutable `ListMap` classes. They return elements in the opposite order in which you inserted them, as though each insert was at the head of the map (like a `List`):

```
scala> import scala.collection.mutable.ListMap
import scala.collection.mutable.ListMap

scala> var states = ListMap("IL" -> "Illinois")
states: scala.collection.mutable.ListMap[String,String] =
  Map(IL -> Illinois)

scala> states += ("KY" -> "Kentucky")
res0: scala.collection.mutable.ListMap[String,String] =
  Map(KY -> Kentucky, IL -> Illinois)

scala> states += ("TX" -> "Texas")
res1: scala.collection.mutable.ListMap[String,String] =
  Map(TX -> Texas, KY -> Kentucky, IL -> Illinois)
```

The `LinkedHashMap` implements a mutable map using a hashtable, whereas a `ListMap` is backed by a list-based data structure. (Personally, I don't use the `List` class very often, so I prefer the `LinkedHashMap`.)

Discussion

Table 11-1 shows a summary of the basic Scala map classes and traits, and provides a brief description of each.

Table 11-1. Basic map classes and traits

Class or trait	Description
<code>collection.immutable.Map</code>	This is the default, general-purpose immutable map you get if you don't import anything.
<code>collection.mutable.Map</code>	A mutable version of the basic map.
<code>collection.mutable.LinkedHashMap</code>	All methods that traverse the elements will visit the elements in their insertion order.
<code>collection.immutable.ListMap</code> <code>collection.mutable.ListMap</code>	Per the Scaladoc , “implements immutable maps using a list-based data structure.” As shown in the examples, elements that are added are prepended to the head of the list.
<code>collection.SortedMap</code>	Keys of the map are returned in sorted order. Therefore, all traversal methods (such as <code>foreach</code>) return keys in that order.

Although those are the most commonly used maps, Scala offers even more map types. They are summarized in **Table 11-2**.

Table 11-2. More map classes and traits

Class or trait	Description
<code>collection.immutable.HashMap</code>	From the Scaladoc , “implements immutable maps using a hash trie.”
<code>collection.mutable.ObservableMap</code>	From the Scaladoc : “This class is typically used as a mixin. It adds a subscription mechanism to the <code>Map</code> class into which this abstract class is mixed in.”

Class or trait	Description
<code>collection.mutable.MultiMap</code>	From the Scaladoc: “A trait for mutable maps with multiple values assigned to a key.”
<code>collection.mutable.SynchronizedMap</code>	From the Scaladoc: This trait “should be used as a mixin. It synchronizes the map functions of the class into which it is mixed in.”
<code>collection.immutable.TreeMap</code>	From the Scaladoc: “implements immutable maps using a tree.”
<code>collection.mutable.WeakHashMap</code>	A wrapper around <code>java.util.WeakHashMap</code> , “a map entry is removed if the key is no longer strongly referenced.”

But wait, there’s still more. Beyond these types, Scala also offers several more map types that have parallel/concurrent implementations built into them:

- `collection.parallel.immutable.ParHashMap`
- `collection.parallel.mutable.ParHashMap`
- `collection.concurrent.TrieMap`

See Also

- [Map methods](#)
- When map performance is important, see [Recipe 10.4, “Understanding the Performance of Collections”](#)
- [Scala’s parallel collections](#)

11.14. Adding, Updating, and Removing Elements with a Mutable Map

Problem

You want to add, remove, or update elements in a *mutable* map.

Solution

Add elements to a mutable map by simply assigning them, or with the `+=` method. Remove elements with `-=` or `--=`. Update elements by reassigning them.

Given a new, mutable Map:

```
scala> var states = scala.collection.mutable.Map[String, String]()
states: scala.collection.mutable.Map[String,String] = Map()
```

You can add an element to a map by assigning a key to a value:

```
scala> states("AK") = "Alaska"
```

You can also add elements with the += method:

```
scala> states += ("AL" -> "Alabama")
res0: scala.collection.mutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska)
```

Add multiple elements at one time with +=:

```
scala> states += ("AR" -> "Arkansas", "AZ" -> "Arizona")
res1: scala.collection.mutable.Map[String,String] =
  Map(AL -> Alabama, AR -> Arkansas, AK -> Alaska, AZ -> Arizona)
```

Add multiple elements from another collection using +=:

```
scala> states ++= List("CA" -> "California", "CO" -> "Colorado")
res2: scala.collection.mutable.Map[String,String] = Map(CO -> Colorado,
  AZ -> Arizona, AL -> Alabama, CA -> California, AR -> Arkansas,
  AK -> Alaska)
```

Remove a single element from a map by specifying its key with the -= method:

```
scala> states -= "AR"
res3: scala.collection.mutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska, AZ -> Arizona)
```

Remove multiple elements by key with the -= or --= methods:

```
scala> states -= ("AL", "AZ")
res4: scala.collection.mutable.Map[String,String] = Map(AK -> Alaska)

// remove multiple with a List of keys
scala> states --= List("AL", "AZ")
res5: scala.collection.mutable.Map[String,String] = Map(AK -> Alaska)
```

Update elements by reassigning their key to a new value:

```
scala> states("AK") = "Alaska, A Really Big State"

scala> states
res6: scala.collection.mutable.Map[String,String] =
  Map(AK -> Alaska, A Really Big State)
```

There are other ways to add elements to maps, but these examples show the most common uses.

Discussion

The methods shown in the Solution demonstrate the most common approaches. You can also use `put` to add an element (or replace an existing element); `retain` to keep only the elements in the map that match the predicate you supply; `remove` to remove an element by its key value; and `clear` to delete all elements in the map. These methods are shown in the following examples:


```

scala> val states = collection.mutable.Map(
  |   "AK" -> "Alaska",
  |   "IL" -> "Illinois",
  |   "KY" -> "Kentucky"
  | )
states: collection.mutable.Map[String,String] =
  Map(KY -> Kentucky, IL -> Illinois, AK -> Alaska)

scala> states.put("CO", "Colorado")
res0: Option[String] = None

scala> states.retain((k,v) => k == "AK")
res1: states.type = Map(AK -> Alaska)

scala> states.remove("AK")
res2: Option[String] = Some(Alaska)

scala> states
res3: scala.collection.mutable.Map[String,String] = Map()

scala> states.clear

scala> states
res4: scala.collection.mutable.Map[String,String] = Map()

```

As shown, the `remove` method returns an `Option` that contains the value that was removed. It's not shown in the example, but if the element put into the collection by `put` replaced another element, that value would be returned. Because this example didn't replace anything, it returned `None`.

See Also

The Scala mutable `Map` trait

11.15. Adding, Updating, and Removing Elements with Immutable Maps

Problem

You want to add, update, or delete elements when working with an *immutable* map.

Solution

Use the correct operator for each purpose, remembering to assign the results to a new map.

To be clear about the approach, the following examples use an immutable map with a series of `val` variables. First, create an immutable map as a `val`:

```
scala> val a = Map("AL" -> "Alabama")
a: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama)
```

Add one or more elements with the `+` method, assigning the result to a new `Map` variable during the process:

```
// add one element
scala> val b = a + ("AK" -> "Alaska")
b: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska)

// add multiple elements
scala> val c = b + ("AR" -> "Arkansas", "AZ" -> "Arizona")
c: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska, AR -> Arkansas, AZ -> Arizona)
```

To update a key/value pair with an immutable map, reassign the key and value while using the `+` method, and the new values replace the old:

```
scala> val d = c + ("AR" -> "banana")
d: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska, AR -> banana, AZ -> Arizona)
```

To remove one element, use the `-` method:

```
scala> val e = d - "AR"
e: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska, AZ -> Arizona)
```

To remove multiple elements, use the `-` or `--` methods:

```
scala> val f = e - "AZ" - "AL"
f: scala.collection.immutable.Map[String,String] =
  Map(AK -> Alaska)
```

Discussion

You can also declare an immutable map as a `var`. Doing so has a dramatic difference on how you can treat the map:

```
scala> var x = Map("AL" -> "Alabama")
x: scala.collection.mutable.Map[String,String] = Map(AL -> Alabama)

// add one element
scala> x += ("AK" -> "Alaska"); println(x)
Map(AL -> Alabama, AK -> Alaska)

// add multiple elements
scala> x += ("AR" -> "Arkansas", "AZ" -> "Arizona"); println(x)
Map(AZ -> Arizona, AL -> Alabama, AR -> Arkansas, AK -> Alaska)

// add a tuple to a map (replacing the previous "AR" key)
scala> x += ("AR" -> "banana"); println(x)
```

```
Map(AZ -> Arizona, AL -> Alabama, AR -> banana, AK -> Alaska)

// remove an element
scala> x -= "AR"; println(x)
Map(AZ -> Arizona, AL -> Alabama, AK -> Alaska)

// remove multiple elements (uses varargs method)
scala> x -= ("AL", "AZ"); println(x)
Map(AK -> Alaska)

// reassign the map that 'x' points to
scala> x = Map("CO" -> "Colorado")
x: scala.collection.mutable.Map[String,String] = Map(CO -> Colorado)
```

It's important to understand that when you create an immutable map as a var, you still have an immutable map. For instance, you can't reassign an element in the map:

```
scala> x("AL") = "foo"
<console>:9: error: value update is not a member of scala.collection.immutable.Map[String,String]
Map[String,String]
      x("AL") = "foo"
      ^
```

What's really happening in the previous examples is that because x was defined as a var, it's being reassigned during each step in the process. This is a subtle but important distinction to understand. See [Recipe 10.6, “Understanding Mutable Variables with Immutable Collections”](#) for more information.

See Also

[The immutable Map class](#)

11.16. Accessing Map Values

Problem

You want to access individual values stored in a map. You may have tried this and run into an exception when a key didn't exist, and want to see how to avoid that exception.

Solution

Given a sample map:

```
scala> val states = Map("AL" -> "Alabama", "AK" -> "Alaska", "AZ" -> "Arizona")
states: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama, AK -> Alaska, AZ -> Arizona)
```

Access the value associated with a key in the same way you access an element in an array:

```
scala> val az = states("AZ")
az: String = Arizona
```

However, be careful, because if the map doesn't contain the requested key, a `java.util.NoSuchElementException` exception is thrown:

```
scala> val s = states("F00")
java.util.NoSuchElementException: key not found: F00
```

One way to avoid this problem is to create the map with the `withDefaultValue` method. As the name implies, this creates a default value that will be returned by the map whenever a key isn't found:

```
scala> val states = Map("AL" -> "Alabama").withDefaultValue("Not found")
states: scala.collection.immutable.Map[String,String] =
  Map(AL -> Alabama)

scala> states("foo")
res0: String = Not found
```

Another approach is to use the `getOrElse` method when attempting to find a value. It returns the default value you specify if the key isn't found:

```
scala> val s = states.getOrElse("F00", "No such state")
s: String = No such state
```

You can also use the `get` method, which returns an `Option`:

```
scala> val az = states.get("AZ")
az: Option[String] = Some(Arizona)

scala> val az = states.get("F00")
az: Option[String] = None
```

To loop over the values in a map, see the next recipe.

See Also

- [Recipe 11.20, “Testing for the Existence of a Key or Value in a Map”](#).
- [Recipe 20.6, “Using the Option/Some/None Pattern”](#), shows how to work with `Option`, `Some`, and `None` values.

11.17. Traversing a Map

Problem

You want to iterate over the elements in a map.

Solution

There are several different ways to iterate over the elements in a map. Given a sample map:

```
val ratings = Map("Lady in the Water" -> 3.0,
                  "Snakes on a Plane" -> 4.0,
                  "You, Me and Dupree" -> 3.5)
```

my preferred way to loop over all of the map elements is with this for loop syntax:

```
for ((k,v) <- ratings) println(s"key: $k, value: $v")
```

Using a match expression with the `foreach` method is also very readable:

```
ratings.foreach {
  case (movie, rating) => println(s"key: $movie, value: $rating")
}
```

The following approach shows how to use the `Tuple` syntax to access the key and value fields:

```
ratings.foreach(x => println(s"key: ${x._1}, value: ${x._2}"))
```

If you just want to use the keys in the map, the `keys` method returns an `Iterable` you can use:

```
ratings.keys.foreach((movie) => println(movie))
```

For simple examples like this, that expression can be reduced as follows:

```
ratings.keys.foreach(println)
```

In the same way, use the `values` method to iterate over the values in the map:

```
ratings.values.foreach((rating) => println(rating))
```

Note: Those are not my movie ratings. They are taken from the book, *Programming Collective Intelligence* (O'Reilly), by Toby Segaran.

Operating on map values

If you want to traverse the map to perform an operation on its values, the `mapValues` method may be a better solution. It lets you perform a function on each map value, and returns the modified map:

```
scala> var x = collection.mutable.Map(1 -> "a", 2 -> "b")
x: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 1 -> a)

scala> val y = x.mapValues(_.toUpperCase)
y: scala.collection.Map[Int,String] = Map(2 -> B, 1 -> A)
```

The `transform` method gives you another way to create a new map from an existing map. Unlike `mapValues`, it lets you use both the key and value to write a transformation method:

```
scala> val map = Map(1 -> 10, 2 -> 20, 3 -> 30)
map: scala.collection.mutable.Map[Int,Int] = Map(2 -> 20, 1 -> 10, 3 -> 30)

scala> val newMap = map.transform((k,v) => k + v)
newMap: map.type = Map(2 -> 22, 1 -> 11, 3 -> 33)
```

11.18. Getting the Keys or Values from a Map

Problem

You want to get all of the keys or values from a map.

Solution

To get the keys, use `keySet` to get the keys as a `Set`, `keys` to get an `Iterable`, or `keysIterator` to get the keys as an iterator:

```
scala> val states = Map("AK" -> "Alaska", "AL" -> "Alabama", "AR" -> "Arkansas")
states: scala.collection.immutable.Map[String,String] =
  Map(AK -> Alaska, AL -> Alabama, AR -> Arkansas)

scala> states.keySet
res0: scala.collection.immutable.Set[String] = Set(AK, AL, AR)

scala> states.keys
res1: Iterable[String] = Set(AK, AL, AR)

scala> states.keysIterator
res2: Iterator[String] = non-empty iterator
```

To get the values from a map, use the `values` method to get the values as an `Iterable`, or `valuesIterator` to get them as an `Iterator`:

```
scala> states.values
res0: Iterable[String] = MapLike(Alaska, Alabama, Arkansas)

scala> states.valuesIterator
res1: Iterator[String] = non-empty iterator
```

As shown in these examples, `keysIterator` and `valuesIterator` return an iterator from the map data. I tend to prefer these methods because they don't create a new collection; they just provide an iterator to walk over the existing elements.

11.19. Reversing Keys and Values

Problem

You want to reverse the contents of a map, so the values become the keys, and the keys become the values.

Solution

You can reverse the keys and values of a map with a *for comprehension*, being sure to assign the result to a new variable:

```
val reverseMap = for ((k,v) <- map) yield (v, k)
```

But be aware that values don't have to be unique and keys must be, so you might lose some content. As an example of this, reversing the following map—where two values are \$5—results in one of the items being dropped when the keys and values are reversed:

```
scala> val products = Map(
  |   "Breadsticks" -> "$5",
  |   "Pizza" -> "$10",
  |   "Wings" -> "$5"
  | )
products: scala.collection.mutable.Map[String,String] =
  Map(Wings -> $5, Pizza -> $10, Breadsticks -> $5)

scala> val reverseMap = for ((k,v) <- products) yield (v, k)
reverseMap: scala.collection.mutable.Map[String,String] =
  Map($5 -> Breadsticks, $10 -> Pizza)
```

As shown, the \$5 wings were lost when the values became the keys, because both the breadsticks and the wings had the String value \$5.

See Also

- [Recipe 3.4, “Creating a for Comprehension \(for/yield Combination\)”](#)
- [Recipe 10.13, “Transforming One Collection to Another with for/yield”](#)

11.20. Testing for the Existence of a Key or Value in a Map

Problem

You want to test whether a map contains a given key or value.

Solution

To test for the existence of a key in a map, use the `contains` method:

```
scala> val states = Map(
  |   "AK" -> "Alaska",
  |   "IL" -> "Illinois",
  |   "KY" -> "Kentucky"
  | )
states: scala.collection.immutable.Map[String,String] =
  Map(AK -> Alaska, IL -> Illinois, KY -> Kentucky)
```

```
scala> if (states.contains("FOO")) println("Found foo") else println("No foo")
No foo
```

To test whether a value exists in a map, use the `valuesIterator` method to search for the value using `exists` and `contains`:

```
scala> states.valuesIterator.exists(_ contains("ucky"))
res0: Boolean = true
```

```
scala> states.valuesIterator.exists(_ contains("yucky"))
res1: Boolean = false
```

This works because the `valuesIterator` method returns an `Iterator`:

```
scala> states.valuesIterator
res2: Iterator[String] = MapLike(Alaska, Illinois, Kentucky)
```

and `exists` returns `true` if the function you define returns `true` for at least one element in the collection. In the first example, because at least one element in the collection contains the `String` literal `ucky`, the `exists` call returns `true`.

Discussion

When chaining methods like this together, be careful about intermediate results. In this example, I originally used the `values` methods to get the values from the map, but this produces a new collection, whereas the `valuesIterator` method returns a lightweight iterator.

See Also

- [Recipe 11.16, “Accessing Map Values”](#), shows how to avoid an exception while accessing a map key.
- [Recipe 11.18, “Getting the Keys or Values from a Map”](#), demonstrates the `values` and `valuesIterator` methods.

11.21. Filtering a Map

Problem

You want to filter the elements contained in a map, either by directly modifying a mutable map, or by applying a filtering algorithm on an immutable map to create a new map.

Solution

Use the `retain` method to define the elements to retain when using a mutable map, and use `filterKeys` or `filter` to filter the elements in a mutable or immutable map, remembering to assign the result to a new variable.

Mutable maps

You can filter the elements in a *mutable* map using the `retain` method to specify which elements should be retained:

```
scala> var x = collection.mutable.Map(1 -> "a", 2 -> "b", 3 -> "c")
x: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 1 -> a, 3 -> c)

scala> x.retain((k,v) => k > 1)
res0: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 3 -> c)

scala> x
res1: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 3 -> c)
```

As shown, `retain` modifies a mutable map in place. As implied by the anonymous function signature used in that example:

```
(k,v) => ...
```

your algorithm can test both the key and value of each element to decide which elements to retain in the map.

In a related note, the `transform` method doesn't filter a map, but it lets you transform the elements in a mutable map:

```
scala> x.transform((k,v) => v.toUpperCase)
res0: scala.collection.mutable.Map[Int,String] = Map(2 -> B, 3 -> C)

scala> x
res1: scala.collection.mutable.Map[Int,String] = Map(2 -> B, 3 -> C)
```

Depending on your definition of “filter,” you can also remove elements from a map using methods like `remove` and `clear`, which are shown in [Recipe 11.15](#).

Mutable and immutable maps

When working with a mutable or immutable map, you can use a predicate with the `filterKeys` methods to define which map elements to retain. When using this method, remember to assign the filtered result to a new variable:

```
scala> val x = Map(1 -> "a", 2 -> "b", 3 -> "c")
x: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 1 -> a, 3 -> c)

scala> val y = x.filterKeys(_ > 2)
y: scala.collection.Map[Int,String] = Map(3 -> c)
```

The predicate you supply should return `true` for the elements you want to keep in the new collection and `false` for the elements you don't want.

If your algorithm is longer, you can define a function (or method), and then use it in the `filterKeys` call, rather than using an anonymous function. First define your method, such as this method, which returns `true` when the value the method is given is 1:

```
scala> def only1(i: Int) = if (i == 1) true else false
only1: (i: Int)Boolean
```

Then pass the method to the `filterKeys` method:

```
scala> val x = Map(1 -> "a", 2 -> "b", 3 -> "c")
x: scala.collection.mutable.Map[Int,String] = Map(2 -> b, 1 -> a, 3 -> c)

scala> val y = x.filterKeys(only1)
y: scala.collection.Map[Int,String] = Map(1 -> a)
```

In an interesting use, you can also use a `Set` with `filterKeys` to define the elements to retain:

```
scala> var m = Map(1 -> "a", 2 -> "b", 3 -> "c")
m: scala.collection.immutable.Map[Int,String] = Map(1 -> a, 2 -> b, 3 -> c)

scala> val newMap = m.filterKeys(Set(2,3))
newMap: scala.collection.immutable.Map[Int,String] = Map(2 -> b, 3 -> c)
```

You can also use all of the filtering methods that are shown in [Chapter 10](#). For instance, the map version of the `filter` method lets you filter the map elements by either key, value, or both. The `filter` method provides your predicate a `Tuple2`, so you can access the key and value as shown in these examples:

```
scala> var m = Map(1 -> "a", 2 -> "b", 3 -> "c")
m: scala.collection.immutable.Map[Int,String] = Map(1 -> a, 2 -> b, 3 -> c)

// access the key
scala> m.filter((t) => t._1 > 1)
res0: scala.collection.immutable.Map[Int,String] = Map(2 -> b, 3 -> c)

// access the value
scala> m.filter((t) => t._2 == "c")
res1: scala.collection.immutable.Map[Int,String] = Map(3 -> c)
```

The `take` method lets you “take” (keep) the first `N` elements from the map:

```
scala> m.take(2)
res2: scala.collection.immutable.Map[Int,String] = Map(1 -> a, 2 -> b)
```

See the filtering recipes in [Chapter 10](#) for examples of other methods that you can use, including `takeWhile`, `drop`, `slice`, and more.

11.22. Sorting an Existing Map by Key or Value

Problem

You have an unsorted map and want to sort the elements in the map by the key or value.

Solution

Given a basic, immutable Map:

```
scala> val grades = Map("Kim" -> 90,
  |   "Al" -> 85,
  |   "Melissa" -> 95,
  |   "Emily" -> 91,
  |   "Hannah" -> 92
  | )
grades: scala.collection.immutable.Map[String,Int] =
  Map(Hannah -> 92, Melissa -> 95, Kim -> 90, Emily -> 91, Al -> 85)
```

You can sort the map by *key*, from low to high, using `sortBy`:

```
scala> import scala.collection.immutable.ListMap
import scala.collection.immutable.ListMap

scala> ListMap(grades.toSeq.sortBy(_._1):_*)
res0: scala.collection.immutable.ListMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)
```

You can also sort the keys in ascending or descending order using `sortWith`:

```
// low to high
scala> ListMap(grades.toSeq.sortWith(_._1 < _._1):_*)
res0: scala.collection.immutable.ListMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)

// high to low
scala> ListMap(grades.toSeq.sortWith(_._1 > _._1):_*)
res1: scala.collection.immutable.ListMap[String,Int] =
  Map(Melissa -> 95, Kim -> 90, Hannah -> 92, Emily -> 91, Al -> 85)
```

You can sort the map by *value* using `sortBy`:

```
scala> ListMap(grades.toSeq.sortBy(_._2):_*)
res0: scala.collection.immutable.ListMap[String,Int] =
  Map(Al -> 85, Kim -> 90, Emily -> 91, Hannah -> 92, Melissa -> 95)
```

You can also sort by value in ascending or descending order using `sortWith`:

```
// low to high
scala> ListMap(grades.toSeq.sortWith(_._2 < _._2):_*)
res0: scala.collection.immutable.ListMap[String,Int] =
  Map(Al -> 85, Kim -> 90, Emily -> 91, Hannah -> 92, Melissa -> 95)
```

```
// high to low
scala> ListMap(grades.toSeq.sortWith(_._2 > _._2):_*)
res1: scala.collection.immutable.ListMap[String,Int] =
  Map(Melissa -> 95, Hannah -> 92, Emily -> 91, Kim -> 90, Al -> 85)
```

In all of these examples, you're not sorting the existing map; the sort methods result in a new sorted map, so the output of the result needs to be assigned to a new variable. Also, you can use either a `ListMap` or a `LinkedHashMap` in these recipes. This example shows how to use a `LinkedHashMap` and assign the result to a new variable:

```
scala> val x = collection.mutable.LinkedHashMap(grades.toSeq.sortBy(_._1):_*)
x: scala.collection.mutable.LinkedHashMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)

scala> x.foreach(println)
(Al,85)
(Emily,91)
(Hannah,92)
(Kim,90)
(Melissa,95)
```

Discussion

To understand these solutions, it's helpful to break them down into smaller pieces. First, start with the basic immutable `Map`:

```
scala> val grades = Map("Kim" -> 90,
  |   "Al" -> 85,
  |   "Melissa" -> 95,
  |   "Emily" -> 91,
  |   "Hannah" -> 92
  | )
grades: scala.collection.immutable.Map[String,Int] =
  Map(Hannah -> 92, Melissa -> 95, Kim -> 90, Emily -> 91, Al -> 85)
```

Next, this is what `grades.toSeq` looks like:

```
scala> grades.toSeq
res0: Seq[(String, Int)] =
  ArrayBuffer((Hannah,92), (Melissa,95), (Kim,90), (Emily,91), (Al,85))
```

You make the conversion to a `Seq` because it has sorting methods you can use:

```
scala> grades.toSeq.sortBy(_._1)
res0: Seq[(String, Int)] =
  ArrayBuffer((Al,85), (Emily,91), (Hannah,92), (Kim,90), (Melissa,95))

scala> grades.toSeq.sortWith(_._1 < _._1)
res1: Seq[(String, Int)] =
  ArrayBuffer((Al,85), (Emily,91), (Hannah,92), (Kim,90), (Melissa,95))
```

Once you have the map data sorted as desired, store it in a `ListMap` to retain the sort order:

```
scala> ListMap(grades.toSeq.sortBy(_._1):_*)
res0: scala.collection.immutable.ListMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)
```

The `LinkedHashMap` also retains the sort order of its elements, so it can be used in all of the examples as well:

```
scala> import scala.collection.mutable.LinkedHashMap
import scala.collection.mutable.LinkedHashMap

scala> LinkedHashMap(grades.toSeq.sortBy(_._1):_*)
res0: scala.collection.mutable.LinkedHashMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)
```

There are both mutable and immutable versions of a `ListMap`, but `LinkedHashMap` is only available as a mutable class. Use whichever is best for your situation.

About that `_*`

The `_*` portion of the code takes a little getting used to. It's used to convert the data so it will be passed as multiple parameters to the `ListMap` or `LinkedHashMap`. You can see this a little more easily by again breaking down the code into separate lines. The `sortBy` method returns a `Seq[(String, Int)]`, i.e., a sequence of tuples:

```
scala> val x = grades.toSeq.sortBy(_._1)
x: Seq[(String, Int)] =
  ArrayBuffer((Al,85), (Emily,91), (Hannah,92), (Kim,90), (Melissa,95))
```

You can't directly construct a `ListMap` with a sequence of tuples, but because the `apply` method in the `ListMap` companion object accepts a `Tuple2` varargs parameter, you can adapt `x` to work with it, i.e., giving it what it wants:

```
scala> ListMap(x: _*)
res0: scala.collection.immutable.ListMap[String,Int] =
  Map(Al -> 85, Emily -> 91, Hannah -> 92, Kim -> 90, Melissa -> 95)
```

Attempting to create the `ListMap` without using this approach results in an error:

```
scala> ListMap(x)
<console>:16: error: type mismatch;
 found   : Seq[(String, Int)]
 required: (?, ?)
      ListMap(x)
              ^
```

Another way to see how `_*` works is to define your own method that takes a varargs parameter. The following `printAll` method takes one parameter, a varargs field of type `String`:

```
def printAll(strings: String*) {
  strings.foreach(println)
}
```

If you then create a `List` like this:

```
// a sequence of strings
val fruits = List("apple", "banana", "cherry")
```

you won't be able to pass that `List` into `printAll`; it will fail like the previous example:

```
scala> printAll(fruits)
<console>:20: error: type mismatch;
 found   : List[String]
 required: String
    printAll(fruits)
           ^
```

But you can use `_*` to adapt the `List` to work with `printAll`, like this:

```
// this works
printAll(fruits: _*)
```

If you come from a Unix background, it may be helpful to think of `_*` as a “splat” operator. This operator tells the compiler to pass each element of the sequence to `printAll` as a separate argument, instead of passing `fruits` as a single `List` argument.

See Also

- [The immutable `ListMap` class](#)
- [The immutable `ListMap` companion object](#)
- [The mutable `ListMap` class](#)
- [The mutable `LinkedHashMap` class](#)

11.23. Finding the Largest Key or Value in a Map

Problem

You want to find the largest value of a key or value in a map.

Solution

Use the `max` method on the map, or use the map's `keysIterator` or `valuesIterator` with other approaches, depending on your needs.

For example, given this map:

```
val grades = Map("Al" -> 80, "Kim" -> 95, "Teri" -> 85, "Julia" -> 90)
```

the key is type `String`, so which key is “largest” depends on your definition. You can find the “largest” key using the natural `String` sort order by calling the `max` method on the map:

```
scala> grades.max  
res0: (String, Int) = (Teri,85)
```

Because the “T” in “Teri” is farthest down the alphabet in the names, it is returned.

You can also call `keysIterator` to get an iterator over the map keys, and call its `max` method:

```
scala> grades.keysIterator.max  
res1: String = Teri
```

You can find the same maximum by getting the `keysIterator` and using `reduceLeft`:

```
scala> grades.keysIterator.reduceLeft((x,y) => if (x > y) x else y)  
res2: String = Teri
```

This approach is flexible, because if your definition of “largest” is the longest string, you can compare string lengths instead:

```
scala> grades.keysIterator.reduceLeft((x,y) => if (x.length > y.length) x else y)  
res3: String = Julia
```

Because the values in the map are of type `Int` in this example, you can use this simple approach to get the largest value:

```
scala> grades.valuesIterator.max  
res4: Int = 95
```

You can also use the `reduceLeft` approach, if you prefer:

```
scala> grades.valuesIterator.reduceLeft(_ max _)  
res5: Int = 95
```

You can also compare the numbers yourself, which is representative of what you may need to do with more complex types:

```
scala> grades.valuesIterator.reduceLeft((x,y) => if (x > y) x else y)  
res6: Int = 95
```

To find minimum keys and values, just reverse the algorithms in these examples.

See Also

[Recipe 11.18, “Getting the Keys or Values from a Map”](#)

11.24. Adding Elements to a Set

Problem

You want to add elements to a mutable set, or create a new set by adding elements to an immutable set.

Solution

Mutable and immutable sets are handled differently, as demonstrated in the following examples.

Mutable set

Add elements to a *mutable* Set with the `+=`, `++=`, and `add` methods:

```
// use var with mutable
scala> var set = scala.collection.mutable.Set[Int]()
set: scala.collection.mutable.Set[Int] = Set()

// add one element
scala> set += 1
res0: scala.collection.mutable.Set[Int] = Set(1)

// add multiple elements
scala> set += (2, 3)
res1: scala.collection.mutable.Set[Int] = Set(2, 1, 3)

// notice that there is no error when you add a duplicate element
scala> set += 2
res2: scala.collection.mutable.Set[Int] = Set(2, 6, 1, 4, 3, 5)

// add elements from any sequence (any TraversableOnce)
scala> set ++= Vector(4, 5)
res3: scala.collection.mutable.Set[Int] = Set(2, 1, 4, 3, 5)

scala> set.add(6)
res4: Boolean = true

scala> set.add(5)
res5: Boolean = false
```

The last two examples demonstrate a unique characteristic of the `add` method on a set: It returns `true` or `false` depending on whether or not the element was added. The other methods silently fail if you attempt to add an element that's already in the set.

You can test to see whether a set contains an element before adding it:

```
set.contains(5)
```

But as a practical matter, I use `+=` and `++=`, and ignore whether the element was already in the set.

Whereas the first example demonstrated how to create an empty set, you can also add elements to a mutable set when you declare it, just like other collections:

```
scala> var set = scala.collection.mutable.Set(1, 2, 3)
set: scala.collection.mutable.Set[Int] = Set(2, 1, 3)
```


Immutable set

The following examples show how to create a new immutable set by adding elements to an existing immutable set.

First, create an immutable set:

```
scala> val s1 = Set(1, 2)
s1: scala.collection.immutable.Set[Int] = Set(1, 2)
```

Create a new set by adding elements to a previous set with the `+` and `++` methods:

```
// add one element
scala> val s2 = s1 + 3
s2: scala.collection.immutable.Set[Int] = Set(1, 2, 3)

// add multiple elements (+ method has a varargs field)
scala> val s3 = s2 + (4, 5)
s3: scala.collection.immutable.Set[Int] = Set(5, 1, 2, 3, 4)

// add elements from another sequence
scala> val s4 = s3 ++ List(6, 7)
s4: scala.collection.immutable.Set[Int] = Set(5, 1, 6, 2, 7, 3, 4)
```

I showed these examples with immutable variables just to be clear about how the approach works. You can also declare your variable as a `var`, and reassign the resulting set back to the same variable:

```
scala> var set = Set(1, 2, 3)
set: scala.collection.immutable.Set[Int] = Set(1, 2, 3)

scala> set += 4

scala> set
res0: scala.collection.immutable.Set[Int] = Set(1, 2, 3, 4)
```

See [Recipe 10.6, “Understanding Mutable Variables with Immutable Collections”](#) for more information on the difference between mutable/immutable *variables* and mutable/immutable *collections*.

11.25. Deleting Elements from Sets

Problem

You want to remove elements from a mutable or immutable set.

Solution

Mutable and immutable sets are handled differently, as demonstrated in the following examples.

Mutable set

When working with a *mutable* Set, remove elements from the set using the `-=` and `--=` methods, as shown in the following examples:

```
scala> var set = scala.collection.mutable.Set(1, 2, 3, 4, 5)
set: scala.collection.mutable.Set[Int] = Set(2, 1, 4, 3, 5)

// one element
scala> set -= 1
res0: scala.collection.mutable.Set[Int] = Set(2, 4, 3, 5)

// two or more elements (-= has a varargs field)
scala> set -= (2, 3)
res1: scala.collection.mutable.Set[Int] = Set(4, 5)

// multiple elements defined in another sequence
scala> set --= Array(4,5)
res2: scala.collection.mutable.Set[Int] = Set()
```

You can also use other methods like `retain`, `clear`, and `remove`, depending on your needs:

```
// retain
scala> var set = scala.collection.mutable.Set(1, 2, 3, 4, 5)
set: scala.collection.mutable.Set[Int] = Set(2, 1, 4, 3, 5)

scala> set.retain(_ > 2)

scala> set
res0: scala.collection.mutable.Set[Int] = Set(4, 3, 5)

// clear
scala> var set = scala.collection.mutable.Set(1, 2, 3, 4, 5)
set: scala.collection.mutable.Set[Int] = Set(2, 1, 4, 3, 5)

scala> set.clear

scala> set
res1: scala.collection.mutable.Set[Int] = Set()

// remove
scala> var set = scala.collection.mutable.Set(1, 2, 3, 4, 5)
set: scala.collection.mutable.Set[Int] = Set(2, 1, 4, 3, 5)

scala> set.remove(2)
res2: Boolean = true

scala> set
res3: scala.collection.mutable.Set[Int] = Set(1, 4, 3, 5)

scala> set.remove(40)
res4: Boolean = false
```

As shown, the `remove` method provides feedback as to whether or not any elements were removed.

Immutable set

By definition, when using an *immutable* Set you can't remove elements from it, but you can use the `-` and `--` operators to remove elements while assigning the result to a new variable:

```
scala> val s1 = Set(1, 2, 3, 4, 5, 6)
s1: scala.collection.immutable.Set[Int] = Set(5, 1, 6, 2, 3, 4)

// one element
scala> val s2 = s1 - 1
s2: scala.collection.immutable.Set[Int] = Set(5, 6, 2, 3, 4)

// multiple elements
scala> val s3 = s2 - (2, 3)
s3: scala.collection.immutable.Set[Int] = Set(5, 6, 4)

// multiple elements defined in another sequence
scala> val s4 = s3 -- Array(4, 5)
s4: scala.collection.immutable.Set[Int] = Set(6)
```

You can also use all of the filtering methods shown in [Chapter 10](#). For instance, you can use the `filter` or `take` methods:

```
scala> val s1 = Set(1, 2, 3, 4, 5, 6)
s1: scala.collection.immutable.Set[Int] = Set(5, 1, 6, 2, 3, 4)

scala> val s2 = s1.filter(_ > 3)
s2: scala.collection.immutable.Set[Int] = Set(5, 6, 4)

scala> val firstTwo = s1.take(2)
firstTwo: scala.collection.immutable.Set[Int] = Set(5, 1)
```

11.26. Using Sortable Sets

Problem

You want to be able to store and retrieve items from a set in a sorted order.

Solution

To retrieve values from a set in sorted order, use a `SortedSet`. To retrieve elements from a set in the order in which elements were inserted, use a `LinkedHashSet`.

A `SortedSet` returns elements in a sorted order:

```
scala> val s = scala.collection.SortedSet(10, 4, 8, 2)
s: scala.collection.SortedSet[Int] = TreeSet(2, 4, 8, 10)
```

```
scala> val s = scala.collection.SortedSet("cherry", "kiwi", "apple")
s: scala.collection.SortedSet[String] = TreeSet(apple, cherry, kiwi)
```

A `LinkedHashSet` saves elements in the order in which they were inserted:

```
scala> var s = scala.collection.mutable.LinkedHashSet(10, 4, 8, 2)
s: scala.collection.mutable.LinkedHashSet[Int] = Set(10, 4, 8, 2)
```

Discussion

The `SortedSet` is available only in an immutable version. If you need a mutable version, use the `java.util.TreeSet`. The `LinkedHashSet` is available only as a mutable collection.

The examples shown in the Solution work because the types used in the sets have an implicit `Ordering`. Custom types won't work unless you also provide an implicit `Ordering`. For example, the following code won't work because the `Person` class is just a basic class:

```
class Person (var name: String)

import scala.collection.SortedSet
val aleka = new Person("Aleka")
val christina = new Person("Christina")
val molly = new Person("Molly")
val tyler = new Person("Tyler")

// this won't work
val s = SortedSet(molly, tyler, christina, aleka)
```

In the REPL, the last line of code fails with this error:

```
scala> val s = SortedSet(molly, tyler, christina, aleka)
<console>:17: error: No implicit Ordering defined for Person.
    val s = SortedSet(molly, tyler, christina, aleka)
                      ^
```

To solve this problem, modify the `Person` class to extend the `Ordered` trait, and implement a `compare` method:

```
class Person (var name: String) extends Ordered [Person]
{
  override def toString = name

  // return 0 if the same, negative if this < that, positive if this > that
  def compare (that: Person) = {
    if (this.name == that.name)
      0
    else if (this.name > that.name)
      1
    else
      -1
  }
}
```

```
}  
  
}
```

With this new `Person` class definition, sorting works as desired:

```
scala> val s = SortedSet(molly, tyler, christina, aleka)  
s: scala.collection.SortedSet[Person] = TreeSet(Aleka, Christina, Molly, Tyler)
```

For more information about the `Ordered` and `Ordering` traits, see [Recipe 10.28, “Sorting a Collection”](#) and the links in the See Also section.

See Also

- [The `SortedSet` trait](#)
- [The `LinkedHashSet` class](#)
- [The `Ordering` trait](#)
- [The `Ordered` trait](#)

11.27. Using a Queue

Problem

You want to use a queue data structure in a Scala application.

Solution

A queue is a first-in, first-out (FIFO) data structure. Scala offers both an immutable queue and mutable queue. This recipe demonstrates the *mutable* queue.

You can create an empty, mutable queue of any data type:

```
import scala.collection.mutable.Queue  
var ints = Queue[Int]()  
var fruits = Queue[String]()  
var q = Queue[Person]()
```

You can also create a queue with initial elements:

```
scala> val q = Queue(1, 2, 3)  
q: scala.collection.mutable.Queue[Int] = Queue(1, 2, 3)
```

Once you have a mutable queue, add elements to it using `+=`, `++=`, and `enqueue`, as shown in the following examples:

```
scala> import scala.collection.mutable.Queue  
import scala.collection.mutable.Queue
```

```

// create an empty queue
scala> var q = new Queue[String]
q: scala.collection.mutable.Queue[String] = Queue()

// add elements to the queue in the usual ways
scala> var q = new Queue[String]
q: scala.collection.mutable.Queue[String] = Queue()

scala> q += "apple"
res0: scala.collection.mutable.Queue[String] = Queue(apple)

scala> q += ("kiwi", "banana")
res1: scala.collection.mutable.Queue[String] = Queue(apple, kiwi, banana)

scala> q ++= List("cherry", "coconut")
res2: scala.collection.mutable.Queue[String] =
  Queue(apple, kiwi, banana, cherry, coconut)

// can also use enqueue
scala> q.enqueue("pineapple")

scala> q
res3: scala.collection.mutable.Queue[String] =
  Queue(apple, kiwi, banana, cherry, coconut, pineapple)

```

Because a queue is a FIFO, you typically remove elements from the head of the queue, one element at a time, using `dequeue`:

```

// take an element from the head of the queue
scala> val next = q.dequeue
next: String = apple

// 'apple' is removed from the queue
scala> q
res0: scala.collection.mutable.Queue[String] = Queue(kiwi, banana, cherry, ↵
coconut, pineapple)

// take the next element
scala> val next = q.dequeue
next: String = kiwi

// 'kiwi' is removed from the queue
scala> q
res1: scala.collection.mutable.Queue[String] = Queue(banana, cherry, coconut, ↵
pineapple)

```

You can also use the `dequeueFirst` and `dequeueAll` methods to remove elements from the queue by specifying a predicate:

```

scala> q.dequeueFirst(_.startsWith("b"))
res2: Option[String] = Some(banana)

scala> q

```

```
res3: scala.collection.mutable.Queue[String] = Queue(cherry, coconut, pineapple)

scala> q.dequeueAll(_.length > 6)
res4: scala.collection.mutable.Seq[String] = ArrayBuffer(coconut, pineapple)

scala> q
res5: scala.collection.mutable.Queue[String] = Queue(cherry)
```

A Queue is a collection class that extends from Iterable and Traversable, so it has all the usual collection methods, including foreach, map, etc. See the Queue Scaladoc for more information.

See Also

- [The mutable Queue class](#)
- [The immutable Queue class](#)

11.28. Using a Stack

Problem

You want to use a stack data structure in a Scala application.

Solution

A stack is a last-in, first-out (LIFO) data structure. In most programming languages you add elements to a stack using a push method, and take elements off the stack with pop, and Scala is no different.

Scala has both immutable and mutable versions of a stack, as well as an ArrayStack (discussed shortly). The following examples demonstrate how to use the *mutable* Stack class.

Create an empty, mutable stack of any data type:

```
import scala.collection.mutable.Stack
var ints = Stack[Int]()
var fruits = Stack[String]()

case class Person(var name: String)
var people = Stack[Person]()
```

You can also populate a stack with initial elements when you create it:

```
val ints = Stack(1, 2, 3)
```

Once you have a mutable stack, push elements onto the stack with push:

```
// create a stack
scala> var fruits = Stack[String]()
fruits: scala.collection.mutable.Stack[String] = Stack()

// add one element at a time
scala> fruits.push("apple")
res0: scala.collection.mutable.Stack[String] = Stack(apple)

scala> fruits.push("banana")
res1: scala.collection.mutable.Stack[String] = Stack(banana, apple)

// add multiple elements
scala> fruits.push("coconut", "orange", "pineapple")
res2: scala.collection.mutable.Stack[String] =
  Stack(pineapple, orange, coconut, banana, apple)
```

To take elements off the stack, pop them off the top of the stack:

```
scala> val next = fruits.pop
next: String = pineapple

scala> fruits
res3: scala.collection.mutable.Stack[String] =
  Stack(orange, coconut, banana, apple)
```

You can peek at the next element on the stack without removing it, using top:

```
scala> fruits.top
res4: String = orange

// 'orange' is still on the top
scala> fruits
res5: scala.collection.mutable.Stack[String] =
  Stack(orange, coconut, banana, apple)
```

Stack extends from Seq, so you can inspect it with the usual methods:

```
scala> fruits.size
res6: Int = 4

scala> fruits.isEmpty
res7: Boolean = false
```

You can empty a mutable stack with clear:

```
scala> fruits.clear

scala> fruits
res8: scala.collection.mutable.Stack[String] = Stack()
```


Discussion

There's also an `ArrayStack` class, and according to the Scala documentation, "It provides fast indexing and is generally slightly more efficient for most operations than a normal mutable stack."

Although I haven't used an immutable `Stack`, I've seen several people recommend using a `List` instead of an immutable `Stack` for this use case. A `List` has at least one less layer of code, and you can push elements onto the `List` with `::` and access the first element with the `head` method.

See Also

- [The mutable `Stack` class](#)
- [The immutable `Stack` class](#)
- [The `ArrayStack` class](#)

11.29. Using a Range

Problem

You want to use a `Range` in a Scala application.

Solution

Ranges are often used to populate data structures, and to iterate over for loops. Ranges provide a lot of power with just a few methods, as shown in these examples:

```
scala> 1 to 10
res0: scala.collection.immutable.Range.Inclusive =
  Range(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)

scala> 1 until 10
res1: scala.collection.immutable.Range = Range(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> 1 to 10 by 2
res2: scala.collection.immutable.Range = Range(1, 3, 5, 7, 9)

scala> 'a' to 'c'
res3: collection.immutable.NumericRange.Inclusive[Char] = NumericRange(a, b, c)
```

You can use ranges to create and populate sequences:

```
scala> val x = (1 to 10).toList
x: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
```

```
scala> val x = (1 to 10).toArray
x: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)

scala> val x = (1 to 10).toSet
x: scala.collection.immutable.Set[Int] = Set(5, 10, 1, 6, 9, 2, 7, 3, 8, 4)
```

Some sequences have a `range` method in their objects to perform the same function:

```
scala> val x = Array.range(1, 10)
x: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val x = Vector.range(1, 10)
x: collection.immutable.Vector[Int] = Vector(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val x = List.range(1, 10)
x: List[Int] = List(1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> val x = List.range(0, 10, 2)
x: List[Int] = List(0, 2, 4, 6, 8)

scala> val x = collection.mutable.ArrayBuffer.range('a', 'd')
x: scala.collection.mutable.ArrayBuffer[Char] = ArrayBuffer(a, b, c)
```

Ranges are also commonly used in for loops:

```
scala> for (i <- 1 to 3) println(i)
1
2
3
```

Discussion

In addition to the approaches shown, a `Range` can be combined with the `map` method to populate a collection:

```
scala> val x = (1 to 5).map { e => (e + 1.1) * 2 }
x: scala.collection.immutable.IndexedSeq[Double] =
  Vector(4.2, 6.2, 8.2, 10.2, 12.2)
```

While discussing ways to populate collections, the `tabulate` method is another nice approach:

```
scala> val x = List.tabulate(5)(_ + 1)
x: List[Int] = List(1, 2, 3, 4, 5)

scala> val x = List.tabulate(5)(_ + 2)
x: List[Int] = List(2, 3, 4, 5, 6)

scala> val x = Vector.tabulate(5)(_ * 2)
x: scala.collection.immutable.Vector[Int] = Vector(0, 2, 4, 6, 8)
```

See Also

The immutable `Range` class

Files and Processes

12.0. Introduction

When it comes to working with files, the `scala.io.Source` class and its companion object offer some nice simplifications compared to Java. Not only does `Source` make it easy to open and read text files, but it also makes it easy to accomplish other tasks, such as downloading content from URLs, or substituting a `String` for a `File`, which is useful for testing. The `Scala Console` class also simplifies console interaction, letting you print to the console (command line) and read from it very easily. In other cases, such as when reading a YAML file or working with directories, you simply fall back to use existing Java libraries.

Scala also makes it *much* easier to execute system commands. When it comes to interacting with system processes, the Scala API designers created a clean and familiar API to let you run external commands. This is useful for applications, and it's terrific for scripts.

12.1. How to Open and Read a Text File

Problem

You want to open a plain-text file in Scala and process the lines in that file.

Solution

There are two primary ways to open and read a text file:

- Use a concise, one-line syntax. This has the side effect of leaving the file open, but can be useful in short-lived programs, like shell scripts.
- Use a slightly longer approach that properly closes the file.

This solution shows both approaches.

Using the concise syntax

In Scala shell scripts, where the JVM is started and stopped in a relatively short period of time, it may not matter that the file is closed, so you can use the Scala `scala.io.Source.fromFile` method as shown in the following examples.

To handle each line in the file as it's read, use this approach:

```
import scala.io.Source

val filename = "fileopen.scala"
for (line <- Source.fromFile(filename).getLines) {
  println(line)
}
```

As a variation of this, use the following approach to get all of the lines from the file as a List or Array:

```
val lines = Source.fromFile("/Users/Al/.bash_profile").getLines.toList
val lines = Source.fromFile("/Users/Al/.bash_profile").getLines.toArray
```

The `fromFile` method returns a `BufferedSource`, and its `getLines` method treats “any of `\r\n`, `\r`, or `\n` as a line separator (longest match),” so each element in the sequence is a line from the file.

Use this approach to get all of the lines from the file as one String:

```
val fileContents = Source.fromFile(filename).getLines.mkString
```

This approach has the side effect of leaving the file open as long as the JVM is running, but for short-lived shell scripts, this shouldn't be an issue; the file is closed when the JVM shuts down.

Properly closing the file

To properly close the file, get a reference to the `BufferedSource` when opening the file, and manually close it when you're finished with the file:

```
val bufferedSource = Source.fromFile("example.txt")
for (line <- bufferedSource.getLines) {
  println(line.toUpperCase)
}
bufferedSource.close
```

For automated methods of closing the file, see the “Loan Pattern” examples in the Discussion.

Discussion

The `getLines` method of the `Source` class returns a `scala.collection.Iterator`. The iterator returns each line without any newline characters. An iterator has many methods for working with a collection, and for the purposes of working with a file, it works well with the `for` loop, as shown.

Leaving files open

As mentioned, the first solution leaves the file open as long as the JVM is running:

```
// leaves the file open
for (line <- io.Source.fromFile("/etc/passwd").getLines) {
  println(line)
}

// also leaves the file open
val contents = io.Source.fromFile("/etc/passwd").mkString
```

On Unix systems, you can show whether a file is left open by executing one of these `fromFile` statements in the REPL with a real file (like `/etc/passwd`), and then running an `lsof` (“list open files”) command like this at the Unix command line:

```
$ sudo lsof -u Al | grep '/etc/passwd'
```

That command lists all the open files for the user named `Al`, and then searches the output for the `/etc/passwd` file. If this filename is in the output, it means that it’s open. On my Mac OS X system I see a line of output like this when the file is left open:

```
java 17148 Al 40r REG 14,2 1475 174214161 /etc/passwd
```

When I shut down the REPL—thereby stopping the JVM process—the file no longer appears in the `lsof` output. So while this approach has this flaw, it can be used in short-lived JVM processes, such as a shell script. (You can demonstrate the same result using a Scala shell script. Just add a `Thread.sleep` call after the `for` loop so you can keep the script running long enough to check the `lsof` command.)

Automatically closing the resource

When working with files and other resources that need to be properly closed, it’s best to use the **Loan Pattern**. According to this website, the pattern “ensures that a resource is deterministically disposed of once it goes out of scope.”

In Scala, this can be ensured with a `try/finally` clause, which the Loan Pattern website shows like this:

```
def using[A](r : Resource)(f : Resource => A) : A =
  try {
    f(r)
  } finally {
    r.dispose()
  }
```

One way to implement the Loan Pattern when working with files is to use Joshua Suereth's [ARM library](#). To demonstrate this library, create an SBT project, and then add the following line to its *build.sbt* file to pull in the required dependencies:

```
libraryDependencies += "com.jsuereth" %% "scala-arm" % "1.3"
```

Next, create a file named *TestARM.scala* in the root directory of your SBT project with these contents:

```
import resource._

object TestARM extends App {

  for (source <- managed(scala.io.Source.fromFile("example.txt"))) {
    for (line <- source.getLines) {
      println(line)
    }
  }

}
```

This code prints all of the lines from the file named *example.txt*. The `managed` method from the ARM library makes sure that the resource is closed automatically when the resource goes out of scope. The ARM website shows several other ways the library can be used.

A second way to demonstrate the Loan Pattern is with the `using` method described on the Loan Pattern website. The best implementation I've seen of a `using` method is in the book *Beginning Scala* (Apress), by David Pollak. The following code is a slight modification of his code:

```
object Control {

  def using[A <: { def close(): Unit }, B](resource: A)(f: A => B): B =
    try {
      f(resource)
    } finally {
      resource.close()
    }

}
```

This `using` method takes two parameters:

- An object that has a `close()` method
- A block of code to be executed, which transforms the input type `A` to the output type `B`

The body of this `using` method does exactly what's shown on the Loan Pattern web page, wrapping the block of code it's given in a `try/finally` block.

The following code demonstrates how to use this method when reading from a file:

```
import Control._

object TestUsing extends App {

  using(io.Source.fromFile("example.txt")) { source => {
    for (line <- source.getLines) {
      println(line)
    }
  }
}
```

Both the ARM library and the using method end up with the same result, implementing the Loan Pattern to make sure your resource is closed automatically.

Handling exceptions

You can generate exceptions any time you try to open a file, and if you want to handle your exceptions, use Scala's try/catch syntax:

```
import scala.io.Source
import java.io.{FileNotFoundException, IOException}

val filename = "no-such-file.scala"
try {
  for (line <- Source.fromFile(filename).getLines) {
    println(line)
  }
} catch {
  case e: FileNotFoundException => println("Couldn't find that file.")
  case e: IOException => println("Got an IOException!")
}
```

The following code demonstrates how the fromFile method can be used with using to create a method that returns the entire contents of a file as a List[String], wrapped in an Option:

```
import Control._

def readTextFile(filename: String): Option[List[String]] = {
  try {
    val lines = using(io.Source.fromFile(filename)) { source =>
      (for (line <- source.getLines) yield line).toList
    }
    Some(lines)
  } catch {
    case e: Exception => None
  }
}
```

This method returns a `Some(List[String])` on success, and `None` if something goes wrong, such as a `FileNotFoundException`. It can be used in the following ways:

```
val filename = "/etc/passwd"

println("--- FOREACH ---")
val result = readTextFile(filename)
result foreach { strings =>
  strings.foreach(println)
}

println("\n--- MATCH ---")
readTextFile(filename) match {
  case Some(lines) => lines.foreach(println)
  case None => println("couldn't read file")
}
```

If the process of opening and reading a file fails, you may prefer to return a `Try` or an empty `List[String]`. See Recipes 20.5 and 20.6 for examples of those approaches.

Multiple fromFile methods

In Scala 2.10, there are eight variations of the `fromFile` method that let you specify a character encoding, buffer size, codec, and URI. For instance, you can specify an expected character encoding for a file like this:

```
// specify the encoding
Source.fromFile("example.txt", "UTF-8")
```

See the Scaladoc for the `scala.io.Source` object (not the `Source` class, which is an abstract class) for more information.



Because Scala works so well with Java, you can use the Java `FileReader` and `BufferedReader` classes, as well as other Java libraries, like the Apache Commons `FileUtils` library.

See Also

- [The Source object.](#)
- [The Loan Pattern.](#)
- [Joshua Suereth's ARM library.](#)
- David Pollak's book, *Beginning Scala*.

- [A detailed discussion of David Pollak's using method.](#)
- The Apache Commons FileUtils project has [many methods](#) for reading and writing files that can be used with Scala.

12.2. Writing Text Files

Problem

You want to write plain text to a file, such as a simple configuration file, text data file, or other plain-text document.

Solution

Scala doesn't offer any special file writing capability, so fall back and use the Java `PrintWriter` or `FileWriter` approaches:

```
// PrintWriter
import java.io._
val pw = new PrintWriter(new File("hello.txt" ))
pw.write("Hello, world")
pw.close

// FileWriter
val file = new File(canonicalFilename)
val bw = new BufferedWriter(new FileWriter(file))
bw.write(text)
bw.close()
```

Discussion

Although I normally use a `FileWriter` to write plain text to a file, a good post at coderanch.com describes some of the differences between `PrintWriter` and `FileWriter`. For instance, while both classes extend from `Writer`, and both can be used for writing plain text to files, `FileWriter` throws `IOExceptions`, whereas `PrintWriter` does not throw exceptions, and instead sets Boolean flags that can be checked. There are a few other differences between the classes; check their Javadoc for more information.

See Also

- [My Java file utilities](#) and my [Scala file utilities](#)
- [The Java `FileWriter` class](#)

- The [Java PrintWriter class](#)
- The [coderanch.com](#) [PrintWriter versus FileWriter page](#)

12.3. Reading and Writing Binary Files

Problem

You want to read data from a binary file or write data to a binary file.

Solution

Scala doesn't offer any special conveniences for reading or writing binary files, so use the Java `FileInputStream` and `FileOutputStream` classes.

To demonstrate this, the following code is a close Scala translation of the `CopyBytes` class on the [Oracle Byte Streams tutorial](#):

```
import java.io._

object CopyBytes extends App {

  var in = None: Option[FileInputStream]
  var out = None: Option[FileOutputStream]

  try {
    in = Some(new FileInputStream("/tmp/Test.class"))
    out = Some(new FileOutputStream("/tmp/Test.class.copy"))
    var c = 0
    while ({c = in.get.read; c != -1}) {
      out.get.write(c)
    }
  } catch {
    case e: IOException => e.printStackTrace
  } finally {
    println("entered finally ...")
    if (in.isDefined) in.get.close
    if (out.isDefined) out.get.close
  }
}
```

In this code, `in` and `out` are populated in the `try` clause. It's safe to call `in.get` and `out.get` in the `while` loop, because if an exception had occurred, flow control would have switched to the `catch` clause, and then the `finally` clause before leaving the method.

Normally I tell people that I think the `get` and `isDefined` methods on `Option` would be deprecated, but this is one of the few times where I think their use is acceptable and they lead to more readable code.

Another difference between this code and Oracle's example is the `while` loop, which is slightly different in Scala. This change is required because a Java statement like `c = in.read` has a type of `Unit` in Scala, and will therefore never be equal to `-1` (or any other value). There are several other ways to work around this difference, but this example shows a fairly direct translation.

See Also

- [The Oracle Byte Streams tutorial](#)
- The [Apache Commons FileUtils project](#) has many methods for reading and writing files that can be used with Scala

12.4. How to Process Every Character in a Text File

Problem

You want to open a text file and process every character in the file.

Solution

If performance isn't a concern, write your code in a straightforward, obvious way:

```
val source = io.Source.fromFile("/Users/Al/.bash_profile")
for (char <- source) {
  println(char.toUpperCase)
}
source.close
```

However, be aware that this code may be slow on large files. For instance, the following method that counts the number of lines in a file takes 100 seconds to run on an Apache access logfile that is ten million lines long:

```
// run time: took 100 secs
def countLines1(source: io.Source): Long = {
  val NEWLINE = 10
  var newlineCount = 0L
  for {
    char <- source
    if char.toByte == NEWLINE
  } newlineCount += 1
  newlineCount
}
```

The time can be significantly reduced by using the `getLines` method to retrieve one line at a time, and then working through the characters in each line. The following line-counting algorithm counts the same ten million lines in just 23 seconds:

```
// run time: 23 seconds
// use getLines, then count the newline characters
// (redundant for this purpose, i know)
def countLines2(source: io.Source): Long = {
  val NEWLINE = 10
  var newlineCount = 0L
  for {
    line <- source.getLines
    c <- line
    if c.toByte == NEWLINE
  } newlineCount += 1
  newlineCount
}
```

Both algorithms work through each byte in the file, but by using `getLines` in the second algorithm, the run time is reduced dramatically.



Notice that there's the equivalent of two for loops in the second example. If you haven't seen this approach before, here's what the code looks like with two explicit for loops:

```
for (line <- source.getLines) {
  for {
    c <- line
    if c.toByte == NEWLINE
  } newlineCount += 1
}
```

The two approaches are equivalent, but the first is more concise.

12.5. How to Process a CSV File

Problem

You want to process the lines in a CSV file, either handling one line at a time or storing them in a two-dimensional array.

Solution

Combine [Recipe 12.1, “How to Open and Read a Text File”](#) with [Recipe 1.3, “Splitting Strings”](#). Given a simple CSV file like this named *finance.csv*:

```
January, 10000.00, 9000.00, 1000.00
February, 11000.00, 9500.00, 1500.00
March, 12000.00, 10000.00, 2000.00
```

you can process the lines in the file with the following code:

```
object CSVDemo extends App {  
  
  println("Month, Income, Expenses, Profit")  
  val bufferedSource = io.Source.fromFile("/tmp/finance.csv")  
  for (line <- bufferedSource.getLines) {  
    val cols = line.split(",").map(_.trim)  
    // do whatever you want with the columns here  
    println(s"${cols(0)}|${cols(1)}|${cols(2)}|${cols(3)}")  
  }  
  bufferedSource.close  
  
}
```

The magic in that code is this line:

```
val cols = line.split(",").map(_.trim)
```

It splits each line using the comma as a field separator character, and then uses the `map` method to trim each field to remove leading and trailing blank spaces. The resulting output looks like this:

```
January|10000.00|9000.00|1000.00  
February|11000.00|9500.00|1500.00  
March|12000.00|10000.00|2000.00
```

If you prefer named variables instead of accessing array elements, you can change the `for` loop to look like this:

```
for (line <- bufferedSource.getLines) {  
  val Array(month, revenue, expenses, profit) = line.split(",").map(_.trim)  
  println(s"$month $revenue $expenses $profit")  
}
```

If the first line of the file is a header line and you want to skip it, just add `drop(1)` after `getLines`:

```
for (line <- bufferedSource.getLines.drop(1)) { // ...
```

If you prefer, you can also write the loop as a `foreach` loop:

```
bufferedSource.getLines.foreach { line =>  
  rows(count) = line.split(",").map(_.trim)  
  count += 1  
}
```

If you'd like to assign the results to a two-dimensional array, there are a variety of ways to do this. One approach is to create a 2D array, and then use a counter while assigning each line to a row. To do this, you need to know the number of rows in the file before creating the array:

```
object CSVDemo2 extends App {
```

```

val nrows = 3
val ncols = 4
val rows = Array.ofDim[String](nrows, ncols)

val bufferedSource = io.Source.fromFile("/tmp/finance.csv")
var count = 0
for (line <- bufferedSource.getLines) {
  rows(count) = line.split(",").map(_.trim)
  count += 1
}
bufferedSource.close

// print the rows
for (i <- 0 until nrows) {
  println(s"${rows(i)(0)} ${rows(i)(1)} ${rows(i)(2)} ${rows(i)(3)}")
}
}

```

Rather than use a counter, you can do the same thing with the `zipWithIndex` method. This changes the loop to:

```

val bufferedSource = io.Source.fromFile("/tmp/finance.csv")
for ((line, count) <- bufferedSource.getLines.zipWithIndex) {
  rows(count) = line.split(",").map(_.trim)
}
bufferedSource.close

```

If you don't know the number of rows ahead of time, read each row as an `Array[String]`, adding each row to an `ArrayBuffer` as the file is read. That approach is shown in this example, which uses the `using` method introduced in the Solution:

```

import scala.collection.mutable.ArrayBuffer

object CSVDemo3 extends App {

  // each row is an array of strings (the columns in the csv file)
  val rows = ArrayBuffer[Array[String]]()

  // (1) read the csv data
  using(io.Source.fromFile("/tmp/finance.csv")) { source =>
    for (line <- source.getLines) {
      rows += line.split(",").map(_.trim)
    }
  }

  // (2) print the results
  for (row <- rows) {
    println(s"${row(0)}|${row(1)}|${row(2)}|${row(3)}")
  }

  def using[A <: { def close(): Unit }, B](resource: A)(f: A => B): B =
    try {

```



```

        f(resource)
    } finally {
        resource.close()
    }
}

```

An `Array[String]` is used for each row because that's what the `split` method returns. You can convert this to a different collection type, if desired.

Discussion

As you can see, there are a number of ways to tackle this problem. Of all the examples shown, the `zipWithIndex` method probably requires some explanation. The `Iterator Scaladoc` denotes that it creates an iterator that pairs each element produced by this iterator with its index, counting from 0.

So the first time through the loop, `line` is assigned the first line from the file, and `count` is 0. The next time through the loop, the second line of the file is assigned to `line`, and `count` is 1, and so on. The `zipWithIndex` method offers a nice solution for when you need a line counter.

In addition to these approaches, a quick search for “scala csv parser” will turn up a number of competing open source projects that you can use.

See Also

- [Recipe 12.1, “How to Open and Read a Text File”](#), shows both manual and automated ways of closing file resources.
- [Recipe 10.11, “Using `zipWithIndex` or `zip` to Create Loop Counters”](#), provides more examples of the `zipWithIndex` method.
- [The `Iterator` trait](#).

12.6. Pretending that a String Is a File

Problem

Typically for the purposes of testing, you want to pretend that a `String` is a file.

Solution

Because `Scala.fromFile` and `Scala.fromString` both extend `scala.io.Source`, they are easily interchangeable. As long as your method takes a `Source` reference, you can pass it the `BufferedSource` you get from calling `Source.fromFile`, or the `Source` you get from calling `Source.fromString`.

For example, the following method takes a `Source` object and prints the lines it contains:

```
import io.Source

def printLines(source: Source) {
  for (line <- source.getLines) {
    println(line)
  }
}
```

It can be called when the source is constructed from a `String`:

```
val s = Source.fromString("foo\nbar\n")
printLines(s)
```

It can also be called when the source is a file:

```
val f = Source.fromFile("/Users/Al/.bash_profile")
printLines(f)
```

Discussion

When writing unit tests, you might have a method like this that you'd like to test:

```
package foo

object FileUtils {

  def getLinesUppercased(source: io.Source): List[String] = {
    (for (line <- source.getLines) yield line.toUpperCase).toList
  }

}
```

As shown in the following `ScalaTest` tests, you can test the `getLinesUppercased` method by passing it either a `Source` from a file or a `String`:

```
package foo

import org.scalatest.{FunSuite, BeforeAndAfter}
import scala.io.Source

class FileUtilTests extends FunSuite with BeforeAndAfter {

  var source: Source = _
  after { source.close }

  // assumes the file has the string "foo" as its first line
  test("1 - foo file") {
    source = Source.fromFile("/Users/Al/tmp/foo")
    val lines = FileUtils.getLinesUppercased(source)
    assert(lines(0) == "FOO")
  }
}
```

```
test("2 - foo string") {
  source = Source.fromString("foo\n")
  val lines = FileUtils.getLinesUppercased(source)
  assert(lines(0) == "FOO")
}
}
```

If you're interested in making your method easily testable with a `String` instead of a file, define your method to take a `Source` instance.

See Also

- [The Source class](#)
- [The Source object](#)
- [The BufferedSource class](#)

12.7. Using Serialization

Problem

You want to serialize a Scala class and save it as a file, or send it across a network.

Solution

The general approach is the same as Java, but the syntax to make a class serializable is different.

To make a Scala class serializable, extend the `Serializable` trait and add the `@SerialVersionUID` annotation to the class:

```
@SerialVersionUID(100L)
class Stock(var symbol: String, var price: BigDecimal)
extends Serializable {
  // code here ...
}
```

Because `Serializable` is a trait, you can mix it into a class, even if your class already extends another class:

```
@SerialVersionUID(114L)
class Employee extends Person with Serializable ...
```

After marking the class serializable, use the same techniques to write and read the objects as you did in Java, including the [Java “deep copy” technique that uses serialization](#).

Discussion

The following code demonstrates the proper approach. The comments in the code explain the process:

```
import java.io._

// create a serializable Stock class
@SerialVersionUID(123L)
class Stock(var symbol: String, var price: BigDecimal)
extends Serializable {
  override def toString = f"$symbol%s is ${price.toDouble}%.2f"
}

object SerializationDemo extends App {

  // (1) create a Stock instance
  val nflx = new Stock("NFLX", BigDecimal(85.00))

  // (2) write the instance out to a file
  val oos = new ObjectOutputStream(new FileOutputStream("/tmp/nflx"))
  oos.writeObject(nflx)
  oos.close

  // (3) read the object back in
  val ois = new ObjectInputStream(new FileInputStream("/tmp/nflx"))
  val stock = ois.readObject().asInstanceOf[Stock]
  ois.close

  // (4) print the object that was read back in
  println(stock)
}
```

This code prints the following output when run:

```
NFLX is 85.00
```

See Also

- [The Serializable trait](#)
- [Recipe 17.3, “Using @SerialVersionUID and Other Annotations”](#)
- [My Java “Deep Copy/Clone” example](#)

12.8. Listing Files in a Directory

Problem

You want to get a list of files that are in a directory, potentially limiting the list of files with a filtering algorithm.

Solution

Scala doesn't offer any different methods for working with directories, so use the `listFiles` method of the Java `File` class. For instance, this method creates a list of all files in a directory:

```
def getListOfFiles(dir: String): List[File] = {  
  val d = new File(dir)  
  if (d.exists && d.isDirectory) {  
    d.listFiles.filter(_.isFile).toList  
  } else {  
    List[File]()  
  }  
}
```

The REPL demonstrates how you can use this method:

```
scala> import java.io.File  
import java.io.File  
  
scala> val files = getListOfFiles("/tmp")  
files: List[java.io.File] = List(/tmp/foo.log, /tmp/Files.scala.swp)
```

Note that if you're sure that the file you're given is a directory and it exists, you can shorten this method to just the following code:

```
def getListOfFiles(dir: File): List[File] =  
  dir.listFiles.filter(_.isFile).toList
```

Discussion

If you want to limit the list of files that are returned based on their filename extension, in Java, you'd implement a `FileFilter` with an `accept` method to filter the filenames that are returned. In Scala, you can write the equivalent code without requiring a `FileFilter`. Assuming that the `File` you're given represents a directory that is known to exist, the following method shows how to filter a set of files based on the filename extensions that should be returned:

```
import java.io.File  
  
def getListOfFiles(dir: File, extensions: List[String]): List[File] =  
{  
  dir.listFiles.filter(_.isFile).toList.filter { file =>
```

```

        extensions.exists(file.getName.endsWith(_))
    }
}

```

You can call this method as follows to list all WAV and MP3 files in a given directory:

```

val okFileExtensions = List("wav", "mp3")
val files = getListOfFiles(new File("/tmp"), okFileExtensions)

```

As long as this method is given a directory that exists, this method will return an empty List if no matching files are found:

```

scala> val files = getListOfFiles(new File("/Users/Al"), okFileExtensions)
files: List[java.io.File] = List()

```

This is nice, because you can use the result normally, without having to worry about a null value:

```

scala> files.foreach(println)

(no output or errors, because an empty List was returned)

```

See Also

[The Java File class](#)

12.9. Listing Subdirectories Beneath a Directory

Problem

You want to generate a list of subdirectories in a given directory.

Solution

Use a combination of the Java File class and Scala collection methods:

```

// assumes that dir is a directory known to exist
def getListOfSubDirectories(dir: File): List[String] =
    dir.listFiles
      .filter(_.isDirectory)
      .map(_.getName)
      .toList

```

This algorithm does the following:

- Uses the `listFiles` method of the `File` class to list all the files in the given directory as an `Array[File]`.
- The `filter` method trims that list to contain only directories.

- `map` calls `getName` on each file to return an array of directory names (instead of `File` instances).
- `toList` converts that to a `List[String]`.

Calling `toList` isn't necessary, but if it isn't used, the method should be declared to return `Array[String]`.

This method can be used like this:

```
getListOfSubDirectories(new File("/Users/Al")).foreach(println)
```

As mentioned, this method returns a `List[String]`. If you'd rather return a `List[File]`, write the method as follows, dropping the `map` method call:

```
dir.listFiles.filter(_.isDirectory).toList
```

Discussion

This problem provides a good way to demonstrate the differences between writing code in a functional style versus writing code in an imperative style.

When a developer first comes to Scala from Java, she might write a more Java-like (imperative) version of that method as follows:

```
def getListOfSubDirectories1(dir: File): List[String] = {
  val files = dir.listFiles
  val dirNames = collection.mutable.ArrayBuffer[String]()
  for (file <- files) {
    if (file.isDirectory) {
      dirNames += file.getName
    }
  }
  dirNames.toList
}
```

After getting more comfortable with Scala, she'd realize the code can be shortened. One simplification is that she can eliminate the need for the `ArrayBuffer` by using a `for` loop with a `yield` expression. Because the method should return a `List[String]`, the `for` loop is made to `yield file.getName`, and the `for` loop result is assigned to the variable `dirs`. Finally, `dirs` is converted to a `List` in the last line of the method, and it's returned from there:

```
def getListOfSubDirectories2(dir: File): List[String] = {
  val files = dir.listFiles
  val dirs = for {
    file <- files
    if file.isDirectory
  } yield file.getName
  dirs.toList
}
```

Although there's nothing wrong with this code—indeed, some programmers prefer writing for comprehensions to using `map`—at some point, as the developer gets more comfortable with the Scala collections and FP style, she'll realize the *intention* of the code is to create a filtered list of files, and using the `filter` method on the collection to return only directories will come to mind. Also, when she sees a `for/yield` combination, she should think, “map method,” and in short order, she'll be at the original solution.

12.10. Executing External Commands

Problem

You want to execute an external (system) command from within a Scala application. You're not concerned about the output from the command, but you are interested in its exit code.

Solution

To execute external commands, use the methods of the `scala.sys.process` package. There are three primary ways to execute external commands:

- Use the `!` method to execute the command and get its exit status.
- Use the `!!` method to execute the command and get its output.
- Use the `lines` method to execute the command in the background and get its result as a `Stream`.

This recipe demonstrates the `!` method, and the next recipe demonstrates the `!!` method. The `lines` method is shown in the Discussion of this recipe.

To execute a command and get its exit status, import the necessary members and run the desired command with the `!` method:

```
scala> import sys.process._
import sys.process._

scala> "ls -al".!
total 64
drwxr-xr-x 10 Al  staff   340 May 18 18:00 .
drwxr-xr-x  3 Al  staff   102 Apr  4 17:58 ..
-rw-r--r--  1 Al  staff   118 May 17 08:34 Foo.sh
-rw-r--r--  1 Al  staff  2727 May 17 08:34 Foo.sh.jar
res0: Int = 0
```

When using the `!` method, you can get the exit code of the command that was run:

```
scala> val exitCode = "ls -al".!
total 64
drwxr-xr-x 10 Al  staff   340 May 18 18:00 .
```



```
drwxr-xr-x  3 Al  staff   102 Apr  4 17:58 ..
-rw-r--r--  1 Al  staff   118 May 17 08:34 Foo.sh
-rw-r--r--  1 Al  staff 2727 May 17 08:34 Foo.sh.jar
result: Int = 0
```

```
scala> println(exitCode)
0
```

Both of those examples work because of an implicit conversion that adds the `!` method to a `String` when you add the import statement shown.

Discussion

I use this technique to execute the `afplay` system command on Mac OS X systems to play sound files in one of my Scala applications, as shown in this method:

```
def playSoundFile(filename: String): Int = {
  val cmd = "afplay " + filename
  val exitCode = cmd.!.
  exitCode
}
```

That method attempts to play the given filename as a sound file with the `afplay` command, and returns the `exitCode` from the command. This method can be shortened to just one line, but I prefer the approach shown because it's easy to read, especially if you don't execute system processes very often.

To execute system commands I generally just use `!` after a `String`, but the `Seq` approach is also useful. The first element in the `Seq` should be the name of the command you want to run, and subsequent elements are considered to be arguments to it, as shown in these examples:

```
val exitCode = Seq("ls", "-al").!.
val exitCode = Seq("ls", "-a", "-l").!.
val exitCode = Seq("ls", "-a", "-l", "/tmp").!.

```

I've omitted the output from each of those examples, but each command provides the same directory listing you'd get at the Unix command line.

You can also create a `Process` object to execute an external command, if you prefer:

```
val exitCode = Process("ls").!.

```

When running these commands, be aware of whitespace around your command and arguments. All of the following examples fail because of extra whitespace:

```
// beware leading whitespace
scala> " ls".!.
java.io.IOException: Cannot run program "": error=2,
  No such file or directory
  at java.lang.ProcessBuilder.start(ProcessBuilder.java:460)
```

```
scala> val exitCode = Seq(" ls ", "-al").!
java.io.IOException: Cannot run program " ls ": error=2,
  No such file or directory

// beware trailing whitespace
scala> val exitCode = Seq("ls", "-al ").!
ls: -al : No such file or directory
exitCode: Int = 1
```

If you enter the strings yourself, leave the whitespace out, and if you get the strings from user input, be sure to trim them.

Using the lines method

The `lines` method is an interesting alternative to the `!` and `!!` commands. With `lines`, you can immediately execute a command in the background. For instance, the following command will run for a long time on a Unix system and result in a large amount of output:

```
val process = Process("find / -print").lines
```

The variable `process` in this example is a `Stream[String]`. With `lines` running the process in the background, you can either work with the result immediately or at some later point. For instance, you can read from the stream like this:

```
process.foreach(println)
```

The `lines` method throws an exception if the exit status of the command is nonzero. You can catch that with a `try/catch` expression, but if this is a problem, or if you also want to retrieve the standard error from the command, use the `lines_!` method instead of `lines`. The `lines_!` method is demonstrated in [Recipe 12.11](#) and discussed in [Table 12-1](#) in [Recipe 12.19](#).

External commands versus built-in commands

As a final note, you can run any external command from Scala that you can run from the Unix command line. However, there's a big difference between an *external* command and a shell *built-in* command. The `ls` command is an external command that's available on all Unix systems, and can be found as a file in the `/bin` directory:

```
$ which ls
/bin/ls
```

Some other commands that can be used at a Unix command line, such as the `cd` or `for` commands in the Bash shell, are actually built into the shell; you won't find them as files on the filesystem. Therefore, these commands can't be executed unless they're executed from within a shell. See [Recipe 12.13](#), “[Building a Pipeline of Commands](#)” for an example of how to execute a shell built-in command.

12.11. Executing External Commands and Using STDOUT

Problem

You want to run an external command and then use the standard output (STDOUT) from that process in your Scala program.

Solution

Use the `!!` method to execute the command and get the standard output from the resulting process as a `String`.

Just like the `!` command in the previous recipe, you can use `!!` after a `String` to execute a command, but `!!` returns the STDOUT from the command rather than the exit code of the command. This returns a multiline string, which you can process in your application:

```
scala> import sys.process._
import sys.process._

scala> val result = "ls -al" !!
result: String =
"total 64
drwxr-xr-x  10 Al  staff   340 May 18 18:00 .
drwxr-xr-x   3 Al  staff   102 Apr  4 17:58 ..
-rw-r--r--   1 Al  staff   118 May 17 08:34 Foo.sh
-rw-r--r--   1 Al  staff  2727 May 17 08:34 Foo.sh.jar
"

scala> println(result)
total 64
drwxr-xr-x  10 Al  staff   340 May 18 18:00 .
drwxr-xr-x   3 Al  staff   102 Apr  4 17:58 ..
-rw-r--r--   1 Al  staff   118 May 17 08:34 Foo.sh
-rw-r--r--   1 Al  staff  2727 May 17 08:34 Foo.sh.jar
```

If you prefer, you can do the same thing with a `Process` or `Seq` instead of a `String`:

```
val result = Process("ls -al").!!
val result = Seq("ls -al").!!
```

As shown in the previous recipe, using a `Seq` is a good way to execute a system command that requires arguments:

```
val output = Seq("ls", "-al").!!
val output = Seq("ls", "-a", "-l").!!
val output = Seq("ls", "-a", "-l", "/tmp").!!
```

The first element in the `Seq` is the name of the command to be run, and subsequent elements are arguments to the command. The following code segment shows how to run a complex Unix `find` command:

```

val dir = "/Users/Al/tmp"
val searchTerm = "dawn"

val results = Seq("find", dir, "-type", "f", "-exec", "grep", "-il",
                  searchTerm, "{}", ";").!!
println(results)

```

This code is the equivalent of running the following `find` command at the Unix prompt:

```
find /Users/Al/tmp -type f -exec grep -il dawn {} \;
```

If you're not familiar with Unix commands, this command can be read as, "Search all files under the `/Users/Al/tmp` directory for the string `dawn`, ignoring case, and print the names of all files where a match is found."

Discussion

Use the `!` method to get the exit code from a process, or `!!` to get the standard output from a process.

Be aware that attempting to get the standard output from a command exposes you to exceptions that can occur. As a simple example, if you write the following statement to get the exit code of a command using the `!` operator, even though a little extra `STDERR` information is printed in the REPL, `out` is just assigned a nonzero exit code:

```

scala> val out = "ls -l fred" !
ls: fred: No such file or directory
out: Int = 1

```

But if you attempt to get the standard output from the same command using the `!!` method, an exception is thrown, and `out` is not assigned:

```

scala> val out = "ls -l fred" !!
ls: fred: No such file or directory
java.lang.RuntimeException: Nonzero exit value: 1
    many more lines of output ...

```

Unexpected newline characters

When running an external command, you may expect a one-line string to be returned, but you can get a newline character as well:

```

scala> val dir = "pwd" !!
dir: String =
"/Users/Al/Temp
"

```

When this happens, just trim the result:

```

scala> val dir = "pwd".!!.trim
dir: java.lang.String = /Users/Al/Temp

```

Using the `lines_!` method

You may want to check to see whether an executable program is available on your system. For instance, suppose you wanted to know whether the `hadoop2` executable is available on a Unix-based system. A simple way to handle this situation is to use the Unix `which` command with the `!` method, where a nonzero exit code indicates that the command isn't available:

```
scala> val executable = "which hadoop2".!  
executable: Int = 1
```

If the value is nonzero, you know that the executable is not available on the current system. More accurately, it may be on the system, but it's not on the `PATH` (or much less likely, the `which` command is not available).

Another way to handle this situation is to use the `lines_!` method. This can be used to return a `Some` or `None`, depending on whether or not the `hadoop` command is found by `which`. The syntax for the `lines_!` method is shown in this example:

```
val executable = "which hadoop2".lines_!.headOption
```

In the Scala REPL, you can see that if the executable isn't available on the current system, this expression returns `None`:

```
scala> val executable = "which hadoop2".lines_!.headOption  
executable: Option[String] = None
```

Conversely, if the command is found, the expression returns a `Some`:

```
scala> val executable = "which ls".lines_!.headOption  
executable: Option[String] = Some(/bin/ls)
```

Note the call to the `headOption` method at the end of this pipeline. You call this method because the `lines_!` method returns a `Stream`, but you want the `Option` immediately.

See [Recipe 12.19](#) for a description of the `lines_!` method.

12.12. Handling STDOUT and STDERR for External Commands

Problem

You want to run an external command and get access to both its `STDOUT` and `STDERR`.

Solution

The simplest way to do this is to run your commands as shown in previous recipes, and then capture the output with a `ProcessLogger`. This Scala shell script demonstrates the approach:

```
#!/bin/sh
exec scala "$@" "$@"
!#

import sys.process._

val stdout = new StringBuilder
val stderr = new StringBuilder

val status = "ls -al FRED" ! ProcessLogger(stdout append _, stderr append _)
println(status)
println("stdout: " + stdout)
println("stderr: " + stderr)
```

When this script is run, the `status` variable contains the exit status of the command. The `stdout` variable contains the STDOUT if the command is successful (such as with `ls -al`), and `stderr` contains the STDERR from the command if there are problems. If the command you're running writes to both STDOUT and STDERR, both `stdout` and `stderr` will contain data.

For instance, assuming you don't run the following command as root, changing the `status` expression in the script to the following code should generate output to both STDOUT and STDERR on a Unix system:

```
val status = Seq("find", "/usr", "-name", "make") ! ⚡
ProcessLogger(stdout append _, stderr append _)
```

Running the script with this command on a Mac OS X (Unix) system, I correctly get the following exit status, STDOUT, and STDERR output:

```
scala> val status = Seq("find", "/usr", "-name", "make") ! ProcessLogger(stdout⚡
append _, stderr append _)
status: Int = 1

scala> println(stdout)
/usr/bin/make

scala> println(stderr)
find: /usr/local/mysql-5.0.67-osx10.5-x86/data: Permission denied
```

Depending on your needs, this can get much more complicated very quickly. The Scaladoc states, “If one desires full control over input and output, then a `ProcessIO` can be used with `run`.” See the `scala.sys.process` API documentation for the `ProcessLogger` and `ProcessIO` classes for more examples.

See Also

The process [package object documentation](#) includes many details and examples.

12.13. Building a Pipeline of Commands

Problem

You want to execute a series of external commands, redirecting the output from one command to the input of another command, i.e., you want to pipe the commands together.

Solution

Use the `#|` method to pipe the output from one command into the input stream of another command. When doing this, use `!` at the end of the pipeline if you want the exit code of the pipeline, or `!!` if you want the output from the pipeline.

The `!!` approach is shown in the following example where the output from the `ps` command is piped as the input to the `wc` command:

```
import sys.process._
val numProcs = ("ps auxw" #| "wc -l").!!.trim
println(s"#procs = $numProcs")
```

Because the output from the `ps` command is piped into a line count command (`wc -l`), that code prints the number of processes running on a Unix system. The following command creates a list of all Java processes running on the current system:

```
val javaProcs = ("ps auxw" #| "grep java").!!.trim
```

There are other ways to write these commands, but because I usually end up trimming the result I get back from commands, I find this syntax to be the most readable approach.

Discussion

If you come from a Unix background, the `#|` command is easy to remember because it's just like the Unix pipe symbol, but preceded by a `#` character (`#|`). In fact, with the exception of the `###` operator (which is used instead of the Unix `;` symbol), the entire library is consistent with the equivalent Unix commands.

Note that attempting to pipe commands together inside a `String` and then execute them with `!` won't work:

```
// won't work
val result = ("ls -al | grep Foo").!!
```

This doesn't work because the piping capability comes from a shell (Bourne shell, Bash, etc.), and when you run a command like this, you don't have a shell.

To execute a series of commands in a shell, such as the Bourne shell, use a `Seq` with multiple parameters, like this:

```
val r = Seq("/bin/sh", "-c", "ls | grep .scala").!!
```

This approach runs the `ls | grep .scala` command inside a Bourne shell instance. A quick run in the REPL demonstrates this:

```
scala> val r = Seq("/bin/sh", "-c", "ls | grep .scala").!!
r: String =
Bar.scala
Baz.scala
Foo.scala
"
```

However, note that when using `!!`, you'll get the following exception if there are no `.scala` files in the directory:

```
java.lang.RuntimeException: Nonzero exit value: 1
```

I've found it best to wrap commands executed with `!!` in a `try/catch` expression.

See Also

My tutorial, [“How to Execute a System Command Pipeline in Java,”](#) discusses the need for a shell when piping commands.

12.14. Redirecting the STDOUT and STDIN of External Commands

Problem

You want to redirect the standard output (STDOUT) and standard input (STDIN) when running external commands. For instance, you may want to redirect STDOUT to log the output of an external command to a file.

Solution

Use `#>` to redirect STDOUT, and `#<` to redirect STDIN.

When using `#>`, place it after your command and before the filename you want to write to, just like using `>` in Unix:

```
import sys.process._
import java.io.File

("ls -al" #> new File("files.txt")).!
("ps aux" #> new File("processes.txt")).!
```

You can also pipe commands together and then write the resulting output to a file:

```
("ps aux" #| "grep http" #> new File("http-processes.out")).!
```


Get the exit status from a command like this:

```
val status = ("cat /etc/passwd" #> new File("passwd.copy")).!  
println(status)
```

You can also download a URL and write its contents to a file:

```
import sys.process._  
import scala.language.postfixOps  
import java.net.URL  
import java.io.File  
  
new URL("http://www.google.com") #> new File("Output.html") !
```

I don't redirect STDIN too often, but this example shows one possible way to read the contents of the `/etc/passwd` file into a variable using `#<` and the Unix `cat` command:

```
import scala.sys.process._  
import java.io.File  
  
val contents = ("cat" #< new File("/etc/passwd")).!!  
println(contents)
```

Discussion

The `#>` and `#<` operators generally work like their equivalent `>` and `<` Unix commands, though you can also use them for other purposes, such as using `#>` to write from one `ProcessBuilder` to another, like a pipeline:

```
val numLines = ("cat /etc/passwd" #> "wc -l").!!.trim  
println(numLines)
```

The `ProcessBuilder` Scaladoc states that `#>` and `#<` “may take as input either another `ProcessBuilder`, or something else such as a `java.io.File` or a `java.lang.InputStream`.”

As mentioned, the Scala process commands are consistent with the standard Unix redirection symbols, so you can also append to a file with the `#>>` method:

```
// append to a file  
("ps aux" #>> new File("ps.out")).!
```

Regarding the use of the `URL` and `File` classes, the Scaladoc states that instances of `java.io.File` and `java.net.URL` can be used as input to other processes, and a `File` instance can also be used as output. This was demonstrated in the Solution with the ability to download the contents from a `URL` and write it to a file with the `#>` operator.

See Also

- [The process package object](#)
- [The Scala ProcessBuilder trait](#)
- [The Scala Process trait](#)

12.15. Using AND (&&) and OR (||) with Processes

Problem

You want to use the equivalent of the Unix `&&` and `||` commands to perform an `if/then/else` operation when executing external commands.

Solution

Use the Scala operators `#&&` and `#||`, which mirror the Unix `&&` and `||` operators:

```
val result = ("ls temp" #&& "rm temp" #|| "echo 'temp' not found").!!.trim
```

This command can be read as, “Run the `ls` command on the file `temp`, and if it’s found, remove it, otherwise, print the ‘not found’ message.”

In practice, this can be a little more difficult than shown, because these commands usually involve the use of a wildcard operator. For instance, even if there are `.scala` files in the current directory, the following attempt to compile them using `#&&` and `#||` will fail because of the lack of wildcard support:

```
scala> ("ls *.scala" #&& "scalac *.scala" #|| "echo no files to compile").!  
ls: *.scala: No such file or directory  
no files to compile  
res0: Int = 0
```

To get around this problem, use the formula shared in [Recipe 12.16, “Handling Wildcard Characters in External Commands”](#) running each command in a shell (and also separating each command to make the `#&&` and `#||` command readable):

```
#!/bin/sh  
exec scala "$@" "$@"  
!  
  
import scala.sys.process._  
  
val filesExist = Seq("/bin/sh", "-c", "ls *.scala")  
val compileFiles = Seq("/bin/sh", "-c", "scalac *.scala")  
(filesExist #&& compileFiles #|| "echo no files to compile").!!
```

That script compiles all *.scala* files in the current directory.

12.16. Handling Wildcard Characters in External Commands

Problem

You want to use a Unix shell wildcard character, such as ***, in an external command.

Solution

In general, the best thing you can do when using a wildcard character like *** is to run your command while invoking a Unix shell. For instance, if you have *.scala* files in the current directory and try to list them with the following command, the command will fail:

```
scala> import scala.sys.process._
import scala.sys.process._

scala> "ls *.scala".!
ls: *.scala: No such file or directory
res0: Int = 1
```

But by running the same command inside a Bourne shell, the command now correctly shows the *.scala* files (and returns the exit status of the command):

```
scala> val status = Seq("/bin/sh", "-c", "ls *.scala").!
AndOrTest.scala
Console.scala
status: Int = 0
```

Discussion

Putting a shell wildcard character like *** into a command doesn't work because the *** needs to be interpreted and expanded by a shell, like the Bourne or Bash shells. In this example, even though there are files in the current directory named *AndOrTest.scala* and *Console.scala*, the first attempt doesn't work. These other attempts will also fail as a result of the same problem:

```
scala> "echo *".!
*
res0: Int = 0

scala> Seq("grep", "-i", "foo", "*.scala").!
grep: *.scala: No such file or directory
res1: Int = 2

scala> Seq("ls", "*.scala").!
```

```
ls: *.scala: No such file or directory
res2: Int = 1
```

In each example, you can make these commands work by invoking a shell in the first two parameters to a Seq:

```
val status = Seq("/bin/sh", "-c", "echo *").!
val status = Seq("/bin/sh", "-c", "ls *.scala").!
val status = Seq("/bin/sh", "-c", "grep -i foo *.scala").!
```

An important part of this recipe is using the `-c` argument of the `/bin/sh` command. The `sh` manpage describes this parameter as follows:

```
-c string
```

```
If the -c option is present, then commands are read from string.
If there are arguments after the string, they are assigned to the
positional parameters, starting with $0.
```

As an exception to this general rule, the `-name` option of the `find` command may work because it treats the `*` character as a wildcard character itself. As a result, the following `find` command finds the two files in the current directory without having to be run in a shell:

```
scala> val status = Seq("find", ".", "-name", "*.scala", "-type", "f").!
./AndOrTest.scala
./Console.scala
status: Int = 0
```

However, as shown, other commands generally require that the `*` wildcard character be interpreted and expanded by a shell.

See Also

- [“How to Execute a Command Pipeline in Java”](#)
- [“Execute System Processes with Java Process and ProcessBuilder”](#)

12.17. How to Run a Process in a Different Directory

Problem

You want to use another directory as the base directory when running an external command.

Solution

Use one of the `Process` factory methods, setting your command and the desired directory, then running the process with the usual `!` or `!!` commands. The following example runs the `ls` command with the `-al` arguments in the `/var/tmp` directory:

```
import sys.process._
import java.io.File

object Test extends App {

  val output = Process("ls -al", new File("/tmp")).!!
  println(output)

}
```

To run that same command in the current directory, just remove the second parameter when creating the `Process`:

```
val p = Process("ls -al")
```

You can use another `Process` factory method to set system environment variables, i.e., those that can be seen at the shell command line with `set` or `env`. See the next recipe for examples of that method.

12.18. Setting Environment Variables When Running Commands

Problem

You need to set one or more environment variables when running an external command.

Solution

Specify the environment variables when calling a `Process` factory method (an `apply` method in the `Process` object).

The following example shows how to run a shell script in a directory named `/home/al/bin` while also setting the `PATH` environment variable:

```
val p = Process("runFoo.sh",
  new File("/Users/Al/bin"),
  "PATH" -> ".:usr/bin:/opt/scala/bin")

val output = p.!!
```

To set multiple environment variables at one time, keep adding them at the end of the `Process` constructor:

```
val output = Process("env",
  None,
  "VAR1" -> "foo",
  "VAR2" -> "bar")
```

These examples work because of the overloaded `apply` methods in the `Process` object. For instance, one method takes a `File` for the directory parameter, and another method takes an `Option[File]` for that parameter. This second approach lets you use `None` to indicate the current directory.

The ability to specify multiple environment variables when calling a `Process` factory method works because the `apply` methods accept a `varargs` argument of the type `(String, String)*` for their last argument. This means “a variable number of tuple arguments.”

See Also

The `Process` object

12.19. An Index of Methods to Execute External Commands

The following tables list the methods of the `scala.sys.process` package that you can use when running external (system) commands.

Table 12-1 lists the methods that you can use to execute system commands.

Table 12-1. Methods to execute system commands

Method	Description
<code>!</code>	Runs the command and returns its exit code. Blocks until all external commands exit. If used in a chain, returns the exit code of the last command in the chain.
<code>!!</code>	Runs the command (or command pipe/chain), and returns the output from the command as a <code>String</code> . Blocks until all external commands exit. Warning: throws exceptions when the command's exit status is nonzero.
<code>run</code>	Returns a <code>Process</code> object immediately while running the process in the background. The <code>Process</code> can't currently be polled to see if it has completed.
<code>lines</code>	Returns immediately, while running the process in the background. The output that's generated is provided through a <code>Stream[String]</code> . Getting the next element of the <code>Stream</code> may block until it becomes available. Throws an exception if the return code is not zero; if this isn't desired, use the <code>lines_!</code> method. Example: <pre>scala> val x = Process("ls").lines x: Stream[String] = Stream(Bar.scala, ?)</pre>
<code>lines_!</code>	Like the <code>lines</code> method, but <code>STDERR</code> output is sent to the <code>ProcessLogger</code> you provide. Per the Scaladoc, “If the process exits with a nonzero value, the <code>Stream</code> will provide all lines up to termination but will not throw an exception.” Demonstrated in Recipe 12.11 .

Table 12-2 lists the methods that you can use to redirect STDIN and STDOUT when external commands are executed.

Table 12-2. *Methods to redirect STDIN and STDOUT*

Methods	Description
#<	Read from STDIN
#>	Write to STDOUT
#>>	Append to STDOUT

Table 12-3 lists the methods that you can use to combine (pipe) external commands.

Table 12-3. *Methods to combine external commands*

Methods	Description
cmd1 # cmd2	The output of the first command is used as input to the second command, like a Unix shell pipe.
cmd1 ### cmd2	cmd1 and cmd2 will be executed in sequence, one after the other. This is like the Unix ; operator, but ; is a reserved keyword in Scala.
cmd1 #> cmd2	Normally used to write to STDOUT but can be used like # to chain commands together. Example: <pre>scala> ("ps aux" #> "grep java" #> "wc -l").!!!.trim res0: String = 2</pre>
cmd1 #&& cmd2	Run cmd2 if cmd1 runs successfully (i.e., it has an exit status of 0).
cmd1 # cmd2	Run cmd2 if cmd1 has an unsuccessful (nonzero) exit status.
cmd1 #&& cmd2 # cmd3	Run cmd2 if cmd1 has a successful exit status, otherwise, run cmd3.

The primary online documentation for the Scala process API is at these URLs:

- [The `scala.sys.process` package object](#)
- [The `ProcessBuilder` trait](#)

Actors and Concurrency

Introduction

In Scala you can still use Java threads, but the Actor model is the preferred approach for concurrency. The Actor model is at a much higher level of abstraction than threads, and once you understand the model, it lets you focus on solving the problem at hand, rather than worrying about the low-level problems of threads, locks, and shared data.

Although earlier versions of Scala included its original Actors library, Scala 2.10.0 began the official transition to the **Akka actor library** from Typesafe, which is more robust than the original library. Scala 2.10.1 then deprecated the original *scala.actors* library.

In general, actors give you the benefit of offering a high level of abstraction for achieving concurrency and parallelism. Beyond that, the Akka actor library adds these additional benefits:

- Lightweight, event-driven processes. The documentation states that there can be approximately 2.7 million actors per gigabyte of RAM.
- Fault tolerance. Akka actors can be used to create “self-healing systems.” (The Akka “team blog” is located at <http://letitcrash.com/>.)
- Location transparency. Akka actors can span multiple JVMs and servers; they’re designed to work in a distributed environment using pure message passing.

A “high level of abstraction” can also be read as “ease of use.” It doesn’t take very long to understand the Actor model, and once you do, you’ll be able to write complex, concurrent applications much more easily than you can with the basic Java libraries. I wrote a speech interaction application (speech recognition input, text-to-speech output) named **SARAH** that makes extensive use of Akka actors, with agents constantly working on tasks in the background. Writing this code with actors was *much* easier than the equivalent threading approach.

I like to think of an actor as being like a web service on someone else's servers that I can't control. I can send messages to that web service to ask it to do something, or I can query it for information, but I can't reach into the web service to directly modify its state or access its resources; I can only work through its API, which is just like sending immutable messages. In one way, this is a little limiting, but in terms of safely writing parallel algorithms, this is very beneficial.

The Actor Model

Before digging into the recipes in this chapter, it can help to understand the Actor model.

The first thing to understand about the Actor model is the concept of an *actor*:

- An actor is the smallest unit when building an actor-based system, like an object in an OOP system.
- Like an object, an actor encapsulates state and behavior.
- You can't peek inside an actor to get its state. You can send an actor a message requesting state information (like asking a person how they're feeling), but you can't reach in and execute one of its methods, or access its fields.
- An actor has a mailbox (an inbox), and its purpose in life is to process the messages in its mailbox.
- You communicate with an actor by sending it an immutable message. These messages go into the actor's mailbox.
- When an actor receives a message, it's like taking a letter out of its mailbox. It opens the letter, processes the message using one of its algorithms, then moves on to the next letter in the mailbox. If there are no more messages, the actor waits until it receives one.

In an application, actors form hierarchies, like a family, or a business organization:

- The Typesafe team recommends thinking of an actor as being like a person, such as a person in a business organization.
- An actor has one parent (supervisor): the actor that created it.
- An actor may have children. Thinking of this as a business, a president may have a number of vice presidents. Those VPs will have many subordinates, and so on.
- An actor may have siblings. For instance, there may be 10 VPs in an organization.
- A best practice of developing actor systems is to “delegate, delegate, delegate,” especially if behavior will block. In a business, the president may want something done, so he delegates that work to a VP. That VP delegates work to a manager, and so on, until the work is eventually performed by one or more subordinates.

- Delegation is important. Imagine that the work takes several man-years. If the president had to handle that work himself, he couldn't respond to other needs (while the VPs and other employees would all be idle).

A final piece of the Actor model is handling failure. When performing work, something may go wrong, and an exception may be thrown. When this happens, an actor suspends itself and all of its children, and sends a message to its supervisor, signaling that a failure has occurred. (A bit like Scotty calling Captain Kirk with a problem.)

Depending on the nature of the work and the nature of the failure, the supervising actor has a choice of four options at this time:

- Resume the subordinate, keeping its internal state
- Restart the subordinate, giving it a clean state
- Terminate the subordinate
- Escalate the failure

In addition to those general statements about actors, there are a few important things to know about Akka's implementation of the Actor model:

- You can't reach into an actor to get information about its state. When you instantiate an Actor in your code, Akka gives you an `ActorRef`, which is essentially a façade between you and the actor.
- Behind the scenes, Akka runs actors on real threads; many actors may share one thread.
- There are different mailbox implementations to choose from, including variations of unbounded, bounded, and priority mailboxes. You can also create your own mailbox type.
- Akka does not let actors scan their mailbox for specific messages.
- When an actor terminates (intentionally or unintentionally), messages in its mailbox go into the system's "dead letter mailbox."

Hopefully these notes about the general Actor model, and the Akka implementation specifically, will be helpful in understanding the recipes in this chapter.

Other Features

Scala offers other conveniences for writing code that performs operations in parallel. A *future* can be used for simple, "one off" tasks that require concurrency. The Scala collections library also includes special *parallel collections*, which can be used to improve the performance of large collections and certain algorithms.



There are interesting debates about what the terms *concurrency* and *parallelism* mean. I tend to use them interchangeably, but for one interesting discussion of their differences—such as concurrency being one vending machine with two lines, and parallelism being two vending machines and two lines—see [this blog post](#).

13.1. Getting Started with a Simple Actor

Problem

You want to begin using actors to build concurrency into your applications.

Solution

Create an actor by extending the `akka.actor.Actor` class and writing a `receive` method in your class. The `receive` method should be implemented with a case statement that allows the actor to respond to the different messages it receives.

To demonstrate this, create an SBT project directory named *HelloAkka*, move into that directory, and then add the necessary Akka resolver and dependency information to your *build.sbt* file:

```
name := "Hello Test #1"

version := "1.0"

scalaVersion := "2.10.0"

resolvers += "Typesafe Repository" at ↵
            "http://repo.typesafe.com/typesafe/releases/"

libraryDependencies += "com.typesafe.akka" %% "akka-actor" % "2.1.2"
```



At the time of this writing, the Akka actor library is being migrated into the Scala distribution, but it's still necessary to include the library as a dependency in your SBT *build.sbt* file (or download the necessary JAR files manually). This may change in the future, in which case the dependencies shown in this chapter may not be necessary.

Next, define an actor that responds when it receives the `String` literal `hello` as a message. To do this, save the following source code to a file named *Hello.scala* in the root directory of your SBT project. Notice how the literal `hello` is used in the first case statement in the `receive` method of the `HelloActor` class:

```

import akka.actor.Actor
import akka.actor.ActorSystem
import akka.actor.Props

class HelloActor extends Actor {
  def receive = {
    case "hello" => println("hello back at you")
    case _      => println("huh?")
  }
}

object Main extends App {

  // an actor needs an ActorSystem
  val system = ActorSystem("HelloSystem")

  // create and start the actor
  val helloActor = system.actorOf(Props[HelloActor], name = "helloactor")

  // send the actor two messages
  helloActor ! "hello"
  helloActor ! "buenos dias"

  // shut down the system
  system.shutdown

}

```

Then run the application like this:

```
$ sbt run
```

After SBT downloads the Akka JAR files and their dependencies, you should see the following output from the `println` statements in the `HelloActor` class:

```

[info] Running Main
hello back at you
huh?

```

Discussion

Here's a step-by-step description of the code:

- The `import` statements import the members that are needed.
- An Actor named `HelloActor` is defined.
- `HelloActor`'s behavior is implemented by defining a `receive` method, which is implemented using a match expression.
- When `HelloActor` receives the String literal `hello` as a message, it prints the first reply, and when it receives any other type of message, it prints the second reply.
- The `Main` object is created to test the actor.

- In `Main`, an `ActorSystem` is needed to get things started, so one is created. The `ActorSystem` takes a `name` as an argument, so give the system a meaningful name. The name must consist of only the `[a-zA-Z0-9]` characters, and zero or more hyphens, and a hyphen can't be used in the leading space.
- Actors can be created at the `ActorSystem` level, or inside other actors. At the `ActorSystem` level, actor instances are created with the `system.actorOf` method. The `helloActor` line shows the syntax to create an Actor with a constructor that takes no arguments.
- Actors are automatically started (asynchronously) when they are created, so there's no need to call any sort of "start" or "run" method.
- Messages are sent to actors with the `!` method, and `Main` sends two messages to the actor with the `!` method: `hello` and `buenos días`.
- `helloActor` responds to the messages by executing its `println` statements.
- The `ActorSystem` is shut down.

That's all you need to create and use your first Akka Actor.

Details

When implementing the behavior of an Akka actor, you should define a `receive` method using a `match` expression, as shown in the example. Your method should handle all potential messages that can be sent to the actor; otherwise, an `UnhandledMessage` will be published to the `ActorSystem`'s `EventStream`. As a practical matter, this means having the catch-all case `_` line in your `match` expression.

In this example, messages were sent to the `HelloActor` class as `String` literals, but other recipes will show how to send messages to actors using other types. Messages should be immutable, so for simple examples, a `String` works well.

ActorSystem

The API documentation describes an `ActorSystem` like this:

"An actor system is a hierarchical group of actors which share common configuration, e.g. dispatchers, deployments, remote capabilities and addresses. It is also the entry point for creating or looking up actors."

An `ActorSystem` is the structure that allocates one or more threads for your application, so you typically create one `ActorSystem` per (logical) application.

As an example, I wrote a "speech interaction" application named `SARAH` that lets me interact with a Mac OS X computer using only voice commands. Besides allowing interactive commands, `SARAH` also runs background tasks to check my email, notify me of Facebook and Twitter events, stock prices, etc.

SARAH uses a plug-in architecture, so there are plug-ins for each major area of functionality (such as an email plug-in, Facebook plug-in, Twitter plug-in, etc.). A plug-in typically has one parent actor that delegates work to child actors as necessary. All of these plug-ins run under one `ActorSystem`. When SARAH starts, it starts the `ActorSystem` using the same method shown in the Solution. Once started, it creates three main actors named `brain`, `ears`, and `mouth`, and then starts its plug-ins.

As an interesting experiment with the `ActorSystem`, remove the `system.shutdown` line at the end of the `Main` object. You'll see that the application doesn't terminate, because the actors and system are still running. (Press Control-C to terminate the application.)

ActorRef

When you call the `actorOf` method on an `ActorSystem`, it starts the actor asynchronously and returns an instance of an `ActorRef`. This reference is a “handle” to the actor, which you can think of as being a façade or broker between you and the actual actor. This façade keeps you from doing things that would break the Actor model, such as reaching into the Actor instance and attempting to directly mutate variables. Tasks like this should only be done by passing messages to the actor, and the hands-off approach of an `ActorRef` helps reinforce proper programming practices.

(Again, think of an actor as a person you can only communicate with by placing messages in his mailbox.)

The Akka documentation states that an `ActorRef` has these qualities:

- It is immutable.
- It has a one-to-one relationship with the Actor it represents.
- It is serializable and network-aware. This lets you pass the `ActorRef` around the network.

See Also

- The introductory [Akka actor documentation](#)
- The `ActorSystem` class
- The `ActorRef` class

13.2. Creating an Actor Whose Class Constructor Requires Arguments

Problem

You want to create an Akka actor, and you want your actor's constructor to have one or more arguments.

Solution

Create the actor using the syntax shown here, where `HelloActor` takes one constructor parameter:

```
val helloActor = system.actorOf(Props(new HelloActor("Fred")), ↵  
name = "helloactor")
```

Discussion

When creating an actor whose constructor takes one or more arguments, you still use the `Props` class to create the actor, but with a different syntax than when creating an actor whose constructor takes no arguments.

The following code demonstrates the difference between creating an actor with a no-args constructor and an actor that takes at least one constructor parameter:

```
// an actor with a no-args constructor  
val helloActor = system.actorOf(Props[HelloActor], name = "helloactor")  
  
// an actor whose constructor takes one argument  
val helloActor = system.actorOf(Props(new HelloActor("Fred")), ↵  
name = "helloactor")
```

To demonstrate these differences, the following source code is a modified version of the example in [Recipe 13.1](#). Comments are included in the code to highlight the changes:

```
import akka.actor._  
  
// (1) constructor changed to take a parameter  
class HelloActor(myName: String) extends Actor {  
  def receive = {  
    // (2) println statements changed to show the name  
    case "hello" => println(s"hello from $myName")  
    case _       => println(s"'huh?', said $myName")  
  }  
}  
  
object Main extends App {  
  val system = ActorSystem("HelloSystem")  
  // (3) use a different version of the Props constructor
```



```

val helloActor = system.actorOf(
    Props(new HelloActor("Fred")), name = "helloactor")
helloActor ! "hello"
helloActor ! "buenos dias"
system.shutdown
}

```

As shown in this example, if your actor takes more than one argument, include those arguments in the constructor call. If the `HelloActor` constructor required both a first and last name, you'd specify them like this:

```

Props(new HelloActor("John", "Doe")), name = "helloactor")

```

Remember that an actor instance is instantiated and started when the `actorOf` method is called, so the only ways to set a property in an actor instance are:

- By sending the actor a message
- In the actor's constructor
- In its `preStart` method

You've already seen how to send a message to an actor and use its constructor. The `preStart` method is demonstrated in [Recipe 13.4, "Understanding the Methods in the Akka Actor Lifecycle"](#).

See Also

The `Props` class

13.3. How to Communicate Between Actors

Problem

You're building an actor-based application and want to send messages between actors.

Solution

Actors should be sent immutable messages with the `!` method.

When an actor receives a message from another actor, it also receives an implicit reference named `sender`, and it can use that reference to send a message back to the originating actor.

The general syntax to send a message to an actor is:

```

actorInstance ! message

```

For example, if you have an actor instance named `car`, you can send it a `start` message like this:

```
car ! "start"
```

In this case, the message is the `String` literal `start`. The `car` actor should receive this message in a match expression in its `receive` method, and from there it can send a message back to whoever sent the `start` message. A simplified version of a `receive` method for `car` might look like this:

```
def receive = {  
  case "start" =>  
    val result = tryToStart()  
    sender ! result  
  case _ => // do nothing  
}
```

As mentioned, the `sender` instance is implicitly made available to your actor. If you just want to send a message back to the code that sent you a message, that's all you have to do.

Discussion

To demonstrate a more complicated example of actors communicating, the following code shows how to send messages back and forth between Akka actors. It was inspired by the “Ping Pong” threading example in the book by James Gosling et al., *The Java Programming Language* (Addison-Wesley Professional):

```
import akka.actor._  
  
case object PingMessage  
case object PongMessage  
case object StartMessage  
case object StopMessage  
  
class Ping(pong: ActorRef) extends Actor {  
  var count = 0  
  def incrementAndPrint { count += 1; println("ping") }  
  def receive = {  
    case StartMessage =>  
      incrementAndPrint  
      pong ! PingMessage  
    case PongMessage =>  
      incrementAndPrint  
      if (count > 99) {  
        sender ! StopMessage  
        println("ping stopped")  
        context.stop(self)  
      } else {  
        sender ! PingMessage  
      }  
  }  
}
```

```

    }
    case _ => println("Ping got something unexpected.")
  }
}

class Pong extends Actor {
  def receive = {
    case PingMessage =>
      println(" pong")
      sender ! PongMessage
    case StopMessage =>
      println("pong stopped")
      context.stop(self)
    case _ => println("Pong got something unexpected.")
  }
}

object PingPongTest extends App {
  val system = ActorSystem("PingPongSystem")
  val pong = system.actorOf(Props[Pong], name = "pong")
  val ping = system.actorOf(Props(new Ping(pong)), name = "ping")
  // start the action
  ping ! StartMessage

  // commented-out so you can see all the output
  //system.shutdown
}

```

Actors should communicate by sending immutable messages between each other. In this case there are four messages, and they're defined using case objects: `PingMessage`, `PongMessage`, `StartMessage`, and `StopMessage`.

The `PingPongTest` object performs the following work:

1. Creates an `ActorSystem`.
2. Creates `pong`, an instance of the `Pong` actor. (The `pong` object is actually an instance of `ActorRef`, though I loosely refer to it as an actor, or actor instance.) The `Pong` actor constructor does not require any arguments, so the noargs `Props` syntax is used.
3. Creates `ping`, an instance of the `Ping` actor. The `Ping` actor constructor takes one argument, an `ActorRef`, so a slightly different version of the `Props` syntax is used.
4. Starts the `ping/pong` action by sending a `StartMessage` to the `ping` actor.

Once `ping` receives the `StartMessage`, the actors send messages back and forth between each other as fast as they can until the counter limit in `ping` is reached. Messages are sent using the usual `!` method.

To get things started, the `Ping` class needs an initial reference to the `Pong` actor, but once the action starts, the two actors just send a `PingMessage` and `PongMessage` to each other using the sender references they implicitly receive, until the `Ping` actor count limit is reached. At that time, it sends a `StopMessage` to the `Pong` actor, and then both actors call their `context.stop` methods. The context object is implicitly available to all actors, and can be used to stop actors, among other uses.

In addition to demonstrating how to communicate between actors using immutable messages, this example provides several examples of an `ActorRef`. The `ping` and `pong` instances are `ActorRef` instances, as is the `sender` variable.

A great thing about an `ActorRef` is that it hides the actor instance from you. For instance, the `Pong` actor can't directly execute `ping.incrementAndPrint`; the two actors can only send messages between each other. Although this seems limiting at first, once you understand the model, you'll see that it's a terrific way to *safely* implement concurrency in your applications.



Messages can also be sent between actors using the `?` or `ask` methods, but those should be used only rarely. See [Recipe 13.10, “Sending a Message to an Actor and Waiting for a Reply”](#) for examples of those methods.

13.4. Understanding the Methods in the Akka Actor Lifecycle

Problem

You're creating more complicated actors, and need to understand when the methods on an `Actor` are called.

Solution

In addition to its constructor, an `Actor` has the following lifecycle methods:

- `receive`
- `preStart`
- `postStop`
- `preRestart`
- `postRestart`

To demonstrate when these methods are called, basic implementations of these methods have been created in the Kenny actor of the following example:

```
import akka.actor._

class Kenny extends Actor {
  println("entered the Kenny constructor")
  override def preStart { println("kenny: preStart") }
  override def postStop { println("kenny: postStop") }
  override def preRestart(reason: Throwable, message: Option[Any]) {
    println("kenny: preRestart")
    println(s" MESSAGE: ${message.getOrElse("")}")
    println(s" REASON: ${reason.getMessage}")
    super.preRestart(reason, message)
  }
  override def postRestart(reason: Throwable) {
    println("kenny: postRestart")
    println(s" REASON: ${reason.getMessage}")
    super.postRestart(reason)
  }
  def receive = {
    case ForceRestart => throw new Exception("Boom!")
    case _ => println("Kenny received a message")
  }
}

case object ForceRestart

object LifecycleDemo extends App {
  val system = ActorSystem("LifecycleDemo")
  val kenny = system.actorOf(Props[Kenny], name = "Kenny")

  println("sending kenny a simple String message")
  kenny ! "hello"
  Thread.sleep(1000)

  println("make kenny restart")
  kenny ! ForceRestart
  Thread.sleep(1000)

  println("stopping kenny")
  system.stop(kenny)

  println("shutting down system")
  system.shutdown
}
```

The output from this program shows when the lifecycle methods are invoked:

```
[info] Running LifecycleDemo
sending kenny a simple String message
entered the Kenny constructor
kenny: preStart
```

```

Kenny received a message
make kenny restart
[ERROR] [05/14/2013 10:21:54.953] [LifecycleDemo-akka.actor.default-dispatcher-4]
[akka://LifecycleDemo/user/Kenny] Boom!
java.lang.Exception: Boom!
    at Kenny$$anonfun$receive$1.applyOrElse(Test.scala:19)
    (many more lines of exception output ...)

kenny: preRestart
    MESSAGE: ForceRestart
    REASON: Boom!
kenny: postStop
entered the Kenny constructor
kenny: postRestart
    REASON: Boom!
kenny: preStart
stopping kenny
shutting down system
kenny: postStop
[success]

```

Discussion

As shown in the `println` statement at the beginning of the Kenny actor, the body of an Akka Actor is a part of the constructor, just like any regular Scala class. Along with an actor's constructor, the `pre*` and `post*` methods can be used to initialize and close resources that your actor requires.

Notice that `preRestart` and `postRestart` call the super versions of their methods. This is because the default implementation of `postRestart` calls `preRestart`, and I want that default behavior in this application.

Table 13-1 provides a description of each lifecycle method, including an actor's constructor.

Table 13-1. Akka actor lifecycle methods

Method	Description
The actor's constructor	An actor's constructor is called just like any other Scala class constructor, when an instance of the class is first created.
<code>preStart</code>	Called right after the actor is started. During restarts it's called by the default implementation of <code>postRestart</code> .
<code>postStop</code>	Called after an actor is stopped, it can be used to perform any needed cleanup work. According to the Akka documentation, this hook "is guaranteed to run after message queuing has been disabled for this actor."
<code>preRestart</code>	According to the Akka documentation, when an actor is restarted, the old actor is informed of the process when <code>preRestart</code> is called with the exception that caused the restart, and the message that triggered the exception. The message may be <code>None</code> if the restart was not caused by processing a message.

Method	Description
<code>postRestart</code>	The <code>postRestart</code> method of the new actor is invoked with the exception that caused the restart. In the default implementation, the <code>preStart</code> method is called.

See Also

The [Akka actors documentation](#)

13.5. Starting an Actor

Problem

You want to start an Akka actor, or attempt to control the start of an actor.

Solution

This is a bit of a tricky problem, because Akka actors are started asynchronously when they're passed into the `actorOf` method using a `Props`. At the `ActorSystem` level of your application, you create actors by calling the `system.actorOf` method. Within an actor, you create a child actor by calling the `context.actorOf` method.

As demonstrated in [Recipe 13.1](#), you can create an actor at the `ActorSystem` level by passing your actor class name (such as `HelloActor`) to the `system.actorOf` method, using the `Props` case class:

```
val system = ActorSystem("HelloSystem")

// the actor is created and started here
val helloActor = system.actorOf(Props[HelloActor], name = "helloactor")

helloActor ! "hello"
```

The process of creating a child actor from within another actor is almost identical. The only difference is that you call the `actorOf` method on the `context` object instead of on an `ActorSystem` instance. The `context` object is implicitly available to your actor instance:

```
class Parent extends Actor {
  val child = context.actorOf(Props[Child], name = "Child")
  // more code here ...
}
```

Discussion

The following complete example demonstrates how to create actors both at the system level and from within another actor:

```

package actortests.parentchild

import akka.actor._

case class CreateChild (name: String)
case class Name (name: String)

class Child extends Actor {
  var name = "No name"
  override def postStop {
    println(s"D'oh! They killed me ($name): ${self.path}")
  }
  def receive = {
    case Name(name) => this.name = name
    case _ => println(s"Child $name got message")
  }
}

class Parent extends Actor {
  def receive = {
    case CreateChild(name) =>
      // Parent creates a new Child here
      println(s"Parent about to create Child ($name) ...")
      val child = context.actorOf(Props[Child], name = s"$name")
      child ! Name(name)
    case _ => println(s"Parent got some other message.")
  }
}

object ParentChildDemo extends App {

  val actorSystem = ActorSystem("ParentChildTest")
  val parent = actorSystem.actorOf(Props[Parent], name = "Parent")

  // send messages to Parent to create to child actors
  parent ! CreateChild("Jonathan")
  parent ! CreateChild("Jordan")
  Thread.sleep(500)

  // lookup Jonathan, then kill it
  println("Sending Jonathan a PoisonPill ...")
  val jonathan = actorSystem.actorSelection("/user/Parent/Jonathan")
  jonathan ! PoisonPill
  println("jonathan was killed")

  Thread.sleep(5000)
  actorSystem.shutdown
}

```

Here's a brief description of that code:

- At the beginning of the code, the `CreateChild` and `Name` case classes are created. They'll be used to send messages to the actors.
- The `Child` actor has a `receive` method that can handle a `Name` message. It uses that message to set its `name` field.
- The `receive` method of the `Parent` actor can handle a `CreateChild` message. When it receives that message, it creates a new `Child` actor with the given name. Notice that it calls `context.actorOf` to do this.
- The `ParentChildDemo` object creates a new `ActorSystem`, and then creates the `Parent` actor using the `ActorSystem` reference. It then sends two `CreateChild` messages to the parent actor reference. After a brief pause, it looks up the `Child` actor named `Jonathan`, and then sends it a `PoisonPill` message. After another pause, it shuts down the system using the `ActorSystem` reference.

Although it isn't required, in this case, the child actor instance is created in the constructor of the `Parent` actor. The `Child` actor could have been created when the `Parent` actor received a message, so in a sense, that gives you a way to control when an actor instance is created.

13.6. Stopping Actors

Problem

You want to stop one or more running Akka actors.

Solution

There are several ways to stop Akka actors. The most common ways are to call `system.stop(actorRef)` at the `ActorSystem` level or `context.stop(actorRef)` from inside an actor.

There are other ways to stop an actor:

- Send the actor a `PoisonPill` message.
- Program a `gracefulStop`.

To demonstrate these alternatives, at the `ActorSystem` level you can stop an actor by using the `ActorSystem` instance:

```
actorSystem.stop(anActor)
```

Within an actor, you can stop a child actor by using the context reference:

```
context.stop(childActor)
```

An actor can also stop itself:

```
context.stop(self)
```

You can stop an actor by sending it a `PoisonPill` message:

```
actor ! PoisonPill
```

The `gracefulStop` is a little more complicated and involves the use of a future. See the [Discussion](#) for a complete example.

Discussion

Table 13-2 provides a summary of the methods that you can use to stop an actor.

Table 13-2. Ways to stop actors

Message	Description
stop method	The actor will continue to process its current message (if any), but no additional messages will be processed. See additional notes in the paragraphs that follow.
PoisonPill message	A <code>PoisonPill</code> message will stop an actor when the message is processed. A <code>PoisonPill</code> message is queued just like an ordinary message and will be handled after other messages queued ahead of it in its mailbox.
gracefulStop method	Lets you attempt to terminate actors gracefully, waiting for them to timeout. The documentation states that this is a good way to terminate actors in a specific order.

As noted in **Table 13-2**, a major difference between calling the `stop` method on an actor and sending it a `PoisonPill` message is in how the actor is stopped. The `stop` method lets the actor finish processing *the current message* in its mailbox (if any), and then stops it. The `PoisonPill` message lets the actors process *all* messages that are in the mailbox ahead of it before stopping it.

Calling `actorSystem.stop(actor)` and `context.stop(actor)` are the most common ways to stop an actor. The following notes on this process are from the [official Akka actor documentation](#):

- Termination of an actor is performed asynchronously; the `stop` method may return before the actor is actually stopped.
- The actor will continue to process its current message, but no additional messages will be processed.
- An actor terminates in two steps. First, it suspends its mailbox and sends a `stop` message to all of its children. Then it processes termination messages from its children until they're all gone, at which point it terminates itself. If one of the actors doesn't respond (because it's blocking, for instance), the process has to wait for that actor and may get stuck.

- When additional messages aren't processed, they're sent to the `deadLetters` actor of the `ActorSystem` (though this can vary depending on the mailbox implementation). You can access these with the `deadLetters` method on an `ActorSystem`.
- As shown in the following examples, the `postStop` lifecycle method is invoked when an actor is fully stopped, which lets you clean up resources, as needed.

The following subsections demonstrate examples of each of these approaches.

system.stop and context.stop

This is a complete example that shows how to stop an actor by using the `stop` method of an `ActorSystem`:

```
package actortests

import akka.actor._

class TestActor extends Actor {
  def receive = {
    case _ => println("a message was received")
  }
}

object SystemStopExample extends App {
  val actorSystem = ActorSystem("SystemStopExample")
  val actor = actorSystem.actorOf(Props[TestActor], name = "test")
  actor ! "hello"

  // stop our actor
  actorSystem.stop(actor)
  actorSystem.shutdown
}
```

As mentioned, using `context.stop(actorRef)` is similar to using `actorSystem.stop(actorRef)`; just use `context.stop(actorRef)` from within an actor. The `context` variable is implicitly available inside an `Actor`. This is demonstrated in [Recipe 13.5, “Starting an Actor”](#).

PoisonPill message

You can also stop an actor by sending it a `PoisonPill` message. This message will stop the actor when the message is processed. The message is queued in the mailbox like an ordinary message.

Here is a `PoisonPill` example:

```
package actortests

import akka.actor._

class TestActor extends Actor {
```

```

def receive = {
  case s:String => println("Message Received: " + s)
  case _ => println("TestActor got an unknown message")
}
override def postStop { println("TestActor::postStop called") }
}

object PoisonPillTest extends App {
  val system = ActorSystem("PoisonPillTest")
  val actor = system.actorOf(Props[TestActor], name = "test")

  // a simple message
  actor ! "before PoisonPill"

  // the PoisonPill
  actor ! PoisonPill

  // these messages will not be processed
  actor ! "after PoisonPill"
  actor ! "hello?!"

  system.shutdown
}

```

As shown in the comments, the second String message sent to the actor won't be received or processed by the actor because it will be in the mailbox after the PoisonPill. The only output from running this program will be:

```

Message Received: before PoisonPill
TestActor::postStop called

```

gracefulStop

As its name implies, you can use the gracefulStop approach if you want to wait for a period of time for the termination process to complete gracefully. The following code shows a complete example of the gracefulStop approach:

```

package actortests.gracefulstop

import akka.actor._
import akka.pattern.gracefulStop
import scala.concurrent.{Await, ExecutionContext, Future}
import scala.concurrent.duration._
import scala.language.postfixOps

class TestActor extends Actor {
  def receive = {
    case _ => println("TestActor got message")
  }
  override def postStop { println("TestActor: postStop") }
}

```

```

object GracefulStopTest extends App {
  val system = ActorSystem("GracefulStopTest")
  val testActor = system.actorOf(Props[TestActor], name = "TestActor")

  // try to stop the actor gracefully
  try {
    val stopped: Future[Boolean] = gracefulStop(testActor, 2 seconds)(system)
    Await.result(stopped, 3 seconds)
    println("testActor was stopped")
  } catch {
    case e:Exception => e.printStackTrace
  } finally {
    system.shutdown
  }
}

```

Per the Scaladoc, `gracefulStop(actorRef, timeout)` “Returns a Future that will be completed with success when existing messages of the target actor has [sic] been processed and the actor has been terminated.” If the actor isn’t terminated within the timeout, the Future results in an `ActorTimeoutException`. To keep this example simple, I use `Await.result`, so the time period it waits for should be just slightly longer than the timeout value given to `gracefulStop`.

If the order in which actors are terminated is important, using `gracefulStop` can be a good way to attempt to terminate them in a desired order. The “[Akka 2 Terminator](#)” example referenced in the See Also section demonstrates a nice technique for killing child actors in a specific order using `gracefulStop` and `flatMap`.

“Killing” an actor

As you dig deeper into Akka actors, you’ll get into a concept called “supervisor strategies.” When you implement a supervisor strategy, you can send an actor a `Kill` message, which can actually be used to restart the actor. The Akka documentation states that sending a `Kill` message to an actor, “will restart the actor through regular supervisor semantics.”

With the default supervisory strategy, the `Kill` message does what its name states, terminating the target actor. The following example shows the semantics for sending a `Kill` message to an actor:

```

package actortests

import akka.actor._

class Number5 extends Actor {
  def receive = {
    case _ => println("Number5 got a message")
  }
  override def preStart { println("Number5 is alive") }
}

```

```

override def postStop { println("Number5::postStop called") }
override def preRestart(reason: Throwable, message: Option[Any]) {
  println("Number5::preRestart called")
}
override def postRestart(reason: Throwable) {
  println("Number5::postRestart called")
}
}

object KillTest extends App {
  val system = ActorSystem("KillTestSystem")
  val number5 = system.actorOf(Props[Number5], name = "Number5")
  number5 ! "hello"
  // send the Kill message
  number5 ! Kill
  system.shutdown
}

```

Running this code results in the following output:

```

Number5 is alive
Number5 got a message
[ERROR] [16:57:02.220] [KillTestSystem-akka.actor.default-dispatcher-2]
[akka://KillTestSystem/user/Number5] Kill (akka.actor.ActorKilledException)
Number5::postStop called

```

This code demonstrates the `Kill` message so you can see an example of it. In general, this approach is used to kill an actor to allow its supervisor to restart it. If you want to stop an actor, use one of the other approaches described in this recipe.

See Also

- The “[Akka 2 Terminator](#)” example.
- This [Google Groups thread](#) discusses how a `Kill` message is turned into an exception that is handled in the default supervision strategy so it doesn’t restart the actor.
- The [Akka actors documentation](#) provides more examples of these approaches.
- The `gracefulStop` method is described on [this Scaladoc page](#).

13.7. Shutting Down the Akka Actor System

Problem

You want to shut down the Akka actor system, typically because your application is finished, and you want to shut it down gracefully.

Solution

Call the shutdown method on your ActorSystem instance:

```
object Main extends App {  
  // create the ActorSystem  
  val system = ActorSystem("HelloSystem")  
  
  // put your actors to work here ...  
  
  // shut down the ActorSystem when the work is finished  
  system.shutdown  
}
```

Discussion

When you're finished using actors in your application, you should call the shutdown method on your ActorSystem instance. As shown in the examples in this chapter, if you comment out the system.shutdown call, your application will continue to run indefinitely.

In my **SARAH application**, which is a Swing application, I call actorSystem.shutdown when the user shuts down the GUI.

If you want to stop your actors before shutting down the actor system, such as to let them complete their current work, see the examples in [Recipe 13.6, "Stopping Actors"](#).

13.8. Monitoring the Death of an Actor with watch

Problem

You want an actor to be notified when another actor dies.

Solution

Use the watch method of an actor's context object to declare that the actor should be notified when an actor it's monitoring is stopped.

In the following code snippet, the Parent actor creates an actor instance named kenny, and then declares that it wants to "watch" kenny:

```
class Parent extends Actor {  
  val kenny = context.actorOf(Props[Kenny], name = "Kenny")  
  context.watch(kenny)  
  // more code here ...  
}
```

(Technically, kenny is an ActorRef instance, but it's simpler to say "actor.")

If kenny is killed or stopped, the Parent actor is sent a `Terminated(kenny)` message. This complete example demonstrates the approach:

```
package actortests.deathwatch

import akka.actor._

class Kenny extends Actor {
  def receive = {
    case _ => println("Kenny received a message")
  }
}

class Parent extends Actor {
  // start Kenny as a child, then keep an eye on it
  val kenny = context.actorOf(Props[Kenny], name = "Kenny")
  context.watch(kenny)

  def receive = {
    case Terminated(kenny) => println("OMG, they killed Kenny")
    case _ => println("Parent received a message")
  }
}

object DeathWatchTest extends App {

  // create the ActorSystem instance
  val system = ActorSystem("DeathWatchTest")

  // create the Parent that will create Kenny
  val parent = system.actorOf(Props[Parent], name = "Parent")

  // lookup kenny, then kill it
  val kenny = system.actorSelection("/user/Parent/Kenny")
  kenny ! PoisonPill

  Thread.sleep(5000)
  println("calling system.shutdown")
  system.shutdown
}
```

When this code is run, the following output is printed:

```
OMG, they killed Kenny
calling system.shutdown
```

Discussion

Using the `watch` method lets an actor be notified when another actor is stopped (such as with the `PoisonPill` message), or if it's killed with a `Kill` message or `gracefulStop`. This can let the watching actor handle the situation, as desired.

An important thing to understand is that if the Kenny actor throws an exception, this doesn't kill it. Instead it will be restarted. You can confirm this by changing the Kenny actor code to this:

```
case object Explode

class Kenny extends Actor {
  def receive = {
    case Explode => throw new Exception("Boom!")
    case _ => println("Kenny received a message")
  }
  override def preStart { println("kenny: preStart") }
  override def postStop { println("kenny: postStop") }
  override def preRestart(reason: Throwable, message: Option[Any]) {
    println("kenny: preRestart")
    super.preRestart(reason, message)
  }
  override def postRestart(reason: Throwable) {
    println("kenny: postRestart")
    super.postRestart(reason)
  }
}
```

Also, change this line of code in the DeathWatchTest object:

```
kenny ! PoisonPill
```

to this:

```
kenny ! Explode
```

When you run this code, in addition to the error messages that are printed because of the exception, you'll also see this output:

```
kenny: preRestart
kenny: postStop
kenny: postRestart
kenny: preStart
calling system.shutdown
kenny: postStop
```

What you won't see is the "OMG, they killed Kenny" message from the Parent actor, because the exception didn't kill kenny, it just forced kenny to be automatically restarted. You can verify that kenny is restarted after it receives the explode message by sending it another message:

```
kenny ! "Hello?"
```

It will respond by printing the "Kenny received a message" string in the default _ case of its receive method.

Looking up actors

This example also showed one way to look up an actor:

```
val kenny = system.actorSelection("/user/Parent/Kenny")
```

As shown, you look up actors with the `actorSelection` method, and can specify a full path to the actor in the manner shown. The `actorSelection` method is available on an `ActorSystem` instance and on the context object in an `Actor` instance.

You can also look up actors using a relative path. If `kenny` had a sibling actor, it could have looked up `kenny` using its own context, like this:

```
// in a sibling actor
val kenny = context.actorSelection("../Kenny")
```

You can also use various implementations of the `actorFor` method to look up actors. The `kenny` instance could be looked up from the `DeathWatchTest` object in these ways:

```
val kenny = system.actorFor("akka://DeathWatchTest/user/Parent/Kenny")
val kenny = system.actorFor(Seq("user", "Parent", "Kenny"))
```

It could also be looked up from a sibling like this:

```
val kenny = system.actorFor(Seq("..", "Kenny"))
```

13.9. Simple Concurrency with Futures

Problem

You want a simple way to run one or more tasks concurrently, including a way to handle their results when the tasks finish. For instance, you may want to make several web service calls in parallel, and then work with their results after they all return.

Solution

A *future* gives you a simple way to run an algorithm concurrently. A future starts running concurrently when you create it and returns a result at some point, well, in the future. In Scala, it's said that a future returns *eventually*.

The following examples show a variety of ways to create futures and work with their eventual results.

Run one task, but block

This first example shows how to create a future and then block to wait for its result. Blocking is not a good thing—you should block only if you really have to—but this is useful as a first example, in part, because it's a little easier to reason about, and it also gets the bad stuff out of the way early.

The following code performs the calculation $1 + 1$ at some time in the future. When it's finished with the calculation, it returns its result:

```

package actors

// 1 - the imports
import scala.concurrent.{Await, Future}
import scala.concurrent.duration._
import scala.concurrent.ExecutionContext.Implicits.global

object Futures1 extends App {

  // used by 'time' method
  implicit val baseTime = System.currentTimeMillis

  // 2 - create a Future
  val f = Future {
    sleep(500)
    1 + 1
  }

  // 3 - this is blocking (blocking is bad)
  val result = Await.result(f, 1 second)
  println(result)

  sleep(1000)
}

```

Here's how this code works:

- The import statements bring the code into scope that's needed.
- The `ExecutionContext.Implicits.global` import statement imports the “default global execution context.” You can think of an *execution context* as being a thread pool, and this is a simple way to get access to a thread pool.
- A `Future` is created after the second comment. Creating a future is simple; you just pass it a block of code you want to run. This is the code that will be executed at some point in the future.
- The `Await.result` method call declares that it will wait for up to one second for the `Future` to return. If the `Future` doesn't return within that time, it throws a `java.util.concurrent.TimeoutException`.
- The `sleep` statement at the end of the code is used so the program will keep running while the `Future` is off being calculated. You won't need this in real-world programs, but in small example programs like this, you have to keep the JVM running.

I created the `sleep` method in my package object while creating my future and concurrency examples, and it just calls `Thread.sleep`, like this:

```
def sleep(time: Long) { Thread.sleep(time) }
```

As mentioned, blocking is bad; you shouldn't write code like this unless you have to. The following examples show better approaches.

The code also shows a time duration of 1 second. This is made available by the `scala.concurrent.duration._` import. With this library, you can state time durations in several convenient ways, such as 100 nanos, 500 millis, 5 seconds, 1 minute, 1 hour, and 3 days. You can also create a duration as `Duration(100, MILLISECONDS)`, `Duration(200, "millis")`.

Run one thing, but don't block—use callback

A better approach to working with a future is to use its callback methods. There are three callback methods: `onComplete`, `onSuccess`, and `onFailure`. The following example demonstrates `onComplete`:

```
import scala.concurrent.{Future}
import scala.concurrent.ExecutionContext.Implicits.global
import scala.util.{Failure, Success}
import scala.util.Random

object Example1 extends App {

  println("starting calculation ...")
  val f = Future {
    sleep(Random.nextInt(500))
    42
  }

  println("before onComplete")
  f.onComplete {
    case Success(value) => println(s"Got the callback, meaning = $value")
    case Failure(e) => e.printStackTrace
  }

  // do the rest of your work
  println("A ..."); sleep(100)
  println("B ..."); sleep(100)
  println("C ..."); sleep(100)
  println("D ..."); sleep(100)
  println("E ..."); sleep(100)
  println("F ..."); sleep(100)

  sleep(2000)
}
```

This example is similar to the previous example, though it just returns the number 42 after a random delay. The important part of this example is the `f.onComplete` method call and the code that follows it. Here's how that code works:

- The `f.onComplete` method call sets up the callback. Whenever the `Future` completes, it makes a callback to `onComplete`, at which time that code will be executed.
- The `Future` will either return the desired result (42), or an exception.

- The `println` statements with the slight delays represent other work your code can do while the `Future` is off and running.

Because the `Future` is off running concurrently somewhere, and you don't know exactly when the result will be computed, the output from this code is nondeterministic, but it can look like this:

```
starting calculation ...
before onComplete
A ...
B ...
C ...
D ...
E ...
Got the callback, meaning = 42
F ...
```

Because the `Future` returns eventually, at some nondeterministic time, the “Got the callback” message may appear anywhere in that output.

The `onSuccess` and `onFailure` callback methods

There may be times when you don't want to use `onComplete`, and in those situations, you can use the `onSuccess` and `onFailure` callback methods, as shown in this example:

```
import scala.concurrent.{Future}
import scala.concurrent.ExecutionContext.Implicits.global
import scala.util.{Failure, Success}
import scala.util.Random

object OnSuccessAndFailure extends App {

  val f = Future {
    sleep(Random.nextInt(500))
    if (Random.nextInt(500) > 250) throw new Exception("Yikes!") else 42
  }

  f onSuccess {
    case result => println(s"Success: $result")
  }

  f onFailure {
    case t => println(s"Exception: ${t.getMessage}")
  }

  // do the rest of your work
  println("A ..."); sleep(100)
  println("B ..."); sleep(100)
  println("C ..."); sleep(100)
  println("D ..."); sleep(100)
  println("E ..."); sleep(100)
  println("F ..."); sleep(100)
```

```

    sleep(2000)
  }
}

```

This code is similar to the previous example, but this `Future` is wired to throw an exception about half the time, and the `onSuccess` and `onFailure` blocks are defined as partial functions; they only need to handle their expected conditions.

Creating a method to return a `Future[T]`

In the real world, you may have methods that return futures. The following example defines a method named `longRunningComputation` that returns a `Future[Int]`. Declaring it is new, but the rest of this code is similar to the previous `onComplete` example:

```

import scala.concurrent.{Await, Future, future}
import scala.concurrent.ExecutionContext.Implicits.global
import scala.util.{Failure, Success}

object Futures2 extends App {

  implicit val baseTime = System.currentTimeMillis

  def longRunningComputation(i: Int): Future[Int] = future {
    sleep(100)
    i + 1
  }

  // this does not block
  longRunningComputation(11).onComplete {
    case Success(result) => println(s"result = $result")
    case Failure(e) => e.printStackTrace
  }

  // keep the jvm from shutting down
  sleep(1000)
}

```

The future method shown in this example is another way to create a future. It starts the computation asynchronously and returns a `Future[T]` that will hold the result of the computation. This is a common way to define methods that return a future.

Run multiple things; something depends on them; join them together

The examples so far have shown how to run one computation in parallel, to keep things simple. You may occasionally do something like this, such as writing data to a database without blocking the web server, but many times you'll want to run several operations concurrently, wait for them all to complete, and then do something with their combined results.

For example, in a stock market application I wrote, I run all of my web service queries in parallel, wait for their results, and then display a web page. This is faster than running them sequentially.

The following example is a little simpler than that, but it shows how to call an algorithm that may be running in the cloud. It makes three calls to `Cloud.runAlgorithm`, which is defined elsewhere to return a `Future[Int]`. For the moment, this algorithm isn't important, other than to know that it prints its result right before returning it.

The code starts those three futures running, then joins them back together in the for comprehension:

```
import scala.concurrent.{Future, future}
import scala.concurrent.ExecutionContext.Implicits.global
import scala.util.{Failure, Success}
import scala.util.Random

object RunningMultipleCalcs extends App {

  println("starting futures")
  val result1 = Cloud.runAlgorithm(10)
  val result2 = Cloud.runAlgorithm(20)
  val result3 = Cloud.runAlgorithm(30)

  println("before for-comprehension")
  val result = for {
    r1 <- result1
    r2 <- result2
    r3 <- result3
  } yield (r1 + r2 + r3)

  println("before onSuccess")
  result onSuccess {
    case result => println(s"total = $result")
  }

  println("before sleep at the end")
  sleep(2000) // keep the jvm alive

}
```

Here's a brief description of how this code works:

- The three calls to `Cloud.runAlgorithm` create the `result1`, `result2`, and `result3` variables, which are of type `Future[Int]`.
- When those lines are executed, those futures begin running, just like the web service calls in my stock market application.
- The for comprehension is used as a way to join the results back together. When all three futures return, their `Int` values are assigned to the variables `r1`, `r2`, and `r3`,

and the sum of those three values is returned from the `yield` expression, and assigned to the `result` variable.

- Notice that `result` can't just be printed after the `for` comprehension. That's because the `for` comprehension returns a new future, so `result` has the type `Future[Int]`. (This makes sense in more complicated examples.) Therefore, the correct way to print the example is with the `onSuccess` method call, as shown.

When this code is run, the output is nondeterministic, but looks something like this:

```
starting futures
before for-comprehension
before onSuccess
before sleep at end
returning result from cloud: 30
returning result from cloud: 20
returning result from cloud: 40
total = 90
```

Notice how all of the `println` statements in the code print before the `total` is printed. That's because they're running in sequential fashion, while the future is off and running in parallel, and returns at some indeterminate time ("eventually").

I mentioned earlier that the `Cloud.runAlgorithm` code wasn't important—it was just something running "in the cloud,"—but for the sake of completeness, here's that code:

```
object Cloud {

  def runAlgorithm(i: Int): Future[Int] = future {
    sleep(Random.nextInt(500))
    val result = i + 10
    println(s"returning result from cloud: $result")
    result
  }

}
```

In my real-world code, I use a future in a similar way to get information from web services. For example, in a Twitter client, I make multiple calls to the Twitter web service API using futures:

```
// get the desired info from twitter
val dailyTrendsFuture = Future { getDailyTrends(twitter) }
val usFuture = Future { getLocationTrends(twitter, woeidUnitedStates) }
val worldFuture = Future { getLocationTrends(twitter, woeidWorld) }
```

I then join them in a `for` comprehension, as shown in this example. This is a nice, simple way to turn single-threaded web service calls into multiple threads.

Discussion

Although using a future is straightforward, there are also many concepts behind it. The following sections summarize the most important concepts.

A future and ExecutionContext

The following statements describe the basic concepts of a future, as well as the `ExecutionContext` that a future relies on.

- A `Future[T]` is a container that runs a computation concurrently, and at some future time may return either (a) a result of type `T` or (b) an exception.
- Computation of your algorithm starts at some nondeterministic time after the future is created, running on a thread assigned to it by the execution context.
- The result of the computation becomes available once the future completes.
- When it returns a result, a future is said to be *completed*. It may either be *successfully completed*, or *failed*.
- As shown in the examples, a future provides an interface for reading the value that has been computed. This includes callback methods and other approaches, such as `for` comprehension, `map`, `flatMap`, etc.
- An `ExecutionContext` executes a task it's given. You can think of it as being like a thread pool.
- The `ExecutionContext.Implicits.global` import statement shown in the examples imports the default global execution context.

Callback methods

The following statements describe the use of the callback methods that can be used with futures.

- Callback methods are called asynchronously when a future completes.
- The callback methods `onComplete`, `onSuccess`, `onFailure`, are demonstrated in the Solution.
- A callback method is executed by some thread, some time after the future is completed. From the Scala [Futures documentation](#), “There is no guarantee that it will be called by the thread that completed the future or the thread that created the callback.”
- The order in which callbacks are executed is not guaranteed.
- `onComplete` takes a callback function of type `Try[T] => U`.
- `onSuccess` and `onFailure` take partial functions. You only need to handle the desired case. (See [Recipe 9.8, “Creating Partial Functions”](#) for more information on partial functions.)

- `onComplete`, `onSuccess`, and `onFailure` have the result type `Unit`, so they can't be chained. This design was intentional, to avoid any suggestion that callbacks may be executed in a particular order.

For comprehensions (combinators: `map`, `flatMap`, `filter`, `foreach`, `recoverWith`, `fallbackTo`, and `andThen`)

As shown in the Solution, callback methods are good for some purposes. But when you need to run multiple computations in parallel, and join their results together when they're finished running, using *combinators* like `map`, `foreach`, and other approaches, like a `for` comprehension, provides more concise and readable code. The `for` comprehension was shown in the Solution.

The `recover`, `recoverWith`, and `fallbackTo` combinators provide ways of handling failure with futures. If the future they're applied to returns successfully, you get that (desired) result, but if it fails, these methods do what their names suggest, giving you a way to recover from the failure.

As a short example, you can use the `fallbackTo` method like this:

```
val meaning = calculateMeaningOfLife() fallbackTo 42
```

The `andThen` combinator gives you a nice syntax for running whatever code you want to run when a future returns, like this:

```
var meaning = 0
future {
  meaning = calculateMeaningOfLife()
} andThen {
  println(s"meaning of life is $meaning")
}
```

See the Scala [Futures documentation](#) for more information on their use.

See Also

- The Scala [Futures documentation](#)
- These examples (and more) are available at my [GitHub repository](#).
- As shown in these examples, you can read a result from a future, and a *promise* is a way for some part of your software to put that result in there. I've linked to the [best article I can find](#).

13.10. Sending a Message to an Actor and Waiting for a Reply

Problem

You have one actor that needs to ask another actor for some information, and needs an immediate reply. (The first actor can't continue without the information from the second actor.)

Solution

Use the `? ask` methods to send a message to an Akka actor and wait for a reply, as demonstrated in the following example:

```
import akka.actor._
import akka.pattern.ask
import akka.util.Timeout
import scala.concurrent.{Await, ExecutionContext, Future}
import scala.concurrent.duration._
import scala.language.postfixOps

case object AskNameMessage

class TestActor extends Actor {
  def receive = {
    case AskNameMessage => // respond to the 'ask' request
                          sender ! "Fred"
    case _ => println("that was unexpected")
  }
}

object AskTest extends App {

  // create the system and actor
  val system = ActorSystem("AskTestSystem")
  val myActor = system.actorOf(Props[TestActor], name = "myActor")

  // (1) this is one way to "ask" another actor for information
  implicit val timeout = Timeout(5 seconds)
  val future = myActor ? AskNameMessage
  val result = Await.result(future, timeout.duration).asInstanceOf[String]
  println(result)
```

```

// (2) a slightly different way to ask another actor for information
val future2: Future[String] = ask(myActor, AskNameMessage).mapTo[String]
val result2 = Await.result(future2, 1 second)
println(result2)

system.shutdown
}

```

Discussion

Both the `?` or `ask` methods use the `Future` and `Await.result` approach demonstrated in [Recipe 13.9, “Simple Concurrency with Futures”](#). The recipe is:

1. Send a message to an actor using either `?` or `ask` instead of the usual `!` method.
2. The `?` and `ask` methods create a `Future`, so you use `Await.result` to wait for the response from the other actor.
3. The actor that’s called should send a reply back using the `!` method, as shown in the example, where the `TestActor` receives the `AskNameMessage` and returns an answer using `sender ! "Fred"`.

To keep the previous example simple, only one actor is shown, but the same approach is used by two actors. Just use the `?` or `ask` method in your actor, like this:

```

class FooActor extends Actor {
  def receive = {
    case GetName =>
      val future: Future[String] = ask(otherActor, AskNameMessage).mapTo[
[String]
      val result = Await.result(future, 1 second)
      case _ => // handle other messages
  }
}

```

Be careful when writing code that waits for immediate responses like this. This causes your actor to block, which means that it can’t respond to anything else while it’s in this state. When you need to perform work like this, the mantra is, “Delegate, delegate, delegate.”

13.11. Switching Between Different States with `become`

Problem

You want a simple mechanism to allow an actor to switch between the different states it can be in at different times.

Solution

Use the Akka “become” approach. To do this, first define the different possible states the actor can be in. Then, in the actor’s receive method, switch between the different states based on the messages it receives.

The following example shows how the actor named `DavidBanner` might switch between its `normalState` and its `angryState` (when he becomes The Hulk):

```
package actortests.becometest

import akka.actor._

case object ActNormalMessage
case object TryToFindSolution
case object BadGuysMakeMeAngry

class DavidBanner extends Actor {
  import context._

  def angryState: Receive = {
    case ActNormalMessage =>
      println("Phew, I'm back to being David.")
      become(normalState)
  }

  def normalState: Receive = {
    case TryToFindSolution =>
      println("Looking for solution to my problem ...")
    case BadGuysMakeMeAngry =>
      println("I'm getting angry...")
      become(angryState)
  }

  def receive = {
    case BadGuysMakeMeAngry => become(angryState)
    case ActNormalMessage => become(normalState)
  }
}

object BecomeHulkExample extends App {
  val system = ActorSystem("BecomeHulkExample")
  val davidBanner = system.actorOf(Props[DavidBanner], name = "DavidBanner")
  davidBanner ! ActNormalMessage // init to normalState
  davidBanner ! TryToFindSolution
  davidBanner ! BadGuysMakeMeAngry
  Thread.sleep(1000)
  davidBanner ! ActNormalMessage
  system.shutdown
}
```

Here’s a description of the code:

1. The `davidBanner` actor instance is created, as shown in previous recipes.
2. The `davidBanner` instance is sent the `ActNormalMessage` to set an initial state.
3. After sending `davidBanner` a `TryToFindSolution` message, it sends a `BadGuysMakeMeAngry` message.
4. When `davidBanner` receives the `BadGuysMakeMeAngry` message, it uses `become` to switch to the `angryState` .
5. In the `angryState` the only message `davidBanner` can process is the `ActNormalMessage` . (In the real world, er, entertainment world, it should be programmed to receive other messages, like `SmashThings` .)
6. When `davidBanner` receives the final `ActNormalMessage` , it switches back to the `normalState` , again using the `become` method.

Discussion

As shown, the general recipe for using the `become` approach to switch between different possible states is:

- Define the different possible states, such as the `normalState` and `angryState` .
- Define the `receive` method in the actor to switch to the different states based on the messages it can receive. As shown in the example, this is handled with a `match` expression.

It's important to note that the different states can only receive the messages they're programmed for, and those messages can be different in the different states. For instance, the `normalState` responds to the messages `TryToFindSolution` and `BadGuysMakeMeAngry` , but the `angryState` can only respond to the `ActNormalMessage` .

See Also

The [Akka actors documentation](#) shows a `become` example.

13.12. Using Parallel Collections

Problem

You want to improve the performance of algorithms by using parallel collections.

Solution

When creating a collection, use one of the Scala's parallel collection classes, or convert an existing collection to a parallel collection. In either case, test your algorithm to make sure you see the benefit you're expecting.

You can convert an existing collection to a parallel collection. To demonstrate this, first create a sequential collection, such as a `Vector`:

```
scala> val v = Vector.range(0, 10)
v: scala.collection.immutable.Vector[Int] = Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)
```

Next, print the sequence, and you'll see that it prints as usual:

```
scala> v.foreach(print)
0123456789
```

As expected, that example prints the string `0123456789`. No matter how many times you print it, you'll always see that same result; that's the linear world you're used to.

Next, call the `par` method on your collection to turn it into a parallel collection, and repeat the experiment:

```
scala> v.par.foreach(print)
5678901234

scala> v.par.foreach(print)
0123456789

scala> v.par.foreach{ e => print(e); Thread.sleep(50) }
0516273894
```

Whoa. Sometimes the collection prints in order, other times it prints in a seemingly random order. That's because it's now using an algorithm that runs concurrently. Welcome to the brave, new, parallel world.

That example showed how to convert a “normal” collection to a parallel collection. You can also create a parallel collection directly:

```
scala> import scala.collection.parallel.immutable.ParVector
import scala.collection.parallel.immutable._

scala> val v = ParVector.range(0, 10)
v: scala.collection.parallel.immutable.ParVector[Int] =
  ParVector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)

scala> v.foreach{ e => Thread.sleep(100); print(e) }
0516273849
```

Discussion

As shown, you can create parallel collections in two ways:

- Convert a “normal” collection to its parallel counterpart
- Instantiate them directly, just like their nonparallel counterparts

You can create a new instance of a parallel collection directly. As with the “normal” collection classes that are discussed in [Chapter 10](#) and [Chapter 11](#), there are both immutable and mutable parallel collections. Here’s a list of some of the immutable parallel collection classes:

ParHashMap	ParHashSet	ParIterable	ParMap
ParRange	ParSeq	ParSet	ParVector

In addition to these, the mutable collection has other classes and traits, including `ParArray`.



For a full list of Scala’s parallel collections, see the [Scala website](#).

Where are parallel collections useful?

To understand where a parallel collection can be useful, it helps to think about how they work. Conceptually, you can imagine a collection being split into different chunks; your algorithm is then applied to the chunks, and at the end of the operation, the different chunks are recombined.

For instance, in the Solution, a `ParVector` was created like this:

```
scala> val v = ParVector.range(0, 10)
v: scala.collection.parallel.immutable.ParVector[Int] =
  ParVector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9)
```

The elements in the `ParVector` were then printed like this:

```
scala> v.foreach{ e => Thread.sleep(100); print(e) }
0516273849
```

This makes sense if you imagine that the original `ParVector` is split into two sequences before the printing operation begins:

```
(0,1,2,3,4)
(5,6,7,8,9)
```

In this case you can imagine the `foreach` method taking (or receiving) the 0 from the first sequence, printing it; getting the 5 from the second sequence, printing it; then getting the 1 from the first sequence, etc.

To summarize the basic concept:

- Collection elements are split into different groups.
- The operation is performed.
- The elements are recombined.

The impact of this approach is that it must be okay that your algorithm receives elements in an arbitrary order. This means that algorithms like `sum`, `max`, `min`, `mean`, and `filter` will all work fine.

Conversely, any algorithm that depends on the collection elements being received in a predictable order should not be used with a parallel collection. A simple demonstration of this is the `foreach` examples that have been shown: if it's important that the collection elements are printed in a particular order, such as the order in which they were placed in the collection, using a parallel collection isn't appropriate.

The official Scala documentation refers to this as “side-effecting operations.” The [Parallel Collections Overview URL](#) in the See Also section discusses this in detail.

Performance

Using parallel collections won't always make your code faster. It's important to test your algorithm with and without a parallel collection to make sure your algorithm is faster with a parallel collection. The [“Measuring Performance” URL](#) in the See Also section has a terrific discussion about how to properly benchmark JVM performance.

For a parallel algorithm to provide a benefit, a collection usually needs to be fairly large. The documentation states:

“As a general heuristic, speed-ups tend to be noticeable when the size of the collection is large, typically several thousand elements.”

Finally, if using a parallel collection won't solve your problem, using Akka actors and futures can give you complete control over your algorithms.

See Also


- [Immutable parallel collections](#)
- [Mutable parallel collections](#)
- [Parallel collections overview](#)
- [Measuring the performance of parallel collections](#)

CHAPTER 14

Command-Line Tasks

14.0. Introduction

Scala offers a number of tools to let you work at the command line, including the Read-Eval-Print-Loop, or *REPL*. As shown in [Figure 14-1](#), the REPL lets you execute Scala expressions in an interactive environment.



```
$ scala
Welcome to Scala version 2.10.0 (Java HotSpot(TM) 64-Bit Server VM, Java 1.6.0_37).
Type in expressions to have them evaluated.
Type :help for more information.

scala> val s = "Hello, command line"
s: String = Hello, command line

scala> s.length
res0: Int = 19

scala> "scala".foreach(println)
s
c
a
l
a

scala>
```

Figure 14-1. The REPL lets you execute Scala expressions in an interactive environment

If you’ve used an interactive interpreter before (such as Ruby’s *irb* tool), the Scala REPL will seem very familiar.

When it comes to building your projects, you’ll be well served to use the Simple Build Tool (SBT), so that’s covered in [Chapter 18](#). But there are still times when you’ll want to use `scalac`, `fsc`, `scaladoc`, and other command-line tools, and this chapter demonstrates all of those tools.

The name “Scala” comes from the word “scalable,” and Scala does indeed scale from small shell scripts to the largest, highest-performance applications in the world. On the low end of that scale, this chapter demonstrates how to create your own shell scripts, prompt for input from your scripts, and then make them run faster.

14.1. Getting Started with the Scala REPL

Problem

You want to get started using the Scala REPL, including understanding some of its basic features, such as tab completion, starting the REPL with different options, and dealing with errors.

Solution

To start the Scala REPL, type `scala` at your operating system command line:

```
$ scala
```

You’ll see a welcome message and Scala prompt:

```
Welcome to Scala version 2.10.0
Type in expressions to have them evaluated.
Type :help for more information.

scala> _
```

Welcome, you’re now using the Scala REPL.

Inside the REPL environment, you can try all sorts of different experiments and expressions:

```
scala> val x, y = 1
x: Int = 1
y: Int = 1

scala> x + y
res0: Int = 2

scala> val a = Array(1, 2, 3)
a: Array[Int] = Array(1, 2, 3)
```

```
scala> a.sum
res1: Int = 6
```

As shown in the second example, if you don't assign the result of an expression to a variable, the REPL creates its own variable, beginning with `res0`, then `res1`, etc. You can use these variable names just as though you had created them yourself:

```
scala> res1.getClass
res2: Class[Int] = int
```

Writing tests like this in the REPL is a great way to run experiments outside of your IDE or editor.

There are a few simple tricks that can make using the REPL more effective. One trick is to use *tab completion* to see the methods that are available on an object. To see how tab completion works, create a `String` object, type a decimal, and then press the Tab key. With Scala 2.10, the REPL shows that more than 30 methods are available on a `String` instance:

```
scala> "foo".[Tab]
+               asInstanceOf      charAt      codePointAt
codePointBefore codePointCount    compareTo

// a total of thirty methods listed here ...
```

If you press the Tab key again, the REPL expands the list to more than 50 methods:

```
scala> "foo".[Tab][Tab]
// 51 methods now listed ...
```

Similarly, the `Int` object expands from 25 to 34 methods when you press the Tab key twice.

When you press the Tab key the first time, the REPL filters out many common methods, but by pressing the Tab key the second time, it removes those filters and increases the verbosity of its output. You can find an explanation of how this works at the [JLineCompletion](#) class link in the See Also section of this recipe.

You can also limit the list of methods that are displayed by typing the first part of a method name and then pressing the Tab key. For instance, if you know that you're interested in the `to*` methods on a Scala `List`, type a decimal and the characters `to` after a `List` instance, and then press Tab:

```
scala> List(1,2,3).to[Tab]
toByte toChar toDouble toFloat toInt toLong toShort toString
```

These are all the `List` methods that begin with the letters `to`.



Although the REPL tab-completion feature is good, it currently doesn't show methods that are available to an object that results from implicit conversions. For instance, when you invoke the tab-completion feature on a `String` instance, the REPL doesn't show the methods that are available to the `String` that come from the implicit conversions defined in the `StringOps` class.

To see methods available from the `StringOps` class, you currently have to do something like this:

```
scala> val s = new collection.immutable.StringOps("")
s: scala.collection.immutable.StringOps = s

scala> s.[Tab]
```

After pressing the Tab key, you'll see dozens of additional methods that are available to a `String` object, such as all the `to*` and collection methods.

The REPL also doesn't show method signatures. Hopefully features like this will be added to future versions of the REPL. In the meantime, these are most easily seen in an IDE.

Discussion

I use the REPL to create many small experiments, and it also helps me understand some type conversions that Scala performs automatically. For instance, when I first started working with Scala and typed the following code into the REPL, I didn't know what type the variable `x` was:

```
scala> val x = (3, "Three", 3.0)
x: (Int, java.lang.String, Double) = (3,Three,3.0)
```

With the REPL, it's easy to run tests like this, and then call `getClass` on a variable to see its type:

```
scala> x.getClass
res0: java.lang.Class[_ <: (Int, java.lang.String, Double)] = class scala.Tuple3
```

Although some of that result line is hard to read when you first start working with Scala, the text on the right side of the `=` lets you know that the type is a `Tuple3`.

Though this is a simple example, when you're working with more complicated code or a new library, you'll find yourself running many small tests like this in the REPL.



A `Tuple3` is a specific instance of a *tuple*. A tuple is a container for heterogeneous objects. A `Tuple3` is simply a tuple that contains three elements. Here's a `Tuple2` that holds a `String` and a `Char`:

```
scala> val y = ("Foo", 'a')
y: (java.lang.String, Char) = (Foo,a)

scala> y.getClass
res1: java.lang.Class[_ <: (java.lang.String, Char)]
      = class scala.Tuple2
```

See [Recipe 10.27, “Tuples, for When You Just Need a Bag of Things”](#) for more information.

REPL command-line options

If you need to set Java properties when starting the Scala interpreter, you can do so like this on Unix systems:

```
$ env JAVA_OPTS="-Xmx512M -Xms64M" scala
```

That command sets the maximum and initial size of the Java memory allocation pool. You can confirm this by looking at the maximum available memory in the REPL:

```
scala> Runtime.getRuntime.maxMemory / 1024
res0: Long = 520064
```

When starting the Scala 2.10 REPL without any options, the same command yields a different result:

```
scala> Runtime.getRuntime.maxMemory / 1024
res0: Long = 258880
```

You can also use the `-J` command-line argument to set parameters. I ran into a `java.lang.OutOfMemoryError` in the REPL while processing a large XML dataset, and fixed the problem by starting the REPL with this command:

```
$ scala -J-Xms256m -J-Xmx512m
```

The `scala` command you're running in these examples is actually a shell script, so if you need to modify these parameters permanently, just edit that script. (On Unix systems, you can also create a wrapper script or an alias.)

Deprecation and feature warnings

From time to time, you may see a message that suggests starting the REPL with the `-deprecation` or `-feature` option enabled. For instance, attempting to create an octal value by entering an integer value with a leading zero generates a deprecation warning:

```
scala> 012
warning: there were 1 deprecation warnings; re-run with -deprecation for details
res0: Int = 10
```

To see the error, you *could* restart the REPL with the `-deprecation` option, like this:

```
$ scala -deprecation
```

Fortunately, restarting the REPL isn't usually necessary. Beginning with Scala 2.10, it's usually easier to ask the REPL to show the message with the `:warning` command:

```
scala> 012
warning: there were 1 deprecation warnings; re-run with -deprecation for details
res0: Int = 10

scala> :warning
<console>:8: warning: Treating numbers with a leading zero as octal is deprecated.
    012
    ^
```

The REPL documentation states that the `:warning` command shows “the suppressed warnings from the most recent line.”

If you run into the similar feature warning message, you can also issue the `:warning` command to see the error. If necessary, you can also restart the REPL with the `-feature` option:

```
$ scala -feature
```

The Scala Worksheet

If you're using Eclipse with the Scala IDE plug-in, you can also run a REPL session in a Scala Console panel. Another alternative is to use the [Scala Worksheet](#). The Worksheet is a plug-in that's available for Eclipse and IntelliJ IDEA. It works like the REPL, but runs inside the IDE. [Figure 14-2](#) shows what the Worksheet looks like in Eclipse.




Figure 14-2. The Scala Worksheet plug-in works like the REPL

See Also

- [Source code for the JLineCompletion class](#)
- The `Tuple3` class

14.2. Pasting and Loading Blocks of Code into the REPL

Problem

You want to experiment with some code in the Scala REPL, and typing it in or trying to paste it into the REPL won't work.

Solution

The REPL is “greedy” and consumes the first full statement you type in, so attempting to paste blocks of code into it can fail. To solve the problem, either use the `:paste` command to paste blocks of code into the REPL, or use the `:load` command to load the code from a file into the REPL.

The `:paste` command

Attempting to paste the following `if/else` block into the REPL will cause an error:

```
if (true)
  print("that was true")
else
  print("that was false")
```

But by issuing the `:paste` command before pasting in the code, the code will be interpreted properly:

```
scala> :paste
// Entering paste mode (ctrl-D to finish)

if (true)
  print("that was true")
else
  print("false")

[Ctrl-D]

// Exiting paste mode, now interpreting.
that was true
```

As shown, follow these steps to paste your code into the REPL:

1. Type the `:paste` command in the REPL.
2. Paste in your block of code (Command-V on a Mac, Ctrl-V on Windows).
3. Press Ctrl-D, and the REPL will evaluate what you pasted in.

The `:load` command

Similarly, if you have source code in a file that you want to read into the REPL environment, you can use the `:load` command. For example, assume you have the following source code in a file named *Person.scala* in the same directory where you started the REPL:

```
case class Person(name: String)
```

You can load that source code into the REPL environment like this:

```
scala> :load Person.scala
Loading /Users/AL/ScalaTests/Person.scala...
defined class Person
```

Once the code is loaded into the REPL, you can create a new *Person* instance:

```
scala> val al = Person("Alvin Alexander")
al: Person = Person(Alvin Alexander)
```

Note, however, that if your source code has a package declaration:

```
// Person.scala source code
package com.alvinalexander.foo
case class Person(name: String)
```

the `:load` command will fail:

```
scala> :load /Users/AL/ProjectX/Person.scala
Loading /Users/AL/ProjectX/Person.scala...
<console>:1: error: illegal start of definition
    package com.alvinalexander.foo
    ^
defined class Person
```

You can't use packages in the REPL, so for situations like this, you'll need to compile your file(s) and then include them on the classpath, as shown in [Recipe 14.3, "Adding JAR Files and Classes to the REPL Classpath"](#).

Discussion

Although the REPL is incredibly helpful, its greedy nature can cause multiline statements to fail. Imagine that you want to type the following block of code into the REPL:

```
if (true)
  't'
else
  'f'
```

If you try typing this code in one line at a time, the REPL will cut you off as soon as it sees a complete statement:

```
scala> if (true)
|   't'
res0: AnyVal = t
```

In this simple example, you can get around the problem by adding curly braces to the expression, in which case the REPL recognizes that the expression isn't finished:

```
scala> if (true) {
|   't'
| } else {
|   'f'
| }
res0: Char = t
```

But you can't always do this. In the cases where this fails, use one of the approaches shown in the Solution.

Scala's -i option

Another approach you can use is to load your source code with the `-i` argument when starting the Scala REPL. See [Recipe 14.4, "Running a Shell Command from the REPL"](#) for more information on that approach.

See Also

[Recipe 14.3, "Adding JAR Files and Classes to the REPL Classpath"](#)

14.3. Adding JAR Files and Classes to the REPL Classpath

Problem

You want to add individual classes or one or more JAR files to the REPL classpath so you can use them in a REPL session.

Solution

If you know that you want to use code from a JAR file when you start the REPL session, add the `-cp` or `-classpath` argument to your `scala` command when you start the session. This example shows how to load and use my *DateUtils.jar* library:

```
$ scala -cp DateUtils.jar

scala> import com.alvinalexander.dateutils._
import com.alvinalexander.dateutils._
```

```
scala> DateUtils.getCurrentDate
res0: String = Saturday, March 16
```

If you realize you need a JAR file on your classpath *after* you've started a REPL session, you can add one dynamically with the `:cp` command:

```
scala> :cp DateUtils.jar
Added '/Users/AL/Projects/Scala/Tests/DateUtils.jar'.
Your new classpath is:
"./Users/AL/Projects/Scala/Tests/DateUtils.jar"
```

```
scala> import com.alvinalexander.dateutils._
import com.alvinalexander.dateutils._
```

```
scala> DateUtils.getCurrentDate
res0: String = Saturday, March 16
```

Compiled class files in the current directory (`*.class`) are automatically loaded into the REPL environment, so if a simple `Person.class` file is in the current directory when you start the REPL, you can create a new `Person` instance without requiring a classpath command:

```
scala> val p = new Person("Bill")
p: Person = Person(Bill)
```

However, if your class files are in a subdirectory, you can add them to the environment when you start the session, just as with JAR files. If all the class files are located in a subdirectory named `classes`, you can include them by starting your REPL session like this:

```
$ scala -cp classes
```

If the class files you want to include are in several different directories, you can add them all to your classpath:

```
$ scala -cp "../Project1/bin:../Project2/classes"
```

(This command works on Unix systems, but it may be slightly different on Windows.)

These approaches let you add JAR files and other compiled classes to your REPL environment, either at startup or as the REPL is running.

14.4. Running a Shell Command from the REPL

Problem

You want to be able to run a shell command from within the Scala REPL, such as listing the files in the current directory.

Solution

Run the command using the `:sh` REPL command, then print the output. The following example shows how to run the Unix `ls -al` command from within the REPL, and then show the results of the command:

```
scala> :sh ls -al
res0: scala.tools.nsc.interpreter.ProcessResult = `ls -al` (6 lines, exit 0)

scala> res0.show
total 24
drwxr-xr-x  5 Al  staff  170 Jul 14 17:14 .
drwxr-xr-x 29 Al  staff  986 Jul 14 15:27 ..
-rw-r--r--  1 Al  staff  108 Jul 14 15:34 finance.csv
-rw-r--r--  1 Al  staff  469 Jul 14 15:38 process.scala
-rw-r--r--  1 Al  staff  412 Jul 14 16:24 process2.scala
```

Alternatively you can import the `scala.sys.process` package, and then use the normal `Process` and `ProcessBuilder` commands described in [Recipe 12.10, “Executing External Commands”](#):

```
scala> import sys.process._
import sys.process._

scala> "ls -al" !
total 24
drwxr-xr-x  5 Al  staff  170 Jul 14 17:14 .
drwxr-xr-x 29 Al  staff  986 Jul 14 15:27 ..
-rw-r--r--  1 Al  staff  108 Jul 14 15:34 finance.csv
-rw-r--r--  1 Al  staff  469 Jul 14 15:38 process.scala
-rw-r--r--  1 Al  staff  412 Jul 14 16:24 process2.scala
res0: Int = 0
```

Scala’s -i option

Although those examples show the correct approach, you can improve the situation by loading your own custom code when you start the Scala interpreter. For instance, I always start the REPL in my `/Users/Al/tmp` directory, and I have a file in that directory named *repl-commands* with these contents:

```
import sys.process._

def clear = "clear".!
def cmd(cmd: String) = cmd.!!
def ls(dir: String) { println(cmd(s"ls -al $dir")) }
def help {
  println("\n=== MY CONFIG ===")
  "cat /Users/Al/tmp/repl-commands".!
}

case class Person(name: String)
val nums = List(1, 2, 3)
```

```
val strings = List("sundance", "rocky", "indigo")

// lets me easily see the methods from StringOps
// with tab completion
val so = new collection.immutable.StringOps("")
```

With this setup, I start the Scala interpreter with the `-i` argument, telling it to load this file when it starts:

```
$ scala -i repl-commands
```

This makes those pieces of code available to me inside the REPL. For instance, I can clear my terminal window by invoking the `clear` method:

```
scala> clear
```

My `ls` method provides a directory listing:

```
scala> ls("/tmp")
```

With my `cmd` method I can run other external commands:

```
scala> cmd("cat /etc/passwd")
```

The `help` method uses the system `cat` command to display this file, which is helpful if I haven't used it in a while. The `nums` and `strings` variables and `Person` class also make it easy to run quick experiments.

This approach is similar to using a startup file to initialize a Unix login session, like a `.bash_profile` file for Bash users, and I highly recommend it. As you use the REPL more and more, use this technique to customize its behavior.

To make this even easier, I created the following Unix alias and put it in my `.bash_profile` file:

```
alias repl="scala -i /Users/Al/tmp/repl-commands"
```

I now use this alias to start a REPL session, rather than starting it by typing `scala`:

```
$ repl
```

See Also

The “Executing external commands” recipes in [Chapter 12](#) for more examples of executing external commands from Scala code

14.5. Compiling with scalac and Running with scala

Problem

Though you normally use the Simple Build Tool (SBT) to build Scala applications, you may want to use more basic tools to compile and run small test programs, in the same way you might use `javac` and `java` with small Java applications.

Solution

Compile programs with `scalac`, and run them with `scala`. For example, given a Scala source code file named *Hello.scala*:

```
object Hello extends App {  
  println("Hello, world")  
}
```

Compile it from the command line with `scalac`:

```
$ scalac Hello.scala
```

Then run it with `scala`:

```
$ scala Hello  
Hello, world
```

Discussion

Compiling and executing classes is basically the same as Java, including concepts like the classpath. For instance, if you have a class named `Pizza` in a file named *Pizza.scala*, it may depend on a `Topping` class:

```
class Pizza (var toppings: Topping*) {  
  override def toString = toppings.toString  
}
```

Assuming that the `Topping` class is compiled to a file named *Topping.class* in a subdirectory named *classes*, compile *Pizza.scala* like this:

```
$ scalac -classpath classes Pizza.scala
```

In a more complicated example, you may have your source code in subdirectories under a *src* folder, one or more JAR files in a *lib* directory, and you want to compile your output class files to a *classes* folder. In this case, your files and directories will look like this:

```
./classes  
./lib/DateUtils.jar  
./src/com/alvinalexander/pizza/Main.scala  
./src/com/alvinalexander/pizza/Pizza.scala  
./src/com/alvinalexander/pizza/Topping.scala
```

The *Main.scala*, *Pizza.scala*, and *Topping.scala* files will also have package declarations corresponding to the directories they are located in, i.e.:

```
package com.alvinalexander.pizza
```

Given this configuration, to compile your source code files to the *classes* directory, use the following command:

```
$ scalac -classpath lib/DateUtils.jar -d classes ↵ src/com/alvinalexander/pizza/*
```

Assuming *Main.scala* is an object that extends *App*, *Pizza.scala* is a regular class file, and *Topping.scala* is a case class, your *classes* directory will contain these files after your *scalac* command:

```
./classes/com/alvinalexander/pizza/Main$.class
./classes/com/alvinalexander/pizza/Main$delayedInit$body.class
./classes/com/alvinalexander/pizza/Main.class
./classes/com/alvinalexander/pizza/Pizza.class
./classes/com/alvinalexander/pizza/Topping$.class
./classes/com/alvinalexander/pizza/Topping.class
```

Once the files have been compiled in this manner, you can run the application like this:

```
$ scala -classpath classes:lib/DateUtils.jar com.alvinalexander.pizza.Main
```

As you can imagine, this process gets more and more difficult as you add new classes and libraries, and it's strongly recommended that you use a tool like SBT, Maven, or Ant to manage your application's build process. The examples shown in this recipe are shown for the “one off” cases where you might want to compile and run a small application or test code.

For other useful command-line options, see the manpages for the *scalac* and *scala* commands.

14.6. Disassembling and Decompiling Scala Code

Problem

In the process of learning Scala, or trying to understand a particular problem, you want to examine the bytecode the Scala compiler generates from your source code.

Solution

You can use several different approaches to see how your Scala source code is translated:

- Use the `javap` command to disassemble a `.class` file to look at its signature.
- Use `scalac` options to see how the compiler converts your Scala source code to Java code.
- Use a decompiler to convert your class files back to Java source code.

All three solutions are shown here.

Using `javap`

Because your Scala source code files are compiled into regular Java class files, you can use the `javap` command to disassemble them. For example, assume that you’ve created a file named `Person.scala` that contains the following source code:

```
class Person (var name: String, var age: Int)
```

If you compile that file with `scalac`, you can disassemble the resulting class file into its signature using `javap`, like this:

```
$ javap Person
Compiled from "Person.scala"
public class Person extends java.lang.Object implements scala.ScalaObject{
    public java.lang.String name();
    public void name_$eq(java.lang.String);
    public int age();
    public void age_$eq(int);
    public Person(java.lang.String, int);
}
```

This shows the signature of the `Person` class, which is basically its public API, or interface. Even in a simple example like this you can see the Scala compiler doing its work for you, creating methods like `name()`, `name_$eq`, `age()`, and `age_$eq`.

Using `scalac` print options

Depending on your needs, another approach is to use the “print” options available with the `scalac` command. These are demonstrated in detail in [Recipe 3.1, “Looping with `for` and `foreach`”](#).

As that recipe shows, you begin with a file named `Main.scala` that has these contents:

```
class Main {
    for (i <- 1 to 10) println(i)
}
```

Next, compile this code with the `scalac -Xprint:parse` command:

```
$ scalac -Xprint:parse Main.scala

[[syntax trees at end of parser]] // Main.scala
package <empty> {
    class Main extends scala.AnyRef {
        def <init>() = {
```

```

        super.<init>();
        ()
    };
    1.to(10).foreach(((i) => println(i)))
}
}

```

Recipe 3.1 demonstrates that the initial Scala for loop is translated into a foreach method call, as shown by this line in the compiler output:

```
1.to(10).foreach(((i) => println(i)))
```

If you want to see more details, use the `-Xprint:all` option instead of `-Xprint:parse`. For this simple class, this command yields more than 200 lines of output. A portion of the code at the end of the output looks like this:

```

class Main extends Object {
  def <init>(): Main = {
    Main.super.<init>();
    RichInt.this.to$extension0(scala.this.Predef.intWrapper(1),
    10).foreach$mVc$sp({
      (new anonymous class anonfun$1(Main.this): Function1)
    });
    ()
  }
};

```

As you can see, your beautiful Scala code gets translated into something quite different, and this is only part of the output.

Whereas `scalac -Xprint:all` prints a *lot* of output, the basic `scalac -print` command only prints the output shown at the very end of the `-Xprint:all` output. The `scalac` manpage states that this print option, “Prints program with all Scala-specific features removed.” View the manpage for the `scalac` command to see other `-Xprint` options that are available.

Use a decompiler

Depending on class versions and legal restrictions, you may be able to take this approach a step further and decompile a class file back to its Java source code representation using a Java decompiler tool, such as *JAD*. Continuing from the previous example, you can decompile the *Main.class* file like this:

```

$ jad Main

Parsing Main...Parsing inner class Main$$anonfun$1.class...
Generating Main.jad

```

The *Main.jad* file that results from this process contains the following Java source code:

```

import scala.*;
import scala.collection.immutable.Range;
import scala.runtime.*;

```

```

public class Main
{

    public Main()
    {
        RichInt$.MODULE$.to$extension0(Predef$.MODULE$.intWrapper(1),
            10).foreach$mVc$sp(new Serializable() {

            public final void apply(int i)
            {
                apply$mCvI$sp(i);
            }

            public void apply$mCvI$sp(int v1)
            {
                Predef$.MODULE$.println(BoxesRunTime.boxToInteger(v1));
            }

            public final volatile Object apply(Object v1)
            {
                apply(BoxesRunTime.unboxToInt(v1));
                return BoxedUnit.UNIT;
            }

            public static final long serialVersionUID = 0L;

        });
    }
}

```

Though you may have to be careful with legal issues when using a decompiler, when you're first learning Scala, a tool like **JAD** or the **Java Decompiler Project** can really help to see how your Scala source code is converted into Java source code. Additionally, both Eclipse and IntelliJ offer decompiler plug-ins that are based on JAD or the Java Decompiler Project.

Discussion

Disassembling class files with `javap` can be a helpful way to understand how Scala works. As you saw in the first example with the `Person` class, defining the constructor parameters `name` and `age` as `var` fields generates quite a few methods for you.

As a second example, take the `var` attribute off both of those fields, so you have this class definition:

```

class Person (name: String, age: Int)

```

Compile this class with `scalac`, and then run `javap` on the resulting class file. You'll see that this results in a much shorter class signature:

```
$ javap Person
Compiled from "Person.scala"
public class Person extends java.lang.Object
implements scala.Serializable{
    public Person(java.lang.String, int);
}
```

Conversely, leaving `var` on both fields and turning the class into a case class significantly expands the amount of code Scala generates on your behalf. To see this, change the code in *Person.scala* so you have this case class:

```
case class Person (var name: String, var age: Int)
```

When you compile this code, it creates two output files, *Person.class* and *Person\$.class*. Disassemble these two files using `javap`:

```
$ javap Person
Compiled from "Person.scala"
public class Person extends java.lang.Object implements scala.Serializable,scala.Product,scala.Serializable{
    public static final scala.Function1 tupled();
    public static final scala.Function1 curry();
    public static final scala.Function1 curried();
    public scala.collection.Iterator productIterator();
    public scala.collection.Iterator productElements();
    public java.lang.String name();
    public void name_$eq(java.lang.String);
    public int age();
    public void age_$eq(int);
    public Person copy(java.lang.String, int);
    public int copy$default$2();
    public java.lang.String copy$default$1();
    public int hashCode();
    public java.lang.String toString();
    public boolean equals(java.lang.Object);
    public java.lang.String productPrefix();
    public int productArity();
    public java.lang.Object productElement(int);
    public boolean canEqual(java.lang.Object);
    public Person(java.lang.String, int);
}
```

```
$ javap Person$
Compiled from "Person.scala"
public final class Person$ extends scala.runtime.AbstractFunction2 implements scala.Serializable,scala.Serializable{
    public static final Person$ MODULE$;
    public static {};
    public final java.lang.String toString();
    public scala.Option unapply(Person);
    public Person apply(java.lang.String, int);
    public java.lang.Object readResolve();
    public java.lang.Object apply(java.lang.Object, java.lang.Object);
}
```

As shown, when you define a class as a *case class*, Scala generates a *lot* of code for you. This output shows the signature for that code. See [Recipe 4.14, “Generating Boilerplate Code with Case Classes”](#) for a detailed discussion of this code.

See Also

- Information on the [JAD decompiler](#)
- The [Java Decompiler project](#)

14.7. Finding Scala Libraries

Problem

Ruby has the RubyGems package manager, which lets developers easily distribute and manage the installation of Ruby libraries; does Scala have anything like this?

Solution

Prior to Scala 2.9.2, a tool named sbaz shipped with Scala, but it wasn't very popular. Instead, most tools are “discovered” by paying attention to the mailing lists, using a search engine, and being aware of a few key websites.

As discussed in [Chapter 18](#), once you've found a tool you want to use, you usually add it as a dependency to your project with SBT. For instance, to include libraries into your project, such as ScalaTest and Mockito, just add lines like this to your SBT *build.sbt* file:

```
resolvers += "Typesafe Repository" at␣  
    "http://repo.typesafe.com/typesafe/releases/"  
  
libraryDependencies += Seq(  
    "org.scalatest" %% "scalatest" % "1.8" % "test",  
    "org.mockito" % "mockito-core" % "1.9.0" % "test"  
)
```

SBT has become the de facto tool for building Scala applications and managing dependencies. Possibly because of this success, a system like RubyGems hasn't evolved, or been necessary.

Some of the top ways of finding Scala libraries are:

- Searching for libraries using a search engine, or ls.implicit.ly.
- Asking questions and searching the scala-tools@googlegroups.com and scala-language@googlegroups.com mailing lists.
- New software is also announced at the “scala-announce” mailing list; you can find a list of Scala mailing lists [online](#).
- Viewing tools listed at the [Scala wiki](#).
- Scala project updates are often noted at <http://notes.implicit.ly/>, the archive is at <http://notes.implicit.ly/archive>, and you can search for tools at <http://ls.implicit.ly/>.
- Asking questions on [StackOverflow.com](http://stackoverflow.com).

The search engine at ls.implicit.ly is interesting. The owners advertise the site as “A card catalog for Scala libraries.” As they state on their website, they make two assumptions regarding their search process:

- The library you’re looking for is an open source library that’s hosted at GitHub.
- You build your projects with [SBT](#).

For instance, if you search for “logging,” the website currently shows tools like the “Grizzled-SLF4J” library.

14.8. Generating Documentation with scaladoc

Problem

You’ve annotated your Scala code with Scaladoc, and you want to generate developer documentation for your API.

Solution

To generate Scaladoc API documentation, document your code using Scaladoc tags, and then create the documentation using an SBT task or the `scaladoc` command.

You can mark up your source code using [Scaladoc tags](#) as well as [a wiki-like syntax](#). The following code shows many of the Scaladoc tags and a few of the wiki-style markup tags:

```
package com.acme.foo
```

```
/**
 * A class to represent a 'human being'.
 *
 * Specify the `name`, `age`, and `weight` when creating a new `Person`,
 * then access the fields like this:
 * {{{
 * val p = Person("Al", 42, 200.0)
 * p.name
 * p.age
 * p.weight
 * }}}
 *
 * Did you know: The [[com.acme.foo.Employee]] extends this class.
 *
 * @constructor Create a new person with a `name`, `age`, and `weight`.
 * @param name The person's name.
 * @param age The person's age.
 * @param weight The person's weight.
 * @author Alvin Alexander
 * @version 1.0
 * @todo Add more functionality.
 * @see See [[http://alvinalexander.com alvinalexander.com]] for more
 * information.
 */
@deprecated("The `weight` field is going away", "1.0")
class Person (var name: String, var age: Int, var weight: Double) {

  /**
   * @constructor This is an auxiliary constructor. Just need a `name` here.
   */
  def this(name: String) {
    this(name, 0, 0.0)
  }

  /**
   * @return Returns a greeting based on the `name` field.
   */
  def greet = s"Hello, my name is $name"
}
```

```

/**
 * @constructor Create a new `Employee` by specifying their `name`, `age`,
 * and `role`.
 * @param name The employee's name.
 * @param age The employee's age.
 * @param role The employee's role in the organization.
 * @example val e = Employee("Al", 42, "Developer")
 */
class Employee(name: String, age: Int, role: String) extends Person(name, age, 0)
{

  /**
   * @throws boom Throws an Exception 100% of the time, be careful.
   */
  @throws(classOf[Exception])
  def boom { throw new Exception("boom") }

  /**
   * @return Returns a greeting based on the `other` and `name` fields.
   * @param other The name of the person we're greeting.
   */
  override def greet(other: String) = s"Hello $other, my name is $name"

}

```

With this code saved to a file named *Person.scala*, generate the Scaladoc documentation with the `scaladoc` command:

```
$ scaladoc Person.scala
```

This generates a root *index.html* file and other related files for your API documentation.

Similarly, if you're using SBT, generate Scaladoc API documentation by running the `sbt doc` command in the root directory of your project:

```
$ sbt doc
```

This generates the same API documentation, and places it under the *target* directory of your SBT project. With Scala 2.10 and SBT 0.12.3, the root file is located at *target/scala-2.10/api/index.html*.

Figure 14-3 shows the resulting Scaladoc for the *Person* class, and Figure 14-4 shows the Scaladoc for the *Employee* class. Notice how the Scaladoc and wiki tags affect the documentation.



com.acme.foo
Person

class **Person** extends AnyRef

A class to represent a *human being*.

Specify the name, age, and weight when creating a new **Person**, then access the fields like this:

```
val p = Person("Al", 42, 200.0)
p.name
p.age
p.weight
```

Did you know: The [com.acme.foo.Employee](#) extends this class.

Annotations	@deprecated
Deprecated	(Since version 1.0) The weight field is going away
Version	1.0
To do	Add more functionality.
See also	See alvinalexander.com for more information.

▶ Linear Supertypes

▶ Known Subclasses

Ordering

Inherited

[Learn more about member selection](#)

Visibility

Instance Constructors

new **Person**(name: String)

▼

new **Person**(name: String, age: Int, weight: Double)

Create a new person with a name, age, and weight.

name

The person's name.

age

The person's age.

weight

The person's weight.

Value Members

var **age**: Int

The person's age.

▼

def **greet**: String

returns

Returns a greeting based on the name field.

var **name**: String

The person's name.

var **weight**: Double

The person's weight.

Figure 14-3. The Scaladoc for the **Person** class

14.8. Generating Documentation with scaladoc | 475


com.acme.foo

Employee

class **Employee** extends [Person](#)

Example:

```
val e = Employee("Rocky", 33, "Developer")
```

▼ Linear Supertypes

[Person](#), AnyRef, Any

Ordering

Alphabetic
By inheritance

Inherited

Employee
Person
AnyRef
Any

Hide All

Show all

[Learn more about member selection](#)

Visibility

Public
All

Instance Constructors

▼

```
new Employee(name: String, age: Int, role: String)
```

Create a new **Employee** by specifying their name, age, and role.

name	The employee's name.
age	The employee's age.
role	The employee's role in the organization.

Value Members

▼

```
def boom: Unit
```

Annotations	@throws(classOf[Exception])
Exceptions thrown	boom Throws an Exception 100% of the time, be careful.

▼

```
def greet(other: String): String
```

other	The name of the person we're greeting.
returns	Returns a greeting based on the other and name fields.

▼

```
def greet: String
```

returns	Returns a greeting based on the name field.
Definition Classes	Person

▼

```
var weight: Double
```

Definition Classes	Person
--------------------	------------------------

Figure 14-4. The Scaladoc for the *Employee* class

Discussion

Most Scaladoc tags are similar to Javadoc tags. Common Scaladoc tags are shown in [Table 14-1](#).

Table 14-1. Common Scaladoc tags

Tag	Description	Number allowed
@author	The author of the class.	Multiple tags are allowed
@constructor	Documentation you want to provide for the constructor.	One (does not currently work on auxiliary constructors)
@example	Provide an example of how to use a method or constructor.	Multiple
@note	Document pre- and post-conditions, and other requirements.	Multiple
@param	Document a method or constructor parameter.	One per parameter
@return	Document the return value of a method.	One
@see	Describe other sources of related information.	Multiple
@since	Used to indicate that a member has been available since a certain version release.	One
@todo	Document “to do” items for a method or class.	Multiple
@throws	Document an exception type that can be thrown by a method or constructor.	Multiple
@version	The version of a class.	One

These are just some of the common tags. Other tags include @define, @migration, @tparam, and @usecase. Other Scala annotation tags like @deprecated and @throws also result in output to your documentation.

As shown in the source code, you can format your documentation using wiki-like tags. Table 14-2 shows the most common wiki *character* formatting tags, and Table 14-3 shows the most common wiki *paragraph* formatting tags.

Table 14-2. Scaladoc wiki character formatting tags

Format	Tag example
Bold	<code>'''foo'''</code>
Italic	<code>''foo''</code>
Monospace (fixed-width)	<code>`foo`</code>
Subscript	<code>,,foo,,</code>
Superscript	<code>^foo^</code>
Underline	<code>__foo__</code>

Table 14-3. Scaladoc wiki paragraph formatting tags

Format	Tag example
Headings	<code>=heading1= ==heading2== ===heading3===</code>
New paragraph	A blank line starts a new paragraph

Format	Tag example
Source code block	<pre>// all on one line {{{ if (foo) bar else baz }}} // multiple lines {{{ val p = Person("Al", 42) p.name p.age }}}</pre>

Table 14-4 shows how to create hyperlinks in Scaladoc.

Table 14-4. Scaladoc hyperlink tags

Link type	Tag example
Link to a Scala type	<code>[[scala.collection.immutable.List]]</code>
Link to an external web page	<code>[[http://alvinalexander.com My website]]</code>

The Scaladoc tags and annotations are described in more detail in the [Scala wiki](#), as well as the [Wiki markup tags](#).

Generating Scaladoc documentation with SBT

SBT has several commands that can be used to generate project documentation. See [Recipe 18.8, “Generating Project API Documentation”](#) for a tabular listing of those commands.

See Also

- [Recipe 5.8, “Declaring That a Method Can Throw an Exception”](#) and [Recipe 17.2, “Add Exception Annotations to Scala Methods to Work with Java”](#) for demonstrations of the `@throws` annotation
- Scaladoc [wiki-like syntax](#)
- Scaladoc [tags](#)
- The Scaladoc page in the [Scala Style Guide](#)
- [Recipe 18.8, “Generating Project API Documentation”](#) for details on generating Scaladoc documentation with SBT

14.9. Faster Command-Line Compiling with fsc

Problem

You're making changes to a project and recompiling it with `scalac`, and you'd like to reduce the compile time.

Solution

Use the `fsc` command instead of `scalac` to compile your code:

```
$ fsc *.scala
```

The `fsc` command works by starting a compilation daemon and also maintains a cache, so compilation attempts after the first attempt run much faster than `scalac`.

Discussion

Although the primary advantage is that compile times are significantly improved when recompiling the same code, it's important to be aware of a few caveats, per the `fsc` manpage:

- “The tool is especially effective when repeatedly compiling with the same class paths, because the compilation daemon can reuse a compiler instance.”
- “The compilation daemon is smart enough to flush its cached compiler when the class path changes. However, if the contents of the class path change, for example due to upgrading a library, then the daemon should be explicitly shut down with `-shutdown`.”

As an example of the second caveat, if the JAR files on the classpath have changed, you should shut down the daemon, and then reissue your `fsc` command:

```
$ fsc -shutdown  
[Compile server exited]
```

```
$ fsc *.scala
```

On Unix systems, running `fsc` creates a background process with the name `CompileServer`. You can see information about this process with the following `ps` command:

```
$ ps auxw | grep CompileServer
```

See the `fsc` manpage for more information.

See Also

- The `fsc` manpage (type `man fsc` at the command line).
- When using SBT, you can achieve similar performance improvements by working in the SBT shell instead of your operating system's command line. See [Recipe 18.2, “Compiling, Running, and Packaging a Scala Project with SBT”](#) for more information.

14.10. Using Scala as a Scripting Language

Problem

You want to use Scala as a scripting language on Unix systems, replacing other scripts you've written in a Unix shell (Bourne Shell, Bash), Perl, PHP, Ruby, etc.

Solution

Save your Scala code to a text file, making sure the first three lines of the script contain the lines shown, which will execute the script using the `scala` interpreter:

```
#!/bin/sh
exec scala "$@" "$@"
!#

println("Hello, world")
```

To test this, save the code to a file named *hello.sh*, make it executable, and then run it:

```
$ chmod +x hello.sh

$ ./hello.sh
Hello, world
```

As detailed in the next recipe, command-line parameters to the script can be accessed via an `args` array, which is implicitly made available to you:

```
#!/bin/sh
exec scala "$@" "$@"
!#

args.foreach(println)
```

Discussion

Regarding the first three lines of a shell script:

- The `#!` in the first line is the usual way to start a Unix shell script. It invokes a Unix Bourne shell.
- The `exec` command is a shell built-in. `$0` expands to the name of the shell script, and `$@` expands to the positional parameters.
- The `!#` characters as the third line of the script is how the header section is closed.

A great thing about using Scala in your scripts is that you can use all of its advanced features, such as the ability to create and use classes in your scripts:

```
#!/bin/sh
exec scala "$0" "$@"
!#

class Person(var firstName: String, var lastName: String) {
  override def toString = firstName + " " + lastName
}

println(new Person("Nacho", "Libre"))
```

Using the App trait or main method

To use an App trait in a Scala script, start the script with the usual first three header lines, and then create an object that extends the App trait:

```
#!/bin/sh
exec scala "$0" "$@"
!#

object Hello extends App {
  println("Hello, world")
  // if you want to access the command line args:
  //args.foreach(println)
}

Hello.main(args)
```

The last line in that example shows how to pass the script's command-line arguments to the implicit `main` method in the `Hello` object. As usual in an App trait object, the arguments are available via a variable named `args`.

You can also define an object with a `main` method to kick off your shell script action:

```
#!/bin/sh
exec scala "$0" "$@"
!#

object Hello {
  def main(args: Array[String]) {
    println("Hello, world")
    // if you want to access the command line args:
    //args.foreach(println)
  }
}
```

```
}
}
```

```
Hello.main(args)
```

Building the classpath

If your shell script needs to rely on external dependencies (such as JAR files), add them to your script's classpath using this syntax:

```
#!/bin/sh
exec scala -classpath "lib/htmlcleaner-2.2.jar:lib/scalaemail_2.10.0-1.0.jar:lib/stockutils_2.10.0-1.0.jar" "$@" "$@"
!#
```

You can then import these classes into your code as usual. The following code shows a complete script I wrote that retrieves stock quotes and mails them to me:

```
#!/bin/sh
exec scala -classpath "lib/htmlcleaner-2.2.jar:lib/scalaemail_2.10.0-1.0.jar:lib/stockutils_2.10.0-1.0.jar" "$@" "$@"
!#

import java.io._
import scala.io.Source
import com.devdaily.stocks.StockUtils
import scala.collection.mutable.ArrayBuffer

object GetStocks {

  case class Stock(symbol: String, name: String, price: BigDecimal)

  val DIR = System.getProperty("user.dir")
  val SLASH = System.getProperty("file.separator")
  val CANON_STOCKS_FILE = DIR + SLASH + "stocks.dat"
  val CANON_OUTPUT_FILE = DIR + SLASH + "quotes.out"

  def main(args: Array[String]) {

    // read the stocks file into a list of strings ("AAPL|Apple")
    val lines = Source.fromFile(CANON_STOCKS_FILE).getLines().toList

    // create a list of Stock from the symbol, name, and by
    // retrieving the price
    var stocks = new ArrayBuffer[Stock]()
    lines.foreach{ line =>
      val fields = line.split("\\|")
      val symbol = fields(0)
      val html = StockUtils.getHtmlFromUrl(symbol)
      val price = StockUtils.extractPriceFromHtml(html, symbol)
      val stock = Stock(symbol, fields(1), BigDecimal(price))
      stocks += stock
    }
  }
}
```



```

    // build a string to output
    var sb = new StringBuilder
    stocks.foreach { stock =>
        sb.append("%s is %s\n".format(stock.name, stock.price))
    }
    val output = sb.toString

    // write the string to the file
    val pw = new PrintWriter(new File(CANON_OUTPUT_FILE))
    pw.write(output)
    pw.close
}
}

GetStocks.main(args)

```

I run this script twice a day through a crontab entry on a Linux server. The *stocks.dat* file it reads has entries like this:

```

AAPL|Apple
KKD|Krispy Kreme
NFLX|Netflix

```

See Also

- More about the first three lines of these shell script examples at [my site](#)
- [Recipe 14.13, “Make Your Scala Scripts Run Faster”](#) for a way to make your scripts run faster

14.11. Accessing Command-Line Arguments from a Script

Problem

You want to access the command-line arguments from your Scala shell script.

Solution

Use the same script syntax as shown in [Recipe 14.8, “Generating Documentation with scaladoc”](#), and then access the command-line arguments using `args`, which is a `List[String]` that is implicitly made available:

```

#!/bin/sh
exec scala "$@" "$@"
!#

args.foreach(println)

```

Save this code to a file named *args.sh*, make the file executable, and run it like this:

```
$ ./args.sh a b c
a
b
c
```

Discussion

Because the implicit field *args* is a `List[String]`, you can perform all the usual operations on it, including getting its size, and accessing elements with the usual syntax.

In a more “real-world” example, you’ll check for the number of command-line arguments, and then assign those arguments to values. This is demonstrated in the following script:

```
#!/bin/sh
exec scala "$@" "$@"
!#

if (args.length != 2) {
  Console.err.println("Usage: replacer <search> <replace>")
  System.exit(1)
}

val searchPattern = args(0)
val replacePattern = args(1)

println(s"Replacing $searchPattern with $replacePattern ...")

// more code here ...
```

When this script is run from the command line without arguments, the result looks like this:

```
$ ./args.sh
Usage: replacer <search> <replace>
```

When it’s run with the correct number of arguments, the result looks like this:

```
$ ./args.sh foo bar
Replacing foo with bar ...
```

If you decide to use the *App* trait in your script, make sure you pass the command-line arguments to your *App* object, as shown in the `Hello.main(args)` line in this example:

```
#!/bin/sh
exec scala "$@" "$@"
!#

object Hello extends App {
  println("Hello, world")
  // if you want to access the command line args:
```

```
//args.foreach(println)
}
```

```
Hello.main(args)
```

Use the same syntax if you use a `main` method instead of an App object.

14.12. Prompting for Input from a Scala Shell Script

Problem

You want to prompt a user for input from a Scala shell script and read her responses.

Solution

Use the `readLine`, `print`, `printf`, and `Console.read*` methods to read user input, as demonstrated in the following script. Comments in the script describe each method:

```
#!/bin/sh
exec scala "$@" "$@"
!#

// write some text out to the user with Console.println
Console.println("Hello")

// Console is imported by default, so it's not really needed, just use println
println("World")

// readLine lets you prompt the user and read their input as a String
val name = readLine("What's your name? ")

// readInt lets you read an Int, but you have to prompt the user manually
print("How old are you? ")
val age = readInt()

// you can also print output with printf
println(s"Your name is $name and you are $age years old.")
```

Discussion

The `readLine` method lets you prompt a user for input, but the other `read*` methods don't, so you need to prompt the user manually with `print`, `println`, or `printf`.

You can list the `Console.read*` methods in the Scala REPL:

```
scala> Console.read
readBoolean  readByte  readChar  readDouble  readFloat
readInt      readLine  readLong  readShort   readf
readf1       readf2    readf3
```

Be careful with the methods that read numeric values; as you might expect, they can all throw a `NumberFormatException`.

Although these methods are thorough, if you prefer, you can also fall back and read input with the Java Scanner class:

```
// you can also use the Java Scanner class, if desired
val scanner = new java.util.Scanner(System.in)
print("Where do you live? ")
val input = scanner.nextLine()
print(s"I see that you live in $input")
```

Reading multiple values from one line

If you want to read multiple values from one line of user input (such as a person's name, age, and weight), there are several approaches to the problem.

To my surprise, I prefer to use the Java Scanner class. The following code demonstrates the Scanner approach:

```
import java.util.Scanner

// simulated input
val input = "Joe 33 200.0"

val line = new Scanner(input)
val name = line.next
val age = line.nextInt
val weight = line.nextDouble
```

To use this approach in a shell script, replace the `input` line with a `readLine()` call, like this:

```
val input = readLine()
```

Of course if the input doesn't match what you expect, an error should be thrown. The `Scanner` `next*` methods throw a `java.util.InputMismatchException` when the data doesn't match what you expect, so you'll want to wrap this code in a `try/catch` block.

I initially assumed that one of the `readf` methods on the `Console` object would be the best solution to this problem, but unfortunately they return their types as `Any`, and then you have to cast them to the desired type. For instance, suppose you want to read the same name, age, and weight information as the previous example. After prompting the user, you read three values with the `readf3` method like this:

```
val(a,b,c) = readf3("{0} {1,number} {2,number}")
```

If the user enters a string followed by two numbers, a result is returned, but if he enters an improperly formatted string, such as `1 a b`, the expression fails with a `ParseException`:

```
java.text.ParseException: MessageFormat parse error!
at java.text.MessageFormat.parse(MessageFormat.java:1010)
```

```
at scala.Console$.readf(Console.scala:413)
at scala.Console$.readf3(Console.scala:445)
```

Unfortunately, even if the user enters the text as desired, you still need to cast the values to the correct type, because the variables `a`, `b`, and `c` are of type `Any`. You can try to cast them with this approach:

```
val name = a
val age = b.asInstanceOf[Long]
val weight = c.asInstanceOf[Double]
```

Or convert them like this:

```
val name = a.toString
val age = b.toString.toInt
val weight = c.toString.toDouble
```

But for me, the `Scanner` is cleaner and easier.

A third approach is to read the values in as a `String`, and then split them into tokens. Here's what this looks like in the REPL:

```
scala> val input = "Michael 54 250.0"
input: String = Michael 54 250.0

scala> val tokens = input.split(" ")
tokens: Array[String] = Array(Michael, 54, 250.0)
```

The `split` method creates an `Array[String]`, so access the array elements and cast them to the desired types to create your variables:

```
val name = tokens(0)
val age = tokens(1).toInt
val weight = tokens(2).toDouble
```

Note that the `age` and `weight` fields in this example can throw a `NumberFormatException`.

A fourth way to read the user's input is by specifying a regular expression to match the input you expect to receive. Using this technique, you again receive each variable as a `String`, and then cast it to the desired type. The process looks like this in the REPL:

```
scala> val ExpectedPattern = "(.*) (\\d+) (\\d*\\.?.?\\d*)".r
ExpectedPattern: scala.util.matching.Regex = (.*) (\\d+) (\\d*\\.?.?\\d*)

// you would use readLine() here
scala> val input = "Paul 36 180.0"
input: String = Paul 36 180.0

scala> val ExpectedPattern(a, b, c) = input
a: String = Paul
b: String = 36
c: String = 180.0
```

Now that you have the variables as strings, cast them to the desired types, as before:

```
val name = a
val age = b.toInt
val weight = c.toDouble
```

The `ExpectedPattern` line in this example will fail with a `scala.MatchError` if the input doesn't match what's expected.

Hopefully with all of these examples you'll find your own preferred way to read in multiple values at one time.

Fun with output

Use `print`, `printf`, or `println` to write output. As shown in the Solution, the `readLine` method also lets you prompt a user for input.

The `Console` object contains a number of fields that you can use with the print methods to control the display. For instance, if you want your entire line of output to be underlined, change the last lines of the script to look like this:

```
val qty = 2
val pizzaType = "Cheese"
val total = 20.10

print(Console.UNDERLINED)
println(f"$qty%d $pizzaType pizzas coming up, $$$total%.2f.")
print(Console.RESET)
```

This prints the following string, underlined:

```
2 Cheese pizzas coming up, $20.10.
```

Other displayable attributes include colors and attributes such as `BLINK`, `BOLD`, `INVISIBLE`, `RESET`, `REVERSED`, and `UNDERLINED`. See the `Console` object Scaladoc page for more options.

See Also

- [Recipe 1.8, “Extracting Parts of a String That Match Patterns”](#) for more examples of the pattern-matching technique shown in this recipe.
- The [Java Scanner class](#)
- The [Java Pattern class](#)
- The [Scala Console object](#) provides the `read*` methods

14.13. Make Your Scala Scripts Run Faster

Problem

You love using Scala as a scripting language, but you'd like to eliminate the lag time in starting up a script.

Solution

Use the `-savecompiled` argument of the Scala interpreter to save a compiled version of your script.

A basic Scala script like this:

```
#!/bin/sh
exec scala "$@" "$@"
!#

println("Hello, world!")
args foreach println
```

consistently runs with times like this on one of my computers:

```
real    0m1.573s
user    0m0.574s
sys     0m0.089s
```

To improve this, add the `-savecompiled` argument to the Scala interpreter line:

```
#!/bin/sh
exec scala -savecompiled "$@" "$@"
!#

println("Hello, world!")
args foreach println
```

Then run the script once. This generates a compiled version of the script. After that, the script runs with a consistently lower real time (wall clock) on all subsequent runs:

```
real    0m0.458s
user    0m0.487s
sys     0m0.075s
```

Precompiling your script shaves a nice chunk of time off the runtime of your script, even for a simple example like this.

Discussion

When you run your script the first time, Scala generates a JAR file that matches the name of your script. For instance, I named my script *test1.sh*, and then ran it like this:

```
$ ./test1.sh
```

After running the script, I looked at the directory contents and saw that Scala created a file named *test1.sh.jar*. Scala creates this new file and also leaves your original script in place.

On subsequent runs, Scala sees that there's a JAR file associated with the script, and if the script hasn't been modified since the JAR file was created, it runs the precompiled code from the JAR file instead of the source code in the script. This results in a faster runtime because the source code doesn't need to be compiled.

You can look at the contents of the JAR file using the `jar` command:

```
$ jar tvf test1.sh.jar
 43 Wed Jul 25 15:44:26 MDT 2012 META-INF/MANIFEST.MF
965 Wed Jul 25 15:44:26 MDT 2012 Main$$anon$1$$anonfun$1.class
725 Wed Jul 25 15:44:26 MDT 2012 Main$$anon$1.class
557 Wed Jul 25 15:44:26 MDT 2012 Main$.class
646 Wed Jul 25 15:44:26 MDT 2012 Main.class
```

In this example, I didn't include a `main` method in an object or use the `App` trait with an object, so Scala assumed the name `Main` for the main/primary object that it created to run my script.

Introduction

Between the Java web services libraries and the newer Scala libraries and frameworks that are available, Scala easily handles web service tasks. You can rapidly create web service clients to send and receive data using these general libraries, or solve problems with more specific libraries, such as creating a Twitter client with the Twitter4J library. There are also several good JSON libraries available, so you can easily convert between data JSON strings and Scala objects.

When it comes to creating your own RESTful web services, you can use lightweight frameworks like **Scalatra** or **Unfiltered** and have web services up and running in a matter of minutes. But you have many choices, so you can also use the **Play Framework (Play)**, **Lift Framework**, or other Scala libraries to create web services, as well as all of the previously available Java web service libraries.

As demonstrated in **Chapter 16**, Scala has nice support for the **MongoDB** database, and this chapter demonstrates how to provide a complete web services solution using Scalatra and MongoDB. This chapter shares a few recipes that are specific to using Play to create web services.

Finally, although the Scala libraries offer some nice convenience classes and methods for connecting to web services, the trusty old Java **Apache HttpClient library** is still very useful, and it's also demonstrated in several recipes.

15.1. Creating a JSON String from a Scala Object

Problem

You're working outside of a specific framework, and want to create a JSON string from a Scala object.

Solution

If you're using the Play Framework, you can use its library to work with JSON, as shown in Recipes 15.13 and 15.14, but if you're using JSON outside of Play, you can use the best libraries that are available for Scala and Java:

- [Lift-JSON](#)
- The [Google Gson library](#) (Java)
- [Json4s](#)
- [spray-json](#)

This recipe demonstrates the Lift-JSON and Gson libraries. (Json4s is a port of Lift-JSON, so it shares the same API.)

Lift-JSON solution

To demonstrate the Lift-JSON library, create an empty SBT test project. With Scala 2.10 and SBT 0.12.x, configure your *build.sbt* file as follows:

```
name := "Basic Lift-JSON Demo"

version := "1.0"

scalaVersion := "2.10.0"

libraryDependencies += "net.liftweb" %% "lift-json" % "2.5+"
```

Next, in the root directory of your project, create a file named *LiftJsonTest.scala*:

```
import scala.collection.mutable._
import net.liftweb.json._
import net.liftweb.json.Serialization.write

case class Person(name: String, address: Address)
case class Address(city: String, state: String)

object LiftJsonTest extends App {

  val p = Person("Alvin Alexander", Address("Talkeetna", "AK"))

  // create a JSON string from the Person, then print it
  implicit val formats = DefaultFormats
  val jsonString = write(p)
  println(jsonString)

}
```

This code creates a JSON string from the *Person* instance, and prints it. When you run the project with the `sbt run` command, you'll see the following JSON output:

```
{"name":"Alvin Alexander","address":{"city":"Talkeetna","state":"AK"}}
```

Gson solution

To demonstrate the Gson library, follow similar steps. Create an empty SBT test project, then download the Gson JAR file from the [Gson website](#), and place it in your project's *lib* directory.

In the root directory of the project, create a file named *GsonTest.scala* with these contents:

```
import com.google.gson.Gson

case class Person(name: String, address: Address)
case class Address(city: String, state: String)

object GsonTest extends App {
  val p = Person("Alvin Alexander", Address("Talkeetna", "AK"))

  // create a JSON string from the Person, then print it
  val gson = new Gson
  val jsonString = gson.toJson(p)
  println(jsonString)
}
```

In a manner similar to the first example, this code converts a *Person* instance to a JSON string and prints the string. When you run the project with `sbt run`, you'll see the same output as before:

```
{"name":"Alvin Alexander","address":{"city":"Talkeetna","state":"AK"}}
```

Discussion

The Lift-JSON project is a subproject of the [Lift Framework](#), which is a complete Scala web framework. Fortunately the library has been created as a separate module you can download and use on its own.

In addition to working with simple classes, it works well with Scala collections. The following example shows how to generate JSON strings from a simple Scala Map:

```
import net.liftweb.json.JsonAST
import net.liftweb.json.JsonDSL._
import net.liftweb.json.Printer.{compact,pretty}

object LiftJsonWithCollections extends App {

  val json = List(1, 2, 3)
  println(compact(JsonAST.render(json)))

  val map = Map("fname" -> "Alvin", "lname" -> "Alexander")
  println(compact(JsonAST.render(map)))

}
```

That program prints the following output:

```
[1,2,3]
{"fname":"Alvin","lname":"Alexander"}
```

When communicating with other computer systems you'll want to use the `compact` method as shown, but when a human needs to look at your JSON strings, use the `pretty` method instead:

```
println(pretty(JsonAST.render(map)))
```

This changes the `map` output to look like this:

```
{
  "fname": "Alvin",
  "lname": "Alexander"
}
```

The Lift-JSON examples in this recipe work well for either objects or collections, but when you have an object that *contains* collections, such as a `Person` class that has a list of friends defined as `List[Person]`, it's best to use the Lift-JSON DSL. This is demonstrated in [Recipe 15.2](#).

Gson is a Java library that you can use to convert back and forth between Scala objects and their JSON representation. From the Gson documentation:

There are a few open-source projects that can convert Java objects to JSON. However, most of them require that you place Java annotations in your classes; something that you can not do if you do not have access to the source-code. Most also do not fully support the use of Java Generics. Gson considers both of these as very important design goals.

I used Gson to generate JSON for a while, but because it's written in Java, it has a few issues when trying to work with Scala collections. One such problem is demonstrated in [Recipe 15.2](#).

See Also

- The [Lift-JSON library](#)
- The [Gson library](#)
- A project named [Json4s](#) aims to provide a unified interface for all Scala JSON projects. The current package is a port of Lift-JSON, with support for using the Java Jackson library as a backend as well.
- [spray-json](#) is another popular Scala JSON library.

15.2. Creating a JSON String from Classes That Have Collections

Problem

You want to generate a JSON representation of a Scala object that contains one or more collections, such as a `Person` class that has a list of friends or addresses.

Solution

Once classes start containing collections, converting them to JSON becomes more difficult. In this situation, I prefer to use the Lift-JSON domain-specific library (DSL) to generate the JSON.

Lift-JSON version 1

The Lift-JSON library uses its own DSL for generating JSON output from Scala objects. As shown in the previous recipe, this isn't necessary for simple objects, but it is necessary once objects become more complex, specifically once they contain collections. The benefit of this approach is that you have complete control over the JSON that is generated.

The following example shows how to generate a JSON string for a `Person` class that has a `friends` field defined as `List[Person]`:

```
import net.liftweb.json._
import net.liftweb.json.JsonDSL._

case class Person(name: String, address: Address) {
  var friends = List[Person]()
}

case class Address(city: String, state: String)

object LiftJsonListsVersion1 extends App {

  //import net.liftweb.json.JsonParser._
  implicit val formats = DefaultFormats

  val merc = Person("Mercedes", Address("Somewhere", "KY"))
  val mel = Person("Mel", Address("Lake Zurich", "IL"))
  val friends = List(merc, mel)
  val p = Person("Alvin Alexander", Address("Talkeetna", "AK"))
  p.friends = friends

  // define the json output
  val json =
    ("person" ->
      ("name" -> p.name) ~
```

```

        ("address" ->
          ("city" -> p.address.city) ~
          ("state" -> p.address.state)) ~
        ("friends" ->
          friends.map { f =>
            ("name" -> f.name) ~
            ("address" ->
              ("city" -> f.address.city) ~
              ("state" -> f.address.state))
          })
      )

    println(pretty(render(json)))
  }
}

```

The JSON output from this code looks like this:

```

{
  "person": {
    "name": "Alvin Alexander",
    "address": {
      "city": "Talkeetna",
      "state": "AK"
    },
    "friends": [ {
      "name": "Mercedes",
      "address": {
        "city": "Somewhere",
        "state": "KY"
      }
    }, {
      "name": "Mel",
      "address": {
        "city": "Lake Zurich",
        "state": "IL"
      }
    }
  ]
}

```

The JSON-generating code is shown after the “define the json output” comment, and is repeated here:

```

val json =
  ("person" ->
    ("name" -> p.name) ~
    ("address" ->
      ("city" -> p.address.city) ~
      ("state" -> p.address.state)) ~
    ("friends" ->
      friends.map { f =>
        ("name" -> f.name) ~

```

```

        ("address" ->
          ("city" -> f.address.city) ~
          ("state" -> f.address.state))
      })
    )
  )
}

```

As you can see, Lift uses a custom DSL to let you generate the JSON, and also have control over how the JSON is generated (as opposed to using reflection to generate the JSON). Although you'll want to read the details of the DSL to take on more difficult tasks, the basics are straightforward.

The first thing to know is that any `Tuple2` generates a JSON field, so a code snippet like `("name" -> p.name)` produces this output:

```
"name": "Alvin Alexander"
```

The other important thing to know is that the `~` operator lets you join fields. You can see from the example code and output how it works.

You can also refer to objects and methods when generating the JSON. You can see this in sections of the code like `p.address.city` and `friends.map { f => ...}`. Writing JSON-generating code like this feels just like writing other Scala code.

Lift-JSON Version 2

As your classes grow, creating a larger JSON generator in one variable becomes hard to read and maintain. Fortunately, with the Lift-JSON DSL you can break your JSON-generating code down into small chunks to keep the code maintainable. The following code achieves the same result as the previous example, but I've broken the JSON-generating code down into small methods that are easier to maintain and reuse:

```

import net.liftweb.json._
import net.liftweb.json.JsonDSL._

object LiftJsonListsVersion2 extends App {

  val merc = Person("Mercedes", Address("Somewhere", "KY"))
  val mel = Person("Mel", Address("Lake Zurich", "IL"))
  val friends = List(merc, mel)
  val p = Person("Alvin Alexander", Address("Talkeetna", "AK"))
  p.friends = friends

  val json =
    ("person" ->
      ("name" -> p.name) ~
      getAddress(p.address) ~
      getFriends(p)
    )

  println(pretty(render(json)))

  def getFriends(p: Person) = {

```

```

    ("friends" ->
      p.friends.map { f =>
        ("name" -> f.name) ~
        getAddress(f.address)
      })
  }

  def getAddress(a: Address) = {
    ("address" ->
      ("city" -> a.city) ~
      ("state" -> a.state))
  }
}

case class Person(name: String, address: Address) {
  var friends = List[Person]()
}

case class Address(city: String, state: String)

```

As shown, this approach lets you create methods that can be reused. The `getAddress` method, for instance, is called several times in the code.

Discussion

As shown in [Recipe 15.1](#), Gson works via reflection, and it works well for simple classes. However, I've found it to be harder to use when your classes have certain collections. For instance, the following code works fine when the list of friends is defined as an `Array[Person]`:

```

import com.google.gson.Gson
import com.google.gson.GsonBuilder

case class Person(name: String, address: Address) {
  var friends: Array[Person] = _
}

case class Address(city: String, state: String)

/**
 * This approach works with Array.
 */
object GsonWithArray extends App {

  val merc = Person("Mercedes", Address("Somewhere", "KY"))
  val mel = Person("Mel", Address("Lake Zurich", "IL"))
  val friends = Array(merc, mel)
  val p = Person("Alvin Alexander", Address("Talkeetna", "AK"))
  p.friends = friends
  val gson = (new GsonBuilder()).setPrettyPrinting.create

```



```
println(gson.toJson(p))

}
```

Because a Scala Array is backed by a Java array, that code works well, generating JSON output that is similar to Lift-JSON. However, if you change the `Array[Person]` to `List[Person]`, Gson removes the list of friends from the output:

```
{
  "name": "Alvin Alexander",
  "address": {
    "city": "Talkeetna",
    "state": "AK"
  },
  "friends": {}
}
```

Changing the Array to an `ArrayBuffer` also causes problems and exposes the internal implementation of an `ArrayBuffer`:

```
{
  "name": "Alvin Alexander",
  "address": {
    "city": "Talkeetna",
    "state": "AK"
  },
  "friends": {
    "initialSize": 16,
    "array": [
      {
        "name": "Mercedes",
        "address": {
          "city": "Somewhere",
          "state": "KY"
        }
      },
      {
        "name": "Mel",
        "address": {
          "city": "Lake Zurich",
          "state": "IL"
        }
      },
      null, // this line is repeated 13 more times
      ...
      null
    ],
    "size0": 2
  }
}
```

An `ArrayBuffer` begins with 16 elements, and when Gson generates the JSON for the list of friends, it correctly includes the two friends, but then outputs the word `null` 14 times, along with including the other output shown.

If you like the idea of generating JSON from your code using reflection, see the Gson User Guide link in the See Also section for information on how to try to resolve these issues by writing custom serializers (creating a JSON string from an object) and deserializers (creating an object from a JSON string).

See Also

- The [Lift-JSON library](#).
- The [Gson User Guide](#) shows how to write serializers and deserializers.

15.3. Creating a Simple Scala Object from a JSON String

Problem

You need to convert a JSON string into a simple Scala object, such as a Scala case class that has no collections.

Solution

Use the Lift-JSON library to convert a JSON string to an instance of a case class. This is referred to as *deserializing* the string into an object.

The following code shows a complete example of how to use Lift-JSON to convert a JSON string into a case class named `MailServer`. As its name implies, `MailServer` represents the information an email client needs to connect to a server:

```
import net.liftweb.json._

// a case class to represent a mail server
case class MailServer(url: String, username: String, password: String)

object JsonParsingExample extends App {

  implicit val formats = DefaultFormats

  // simulate a json string
  val jsonString = """
{
  "url": "imap.yahoo.com",
  "username": "myusername",
  "password": "mypassword"
}
```

"""

```
// convert a String to a JValue object
val jValue = parse(jsonString)

// create a MailServer object from the string
val mailServer = jValue.extract[MailServer]
println(mailServer.url)
println(mailServer.username)
println(mailServer.password)

}
```

In this example, the `jsonString` contains the text you'd expect to receive if you called a web service asking for a `MailServer` instance. That string is converted into a Lift-JSON `JValue` object with the `parse` function:

```
val jValue = parse(jsonString)
```

Once you have a `JValue` object, use its `extract` method to create a `MailServer` object:

```
val mailServer = jValue.extract[MailServer]
```

The `JValue` class is the root class in the Lift-JSON abstract syntax tree (AST), and its `extract` method builds a case class instance from a JSON string.

Working with objects that have collections is a little more difficult, and that process is covered in the next recipe.

See Also

- The [Lift-JSON library](#)
- [Lift-JSON documentation](#)

15.4. Parsing JSON Data into an Array of Objects

Problem

You have a JSON string that represents an array of objects, and you need to deserialize it into objects you can use in your Scala application.

Solution

Use a combination of methods from the Lift-JSON library. The following example demonstrates how to deserialize the string `jsonString` into a series of `EmailAccount` objects, printing each object as it is deserialized:

```

import net.liftweb.json.DefaultFormats
import net.liftweb.json._

// a case class to match the JSON data
case class EmailAccount(
  accountName: String,
  url: String,
  username: String,
  password: String,
  minutesBetweenChecks: Int,
  usersOfInterest: List[String]
)

object ParseJsonArray extends App {

  implicit val formats = DefaultFormats

  // a JSON string that represents a list of EmailAccount instances
  val jsonString = """
{
  "accounts": [
    { "emailAccount": {
      "accountName": "YMail",
      "username": "USERNAME",
      "password": "PASSWORD",
      "url": "imap.yahoo.com",
      "minutesBetweenChecks": 1,
      "usersOfInterest": ["barney", "betty", "wilma"]
    }},
    { "emailAccount": {
      "accountName": "Gmail",
      "username": "USER",
      "password": "PASS",
      "url": "imap.gmail.com",
      "minutesBetweenChecks": 1,
      "usersOfInterest": ["pebbles", "bam-bam"]
    }}
  ]
}
"""

  // json is a JValue instance
  val json = parse(jsonString)

  val elements = (json \\ "emailAccount").children
  for (acct <- elements) {
    val m = acct.extract[EmailAccount]
    println(s"Account: ${m.url}, ${m.username}, ${m.password}")
    println("  Users: " + m.usersOfInterest.mkString(", "))
  }
}

```

Running this program results in the following output:

```
Account: imap.yahoo.com, USERNAME, PASSWORD
Users: barney,betty,wilma
Account: imap.gmail.com, USER, PASS
Users: pebbles,bam-bam
```

Discussion

I use code like this in my [SARAH application](#) to notify me when I receive an email message from people in the `usersOfInterest` list. SARAH scans my email inbox periodically, and when it sees an email message from people in this list, it speaks, “You have new email from Barney and Betty.”

This example begins with some sample JSON stored in a string named `jsonString`. This string is turned into a `JValue` object named `json` with the `parse` function. The `json` object is then searched for all elements named `emailAccount` using the `\\` method. This syntax is nice, because it’s consistent with the XPath-like methods used in Scala’s XML library.

The `for` loop iterates over the elements that are found, and each element is extracted as an `EmailAccount` object, and the data in that object is then printed.

Notice that the `EmailAccount` class has the `usersOfInterest` field, which is defined as `List[String]`. The Lift-JSON library converts this sequence easily, with no additional coding required.

See Also

- The Lift-JSON library is well-documented on [GitHub](#) and [Assembla](#).
- [SARAH](#) is a voice-interaction application written in Scala.

15.5. Creating Web Services with Scalatra

Problem

You want to be able to build new web services with [Scalatra](#), a lightweight Scala web framework similar to the Ruby [Sinatra](#) library.

Solution

The recommended approach to create a new Scalatra project is to use [Giter8](#), a great tool for building SBT directories for new projects.

Assuming you have Giter8 installed, use the `g8` command to create a new project with a Scalatra template:

```
$ g8 scalatra/scalatra-sbt
organization [com.example]: com.alvinalexander
package [com.example.app]: com.alvinalexander.app
name [My Scalatra Web App]:
scalatra_version [2.2.0]:
servlet_name [MyScalatraServlet]:
scala_version [2.10.0]:
version [0.1.0-SNAPSHOT]:
```

Template applied in `./my-scalatra-web-app`

When Giter8 finishes, move into the new directory it created:

```
$ cd my-scalatra-web-app
```

Start SBT in that directory, and then issue the `container:start` command to start the Jetty server:

```
$ sbt

> container:start
// a lot of output here ...
[info] Started SelectChannelConnector@0.0.0.0:8080
[success] Total time: 11 s, completed May 13, 2013 4:32:08 PM
```

Then use the following command to enable continuous compilation:

```
> ~ ;copy-resources;aux-compile
1. Waiting for source changes... (press enter to interrupt)
```

That command is nice; it automatically recompiles your source code when it changes.

The Jetty server starts on port 8080 by default. If you switch to a browser and go to the URL `http://localhost:8080/`, you should see some default “Hello, world” output, indicating that Scalatra is running.

The content displayed at this URL comes from a class named `MyScalatraServlet`, located in the project’s `src/main/scala/com/alvinalexander/app` directory:

```
package com.alvinalexander.app

import org.scalatra._
import scalate.ScalateSupport

class MyScalatraServlet extends MyScalatraWebAppStack {

  get("/") {
    <html>
    <body>
      <h1>Hello, world!</h1>
      Say <a href="hello-scalate">hello to Scalate</a>.
  }
```

```

    </body>
  </html>
}

}

```

That's the entire servlet. If you're used to building web services with "heavier" tools, this can be quite a shock.

The `get` method shown declares that it's listening to GET requests at the `/` URI. If you try accessing another URL like `http://localhost:8080/foo` in your browser, you'll see output like this in the browser:

```

Requesting "GET /foo" on servlet "" but only have:
GET /

```

This is because `MyScalaatraServlet` only has one method, and it's programmed to listen for a GET request at the `/` URI.

Add a new service

To demonstrate how the process of adding a new web service works, add a new method that listens to GET requests at the `/hello` URI. To do this, just add the following method to the servlet:

```

get("/hello") {
  <p>Hello, world!</p>
}

```

A few moments after saving this change to `MyScalaatraServlet`, you should see some output in your SBT console. An abridged version of the output looks like this:

```

[info] Compiling 1 Scala source to target/scala-2.10/classes...
[success] Total time: 8 s
[info] Generating target/scala-2.10/resource_managed/main/rebel.xml.
[info] Compiling Templates in Template Directory:
      src/main/webapp/WEB-INF/templates
[success] Total time: 1 s, completed May 28, 2013 1:56:36 PM
2. Waiting for source changes... (press enter to interrupt)

```

As a result of the `~ aux-compile` command, SBT automatically recompiles your source code. Once the code is recompiled, you can go to the `http://localhost:8080/hello` URL in your browser, where you'll see the new output.

Congratulations. By following the steps in this recipe, you should have a web service up and running in a matter of minutes.

Discussion

Giter8 is a tool for creating SBT project directory structures based on templates. The template used in this example is just one of many [Giter8 templates](#). Giter8 requires [SBT](#)

and another tool named **Conscript**. Despite these requirements, the overall installation process is simple, and is described in **Recipe 18.1**.

In addition to the `MyScalatraServlet` class, this list shows some of the most important files in your project:

```
project/build.scala
project/plugins.sbt
src/main/resources/logback.xml
src/main/scala/com/alvinalexander/app/MyScalatraServlet.scala
src/main/scala/com/alvinalexander/app/MyScalatraWebAppStack.scala
src/main/scala/ScalatraBootstrap.scala
src/main/webapp/WEB-INF/web.xml
src/main/webapp/WEB-INF/templates/layouts/default.jade
src/main/webapp/WEB-INF/templates/views/hello-scalate.jade
src/test/scala/com/alvinalexander/app/MyScalatraServletSpec.scala
```

Notice that this includes a `WEB-INF/web.xml` file. If you're used to the Java web programming world, you'll find that this is a normal `web.xml` file, albeit a very small one. Excluding the boilerplate XML, it has only this entry:

```
<listener>
  <listener-class>org.scalatra.servlet.ScalatraListener</listener-class>
</listener>
```

You'll rarely need to edit this file. **Recipe 15.6, "Replacing XML Servlet Mappings with Scalatra Mounts"** shows one instance where you'll need to make a small change to it, but that's it.

As shown in the list of files, an interesting thing about the current Giter8 template for Scalatra is that it uses a `project/build.scala` file rather than a `build.sbt` file. You can find all of Scalatra's dependencies in that file, including the use of tools such as the **Scalate template engine**, **specs2**, **Logback**, and **Jetty**.

See Also

- The **Scalatra website**
- The **Giter8 website**
- **Recipe 18.1, "Creating a Project Directory Structure for SBT"** for how to install Giter8, and use it in other scenarios

15.6. Replacing XML Servlet Mappings with Scalatra Mounts

Problem

You want to add new servlets to your Scalatra application, and need to know how to add them, including defining their URI namespace.

Solution

Scalatra provides a nice way of getting you out of the business of declaring your servlets and servlet mappings in the *web.xml* file. Simply create a boilerplate *web.xml* file like this in the *src/main/webapp/WEB-INF* directory:

```
<?xml version="1.0" encoding="UTF-8"?>
<web-app xmlns="http://java.sun.com/xml/ns/javaee"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://java.sun.com/xml/ns/javaee
http://java.sun.com/xml/ns/javaee/web-app_3_0.xsd"
  version="3.0">
  <listener>
    <listener-class>org.scalatra.servlet.ScalatraListener</listener-class>
  </listener>

  <servlet-mapping>
    <servlet-name>default</servlet-name>
    <url-pattern>/img/*</url-pattern>
    <url-pattern>/css/*</url-pattern>
    <url-pattern>/js/*</url-pattern>
    <url-pattern>/assets/*</url-pattern>
  </servlet-mapping>
</web-app>
```

Next, assuming that you're working with the application created in [Recipe 15.5](#), edit the *src/main/scala/ScalatraBootstrap.scala* file so that it has these contents:

```
import org.scalatra._
import javax.servlet.ServletContext
import com.alvinalexander.app._

class ScalatraBootstrap extends Lifecycle {
  override def init(context: ServletContext) {

    // created by default
    context.mount(new MyScalatraServlet, "/*")

    // new
    context.mount(new StockServlet, "/stocks/*")
    context.mount(new BondServlet, "/bonds/*")
  }
}
```

```
}
}
```

The two new `context.mount` lines shown tell Scalatra that a class named `StockServlet` should handle all URI requests that begin with `/stocks/`, and another class named `BondServlet` should handle all URI requests that begin with `/bonds/`.

Next, create a file named `src/main/scala/com/alvinalexander/app/OtherServlets.scala` to define the `StockServlet` and `BondServlet` classes:

```
package com.alvinalexander.app

import org.scalatra._
import scalate.ScalateSupport

class StockServlet extends MyScalatraWebAppStack {
  get("/") {
    <p>Hello from StockServlet</p>
  }
}

class BondServlet extends MyScalatraWebAppStack {
  get("/") {
    <p>Hello from BondServlet</p>
  }
}
```

Assuming your project is still configured to recompile automatically, when you access the `http://localhost:8080/stocks/` and `http://localhost:8080/bonds/` URLs, you should see the content from your new servlets.

Discussion

Scalatra refers to this configuration process as “mounting” the servlets, and if you’ve used a filesystem technology like NFS, it does indeed feel similar to the process of mounting a remote filesystem.

As a result of the configuration, new methods in the `StockServlet` and `BondServlet` will be available under the `/stocks/` and `/bonds/` URIs. For example, if you define a new method like this in the `StockServlet`:

```
get("/foo") {
  <p>Foo!</p>
}
```

you’ll be able to access this method at the `/stocks/foo` URI, e.g., the `http://localhost:8080/stocks/foo` URL, if you’re running on port 8080 on your local computer.

In the end, this approach provides the same functionality as servlet mappings, but it’s more concise, with the added benefit that you’re working in Scala code instead of XML, and you can generally forget about the `web.xml` file after the initial configuration.

See Also

Scalatra configuration and deployment guide

15.7. Accessing Scalatra Web Service GET Parameters

Problem

When creating a Scalatra web service, you want to be able to handle parameters that are passed into a method using a GET request.

Solution

If you want to let parameters be passed into your Scalatra servlet with a URI that uses traditional ? and & characters to separate data elements, like this:

```
http://localhost:8080/saveName?fname=Alvin&lname=Alexander
```

you can access them through the implicit `params` variable in a `get` method:

```
/**
 * The URL
 * http://localhost:8080/saveName?fname=Alvin&lname=Alexander
 * prints: Some(Alvin), Some(Alexander)
 */
get("/saveName") {
  val firstName = params.get("fname")
  val lastName = params.get("lname")
  <p>{firstName}, {lastName}</p>
}
```

However, Scalatra also lets you use a “named parameters” approach, which can be more convenient, and also documents the parameters your method expects to receive. Using this approach, callers can access a URL like this:

```
http://localhost:8080/hello/Alvin/Alexander
```

You can handle these parameters in a `get` method like this:

```
get("/hello/:fname/:lname") {
  <p>Hello, {params("fname")}, {params("lname")}</p>
}
```

As mentioned, a benefit of this approach is that the method signature documents the expected parameters.

With this approach, you can use wildcard characters for other needs, such as when a client needs to pass in a filename path, where you won’t know the depth of the path beforehand:

```
get("/getFilename/*.") {
  val data = multiParams("splat")
  <p>{data.mkString("[", ", ", "]" )}</p>
}
```

You can understand this method by looking at a specific example. Imagine a GET request to the `http://localhost:8080/getFilename/Users/Al/tmp/file.txt` URL. The comments in the following code show how this URL is handled:

```
/**
 * (1) GET http://localhost:8080/getFilename/Users/Al/tmp/file.txt
 */
get("/getFilename/*.") {

  // (2) creates a Vector(Users/Al/tmp/file, txt)
  val data = multiParams("splat")

  // (3) prints: [Users/Al/tmp/file, txt]
  <pre>{data.mkString("[", ", ", "]" )}</pre>

}
```

This code works because the `multiParams` method with the `splat` argument creates a `Vector` that contains two elements: the strings `Users/Al/tmp/file` and `txt`. Next, the information is printed back to the browser with the `data.mkString` line. In a real-world program, you can put the filename back together by merging `data(0)` and `data(1)`, and then using the filename as needed.

There are more methods for parsing GET request parameters with Scalatra, including additional uses of wildcard characters, and Rails-like pattern matching. See the latest [Scalatra documentation](#) for more information.

15.8. Accessing POST Request Data with Scalatra

Problem

You want to write a Scalatra web service method to handle POST data, such as handling JSON data sent as a POST request.

Solution

To handle a POST request, write a `post` method in your Scalatra servlet, specifying the URI the method should listen at:

```
post("/saveJsonStock") {
  val jsonString = request.body
  // deserialize the JSON ...
}
```

As shown, access the data that's passed to the POST request by calling the `request.body` method.

The Discussion shows an example of how to process JSON data received in a `post` method, and two clients you can use to test a `post` method: a Scala client, and a command-line client that uses the Unix `curl` command.

Discussion

[Recipe 15.3](#) shows how to convert a JSON string into a Scala object using the Lift-JSON library, in a process known as deserialization. In a Scalatra `post` method, you access a JSON string that has been POSTed to your method by calling `request.body`. Once you have that string, deserialize it using the approach shown in [Recipe 15.3](#).

For instance, the `post` method in the following `StockServlet` shows how to convert the JSON string it receives as a POST request and deserialize it into a `Stock` object. The comments in the code explain each step:

```
package com.alvinalexander.app

import org.scalatra._
import scalate.ScalateSupport
import net.liftweb.json._

class StockServlet extends MyScalatraWebAppStack {

  /**
   * Expects an incoming JSON string like this:
   * {"symbol": "GOOG", "price": "600.00"}
   */
  post("/saveJsonStock") {

    // get the POST request data
    val jsonString = request.body

    // needed for Lift-JSON
    implicit val formats = DefaultFormats

    // convert the JSON string to a JValue object
    val jValue = parse(jsonString)

    // deserialize the string into a Stock object
    val stock = jValue.extract[Stock]

    // for debugging
    println(stock)

    // you can send information back to the client
    // in the response header
    response.addHeader("ACK", "GOT IT")
  }
}
```

```

    }

}

// a simple Stock class
class Stock (var symbol: String, var price: Double) {
  override def toString = symbol + ", " + price
}

```

The last step to get this working is to add the Lift-JSON dependency to your project. Assuming that you created your project as an SBT project as shown in [Recipe 15.1](#), add this dependency to the `libraryDependencies` declared in the `project/build.scala` file in your project:

```
"net.liftweb" %% "lift-json" % "2.5+"
```

Test the POST method with Scala code

As shown in the code comments, the `post` method expects a JSON string with this form:

```
{"symbol": "GOOG", "price": 600.00}
```

You can test your `post` method in a variety of ways, including (a) a Scala POST client or (b) a simple shell script. The following `PostTester` object shows how to test the `post` method with a Scala client:

```

import net.liftweb.json._
import net.liftweb.json.Serialization.write
import org.apache.http.client.methods.HttpPost
import org.apache.http.entity.StringEntity
import org.apache.http.impl.client.DefaultHttpClient

object PostTester extends App {

  // create a Stock and convert it to a JSON string
  val stock = new Stock("AAPL", 500.00)
  implicit val formats = DefaultFormats
  val stockAsJsonString = write(stock)

  // add the JSON string as a StringEntity to a POST request
  val post = new HttpPost("http://localhost:8080/stocks/saveJsonStock")
  post.setHeader("Content-type", "application/json")
  post.setEntity(new StringEntity(stockAsJsonString))

  // send the POST request
  val response = (new DefaultHttpClient).execute(post)

  // print the response
  println("--- HEADERS ---")
  response.getAllHeaders.foreach(arg => println(arg))

}

```

```
class Stock (var symbol: String, var price: Double)
```

The code starts by creating a `Stock` object and converting the object to a JSON string using Lift-JSON. It then uses the methods of the Apache HttpClient library to send the JSON string as a POST request: it creates an `HttpPost` object, sets the header content type, then wraps the JSON string as a `StringEntity` object before sending the POST request and waiting for the response.

When this test object is run against the Scalatra `saveJsonStock` method, it results in the following output:

```
--- HEADERS ---
ACK: GOT IT
Content-Type: text/html; charset=UTF-8
Content-Length: 0
Server: Jetty(8.1.8.v20121106)
```

Note that it receives the ACK message that was returned by the Scalatra `post` method. This isn't required, but it gives the client a way to confirm that the data was properly received and processed by the server method (or that it failed).

Test the POST method with a curl command

Another way to test the `post` method is with a Unix shell script. The following `curl` command sets the `Content-type` header, and sends a sample JSON string to the Scalatra `StockServlet` `post` method as a POST request:

```
curl \
  --header "Content-type: application/json" \
  --request POST \
  --data '{"symbol": "GOOG", "price": 600.00}' \
  http://localhost:8080/stocks/saveJsonStock
```

On Unix systems, save this command to a file named *postJson.sh*, and then make it executable:

```
$ chmod +x postJson.sh
```

Then run it to test your Scalatra web service:

```
$ ./postJson.sh
```

You won't see any output from this command, but you should see the correct debugging output printed by the `StockServlet` in its output window. Assuming that you're running your Scalatra web service using SBT, the debug output will appear there.

Notes

Recent versions of Scalatra use the `Json4s` library to deserialize JSON. This library is currently based on Lift-JSON, so the deserialization code will be similar, if not exactly the same. Either library will have to be added as a dependency.

The other important parts about this recipe are:

- Knowing to use the `post` method to handle a `POST` request
- Using `request.body` to get the `POST` data
- Using `response.addHeader("ACK", "GOT IT")` to return a success or failure message to the client (though this is optional)
- Having `POST` request client programs you can use

15.9. Creating a Simple GET Request Client

Problem

You want an HTTP client you can use to make GET request calls.

Solution

There are many potential solutions to this problem. This recipe demonstrates three approaches:

- A simple use of the `scala.io.Source.fromURL` method
- Adding a timeout wrapper around `scala.io.Source.fromURL` to make it more robust
- Using the Apache `HttpClient` library

These solutions are demonstrated in the following sections.

A simple use of `scala.io.Source.fromURL`

If it doesn't matter that your web service client won't time out in a controlled manner, you can use this simple method to download the contents from a URL:

```
/**
 * Returns the text (content) from a URL as a String.
 * Warning: This method does not time out when the service is non-responsive.
 */
def get(url: String) = scala.io.Source.fromURL(url).mkString
```

This GET request method lets you call the given RESTful URL to retrieve its content. You can use it to download web pages, RSS feeds, or any other content using an HTTP GET request.

Under the covers, the `Source.fromURL` method uses classes like `java.net.URL` and `java.io.InputStream`, so this method can throw exceptions that extend from

`java.io.IOException`. As a result, you may want to annotate your method to indicate that:

```
@throws(classOf[java.io.IOException])
def get(url: String) = io.Source.fromURL(url).mkString
```

Setting the timeout while using `scala.io.Source.fromURL`

As mentioned, that simple solution suffers from a significant problem: it doesn't time out if the URL you're calling is unresponsive. If the web service you're calling isn't responding, your code will become unresponsive at this point as well.

Therefore, a better approach is to write a similar method that allows the setting of a timeout value. By using a combination of *java.net* classes and the method `io.Source.fromInputStream`, you can create a more robust method that lets you control both the *connection* and *read* timeout values:

```
/**
 * Returns the text (content) from a REST URL as a String.
 * Inspired by http://matthewkwong.blogspot.com/2009/09/scala-scalaiosource-fromurl-blockshangs.html
 * and http://alvinalexander.com/blog/post/java/how-open-url-read-contents-httpurl-connection-java
 *
 * The `connectTimeout` and `readTimeout` comes from the Java URLConnection
 * class Javadoc.
 * @param url The full URL to connect to.
 * @param connectTimeout Sets a specified timeout value, in milliseconds,
 * to be used when opening a communications link to the resource referenced
 * by this URLConnection. If the timeout expires before the connection can
 * be established, a java.net.SocketTimeoutException
 * is raised. A timeout of zero is interpreted as an infinite timeout.
 * Defaults to 5000 ms.
 * @param readTimeout If the timeout expires before there is data available
 * for read, a java.net.SocketTimeoutException is raised. A timeout of zero
 * is interpreted as an infinite timeout. Defaults to 5000 ms.
 * @param requestMethod Defaults to "GET". (Other methods have not been tested.)
 *
 * @example get("http://www.example.com/getInfo")
 * @example get("http://www.example.com/getInfo", 5000)
 * @example get("http://www.example.com/getInfo", 5000, 5000)
 */
@throws(classOf[java.io.IOException])
@throws(classOf[java.net.SocketTimeoutException])
def get(url: String,
        connectTimeout: Int = 5000,
        readTimeout: Int = 5000,
        requestMethod: String = "GET") = {
  import java.net.{URL, HttpURLConnection}
  val connection = (new URL(url)).openConnection.asInstanceOf[HttpURLConnection]
  connection.setConnectTimeout(connectTimeout)
  connection.setReadTimeout(readTimeout)
```

```

    connection.setRequestMethod(requestMethod)
    val inputStream = connection.getInputStream
    val content = io.Source.fromInputStream(inputStream).mkString
    if (inputStream != null) inputStream.close
    content
}

```

As the Scaladoc shows, this method can be called in a variety of ways, including this:

```

try {
    val content = get("http://localhost:8080/waitForever")
    println(content)
} catch {
    case ioe: java.io.IOException => // handle this
    case ste: java.net.SocketTimeoutException => // handle this
}

```

I haven't tested this method with other request types, such as PUT or DELETE, but I have allowed for this possibility by making the `requestMethod` an optional parameter.

Using the Apache HttpClient

Another approach you can take is to use the Apache HttpClient library. Before I learned about the previous approaches, I wrote a `getRestContent` method using this library like this:

```

import java.io._
import org.apache.http.{HttpEntity, HttpResponse}
import org.apache.http.client._
import org.apache.http.client.methods.HttpGet
import org.apache.http.impl.client.DefaultHttpClient
import scala.collection.mutable.StringBuilder
import scala.xml.XML
import org.apache.http.params.HttpConnectionParams
import org.apache.http.params.HttpParams

/**
 * Returns the text (content) from a REST URL as a String.
 * Returns a blank String if there was a problem.
 * This function will also throw exceptions if there are problems trying
 * to connect to the url.
 *
 * @param url A complete URL, such as "http://foo.com/bar"
 * @param connectionTimeout The connection timeout, in ms.
 * @param socketTimeout The socket timeout, in ms.
 */
def getRestContent(url: String,
                  connectionTimeout: Int,
                  socketTimeout: Int): String = {
    val httpClient = buildHttpClient(connectionTimeout, socketTimeout)
    val httpResponse = httpClient.execute(new HttpGet(url))
    val entity = httpResponse.getEntity
    var content = ""

```

```

    if (entity != null) {
        val inputStream = entity.getContent
        content = io.Source.fromInputStream(inputStream).getLines.mkString
        inputStream.close
    }
    httpClient.getConnectionManager.shutdown
    content
}

private def buildHttpClient(connectionTimeout: Int, socketTimeout: Int):
DefaultHttpClient = {
    val httpClient = new DefaultHttpClient
    val httpParams = httpClient.getParams
    HttpConnectionParams.setConnectionTimeout(httpParams, connectionTimeout)
    HttpConnectionParams.setSoTimeout(httpParams, socketTimeout)
    httpClient.setParams(httpParams)
    httpClient
}

```

This requires significantly more code than the `Source.fromURL` approaches, as well as the `HttpClient` library. If you're already using the Apache `HttpClient` library for other purposes, this is a viable alternative. As shown in Recipes 15.11 and 15.12, the `HttpClient` library definitely has advantages in situations such as working with request headers.

Discussion

There are several other approaches you can take to handle this timeout problem. One is to use the Akka Futures as a wrapper around the `Source.fromURL` method call. See Recipe 13.9, “Simple Concurrency with Futures” for an example of how to use that approach.

Also, new libraries are always being released. A library named **Newman** was released by StackMob as this book was in the production process, and looks promising. The Newman DSL was inspired by the Dispatch library, but uses method names instead of symbols, and appears to be easier to use as a result. It also provides separate methods for the GET, POST, PUT, DELETE, and HEAD request methods.

See Also

- Matthew Kwong's `Source.fromURL` **timeout approach**.
- If you prefer asynchronous programming, you can mix this recipe with Scala Futures, which are demonstrated in Chapter 13. Another option is the Dispatch library. As its **documentation** states, “Dispatch is a library for asynchronous HTTP interaction. It provides a Scala vocabulary for Java's `async-http-client`.”
- **Newman**, from StackMob.

15.10. Sending JSON Data to a POST URL

Problem

You want to send JSON data (or other data) to a POST URL, either from a standalone client, or when using a framework that doesn't provide this type of service.

Solution

Create a JSON string using your favorite JSON library, and then send the data to the POST URL using the Apache HttpClient library. In the following example, the Gson library is used to construct a JSON string, which is then sent to a server using the methods of the HttpClient library:

```
import java.io._
import org.apache.commons._
import org.apache.http._
import org.apache.http.client._
import org.apache.http.client.methods._
import org.apache.http.impl.client._
import java.util._
import org.apache.http.message._
import org.apache.http.entity._
import com.google.gson._

case class Person(firstName: String, lastName: String, age: Int)

object HttpJsonPostTest extends App {

  // create our object as a json string
  val spock = new Person("Leonard", "Nimoy", 82)
  val spockAsJson = new Gson().toJson(spock)

  // add name value pairs to a post object
  val post = new HttpPost("http://localhost:8080/posttest")
  val nameValuePairs = new ArrayList<NameValuePair>()
  nameValuePairs.add(new BasicNameValuePair("JSON", spockAsJson))
  post.setEntity(new UrlEncodedFormEntity(nameValuePairs))

  // send the post request
  val client = new DefaultHttpClient
  val response = client.execute(post)
  println("--- HEADERS ---")
  response.getAllHeaders().foreach(arg => println(arg))

}
```

Discussion

The Gson library is used to construct a JSON string in this code because this is a simple example. For more complex cases, you'll probably want to use the Lift-JSON library, as discussed in the first several recipes of this chapter.

In this example, once you've constructed a JSON string from a Scala object, the Apache `HttpClient` `NameValuePair`, `BasicNameValuePair`, and `UrlEncodedFormEntity` classes are used to set an `Entity` on an `HttpPost` object. In the last lines of the code, the `POST` request is sent using the `client.execute` call, the response is received, and the response headers are printed (though this isn't necessary).

See Also

- [Recipe 15.1, “Creating a JSON String from a Scala Object”](#) and [Recipe 15.2, “Creating a JSON String from Classes That Have Collections”](#).
- The [Lift-JSON library](#).
- The [Google Gson library](#).
- [Dispatch](#) is a “library for asynchronous HTTP interaction.”

15.11. Getting URL Headers

Problem

You need to access the HTTP response headers after making an HTTP request.

Solution

Use the Apache `HttpClient` library, and get the headers from the `HttpResponse` object after making a request:

```
import org.apache.http.client.methods.HttpGet
import org.apache.http.impl.client.DefaultHttpClient

object FetchUrlHeaders extends App {

  val get = new HttpGet("http://alvinalexander.com/")
  val client = new DefaultHttpClient
  val response = client.execute(get)
  response.getAllHeaders().foreach(header => println(header))

}
```

Running that program prints the following header output:

```
Server: nginx/1.0.10
Date: Sun, 15 Jul 2012 19:10:19 GMT
Content-Type: text/html; charset=utf-8
Connection: keep-alive
Keep-Alive: timeout=20
Content-Length: 28862
Cache-Control: no-store, no-cache, must-revalidate, post-check=0, pre-check=0
Expires: Sun, 19 Nov 1978 05:00:00 GMT
Vary: Accept-Encoding
```

Discussion

When I worked with a Single Sign-On (SSO) system named OpenSSO from Sun (now known as [OpenAM](#)), much of the work in the sign-on process involved setting and reading header information. The `HttpClient` library greatly simplifies this process.

See Also

- [Apache HttpClient library](#).
- You may also be able to use the [Dispatch library](#) for this purpose.

15.12. Setting URL Headers When Sending a Request

Problem

You need to set URL headers when making an HTTP request.

Solution

Use the Apache `HttpClient` library to set the headers before making the request, as shown in this example:

```
import org.apache.http.client.methods.HttpGet
import org.apache.http.impl.client.DefaultHttpClient

object SetUrlHeaders extends App {

  val url = "http://localhost:9001/baz"
  val httpGet = new HttpGet(url)

  // set the desired header values
  httpGet.setHeader("KEY1", "VALUE1")
  httpGet.setHeader("KEY2", "VALUE2")

  // execute the request
  val client = new DefaultHttpClient
  client.execute(httpGet)
```

```
        client.getConnectionManager().shutdown();
    }
}
```

Discussion

If you don't have a web server to test against, you can use a tool like [HttpTea](#) to see the results of running this program. HttpTea helps to simulate a server in a test environment.

Start HttpTea at the command line to listen on port 9001 like this:

```
$ java -jar HttpTea.jar -l 9001
```

Now when you run your client program—such as the program shown in the Solution—you should see the following output from HttpTea, including the headers that were set:

```
Client>>>
GET /baz HTTP/1.1
KEY1: VALUE1
KEY2: VALUE2
Host: localhost:9001
Connection: Keep-Alive
User-Agent: Apache-HttpClient/4.1.3 (java 1.5)
```

See Also

- [HttpTea](#).
- [Apache HttpClient library](#).
- You may also be able to use the [Dispatch library](#) for this purpose.

15.13. Creating a GET Request Web Service with the Play Framework

Problem

You want to create a GET request web service using the Play Framework, such as returning a JSON string when the web service URI is accessed.

Solution

When working with RESTful web services, you'll typically be converting between one or more model objects and their JSON representation.

To demonstrate how a GET request might be used to return the JSON representation of an object, create a new Play project with the `play new` command:

```
$ play new WebServiceDemo
```

Respond to the prompts to create a new Scala application, and then move into the *WebServiceDemo* directory that's created.

Next, assume that you want to create a web service to return an instance of a *Stock* when a client makes a GET request at the */getStock* URI. To do this, first add this line to your *conf/routes* file:

```
GET    /getStock    controllers.Application.getStock
```

Next, create a method named *getStock* in the default *Application* controller (*apps/controllers/Application.scala*), and have it return a JSON representation of a *Stock* object:

```
package controllers

import play.api._
import play.api.mvc._
import play.api.libs.json._
import models.Stock

object Application extends Controller {

  def index = Action {
    Ok(views.html.index("Your new application is ready."))
  }

  def getStock = Action {
    val stock = Stock("GOOG", 650.0)
    Ok(Json.toJson(stock))
  }

}
```

That code uses the Play *Json.toJson* method. Although the code looks like you can create *Stock* as a simple case class, attempting to use only a case class will result in this error when you access the */getStock* URI:

```
No Json deserializer found for type models.Stock. Try to implement an implicit Writes
or Format for this type.
```

To get this controller code to work, you need to create an instance of a *Format* object to convert between the *Stock* model object and its JSON representation. To do this, create a model file named *Stock.scala* in the *app/models* directory of your project. (Create the directory if it doesn't exist.)

In that file, define the *Stock* case class, and then implement a *play.api.libs.json.Format* object. In that object, define a *reads* method to convert from a JSON string to a *Stock* object and a *writes* method to convert from a *Stock* object to a JSON string:


```

package models

case class Stock(symbol: String, price: Double)

object Stock {

  import play.api.libs.json._

  implicit object StockFormat extends Format[Stock] {

    // convert from JSON string to a Stock object (de-serializing from JSON)
    def reads(json: JsValue): JsResult[Stock] = {
      val symbol = (json \ "symbol").as[String]
      val price = (json \ "price").as[Double]
      JsSuccess(Stock(symbol, price))
    }

    // convert from Stock object to JSON (serializing to JSON)
    def writes(s: Stock): JsValue = {
      // JsObject requires Seq[(String, play.api.libs.json.JsValue)]
      val stockAsList = Seq("symbol" -> JsString(s.symbol),
                           "price" -> JsNumber(s.price))
      JsObject(stockAsList)
    }
  }
}

```

The comments in that code help to explain how the reads and writes methods work.

With this code in place, you can now access the `getStock` web service. If you haven't already done so, start the Play console from within the root directory of your project, then issue the run command:

```

$ play

[WebServiceDemo] $ run 8080

```

Play runs on port 9000 by default, but this collides with other services on my system, so I run it on port 8080, as shown. Assuming that you're running on port 8080, access the `http://localhost:8080/getStock` URL from a web browser. You should see this result in the browser:

```
{"symbol": "GOOG", "price": 650.0}
```

Discussion

When converting from a `Stock` object to its JSON representation, the `writes` method of your `Format` object is implicitly used in this line of code:

```
Json.toJson(stock)
```

Although there are other approaches to converting between objects and their JSON representation, implementing the `reads` and `writes` methods of a `Format` object provides a straightforward means for this serialization and deserialization process.

See Also

The `Play json` package object

15.14. POSTing JSON Data to a Play Framework Web Service

Problem

You want to create a web service using the Play Framework that lets users send JSON data to the service using the POST request method.

Solution

Follow the steps from the previous recipe to create a new Play project, controller, and model.

Whereas the previous recipe used the `writes` method of the `Format` object in the model, this recipe uses the `reads` method. When JSON data is received in a POST request, the `reads` method is used to convert from the JSON string that's received to a `Stock` object. Here's the code for the `reads` method:

```
def reads(json: JsValue): JsResult[Stock] = {  
  val symbol = (json \ "symbol").as[String]  
  val price = (json \ "price").as[Double]  
  JsSuccess(Stock(symbol, price))  
}
```

This method creates a `Stock` object from the JSON value it's given. (The complete code for the model object is shown in the previous recipe.)

With this method added to the model, create a `saveStock` method in the `Application` controller:

```
import play.api._  
import play.api.mvc._  
  
object Application extends Controller {  
  
  import play.api.libs.json.Json  
  
  def saveStock = Action { request =>  
    val json = request.body.asJson.get
```

```

    val stock = json.as[Stock]
    println(stock)
    Ok
  }
}

```

The `saveStock` method gets the JSON data sent to it from the `request` object, and then converts it with the `json.as` method. The `println` statement in the method is used for debugging purposes, and prints to the Play command line (the Play console).

Finally, add a route that binds a POST request to the desired URI and the `saveStock` method in the `Application` controller by adding this line to the `conf/routes` file:

```
POST    /saveStock    controllers.Application.saveStock
```

If you haven't already done so, start the Play console from within the root directory of your project, and issue the `run` command:

```

$ play

[WebServicesDemo] $ run 8080

```

With the Play server running, use the following Unix `curl` command to POST a sample JSON string to your `saveStock` web service:

```

curl \
  --header "Content-type: application/json" \
  --request POST \
  --data '{"symbol": "GOOG", "price": 900.00}' \
  http://localhost:8080/saveStock

```

If everything works properly, you should see this output in your Play console window:

```
STOCK: Stock(GOOG,900.0)
```

Discussion

A few notes about the code:

- The `request` object is a `play.api.mvc.AnyContent` object.
- The `request.body` is also a `play.api.mvc.AnyContent` object.
- The `request.body.asJson` returns an instance of the following: `Option[play.api.libs.json.JsValue]`.
- `request.body.asJson.get` returns a `JsValue`.

In a real-world web service, once you've converted the JSON string to an object, you can do anything else you need to do with it, such as saving it to a database.

See Also

- The [Play json](#) package object
- The [Play Request](#) trait

Databases and Persistence

Introduction

With Scala, you can interact with traditional relational databases using their JDBC drivers, just like you do in Java. As an example of this, the first recipe in this chapter demonstrates how to connect to a MySQL database using the “plain old JDBC” approach.

In the real world, once applications grow in size, few people use plain old JDBC to work with databases. Typically on those projects you use a library, such as the **Spring Framework**, to make development easier and handle issues like connection pooling. Therefore, this chapter also demonstrates the few changes you’ll need to make to use the Spring JDBC library with Scala. As an added benefit, by showing the changes needed to instantiate a bean from a Spring application context file, this recipe will help you use other Spring libraries with Scala as well. You can use other technologies with Scala, such as the Java Persistence API (JPA) and Hibernate, with just a few changes.

The Scala community is also developing new approaches to database development. The **Squeryl** and **Slick** libraries both take “type-safe” approaches to writing database code. The Squeryl documentation states that it’s a “Scala ORM and DSL.” In a manner similar to Hibernate, Squeryl lets you write database code like this:

```
// insert
val bill = people.insert(new Person("Bill"))
val candy = people.insert(new Person("Candy"))

// update
stock.price = 500.00
stocks.update(stock)
```

With Squeryl’s DSL, you can also write statements like this:

```
update(stocks)(s =>
  where(s.symbol === "AAPL")
```

```
    set(s.price := 500.00)
  )
```

Slick isn't an object-relational mapping (ORM) tool, but with its type-safe approach, it lets you write database access code almost like you're working with a collection. This approach is demonstrated in the last recipe in this chapter.

When you get to “big data” projects, it's nice to know that Scala works there as well. There are several Scala drivers available for the [MongoDB](#) database, including [Casbah](#) and [ReactiveMongo](#). The recipes in this chapter demonstrate how to use the Casbah driver to insert, update, read, and delete objects in a MongoDB collection with Scala.

If you want to use Scala to work with Hadoop, Twitter has created a project named [Scalding](#) that “makes it easy to specify Hadoop MapReduce jobs.” Scalding is analogous to the [Apache Pig project](#), but is tightly integrated with Scala. Scalding and Hadoop are not covered in this chapter, but the [Scalding source code tutorials](#) can help you quickly get up and running with Scalding.

16.1. Connecting to MySQL with JDBC

Problem

You want to connect to a MySQL database (or any other database with a JDBC driver) from a Scala application using “plain old JDBC.”

Solution

Use JDBC just like you would in a Java application. Download the MySQL JDBC driver, and then access your database with code like this:

```
package tests

import java.sql.{Connection, DriverManager}

object ScalaJdbcConnectSelect extends App {

  // connect to the database named "mysql" on port 8889 of localhost
  val url = "jdbc:mysql://localhost:8889/mysql"
  val driver = "com.mysql.jdbc.Driver"
  val username = "root"
  val password = "root"
  var connection:Connection = _
```

```

try {
  Class.forName(driver)
  connection = DriverManager.getConnection(url, username, password)
  val statement = connection.createStatement
  val rs = statement.executeQuery("SELECT host, user FROM user")
  while (rs.next) {
    val host = rs.getString("host")
    val user = rs.getString("user")
    println("host = %s, user = %s".format(host,user))
  }
} catch {
  case e: Exception => e.printStackTrace
}
connection.close
}

```

That code shows how to query a database table named `user` in a database named `mysql`. That database name and table name are standard in any MySQL installation.

As shown in the example, the format of the MySQL JDBC URL string is:

```
jdbc:mysql://HOST:PORT/DATABASE
```

In this code I have MySQL running on port 8889 on my computer because it's the default port when using **MAMP**, a tool that makes it easy to run MySQL, Apache, and PHP on Mac OS X systems. If you have MySQL running on its standard port (3306), just drop the port off the URL string.

Discussion

The easiest way to run this example is to use the **Simple Build Tool** (SBT). To do this, create an SBT directory structure as described in [Recipe 18.1, “Creating a Project Directory Structure for SBT”](#), then add the MySQL JDBC dependency to the *build.sbt* file:

```
libraryDependencies += "mysql" % "mysql-connector-java" % "5.1.24"
```

Copy and paste the code shown in this recipe into a file named *Test1.scala* in the root directory of your project, and then run the program:

```
$ sbt run
```

You should see some output like this:

```
host = localhost, user =
host = localhost, user = fred
```

That output will vary depending on the users actually defined in your MySQL database.

This recipe works well for small applications where you want one connection to a database, and you don't mind running simple JDBC SQL queries using the `Statement`, `PreparedStatement`, and `ResultSet` classes. For larger applications, you'll want to use

a tool that gives you connection pooling capabilities, and possibly DSL or ORM capabilities to simplify your SQL queries.

If you're using a different relational database, the approach is the same as long as the database provides a JDBC driver. For instance, to use **PostgreSQL**, just use the **PostgreSQL JDBC driver** and this information to create a connection:

```
Class.forName("org.postgresql.Driver")
val url = "jdbc:postgresql://HOST/DATABASE"
val conn = DriverManager.getConnection(url, "username", "password")
```

Of course your database tables will be different, but the process of connecting to the database is the same.

See Also

- The **MySQL JDBC driver**.
- **MAMP**.
- The **Simple Build Tool (SBT)**.
- **Recipe 18.1, “Creating a Project Directory Structure for SBT”** shows how to create an SBT directory structure.
- If you're new to MySQL and JDBC, I wrote a series of **MySQL JDBC tutorials** that can help you get started.

16.2. Connecting to a Database with the Spring Framework

Problem

You want to connect to a database using the **Spring Framework**. This gives you a nice way to add connection pooling and other capabilities to your SQL code.

Solution

Use the same Spring Framework configuration you've used in Java applications, but convert your Java source code to Scala. The biggest changes involve the differences in class casting between Java and Scala, and conversions between Java and Scala collections.

Discussion

To demonstrate this, create a basic Spring JDBC example. Start by creating a simple SBT project directory structure as demonstrated in [Recipe 18.1, “Creating a Project Directory Structure for SBT”](#).

Once the SBT directory structure is created, place this Spring *applicationContext.xml* file in the *src/main/resources* directory:

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE beans PUBLIC "-//SPRING//DTD BEAN//EN"
"http://www.springframework.org/dtd/spring-beans.dtd">

<beans>

  <bean id="testDao" class="springtests.TestDao">
    <property name="dataSource" ref="basicDataSource"/>
  </bean>

  <bean id="basicDataSource" class="org.apache.commons.dbcp.BasicDataSource">
    <property name="driverClassName" value="com.mysql.jdbc.Driver" />
    <property name="url" value="jdbc:mysql://localhost/mysql" />
    <property name="username" value="root" />
    <property name="password" value="root" />
    <property name="initialSize" value="1" />
    <property name="maxActive" value="5" />
  </bean>

</beans>
```

This file declares that you'll have a class named *TestDao* in a package named *springtests*. This bean declaration will be used in the *Main* object, which you'll create shortly.

This file also lets you connect to a MySQL database named *mysql*, on the default port (3306) of the *localhost* server, with the username and password both set to *root*. The *initialSize* and *maxActive* settings let you control the database connection pool settings. Change those properties as needed for your system.

You'll need to add a number of dependencies to your *build.sbt* file to get Spring to work:

```
name := "MySQLTest1"

version := "1.0"

scalaVersion := "2.10.1"

libraryDependencies += Seq(
  "mysql" % "mysql-connector-java" % "5.1.+",
  "commons-dbcp" % "commons-dbcp" % "1.4",
  "org.springframework" % "spring-core" % "3.1+",
```

```

    "org.springframework" % "spring-beans" % "3.1+",
    "org.springframework" % "spring-jdbc" % "3.1+",
    "org.springframework" % "spring-tx" % "3.1+"
  )

```

Alternatively, you can manually download the JAR files that are needed and put them in your *lib* directory.

Next, create a file named *Main.scala* in your root SBT directory with the following contents:

```

package springtests

import org.springframework.context.support.ClassPathXmlApplicationContext

object Main extends App {

  // read the application context file
  val ctx = new ClassPathXmlApplicationContext("applicationContext.xml")

  // get a testDao instance
  val testDao = ctx.getBean("testDao").asInstanceOf[TestDao]
  val numUsers = testDao.getNumUsers
  println("You have this many users: " + numUsers)

}

```

Note how an instance of the *TestDao* is instantiated in this object. This code is similar to Java, except for the way class casting is handled. As shown, Scala uses the *asInstanceOf* method to declare that the *testDao* bean is of the type *TestDao*.

Next, create another file in the root directory of the project named *TestDao.scala* with these contents:

```

package springtests

import org.springframework.jdbc.core.simple._

class TestDao
extends SimpleJdbcDaoSupport {

  def getNumUsers: Int = {
    val query = "select count(1) from user"
    return getJdbcTemplate.queryForInt(query)
  }

}

```

Now run the project with the `sbt run` command. You should see some simple output, including the number of records in your MySQL user database table.

Although this example was created to demonstrate how to use the Spring JDBC support with Scala, you can use the steps in this recipe to use other Spring libraries in your Scala applications.

See Also

- The [Spring Framework](#).
- [MAMP](#).
- [Recipe 18.1, “Creating a Project Directory Structure for SBT”](#).
- A project named “Spring Scala” is being created to make it easier to use [Spring in Scala applications](#).

16.3. Connecting to MongoDB and Inserting Data

Problem

You want to use the MongoDB database with a Scala application, and want to learn how to connect to it, and insert and retrieve data.

Solution

If you don’t already have a [MongoDB](#) installation, download and install the MongoDB software per the instructions on its website. (It’s simple to install.) Once it’s running, use the [Casbah driver](#) with your Scala application to interact with MongoDB.

In development, I start my test instance of MongoDB from its installation directory with this command:

```
$ bin/mongod -vvvv --dbpath /Users/AL/data/mongodatabases
```

This starts the MongoDB server in a verbose mode, using the directory shown for its databases. After a lot of output, the last few lines from the `mongod` command look like this:

```
Sun Sep 16 14:27:34 [websvr] admin web console waiting for connections
on port 28017
Sun Sep 16 14:27:34 [initandlisten] waiting for connections on port 27017
```

To demonstrate Casbah, build a small application. First, create a simple SBT project directory structure, as demonstrated in [Recipe 18.1, “Creating a Project Directory Structure for SBT”](#).



You can follow along by cloning my [GitHub project](#).

Second, create your *build.sbt* file, specifically including the Casbah driver dependency:

```
name := "MongoDBDemo1"

version := "1.0"

scalaVersion := "2.10.0"

libraryDependencies += Seq(
  "org.mongodb" %% "casbah" % "2.6.0",
  "org.slf4j" % "slf4j-simple" % "1.6.4"
)

scalacOptions += "-deprecation"
```

The SLF4J library shown isn't necessary for a simple example, but including it gets rid of a few warning messages.

Next, put the following code in a file named *MongoFactory.scala* in the root directory of your SBT project:

```
import com.mongodb.casbah.MongoCollection
import com.mongodb.casbah.MongoConnection

object MongoFactory {

  private val SERVER = "localhost"
  private val PORT   = 27017
  private val DATABASE = "portfolio"
  private val COLLECTION = "stocks"

  val connection = MongoConnection(SERVER)
  val collection = connection(DATABASE)(COLLECTION)

}
```

This object helps to simplify the interactions with a MongoDB database. You won't need all of its functionality for this recipe, but it will be used completely in other recipes. If your MongoDB instance is running on the default port on localhost, those settings will work fine. If you already have a database named `portfolio`, be sure to use a different name.

Next, put the following code in a file named *Common.scala*, also in the root directory of your SBT project:

```

import com.mongodb.casbah.Imports._

case class Stock (symbol: String, price: Double)

object Common {

  /**
   * Convert a Stock object into a BSON format that MongoDB can store.
   */
  def buildMongoDBObject(stock: Stock): MongoDBObject = {
    val builder = MongoDBObject.newBuilder
    builder += "symbol" -> stock.symbol
    builder += "price" -> stock.price
    builder.result
  }
}

```

That code includes a simple case class to represent a `Stock`, and the `buildMongoDBObject` method in the `Common` object does the work of converting a `Stock` into a `MongoDBObject` that can be stored in a MongoDB database. The method converts the fields in the `Stock` object into key/value pairs that correspond to the MongoDB “document” paradigm. The `MongoDBObject` from the Casbah driver simplifies the conversion process.

With this code in place, it’s time to create a simple test program to insert several `Stock` instances into the database. Put the following code into a file named *Insert.scala* in the root directory of your SBT project:

```

import com.mongodb.casbah.Imports._
import Common._

object Insert extends App {

  // create some Stock instances
  val apple = Stock("AAPL", 600)
  val google = Stock("GOOG", 650)
  val netflix = Stock("NFLX", 60)

  // save them to the mongodb database
  saveStock(apple)
  saveStock(google)
  saveStock(netflix)

  // our 'save' method
  def saveStock(stock: Stock) {
    val mongoObj = buildMongoDBObject(stock)
    MongoFactory.collection.save(mongoObj)
  }
}

```

The interesting part of this code is the `saveStock` method. It does the following work:

- It takes a `Stock` object as an input parameter.
- It converts the `Stock` object to a `MongoDBObject` with the `buildMongoDBObject` method.
- It saves the `mongoObj` object to the database collection with the `save` method of the collection instance. The collection is an instance of `MongoCollection`, which is obtained from the `MongoFactory`.

With everything in place, run this object with `sbt run`, and it will quietly insert the data into the collection.

Discussion

In [Recipe 16.5, “Searching a MongoDB Collection”](#), you’ll see how to search a MongoDB collection using Scala and Casbah, but for the time being, if you open up the MongoDB command-line client and switch to the `portfolio` database, you can see the new documents in the `stocks` collection.

To do this, move to your MongoDB installation *bin* directory, start the `mongo` command-line client, move to the `portfolio` database, and list all the documents in the `stocks` collection, using these commands:

```
$ mongo
> use portfolio
> db.stocks.find()
{"_id" : ObjectId("5023fad43004f32afda0b550"), "symbol" : "AAPL", "price" : 600 }
{"_id" : ObjectId("5023fad43004f32afda0b551"), "symbol" : "GOOG", "price" : 650 }
{"_id" : ObjectId("5023fad43004f32afda0b552"), "symbol" : "NFLX", "price" : 60 }
```

This shows the three objects the `Insert` application inserted. You can remove those objects with the following command if you’d like to modify and run the program again:

```
> db.stocks.remove()
```

To help you work with MongoDB, I’ve created a Scala + MongoDB + Casbah example project on [GitHub](#) that includes the source code shown in this recipe, as well as additional code from the `Find`, `Update`, and `Delete` recipes in this chapter.

See Also

- [MongoDB](#)
- [Casbah](#)
- The [MongoCollection](#) API

16.4. Inserting Documents into MongoDB with insert, save, or +=

Problem

You want to save documents to a MongoDB collection from a Scala application.

Solution

Use the `insert`, `save`, or `+=` methods of the Casbah `MongoCollection` class.

In order to save a document to your MongoDB collection, you can use the `MongoCollection` `insert` method:

```
collection.insert(buildMongoDBObject(apple))
collection.insert(buildMongoDBObject(google))
collection.insert(buildMongoDBObject(netflix))
```

You can also use the `save` method:

```
collection.save(buildMongoDBObject(apple))
collection.save(buildMongoDBObject(google))
collection.save(buildMongoDBObject(netflix))
```

And you can also use the `+=` method:

```
collection += buildMongoDBObject(apple)
collection += buildMongoDBObject(google)
collection += buildMongoDBObject(netflix)
collection += buildMongoDBObject(amazon)
```

The intention of the `insert` and `save` methods is obvious; you're inserting/saving data to your MongoDB collection. The third approach is a little different; it looks like what you're doing is adding an object to a collection. In fact, you're saving your object to the database collection with each `+=` call.

Here's what the `+=` approach looks like in a complete program:

```
import com.mongodb.casbah.Imports._
import Common._
```

```
object Insert2 extends App {

    val collection = MongoFactory.collection

    // create some Stock instances
    val apple = Stock("AAPL", 600)
    val google = Stock("GOOG", 650)
    val netflix = Stock("NFLX", 60)
    val amazon = Stock("AMZN", 220)

    // add them to the collection (+= does the save)
    collection += buildMongoDBObject(apple)
    collection += buildMongoDBObject(google)
    collection += buildMongoDBObject(netflix)
    collection += buildMongoDBObject(amazon)

}
```

To use the insert or save methods, simply replace the += lines with their equivalent lines.

Discussion

If you'd like to experiment with this code, just add it to the SBT project that you started in [Recipe 16.3](#). The buildMongoDBObject method in the Common class of that recipe converts a Scala object to a MongoDBObject that can be saved to the database using save, insert, or +=.

When choosing between save, insert, or +=, there's obviously a big difference in style between += and the other methods. The save and insert methods accept a variety of different parameters and both return a WriteResult, so you have a number of options to choose from.

You'll encounter the WriteResult and WriteConcern classes while working with the Casbah driver. These classes come from the MongoDB Java driver, which Casbah wraps. WriteResult lets you access the results of the previous write, and has methods like getField, getError, and getLastError.

WriteConcern provides options to let you control the write behavior, including behavior about network errors, slaves, timeouts, and forcing fsync to disk.

See Also

- The [WriteResult Javadoc](#)
- The [WriteConcern Javadoc](#)

16.5. Searching a MongoDB Collection

Problem

You want to find objects in your MongoDB collection using Scala and the Casbah driver.

Solution

Use the `find*` methods of the `MongoCollection` class to get the elements you want, specifically the `find` and `findOne` methods.

Assuming that you have everything set up as shown in [Recipe 16.3](#), the following code demonstrates these techniques:

- How to find all the documents in a collection
- How to find one document that matches your search criteria
- How to find all documents that match your search criteria
- How to limit the number of results returned by a `find` query

Here's the code:

```
import com.mongodb.casbah.Imports._

object Find extends App {

  val collection = MongoFactory.collection

  // (1) find all stocks with find()
  // -----
  println("\n__ all stocks __")
  val stocks = collection.find
  stocks.foreach(println)

  // (2) search for an individual stock
  // -----
  println("\n__ .findOne(query) __")
  val query = MongoDBObject("symbol" -> "GOOG")
  val result = collection.findOne(query) // Some
  val stock = convertDBObjectToStock(result.get) // convert it to a Stock
  println(stock)

  // (3) find all stocks that meet a search criteria
  // -----
  println("\n__ price $gt 500 __")
  stocks = collection.find("price" $gt 500)
  stocks.foreach(println)

  // (4) find all stocks that match a search pattern
```

```
// -----
println("\n___ stocks that begin with 'A' ___")
stocks = collection.find(MongoDBObject("symbol" -> "A.*".r))
stocks.foreach(println)

// (5) find.limit(2)
// -----
println("\n___ find.limit(2) ___")
stocks = collection.find.limit(2)
stocks.foreach(println)

// warning: don't use the 'get' method in real-world code
def convertDBObjectToStock(obj: MongoDBObject): Stock = {
  val symbol = obj.getAs[String]("symbol").get
  val price = obj.getAs[Double]("price").get
  Stock(symbol, price)
}

}
```

Save that code to a file named *Find.scala* in the root directory of your SBT project, and then run the object with SBT:

```
$ sbt run
```

If you've been working through the MongoDB recipes in this chapter, or you cloned [my Scala + Casbah + MongoDB project from GitHub](#), you may have multiple main methods in your project. If so, SBT detects those main methods and asks which one you want to run. To run the Find object, select it from the list SBT displays:

```
Multiple main classes detected, select one to run:
```

```
[1] Find
[2] Insert
[3] Insert2
```

```
Enter number: 1
```

Running the Find object after populating the database in the earlier recipes results in the following output:

```
___ all stocks ___
{ "_id" : { "$oid" : "502683283004b3802ec47df2" } , "symbol" : "AAPL" ,
  "price" : 600.0}
{ "_id" : { "$oid" : "502683283004b3802ec47df3" } , "symbol" : "GOOG" ,
  "price" : 650.0}
{ "_id" : { "$oid" : "502683283004b3802ec47df4" } , "symbol" : "NFLX" ,
  "price" : 60.0}
{ "_id" : { "$oid" : "502683283004b3802ec47df5" } , "symbol" : "AMZN" ,
  "price" : 220.0}

___ .findOne(query) ___
Stock(GOOG,650.0)
```

```

__ price $gt 500 __
{ "_id" : { "$oid" : "502683283004b3802ec47df2" } , "symbol" : "AAPL" ,
  "price" : 600.0}
{ "_id" : { "$oid" : "502683283004b3802ec47df3" } , "symbol" : "GOOG" ,
  "price" : 650.0}

__ stocks that begin with 'A' __
{ "_id" : { "$oid" : "502683283004b3802ec47df2" } , "symbol" : "AAPL" ,
  "price" : 600.0}
{ "_id" : { "$oid" : "502683283004b3802ec47df5" } , "symbol" : "AMZN" ,
  "price" : 220.0}

__ find.limit(2) __
{ "_id" : { "$oid" : "502683283004b3802ec47df2" } , "symbol" : "AAPL" ,
  "price" : 600.0}
{ "_id" : { "$oid" : "502683283004b3802ec47df3" } , "symbol" : "GOOG" ,
  "price" : 650.0}

```

Discussion

In the first query, the `find` method returns all documents from the specified collection. This method returns a `MongoCursor`, and the code iterates over the results using that cursor.

In the second query, the `findOne` method is used to find one stock that matches the search query. The query is built by creating a `MongoDBObject` with the desired attributes. In this example, that's a stock whose symbol is `GOOG`. The `findOne` method is called to get the result, and it returns an instance of `Some[MongoDBObject]`.

In this example, `result.get` is called on the next line, but in the real world, it's a better practice to use a `for` loop or a match expression:

```

collection.findOne(query) match {
  case Some(Stock) =>
    // convert it to a Stock
    println(convertDBObjectToStock(result.get))
  case None =>
    println("Got something else")
}

```

Of course, how you implement that will vary depending on your needs.

The `convertDBObjectToStock` method does the reverse of the `buildMongoDBObject` method shown in the earlier recipes, and converts a `MongoDBObject` to a `Stock` instance.

The third query shows how to search for all stocks whose price is greater than 500:

```
stocks = collection.find("price" $gt 500)
```

This again returns a `MongoCursor`, and all matches are printed.

Casbah includes other methods besides `$gt`, such as `$gte`, `$lt`, and `$lte`. You can use multiple operators against one field like this:

```
"price" $gt 50 $lte 100
```

You can also query against multiple fields by joining tuples:

```
val query: DBObject = ("price" $gt 50 $lte 100) ++ ("priceToBook" $gt 1)
```

See the [Casbah documentation](#) for more examples of creating Casbah-style queries.

In the fourth query, a simple regular expression pattern is used to search for all stocks whose symbol begins with the letter A:

```
stocks = collection.find(MongoDBObject("symbol" -> "A.*".r))
```

Notice that the `r` method is called on a `String` to create the query. This converts the `String` to a `Regex`, as demonstrated in the REPL:

```
scala> "A.*".r  
res0: scala.util.matching.Regex = A.*
```

The fifth query demonstrates how to use the `limit` method to limit the number of results that are returned:

```
stocks = collection.find.limit(2)
```

Because MongoDB is typically used to store a *lot* of data, you'll want to use `limit` to control the amount of data you get back from a query.

The `MongoCollection` class also has a `findById` method that you can use when you know the ID of your object. Additionally, there are `findAndModify` and `findAndRemove` methods, which are discussed in other recipes in this chapter.

See Also

- [Casbah documentation](#)
- The [MongoCollection class](#)
- The [MongoDB tutorial](#)

16.6. Updating Documents in a MongoDB Collection

Problem

You want to update one or more documents in a MongoDB collection.

Solution

Use either the `findAndModify` or `update` methods from the `Casbah MongoClient` class, as shown in this example:

```
import com.mongodb.casbah.Imports._
import Common._

object Update extends App {

  val collection = MongoFactory.collection

  // findAndModify
  // -----

  // create a new Stock object
  val google = Stock("GOOG", 500)
  // search for an existing document with this symbol
  var query = MongoDBObject("symbol" -> "GOOG")
  // replace the old document with one based on the 'google' object
  val res1 = collection.findAndModify(query, buildMongoDBObject(google))
  println("findAndModify: " + res1)

  // update
  // -----

  // create a new Stock
  var apple = Stock("AAPL", 1000)
  // search for a document with this symbol
  query = MongoDBObject("symbol" -> "AAPL")
  // replace the old document with the 'apple' instance
  val res2 = collection.update(query, buildMongoDBObject(apple))
  println("update: " + res2)

}
```

In both cases, you build a document object to replace the existing document in the database, and then create a query object, which lets you find what you want to replace. Then you call either `findAndModify` or `update` to perform the update.

For instance, in the `findAndModify` example, a new `Stock` instance named `google` is created, and it's used to replace the old document in the database whose symbol is `GOOG`. The `buildMongoDBObject` method is used to convert the `google` instance into a MongoDB document before the `update` method is called.

The difference between the two methods can be seen in the output:

```
findAndModify: Some({ "_id" : { "$oid" : "502683283004b3802ec47df3" } ,
  "symbol" : "GOOG" , "price" : 500.0})
update: N/A
```

Whereas the `findAndModify` method returns the old document (the document that was replaced), the `update` method returns a `WriteResult` instance.

If you've been following along with the MongoDB recipes in this chapter, save that file as *Update.scala* in the root directory of your project, and run it with `sbt run`.

16.7. Accessing the MongoDB Document ID Field

Problem

You want to get the ID field for a document you've inserted into a MongoDB collection.

Solution

Perform a query to get the document you want, and then call `get("_id")` on the resulting `MongoDBObject`, like this:

```
basicDBObject.get("_id")
```

The following example shows how to get the ID field from a `DBObject` after inserting the object into the database. I first create a `Stock` as usual, convert the `Stock` to a `MongoDBObject`, perform the insert, and then get the ID value, which is added to the `MongoDBObject` after the insert operation is performed:

```
import com.mongodb.casbah.Imports._
import Common._

object InsertAndGetId extends App {

  val coll = MongoFactory.collection

  // get the _id field after an insert
  val amazon = Stock("AMZN", 220)
  val amazonMongoObject = buildMongoDBObject(amazon)
  coll.insert(amazonMongoObject)
  println("ID: " + amazonMongoObject.get("_id"))

}
```

If you just need to get the ID field from a `MongoDBObject` after performing a query, the following complete example shows how to do that with a match expression:

```
import com.mongodb.casbah.Imports._

object GetId extends App {

  val collection = MongoFactory.collection

  val query = MongoDBObject("symbol" -> "GOOG")
  collection.findOne(query) match {
```

```

    case Some(result) => println("ID: " + result.get("_id"))
    case None => println("Stock not found")
  }
}

```

A match expression is used in this example because the `findOne(query)` will return `None` if no matching documents are found in the collection. You can also use the usual `getOrElse` and `foreach` techniques to work with an `Option`.

If you've been following along with the MongoDB recipes in this chapter, save those files with the names *InsertAndGetId.scala* and *GetId.scala* in the root directory of your project, and run them with `sbt run`.

See Also

Recipe 20.6, “Using the Option/Some/None Pattern” for many examples of working with methods that return an `Option`

16.8. Deleting Documents in a MongoDB Collection

Problem

You want to delete one or more documents in a MongoDB collection.

Solution

Use the `findAndRemove` method of the Casbah `MongoCollection` class to delete one document at a time, or use the `remove` method to delete one or more documents at a time.

The following code uses `findAndRemove` to delete the document whose `symbol` field is AAPL:

```

val query = MongoDBObject("symbol" -> "AAPL")
val result = collection.findAndRemove(query)
println("result: " + result)

```

When a document is deleted, the `findAndRemove` method returns the document that was deleted, wrapped in a `Some`:

```

result: Some({ "_id" : { "$oid" : "50255d1d03644925d83b3d07" } ,
  "symbol" : "AAPL" , "price" : 600.0})

```

If nothing is deleted, such as when you try to delete a document that doesn't exist, the result is `None`:

```

result: None

```

Therefore, you'll probably want to handle this using a match expression, as shown in the previous recipe.

To delete multiple documents from the collection, specify your search criteria when using the `remove` method, such as deleting all documents whose `price` field is greater than 500:

```
collection.remove("price" $gt 500)
```

The following method is dangerous: it shows how to delete all documents in the current collection:

```
// removes all documents
def deleteAllObjectsFromCollection(coll: MongoClient) {
  coll.remove(MongoDBObject.newBuilder.result)
}
```

(Be careful with that one.)

Discussion

If you've been following along with the MongoDB recipes in this chapter, you can experiment with these approaches by saving the following code to a file named *DeleteApple.scala* in the root directory of your SBT project:

```
import com.mongodb.casbah.Imports._

object DeleteApple extends App {

  var collection = MongoFactory.collection

  // delete AAPL
  val query = MongoDBObject("symbol" -> "AAPL")
  val result = collection.findAndRemove(query)
  println("result: " + result)

}
```



You can also clone my complete Scala + Casbah + MongoDB project from [GitHub](#).

If your database has a document whose `symbol` field is `AAPL`, when you run this object with `sbt run`, the result will show the document that was deleted:

```
result: Some({ "_id" : { "$oid" : "5026b22c300478e85a145d43" } ,
  "symbol" : "AAPL" , "price" : 600.0 })
```

The following complete code shows how to delete multiple documents:


```
import com.mongodb.casbah.Imports._

object DeleteMultiple extends App {

  var collection = MongoFactory.collection

  // delete all documents with price > 200
  collection.remove("price" $gt 200)

}
```

In this case, the `remove` method doesn't return anything interesting, so I don't assign it to a result.

See Also

My [Scala + Casbah + MongoDB sample project](#)

16.9. A Quick Look at Slick

When it comes to working with relational databases, you can use the wealth of Java solutions that are available, but other tools are emerging to provide a “Scala way” of working with databases. One of these solutions is a library named **Slick**, from **Type-safe**, a company that was founded by the creators of the Scala language. According to their documentation, Slick provides a “modern database query and access library.”

This recipe doesn't cover Slick in depth because it's well documented on the Typesafe website, but instead offers a quick look at what Slick offers.

In short, Slick lets you define database table objects in your code like this:

```
object Authors extends Table[(Int, String, String)]("AUTHORS") {
  def id       = column[Int]("ID", 0.PrimaryKey)
  def firstName = column[String]("FIRST_NAME")
  def lastName  = column[String]("LAST_NAME")
  def * = id ~ firstName ~ lastName
}

object Books extends Table[(Int, String)]("BOOKS") {
  def id       = column[Int]("ID", 0.PrimaryKey)
  def title    = column[String]("TITLE")
  def * = id ~ title
}

object BookAuthors extends Table[(Int, Int, Int)]("BOOK_AUTHORS") {
  def id       = column[Int]("ID", 0.PrimaryKey)
  def bookId   = column[Int]("BOOK_ID")
  def authorId = column[Int]("AUTHOR_ID")
  def bookFk   = foreignKey("BOOK_FK", bookId, Books)(_.id)
  def authorFk = foreignKey("AUTHOR_FK", authorId, Authors)(_.id)
}
```

```
def * = id ~ bookId ~ authorId
}
```

Having defined your tables in Scala code, you can refer to the fields in the tables in a type-safe manner. You can create your database tables using Scala code, like this:

```
(Books.ddl ++ Authors.ddl ++ BookAuthors.ddl).create
```

A simple query to retrieve all records from the resulting books database table looks like this:

```
val q = Query(Books)
q.list.foreach(println)
```

You can filter queries using a `filter` method:

```
val q = Query(Books).filter(_.title.startsWith("Zen"))
q.list.foreach(println)
```

You can write a join like this:

```
val q = for {
  b <- Books
  a <- Authors
  ba <- BookAuthors if b.id === ba.bookId && a.id === ba.authorId
} yield (b.title, a.lastName)
q.foreach(println)
```

Insert, update, and delete expressions follow the same pattern. Because you declared the database design in Scala code, Slick makes working with a database feel like working with collections.

Though I appreciate a good DSL, one thing I always look for in a database library is a way to break out of the library to let me write my own SQL queries, and Slick allows this as well.

As mentioned, the [Slick documentation](#) is thorough, so it's not covered in this chapter. See the [Slick website](#) for more information.

Interacting with Java

Introduction

In general, the ability to easily mix Scala and Java code is pretty seamless and amazing. In most cases, you can create an SBT project, put your Scala code in *src/main/scala*, put your Java code in *src/main/java*, and it “just works.” For instance, the recipes on web services in [Chapter 15](#) provide many examples of calling existing Java libraries from Scala code.

In my Scala/Java interactions, the biggest issues I’ve run into deal with the differences between their collections libraries. However, I’ve always been able to work through those problems with Scala’s `JavaConversions` object.

If you’re going to be accessing Scala code from Java, the other problem you can run into is that there are things you can do in Scala that don’t map well to Java. If you’re going to use Scala features like implicit conversions and parameters, currying, traits that have implemented methods, and other advanced features, you’ll want to keep that Scala code away from your public API.

Finally, for some cases such as serialization, methods with `varargs` parameters, and creating JavaBean-like classes in Scala, it’s important to know the annotations that are available to you.

17.1. Going to and from Java Collections

Problem

You’re using Java classes in a Scala application, and those classes either return Java collections, or require Java collections in their method calls.

Solution

Use the methods of Scala's `JavaConversions` object to make the conversions work.

For instance, the `java.util.ArrayList` class is commonly used in Java applications, and you can simulate receiving an `ArrayList` from a method in the REPL, like this:

```
scala> def nums = {  
  |   var list = new java.util.ArrayList[Int]()  
  |   list.add(1)  
  |   list.add(2)  
  |   list  
  | }  
nums: java.util.ArrayList[Int]
```

Even though this method is written in Scala, when it's called, it acts just as though it was returning an `ArrayList` from a Java method:

```
scala> val list = nums  
list: java.util.ArrayList[Int] = [1, 2]
```

However, because it's a Java collection, I can't call the `foreach` method on it that I've come to know and love in Scala, because it isn't there:

```
scala> list.foreach(println)  
<console>:10: error:  
value foreach is not a member of java.util.ArrayList[Int]  
    list.foreach(println)  
      ^
```

But by importing the methods from the `JavaConversions` object, the `ArrayList` magically acquires a `foreach` method:

```
scala> import scala.collection.JavaConversions._  
import scala.collection.JavaConversions._  
  
scala> list.foreach(println)  
1  
2
```

This “magic” comes from the power of Scala's implicit conversions, which are demonstrated in [Recipe 1.10, “Add Your Own Methods to the String Class”](#).

Discussion

When you get a reference to a Java collections object, you can either use that object as a Java collection (such as using its `Iterator`), or you can convert that collection to a Scala collection. Once you become comfortable with Scala collection methods like `foreach`, `map`, etc., you'll definitely want to treat it as a Scala collection, and the way to do that is to use the methods of the `JavaConversions` object.

As a more thorough example of how the `JavaConversions` methods work, assume you have a Java class named `JavaExamples` with the following `getNumbers` method:

```
// java
public static List<Integer> getNumbers() {
    List<Integer> numbers = new ArrayList<Integer>();
    numbers.add(1);
    numbers.add(2);
    numbers.add(3);
    return numbers;
}
```

You can attempt to call that method from Scala code, as shown in this example:

```
val numbers = JavaExamples.getNumbers()
numbers.foreach(println) // this won't work
```

But this code will fail with the following compiler error:

```
value 'foreach' is not a member of java.util.List[Integer]
```

To solve this problem, you need to import the `JavaConversions.asScalaBuffer` method. When you do this, you can either explicitly call the `asScalaBuffer` method, or let it be used implicitly. The *explicit* call looks like this:

```
import scala.collection.JavaConversions.asScalaBuffer

val numbers = asScalaBuffer(JavaExamples.getNumbers)
numbers.foreach(println)

// prints 'scala.collection.convert.Wrappers$JListWrapper'
println(numbers.getClass)
```

The *implicit* use looks like this:

```
import scala.collection.JavaConversions.asScalaBuffer

val numbers = JavaExamples.getNumbers
numbers.foreach(println)

// prints 'java.util.ArrayList'
println(numbers.getClass)
```

The `println(numbers.getClass)` calls show that there's a slight difference in the result between the explicit and implicit uses. Using the explicit `asScalaBuffer` method call makes the `numbers` object an instance of `collection.convert.Wrappers$JListWrapper`, whereas the implicit use shows that `numbers` is an `ArrayList`. As a practical matter, you can use either approach, depending on your preferences about working with implicit conversions; they both let you call `foreach`, `map`, and other Scala sequence methods.

You can repeat the same example using a Java `Map` and `HashMap`. First, create this method in a `JavaExamples` class:

```
// java
public static Map<String, String> getPeeps() {
    Map<String, String> peeps = new HashMap<String, String>();
    peeps.put("captain", "Kirk");
    peeps.put("doctor", "McCoy");
    return peeps;
}
```

Then, before calling this method from your Scala code, import the appropriate `JavaConversions` method:

```
import scala.collection.JavaConversions.mapAsScalaMap
```

You can then call the `mapAsScalaMap` method explicitly, or allow it to be called implicitly:

```
// explicit call
val peeps1 = mapAsScalaMap(JavaExamples.getPeeps)

// implicit conversion
val peeps2 = JavaExamples.getPeeps
```

Again there is a difference between the types of the map objects. In this case, `peeps1`, which used the explicit method call, has a type of `collection.convert.Wrappers$JMapWrapper`, whereas `peeps2` has a type of `java.util.HashMap`.

Note that the `JavaConversions` class has been through a number of changes, and although you'll see a large number of conversion methods in your IDE, many of them are deprecated. See the latest Scaladoc for [the `JavaConversions` object](#) for up-to-date information.

Conversion tables

One interesting thing that happens during the process of converting Java collections is that you learn more about the Scala collections. For instance, given their names, you might expect a Scala `List` and a Java `List` to convert back and forth between each other, but that isn't the case. A Java `List` is much more like a Scala `Seq` or a mutable `Buffer`.

This is shown in [Table 17-1](#), which shows the two-way conversions that the `JavaConversions` object allows between Java and Scala collections. This table is adapted from [the `JavaConversions` documentation](#).

Table 17-1. The two-way conversions provided by the `JavaConversions` object

Scala collection	Java collection
<code>collection.Iterable</code>	<code>java.lang.Iterable</code>
<code>collection.Iterable</code>	<code>java.util.Collection</code>
<code>collection.Iterator</code>	<code>java.util.{Iterator, Enumeration}</code>
<code>collection.mutable.Buffer</code>	<code>java.util.List</code>
<code>collection.mutable.Set</code>	<code>java.util.Set</code>
<code>collection.mutable.Map</code>	<code>java.util.{Map, Dictionary}</code>

Scala collection	Java collection
<code>collection.mutable.ConcurrentMap</code>	<code>java.util.concurrent.ConcurrentMap</code>

As an example of the two-way conversions shown in [Table 17-1](#), the `JavaConversions` object provides methods that convert between a Java `List` and a Scala `Buffer`. The `asScalaBuffer` method converts a Java `List` to a Scala `Buffer`, and `bufferAsJavaList` converts in the opposite direction, from a `Buffer` to a `List`.

Going from Scala collections to Java collections

So far you’ve looked primarily at converting Java collections to Scala collections. You may also need to go in the other direction, from a Scala collection to a Java collection.

If you’re converting a Scala collection to a Java collection, in addition to the two-way conversions shown in [Table 17-1](#), the one-way conversions shown in [Table 17-2](#) are available. Again, these have been adapted from the `JavaConversions` Scaladoc.

Table 17-2. The Scala to Java one-way conversions provided by the `JavaConversions` class

Scala collection	Java collection
<code>collection.Seq</code>	<code>java.util.List</code>
<code>collection.mutable.Seq</code>	<code>java.util.List</code>
<code>collection.Set</code>	<code>java.util.Set</code>
<code>collection.Map</code>	<code>java.util.Map</code>
<code>collection.mutable.Map[String,String]</code>	<code>java.util.Properties</code>

For example, assume you want to call the following `sum` method declared in a Java class named `ConversionExamples`, which expects a `java.util.List<Integer>`:

```
// java
public static int sum(List<Integer> list) {
    int sum = 0;
    for (int i: list) { sum = sum + i; }
    return sum;
}
```

Putting the conversion tables to work, the following examples show how to pass a `Seq`, `ArrayBuffer`, and `ListBuffer` to the `sum` method:

```
import scala.collection.JavaConversions._
import scala.collection.mutable._

val sum1 = ConversionExamples.sum(seqAsJavaList(Seq(1, 2, 3)))
val sum2 = ConversionExamples.sum(bufferAsJavaList(ArrayBuffer(1,2,3)))
val sum3 = ConversionExamples.sum(bufferAsJavaList(ListBuffer(1,2,3)))
```

There are many other collection conversion possibilities, and hopefully these examples will get you started on the right path.

The JavaConverters object

The `Scala JavaConverters object` lets you perform conversions in a manner similar to the examples shown, though they don't offer implicit conversions. Instead they require you to explicitly call `asJava` or `asScala` methods to perform the conversions. Be careful, because the object also contains many deprecated methods.

See Also

- The `JavaConversions object`
- The `JavaConverters object`

17.2. Add Exception Annotations to Scala Methods to Work with Java

Problem

You want to let Java users know that a method can throw one or more exceptions so they can handle those exceptions with try/catch blocks.

Solution

Add the `@throws` annotation to your Scala methods so Java consumers will know which methods can throw exceptions and what exceptions they throw.

For example, the following Scala code shows how to add an `@throws` annotation to let callers know that the `exceptionThrower` method can throw an `Exception`:

```
// scala
class Thrower {

  @throws(classOf[Exception])
  def exceptionThrower {
    throw new Exception("Exception!")
  }

}
```

With your Scala method annotated like that, it will work just like a Java method that throws an exception. If you attempt to call `exceptionThrower` from a Java class without wrapping it in a try/catch block, or declaring that your Java method throws an exception, the compiler (or your IDE) will give you the following error:

```
unreported exception java.lang.Exception; must be caught or declared to be thrown
```

In your Java code, you'll write a try/catch block as usual to handle the exception:


```
// java
Thrower t = new Thrower();
try {
    t.exceptionThrower();
} catch (Exception e) {
    System.err.println("Caught the exception.");
    e.printStackTrace();
}
```

If you want to declare that your Scala method throws multiple exceptions, add an annotation for each exception:

```
@throws(classOf[IOException])
@throws(classOf[LineUnavailableException])
@throws(classOf[UnsupportedAudioFileException])
def playSoundFileWithJavaAudio {
    // exception throwing code here ...
}
```

Discussion

If you don't mark the Scala `exceptionThrower` method with the `@throws` annotation, a Java developer *can* call it without using a try/catch block in her method, or declaring that her method throws an exception. For example, you can define the Scala method as follows, without declaring that it throws an exception:

```
//scala
def exceptionThrower {
    throw new Exception("Exception!")
}
```

This method can then be called from Java:

```
// java
public static void main(String[] args) {
    Thrower t = new Thrower();
    t.exceptionThrower();
}
```

However, when the Java developer calls `exceptionThrower`, the uncaught exception will cause the Java method to fail:

```
[error] (run-main) java.lang.Exception: Exception!
java.lang.Exception: Exception!
    at Thrower.exceptionThrower(Thrower.scala:6)
    at Main.main(Main.java:9)
```

As shown, if a Java consumer doesn't know an exception can be thrown, it can wreak havoc on her application.

17.3. Using @SerialVersionUID and Other Annotations

Problem

You want to specify that a class is serializable, and set the serialVersionUID. More generally, you want to know the syntax for using annotations in your Scala code, and know which annotations are available.

Solution

Use the Scala @SerialVersionUID annotation while also having your class extend the Serializable trait:

```
@SerialVersionUID(1000L)
class Foo extends Serializable {
  // class code here
}
```

Note that Scala has a serializable annotation, but it has been deprecated since version 2.9.0. The serializable annotation Scaladoc includes the following note:

```
instead of @serializable class C, use class C extends Serializable
```

Discussion

In addition to the @SerialVersionUID annotation and the Serializable trait, Scala has other annotations that should be used for various purposes, including the cloneable, remote, transient, and volatile annotations. Based primarily on the “[A Tour of Scala Annotations](#)” web page, Table 17-3 shows a mapping of Scala annotations to their Java equivalents.

Table 17-3. Scala annotations and their Java equivalents

Scala	Java
scala.beans.BeanProperty	No equivalent. When added to a class field, it results in getter and setter methods being generated that match the JavaBean specification.
scala.cloneable	java.lang.Cloneable
scala.deprecated	java.lang.Deprecated
scala.inline	Per the Scaladoc, @inline “requests that the compiler should try especially hard to inline the annotated method.”
scala.native	The Java native keyword.
scala.remote	java.rmi.Remote
scala.serializable	java.io.Serializable
scala.SerialVersionUID	serialVersionUID field.
scala throws	throws keyword.

Scala	Java
<code>scala.transient</code>	<code>transient</code> keyword.
<code>scala.unchecked</code>	No equivalent. According to its Scaladoc, it designates that “the annotated entity should not be considered for additional compiler checks.”
<code>scala.annotation.varargs</code>	Used on a field in a method, it instructs the compiler to generate a Java varargs-style parameter.
<code>scala.volatile</code>	<code>volatile</code> keyword.

As one example of these annotations, the current nightly version of [the Scala Remote Scaladoc](#) states that the following Scala code and Java code are equivalent:

```
// scala
@remote trait Hello {
  def sayHello(): String
}

// java
public interface Hello extends java.rmi.Remote {
  String sayHello() throws java.rmi.RemoteException;
}
```

[Recipe 17.6](#), “When Java Code Requires JavaBeans” provides examples of the `BeanProperty` annotation.

See Also

- The `Serializable` trait is [deprecated](#)
- “[A Tour of Scala Annotations](#)”
- [Recipe 17.5](#) discusses the `@varargs` annotation, and [Recipe 17.6](#) discusses JavaBeans

17.4. Using the Spring Framework

Problem

You want to use the Java Spring Framework library in your Scala application.

Solution

In my experience, the only real changes in using the Spring Framework in Scala applications involve how you cast the objects you instantiate from your Spring application context file, and that’s only because the casting processes in Scala and Java are different.

To demonstrate this, create an empty SBT project. (See [Recipe 18.1](#), if necessary.) Within that project, create a Spring *applicationContext.xml* file in the *src/main/resources* directory with the following contents:

```
<?xml version="1.0" encoding="UTF-8"?>
<!DOCTYPE beans PUBLIC "-//SPRING//DTD BEAN//EN" ⚡
    "http://www.springframework.org/dtd/spring-beans.dtd">

<beans>

    <bean id="dog" class="scalaspring.Dog">
        <constructor-arg value="Fido" />
    </bean>

    <bean id="cat" class="scalaspring.Cat">
        <constructor-arg value="Felix" />
    </bean>

</beans>
```

This file declares that there are two classes, one named *Dog* and the other named *Cat*, in a package named *scalaspring*. You can't tell it from looking at this file, but as you'll see shortly, both the *Dog* and *Cat* classes extend a base *Animal* class.

Next, create a simple Scala object in a file named *SpringExample.scala* in the root directory of your project with a *main* method to read the *applicationContext.xml* file and create instances of the *Dog* and *Cat* classes:

```
package scalaspring

import org.springframework.context.support.ClassPathXmlApplicationContext

object ScalaSpringExample extends App {

    // open & read the application context file
    val ctx = new ClassPathXmlApplicationContext("applicationContext.xml")

    // instantiate the dog and cat objects from the application context
    val dog = ctx.getBean("dog").asInstanceOf[Animal]
    val cat = ctx.getBean("cat").asInstanceOf[Animal]

    // let them speak
    dog.speak
    cat.speak

}
```

In this code, the *applicationContext.xml* file is loaded, the *dog* and *cat* instances are created from their bean definitions in the application context, and their *speak* methods are executed.

Next, define the Dog and Cat classes in a file named *Animals.scala*, along with their abstract parent class *Animal*. You can also save this file in the root directory of your SBT project:

```
package scalaspring

abstract class Animal(name: String) {
  def speak: Unit    // abstract
}

class Dog(name: String) extends Animal(name) {
  override def speak {
    println(name + " says Woof")
  }
}

class Cat(name: String) extends Animal(name) {
  override def speak {
    println(name + " says Meow")
  }
}
```

The base *Animal* class requires that the concrete classes have a *speak* method, and the *Dog* and *Cat* classes define their *speak* methods in different ways. The *Dog* and *Cat* classes are defined using a one-argument constructor, and if you look back at the application context file, you'll see that the names *Fido* and *Felix* are used in their Spring bean definitions.

Next, add Spring as a dependency to your *build.sbt* file. A basic file looks like this:

```
name := "Scala Spring Example"

version := "1.0"

scalaVersion := "2.10.0"

libraryDependencies += "org.springframework" % "spring" % "2.5.6"
```

As mentioned, you should place the *applicationContext.xml* file in your project's *src/main/resources* folder. This listing shows all the files in my project:

```
./Animals.scala
./build.sbt
./SpringExample.scala
./src/main/resources/applicationContext.xml
```

With everything in place, run the project with the usual `sbt run` command. You'll see a lot of output, including these lines, showing that the program ran successfully:

```
$ sbt run
```

```
Fido says Woof  
Felix says Meow
```



You can put the two Scala source files under the *src/main/scala* directory if you prefer, but for simple examples like this, I put them in the root directory of my SBT project.

Discussion

Although there was a bit of boilerplate work in this example, the only major differences between using Scala and Java are these two lines of code in the `ScalaSpringExample` object:

```
val dog = ctx.getBean("dog").asInstanceOf[Animal]  
val cat = ctx.getBean("cat").asInstanceOf[Animal]
```

That's because this is how you cast classes in Scala. In Java, these same lines of code would look like this:

```
Animal dog = (Animal)ctx.getBean("dog");  
Animal cat = (Animal)ctx.getBean("cat");
```

See Also

- [Recipe 6.1](#) provides other examples of casting in Scala
- [Recipe 16.2, “Connecting to a Database with the Spring Framework”](#) shows another Scala Spring example
- The “Spring Scala” project aims to make it easier to use the [Spring Framework in Scala](#)

17.5. Annotating varargs Methods

Problem

You've created a Scala method with a varargs field, and would like to be able to call that method from Java code.

Solution

When a Scala method has a field that takes a variable number of arguments, mark it with the `@varargs` annotation.

For example, the `printAll` method in the following Scala class is marked with `@varargs` so it can be called as desired from Java:

```
package varargs

import scala.annotation.varargs

class Printer {

  @varargs def printAll(args: String*) {
    args.foreach(print)
    println
  }

}
```

The `printAll` method can now be called from a Java program with a variable number of parameters, as shown in this example:

```
package varargs;

public class Main {

  public static void main(String[] args) {
    Printer p = new Printer();
    p.printAll("Hello");
    p.printAll("Hello, ", "world");
  }

}
```

When this code is run, it results in the following output:

```
Hello
Hello, world
```

Discussion

If the `@varargs` annotation isn't used on the `printAll` method, the Java code shown won't even compile, failing with the following compiler errors:

```
Main.java:7: printAll(scala.collection.Seq<java.lang.String>) in
varargs.Printer cannot be applied to (java.lang.String)
[error]      p.printAll("Hello");
[error]      ^

Main.java:8: printAll(scala.collection.Seq<java.lang.String>) in
varargs.Printer cannot be applied to (java.lang.String,java.lang.String)
[error]      p.printAll("Hello, ", "world");
[error]      ^
```

Without the `@varargs` annotation, from a Java perspective, the `printAll` method appears to take a `scala.collection.Seq<java.lang.String>` as its argument.

17.6. When Java Code Requires JavaBeans

Problem

You need to interact with a Java class or library that accepts only classes that conform to the JavaBean specification.

Solution

Use the `@BeanProperty` annotation on your fields, also making sure you declare each field as a `var`.

The `@BeanProperty` annotation can be used on fields in a Scala class constructor:

```
import scala.reflect.BeanProperty

class Person(@BeanProperty var firstName: String,
             @BeanProperty var lastName: String) {
  override def toString = s"Person: $firstName $lastName"
}
```

It can also be used on the fields in a Scala class:

```
import scala.reflect.BeanProperty

class EmailAccount {
  @BeanProperty var username: String = ""
  @BeanProperty var password: String = ""
  override def toString = s"Email Account: ($username, $password)"
}
```

To demonstrate this, create an SBT project, then save the following code to a file named *Test.scala* in the root directory of the project:

```
package foo

import scala.reflect.BeanProperty

class Person(@BeanProperty var firstName: String,
             @BeanProperty var lastName: String) {
}

class EmailAccount {
  @BeanProperty var username: String = ""
  @BeanProperty var password: String = ""
}
```

This code shows how to use the `@BeanProperty` annotation on class constructor parameters, as well as the fields in a class.

Next, create a directory named *src/main/java/foo*, and save the following Java code in a file named *Main.java* in that directory:

```
package foo;

public class Main {

    public static void main(String[] args) {

        // create instances
        Person p = new Person("Regina", "Goode");
        EmailAccount acct = new EmailAccount();

        // demonstrate 'setter' methods
        acct.setUsername("regina");
        acct.setPassword("secret");

        // demonstrate 'getter' methods
        System.out.println(p.getFirstName());
        System.out.println(p.getLastName());
        System.out.println(acct.getUsername());
        System.out.println(acct.getPassword());

    }

}
```

This Java code demonstrates how to create instances of the Scala *Person* and *EmailAccount* classes, and access the JavaBean methods of those classes. When the code is run with `sbt run`, you'll see the following output, showing that all the getter and setter methods work:

```
$ sbt run
[info] Running foo.Main
Regina
Goode
regina
secret
```

Discussion

You can see how the `@BeanProperty` annotation works by compiling a simple class and then disassembling it. First, save these contents to a file named *Person.scala*:

```
import scala.reflect.BeanProperty

class Person(@BeanProperty var name: String,
             @BeanProperty var age: Int) {

}
```

Then compile the class:

```
$ scalac Person.scala
```

After it's compiled, disassemble it with the `javap` command:

```
$ javap Person
```

```
Compiled from "Person.scala"
public class Person extends java.lang.Object implements scala.ScalaObject{
  public java.lang.String name();
  public void name_$eq(java.lang.String);
  public void setName(java.lang.String);
  public int age();
  public void age_$eq(int);
  public void setAge(int);
  public int getAge();
  public java.lang.String getName();
  public Person(java.lang.String, int);
}
```

As you can see from the disassembled code, the methods `getName`, `setName`, `getAge`, and `setAge` have all been generated because of the `@BeanProperty` annotation.

Note that if you declare your fields as type `val`, the “setter” methods (`setName`, `setAge`) won't be generated:

```
Compiled from "Person.scala"
public class Person extends java.lang.Object implements scala.ScalaObject{
  public java.lang.String name();
  public int age();
  public int getAge();
  public java.lang.String getName();
  public Person(java.lang.String, int);
}
```

Without these methods, your class will not follow the JavaBean specification.

As a final example, if the `@BeanProperty` annotation is removed from all fields, you're left with this code:

```
class Person(var firstName: String, var lastName: String)
```

When you compile this code with `scalac` and then disassemble it with `javap`, you'll see that no getter or setter methods are generated (except for those that follow the Scala convention):

```
Compiled from "Person.scala"
public class Person extends java.lang.Object{
  public java.lang.String firstName();
  public void firstName_$eq(java.lang.String);
  public java.lang.String lastName();
  public void lastName_$eq(java.lang.String);
  public Person(java.lang.String, java.lang.String);
}
```

See Also

My [tutorial about using the Java SnakeYaml library in Scala](#) shows more examples of the `@BeanProperty` annotation.

17.7. Wrapping Traits with Implementations

Problem

You've written a Scala trait with implemented methods and need to be able to use those methods from a Java application.

Solution

You can't use the implemented methods of a Scala trait from Java, so wrap the trait in a class.

Assuming you have a Scala trait named `MathTrait` with a method named `sum` that you want to access from Java code, create a Scala class named `MathTraitWrapper` that wraps `MathTrait`:

```
// scala
package foo

// the original trait
trait MathTrait {
  def sum(x: Int, y: Int) = x + y
}

// the wrapper class
class MathTraitWrapper extends MathTrait
```

In your Java code, extend the `MathTraitWrapper` class, and access the `sum` method:

```
// java
package foo;

public class JavaMath extends MathTraitWrapper {

  public static void main(String[] args) {
    new JavaMath();
  }

  public JavaMath() {
    System.out.println(sum(2,2));
  }

}
```

This code works as expected, printing the number 4 when it is run.

Discussion

A Java class can't extend a Scala trait that has implemented methods. To demonstrate the problem, first create a trait with a simple implemented method named `sum`:

```
package foo

trait MathTrait {
  def sum(x: Int, y: Int) = x + y
}
```

Next, to attempt to use this trait from Java, you have a choice of trying to extend it or implement it. Let's first try to *extend* it:

```
package foo;

public class JavaMath extends MathTrait {}
```

By the time you finish typing that code, you see the following compiler error message:

```
The type MathTrait cannot be the superclass of JavaMath;
a superclass must be a class
```

Next, you can attempt to *implement* the trait, but intuitively you know that won't work, because in Java you implement *interfaces*, and this trait has implemented behavior, so it's not a regular Java interface:

```
package foo;

public class JavaMath implements MathTrait {}
```

This code leads to the following compiler error:

```
The type JavaMath must implement the inherited abstract method
MathTrait.sum(int, int)
```

You *could* implement a `sum` method in your `JavaMath` class, but that defeats the purpose of trying to use the `sum` method that's already written in the Scala `MathTrait`.

Other attempts

You can try other things, such as attempting to create an instance of the `MathTrait` and trying to use the `sum` method, but this won't work either:

```
// java
package foo;

public static void main(String[] args) {
  MathTrait trait = new MathTrait(); // error, won't compile
  int sum = trait.sum(1,2);
  System.out.println("SUM = " + sum);
}
```

Trying to instantiate a `MathTrait` instance results in this compiler error:

```
foo.MathTrait is abstract; cannot be instantiated
[error] MathTrait trait = new MathTrait();
[error]           ^
```

You may already know what the problem is, but to be clear, let's see what class files are generated on the Scala side. In an SBT project, the class files are located in the following directory:

```
$PROJECT/target/scala-2.10.0/classes/foo
```

If you move into that directory and list the files, you'll see that two files related to the Scala `MathTrait` trait have been created:

```
MathTrait.class
MathTrait$class.class
```

You can see the problem by disassembling these files with `javap`. First, the *MathTrait.class* file:

```
$ javap MathTrait
Compiled from "MathTrait.scala"
public interface foo.MathTrait{
    public abstract int sum(int, int);
}
```

Next, the *MathTrait\$class.class* file:

```
$ javap MathTrait$class
Compiled from "MathTrait.scala"
public abstract class foo.MathTrait$class extends java.lang.Object{
    public static int sum(foo.MathTrait, int, int);
    public static void $init$(foo.MathTrait);
}
```

The problem with trying to work with the Scala `MathTrait` from a Java perspective is that `MathTrait` looks like an interface, and `MathTrait$class` looks like an abstract class. Neither one will let you use the logic in the `sum` method.

Because `MathTrait` looks like just an interface, you realize you might be able to create a Java class that implements that interface, and then override the `sum` method:

```
// java
package foo;

public class JavaMath implements MathTrait {

    public int sum(int x, int y) {
        return x + y;
    }
}
```

```

    public static void main(String[] args) {
        JavaMath math = new JavaMath();
        System.out.println(math.sum(1,1));
    }
}

```

This does indeed work, but for the purposes of this recipe, it doesn't really matter. The purpose of trying to use the trait was to use the behavior of the trait's `sum` method, and there's no way to do this from Java without creating a Scala wrapper class.

In a last desperate attempt, you might try to call `super.sum(x,y)` from your Java method, like this:

```

// java
public int sum(int x, int y) {
    return super.sum(x, y);
}

```

But that won't work either, failing with the following error message:

```

cannot find symbol
[error] symbol : method sum(int,int)
[error] location: class java.lang.Object
[error]     return super.sum(x,y);
[error]                   ^

```

The only way to solve the problem is to wrap the trait with a class on the Scala side, which was demonstrated in the Solution.

To summarize: If you're writing a Scala API that will be used by Java clients, don't expose traits that have implemented behavior in your public API. If you have traits like that, wrap them in a class for your Java consumers.

The Simple Build Tool (SBT)

Introduction

Although you can use Ant and Maven to build your Scala projects, SBT, or the **Simple Build Tool**, is the de facto build tool for Scala applications. SBT makes the basic build and dependency management tasks simple, and lets you use the Scala language itself to conquer more difficult tasks.

SBT uses the same directory structure as Maven, and like Maven, it uses a “convention over configuration” approach that makes the build process incredibly easy for basic projects. Because it provides a well-known, standard build process, if you work on one Scala project that’s built with SBT, it’s easy to move to another project that also uses SBT. The project’s directory structure will be the same, and you’ll know that you should look at the *build.sbt* file and the optional *project/*.scala* files to see how the build process is configured.

Like Maven, under the covers, SBT’s dependency management system is handled by **Apache Ivy**. This means that all those Java projects that have been created and packaged for use with Maven over the years can easily be used by SBT. Additionally, other JAR files not in an Ivy/Maven repository can simply be placed in your project’s *lib* folder, and SBT will automatically find them.

As a result of all these features, with very little effort on your part, SBT lets you build projects that contain both Scala and Java code, unit tests, and both managed and unmanaged dependencies.



All examples in this chapter were tested with SBT version 0.12.3.

18.1. Creating a Project Directory Structure for SBT

Problem

SBT doesn't include a command to create a new project, and you'd like to quickly and easily create the directory structure for a new project.

Solution

Use either a shell script or a tool like [Giter8](#) to create your project's directory structure. Both approaches are shown here.

Use a shell script

SBT uses the same directory structure as Maven, and for simple needs, you can generate a compatible structure using a shell script. For example, the following Unix shell script creates the initial set of files and directories you'll want for most projects:

```
#!/bin/sh
mkdir -p src/{main,test}/{java,resources,scala}
mkdir lib project target

# create an initial build.sbt file
echo 'name := "MyProject"'

version := "1.0"

scalaVersion := "2.10.0" > build.sbt
```

Just save that code as a shell script on Unix systems (or Cygwin on Windows), make it executable, and run it inside a new project directory to create all the subdirectories SBT needs, as well as an initial *build.sbt* file.

Assuming this script is named `mkdirs4sbt`, and it's on your path, the process looks like this:

```
/Users/Al/Projects> mkdir MyNewProject

/Users/Al/Projects> cd MyNewProject

/Users/Al/Projects/MyNewProject> mkdirs4sbt
```

If you have the `tree` command on your system and run it from the current directory, you'll see that the basic directory structure looks like this:


```

.
|-- build.sbt
|-- lib
|-- project
|-- src
|   |-- main
|   |   |-- java
|   |   |-- resources
|   |   |-- scala
|   |-- test
|       |-- java
|       |-- resources
|       |-- scala
|-- target

```

This is just a simple starter script, but it helps to show how easy it is to create a basic SBT directory structure.



The *build.sbt* file is SBT's basic configuration file. You define most settings that SBT needs in this file, including specifying library dependencies, repositories, and any other basic settings your project requires. I'll demonstrate many examples of it in the recipes in this chapter.

Use Giter8

Although that script shows how simple building a basic directory structure is, **Giter8** is an excellent tool for creating SBT directory structures with specific project needs. It's based on a template system, and the templates usually contain everything you need to create a skeleton SBT project that's preconfigured to use one or more Scala tools, such as ScalaTest, Scalatra, and many others.

The Giter8 templates that you can use are listed on **GitHub**. As a demonstration of how this works, the following example shows how to use the *scalatra/scalatra-sbt* template.

To create a project named “NewApp,” Giter8 prompts you with a series of questions, and then creates a *newapp* folder for your project. To demonstrate this, move to a directory where you normally create your projects, then start Giter8 with the *g8* command, giving it the name of the template you want to use:

```

/Users/Al/Projects> g8 scalatra/scalatra-sbt

organization [com.example]: com.alvinalexander
package [com.example.app]: com.alvinalexander.newapp
name [My Scalatra Web App]: NewApp
scalatra_version [2.2.0]:
servlet_name [MyScalatraServlet]: NewAppServlet

```

```
scala_version [2.10.0]:  
version [0.1.0-SNAPSHOT]:
```

```
Template applied in ./newapp
```

Because I answered the name prompt with `NewApp`, Giter8 creates a subdirectory under the current directory named *newapp*. It contains the following files and directories:

```
newapp/.gitignore  
newapp/project/build.properties  
newapp/project/build.scala  
newapp/project/plugins.sbt  
newapp/README.md  
newapp/sbt  
newapp/src/main/resources/logback.xml  
newapp/src/main/scala/com/alvinalexander/newapp/NewAppServlet.scala  
newapp/src/main/scala/com/alvinalexander/newapp/NewAppStack.scala  
newapp/src/main/scala/ScalatraBootstrap.scala  
newapp/src/main/webapp/WEB-INF/templates/layouts/default.jade  
newapp/src/main/webapp/WEB-INF/templates/views/hello-scalate.jade  
newapp/src/main/webapp/WEB-INF/web.xml  
newapp/src/test/scala/com/alvinalexander/newapp/NewAppServletSpec.scala
```

In this example, Giter8 creates all the configuration files and Scalatra stub files you need to have a skeleton Scalatra project up and running in just a minute or two.

Giter8 notes

At the time of this writing, I had a problem with the current Scalatra template, and had to add this line to the *build.sbt* file in my root project directory to get the template to work:

```
scalaVersion := "2.10.0"
```

To run a Scalatra project, enter the SBT shell from your operating system command line, and then run the `container:start` command:

```
/Users/Al/Projects/newapp> sbt
```

```
> container:start
```

This command starts the Jetty server running on port 8080 on your computer, so you can easily test your installation by accessing the *http://localhost:8080/* URL from a browser.

In the case of using a new template like this, SBT may have a *lot* of files to download. Fear not—once they're downloaded, your new Scalatra project should be up and running, and all of these downloads are required only during the first sbt run.

Discussion

As shown in the Solution, because the SBT directory structure is standard and based on the Maven directory structure, you can create your own tool, or use other tools that are built for this purpose.

For simple SBT projects, I've created an improved version of the shell script shown in the Solution. I named it *sbtmkdirs*, and you can download it from the URL shown in the See Also section. Like Giter8, this script prompts you with several questions, and optionally creates *.gitignore* and *README.md* files, along with a full *build.sbt* file. I use this script whenever I want to create a Scala project where I don't need a template.

As demonstrated, Giter8 works by downloading project templates from GitHub. Giter8 requires SBT and another tool named Conscript, so to install and use Giter8, you'll need to follow these steps:

1. Install **SBT**.
2. Install **Conscript**.
3. Install **Giter8**.

Fortunately those projects are well documented, and it takes just a few minutes to install all three tools.

There have been a couple of times when I've used Giter8 and it failed to download a project template. I don't remember the exact error messages, but this was the most recent one:

```
$ g8 scalatra/scalatra-sbt
```

```
Unable to find github repository: scalatra/scalatra-sbt.g8 (master)
```

Each time this has happened, I've upgraded Conscript and Giter8 to their latest versions, and the problem has gone away.



Conscript is an interesting tool that works with GitHub projects for the purpose of “installing and updating Scala programs.” Its purpose and installation process are well documented at its website.

Giter8 currently uses a Java installer. Installing it on a Mac OS X system failed when I double-clicked the JAR file, but when I ran it from the command line (using the `java -jar` approach), it installed successfully.

See Also

- The [SBT website](#)
- Information about [installing SBT](#)
- My [sbtkmdirs script](#)
- The [Giter8 website](#)
- There are currently over [thirty Giter8 templates](#)
- The [Conscript website](#)

18.2. Compiling, Running, and Packaging a Scala Project with SBT

Problem

You want to use SBT to compile and run a Scala project, and package the project as a JAR file.

Solution

Create a directory layout to match what SBT expects, then run `sbt compile` to compile your project, `sbt run` to run your project, and `sbt package` to package your project as a JAR file.

To demonstrate this, create a new SBT project directory structure as shown in [Recipe 18.1](#), and then create a file named *Hello.scala* in the `src/main/scala` directory with these contents:

```
package foo.bar.baz

object Main extends App {
  println("Hello, world")
}
```

Unlike Java, in Scala, the file's package name doesn't have to match the directory name. In fact, for simple tests like this, you can place this file in the root directory of your SBT project, if you prefer.

From the root directory of the project, you can compile the project:

```
$ sbt compile
```

Run the project:

```
$ sbt run
```

Package the project:

```
$ sbt package
```

Discussion

The first time you run SBT, it may take a while to download all the dependencies it needs, but after that first run, it will download new dependencies only as needed. The commands executed in the Solution, along with their output, are shown here:

```
$ sbt compile
```

```
[info] Loading global plugins from /Users/Al/.sbt/plugins
[info] Set current project to Basic (in build file:/Users/Al/SbtTests/)
[success] Total time: 0 s
```

```
$ sbt run
```

```
[info] Loading global plugins from /Users/Al/.sbt/plugins
[info] Set current project to Basic (in build file:/Users/Al/SbtTests/)
[info] Running foo.bar.baz.Main
Hello, world
[success] Total time: 1 s
```

```
$ sbt package
```

```
[info] Loading global plugins from /Users/Al/.sbt/plugins
[info] Set current project to Basic (in build file:/Users/Al/SbtTests/)
[info] Packaging /Users/Al/SbtTests/target/scala-2.10/basic_2.10-1.0.jar ...
[info] Done packaging.
[success] Total time: 0 s
```

Because `compile` is a dependency of `run`, you don't have to run `compile` before each run; just type `sbt run`.

The JAR file created with `sbt package` is a normal Java JAR file. You can list its contents with the usual `jar tvf` command:

```
$ jar tvf target/scala-2.10/basic_2.10-1.0.jar
 261 Sat Apr 13 13:58:44 MDT 2013 META-INF/MANIFEST.MF
    0 Sat Apr 13 13:58:44 MDT 2013 foo/
    0 Sat Apr 13 13:58:44 MDT 2013 foo/bar/
    0 Sat Apr 13 13:58:44 MDT 2013 foo/bar/baz/
2146 Sat Apr 13 13:57:52 MDT 2013 foo/bar/baz/Main$.class
1003 Sat Apr 13 13:57:52 MDT 2013 foo/bar/baz/Main.class
 759 Sat Apr 13 13:57:52 MDT 2013 foo/bar/baz/Main$delayedInit$body.class
```

You can also execute the main method in the JAR file with the Scala interpreter:

```
$ scala target/scala-2.10/basic_2.10-1.0.jar
Hello, world
```

SBT commands

As with any Java-based command, there can be a little startup lag time involved with running SBT commands, so when you're using SBT quite a bit, it's common to run these commands in interactive mode from the SBT shell prompt to improve the speed of the process:

```
$ sbt

> compile
> run
> package
```

You can run multiple commands at one time, such as running `clean` before `compile`:

```
> clean compile
```

At the time of this writing, there are 247 SBT commands available (which I just found out by hitting the Tab key at the SBT shell prompt, which triggered SBT's tab completion). [Table 18-1](#) shows a list of the most common commands.

Table 18-1. Descriptions of the most common SBT commands

Command	Description
<code>clean</code>	Removes all generated files from the <i>target</i> directory.
<code>compile</code>	Compiles source code files that are in <i>src/main/scala</i> , <i>src/main/java</i> , and the root directory of the project.
<code>~ compile</code>	Automatically recompiles source code files while you're running SBT in interactive mode (i.e., while you're at the SBT command prompt).
<code>console</code>	Compiles the source code files in the project, puts them on the classpath, and starts the Scala interpreter (REPL).
<code>doc</code>	Generates API documentation from your Scala source code using <code>scaLadoc</code> .
<code>help <command></code>	Issued by itself, the <code>help</code> command lists the common commands that are currently available. When given a command, <code>help</code> provides a description of that command.
<code>inspect <setting></code>	Displays information about <code><setting></code> . For instance, <code>inspect library-dependencies</code> displays information about the <code>libraryDependencies</code> setting. (Variables in <i>build.sbt</i> are written in camelCase, but at the SBT prompt, you type them using this hyphen format instead of camelCase.)
<code>package</code>	Creates a JAR file (or WAR file for web projects) containing the files in <i>src/main/scala</i> , <i>src/main/java</i> , and resources in <i>src/main/resources</i> .
<code>package-doc</code>	Creates a JAR file containing API documentation generated from your Scala source code.
<code>publish</code>	Publishes your project to a remote repository. See Recipe 18.15, "Publishing Your Library" .
<code>publish-local</code>	Publishes your project to a local Ivy repository. See Recipe 18.15, "Publishing Your Library" .

Command	Description
reload	Reloads the build definition files (<i>build.sbt</i> , <i>project/*.scala</i> , and <i>project/*.sbt</i>), which is necessary if you change them while you're in an interactive SBT session.
run	Compiles your code, and runs the <code>main</code> class from your project, in the same JVM as SBT. If your project has multiple <code>main</code> methods (or objects that extend <code>App</code>), you'll be prompted to select one to run.
test	Compiles and runs all tests.
update	Updates external dependencies.

There are many other SBT commands available, and when you use plug-ins, they can also make their own commands available. For instance, [Recipe 18.7, “Using SBT with Eclipse”](#) shows that the `sbtclipse` plug-in adds an `eclipse` command. See [the SBT documentation](#) for more information.

Continuous compiling

As mentioned, you can eliminate the SBT startup lag time by starting the SBT interpreter in “interactive mode.” To do this, type `sbt` at your operating system command line:

```
$ sbt
>
```

When you issue your commands from the SBT shell, they'll run noticeably faster.

As shown in the Solution, you can issue the `compile` command from within the SBT shell, but you can also take this a step further and continuously compile your source code by using the `~ compile` command instead. When you issue this command, SBT watches your source code files, and automatically recompiles them whenever it sees the code change.

To demonstrate this, start the SBT shell from the root directory of your project:

```
$ sbt
```

Then issue the `~ compile` command:

```
> ~ compile
[info] Compiling 1 Scala source to /Users/Al/SbtTests/target/scala-2.10/classes
[success] Total time: 4 s, completed Apr 13, 2013 2:34:23 PM
1. Waiting for source changes... (press enter to interrupt)
```

Now, any time you change and save a source code file, SBT automatically recompiles it. You'll see these new lines of output when SBT recompiles the code:

```
[info] Compiling 1 Scala source to /Users/Al/SbtTests/target/scala-2.10/classes
[success] Total time: 2 s, completed Apr 13, 2013 2:34:32 PM
2. Waiting for source changes... (press enter to interrupt)
```

Use `last` to get more information on the last command

From time to time when working in the SBT shell you may have a problem, such as with incremental compiling. When issues like this come up, you may be able to use the shell's `last` command to see what happened.

For instance, you may issue a `compile` command, and then see something wrong in the output:

```
> compile
[info] Updating ...
[info] Resolving com.typesafe#config;1.0.0 ...
[info] Compiling 1 Scala source to
YIKES!
```

I made up the YIKES! part, but you get the idea; something goes wrong. To see what happened, issue the `last compile` command:

```
> last compile
[debug]
[debug] Initial source changes:
[debug]   removed:Set()
[debug]   added: Set(/Users/Al/Projects/Scala/Foo/Test.scala)
[debug]   modified: Set()
[debug] Removed products: Set()
[debug] Modified external sources: Set()
```

many more lines of debug output here ...

The `last` command prints logging information for the last command that was executed. This can help you understand what's happening, including understanding why something is being recompiled over and over when using incremental compilation.

Typing `help last` in the SBT interpreter shows a few additional details, including a note about the `last-grep` command, which can be useful when you need to filter a large amount of output.

See Also

- The [SBT command-line reference](#).
- Information on [publishing an SBT project](#).
- Incremental compiling can often be much (much!) faster than compiling an entire project. See [the Scala website](#) for more details on how it works in SBT.
- Typesafe has made SBT's incremental compiler available as a standalone tool named [Zinc](#), which can be used with other tools, like Maven.

18.3. Running Tests with SBT and ScalaTest

Problem

You want to set up an SBT project with ScalaTest, and run the tests with SBT.

Solution

Create a new SBT project directory structure as shown in [Recipe 18.1](#), and then add the ScalaTest library dependency to your *build.sbt* file, as shown here:

```
name := "BasicProjectWithScalaTest"

version := "1.0"

scalaVersion := "2.10.0"

libraryDependencies += "org.scalatest" % "scalatest_2.10" % "1.9.1" % "test"
```

Add your Scala source code under the *src/main/scala* folder, add your tests under the *src/test/scala* folder, and then run the tests with the SBT test command:

```
$ sbt test
```

Discussion

The `libraryDependencies` tag in the *build.sbt* file shows the standard way of adding new dependencies to an SBT project:

```
libraryDependencies += "org.scalatest" % "scalatest_2.10" % "1.9.1" % "test"
```

You can write that line as shown, or this way:

```
libraryDependencies += "org.scalatest" %% "scalatest" % "1.9.1" % "test"
```

In the second example, I used the `%%` method to automatically append the project's Scala version (2.10) to the end of the artifact name (`scalatest`). These two options are explained more in [Recipe 18.4, “Managing Dependencies with SBT”](#), but hopefully the way they work is clear from those examples.

To demonstrate how ScalaTest integrates seamlessly with SBT, create a source file named *Hello.scala* with the following contents in the *src/main/scala* directory of your project:

```
package com.alvinalexander.testproject

object Hello extends App {
  val p = Person("Alvin Alexander")
  println("Hello from " + p.name)
}

case class Person(var name: String)
```

Then create a test file named *HelloTests.scala* in the *src/test/scala* directory of your project with these contents:

```
package com.alvinalexander.testproject

import org.scalatest.FunSuite

class HelloTests extends FunSuite {

  test("the name is set correctly in constructor") {
    val p = Person("Barney Rubble")
    assert(p.name == "Barney Rubble")
  }

  test("a Person's name can be changed") {
    val p = Person("Chad Johnson")
    p.name = "Ochocinco"
    assert(p.name == "Ochocinco")
  }
}
```

Next, run your tests from your project's root directory with SBT:

```
$ sbt test

[info] Loading global plugins from /Users/Al/.sbt/plugins
[info] Set current project to BasicProjectWithScalaTest (in build
      file:/Users/Al/Projects/BasicProjectWithScalaTest/)
[info] HelloTests:
[info] - the name is set correctly in constructor
[info] - a Person's name can be changed
[info] Passed: : Total 2, Failed 0, Errors 0, Passed 2, Skipped 0
[success] Total time: 0 s
```

This output shows that the two tests in the *HelloTests* test class were run.

As shown in these examples, there's nothing special you have to do to make *ScalaTest* work with SBT, other than adding it as a dependency in the *build.sbt* file; it just works.



If you reused an existing SBT project folder to test this recipe, you may need to issue the SBT reload command. As described in [Table 18-1](#), this command tells SBT to reload the project definition files, including the *build.sbt* file.

See Also

- The [ScalaTest “quick start” page](#).
- If you'd like a simple way to test this, you can download the code for this recipe from [GitHub](#).

- [specs2](#) is another popular Scala testing framework that integrates easily with SBT. It compares well to ScalaTest, and is also the default testing library for the Play Framework.
- The [SBT Quick Configuration documentation](#) shows dozens of *build.sbt* examples.

18.4. Managing Dependencies with SBT

Problem

You want to use one or more external libraries in your Scala/SBT projects.

Solution

You can use both managed and unmanaged dependencies in your SBT projects.

If you have JAR files (*unmanaged* dependencies) that you want to use in your project, simply copy them to the *lib* folder in the root directory of your SBT project, and SBT will find them automatically. If those JARs depend on other JAR files, you'll have to download those other JAR files and copy them to the *lib* directory as well.

If you have a single *managed* dependency, such as wanting to use the Java [HtmlCleaner library](#) in your project, add a `libraryDependencies` line like this to your *build.sbt* file:

```
libraryDependencies += "net.sourceforge.htmlcleaner" % "htmlcleaner" % "2.4"
```

Because configuration lines in *build.sbt* must be separated by blank lines, a simple but complete file with one dependency looks like this:

```
name := "BasicProjectWithScalaTest"

version := "1.0"

scalaVersion := "2.10.0"

libraryDependencies += "org.scalatest" %% "scalatest" % "1.9.1" % "test"
```

To add multiple managed dependencies to your project, define them as a `Seq` in your *build.sbt* file:

```
libraryDependencies += Seq(
  "net.sourceforge.htmlcleaner" % "htmlcleaner" % "2.4",
  "org.scalatest" % "scalatest_2.10" % "1.9.1" % "test",
  "org.foobar" %% "foobar" % "1.8"
)
```

Or, if you prefer, you can add them one line at a time to the file, separating each line by a blank line:

```
libraryDependencies += "net.sourceforge.htmlcleaner" % "htmlcleaner" % "2.4"

libraryDependencies += "org.scalatest" % "scalatest_2.10" % "1.9.1" % "test"

libraryDependencies += "org.foobar" %% "foobar" % "1.6"
```

As you might infer from these examples, entries in *build.sbt* are simple key/value pairs. SBT works by creating a large map of key/value pairs that describe the build, and when it parses this file, it adds the pairs you define to its map. The fields in this file named `version`, `name`, `scalaVersion`, and `libraryDependencies` are all SBT keys (and in fact are probably the most common keys).

Discussion

A *managed dependency* is a dependency that's managed by your build tool, in this case, SBT. In this situation, if library *a.jar* depends on *b.jar*, and that library depends on *c.jar*, and those JAR files are kept in an Ivy/Maven repository along with this relationship information, then all you have to do is add a line to your *build.sbt* file stating that you want to use *a.jar*. The other JAR files will be downloaded and included into your project automatically.

When using a library as an *unmanaged dependency*, you have to manage this situation yourself. Given the same situation as the previous paragraph, if you want to use the library *a.jar* in your project, you must manually download *a.jar*, and then know about the dependency on *b.jar*, and the transitive dependency on *c.jar*, then download all those files yourself, and place them in your project's *lib* directory.

I've found that manually managing JAR files in the *lib* directory works fine for small projects, but as shown in [Recipe 16.2, "Connecting to a Database with the Spring Framework"](#), a few lines of managed dependency declarations can quickly explode into a large number of JAR files you'll need to manually track down and add to your *lib* folder.

Under the covers, SBT uses Apache Ivy as its dependency manager. Ivy is also used by Ant and Maven, and as a result, you can easily use the wealth of Java libraries that have been created over the years in your Scala projects.

There are two general forms for adding a managed dependency to a *build.sbt* file. In the first form, you specify the `groupId`, `artifactID`, and `revision`:

```
libraryDependencies += groupId % artifactID % revision
```

In the second form, you add an optional configuration parameter:

```
libraryDependencies += groupId % artifactID % revision % configuration
```

The `groupId`, `artifactID`, `revision`, and `configuration` strings correspond to what Ivy requires to retrieve the module you want. Typically, the module developer will give you the information you need. For instance, the `specs2` website provides this string:

```
libraryDependencies += "org.specs2" %% "specs2" % "1.14" % "test"
```

It also provides this information, which shows how to use the same library with Maven:

```
<dependency>
  <groupId>org.specs2</groupId>
  <artifactId>specs2_2.10</artifactId>
  <version>1.14</version>
  <scope>test</scope>
</dependency>
```

The symbols `+=`, `%`, and `%%` used in *build.sbt* are part of the DSL defined by SBT. They’re described in [Table 18-2](#).

Table 18-2. Common methods used in a *build.sbt* file

Method	Description
<code>+=</code>	Appends to the key’s value. The <i>build.sbt</i> file works with settings defined as key/value pairs. In the examples shown, <code>libraryDependencies</code> is a key, and it’s shown with several different values.
<code>%</code>	A method used to construct an Ivy Module ID from the strings you supply.
<code>%%</code>	When used after the <code>groupId</code> , it automatically adds your project’s Scala version (such as <code>_2.10</code>) to the end of the artifact name.

As shown in the examples, you can use `%` or `%%` after the `groupId`. This example shows the `%` method:

```
libraryDependencies += "org.scalatest" % "scalatest_2.10" % "1.9.1" % "test"
```

This example shows the `%%` method:

```
libraryDependencies += "org.scalatest" %% "scalatest" % "1.9.1" % "test"
```

When using Scala 2.10, these two lines are equivalent. The `%%` method adds your project’s Scala version to the end of the artifact name. The practice of adding the Scala version (in the format `_2.10.0`) to the `artifactID` is used because modules may be compiled for different Scala versions.

Note that in some of the examples, the string `test` is added after the `revision`:

```
"org.scalatest" % "scalatest_2.10" % "1.9.1" % "test"
```

This demonstrates the use of the “configuration” form for adding a dependency that was shown earlier:

```
libraryDependencies += groupId % artifactID % revision % configuration
```

As the SBT documentation states, this means that the dependency you’re defining “will be added to the classpath only for the Test configuration, and won’t be added in the

Compile configuration.” This is useful for adding dependencies like `ScalaTest`, `specs2`, `Mockito`, etc., that will be used when you want to test your application, but not when you want to compile and run the application.

If you’re not familiar with Apache Ivy, it can be helpful to know that managed dependencies are downloaded beneath a `.ivy2` directory in your home directory (`~/.ivy2/`) on your filesystem. See the [Ivy documentation](#) for more information.

Repositories

SBT uses the standard Maven2 repository by default, so it can locate most libraries when you add a `libraryDependencies` line to a `build.sbt` file. In these cases, there’s no need for you to tell SBT where to look for the file. However, when a library is not in a standard repository, you can tell SBT where to look for it. This process is referred to as adding a resolver, and it’s covered in [Recipe 18.11, “Telling SBT How to Find a Repository \(Working with Resolvers\)”](#).

See Also

- [Apache Ivy](#).
- The [SBT Quick Configuration documentation](#) shows dozens of `build.sbt` examples.
- [Recipe 18.11, “Telling SBT How to Find a Repository \(Working with Resolvers\)”](#).

18.5. Controlling Which Version of a Managed Dependency Is Used

Problem

You want to make sure you always have the desired version of a managed dependency, including the latest integration release, milestone release, or other versions.

Solution

The revision field in the `libraryDependencies` setting isn’t limited to specifying a single, fixed version. According to the Apache Ivy documentation, you can specify terms such as `latest.integration`, `latest.milestone`, and other terms.

As one example of this flexibility, rather than specifying version 1.8 of a `foobar` module, as shown here:

```
libraryDependencies += "org.foobar" %% "foobar" % "1.8"
```

you can request the `latest.integration` version like this:

```
libraryDependencies += "org.foobar" %% "foobar" % "latest.integration"
```

The module developer will often tell you what versions are available or should be used, and Ivy lets you specify tags to control which version of the module will be downloaded and used. The Ivy “dependency” documentation states that the following tags can be used:

- `latest.integration`
- `latest.[any status]`, such as `latest.milestone`
- You can end the revision with a `+` character. This selects the latest subrevision of the dependency module. For instance, if the dependency module exists in revisions `1.0.3`, `1.0.7`, and `1.1.2`, specifying `1.0.+` as your dependency will result in `1.0.7` being selected.
- You can use “version ranges,” as shown in the following examples:

```
[1.0,2.0] matches all versions greater or equal to 1.0 and lower or equal to 2.0
[1.0,2.0[ matches all versions greater or equal to 1.0 and lower than 2.0
]1.0,2.0] matches all versions greater than 1.0 and lower or equal to 2.0
]1.0,2.0[ matches all versions greater than 1.0 and lower than 2.0
[1.0,) matches all versions greater or equal to 1.0
]1.0,) matches all versions greater than 1.0
(,2.0] matches all versions lower or equal to 2.0
(,2.0[ matches all versions lower than 2.0
```

(These configuration examples are courtesy of the Apache Ivy documentation. See the link in the See Also section for more information.)

To demonstrate a few of these tags, this example shows the `latest.milestone` tag:

```
libraryDependencies += "org.scalatest" %% "scalatest" % "latest.milestone" % ↵
"test"
```

At the time of this writing, it retrieves this file:

```
scalatest_2.10-2.0.M6-SNAP13.jar
```

This specification demonstrates the `+` tag:

```
libraryDependencies += "org.scalatest" %% "scalatest" % "1.9.+" % "test"
```

It currently retrieves this file:

```
scalatest_2.10-1.9.2-SNAP1.jar
```

See Also

[Apache Ivy revision documentation](#)

18.6. Creating a Project with Subprojects

Problem

You want to configure SBT to work with a main project that depends on other subprojects you're developing.

Solution

Create your subproject as a regular SBT project, but without a *project* subdirectory. Then, in your main project, define a *project/Build.scala* file that defines the dependencies between the main project and subprojects.

This is demonstrated in the following example, which I created based on [the SBT Multi-Project documentation](#):

```
import sbt._
import Keys._

/**
 * based on http://www.scala-sbt.org/release/docs/Getting-Started/Multi-Project
 */
object HelloBuild extends Build {

  // aggregate: running a task on the aggregate project will also run it
  // on the aggregated projects.
  // dependsOn: a project depends on code in another project.
  // without dependsOn, you'll get a compiler error: "object bar is not a
  // member of package com.alvinalexander".
  lazy val root = Project(id = "hello",
    base = file(".")) aggregate(foo, bar) dependsOn(foo, bar)

  // sub-project in the Foo subdirectory
  lazy val foo = Project(id = "hello-foo",
    base = file("Foo"))

  // sub-project in the Bar subdirectory
  lazy val bar = Project(id = "hello-bar",
    base = file("Bar"))
}
```

To create your own example, you can either follow the instructions in the [SBT Multi-Project documentation](#) to create a main project with subprojects, or clone [my SBT Subproject Example](#) on GitHub, which I created to help you get started quickly.

Discussion

Creating a main project with subprojects is well documented on the SBT website, and the primary glue that defines the relationships between projects is the *project/Build.scala* file you create in your main project.

In the example shown, my main project depends on two subprojects, which are in directories named *Foo* and *Bar* beneath my project's main directory. I reference these projects in the following code in my main project, so it's necessary to tell SBT about the relationship between the projects:

```
package com.alvinalexander.subprojecttests

import com.alvinalexander.bar._
import com.alvinalexander.foo._

object Hello extends App {
  println(Bar("I'm a Bar"))
  println(Bar("I'm a Foo"))
}
```

The following output from the Unix `tree` command shows what the directory structure for my project looks like, including the files and directories for the main project, and the two subprojects:

```
|-- Bar
|   |-- build.sbt
|   +-- src
|       |-- main
|       |   |-- java
|       |   |-- resources
|       |   +-- scala
|       |       +-- Bar.scala
|       +-- test
|           |-- java
|           +-- resources
|-- Foo
|   |-- build.sbt
|   +-- src
|       |-- main
|       |   |-- java
|       |   |-- resources
|       |   +-- scala
|       |       +-- Foo.scala
|       +-- test
|           |-- java
|           +-- resources
|-- build.sbt
|-- project
|   |-- Build.scala
|
```

```

+-- src
|  |-- main
|  |  |-- java
|  |  |-- resources
|  |  +-- scala
|  |      +-- Hello.scala
+-- test
|  |-- java
|  |-- resources
|  +-- scala
|      +-- HelloTest.scala

```

To experiment with this yourself, I encourage you to clone [my GitHub project](#).

See Also

- [SBT Multi-Project documentation](#)
- My example “SBT Subprojects” code at [GitHub](#)

18.7. Using SBT with Eclipse

Problem

You want to use Eclipse with a project you’re managing with SBT.

Solution

Use the [Scala IDE for Eclipse project](#) so you can work on Scala projects in Eclipse, and use the [sbteclipse plug-in](#) to enable SBT to generate files for Eclipse.

The Scala IDE for Eclipse project lets you edit Scala code in Eclipse. With syntax highlighting, code completion, debugging, and many other features, it makes Scala development in Eclipse a pleasure.

To use the sbteclipse plug-in, download it per the instructions on the website. Once installed, when you’re in the root directory of an SBT project, type `sbt eclipse` to generate the files Eclipse needs. You may see a lot of output the first time you run the command as SBT checks everything it needs, but at the end of the output you should see a “success” message, like this:

```
$ sbt eclipse
```

```
[info] Successfully created Eclipse project files for project(s):
[info] YourProjectNameHere
```

The plug-in generates the two files Eclipse needs, the `.classpath` and `.project` files.

Once these files are generated, go to Eclipse and follow the usual steps to import a project into the Eclipse workspace: File → Import → Existing Projects into Workspace. Your project will then appear in the Eclipse Navigator, Project Explorer, Package Explorer, and other views.

Discussion

The *.classpath* file is an XML file that contains a number of `<classpathentry>` tags, like this:

```
<classpath>
  <classpathentry output="target/scala-2.10/classes"
    path="src/main/scala" kind="src"></classpathentry>
  <classpathentry output="target/scala-2.10/classes"
    path="src/main/java" kind="src"></classpathentry>
  <classpathentry output="target/scala-2.10/test-classes"
    path="src/test/scala" kind="src"></classpathentry>
  <classpathentry output="target/scala-2.10/test-classes"
    path="src/test/java" kind="src"></classpathentry>
  <classpathentry kind="con"
    path="org.scala-ide.sdt.launching.SCALA_CONTAINER"></classpathentry>
  <classpathentry
    path="/Users/AL/.ivy2/cache/com.typesafe/config/bundles/config-1.0.0.jar"
    kind="lib"></classpathentry>
  <classpathentry path="org.eclipse.jdt.launching.JRE_CONTAINER"
    kind="con"></classpathentry>
  <classpathentry path="bin" kind="output"></classpathentry>
</classpath>
```

The *.project* file is an XML file that describes your project and looks like this:

```
<projectDescription>
  <name>YourProjectName</name>
  <buildSpec>
    <buildCommand>
      <name>org.scala-ide.sdt.core.scalabuilder</name>
    </buildCommand>
  </buildSpec>
  <natures>
    <nature>org.scala-ide.sdt.core.scalanature</nature>
    <nature>org.eclipse.jdt.core.javanature</nature>
  </natures>
</projectDescription>
```

Any time you update your SBT build definition files (*build.sbt*, *project/*.scala*, *project/*.sbt*) you should rerun the `sbt eclipse` command to update the *.classpath* and *.project* files. Eclipse will also need to know that these files were regenerated, so this is really a two-step process:

- Run `sbt eclipse` from the command line.
- In Eclipse, select your project and then refresh it (using the F5 function key, or refreshing it with the menu commands).

See Also

- The [Scala IDE for Eclipse](#)
- The [sbteclipse plug-in](#)
- [JetBrains](#) also has plug-ins for IntelliJ IDEA

18.8. Generating Project API Documentation

Problem

You’ve marked up your source code with Scaladoc comments, and want to generate the API documentation for your project.

Solution

Use any of the commands listed in [Table 18-3](#), depending on your needs.

Table 18-3. Descriptions of SBT commands that generate project documentation

SBT command	Description
<code>doc</code>	Creates Scaladoc API documentation from the Scala source code files located in <i>src/main/scala</i> .
<code>test:doc</code>	Creates Scaladoc API documentation from the Scala source code files located in <i>src/test/scala</i> .
<code>package-doc</code>	Creates a JAR file containing the API documentation created from the Scala source code in <i>src/main/scala</i> .
<code>test:package-doc</code>	Creates a JAR file containing the API documentation created from the Scala source code in <i>src/test/scala</i> .
<code>publish</code>	Publishes artifacts to the repository defined by the <code>publish - to</code> setting. See Recipe 18.15 , “Publishing Your Library”.
<code>publish-local</code>	Publishes artifacts to the local Ivy repository as described. See Recipe 18.15 , “Publishing Your Library”.

For example, to generate API documentation, use the `doc` command:

```
$ sbt doc
```

At the time of this writing, SBT doesn't show where the output from this command is written to, but with Scala 2.10.0, SBT 0.12.3 places the root *index.html* Scaladoc file at *target/scala-2.10/api/index.html* under the root directory of your project. Other commands, including `package-doc` and `publish`, do indicate where their output is located.

The following example shows that `publish-local` generates its output for a project named “Hello” to the *.ivy2* directory under your *\$HOME* directory:

```
> sbt publish-local
[info] Loading global plugins from /Users/Al/.sbt/plugins
$HOME/.ivy2/local/hello/hello_2.10/1.0/poms/hello_2.10.pom
$HOME/.ivy2/local/hello/hello_2.10/1.0/jars/hello_2.10.jar
$HOME/.ivy2/local/hello/hello_2.10/1.0/srcs/hello_2.10-sources.jar
$HOME/.ivy2/local/hello/hello_2.10/1.0/docs/hello_2.10-javadoc.jar
$HOME/.ivy2/local/hello/hello_2.10/1.0/ivys/ivy.xml
```

See [Recipe 18.15, “Publishing Your Library”](#) for examples of how to use `publish` and `publish-local`.

For a detailed example of how to use Scaladoc, see [Recipe 14.8, “Generating Documentation with scaladoc”](#).

See Also

- The [SBT command-line reference](#) has more information on these commands
- When writing Scaladoc, you can use a [Wiki-like syntax](#)
- The Scaladoc tags (`@see`, `@param`, etc.) are listed in the [Scala wiki](#)
- [Recipe 14.8, “Generating Documentation with scaladoc”](#) provides more examples of the documentation publishing commands
- See [Recipe 18.15, “Publishing Your Library”](#) for examples of using `publish` and `publish-local`

18.9. Specifying a Main Class to Run

Problem

You have multiple `main` methods in objects in your project, and you want to specify which `main` method should be run when you type `sbt run`, or specify the `main` method that should be invoked when your project is packaged as a JAR file.

Solution

If you have multiple `main` methods in your project and want to specify which `main` method to run when typing `sbt run`, add a line like this to your *build.sbt* file:

```
// set the main class for 'sbt run'
mainClass in (Compile, run) := Some("com.alvinalexander.Foo")
```

This class can either contain a `main` method, or extend the `App` trait.

To specify the class that will be added to the manifest when your application is packaged as a JAR file, add this line to your *build.sbt* file:

```
// set the main class for packaging the main jar
mainClass in (Compile, packageBin) := Some("com.alvinalexander.Foo")
```

That setting tells SBT to add the following line to the *META-INF/MANIFEST.MF* file in your JAR when you run `sbt package`:

```
Main-Class: com.alvinalexander.Foo
```

Using run-main

When running your application with SBT, you can also use SBT's `run-main` command to specify the class to run. Invoke it like this from your operating system command line:

```
$ sbt "run-main com.alvinalexander.Foo"
[info] Loading global plugins from /Users/Al/.sbt/plugins
[info] Running com.alvinalexander.Foo
hello
[success] Total time: 1 s
```

Invoke it like this from inside the SBT shell:

```
$ sbt

> run-main com.alvinalexander.Foo
[info] Running com.alvinalexander.Foo
hello
[success] Total time: 1 s
```

Discussion

If you have only one `main` method in an object in your project (or one object that extends the `App` trait), SBT can automatically determine the location of that `main` method. In that case, these configuration lines aren't necessary.

If you have multiple `main` methods in your project and don't use any of the approaches shown in the Solution, SBT will prompt you with a list of objects it finds that have a `main` method or extend the `App` trait when you execute `sbt run`:

Multiple main classes detected, select one to run:

```
[1] com.alvinalexander.testproject.Foo
[2] com.alvinalexander.testproject.Bar
```

The following code shows what a *build.sbt* file with both of the `mainClass` settings looks like:

```
name := "Simple Test Project"

version := "1.0"

scalaVersion := "2.10.0"

// set the main class for packaging the main jar
mainClass in (Compile, packageBin) := Some("com.alvinalexander.testproject.Foo")

// set the main class for the main 'sbt run' task
mainClass in (Compile, run) := Some("com.alvinalexander.testproject.Foo")
```

See Also

The [SBT Quick Configuration documentation](#) shows dozens of *build.sbt* examples.

18.10. Using GitHub Projects as Project Dependencies

Problem

You want to use a Scala library project on GitHub as an SBT project dependency.

Solution

Reference the GitHub project you want to include in your *project/Build.scala* file as a `RootProject`.

For example, assuming you want to use the Scala project at <https://github.com/alvinj/SoundFilePlayer> as a dependency, put the following contents in a file named *project/Build.scala* in your SBT project:

```
import sbt._

object MyBuild extends Build {

  lazy val root = Project("root", file(".")) dependsOn(soundPlayerProject)
  lazy val soundPlayerProject =
    RootProject(uri("git://github.com/alvinj/SoundFilePlayer.git"))

}
```

You can now use that library in your code, as shown in this little test program:

```
package githubtest

import com.alvinalexander.sound._
import javazoom.jlgui.basicplayer._
import scala.collection.JavaConversions._
import java.util.Map

object TestJavaSound extends App {

  val testClip = "/Users/al/Sarah/Sounds/HAL-mission-too-important.wav"
  val player = SoundFilePlayer.getSoundFilePlayer(testClip)
  player.play

}
```

With this configuration and a basic *build.sbt* file, you can run this code as usual with the `sbt run` command.

Including this GitHub project is interesting, because it has a number of JAR files in its own *lib* folder, and compiling and running this example works fine.

Note that although this works well for compiling and running your project, you can't package all of this code into a JAR file by just using the `sbt package` command. Unfortunately, SBT doesn't include the code from the GitHub project for you. However, a plug-in named `sbt-assembly` does let you package all of this code together as a single JAR file. See [Recipe 18.14, "Deploying a Single, Executable JAR File"](#) for information on how to configure and use `sbt-assembly`.

Discussion

Whereas the *build.sbt* file is used to define simple settings for your SBT project, the *project/Build.scala* file is used for "everything else." In this file you write Scala code using the SBT API to accomplish any other task you want to achieve, such as including GitHub projects like this.

To use multiple GitHub projects as dependencies, add additional `RootProject` instances to your *project/Build.scala* file:

```
import sbt._

object MyBuild extends Build {

  lazy val root = Project("root", file("."))
    .dependsOn(soundPlayerProject)
    .dependsOn(appleScriptUtils)

  lazy val soundPlayerProject =
    RootProject(uri("git://github.com/alvinj/SoundFilePlayer.git"))
```



```

lazy val appleScriptUtils =
  RootProject(uri("git://github.com/alvinj/AppleScriptUtils.git"))
}

```

See Also

Recipe 18.6, “Creating a Project with Subprojects”, and Recipe 18.16, “Using *Build.scala* Instead of *build.sbt*”, show other examples of the *project/Build.scala* file.

18.11. Telling SBT How to Find a Repository (Working with Resolvers)

Problem

You want to add a managed dependency to your project from an Ivy repository that SBT doesn’t know about by default.

Solution

Use the `resolvers` key in the *build.sbt* file to add any unknown Ivy repositories. Use this syntax to add one resolver:

```

resolvers += "Java.net Maven2 Repository" at ↵
"http://download.java.net/maven/2/"

```

You can use a `Seq` to add multiple resolvers:

```

resolvers += Seq(
  "Typesafe" at "http://repo.typesafe.com/typesafe/releases/",
  "Java.net Maven2 Repository" at "http://download.java.net/maven/2/"
)

```

Or, if you prefer, you can also add them one line at a time, making sure to separate them by a blank line:

```

resolvers += "Typesafe" at "http://repo.typesafe.com/typesafe/releases/"

resolvers += "Java.net Maven2 Repository" at ↵
"http://download.java.net/maven/2/"

```

Discussion

If the module you’re requesting is in the default **Maven2 repository SBT knows about**, adding a managed dependency “just works.” But if the module isn’t there, the library’s author will need to provide you with the repository information.

You define a new repository in the *build.sbt* file with this general format:

```
resolvers += "repository name" at "location"
```

As shown in the Solution, you can enter one resolver at a time with the += method, and you can add multiple resolvers with ++= and a Seq.

In addition to the default Maven2 repository, SBT is configured to know about the **JavaNet1Repository**. To use this repository in your SBT project, add this line to your *build.sbt* file:

```
resolvers += JavaNet1Repository
```

18.12. Resolving Problems by Getting an SBT Stack Trace

Problem

You're trying to use SBT to compile, run, or package a project, and it's failing, and you need to be able to see the stack trace to understand why it's failing.

Solution

When an SBT command silently fails (typically with a “Nonzero exit code” message), but you can't tell why, run your command from within the SBT shell, then use the `last run` command after the command that failed.

This pattern typically looks like this:

```
$ sbt run           // something fails here, but you can't tell what

$ sbt
> run              // failure happens again
> last run          // this shows the full stack trace
```

I've run into this on several projects where I was using JAR files and managing their dependencies myself, and in one specific case, I didn't know I needed to include the Apache Commons Logging JAR file. This was causing the “Nonzero exit code” error message, but I couldn't tell that until I issued the `last run` command from within the SBT shell. Once I ran that command, the problem was obvious from the stack trace.

Depending on the problem, another approach that can be helpful is to set the SBT logging level. See [Recipe 18.13, “Setting the SBT Log Level”](#) for more information.

18.13. Setting the SBT Log Level

Problem

You're having a problem compiling, running, or packaging your project with SBT and need to adjust the SBT logging level to debug the problem. (Or, you're interested in learning about how SBT works.)

Solution

Set the SBT logging level in your *build.sbt* file with this setting:

```
logLevel := Level.Debug
```

Or, if you're working interactively from the SBT command line and don't want to add this to your *build.sbt* file, use this syntax:

```
> set logLevel := Level.Debug
```

Changing the logging levels significantly changes the output SBT produces, which can help you debug problems. If you're just starting out with SBT, the output can also help you learn how SBT works.

Other logging levels are:

- `Level.Info`
- `Level.Warning`
- `Level.Error`

See Also

The [SBT FAQ](#) shows the logging levels.

18.14. Deploying a Single, Executable JAR File

Problem

You're building a Scala application, such as a Swing application, and want to deploy a single, executable JAR file to your users.

Solution

The `sbt package` command creates a JAR file that includes the class files it compiles from your source code, along with the resources in your project (from `src/main/resources`), but there are two things it doesn't include in the JAR file:

- Your project dependencies (JAR files in your project's *lib* folder, or managed dependencies declared in *build.sbt*).
- Libraries from the Scala distribution that are needed to execute the JAR file with the `java` command.

This makes it difficult to distribute a single, executable JAR file for your application. There are three things you can do to solve this problem:

- Distribute all the JAR files necessary with a script that builds the classpath and executes the JAR file with the `scala` command. This requires that Scala be installed on client systems.
- Distribute all the JAR files necessary (including Scala libraries) with a script that builds the classpath and executes the JAR file with the `java` command. This requires that Java is installed on client systems.
- Use an SBT plug-in such as `sbt-assembly` to build a single, complete JAR file that can be executed with a simple `java` command. This requires that Java is installed on client systems.

This solution focuses on the third approach. The first two approaches are examined in the Discussion.

Using `sbt-assembly`

The installation instructions for `sbt-assembly` may change, but at the time of this writing, just add these two lines of code to a *plugins.sbt* file in the *project* directory of your SBT project:

```
resolvers += Resolver.url("artifactory",  
  url("http://scalasbt.artifactoryonline.com/scalasbt/sbt-plugin-releases"))←  
(Resolver.ivyStylePatterns)  
  
addSbtPlugin("com.eed3si9n" % "sbt-assembly" % "0.8.4")
```

You'll need to create that file if it doesn't already exist.

Then add these two lines to the top of your *build.sbt* file:

```
import AssemblyKeys._  
  
// sbt-assembly  
assemblySettings
```

That's the only setup work that's required. Now run `sbt assembly` to create your single, executable JAR file:

```
$ sbt assembly
```

When the `assembly` task finishes running it will tell you where the executable JAR file is located. For instance, when packaging **my Blue Parrot application**, SBT prints the following lines of output that show the dependencies `sbt-assembly` is including, and the location of the final JAR file:

```
[info] Including akka-actor-2.0.1.jar
[info] Including scala-library.jar
[info] Including applescriptutils_2.9.1-1.0.jar
[info] Including forms-1.0.7.jar
[info] Including sounds_2.9.1-1.0.jar
[info] Packaging target/BlueParrot-assembly-1.0.jar ...
[info] Done packaging.
```

The `sbt-assembly` plug-in works by copying the class files from your source code, the class files from your dependencies, and the class files from the Scala library into one single JAR file that can be executed with the `java` interpreter. This can be important if there are license restrictions on a JAR file, for instance.

As noted, there are other plug-ins to help solve this problem, including `One-JAR`, but `sbt-assembly` worked best with several applications I've deployed as single, executable JAR files.

Discussion

A JAR file created by SBT can be run by the Scala interpreter, but not the Java interpreter. This is because class files in the JAR file created by `sbt package` have dependencies on Scala class files (Scala libraries), which aren't included in the JAR file SBT generates. This is easily demonstrated.

First, create an empty SBT project directory. (See [Recipe 18.1](#) for easy ways to do this.)

Then place the following code in a file named *Main.scala* in the root directory of the project:

```
package foo.bar.baz

object Main extends App {
  println("Hello, world")
}
```

Next, run `sbt package` to create the JAR file:

```
$ sbt package
[info] Loading global plugins from /Users/Al/.sbt/plugins
[info] Done updating.
[info] Compiling 1 Scala source to target/scala-2.10/classes...
```

```
[info] Packaging target/scala-2.10/basic_2.10-1.0.jar ...
[info] Done packaging.
[success] Total time: 6 s
```

Now attempt to run the JAR file with the `java -jar` command. This will fail:

```
$ java -jar target/scala-2.10/basic_2.10-1.0.jar
Exception in thread "main" java.lang.NoClassDefFoundError: scala/App
  at java.lang.ClassLoader.defineClass1(Native Method)
  ... 32 more
```

This fails because the Java interpreter doesn't know where the `scala/App` trait is.

Next, demonstrate that you can run the same JAR file with the Scala interpreter:

```
$ scala target/scala-2.10/basic_2.10-1.0.jar
Hello, world
```

This works fine.

For the Java interpreter to run your JAR file, it needs the `scala-library.jar` file from your Scala installation to be on its classpath. You can get this example to work with Java by including that JAR file on its classpath with this command:

```
$ java -cp "${CLASSPATH}:${SCALA_HOME}/lib/scala-library.jar:target/scala-2.10-1.0/basic_2.10-1.0.jar" foo.bar.baz.Main
Hello, world
```

As shown, adding the `scala-library.jar` file lets the Java interpreter find the `scala/App` trait (which is a normal `.class` file), which lets it run the application successfully for you.

This is part of the work that `sbt-assembly` performs for you. It repackages the class files from `${SCALA_HOME}/lib/scala-library.jar` into your single, executable JAR file, and does the same thing with your other project dependencies. Note that if your application is more complicated, it may need additional JAR files from the `${SCALA_HOME}/lib` directory.

See Also

- The [sbt-assembly project](#).
- My [Blue Parrot application](#) is written in Scala, and packaged with SBT and `sbt-assembly`.
- The [One-JAR project](#).

18.15. Publishing Your Library

Problem

You've created a Scala project or library with SBT that you want to share with other users, creating all the files you need for an Ivy repository.

Solution

Define your repository information, then publish your project with `sbt publish` or `sbt publish-local`.

For my `SoundFilePlayer` library, I added this setting to my `build.sbt` file to define the location of my local repository:

```
publishTo := Some(Resolver.file("file", new File("/Users/al/tmp")))
```

I then ran `sbt publish`, and SBT generated the following files:

```
$ sbt publish
```

```
[info] Wrote
/Users/al/SoundFilePlayer/target/scala-2.10.0/sounds_2.10.0-1.0.pom
[info] :: delivering :: default#sounds_2.10.0;1.0 :: 1.0 :: release ::
[info] delivering ivy file to
/Users/al/SoundFilePlayer/target/scala-2.10.0/ivy-1.0.xml
[info] published sounds_2.10.0 to
/Users/al/tmp/default/sounds_2.10.0/1.0/sounds_2.10.0-1.0.pom
[info] published sounds_2.10.0 to
/Users/al/tmp/default/sounds_2.10.0/1.0/sounds_2.10.0-1.0.jar
[info] published sounds_2.10.0 to
/Users/al/tmp/default/sounds_2.10.0/1.0/sounds_2.10.0-1.0-sources.jar
[info] published sounds_2.10.0 to
/Users/al/tmp/default/sounds_2.10.0/1.0/sounds_2.10.0-1.0-javadoc.jar
[success] Total time: 1s
```

Without doing anything to define a “local Ivy repository,” I get the following results when running the `publish-local` task:

```
$ sbt publish-local
```

```
[info] Wrote /Users/al/SoundFilePlayer/target/scala-2.10.0/sounds_2.10.0-1.0.pom
[info] :: delivering :: default#sounds_2.10.0;1.0 :: 1.0 :: release ::
[info] delivering ivy file to
/Users/al/SoundFilePlayer/target/scala-2.10.0/ivy-1.0.xml
[info] published sounds_2.10.0 to
/Users/al/.ivy2/local/default/sounds_2.10.0/1.0/poms/sounds_2.10.0.pom
[info] published sounds_2.10.0 to
/Users/al/.ivy2/local/default/sounds_2.10.0/1.0/jars/sounds_2.10.0.jar
[info] published sounds_2.10.0 to
/Users/al/.ivy2/local/default/sounds_2.10.0/1.0/srcs/sounds_2.10.0-sources.jar
```

```
[info] published sounds_2.10.0 to
/Users/al/.ivy2/local/default/sounds_2.10.0/1.0/docs/sounds_2.10.0-javadoc.jar
[info] published ivy to
/Users/al/.ivy2/local/default/sounds_2.10.0/1.0/ivys/ivy.xml
[success] Total time: 1 s,
```

The “SBT Publishing” documentation provides these descriptions of the `publish` and `publish-local` tasks:

- The `publish` action is used to publish your project to a remote repository. To use publishing, you need to specify the repository to publish to and the credentials to use. Once these are set up, you can run `publish`.
- The `publish-local` action is used to publish your project to a local Ivy repository. You can then use this project from other projects on the same machine.

For more information on publishing to remote servers, repositories, and artifacts, see the [SBT Publishing documentation](#).

18.16. Using *Build.scala* Instead of *build.sbt*

Problem

You want to use the *project/Build.scala* file instead of *build.sbt* to define your Scala project, or you need some examples of how to use *Build.scala* to solve build problems that can't be handled in *build.sbt*.

Solution

The recommended approach when using SBT is to define all your simple settings (key/value pairs) in the *build.sbt* file, and handle all other work, such as build logic, in the *project/Build.scala* file. However, it can be useful to use only the *project/Build.scala* file to learn more about how it works.

To demonstrate this, don't create a *build.sbt* file in your project, and then do create a *Build.scala* file in the *project* subdirectory by extending the SBT Build object:

```
import sbt._
import Keys._

object ExampleBuild extends Build {

  val dependencies = Seq(
    "org.scalatest" %% "scalatest" % "1.9.1" % "test"
  )

  lazy val exampleProject = Project("SbtExample", file(".")) settings(
    version      := "0.2",
```



```

    scalaVersion := "2.10.0",
    scalacOptions := Seq("-deprecation"),
    libraryDependencies ++= dependencies
  )
}

```

With just this *Build.scala* file, you can now run all the usual SBT commands in your project, including `compile`, `run`, `package`, and so on.

Discussion

The *Build.scala* file shown in the Solution is equivalent to the following *build.sbt* file:

```

name := "SbtExample"

version := "0.2"

scalaVersion := "2.10.0"

scalacOptions += "-deprecation"

libraryDependencies += "org.scalatest" %% "scalatest" % "1.9.1" % "test"

```

As mentioned, the recommended approach when working with SBT is to define your basic settings in the *build.sbt* file, and perform all other work in a *Build.scala* file, so creating a *Build.scala* file with only settings in it is not a best practice. However, when you first start working with a *Build.scala* file, it's helpful to see a “getting started” example like this.

Also, although the convention is to name this file *Build.scala*, this is only a convention, which I use here for simplicity. You can give your build file any legal Scala filename, as long as you place the file in the *project* directory with a *.scala* suffix. Another convention is to name this file after the name of your project, so the **Scalaz project** uses the name *ScalazBuild.scala*.

The Full Configuration Example in the SBT documentation

The **Full Configuration Example** in the SBT documentation and the *ScalazBuild.scala* build file both show *many* more examples of what can be put in a *Build.scala* file. For instance, the Full Configuration Example shows how to add a series of resolvers to a project:

```

// build 'oracleResolvers'
object Resolvers {
  val sunrepo = "Sun Maven2 Repo" at "http://download.java.net/maven/2"
  val sunrepoGF = "Sun GF Maven2 Repo" at
    "http://download.java.net/maven/glassfish"
  val oraclerepo = "Oracle Maven2 Repo" at "http://download.oracle.com/maven"
  val oracleResolvers = Seq (sunrepo, sunrepoGF, oraclerepo)
}

```

```

}

object CDAP2Build extends Build {
  import Resolvers._
  // more code here ...

  // use 'oracleResolvers' here
  lazy val server = Project (
    "server",
    file ("cdap2-server"),
    settings = buildSettings ++ Seq (resolvers := oracleResolvers,
                                     libraryDependencies += serverDeps)
  ) dependsOn (common)
}

```

This code is similar to the example shown in [Recipe 18.11, “Telling SBT How to Find a Repository \(Working with Resolvers\)”](#), where the following configuration line is added to a *build.sbt* file:

```
resolvers += "Java.net Maven2 Repository" at "http://download.java.net/maven/2/"
```

The *ScalazBuild.scala* file also shows many examples of using *TaskKey* and *SettingKey*, which are different types of keys that can be used in SBT project definition files.

See Also

- The [Full Configuration Example in the SBT documentation](#).
- The [ScalazBuild.scala](#) file.
- For more examples of using *Build.scala* files, see [Recipe 18.6, “Creating a Project with Subprojects”](#); [Recipe 18.10, “Using GitHub Projects as Project Dependencies”](#); and [Recipe 18.11, “Telling SBT How to Find a Repository \(Working with Resolvers\)”](#).

18.17. Using a Maven Repository Library with SBT

Problem

You want to use a Java library that’s in a Maven repository, but the library doesn’t include information about how to use it with Scala and SBT.

Solution

Translate the Maven *groupId*, *artifactId*, and *version* fields into an SBT *libraryDependencies* string.

For example, I wanted to use the Java HTMLCleaner project in a Scala/SBT project. [The HTMLCleaner website](#) provided the following Maven information, but no SBT information:

```
<dependency>
  <groupId>net.sourceforge.htmlcleaner</groupId>
  <artifactId>htmlcleaner</artifactId>
  <version>2.2</version>
</dependency>
```

Fortunately this translates into the following SBT `libraryDependencies` string:

```
libraryDependencies += "net.sourceforge.htmlcleaner" % "htmlcleaner" % "2.2"
```

After adding this line to my *build.sbt* file, I ran `sbt compile`, and watched as it downloaded the HTMLCleaner JAR file and dependencies:

```
[info] downloading http://repo1.maven.org/maven2/net/sourceforge/htmlcleaner/
htmlcleaner/2.2/htmlcleaner-2.2.jar ...
[info] [SUCCESSFUL ] net.sourceforge.htmlcleaner#htmlcleaner;2.2!htmlcleaner.jar
(864ms)
[info] downloading http://repo1.maven.org/maven2/org/jdom/jdom/1.1/jdom-1.1.jar ...
[info] [SUCCESSFUL ] org.jdom#jdom;1.1!jdom.jar (514ms)
[info] downloading
http://repo1.maven.org/maven2/org/apache/ant/ant/1.7.0/ant-1.7.0.jar ...
[info] [SUCCESSFUL ] org.apache.ant#ant;1.7.0!ant.jar (1997ms)
[info] downloading http://repo1.maven.org/maven2/org/apache/ant/ant-launcher/
1.7.0/ant-launcher-1.7.0.jar ...
[info] [SUCCESSFUL ] org.apache.ant#ant-launcher;1.7.0!ant-launcher.jar (152ms)
[info] Done updating.
[info] Compiling 1 Scala source to target/scala-2.10.0/classes...
[success] Total time: 13 s, completed Aug 10, 2012 9:22:38 PM
```

As mentioned in other recipes, because SBT and Maven both use Apache Ivy under the covers, and SBT also uses the standard Maven2 repository as a default resolver, SBT users can easily use Java libraries that are packaged for Maven.

As shown in [Recipe 18.4, “Managing Dependencies with SBT”](#), there are two formats for adding a `libraryDependencies` line to a *build.sbt* file. The first form was used in the Solution, and its general format looks like this:

```
libraryDependencies += groupId % artifactID % revision
```

As shown with the HTMLCleaner example, the `groupId`, `artifactID`, and `revision` fields correspond directly to the information you’ll find in the documentation for a Maven library.

The second `libraryDependencies` form lets you add an optional configuration parameter:

```
libraryDependencies += groupId % artifactID % revision % configuration
```

Maven doesn't use the term *configuration*, instead using a `<scope>` tag for the same information. This field is optional, and is typically used for testing libraries such as `ScalaTest` and `specs2`, so when it's needed, the value is usually just `test`.

See Also

The [Java HTMLCleaner website](#)

18.18. Building a Scala Project with Ant

Problem

You want to use Ant to build your Scala project.

Solution

Assuming you have a Maven- and SBT-like project directory structure as described in [Recipe 18.1](#), create the following Ant *build.xml* file in the root directory of your project:

```
<project name="AntCompileTest" default="compile" basedir=".">

  <!-- mostly from: http://www.scala-lang.org/node/98 -->

  <property name="sources.dir" value="src" />
  <property name="scala-source.dir" value="main/scala" />
  <property name="scala-test.dir" value="main/test" />
  <property name="build.dir" value="classes" />

  <!-- set scala.home -->
  <property environment="env" />
  <property name="scala.home" value="${env.SCALA_HOME}" />

  <target name="init">
    <property name="scala-library.jar"
      value="${scala.home}/lib/scala-library.jar" />
    <property name="scala-compiler.jar"
      value="${scala.home}/lib/scala-compiler.jar" />
    <property name="scala.reflect"
      value="${scala.home}/lib/scala-reflect.jar" />
    <path id="build.classpath">
      <pathelement location="${scala-library.jar}" />
      <pathelement location="${build.dir}" />
    </path>
    <taskdef resource="scala/tools/ant/antlib.xml">
      <classpath>
        <pathelement location="${scala-compiler.jar}" />
        <pathelement location="${scala-library.jar}" />
        <pathelement location="${scala.reflect}" />
      </classpath>
    </taskdef>
  </target>
</project>
```

```

    </taskdef>
  </target>

  <target name="compile" depends="init">
    <mkdir dir="${build.dir}" />
    <scalac srcdir="${sources.dir}"
            destdir="${build.dir}"
            classpathref="build.classpath"
            deprecation="on">
      <include name="${scala-source.dir}/**/*.scala" />
      <exclude name="${scala-test.dir}/**/*.scala" />
    </scalac>
  </target>

</project>

```

You can then run the usual ant command, which by default will compile your files to a new *classes* folder under the root directory of your project. Running ant on a small project produces output like this:

```

$ ant
Buildfile: /Users/Al/Projects/AntExample/build.xml

init:

compile:
[scalac] Compiling 1 source file to /Users/Al/Projects/AntExample/classes

BUILD SUCCESSFUL
Total time: 5 seconds

```

Discussion

In general, when learning a new technology, it's best to learn the tools of that technology, and in this case, the preferred build tool for Scala projects is SBT. (As a friend once said, when we went from C to Java, we didn't attempt to bring make along with us.) Once you grasp the SBT concepts, you'll find that it's both a simple and powerful tool, and you can find a lot of support in the Scala community.

That being said, you're also hit with a lot of changes when first learning a new technology, and at the beginning, it can be helpful to use the tools you're already comfortable with, so this recipe demonstrates how to use Ant to compile a Scala project to help you get into Scala in a comfortable way.



Recommendation: If someone brought me into their organization to help them adopt Scala, SBT is one of the first things I'd teach. In this case, I think you're better off just diving into the water, so to speak. It doesn't take that long to grasp the SBT basics.

The build.xml code

The secret sauce to this recipe is the `init` target, whose source code can be found on the [official Scala website](#). This target does the work necessary to make the `scalac` Ant task available to you.

As you can see from the code, the `build` target depends on the `init` target, and uses `scalac` to compile all the files in the source directory, while skipping the files in the test directory. Of course that approach is completely optional, and you can adjust it to meet your needs.

The `antlib.xml` file referred to in the `taskdef` tag is shipped with the Scala distribution. You can demonstrate this on a Unix system with the following command:

```
$ jar tvf ${SCALA_HOME}/lib/scala-compiler.jar | grep -i antlib
```

The `build.xml` file shown here is slightly different than the file shown on the Scala website. Specifically, I found that the `scala.home` property needed to be set manually, and with Scala 2.10, it's also necessary to add the `scala.reflect` lines to the build file. The compilation process worked fine with Ant 1.8.4 once I made those changes.

In addition to this `scalac` Ant task, there are `fsc` and `scaladoc` tasks. See [the Scala Ant Tasks page](#) on the official Scala website for more information.

Creating a JAR file with Ant

Once you've compiled your Scala classes, you can treat them as normal Java class files. For instance, you can create a JAR file from them using the following simplified Ant task. This task shows how to create a JAR file named `hello.jar` from the compiled classes in the `classes` directory, and a simple manifest in a `Manifest.txt` file. Here's the `create-jar` task, which you can add to the earlier `build.xml` file:

```
<target name="create-jar" depends="compile">
  <jar basedir="classes"
       jarfile="hello.jar"
       manifest="Manifest.txt"/>
</target>
```

Assuming the `Hello` class in the `hello` package has the `main` method for your application (or extends the `App` trait), place this line in the `Manifest.txt` file:

```
Main-Class: hello.Hello
```

After adding this task to your `build.xml` file, you can run it as follows from the root directory of your project:

```
$ ant create-jar
```

That command creates a JAR file named `hello.jar` in the root directory. You can then run the JAR file with this Scala command:

```
$ scala hello.jar
```

This is similar to running `java -jar` on a JAR file created by a Java application, but because a Scala application has dependencies on its own JAR files, such as `$$SCALA_HOME/lib/scala-library.jar`, you need to run the JAR file with the `scala` interpreter, as shown. You can run the JAR file with the Java interpreter, but this takes a bit more work. See [Recipe 18.14, “Deploying a Single, Executable JAR File”](#) for details on that process.

See Also

The [Scala Ant Tasks documentation](#)

Introduction

As you can tell from one look at the Scaladoc for the collections classes, Scala has a powerful type system. However, unless you're the creator of a library, you can go a long way in Scala without having to go too far down into the depths of Scala types. But once you start creating collections-style APIs for other users, you will need to learn them.

This chapter provides recipes for the most common problems you'll encounter, but when you need to go deeper, I highly recommend the book, *Programming in Scala*, by Odersky, Spoon, and Venners. (Martin Odersky is the creator of the Scala programming language, and I think of that book as “the reference” for Scala.)

Scala's type system uses a collection of symbols to express different generic type concepts, including variance, bounds, and constraints. The most common of these symbols are summarized in the next sections.

Variance

Type variance is a generic type concept, and defines the rules by which parameterized types can be passed into methods. The type variance symbols are briefly summarized in [Table 19-1](#).

Table 19-1. Descriptions of type variance symbols

Symbols	Name	Description
Array[T]	Invariant	Used when elements in the container are mutable. Example: Can <i>only</i> pass Array[String] to a method expecting Array[String].
Seq[+A]	Covariant	Used when elements in the container are immutable. This makes the container more flexible. Example: Can pass a Seq[String] to a method expected Seq[Any].

Symbols	Name	Description
Foo[-A] Function1[-A, +B]	Contravariant	Contravariance is essentially the opposite of covariance, and is rarely used. See Scala's Function1 trait for an example of how it is used.

The following examples, showing what code will and won't compile with the Grandparent, Parent, and Child classes, can also be a helpful reference to understanding variance:

```
class Grandparent
class Parent extends Grandparent
class Child extends Parent

class InvariantClass[A]
class CovariantClass[+A]
class ContravariantClass[-A]

class VarianceExamples {

  def invarMethod(x: InvariantClass[Parent]) {}
  def covarMethod(x: CovariantClass[Parent]) {}
  def contraMethod(x: ContravariantClass[Parent]) {}

  invarMethod(new InvariantClass[Child])           // ERROR - won't compile
  invarMethod(new InvariantClass[Parent])           // success
  invarMethod(new InvariantClass[Grandparent])      // ERROR - won't compile

  covarMethod(new CovariantClass[Child])            // success
  covarMethod(new CovariantClass[Parent])           // success
  covarMethod(new CovariantClass[Grandparent])      // ERROR - won't compile

  contraMethod(new ContravariantClass[Child])       // ERROR - won't compile
  contraMethod(new ContravariantClass[Parent])      // success
  contraMethod(new ContravariantClass[Grandparent]) // success

}
```

Bounds

Bounds let you place restrictions on type parameters. [Table 19-2](#) shows the common bounds symbols.

Table 19-2. Descriptions of Scala's bounds symbols

	Name	Description
A <: B	Upper bound	A must be a subtype of B. See Recipe 19.6 .
A >: B	Lower bound	A must be a supertype of B. Not commonly used. See Recipe 19.8 .
A <: Upper >: Lower	Lower and upper bounds used together	The type A has both an upper and lower bound.

Programming Scala (O'Reilly) had a nice tip that helps me remember these symbols. The authors state that in UML diagrams, subtypes are shown below supertypes, so when I see `A <: B`, I think, “A is less than B ... A is under B ... A is a subtype of B.”

Lower bounds are demonstrated in several methods of the collections classes. To find some lower bound examples, search the Scaladoc of classes like `List` for the `>:` symbol.

There are several additional symbols for bounds. For instance, a *view bound* is written as `A <% B`, and a *context bound* is written as `T : M`. These symbols are not covered in this book; see *Programming in Scala* for details and examples of their use.

Type Constraints

Scala lets you specify additional type constraints. These are written with these symbols:

```
A ::= B    // A must be equal to B
A <:: B    // A must be a subtype of B
A <%% B    // A must be viewable as B
```

These symbols are not covered in this book. See *Programming in Scala* for details and examples. Twitter's Scala School [Advanced Types page](#) also shows brief examples of their use, where they are referred to as “type relation operators.”

Type Examples in Other Chapters

Because types are naturally used in many solutions, you can find some recipes related to types in other chapters:

- [Recipe 2.2, “Converting Between Numeric Types \(Casting\)”](#) and [Recipe 2.3](#) demonstrate ways to convert between types.
- [Recipe 5.9, “Supporting a Fluent Style of Programming”](#) demonstrates how to return `this.type` from a method.
- Implicit conversions let you add new behavior to closed types like `String`, which is declared `final` in Java. They are demonstrated in [Recipe 1.10, “Add Your Own Methods to the String Class”](#) and [Recipe 2.1, “Parsing a Number from a String”](#).
- [Recipe 6.1, “Object Casting”](#) demonstrates how to cast objects from one type to another.

Finally, [Recipe 19.8, “Building Functionality with Types”](#) combines several of the concepts described in this chapter, and also helps to demonstrate Scala's *call-by-name* feature.

19.1. Creating Classes That Use Generic Types

Problem

You want to create a class (and associated methods) that uses a generic type.

Solution

As a library writer, creating a class (and methods) that takes a generic type is similar to Java. For instance, if Scala didn't have a linked list class and you wanted to write your own, you could write the basic functionality like this:

```
class LinkedList[A] {  
  
  private class Node[A] (elem: A) {  
    var next: Node[A] = _  
    override def toString = elem.toString  
  }  
  
  private var head: Node[A] = _  
  
  def add(elem: A) {  
    val n = new Node(elem)  
    n.next = head  
    head = n  
  }  
  
  private def printNodes(n: Node[A]) {  
    if (n != null) {  
      println(n)  
      printNodes(n.next)  
    }  
  }  
  
  def printAll() { printNodes(head) }  
  
}
```

Notice how the generic type `A` is sprinkled throughout the class definition. This is similar to Java, but Scala uses `[A]` everywhere, instead of `<T>` as Java does. (More on the characters `A` versus `T` shortly.)

To create a list of integers with this class, first create an instance of it, declaring its type as `Int`:

```
val ints = new LinkedList[Int]()
```

Then populate it with `Int` values:

```
ints.add(1)  
ints.add(2)
```

Because the class uses a generic type, you can also create a `LinkedList` of `String`:

```
val strings = new LinkedList[String]()
strings.add("Nacho")
strings.add("Libre")
strings.printAll()
```

Or any other type you want to use:

```
val doubles = new LinkedList[Double]()
val frogs = new LinkedList[Frog]()
```

At this basic level, creating a generic class in Scala is just like creating a generic class in Java, with the exception of the brackets.

Discussion

When using generics like this, the container can take subtypes of the base type you specify in your code. For instance, given this class hierarchy:

```
trait Animal
class Dog extends Animal { override def toString = "Dog" }
class SuperDog extends Dog { override def toString = "SuperDog" }
class FunnyDog extends Dog { override def toString = "FunnyDog" }
```

you can define a `LinkedList` that holds `Dog` instances:

```
val dogs = new LinkedList[Dog]
```

You can then add `Dog` subtypes to the list:

```
val fido = new Dog
val wonderDog = new SuperDog
val scooby = new FunnyDog

dogs.add(fido)
dogs.add(wonderDog)
dogs.add(scooby)
```

So far, so good: you can add `Dog` subtypes to a `LinkedList[Dog]`. Where you might run into a problem is when you define a method like this:

```
def printDogTypes(dogs: LinkedList[Dog]) {
  dogs.printAll()
}
```

You can pass your current `dogs` instance into this method, but you won't be able to pass the following `superDogs` collection into `makeDogsSpeak`:

```
val superDogs = new LinkedList[SuperDog]
superDogs.add(wonderDog)

// error: this line won't compile
printDogTypes(superDogs)
```

The last line won't compile because (a) `makeDogsSpeak` wants a `LinkedList[Dog]`, (b) `LinkedList` elements are mutable, and (c) `superDogs` is a `LinkedList[SuperDog]`. This creates a conflict the compiler can't resolve. This situation is discussed in detail in [Recipe 19.5, "Make Immutable Collections Covariant"](#).

In Scala 2.10, the compiler is even nice enough to tell you what's wrong in this situation, and points you toward a solution:

```
[error] Note: SuperDog <: Dog, but class LinkedList is invariant in type A.  
[error] You may wish to define A as +A instead. (SLS 4.5)
```

Type parameter symbols

If a class requires more than one type parameter, use the symbols shown in [Table 19-3](#). For instance, in the [official Java Generics documentation](#), Oracle shows an interface named `Pair`, which takes two types:

```
// from http://docs.oracle.com/javase/tutorial/java/generics/types.html  
public interface Pair<K, V> {  
    public K getKey();  
    public V getValue();  
}
```

You can port that interface to a Scala trait, as follows:

```
trait Pair[A, B] {  
    def getKey: A  
    def getValue: B  
}
```

If you were to take this further and implement the body of a `Pair` class (or trait), the type parameters `A` and `B` would be spread throughout your class, just as the symbol `A` was used in the `LinkedList` example.

The same Oracle document lists the Java type parameter naming conventions. These are mostly the same in Scala, except that Java starts naming simple type parameters with the letter `T`, and then uses the characters `U` and `V` for subsequent types. The Scala standard is that simple types should be declared as `A`, the next with `B`, and so on. This is shown in [Table 19-3](#).

Table 19-3. Standard symbols for generic type parameters

Symbol	Description
A	Refers to a simple type, such as <code>List[A]</code> .
B, C, D	Used for the 2 nd , 3 rd , 4 th types, etc. <pre>// from the Scala Styleguide class List[A] { def map[B](f: A => B): List[B] = ... }</pre>
K	Typically refers to a key in a Java map. Scala collections use <code>A</code> in this situation.

Symbol	Description
N	Refers to a numeric value.
V	Typically refers to a value in a Java map. Scala collections use B in this situation.

See Also

- Oracle’s Java “Generic Types” [documentation](#).
- [Recipe 19.4, “Make Mutable Collections Invariant”](#).
- [Recipe 19.5, “Make Immutable Collections Covariant”](#).
- You can find a little more information on Scala’s generic type naming conventions at the Scala Style Guide’s [Naming Conventions page](#).

19.2. Creating a Method That Takes a Simple Generic Type

Problem

You’re not concerned about type variance, and want to create a method (or function) that takes a generic type, such as a method that accepts a `Seq[A]` parameter.

Solution

As with Scala classes, specify the generic type parameters in brackets, like `[A]`.

For example, when creating a lottery-style application to draw a random name from a list of names, you might follow the “Do the simplest thing that could possibly work” credo, and initially create a method without using generics:

```
def randomName(names: Seq[String]): String = {  
    val randomNum = util.Random.nextInt(names.length)  
    names(randomNum)  
}
```

As written, this works with a sequence of `String` values:

```
val names = Seq("Aleka", "Christina", "Tyler", "Molly")  
val winner = randomName(names)
```

Then, at some point in the future you realize that you could really use a general-purpose method that returns a random element from a sequence of any type. So, you modify the method to use a generic type parameter, like this:

```
def randomElement[A](seq: Seq[A]): A = {  
    val randomNum = util.Random.nextInt(seq.length)  
    seq(randomNum)  
}
```

With this change, the method can now be called on a variety of types:

```
randomElement(Seq("Aleka", "Christina", "Tyler", "Molly"))
randomElement(List(1,2,3))
randomElement(List(1.0,2.0,3.0))
randomElement(Vector.range('a', 'z'))
```

Note that specifying the method's return type isn't necessary, so you can simplify the signature slightly, if desired:

```
// change the return type from ':A =' to just '='
def randomElement[A](seq: Seq[A]) = { ... }
```

Discussion

This is a simple example that shows how to pass a generic collection to a method that doesn't attempt to mutate the collection. See Recipes 19.4 and 19.5 for more complicated situations you can run into.

19.3. Using Duck Typing (Structural Types)

Problem

You're used to "Duck Typing" (structural types) from another language like Python or Ruby, and want to use this feature in your Scala code.

Solution

Scala's version of "Duck Typing" is known as using a *structural type*. As an example of this approach, the following code shows how a `callSpeak` method can require that its `obj` type parameter have a `speak()` method:

```
def callSpeak[A <: { def speak(): Unit }](obj: A) {
  // code here ...
  obj.speak()
}
```

Given that definition, an instance of any class that has a `speak()` method that takes no parameters and returns nothing can be passed as a parameter to `callSpeak`. For example, the following code demonstrates how to invoke `callSpeak` on both a `Dog` and a `Klingon`:

```
class Dog { def speak() { println("woof") } }
class Klingon { def speak() { println("Qapla!") } }

object DuckTyping extends App {

  def callSpeak[A <: { def speak(): Unit }](obj: A) {
    obj.speak()
  }
}
```



```

    }

    callSpeak(new Dog)
    callSpeak(new Klingon)

}

```

Running this code prints the following output:

```

woof
Qapla!

```

The class of the instance that's passed in doesn't matter at all. The only requirement for the parameter `obj` is that it's an instance of a class that has a `speak()` method.

Discussion

The structural type syntax is necessary in this example because the `callSpeak` method invokes a `speak` method on the object that's passed in. In a statically typed language, there must be some guarantee that the object that's passed in will have this method, and this recipe shows the syntax for that situation.

Had the method been written as follows, it wouldn't compile, because the compiler can't guarantee that the type `A` has a `speak` method:

```

// won't compile
def callSpeak[A](obj: A) {
    obj.speak()
}

```

This is one of the great benefits of type safety in Scala.

It may help to break down the structural type syntax. First, here's the entire method:

```

def callSpeak[A <: { def speak(): Unit }](obj: A) {
    obj.speak()
}

```

The type parameter `A` is defined as a structural type like this:

```

[A <: { def speak(): Unit }]

```

The `<:` symbol in the code is used to define something called an *upper bound*. This is described in detail in [Recipe 19.5, “Make Immutable Collections Covariant”](#). As shown in that recipe, an upper bound is usually defined like this:

```

class Stack[A <: Animal] (val elem: A)

```

This states that the type parameter `A` must be a subtype of `Animal`.

However, in this recipe, a variation of that syntax is used to state that `A` must be a subtype of a type that has a `speak` method. Specifically, this code can be read as, “`A` must be a

subtype of a type that has a `speak` method. The `speak` method (or function) can't take any parameters and must not return anything.”

To demonstrate another example of the structural type signature, if you wanted to state that the `speak` method must take a `String` parameter and return a `Boolean`, the structural type signature would look like this:

```
[A <: { def speak(s: String): Boolean }]
```

As a word of warning, this technique uses reflection, so you may not want to use it when performance is a concern.

19.4. Make Mutable Collections Invariant

Problem

You want to create a collection whose elements can be mutated, and want to know how to specify the generic type parameter for its elements.

Solution

When creating a collection of elements that can be changed (mutated), its generic type parameter should be declared as `[A]`, making it *invariant*.

For instance, elements in a Scala `Array` or `ArrayBuffer` can be mutated, and their signatures are declared like this:

```
class Array[A] ...
class ArrayBuffer[A] ...
```

Declaring a type as invariant has several effects. First, the container can hold both the specified types as well as its subtypes. For example, the following class hierarchy states that the `Dog` and `SuperDog` classes both extend the `Animal` trait:

```
trait Animal {
  def speak
}

class Dog(var name: String) extends Animal {
  def speak { println("woof") }
  override def toString = name
}

class SuperDog(name: String) extends Dog(name) {
  def useSuperPower { println("Using my superpower!") }
}
```

With these classes, you can create a `Dog` and a `SuperDog`:

```
val fido = new Dog("Fido")
val wonderDog = new SuperDog("Wonder Dog")
val shaggy = new SuperDog("Shaggy")
```

When you later declare an `ArrayBuffer[Dog]`, you can add both `Dog` *and* `SuperDog` instances to it:

```
val dogs = ArrayBuffer[Dog]()
dogs += fido
dogs += wonderDog
```

So a collection with an invariant type parameter can contain elements of the base type, and subtypes of the base type.

The second effect of declaring an invariant type is the primary purpose of this recipe. Given the same code, you can define a method as follows to accept an `ArrayBuffer[Dog]`, and then have each `Dog` speak:

```
import collection.mutable.ArrayBuffer
def makeDogsSpeak(dogs: ArrayBuffer[Dog]) {
  dogs.foreach(_.speak)
}
```

Because of its definition, this works fine when you pass it an `ArrayBuffer[Dog]`:

```
val dogs = ArrayBuffer[Dog]()
dogs += fido
makeDogsSpeak(dogs)
```

However, the `makeDogsSpeak` call won't compile if you attempt to pass it an `ArrayBuffer[SuperDog]`:

```
val superDogs = ArrayBuffer[SuperDog]()
superDogs += shaggy
superDogs += wonderDog
makeDogsSpeak(superDogs) // ERROR: won't compile
```

This code won't compile because of the conflict built up in this situation:

- Elements in an `ArrayBuffer` can be mutated.
- `makeDogsSpeak` is defined to accept a parameter of type `ArrayBuffer[Dog]`.
- You're attempting to pass in `superDogs`, whose type is `ArrayBuffer[SuperDog]`.
- If the compiler allowed this, `makeDogsSpeak` could replace `SuperDog` elements in `superDogs` with plain old `Dog` elements. This can't be allowed.

One of the reasons this problem occurs is that `ArrayBuffer` elements can be mutated. If you want to write a method to make all `Dog` types *and* subtypes speak, define it to accept a collection of immutable elements, such as a `List`, `Seq`, or `Vector`.

Discussion

The elements of the `Array`, `ArrayBuffer`, and `ListBuffer` classes can be mutated, and they're all defined with invariant type parameters:

```
class Array[T]
class ArrayBuffer[A]
class ListBuffer[A]
```

Conversely, collections classes that are immutable identify their generic type parameters differently, with the `+` symbol, as shown here:

```
class List[+T]
class Vector[+A]
trait Seq[+A]
```

The `+` symbol used on the type parameters of the immutable collections defines their parameters to be *covariant*. Because their elements can't be mutated, adding this symbol makes them more flexible, as discussed in the next recipe.

See Also

You can find the source code for Scala classes by following the “Source code” links in their Scaladoc. The source code for the `ArrayBuffer` class isn't too long, and it shows how the type parameter `A` ends up sprinkled throughout the class:

```
ArrayBuffer class Scaladoc
```

19.5. Make Immutable Collections Covariant

Problem

You want to create a collection whose elements can't be changed (they're immutable), and want to understand how to specify it.

Solution

You *can* define a collection of immutable elements as invariant, but your collection will be much more flexible if you declare that your type parameter is *covariant*. To make a type parameter covariant, declare it with the `+` symbol, like `[+A]`.

Covariant type parameters are shown in the Scaladoc for immutable collection classes like `List`, `Vector`, and `Seq`:

```
class List[+T]
class Vector[+A]
trait Seq[+A]
```

By defining the type parameter to be covariant, you create a situation where the collection can be used in a more flexible manner.

To demonstrate this, modify the example from the previous recipe slightly. First, define the class hierarchy:

```
trait Animal {
  def speak
}

class Dog(var name: String) extends Animal {
  def speak { println("Dog says woof") }
}

class SuperDog(name: String) extends Dog(name) {
  override def speak { println("I'm a SuperDog") }
}
```

Next, define a `makeDogsSpeak` method, but instead of accepting a mutable `ArrayBuffer[Dog]` as in the previous recipe, accept an immutable `Seq[Dog]`:

```
def makeDogsSpeak(dogs: Seq[Dog]) {
  dogs.foreach(_.speak)
}
```

As with the `ArrayBuffer` in the previous recipe, you can pass a sequence of type `[Dog]` into `makeDogsSpeak` without a problem:

```
// this works
val dogs = Seq(new Dog("Fido"), new Dog("Tanner"))
makeDogsSpeak(dogs)
```

However, in this case, you can also pass a `Seq[SuperDog]` into the `makeDogsSpeak` method successfully:

```
// this works too
val superDogs = Seq(new SuperDog("Wonder Dog"), new SuperDog("Scooby"))
makeDogsSpeak(superDogs)
```

Because `Seq` is immutable and defined with a covariant parameter type, `makeDogsSpeak` can now accept collections of both `Dog` and `SuperDog`.

Discussion

You can demonstrate this by creating a collection class with a covariant type parameter. To do this, create a collection class that can hold one element. Because you don't want the collection element to be mutated, define the element as a `val`, and make the type parameter covariant with `+A`:

```
class Container[+A] (val elem: A)
```

Using the same type hierarchy as shown in the Solution, modify the `makeDogsSpeak` method to accept a `Container[Dog]`:

```
def makeDogsSpeak(dogHouse: Container[Dog]) {  
  dogHouse.elem.speak()  
}
```

With this setup, you can pass a `Container[Dog]` into `makeDogsSpeak`:

```
val dogHouse = new Container(new Dog("Tanner"))  
makeDogsSpeak(dogHouse)
```

Finally, to demonstrate the point of adding the `+` symbol to the parameter, you can also pass a `Container[SuperDog]` into `makeDogsSpeak`:

```
val superDogHouse = new Container(new SuperDog("Wonder Dog"))  
makeDogsSpeak(superDogHouse)
```

Because the `Container` element is immutable and its mutable type parameter is marked as covariant, all of this code works successfully. Note that if you change the `Container`'s type parameter from `+A` to `A`, the last line of code won't compile.

As demonstrated in these examples, defining an immutable collection to take a covariant generic type parameter makes the collection more flexible and useful throughout your code.

19.6. Create a Collection Whose Elements Are All of Some Base Type

Problem

You want to specify that a class or method takes a type parameter, and that parameter is limited so it can only be a base type, or a subtype of that base type.

Solution

Define the class or method by specifying the type parameter with an *upper bound*.

To demonstrate this, create a simple type hierarchy:

```
trait CrewMember  
class Officer extends CrewMember  
class RedShirt extends CrewMember  
trait Captain  
trait FirstOfficer  
trait ShipsDoctor  
trait StarfleetTrained
```

Then create a few instances:

```
val kirk = new Officer with Captain
val spock = new Officer with FirstOfficer
val bones = new Officer with ShipsDoctor
```

Given this setup, imagine that you want to create a collection of officers on a ship, like this:

```
val officers = new Crew[Officer]()
officers += kirk
officers += spock
officers += bones
```

The first line lets you create `officers` as a collection that can only contain types that are an `Officer`, or subtype of an `Officer`.

In this example, those who are of type `RedShirt` won't be allowed in the collection, because they don't extend `Officer`:

```
val redShirt = new RedShirt
officers += redShirt // ERROR: this won't compile
```

To enable this functionality and let `Crew` control which types are added to it, define it with an *upper bound* while extending `ArrayBuffer`:

```
class Crew[A <: CrewMember] extends ArrayBuffer[A]
```

This states that any instance of `Crew` can only ever have elements that are of type `CrewMember`. In this example, this lets you define `officers` as a collection of `Officer`, like this:

```
val officers = new Crew[Officer]()
```

It also prevents you from writing code like this, because `String` does not extend `CrewMember`:

```
// error: won't compile
val officers = new Crew[String]()
```

In addition to creating a collection of officers, you can create a collection of `RedShirts`, if desired:

```
val redshirts = new Crew[RedShirt]()
```

(I don't know the names of any redshirts, otherwise I'd add a few to this collection.)

Typically you'll define a class like `Crew` so you can create specific instances as shown. You'll also typically add methods to a class like `Crew` that are specific to the type (`CrewMember`, in this case). By controlling what types are added to `Crew`, you can be assured that your methods will work as desired. For instance, `Crew` could have methods like `beamUp`, `beamDown`, `goWhereNoOneElseHasGone`, etc.—any method that makes sense for a `CrewMember`.

Discussion

This type is referred to as a bound, specifically an *upper bound*.

(If you're working with an implicit conversion, you'll want to use a *view bound* instead of an upper bound. To do this, use the `<%` symbol instead of the `<:` symbol.)

You can use the same technique when you need to limit your class to take a type that extends multiple traits. For example, to create a `Crew` that only allows types that extend `CrewMember` and `StarfleetTrained`, declare the `Crew` like this:

```
class Crew[A <: CrewMember with StarfleetTrained] extends ArrayBuffer[A]
```

If you adapt the officers to work with this new trait:

```
val kirk = new Officer with Captain with StarfleetTrained
val spock = new Officer with FirstOfficer with StarfleetTrained
val bones = new Officer with ShipsDoctor with StarfleetTrained
```

you can still construct a list of officers, with a slight change to the `Crew` definition:

```
val officers = new Crew[Officer with StarfleetTrained]()
officers += kirk
officers += spock
officers += bones
```

This approach works as long as the instances have those types somewhere in their lineage (class hierarchy). For instance, you can define a new `StarfleetOfficer` like this:

```
class StarfleetOfficer extends Officer with StarfleetTrained
```

You could then define the `kirk` instance like this:

```
val kirk = new StarfleetOfficer with Captain
```

With this definition, `kirk` can still be added to the `officers` collection; the instance still extends `Officer` and `StarfleetTrained`.

Methods

Methods can also take advantage of this syntax. For instance, you can add a little behavior to `CrewMember` and `RedShirt`:

```
trait CrewMember {
  def beamDown { println("beaming down") }
}
class RedShirt extends CrewMember {
  def putOnRedShirt { println("putting on my red shirt") }
}
```

With this behavior, you can write methods to work specifically on their types. This method works for any `CrewMember`:


```
def beamDown[A <: CrewMember](crewMember: Crew[A]) {
  crewMember.foreach(_.beamDown)
}
```

But this method will only work for `RedShirt` types:

```
def getReadyForDay[A <: RedShirt](redShirt: Crew[A]) {
  redShirt.foreach(_.putOnRedShirt)
}
```

In both cases, you control which type can be passed into the method using an appropriate upper bound definition on the method's type parameter.

See Also

- [Recipe 19.3, “Using Duck Typing \(Structural Types\)”](#).
- Scala also includes a *lower type bound*, though it is used less frequently. A lower bound is briefly demonstrated in [Recipe 19.8, “Building Functionality with Types”](#). The page titled [“A Tour of Scala: Lower Type Bounds”](#) also describes a situation where a lower type bound might be used.

19.7. Selectively Adding New Behavior to a Closed Model

Problem

You have a closed model, and want to add new behavior to certain types within that model, while potentially excluding that behavior from being added to other types.

Solution

Implement your solution as a *type class*.

To demonstrate the problem and solution, when I first came to Scala, I thought it would be easy to write a single `add` method that would add any two numeric parameters, regardless of whether they were an `Int`, `Double`, `Float`, or other numeric value. Unfortunately I couldn't get this to work—until I learned about type classes.

Because a `Numeric` type class already exists in the Scala library, it turns out that you can create an `add` method that accepts different numeric types like this:

```
def add[A](x: A, y: A)(implicit numeric: Numeric[A]): A = numeric.plus(x, y)
```

Once defined, this method can be used with different numeric types like this:

```
println(add(1, 1))
println(add(1.0, 1.5))
println(add(1, 1.5F))
```

The add method works because of some magic in the `scala.math.Numeric` trait. To see how this magic works, create your own type class.

Creating a type class

The process of creating a type class is a little complicated, but there is a formula:

- Usually you start with a need, such as having a closed model to which you want to add new behavior.
- To add the new behavior, you define a type class. The typical approach is to create a base trait, and then write specific implementations of that trait using implicit objects.
- Back in your main application, create a method that uses the type class to apply the behavior to the closed model, such as writing the add method in the previous example.

To demonstrate this, assume that you have a closed model that contains `Dog` and `Cat` types, and you want to make a `Dog` more human-like by giving it the capability to speak. However, while doing this, you don't want to make a `Cat` more human-like. (Everyone knows that dogs are human-like and can speak; see YouTube for examples.)

The closed model is defined in a class named *Animals.scala*, and looks like this:

```
package typeclassdemo

// an existing, closed model
trait Animal
final case class Dog(name: String) extends Animal
final case class Cat(name: String) extends Animal
```

To begin making a new speak behavior available to a `Dog`, create a type class that implements the speak behavior for a `Dog`, but not a `Cat`:

```
package typeclassdemo

object Humanish {

  // the type class.
  // defines an abstract method named 'speak'.
  trait HumanLike[A] {
    def speak(speaker: A): Unit
  }

  // companion object
  object HumanLike {
    // implement the behavior for each desired type. in this case,
    // only for a Dog.
    implicit object DogIsHumanLike extends HumanLike[Dog] {
      def speak(dog: Dog) { println(s"I'm a Dog, my name is ${dog.name}") }
    }
  }
}
```

```
}
}
```

With this behavior defined, use the new functionality back in your main application:

```
package typeclassdemo

object TypeClassDemo extends App {

  import Humanish.HumanLike

  // create a method to make an animal speak
  def makeHumanLikeThingSpeak[A](animal: A)(implicit humanLike: HumanLike[A]) {
    humanLike.speak(animal)
  }

  // because HumanLike implemented this for a Dog, it will work
  makeHumanLikeThingSpeak(Dog("Rover"))

  // however, the method won't compile for a Cat (as desired)
  //makeHumanLikeThingSpeak(Cat("Morris"))

}
```

The comments in the code explain why this approach works for a Dog, but not a Cat.

There are a few other things to notice from this code:

- The `makeHumanLikeThingSpeak` is similar to the `add` method in the first example.
- In the first example, the `Numeric` type class already existed, so you could just use it to create the `add` method. But when you're starting from scratch, you need to create your own type class (the code in the `HumanLike` trait).
- Because a `speak` method is defined in the `DogIsHumanLike` implicit object, which extends `HumanLike[Dog]`, a `Dog` can be passed into the `makeHumanLikeThingSpeak` method. But because a similar implicit object has not been written for the `Cat` class, it can't be used.

Discussion

Despite the name “class,” a type class doesn't come from the OOP world; it comes from the FP world, specifically Haskell. As shown in the examples, one benefit of a type class is that you can add behavior to a closed model.

Another benefit is that it lets you define methods that take generic types, and provide control over what those types are. For instance, in the first example, the `add` method takes `Numeric` types:

```
def add[A](x: A, y: A)(implicit numeric: Numeric[A]): A = numeric.plus(x, y)
```

Because the `numeric.plus` method is implemented for all the different numeric types, you can create an `add` method that works for `Int`, `Double`, `Float`, and other types:

```
println(add(1, 1))
println(add(1.0, 1.5))
println(add(1, 1.5F))
```

This is great; it works for all numeric types, as desired. As an additional benefit, the `add` method is type safe. If you attempted to pass a `String` into it, it won't compile:

```
// won't compile
add("1", 2.0)
```

In the second example, the `makeHumanLikeThingSpeak` method is similar to the `add` method. However, in this case, it lets a `Dog` type speak, but because the `HumanLike` trait didn't define a similar behavior for a `Cat`, a `Cat` instance can't currently be used by the method. You can resolve this by adding a `speak` method for a `Cat` type as another implicit object, or keep the code as it's currently written to prevent a `Cat` from speaking.

See Also

- If you dig into the source code for Scala's `Numeric` trait, you'll find that it's implemented in a manner similar to what's shown here. You can find the source code for Scala's `Numeric` trait by following the "Source code" link on [its Scaladoc page](#).
- [Recipe 1.10, "Add Your Own Methods to the String Class"](#) demonstrates how to add new functionality to closed classes using implicit conversions.

19.8. Building Functionality with Types

To put what you've learned in this chapter to use, let's create two examples. First, you'll create a "timer" that looks like a control structure and works like the Unix `time` command. Second, you'll create another control structure that works like the Scala 2.10 `Try/Success/Failure` classes.

Example 1: Creating a Timer

On Unix systems you can run a `time` command (`timex` on some systems) to see how long commands take to execute:

```
$ time find . -name "*.scala"
```

That command returns the results of the `find` command it was given, along with the time it took to run. This can be a helpful way to troubleshoot performance problems.

You can create a similar `timer` method in Scala to let you run code like this:

```
val (result, time) = timer(someLongRunningAlgorithm)
println(s"result: $result, time: $time")
```

In this example, the `timer` runs a method named `longRunningAlgorithm`, and then returns the result from the algorithm, along with the algorithm's execution time. You can see how this works by running a simple example in the REPL:

```
scala> val (result, time) = timer{ Thread.sleep(500); 1 }
result: Int = 1
time: Double = 500.32
```

As expected, the code block returns the value 1, with an execution time of about 500 ms.

The `timer` code is surprisingly simple, and involves the use of a generic type parameter:

```
def timer[A](blockOfCode => A) = {
  val startTime = System.nanoTime
  val result = blockOfCode
  val stopTime = System.nanoTime
  val delta = stopTime - startTime
  (result, delta/1000000d)
}
```

The `timer` method uses Scala's *call-by-name* syntax to accept a block of code as a parameter. Rather than declare a specific return type from the method (such as `Int`), you declare the return type to be a generic type parameter. This lets you pass all sorts of algorithms into `timer`, including those that return nothing:

```
scala> val (result, time) = timer{ println("Hello") }
Hello
result: Unit = ()
time: Double = 0.544
```

Or an algorithm that reads a file and returns an iterator:

```
scala> def readFile(filename: String) = io.Source.fromFile(filename).getLines
readFile: (filename: String)Iterator[String]

scala> val (result, time) = timer{ readFile("/etc/passwd") }
result: Iterator[String] = non-empty iterator
time: Double = 32.119
```

This is a simple use of specifying a generic type in a noncollection class, and helps you get ready for the next example.

Example 2: Writing Your Own “Try” Classes

Imagine the days back before Scala 2.10 when there was no such thing as the `Try`, `Success`, and `Failure` classes in *scala.util*. (They were available from Twitter, but just ignore that for now.) In those days, you might have come up with your own solution that you called `Attempt`, `Succeeded`, and `Failed` that would let you write code like this:

```
val x = Attempt("10".toInt) // Succeeded(10)
val y = Attempt("10A".toInt) // Failed(Exception)
```

To enable this basic API, you realize you'll need a class named `Attempt`, and because you know you don't want to use the `new` keyword to create a new instance, you know that you need a companion object with an `apply` method. You further realize that you need to define `Succeeded` and `Failed`, and they should extend `Attempt`. Therefore, you begin with this code, placed in a file named *Attempt.scala*:

```
// version 1
sealed class Attempt[A]

object Attempt {

  def apply[A](f: => A): Attempt[A] =
    try {
      val result = f
      return Succeeded(result)
    } catch {
      case e: Exception => Failed(e)
    }

}

final case class Failed[A](val exception: Throwable) extends Attempt[A]
final case class Succeeded[A](value: A) extends Attempt[A]
```

In a manner similar to the previous `timer` code, the `apply` method takes a call-by-name parameter, and the return type is specified as a generic type parameter. In this case, the type parameter ends up sprinkled around in other areas. Because `apply` returns a type of `Attempt`, it's necessary there; because `Failed` and `Succeeded` extend `Attempt`, it's propagated there as well.

This first version of the code lets you write the basic `x` and `y` examples. However, to be really useful, your API needs a new method named `getOrElse` that lets you get the information from the result, whether that result happens to be a type of `Succeeded` or `Failed`.

The `getOrElse` method will be called like this:

```
val x = Attempt(1/0)
val result = x.getOrElse(0)
```

Or this:

```
val y = Attempt("foo".toInt).getOrElse(0)
```

To enable a `getOrElse` method, make the following changes to the code:

```
// version 2
sealed abstract class Attempt[A] {
  def getOrElse[B >: A](default: => B): B = if (isSuccess) get else default
}
```

```

    var isSuccess = false
    def get: A
  }

  object Attempt {
    def apply[A](f: => A): Attempt[A] =
      try {
        val result = f
        Succeeded(result)
      } catch {
        case e: Exception => Failed(e)
      }
  }

  final case class Failed[A](val exception: Throwable) extends Attempt[A] {
    isSuccess = false
    def get: A = throw exception
  }

  final case class Succeeded[A](result: A) extends Attempt[A] {
    isSuccess = true
    def get = result
  }

```

The variable `isSuccess` is added to `Attempt` so it can be set in `Succeeded` or `Failed`. An abstract method named `get` is also declared in `Attempt` so it can be implemented in the two subclasses. These changes let the `getOrElse` method in `Attempt` work.

The `getOrElse` method signature is the most interesting thing about this new code:

```
def getOrElse[B >: A](default: => B): B = if (isSuccess) get else default
```

Because of the way `getOrElse` works, it can either return the type `A`, which is the result of the expression, or type `B`, which the user supplies, and is presumably a substitute for `A`. The expression `B >: A` is a *lower bound*. Though it isn't commonly used, a lower bound declares that a type is a supertype of another type. In this code, the term `B >: A` expresses that the type parameter `B` is a supertype of `A`.

The Scala 2.10 Try classes

You could keep developing your own classes, but the `Try`, `Success`, and `Failure` classes in the `scala.util` package were introduced in Scala 2.10, so this is a good place to stop.

However, it's worth noting that these classes can be a great way to learn about Scala types. For instance, the `getOrElse` method in the `Attempt` code is the same as the `getOrElse` method declared in `Try`:

```
def getOrElse[U >: T](default: => U): U = if (isSuccess) get else default
```

The `map` method declared in `Success` shows how to define a call-by-name parameter that transforms a type `T` to a type `U`:

```
def map[U](f: T => U): Try[U] = Try[U](f(value))
```

Its `flatten` method uses the `<:<` symbol that wasn't covered in this chapter. When used as `A <:< B`, it declares that "A must be a subtype of B." Here's how it's used in the `Success` class:

```
def flatten[U](implicit ev: T <:< Try[U]): Try[U] = value
```

When it comes to learning about generic parameter types, these classes are very interesting to study. They're self-contained and surprisingly short. The Scala collections classes also demonstrate many more uses of generics.

Introduction

When I first came to Scala from Java, I was happy with the small things, including eliminating a lot of `;`, `()`, and `{}` characters, and writing more concise, Ruby-like code. These were nice little wins that made for “a better Java.”

Over time, I wanted to add more to my repertoire and use Scala the way it’s intended to be used. As Ward Cunningham said in the book, *Clean Code* (Prentice Hall), I wanted to write code that “makes it look like the language was made for the problem.”

That’s what this chapter is about—trying to share some of the best practices of Scala programming so you can write code in “the Scala way.”

Before digging into the recipes in this chapter, here’s a short summary of Scala’s best practices.

At the application level:

- At the big-picture, application-design level, follow the 80/20 rule, and try to write 80% of your application as pure functions, with a thin layer of other code on top of those functions for things like I/O.
- Learn “Expression-Oriented Programming” ([Recipe 20.3](#)).
- Use the Actor classes to implement concurrency ([Chapter 13](#)).
- Move behavior from classes into more granular traits. This is best described in [the Scala Stackable Trait pattern](#).

At the coding level:

- Learn how to write pure functions. At the very least, they simplify testing.
- Learn how to pass functions around as variables ([Recipes 9.2 to 9.4](#)).

- Learn how to use the Scala collections API. Know the most common classes and methods ([10](#) and [11](#)).
- Prefer immutable code. Use `vals` and immutable collections first ([Recipe 20.2](#)).
- Drop the `null` keyword from your vocabulary. Use the `Option/Some/None` and `Try/Success/Failure` classes instead ([Recipe 20.6](#)).
- Use TDD and/or BDD testing tools like `ScalaTest` and `specs2`.

Outside the code:

- Learn how to use SBT. It's the de-facto Scala build tool ([Chapter 18](#)).
- Keep a REPL session open while you're coding (or use [the Scala Worksheet](#)), and constantly try small experiments ([Recipes 14.1](#) to [14.4](#), and many examples throughout the book).

Other Resources

In addition to the practices shared in this chapter, I highly recommend reading Twitter's [Effective Scala document](#). The Twitter team has been a big user and proponent of Scala, and this document summarizes their experiences.

The [Scala Style Guide](#) is a good resource that shares examples of how to write code in the Scala "style."

20.1. Create Methods with No Side Effects (Pure Functions)

Problem

In keeping with the best practices of Functional Programming (FP), you want to write "pure functions."

Solution

In general, when writing a function (or method), your goal should be to write it as a pure function. This raises the question, "What is a pure function?" Before we tackle that question we need to look at another term, *referential transparency*, because it's part of the description of a pure function.

Referential transparency

If you like algebra, you'll like referential transparency. An expression is referentially transparent (RT) if it can be replaced by its resulting value without changing the behavior of the program. This must be true regardless of where the expression is used in the program.

For instance, assume that `x` and `y` are immutable variables within some scope of an application, and within that scope they're used to form this expression:

```
x + y
```

You can assign this expression to a third variable, like this:

```
val z = x + y
```

Now, throughout the given scope of your program, anywhere the expression `x + y` is used, it can be replaced by `z` without affecting the result of the program.

Note that although I stated that `x` and `y` are immutable variables, they can also be the result of RT functions. For instance, `"hello".length + "world".length` will always be 10. This result could be assigned to `z`, and then `z` could be used everywhere instead of this expression.

Although this is a simple example, this is referential transparency in a nutshell.

Pure functions

Wikipedia defines a pure function as follows:

1. The function always evaluates to the same result value given the same argument value(s). It cannot depend on any hidden state or value, and it cannot depend on any I/O.
2. Evaluation of the result does not cause any semantically observable side effect or output, such as mutation of mutable objects or output to I/O devices.

The book *Functional Programming in Scala* by Chiusano and Bjarnason (Manning Publications), states this a little more precisely:

“A function `f` is *pure* if the expression `f(x)` is referentially transparent for all referentially transparent values `x`.”

To summarize, a pure function is referentially transparent and has no side effects.

Regarding side effects, the authors of the book, *Programming in Scala*, make a great observation:

“A telltale sign of a function with side effects is that its result type is `Unit`.”

From these definitions, we can make these statements about pure functions:

- A pure function is given one or more input parameters.
- Its result is based solely off of those parameters and its algorithm. The algorithm will not be based on any hidden state in the class or object it's contained in.
- It won't mutate the parameters it's given.
- It won't mutate the state of its class or object.
- It doesn't perform any I/O operations, such as reading from disk, writing to disk, prompting for input, or reading input.

These are some examples of pure functions:

- Mathematical functions, such as addition, subtraction, multiplication.
- Methods like `split` and `length` on the `String` class.
- The `to*` methods on the `String` class (`toInt`, `toDouble`, etc.)
- Methods on immutable collections, including `map`, `drop`, `take`, `filter`, etc.
- The functions that extract values from an HTML string in [Recipe 20.3](#).

The following functions are not pure functions:

- Methods like `getDayOfWeek`, `getHour`, or `getMinute`. They return a different value depending on when they are called.
- A `getRandomNumber` function.
- A function that reads user input or prints output.
- A function that writes to an external data store, or reads from a data store.

If you're coming to Scala from a pure OOP background, it can be difficult to write pure functions. Speaking for myself, historically my code has followed the OOP paradigm of encapsulating data and behavior in classes, and as a result, my methods often mutated the internal state of objects.



At this point you may be wondering how you can get anything done in a program consisting only of pure functions. If you can't read input from a user or database, and can't write output, how will your application ever work?

The best advice I can share about FP is to follow the 80/20 rule: write 80% of your program using pure functions (the “cake”), then create a 20% layer of other code on top of the functional base (the “icing”) to handle the user interface, printing, database interactions, and other methods that have “side effects”.

Obviously any interesting application will have I/O, and this balanced approach lets you have the best of both worlds.

The Java approach

To look at how to write pure functions, you'll convert the methods in an OOP class into pure functions. The following code shows how you might create a `Stock` class that follows the Java/OOP paradigm. The following class intentionally has a few flaws. It not only has the ability to store information about a `Stock`, but it can also access the Internet to get the current stock price, and further maintains a list of historical prices for the stock:

```
// a poorly written class

class Stock (var symbol: String, var company: String,
             var price: BigDecimal, var volume: Long) {

    var html: String = _
    def buildUrl(stockSymbol: String): String = { ... }
    def getUrlContent(url: String):String = { ... }

    def setPriceFromHtml(html: String) { this.price = ... }
    def setVolumeFromHtml(html: String) { this.volume = ... }
    def setHighFromHtml(html: String) { this.high = ... }
    def setLowFromHtml(html: String) { this.low = ... }

    // some dao-like functionality
    private val _history: ArrayBuffer[Stock] = { ... }
    val getHistory = _history

}
```

Beyond attempting to do too many things, from an FP perspective, it has these other problems:

- All of its fields are mutable.
- All of the set methods mutate the class fields.
- The getHistory method returns a mutable data structure.

The getHistory method is easily fixed by only sharing an immutable data structure, but this class has deeper problems. Let's fix them.

Fixing the problems

The first fix is to separate two concepts that are buried in the class. First, there should be a concept of a Stock, where a Stock consists only of a symbol and company name. You can make this a case class:

```
case class Stock(symbol: String, company: String)
```

Examples of this are Stock("AAPL", "Apple") and Stock("GOOG", "Google").

Second, at any moment in time there is information related to a stock's performance on the stock market. You can call this data structure a StockInstance, and also define it as a case class:

```
case class StockInstance(symbol: String,
                        datetime: String,
                        price: BigDecimal,
                        volume: Long)
```

A StockInstance example looks like this:

```
StockInstance("AAPL", "Nov. 2, 2012 5:00pm", 576.80, 20431707)
```

Going back to the original class, the getUrlContent method isn't specific to a stock, and should be moved to a different object, such as a general-purpose NetworkUtils object:

```
object NetworkUtils {
  def getUrlContent(url: String): String = { ... }
}
```

This method takes a URL as a parameter and returns the HTML content from that URL.

Similarly, the ability to build a URL from a stock symbol should be moved to an object. Because this behavior is specific to a stock, you'll put it in an object named StockUtils:

```
object StockUtils {
  def buildUrl(stockSymbol: String): String = { ... }
}
```

The ability to extract the stock price from the HTML can also be written as a pure function and should be moved into the same object:

```
object StockUtils {
  def buildUrl(stockSymbol: String): String = { ... }
  def extractPrice(html: String): BigDecimal = { ... }
}
```

```

    def getPrice(html: String): String = { ... }
}

```

In fact, all of the methods named `set*` in the previous class should be `get*` methods in `StockUtils`:

```

object StockUtils {
  def buildUrl(stockSymbol: String): String = { ... }
  def getPrice(symbol: String, html: String): String = { ... }
  def getVolume(symbol: String, html: String): String = { ... }
  def getHigh(symbol: String, html: String): String = { ... }
  def getLow(symbol: String, html: String): String = { ... }
}

```

The methods `getPrice`, `getVolume`, `getHigh`, and `getLow` are all pure functions: given the same HTML string and stock symbol, they will always return the same values, and they don't have side effects.

Following this thought process, the date and time are moved to a `DateUtils` object:

```

object DateUtils {
  def currentDate: String = { ... }
  def currentTime: String = { ... }
}

```

With this new design, you create an instance of a `Stock` for the current date and time as a simple series of expressions. First, retrieve the HTML that describes the stock from a web page:

```

val stock = new Stock("AAPL", "Apple")
val url = StockUtils.buildUrl(stock.symbol)
val html = NetUtils.getUrlContent(url)

```

Once you have the HTML, extract the desired stock information, get the date, and create the `Stock` instance:

```

val price = StockUtils.getPrice(html)
val volume = StockUtils.getVolume(html)
val high = StockUtils.getHigh(html)
val low = StockUtils.getLow(html)
val date = DateUtils.currentDate
val stockInstance = StockInstance(symbol, date, price, volume, high, low)

```

Notice that all of the variables are immutable, and each line is an expression.

The code is simple, so you can eliminate all the intermediate variables, if desired:

```

val html = NetUtils.getUrlContent(url)
val stockInstance = StockInstance(
  symbol,
  DateUtils.currentDate,
  StockUtils.getPrice(html),
  StockUtils.getVolume(html),

```

```
StockUtils.getHigh(html),
StockUtils.getLow(html))
```

As mentioned earlier, the methods `getPrice`, `getVolume`, `getHigh`, and `getLow` are all pure functions. But what about methods like `getDate`? It's not a pure function, but the fact is, you need the date and time to solve the problem. This is part of what's meant by having a healthy, balanced attitude about pure functions.

As a final note about this example, there's no need for the `Stock` class to maintain a mutable list of stock instances. Assuming that the stock information is stored in a database, you can create a `StockDao` to retrieve the data:

```
object StockDao {
  def getStockInstances(symbol: String): Vector[StockInstance] = { ... }
  // other code ...
}
```

Though `getStockInstances` isn't a pure function, the `Vector` class is immutable, so you can feel free to share it without worrying that it might be modified somewhere else in your application.



Although I use the prefix `get` in many of those method names, it's not at all necessary to follow a JavaBeans-like naming convention. In fact, in part because you write “setter” methods in Scala without beginning their names with `set`, and also to follow the **Uniform Access Principle**, many Scala APIs don't use `get` or `set` at all.

For example, think of case classes. The accessors and mutators they generate don't use `get` or `set`:

```
case class Person(name: String)
val p = Person("Mark")
p.name           // accessor
p.name = "Bubba" // mutator
```

That being said, although it's best to follow the Scala standards, use whatever method names best fit your API.

Discussion

A benefit of this coding style is that pure functions are easier to test. For instance, attempting to test the `set*` methods in the original code is harder than it needs to be. For each field (`price`, `volume`, `high`, and `low`), you have to follow these steps:

1. Set the `html` field in the object.
2. Call the current `set` method, such as `setPriceFromHtml`.
3. Internally, this method reads the private `html` class field.
4. When the method runs, it mutates a field in the class (`price`).

5. You have to “get” that field to verify that it was changed.
6. In more complicated classes, it’s possible that the `html` and `price` fields may be mutated by other methods in the class.

The test code for the original class looks like this:

```
val stock = new Stock("AAPL", "Apple", 0, 0)
stock.buildUrl
val html = stock.getUrlContent
stock.getPriceFromHtml(html)
assert(stock.getPrice == 500.0)
```

This is a simple example of testing one method that has side effects, but of course this can get much more complicated in a large application.

By contrast, testing a pure function is easier:

1. Call the function, passing in a known value.
2. Get a result back from the function.
3. Verify that the result is what you expected.

The functional approach results in test code like this:

```
val url = NetUtils.buildUrl("AAPL")
val html = NetUtils.getUrlContent(url)
val price = StockUtils.getPrice(html)
assert(price == 500.0)
```

Although the code shown isn’t much shorter, it is much simpler.

StockUtils or Stock object?

The methods that were moved to the `StockUtils` class in the previous examples could be placed in the companion object of the `Stock` class. That is, you could have placed the `Stock` class and object in a file named *Stock.scala*:

```
case class Stock(symbol: String, company: String)

object Stock {
  def buildUrl(stockSymbol: String): String = { ... }
  def getPrice(symbol: String, html: String): String = { ... }
  def getVolume(symbol: String, html: String): String = { ... }
  def getHigh(symbol: String, html: String): String = { ... }
  def getLow(symbol: String, html: String): String = { ... }
}
```

For the purposes of this example, I put these methods in a `StockUtils` class to be clear about separating the concerns of the *Stock* class and *object*. In your own practice, use whichever approach you prefer.

See Also

- [Pure Functions](#)
- [Referential Transparency](#)
- The [Uniform Access Principle](#)

20.2. Prefer Immutable Objects

Problem

You want to reduce the use of mutable objects and data structures in your code.

Solution

Begin with this simple philosophy, stated in the book, *Programming in Scala*:

“Prefer `vals`, immutable objects, and methods without side effects. Reach for them first.”

Then use other approaches with justification.

There are two components to “prefer immutability”:

- Prefer immutable collections. For instance, use immutable sequences like `List` and `Vector` before reaching for the mutable `ArrayBuffer`.
- Prefer immutable variables. That is, prefer `val` to `var`.

In Java, mutability is the default, and it can lead to unnecessarily dangerous code and hidden bugs. In the following example, even though the `List` parameter taken by the `trustMeMuHaHa` method is marked as `final`, the method can still mutate the collection:

```
// java

class EvilMutator {

    // trust me ... mu ha ha (evil laughter)
    public static void trustMeMuHaHa(final List<Person> people) {
        people.clear();
    }

}
```

Although Scala treats method arguments as `vals`, you leave yourself open to the exact same problem by passing around a mutable collection, like an `ArrayBuffer`:

```
def evilMutator(people: ArrayBuffer[Person]) {
    people.clear()
}
```

Just as with the Java code, the `evilMutator` method can call `clear` because the contents of an `ArrayBuffer` are mutable.

Though nobody would write malicious code like this intentionally, accidents do happen. To make your code safe from this problem, if there's no reason for a collection to be changed, don't use a mutable collection class. By changing the collection to a `Vector`, you eliminate the possibility of this problem, and the following code won't even compile:

```
def evilMutator(people: Vector[Person]) {  
  // ERROR - won't compile  
  people.clear()  
}
```

Because `Vector` is immutable, any attempt to add or remove elements will fail.

Discussion

There are at least two major benefits to using immutable variables (`val`) and immutable collections:

- They represent a form of defensive coding, keeping your data from being changed accidentally.
- They're easier to reason about.

The examples shown in the Solution demonstrate the first benefit: if there's no need for other code to mutate your reference or collection, don't let them do it. Scala makes this easy.

The second benefit can be thought of in many ways, but I like to think about it when using actors and concurrency. If I'm using immutable collections, I can pass them around freely. There's no concern that another thread will modify the collection.

Using `val` + mutable, and `var` + immutable

As mentioned several times in this chapter, it's important to have a balanced attitude. I generally use that expression in regards to pure functions, but it also has meaning when discussing "prefer immutability."

For instance, some developers like to use these combinations:

- A mutable collection field declared as a `val`.
- An immutable collection field declared as a `var`.

These approaches generally seem to be used as follows:

- A mutable collection field declared as a `val` is typically made private to its class (or method).
- An immutable collection field declared as a `var` in a class is more often made publicly visible, that is, it's made available to other classes.

As an example of the first approach, the current Akka FSM class (*scala.akka.actor.FSM*) defines several mutable collection fields as private `vals`, like this:

```
private val timers = mutable.Map[String, Timer]()

// some time later ...
timers -= name
timers.clear()
```

This is safe to do, because the `timers` field is private to the class, so its mutable collection isn't shared with others.

An approach I used on a recent project is a variation of this theme:

```
class Pizza {

  private val _toppings = new collection.mutable.ArrayBuffer[Topping]()

  def toppings = _toppings.toList
  def addTopping(t: Topping) { _toppings += t }
  def removeTopping(t: Topping) { _toppings -= t }

}
```

This code defines `_toppings` as a mutable `ArrayBuffer`, but makes it a `val` that's private to the `Pizza` class. Here's my rationale for this approach:

- I made `_toppings` an `ArrayBuffer` because I knew that elements (toppings) would often be added and removed.
- I made `_toppings` a `val` because there was no need for it to ever be reassigned.
- I made it private so its accessor wouldn't be visible outside of my class.
- I created the methods `toppings`, `addTopping`, and `removeTopping` to let other code manipulate the collection.
- When other code calls the `toppings` method, I can give them an immutable copy of the toppings.

I intentionally didn't use the "val + mutable collection" approach, which would have looked like this:

```
// did not do this
val toppings = new collection.mutable.ArrayBuffer[Topping]()
```

I didn't use this approach because I didn't want to expose `toppings` as an immutable collection outside of my `Pizza` class, which would have happened here, because the `val` would have generated an accessor method. In using an OOP design, you think, "Who should be responsible for managing the toppings on the pizza?" and `Pizza` clearly has the responsibility of maintaining its toppings.

I also didn't choose this "var + immutable collection" design:

```
var toppings = Vector[Topping]()
```

The benefits of this approach are (a) it automatically shares toppings as an immutable collection, and (b) it lets me add toppings like this:

```
def addTopping(t: Topping) = toppings :+ t
```

But the approach suffers, because it's a little cumbersome to remove an element from a `Vector` (you have to filter the undesired toppings out of the originating `Vector` while assigning the result to a new `Vector`), and it lets `toppings` be reassigned outside of the `Pizza` class, which I don't want:

```
// bad: other code can mutate 'toppings'
pizza.toppings = Vector(Cheese)
```

You *can* remove elements with this approach by using the `filter` method and then reassigning the result back to `toppings`, like this:

```
toppings = toppings.filter(_ != Pepperoni)
```

But if you create a "double pepperoni" pizza by having two instances of `Pepperoni` in `toppings`, and then want to change it to a regular pepperoni pizza, the earlier `ArrayBuffer` approach is simpler.

Summary

In summary, always begin with the "prefer immutability" approach, and relax that philosophy when it makes sense for the current situation, that is, when you can properly rationalize your decision.

See Also

Recipe 10.6, "Understanding Mutable Variables with Immutable Collections"

20.3. Think "Expression-Oriented Programming"

Problem

You're used to writing *statements* in another programming language, and want to learn how to write *expressions* in Scala, and the benefits of the expression-oriented programming (EOP) philosophy.

Solution

To understand EOP, you have to understand the difference between a *statement* and an *expression*. Wikipedia provides a concise distinction between the two:

“Statements do not return results and are executed solely for their side effects, while expressions always return a result and often do not have side effects at all.”

So statements are like this:

```
order.calculateTaxes()  
order.updatePrices()
```

Expressions are like this:

```
val tax = calculateTax(order)  
val price = calculatePrice(order)
```

On Wikipedia’s EOP page, it also states:

“An expression-oriented programming language is a programming language where every (or nearly every) construction is an expression, and thus yields a value.”

As you might expect, it further states that all pure FP languages are expression-oriented.

The following example helps to demonstrate EOP. This recipe is similar to [Recipe 20.1](#), so it reuses the class from that recipe to show a poor initial design:

```
// an intentionally bad example  
  
class Stock (var symbol: String,  
             var company: String,  
             var price: String,  
             var volume: String,  
             var high: String,  
             var low: String) {  
  
    var html: String = _  
    def buildUrl(stockSymbol: String): String = { ... }  
    def getUrlContent(url: String):String = { ... }  
    def setPriceUsingHtml() { this.price = ... }  
    def setVolumeUsingHtml() { this.volume = ... }  
    def setHighUsingHtml() { this.high = ... }  
    def setLowUsingHtml() { this.low = ... }  
  
}
```

Although I didn’t show it in that recipe, using this class would result in code like this:

```
val stock = new Stock("GOOG", "Google", "", "", "", "")  
val url = buildUrl(stock.symbol)  
stock.html = stock.getUrlContent(url)  
  
// a series of calls on an object ('statements')  
stock.setPriceUsingHtml
```

```
stock.setVolumeUsingHtml
stock.setHighUsingHtml
stock.setLowUsingHtml
```

Although the implementation code isn't shown, all of these "set" methods extract data from the HTML that was downloaded from a [Yahoo Finance](#) page for a given stock, and then update the fields in the current object.

After the first two lines, this code is not expression-oriented at all; it's a series of calls on an object to populate (mutate) the class fields, based on other internal data. These are statements, not expressions; they don't yield values.

Recipe 20.1 showed that by refactoring this class into several different components, you would end up with the following code:

```
// a series of expressions
val url = StockUtils.buildUrl(symbol)
val html = NetUtils.getUrlContent(url)
val price = StockUtils.getPrice(html)
val volume = StockUtils.getVolume(html)
val high = StockUtils.getHigh(html)
val low = StockUtils.getLow(html)
val date = DateUtils.getDate
val stockInstance = StockInstance(symbol, date, price, volume, high, low)
```

This code is expression-oriented. It consists of a series of simple expressions that pass values into pure functions (except for `getDate`), and each function returns a value that's assigned to a variable. The functions don't mutate the data they're given, and they don't have side effects, so they're easy to read, easy to reason about, and easy to test.

Discussion

In Scala, most expressions are obvious. For instance, the following two expressions both return results, which you expect:

```
scala> 2 + 2
res0: Int = 4

scala> List(1,2,3,4,5).filter(_ > 2)
res1: List[Int] = List(3, 4, 5)
```

However, it can be more of a surprise that an `if/else` expression returns a value:

```
val greater = if (a > b) a else b
```

Match expressions also return a result:

```
val evenOrOdd = i match {
  case 1 | 3 | 5 | 7 | 9 => println("odd")
  case 2 | 4 | 6 | 8 | 10 => println("even")
}
```

Even a try/catch block returns a value:

```
val result = try {
  "1".toInt
} catch {
  case _ => 0
}
```

Writing expressions like this is a feature of functional programming languages, and Scala makes using them feel natural and intuitive, and also results in concise, expressive code.

Benefits

Because expressions always return a result, and generally don't have side effects, there are several benefits to EOP:

- The code is easier to reason about. Inputs go in, a result is returned, and there are no side effects.
- The code is easier to test.
- Combined with Scala's syntax, EOP also results in concise, expressive code.
- Although it has only been hinted at in these examples, expressions can often be executed in any order. This subtle feature lets you execute expressions in parallel, which can be a big help when you're trying to take advantage of modern multicore CPUs.

See Also

- [The Wikipedia definition of a statement, and the difference between a statement and an expression](#)
- [Expression-Oriented Programming \(EOP\)](#)

20.4. Use Match Expressions and Pattern Matching

Problem

Match expressions (and pattern matching) are a major feature of the Scala programming language, and you want to see examples of the many ways to use them.

Solution

Match expressions (match/case statements) and pattern matching are a major feature of the Scala language. If you're coming to Scala from Java, the most obvious uses are:

- As a replacement for the Java `switch` statement
- To replace unwieldy `if/then` statements

However, pattern matching is so common, you'll find that match expressions are used in many more situations:

- In `try/catch` expressions
- As the body of a function or method
- With the `Option/Some/None` coding pattern
- In the `receive` method of actors

The following examples demonstrate these techniques.

Replacement for the Java `switch` statement and unwieldy `if/then` statements

Recipe 3.8 showed that a match expression can be used like a Java `switch` statement:

```
val month = i match {  
  case 1 => "January"  
  case 2 => "February"  
  
  // more months here ...  
  
  case 11 => "November"  
  case 12 => "December"  
  case _ => "Invalid month" // the default, catch-all  
}
```

It can be used in the same way to replace unwieldy `if/then/else` statements:

```
i match {  
  case 1 | 3 | 5 | 7 | 9 => println("odd")  
  case 2 | 4 | 6 | 8 | 10 => println("even")  
}
```

These are simple uses of match expressions, but they're a good start.

In `try/catch` expressions

It helps to become comfortable with match expressions, because you'll use them with Scala's `try/catch` syntax. The following example shows how to write a `try/catch` expression that returns an `Option` when lines are successfully read from a file, and `None` if an exception is thrown during the file-reading process:

```
def readTextFile(filename: String): Option[List[String]] = {
  try {
    Some(Source.fromFile(filename).getLines.toList)
  } catch {
    case e: Exception => None
  }
}
```

To catch multiple exceptions in a try/catch expression, list the exception types in the catch clause, just like a match expression:

```
def readTextFile(filename: String): Option[List[String]] = {
  try {
    Some(Source.fromFile(filename).getLines.toList)
  } catch {
    case ioe: IOException =>
      logger.error(ioe)
      None
    case fnf: FileNotFoundException =>
      logger.error(fnf)
      None
  }
}
```

Note that if the specific error is important in a situation like this, use the Try/Success/Failure approach to return the error information to the caller, instead of Option/Some/None. See [Recipe 20.6](#) for both Option and Try examples.

As the body of a function or method

As you get comfortable with match expressions, you'll use them as the body of your methods, such as this method that determines whether the value it's given is true, using the Perl definition of "true":

```
def isTrue(a: Any) = a match {
  case 0 | "" => false
  case _ => true
}
```

In general, a match expression used as the body of a function will accept a parameter as input, match against that parameter, and then return a value:

```
def getClassAsString(x: Any):String = x match {
  case s: String => "String"
  case i: Int => "Int"
  case l: List[_] => "List"
  case p: Person => "Person"
  case Dog() => "That was a Dog"
  case Parrot(name) => s"That was a Parrot, name = $name"
  case _ => "Unknown"
}
```

As shown in [Recipe 9.8](#), a match expression can also be used to create a partial function (i.e., working only for a subset of possible inputs):

```
val divide: PartialFunction[Int, Int] = {  
  case d: Int if d != 0 => 42 / d  
}
```

See that recipe for more details on this approach.

Use with Option/Some/None

Match expressions work well with the Scala Option/Some/None types. For instance, given a method that returns an Option:

```
def toInt(s: String): Option[Int] = {  
  try {  
    Some(s.toInt)  
  } catch {  
    case e: Exception => None  
  }  
}
```

You can handle the result from toInt with a match expression:

```
toInt(aString) match {  
  case Some(i) => println(i)  
  case None => println("Error: Could not convert String to Int.")  
}
```

In a similar way, match expressions are a popular way of handling form verifications with the Play Framework:

```
verifying("If age is given, it must be greater than zero",  
  model =>  
    model.age match {  
      case Some(age) => age < 0  
      case None => true  
    }  
)
```

In actors

Match expressions are baked into actors as *the way* to handle incoming messages:

```
class SarahsBrain extends Actor {  
  def receive = {  
    case StartMessage => handleStartMessage  
    case StopMessage => handleStopMessage  
    case SetMaxWaitTime(time) => helper ! SetMaxWaitTime(time)  
    case SetPhrasesToSpeak(phrases) => helper ! SetPhrasesToSpeak(phrases)  
    case _ => log.info("Got something unexpected.")  
  }  
  
  // other code here ...  
}
```

Summary

Match expressions are an integral part of the Scala language, and as shown, they can be used in many ways. The more you use them, the more uses you'll find for them.

See Also

- Match expressions are demonstrated in many examples in [Chapter 3](#).
- [Chapter 13](#) demonstrates the use of match expressions when writing actors.

20.5. Eliminate null Values from Your Code

Problem

Tony Hoare, inventor of the null reference way back in 1965, refers to the creation of the null value as his “[billion dollar mistake](#).” In keeping with modern best practices, you want to eliminate null values from your code.

Solution

David Pollak, author of the book *Beginning Scala*, offers a wonderfully simple rule about null values:

“Ban null from any of your code. Period.”

Although I've used null values in this book to make some examples easier, in my own practice, I no longer use them. I just imagine that there is no such thing as a null, and write my code in other ways.

There are several common situations where you may be tempted to use null values, so this recipe demonstrates how *not* to use null values in those situations:

- When a var field in a class or method doesn't have an initial default value, initialize it with Option instead of null.
- When a method doesn't produce the intended result, you may be tempted to return null. Use an Option or Try instead.
- If you're working with a Java library that returns null, convert it to an Option, or something else.

Let's look at each of these techniques.

Initialize var fields with Option, not null

Possibly the most tempting time to use a null value is when a field in a class or method won't be initialized immediately. For instance, imagine that you're writing code for the next great social network app. To encourage people to sign up, during the registration process, the only information you ask for is an email address and a password. Because everything else is initially optional, you might write some code like this:

```
case class Address (city: String, state: String, zip: String)

class User(email: String, password: String) {
  var firstName: String = _
  var lastName: String = _
  var address: Address = _
}
```

This is bad news, because `firstName`, `lastName`, and `address` are all declared to be null, and can cause problems in your application if they're not assigned before they're accessed.

A better approach is to define each field as an Option:

```
case class Address (city: String, state: String, zip: String)

class User(email: String, password: String) {
  var firstName = None: Option[String]
  var lastName = None: Option[String]
  var address = None: Option[Address]
}
```

Now you can create a User like this:

```
val u = new User("al@example.com", "secret")
```

At some point later you can assign the other values like this:

```
u.firstName = Some("Al")
u.lastName = Some("Alexander")
u.address = Some(Address("Talkeetna", "AK", "99676"))
```

Later in your code, you can access the fields like this:

```
println(firstName.getOrElse("<not assigned>"))
```

Or this:

```
u.address.foreach { a =>
  println(a.city)
  println(a.state)
  println(a.zip)
}
```

In both cases, if the values are assigned, they'll be printed. With the example of printing the `firstName` field, if the value isn't assigned, the string `<not assigned>` is printed.

In the case of the address, if it's not assigned, the `foreach` loop won't be executed, so the `print` statements are never reached. This is because an `Option` can be thought of as a collection with zero or one elements. If the value is `None`, it has zero elements, and if it is a `Some`, it has one element—the value it contains.

On a related note, you should also use an `Option` in a constructor when a field is optional:

```
case class Stock (id: Long,
                  var symbol: String,
                  var company: Option[String])
```

Don't return null from methods

Because you should never use `null` in your code, the rule for returning `null` values from methods is easy: don't do it.

If you can't return `null`, what can you do? Return an `Option`. Or, if you need to know about an error that may have occurred in the method, use `Try` instead of `Option`.

With an `Option`, your method signatures should look like this:

```
def doSomething: Option[String] = { ... }
def toInt(s: String): Option[Int] = { ... }
def lookupPerson(name: String): Option[Person] = { ... }
```

For instance, when reading a file, a method could return `null` if the process fails, but this code shows how to read a file and return an `Option` instead:

```
def readTextFile(filename: String): Option[List[String]] = {
  try {
    Some(io.Source.fromFile(filename).getLines.toList)
  } catch {
    case e: Exception => None
  }
}
```

This method returns a `List[String]` wrapped in a `Some` if the file can be found and read, or `None` if an exception occurs.

As mentioned, if you want the error information instead of a `Some` or `None`, use the `Try/Success/Failure` approach instead:

```
import scala.util.{Try, Success, Failure}

object Test extends App {

  def readTextFile(filename: String): Try[List[String]] = {
    Try(io.Source.fromFile(filename).getLines.toList)
  }

  val filename = "/etc/passwd"
  readTextFile(filename) match {
    case Success(lines) => lines.foreach(println)
  }
}
```

```

    case Failure(f) => println(f)
  }
}

```

This code prints the lines from the `/etc/passwd` file if the code succeeds, or prints an error message like this if the code fails:

```
java.io.FileNotFoundException: Foo.bar (No such file or directory)
```

As a word of caution (and balance), the Twitter *Effective Scala* page recommends not overusing `Option`, and using the **Null Object Pattern** where it makes sense. As usual, use your own judgment, but try to eliminate all null values using one of these approaches.



A *Null Object* is an object that extends a base type with a “null” or neutral behavior. Here’s a Scala implementation of Wikipedia’s Java example of a Null Object:

```

trait Animal {
  def makeSound()
}

class Dog extends Animal {
  def makeSound() { println("woof") }
}

class NullAnimal extends Animal {
  def makeSound() {}
}

```

The `makeSound` method in the `NullAnimal` class has a neutral, “do nothing” behavior. Using this approach, a method defined to return an `Animal` can return `NullAnimal` rather than `null`.

This is arguably similar to returning `None` from a method declared to return an `Option`, especially when the result is used in a `foreach` loop.

Converting a null into an Option, or something else

The third major place you’ll run into null values is in working with legacy Java code. There is no magic formula here, other than to capture the null value and return something else from your code. That may be an `Option`, a Null Object, an empty list, or whatever else is appropriate for the problem at hand.

For instance, the following `getName` method converts a result from a Java method that may be null and returns an `Option[String]` instead:

```

def getName: Option[String] = {
  var name = javaPerson.getName
  if (name == null) None else Some(name)
}

```

Benefits

Following these guidelines leads to these benefits:

- You'll eliminate `NullPointerException`s.
- Your code will be safer.
- You won't have to write `if` statements to check for `null` values.
- Adding an `Option[T]` return type declaration to a method is a terrific way to indicate that something is happening in the method such that the caller may receive a `None` instead of a `Some[T]`. This is a much better approach than returning `null` from a method that is expected to return an object.
- You'll become more comfortable using `Option`, and as a result, you'll be able to take advantage of how it's used in the collection libraries and other frameworks.

See Also

- [Tony Hoare's Billion Dollar Mistake](#)
- The [“Null Object Pattern”](#)

20.6. Using the Option/Some/None Pattern

Problem

For a variety of reasons, including removing `null` values from your code, you want to use what I call the `Option/Some/None` pattern. Or, if you're interested in a problem (exception) that occurred while processing code, you may want to return `Try/Success/Failure` from a method instead of `Option/Some/None`.

Solution

There is some overlap between this recipe and the previous recipe, [“Eliminate null Values from Your Code”](#). That recipe shows how to use `Option` instead of `null` in the following situations:

- Using `Option` in method and constructor parameters
- Using `Option` to initialize class fields (instead of using `null`)
- Converting `null` results from other code (such as Java code) into an `Option`

See that recipe for examples of how to use an `Option` in those situations.

This recipe adds these additional solutions:

- Returning an `Option` from a method
- Getting the value from an `Option`
- Using `Option` with collections
- Using `Option` with frameworks
- Using `Try/Success/Failure` when you need the error message (Scala 2.10 and newer)
- Using `Either/Left/Right` when you need the error message (pre-Scala 2.10)

Returning an `Option` from a method

The `toInt` method used in this book shows how to return an `Option` from a method. It takes a `String` as input and returns a `Some[Int]` if the `String` is successfully converted to an `Int`, otherwise it returns a `None`:

```
def toInt(s: String): Option[Int] = {  
  try {  
    Some(Integer.parseInt(s.trim))  
  } catch {  
    case e: Exception => None  
  }  
}
```

Although this is a simple method, it shows the common pattern, as well as the syntax. For a more complicated example, see the `readTextFile` example in [Recipe 20.5](#).

This is what `toInt` looks like in the REPL when it succeeds and returns a `Some`:

```
scala> val x = toInt("1")  
x: Option[Int] = Some(1)
```

This is what it looks like when it fails and returns a `None`:

```
scala> val x = toInt("foo")  
x: Option[Int] = None
```

Getting the value from an `Option`

The `toInt` example shows how to declare a method that returns an `Option`. As a consumer of a method that returns an `Option`, there are several good ways to call it and access its result:

- Use `getOrElse`
- Use `foreach`
- Use a match expression

To get the actual value if the method succeeds, or use a default value if the method fails, use `getOrElse`:

```
scala> val x = toInt("1").getOrElse(0)
x: Int = 1
```

Because an `Option` is a collection with zero or one elements, the `foreach` method can be used in many situations:

```
toInt("1").foreach{ i =>
  println(s"Got an int: $i")
}
```

That example prints the value if `toInt` returns a `Some`, but bypasses the `println` statement if `toInt` returns a `None`.

Another good way to access the `toInt` result is with a match expression:

```
toInt("1") match {
  case Some(i) => println(i)
  case None => println("That didn't work.")
}
```

Using Option with Scala collections

Another great feature of `Option` is that it plays well with Scala collections. For instance, starting with a list of strings like this:

```
val bag = List("1", "2", "foo", "3", "bar")
```

imagine you want a list of all the integers that can be converted from that list of strings. By passing the `toInt` method into the `map` method, you can convert every element in the collection into a `Some` or `None` value:

```
scala> bag.map(toInt)
res0: List[Option[Int]] = List(Some(1), Some(2), None, Some(3), None)
```

This is a good start. Because an `Option` is a collection of zero or one elements, you can convert this list of `Int` values by adding `flatten` to `map`:

```
scala> bag.map(toInt).flatten
res1: List[Int] = List(1, 2, 3)
```

As shown in [Recipe 10.16, “Combining map and flatten with flatMap”](#), this is the same as calling `flatMap`:

```
scala> bag.flatMap(toInt)
res2: List[Int] = List(1, 2, 3)
```

The collect method provides another way to achieve the same result:

```
scala> bag.map(toInt).collect{case Some(i) => i}
res3: List[Int] = List(1, 2, 3)
```

That example works because the collect method takes a partial function, and the anonymous function that's passed in is only defined for Some values; it ignores the None values.

These examples work for several reasons:

- toInt is defined to return Option[Int].
- Methods like flatten, flatMap, and others are built to work well with Option values.
- You can pass anonymous functions into the collection methods.

Using Option with other frameworks

Once you begin working with third-party Scala libraries, you'll see that Option is used to handle situations where a variable may not have a value. For instance, they're baked into the Play Framework's Anorm database library, where you use Option/Some/None for database table fields that can be null. In the following example, the third field may be null in the database, so it's handled using Some and None, as shown:

```
def getAll() : List[Stock] = {
  DB.withConnection { implicit connection =>
    sqlQuery().collect {
      // the 'company' field has a value
      case Row(id: Int, symbol: String, Some(company: String)) =>
        Stock(id, symbol, Some(company))
      // the 'company' field does not have a value
      case Row(id: Int, symbol: String, None) =>
        Stock(id, symbol, None)
    }.toList
  }
}
```

The Option approach is also used extensively in Play validation methods:

```
verifying("If age is given, it must be greater than zero",
model =>
  model.age match {
    case Some(age) => age < 0
    case None => true
  }
)
```



The `scala.util.control.Exception` object gives you another way to use an `Option`, depending on your preferences and needs. For instance, the `try/catch` block was removed from the following method and replaced with an `allCatch` method:

```
import scala.util.control.Exception._

def readTextFile(f: String): Option[List[String]] =
  allCatch.opt(Source.fromFile(f).getLines.toList)
```

`allCatch` is described as a `Catch` object “that catches everything.” The `opt` method returns `None` if an exception is caught (such as a `FileNotFoundException`), and a `Some` if the block of code succeeds.

Other `allCatch` methods support the `Try` and `Either` approaches. See the [Exception object Scaladoc](#) for more information.

If you like the `Option/Some/None` approach, but want to write a method that returns error information in the failure case (instead of `None`, which doesn’t return any error information), there are two similar approaches:

- `Try`, `Success`, and `Failure` (introduced in Scala 2.10)
- `Either`, `Left`, and `Right`

I prefer the new `Try/Success/Failure` approach, so let’s look at it next.

Using Try, Success, and Failure

Scala 2.10 introduced `scala.util.Try` as an approach that’s similar to `Option`, but returns failure information rather than a `None`.

The result of a computation wrapped in a `Try` will be one of its subclasses: `Success` or `Failure`. If the computation succeeds, a `Success` instance is returned; if an exception was thrown, a `Failure` will be returned, and the `Failure` will hold information about what failed.

To demonstrate this, first import the new classes:

```
import scala.util.{Try, Success, Failure}
```

Then create a simple method:

```
def divideXByY(x: Int, y: Int): Try[Int] = {
  Try(x / y)
}
```

This method returns a successful result as long as `y` is not zero. When `y` is zero, an `ArithmeticException` happens. However, the exception isn’t thrown out of the method; it’s caught by the `Try`, and a `Failure` object is returned from the method.

The method looks like this in the REPL:

```
scala> divideXByY(1,1)
res0: scala.util.Try[Int] = Success(1)

scala> divideXByY(1,0)
res1: scala.util.Try[Int] = Failure(java.lang.ArithmeticException: / by zero)
```

As with an `Option`, you can access the `Try` result using `getOrElse`, a `foreach` method, or a match expression. If you don't care about the error message and just want a result, use `getOrElse`:

```
// Success
scala> val x = divideXByY(1, 1).getOrElse(0)
x: Int = 1

// Failure
scala> val y = divideXByY(1, 0).getOrElse(0)
y: Int = 0
```

Using a `foreach` method also works well in many situations:

```
scala> divideXByY(1, 1).foreach(println)
1

scala> divideXByY(1, 0).foreach(println)
(no output printed)
```

If you're interested in the `Failure` message, one way to get it is with a match expression:

```
divideXByY(1, 1) match {
  case Success(i) => println(s"Success, value is: $i")
  case Failure(s) => println(s"Failed, message is: $s")
}
```

Another approach is to see if a `Failure` was returned, and then call its `toString` method (although this doesn't really follow the "Scala way"):

```
scala> if (x.isFailure) x.toString
res0: Any = Failure(java.lang.ArithmeticException: / by zero)
```

The `Try` class has the added benefit that you can chain operations together, catching exceptions as you go. For example, the following code won't throw an exception, regardless of what the values of `x` and `y` are:

```
val z = for {
  a <- Try(x.toInt)
  b <- Try(y.toInt)
} yield a * b

val answer = z.getOrElse(0) * 2
```

If `x` and `y` are `String` values like "1" and "2", this code works as expected, with `answer` resulting in an `Int` value. If `x` or `y` is a `String` that can't be converted to an `Int`, `z` will have this value:

```
z: scala.util.Try[Int] =  
  Failure(java.lang.NumberFormatException: For input string: "one")
```

If `x` or `y` is `null`, `z` will have this value:

```
z: scala.util.Try[Int] = Failure(java.lang.NumberFormatException: null)
```

In either `Failure` case, the `getOrElse` method protects us, returning the default value of `0`.

The `readTextFile` method in [Recipe 20.5](#) shows another `Try` example. The method from that example is repeated here:

```
def readTextFile(filename: String): Try[List[String]] = {  
  Try(Source.fromFile(filename).getLines.toList)  
}
```

If the `readTextFile` method runs successfully, the lines from the `/etc/passwd` file are printed, but if an exception happens while trying to open and read the file, the `Failure` line in the match expression prints the error, like this:

```
java.io.FileNotFoundException: Foo.bar (No such file or directory)
```

The `Try` class includes a nice collection of methods that let you handle situations in many ways, including:

- Collection-like implementations of `filter`, `flatMap`, `flatten`, `foreach`, and `map`
- `get`, `getOrElse`, and `orElse`
- `toOption`, which lets you treat the result as an `Option`
- `recover`, `recoverWith`, and `transform`, which let you gracefully handle `Success` and `Failure` results

As you can see, `Try` is a powerful alternative to using `Option/Some/None`.

Using Either, Left, and Right

Prior to Scala 2.10, an approach similar to `Try` was available with the `Either`, `Left`, and `Right` classes. With these classes, `Either` is analogous to `Try`, `Right` is similar to `Success`, and `Left` is similar to `Failure`.

The following method demonstrates how to implement the `Either` approach:

```
def divideXByY(x: Int, y: Int): Either[String, Int] = {  
  if (y == 0) Left("Dude, can't divide by 0")  
  else Right(x / y)  
}
```

As shown, your method should be declared to return an `Either`, and the method body should return a `Right` on success and a `Left` on failure. The `Right` type is the type your method returns when it runs successfully (an `Int` in this case), and the `Left` type is typically a `String`, because that's how the error message is returned.

As with `Option` and `Try`, a method returning an `Either` can be called in a variety of ways, including `getOrElse` or a `match` expression:

```
val x = divideXByY(1, 1).right.getOrElse(0) // returns 1
val x = divideXByY(1, 0).right.getOrElse(0) // returns 0

// prints "Answer: Dude, can't divide by 0"
divideXByY(1, 0) match {
  case Left(s) => println("Answer: " + s)
  case Right(i) => println("Answer: " + i)
}
```

You can also access the error message by testing the result with `isLeft`, and then accessing the `left` value, but this isn't really the Scala way:

```
scala> val x = divideXByY(1, 0)
x: Either[String,Int] = Left(Dude, can't divide by 0)

scala> x.isLeft
res0: Boolean = true

scala> x.left
res1: scala.util.Either.LeftProjection[String,Int] =
  LeftProjection(Left(Dude, can't divide by 0))
```

Although the `Either` classes offered a potential solution prior to Scala 2.10, I now use the `Try` classes in all of my code instead of `Either`.

Discussion

As shown in the Solution, if there's a weakness of using `Option`, it's that it doesn't tell you *why* something failed; you just get a `None` instead of a `Some`. If you need to know why something failed, use `Try` instead of `Option`.

Don't use the `get` method with `Option`

When you first come to Scala from Java, you may be tempted to use the `get` method to access the result:

```
scala> val x = toInt("5").get
x: Int = 5
```

However, this isn't any better than a `NullPointerException`:

```
scala> val x = toInt("foo").get
java.util.NoSuchElementException: None.get
// long stack trace omitted ...
```

Your next thought might be to test the value before trying to access it:

```
// don't do this
scala> val x = if (toInt("foo").isDefined) toInt("foo") else 0
x: Any = 0
```

As the comment says, don't do this. In short, it's a best practice to never call `get` on an `Option`. The preferred approaches are to use `getOrElse`, a `match` expression, or `foreach`. (As with `null` values, I just imagine that `get` doesn't exist.)

See Also

- The `Option` class
- The `Try` class
- The `Either` class

Symbols

! method

 executing command and getting exit code,
 394, 398

 sending messages to actors, 416

!! method, executing command and getting out-
put, 394, 397, 398, 401

" (quotation marks, double), "" surrounding
multiline strings, 6

#&& (AND) operator, 404

#:: method, constructing a Stream, 331

#< operator, redirecting STDIN, 402

#> operator, redirecting STDOUT, 402

#>> method, appending to a file, 403

#| method, pipelining commands, 401

#|| (OR) operator, 404

\$ (dollar sign)

 preceding variable names, 9

\${ }, including expressions inside a string, 10

% method, in build.sbt file, 583

%% method, in build.sbt file, 583

() (parentheses), forcing callers to leave off ac-
cessor methods, 161

*= method, 40

+ method

 adding elements to immutable Map, 348

 adding elements to immutable Set, 363

++ method, 302, 304

 adding elements to immutable Set, 363

 merging two lists, 330

++= method, 302

 adding elements to ArrayBuffer, 335

 adding elements to mutable Queue, 367

 adding elements to mutable Set, 362

 adding multiple elements to mutable maps,
 346

+= method

 adding elements to ArrayBuffer, 335

 adding elements to mutable maps, 346

 adding elements to mutable Queue, 367

 adding elements to mutable Set, 362

 inserting documents into MongoDB, 537

 on Int and other numeric types, 40

 use in build.sbt file, 583

, (comma), adding to numbers, 50

- method

 removing elements from immutable Map,
 348

 removing elements from immutable Set, 365

-- method, 304

 removing elements from immutable Map,
 348

 removing elements from immutable Set, 365

-= method

 deleting elements from mutable Set, 364

 deleting multiple elements from a ListBuffer,
 329

We'd like to hear your suggestions for improving our indexes. Send email to index@oreilly.com.

- removing elements from collections extending TraversableOnce, 336
 - removing elements from mutable maps, 346
 - `--` method
 - deleting elements from ListBuffer, 328
 - deleting elements from mutable Set, 364
 - on Int and other numeric types, 40
 - removing ArrayBuffer elements, 336
 - removing elements from mutable maps, 346
 - `/=` method, 40
 - `/bin/sh` command, `-c` argument, 406
 - `:` (colon), Scala methods ending with, evaluation of, 327
 - `:+` method, 265
 - `::` (cons) method
 - creating a List, 323, 325
 - prepending elements to a list, 326
 - `:::` method, merging or concatenating lists, 303, 330
 - `<-` symbol in for loops, 275
 - ranges created with, 61
 - `<=`, `<`, `>`, and `>=` methods, implemented by Ordered trait, 318
 - `=` (equals sign)
 - `==` (equality) operator
 - testing object equality in Scala, 6
 - testing string equality, 4
 - `==` method, comparing object instances with, 140
 - `=>` symbol as transformer, 218, 220
 - `?` (question mark)
 - `? method`, used with Actors, 445
 - in lazy collections, 331
 - `? :` (ternary operator in Java), 71
 - `@BeanProperty` annotation, 117, 562–565
 - `@SerialVersionUID` annotation, 389, 556
 - `@switch` annotation, 73
 - `@tailrec` annotation, 70
 - `@throws` annotation, 34, 92, 165, 554
 - `@varargs` annotation, 560
 - `\` and `\\` methods, 503
 - `_` (underscore)
 - accessing tuple elements, 313
 - standing for an element in a sequence, 83
 - using to shorten code, 13
 - wildcard character in Scala, xvi, 83, 191, 194
 - catching default in match expression, 75
 - in imports, 196
 - shortening expressions with, 271
 - using variable name for default match, 78
 - `_Seq` methods, 103
 - `_*` operator
 - adapting a sequence to be used as argument for varargs field, 163
 - converting data to pass as multiple parameters to ListMap or LinkedHashMap, 359
 - `{ }` (curly braces)
 - around function body, 221
 - in nested style package notation, 192
 - combining with import statements, 200
 - using to include expressions in strings, 10
 - using to print object fields, 10
 - using to write longer for loops, 61
 - `|` (pipe symbol)
 - separating match conditions for case statements, 76
 - using to create multiline strings, 7
 - `~` operator
 - joining fields with, 497
- ## A
- abstract classes
 - extending traits, 204
 - using traits as, 207
 - when to use, 129
 - abstract fields, using in traits, 206
 - abstract methods, 207
 - (see also abstract classes)
 - implementation by classes implementing traits, 204
 - access modifiers, 148
 - (see also methods, controlling scope)
 - descriptions of Scala modifiers, 152
 - accessor and mutator methods, 99
 - accessor method for constructor val field, 106
 - automatic generation of, in primary constructor, 109
 - case class constructor parameters, 137
 - effect of constructor parameter settings, 107
 - for case classes, 188
 - for constructor var fields, 105
 - forcing callers to leave parentheses off accessors, 161
 - mutator method names for var fields, 103
 - overriding default, 116–119
 - preventing generation of, 119

- Actor model, 412
 - Akka's implementation of, 413
- actorFor method, looking up actors, 436
- actorOf method, 416
 - calling on ActorSystem, 417, 425
- ActorRef, 413, 417, 421, 433
- actors, 411–436
 - benefits of, 411
 - communication between, 419–422
 - creating Actor with class constructor requiring arguments, 418
 - defined, 412
 - getting started with, 414
 - match expressions in, 653
 - methods in Akka actor lifecycle, 422–425
 - monitoring actor's death with watch, 433–436
 - sending message to actor and waiting for reply, 445
 - shutting down Akka actor system, 432
 - Singleton objects as messages for, 179
 - starting, 425
 - stopping, 427–432
 - using gracefulStop method, 430
 - using Kill message, 431
 - using PoisonPill message, 429
 - using stop method, 429
 - switching between different states with become, 446
- actorSelection method, 435
- ActorSystem, 416
 - actorSelection method, 436
 - defined, 416
 - shutdown method, 433
 - stop method, 427
- ActorSystem class, creating actors, 425
- ActorTimeoutException, 431
- add method, using with mutable Set, 362
- afplay system command on Mac OS X, 395
- Akka actor library, 411
- andThen method, 444
 - PartialFunction trait, 240
- annotations
 - adding exception annotations to Scala methods to work with Java, 554
 - annotating variables with a type, 37
 - for return type of implicit methods, 28
 - Scala annotations and Java equivalents, 556
 - using @SerialVersionUID and other annotations, 556
- anonymous functions, 218, 246
 - defining, 218
 - passed into map method, 283
 - returning from body of another function, 236
 - using methods like, 221
- Anorm, using Option, 661
- Ant, 466
 - building a Scala project with, 606–609
- Any class, asInstanceOf method, 172
- AnyRef class, == method, 6
- AnyVal type, specifying for collections, 264
- Apache Commons FileUtils library, 380
- Apache HttpClient library, 491, 513, 514, 516
 - getting URL headers from HttpResponse object, 519
 - sending JSON data to POST URL, 518
 - setting URL headers before making HTTP request, 520
- Apache Ivy, 569, 584
- APIs for libraries, traits commonly used in, 252
- App trait, 176
 - caveats, 178
 - using in Scala shell scripts
 - basic use and header lines, 481
 - passing command-line arguments to, 484
- apply method
 - case classes, 136
 - creating in companion object of regular class, 139
 - defining in companion objects, 110, 185
 - implementing factory method with, 189–190
 - providing multiple case class constructors with multiple apply methods, 186
 - treatment by Scala compiler, 186
- args object, 177
- ARM library, 378
- array notation, using to access a character in a String, 24
- Array.ofDim method, 339, 385
 - decompiling, 340
- ArrayBuffer class, 322
 - deleting elements, 336
 - Gson library working with, 499
 - making your go to mutable sequence, 268
 - methods to add or remove elements, 269
 - reading CSV file rows to, 386

- using to create array whose size can change, 335
 - ArrayList class (Java), 550
 - arrays
 - accessing and changing values by index position, 334
 - Array and ArrayBuffer classes, 321
 - assigning CSV file processing results to 2D array, 385
 - converting findAllIn method results to Array, 18
 - converting Range object to Array, 48
 - creating array whose size can change (ArrayBuffer), 335
 - creating multidimensional arrays, 338
 - using array of arrays, 340
 - deleting elements from an Array, 337
 - different ways to define and populate an Array, 333
 - extracting sequence of elements from, 291
 - iterating over, using for loop, 54
 - looping over elements, xiv
 - manually specifying the type, 264
 - of objects, parsing JSON data into, 501
 - performance of an Array, 261
 - sorting, 337
 - ArrayStack class, 371
 - asInstanceOf method
 - casting numeric types, 173
 - object casting with, 172
 - ask method (of Akka actors), 445
 - assignment, adding elements to mutable maps, 345
 - auxiliary constructors
 - defining, 108–111
 - eliminating need for, 114
 - subclass, inability to call superclass constructor, 128
 - Await.result method, 431, 437, 446
- ## B
- base and radix, handling in conversions of strings to integers, 33
 - @BeanProperty annotation, 117, 562–565
 - “become” approach for actors to switch states, 446
 - best practices in Scala, 635–666
 - creating methods with no side effects (pure functions), 636–644
 - eliminating null values from code, 654–658
 - expression-oriented programming (EOP), 647–650
 - prefer immutable objects, 644–647
 - using match expressions (and pattern matching), 650–654
 - using Option/Some/None pattern, 658–666
 - BigDecimal, 43
 - conversions to other numeric types, 44
 - creating directly from strings, 33
 - BigDecimal class (Java), 44
 - currency formatting with, 51
 - BigInt, 43
 - conversions to other numeric types, 44
 - creating directly from strings, 33
 - BigInteger class (Java), 44
 - binary files, reading and writing, 382
 - Boolean test to break out of loops, 69
 - bounds, 612
 - upper bound, 624, 626
 - view bound, 626
 - Bourne shell, 481
 - break and continue, implementing in Scala, 65–71
 - break example, 66
 - continue example, 67
 - better way to solve problem, 68
 - general syntax, 67
 - nested loops and labeled breaks, 68
 - ways to avoid using break and continue, 69
 - break method, 66
 - breakable method, 66
 - BreakControl exception, 66
 - BufferedSource class, 376
 - close method, 376
 - build tools, 466
 - build.sbt file, 579
 - adding multiple dependencies, 581
 - common methods used in, 583
 - entries as key/value pairs, 582
 - entry for sbt-assembly plug-in, 598
 - example file with one dependency, 581
 - logLevel setting, 597
 - mainClass settings, 593
 - resolvers key, 595
 - specifying main method to run, 592
 - using Build.scala file instead of, 602–604
 - Build.scala file (see project/Build.scala file)

build.xml file (Ant), 606

init target, 608

by method, setting step in ranges, 47

Byte type, 31

C

callback methods, futures, 443

onSuccess and onFailure, 439

canEqual method, 142

Casbah driver, 533

creating Casbah-style queries, 542

case classes, 99, 640

adding uninitialized var field to, 122

as inner class, 143

constructor parameters, val by default, 107

creating new class instances without new

keyword, 186

extracting case class fields in match expressions, 88

final, 644

generating auxiliary constructors for, 110

generating boilerplate code with, 136–140

in match expressions, 86–87

with variable-binding pattern, 84

MailServer, generating from JSON string, 500

methods generated by, 140–143

providing multiple constructors for, 187–189

toString implementation, 137

used to match JSON data with Lift-JSON, 492

case statements

adding if expressions (guards) to, 87

matching multiple conditions with single statement, 76

pattern matching examples, 79

case-insensitive string comparisons, 5

casting

from one numeric type to another, 36

object, 172

catch clause, 92

(see also try/catch/finally)

throwing exception from, 92

chaining methods, 4

Char type, 31

charAt method, 24

checked exceptions, 165

in Java, 166

Child actor, 427

class members

renaming on import, 196

static imports of, 197

ClassCastException, 173

classes, 99

adding to REPL classpath, 462

assigning field to a block or function, 121

calling superclass constructor, 127

controlling visibility of constructor fields, 104–108

creating inner classes, 143

creating primary constructor for, 100–104

default values for constructor parameters, 114

defining auxiliary constructors, 108–111

defining equals method to compare object instances, 140–143

defining properties in abstract base class, 131

determining class of an object, 174

extending traits, 203

generating boilerplate code with case classes, 136–140

handling constructor parameters when extending, 124–127

hiding during import process, 196

importing into current program scope, 193

making serializable, 389

mixing traits in, 208

multiple, in a single file, 201

overriding default accessors and mutators, 116–119

preventing accessor and mutator methods from being generated, 119

providing private primary constructor, 112

renaming on import, 195

setting uninitialized var field types, 122

traits can be mixed into, limiting, 209

traits inheriting from, 210

use of @BeanProperty annotation on fields, 562

using a trait by inheritance, limiting, 209

using generic types, 614–617

when to use abstract class, 129

classOf function, 174

classpath

building for Scala shell script, 482

REPL, adding JAR files and classes to, 461

- clear method
 - deleting all elements in a map, 346
 - emptying mutable Stack with, 370
 - removing all elements from ArrayBuffer, 336
 - using with mutable Set, 364
- closing files
 - automatically, 377
 - manually closing BufferedSource, 376
- closures, 229–234
 - comparison to techniques in Java, 233
 - defined, 230
 - using with other data types, 232
- code blocks
 - assigning class field to, 121
 - for/yield processing after yield keyword, 274
 - import statements in, 199
 - pasting and loading into REPL, 459–461
- code examples, using, xxii
- code listings in this book, xxi
- collect method, taking partial function as input, 241
- collections, xv, 245–319
 - building new collection from input collection, using for/yield, 55
 - choosing a collection class, 250
 - main categories of classes, 250
 - maps, 253
 - sequences, 250
 - sets, 253
 - strict and lazy collections, 254
 - types acting like collections, 254
 - choosing a collection method to solve a problem, 255
 - common collection methods, 256
 - common methods for maps, 259
 - immutable collection operators, 259
 - methods organized by category, 255
 - mutable collection methods, 258
 - classes containing, generating JSON string from, 495–500
 - collection features in String, 3
 - combining map and flatten with flatMap method, 286–289
 - converting iterators to, 279
 - converting to string with mkString, 318
 - creating and using enumerations, 311
 - creating lazy view on, 306–309
 - creating whose elements are all of some base type, 624–627
 - declaring type when creating, 264
 - deleting documents from MongoDB collection, 545
 - extracting sequence of elements from, 291–293
 - extracting unique elements from a sequence, 300
 - filtering with filter method, 289
 - flattening list of lists with flatten, 285
 - important concepts when working with
 - methods of classes, 245
 - iterating over, using collection methods, 57
 - iterating over, using for and for/yield loops, 56
 - Java, going to and from, 549–554
 - conversion tables, 552
 - going from Scala collections to Java collections, 553
 - looping over, 54
 - looping over with for loop, 272–276
 - looping over with foreach, 270
 - making ArrayBuffer your go to mutable sequence, 268
 - making immutable collections covariant, 622
 - making mutable collections invariant, 620
 - making Vector your go to immutable sequence, 266
 - merging sequential collections, 302
 - merging two sequential collections into pairs with zip, 304
 - mutable and immutable, 266
 - mutable variables with immutable collections, 265
 - performance of, 261–264
 - maps and sets, 263
 - performance characteristic keys, 261
 - sequential collections, 262
 - populating
 - using Range with map method, 372
 - using tabulate method, 372
 - with a Range, 309
 - random-length collection of printable characters, 46
 - Scala collections hierarchy, 246
 - high-level view of collections, 247
 - maps, 249
 - more collection classes, 250
 - sequences, 248
 - sets, 249

- strict and lazy collections, 250
- searching MongoDB collection, 539–542
- sorting, 315–318
- splitting sequences into subsets, 293
- transforming one to another with `for/yield`, 279–282
- transforming one to another with `map` method, 282–284
- tuples, 312–315
- updating documents in MongoDB collection, 542
- using collection methods on strings, 14
- using iterators, 278
- using `Option` with, 660
- using parallel collections, 448
- using `zipWithIndex` or `zip` to creating loop counters, 276–278
- walking through with `reduce` and `fold` methods, 295–300
- working with, using `Lift-JSON` library, 493
- combinators, 444
- command-line tasks, 453–490
 - accessing command-line arguments from Scala shell script, 483
 - adding JAR files and classes to REPL class-path, 461
 - compiling with `scalac` and running with `scala`, 465
 - disassembling and decompiling Scala code, 466–471
 - faster compiling with `fsc`, 479
 - finding Scala libraries, 471
 - generating documentation with `Scaladoc`, 472–479
 - getting started with REPL, 454
 - making Scala scripts run faster, 489
 - pasting and loading blocks of code into REPL, 459–461
 - prompting for input from Scala shell script, 485–489
 - REPL (Read-Eval-Print-Loop) tool, 453
 - running shell command from REPL, 462–464
- commas (`,`), adding to numbers, 50
- companion objects
 - `apply` method, 189–190
 - `apply` methods, 110
 - defining, 185
 - providing multiple constructors with, 186
 - creating, 185
 - creating static members with, 180–182
 - defined, 113
 - `getInstance` method in, 112
- `compare` method, 317
- compiling
 - faster command-line compiling with `fsc`, 479
 - Scala projects in SBT, 574
 - continuous compiling, 577
 - using `scalac`, 465
- complex numbers, working with, 32
- `concat` method, 303
 - `List` class, 330
- concrete fields, using in traits, 206
- concurrency
 - simple concurrency with futures, 436–445
 - using parallel collections, 448
- conditional operator in Java (`? :`), 71
- `cons` method (`::`) for creating a `List`, 323
- `Conscript`, 506, 573
- Console, 375
 - fields to use with `print` methods to control display, 488
 - `read*` methods, 485
 - listing in Scala REPL, 485
 - `readf` methods, 486
- constant patterns, 83
- constraints, type constraints, 613
- constructor patterns, 83
- constructors
 - actor's constructor, 424
 - taking arguments, 418
 - calling superclass constructor, 127
 - controlling visibility of constructor fields, 104–108
 - creating primary constructor, 100–104
 - comparison with Java, 102
 - mutator methods for `var` fields, 103
 - defining auxiliary constructors, 108–111
 - handling parameters when extending a class, 124–127
 - private primary constructor, 112
 - providing default values for parameters, 114
 - providing multiple constructors for case classes, 187–189
 - Scala class, use of `@BeanProperty` annotation, 562

- container:start command (SBT), 572
 - contains method
 - testing for existence of a key in a map, 353
 - testing if set contains an element, 362
 - using with valuesIterator on a map, 354
 - context object, 425
 - actorOf method, 425
 - actorSelection method, 436
 - stop method, 427
 - watch method, 433, 434
 - continue, implementing in Scala, 65
 - continue example, better ways to solve problem, 68
 - example code, 67
 - general syntax for, 68
 - control structures, 53–97
 - accessing value of default case in match expression, 78
 - adding if expressions (guards) to case statements, 87
 - assigning match expression result to variable, 77
 - creating your own, 95
 - declaring variable before using in try/catch/finally block, 92–95
 - for and foreach, looping with, 54–60
 - for comprehension, 63–65
 - implementing break and continue, 65–71
 - matching multiple conditions with one case statement, 76
 - matching one or more exceptions with try/catch, 91
 - using case classes in match expressions, 86
 - using if expression like ternary operator, 71
 - using match expression instead of isInstanceOf, 88
 - using match expression like switch statement, 72–76
 - using pattern matching in match expressions, 79–86
 - working with a List in a match expression, 89
 - copy method, case classes, 137
 - CopyBytes class, 382
 - count method, using with anonymous function instead of continue construct, 68
 - counters
 - in for loops, 55, 273
 - multiple counters, 60–62
 - using zipWithIndex or zip to create loop counters for sequential collections, 276
 - covariant immutable collections, 622
 - :cp command (REPL), 462
 - CreateChild case class, 427
 - CSV files, processing, 384
 - CSV strings, 318
 - curl command, 513, 525
 - currency, formatting
 - commas, 50
 - decimals, 50
- ## D
- data types (see types)
 - database views, 308
 - databases, 527–548
 - accessing MongoDB document ID field, 544
 - connecting to database with Spring Framework, 530
 - connecting to MongoDB and inserting data, 533–537
 - connecting to MySQL with JDBC, 528
 - deleting documents from MongoDB collection, 545
 - inserting documents into MongoDB with insert, save, or +=, 537
 - searching a MongoDB collection, 539–542
 - Slick library, 547
 - updating documents in MongoDB collection, 542
 - dates, processing, 32
 - decimals, in numbers or currency, 50
 - decompiling code, using decompiler, 468
 - default global execution context, 437
 - DelayedInit trait, 178
 - dependencies
 - adding to SBT project, 579
 - controlling version of managed dependency to use, 584
 - managing with SBT, 581–584
 - using Scala project on GitHub as SBT project dependency, 593
 - deprecation warnings, 457
 - dequeue method, 368
 - dequeueAll method, 368
 - dequeueFirst method, 368
 - deserializing JSON string into an object, 500

- diff method, 303
 - using to get relative complement of two sets, 303
- directories
 - creating project directory structure for SBT, 570–574
 - listing files in, 391
 - listing subdirectories of, 392
 - running a process in a different directory, 406
- disassembling Scala code, 466, 469
- distinct method
 - calling on collections, 300
 - using with your own classes, 300
- documentation
 - generating project API documentation with SBT, 590
 - generating with Scaladoc, 472–479
- domain-specific library (DSL) (see Lift-JSON DSL)
- Double type, 31
- drop method, 291
- dropRight method, 292
- dropWhile method, 291
- DSL (domain-specific library) (see Lift-JSON DSL)
- Duck Typing (structural types), 618
- durations, 437

E

- Eclipse
 - using SBT with, 588
 - .classpath file, 589
 - .project file, 589
- Either, Left, and Right classes, 664
- enqueue method, 367
- Enumeration class, 311
- enumerations
 - creating and using, 311
 - static imports of, 198
- environment variables, setting when running
 - external commands, 407
- EOP (expression-oriented programming), 647–650
 - expressions versus statements, 648
 - languages, 78
 - model, 217
- equality, 4
 - (see also == operator, under Symbols)

- testing for String instances, 4
- equals method
 - case classes, 137
 - defining for a class, 140–143
 - using with distinct, 300
- equalsIgnoreCase method, String class, 5
- equivalence relation, 142
- Exception object, 662
- exceptions
 - access to information, other possibilities for, 35
 - adding exception annotations to Scala methods, 554
 - declaring that a method can throw, 165
 - in Java, 166
 - matching one or more with try/catch, 91
 - opening and reading files, 379
 - Scala methods throwing, 34
- execution context, 437
- ExecutionContext, futures and, 443
- ExecutionContext.Implicits.global, 437
- exists method, using with valuesIterator on
 - map, 354
- expression-oriented programming (EOP), 647–650
 - expressions versus statements, 648
 - languages, 78
 - model, 217
- extends keyword, 126, 127, 204
 - specifying superclass constructor to call, 127
 - using to implement Java interfaces, 203
- extends with keywords, 203, 204
 - using to implement Java interfaces, 216
- external commands
 - building a pipeline of commands, 401
 - executing and using STDOUT, 397–399
 - executing from Scala application, 394
 - handling STDOUT and STDERR for, 399
 - handling wildcard characters in, 405
 - methods to execute, 408
 - redirecting STDOUT and STDIN of, 402
 - setting environment variables when running, 407
 - versus shell built-in commands, 396

F

- f string interpolator, 11
 - formatting numbers, 49
- factorial function, recursive, 70

- factory method, 110
 - implementing in Scala with `apply`, 189–190
- Failure class, 633, 662
- failure, handling by actors, 413
- `fallbackTo` combinator method, 444
- feature warnings, 458
- fields
 - abstract, in abstract base class or trait, 131, 132
 - constructor, 104
 - adding private to `val` or `var`, 106
 - case classes, 107
 - effect of settings, 107
 - `val` fields, 106
 - `var` fields, 105
 - without `val` or `var` designation, 106
 - defining with private or private[this] access modifiers, 119
 - initializing class field using block of code or calling function, 121
 - object-private, 120
 - setting uninitialized `var` field types in a class, 122
 - static, creating on a class, 180
 - using abstract and concrete fields in traits, 206
- File class (Java), 403
 - `listFiles` method, 391
- FileFilter object (Java), `accept` method, 391
- FileInputStream class, 382
- FileOutputStream class, 382
- files, 375–394
 - listing in a directory, 391
 - logging external command output to, 402
 - opening and reading text file, 375–381
 - pretending that a String is a File, 387
 - processing CSV files, 384
 - processing every character in a text file, 383
 - writing text files, 381
- FileWriter class, 381
- filter expressions, 62
 - (see also guards)
- filter method, xvi
 - calling on String to create new String, 14
 - comparison to other collection filtering methods, 290
 - filtering a collection, 289
 - filtering queries in Slick, 548
 - keys to using, 291
 - on collections, combined with anonymous function, 246
 - trimming list to contain only directories, 392
 - using on immutable Set, 365
 - using on strings, 2
 - using with maps, 355, 356
- filtering collections
 - main methods for, 290
 - maps, 354–356
 - mutable and immutable maps, 355
 - mutable maps, 355
 - methods for, 255
 - methods returning contiguous sequence of elements, 291
 - using filter methods on immutable Set, 365
 - using filter methods to remove elements from collections, 337
 - using filter methods with Array, 337
- filtering files, 391
- `filterKeys` method, 355
 - using a predicate, 355
- filters, 62
 - (see also guards)
- `find` command (Unix), 398
- `findAllIn` method, 18
- `findFirstIn` method, 18
 - handling Option returned by, 19
- `flatMap` method, 286–289
- `flatten` method, 256
 - combining with `map` in `flatMap`, 286–289
 - flattening list of lists, 285
 - flattening list of strings into list of characters, 286
 - using on sequence of Some and None elements, 286
 - using with `mkString`, 319
 - using with other sequences, 285
- Float type, 31
- floating-point numbers, comparing, 41–43
- fluent programming style, 4, 167–170
- `foldLeft` method, 295, 298
- `foldRight` method, 295, 298
- for comprehension, 54, 56, 274, 281
 - (see also for/yield loops)
 - creating, 63–65
- for loops, 54–60
 - <- symbol in, 275
 - alternatives to break and continue, 69
 - Boolean test to break out of, 69

- counters in, 55
- generators and guards, 56
- iterating over keys and values in a Map, 56
- looping over collections, 272–275
 - for/yield construct, 274
 - Maps, 275
- ranges in, 47, 372
- returning values from, 55
- Scala versus Java versions, 53
- translation by compiler, 57–60
 - for yield/construct, 59
 - with guard condition, 59
- treating String as sequence of bytes, 2
- treating String as sequence of characters, 2
- using with multiple counters, 60–62
- using with strings, 14
- using zipWithIndex as loop counter, 276
- with embedded if statements (guards), 62, 275
- yield statement with, 63
 - (see also for comprehension; for/yield loops)
- for/yield loops, 14
 - building new collection from input collection, 274, 279–282
 - using guards, 281
- getting list of directory names, 393
- joining three futures together, 441
- returning values, 55
- reversing keys and values of a map, 353
- using to modify random-length range of numbers, 46
- foreach method, 176
 - adding after getBytes called on String, 17
 - collections, 256
 - for loops' translation by compiler, 57, 58
 - iterating over collections, 57, 270
 - functions or methods called from foreach, 271
- JavaConversions object, 550
- operating on each element without returning result, 14
- printing fields of an Address, 123
- Traversable trait, 248
- using match expression with to iterate over map elements, 351
- using Option in, 20
- using to iterate over iterator's elements, 278
- using with zipWithIndex, 276, 277
- format method
 - calling on strings, 12
 - formatting numbers, 50
- Format object (Play Framework), 522
 - reads method, 524
- forms, verification in Play Framework, using match expressions, 653
- free variables, 230
- fromURL method, Source object, 514
- fsc command, 479
- function literals, 218–219
 - declaring, 220
- function value, 235
- functional programming (FP), 217–244
 - case classes, 108
 - creating a function that returns a function, 236–238
 - creating partial functions, 238
 - defining method that accepts more complex function as parameter, 226–229
 - defining method that accepts simple function as parameter, 223–225
 - functional versus imperative style, 393
 - having no side effects, 162
 - real-world example of use, 242–244
 - using closures, 229–234
 - using function literals (anonymous functions), 218
 - using functions as variables, 219–223
 - using partially applied functions, 234
 - using recursive algorithms, 70
 - writing pure functions, 636–644
- FunctionN traits, 222
- functions
 - assigning class field to results of, 121
 - assigning existing function or method to a function variable, 222
 - defining as method parameter, 223–225
 - general syntax, 224
 - more complex functions, 226–229
 - match expressions as body of, 652
 - methods versus, 16
 - partially applied, 222, 234
 - returning a function, 236–238
 - using as variables, 219–223
 - without side effects (see pure functions)
- future method, 440
- futures, 436–445
 - and ExecutionContext, 443

- callback methods, 438, 443
 - onSuccess and onFailure, 439
- creating a method to return a Future[T], 440
- creating and blocking to wait for its result, 436
- creation of Future object, 437
- joining together with for comprehensions or combinators, 444
- using as wrapper around Source.fromURL method call, 517

G

- generators, 56
 - writing for loops with, 61
- generic types
 - creating classes that use, 614–617
 - generic type parameter symbols, 616
 - creating method that takes, 617
- GenSeqLike, distinct method, 300
- get method
 - finding map values, 350
 - not using with Option, 665
- GET requests
 - accessing Scalatra web service GET parameters, 509
 - creating GET request web service with Play Framework, 521–524
 - creating simple client for, 514
- getBytes method, calling on String and passing block of code to foreach method, 16
- getClass method, 175
- getCurrencyInstance method, NumberFormat, 50
- getInstance method, 112
- getIntegerInstance method, NumberFormat, 50
- getLinesUppercased method, 388
- getOrElse method
 - calling on Option value, 20
 - using to find map value, 350
- getter and setter methods, 100
- Giter8, 503, 506
 - requirements for, 573
 - using to create SBT project directory structure, 571
 - notes, 572
- GitHub
 - Giter8 templates, 571
 - Scala project on, using as SBT project dependency, 593

- Google, Gson library, 492
- gracefulStop method, 428
 - stopping an actor with (example), 430
- groupBy method, 293
- grouping methods, 255
- Gson library, 492, 494
 - constructing JSON string, 518
 - creating JSON string from Scala object, 493
 - generating JSON strings from classes containing collections, 498
- guards, 56
 - adding to case statements, 87
 - attempting to use with map method, 284
 - for loop with, translation by compiler, 59
 - in for/yield loops, 280, 281
 - using for loop with, 62, 275

H

- hashCode method
 - case classes, 137
 - defining, 140
 - if expression in, 71
 - using with distinct, 300
- hex values, storing as Int or Long, 38
- HTMLCleaner library
 - adding library dependency to SBT build.sbt file, 581
 - using Maven repository with SBT, 605
- HTTP requests, setting URL headers for, 520
- HTTP response headers, accessing after making a request, 519
- HttpClient library (see Apache HttpClient library)
- HttpPost object, 513
- HttpTea, 521
- hyperlinks, Scaladoc hyperlink tags, 478

I

- i argument, using to load code into Scala REPL, 461
- if expressions
 - adding to case statements, 87
 - attempting to use as guard with map method, 284
 - guards in for loops, 62, 275
 - guards in for/yield loops, 280, 281
 - using in for loops, 56
 - using like ternary operator, 71

- if/then/else statements, match expressions as replacement for, [651](#)
- immutable collections
 - common operators (methods) on, [259](#)
 - maps
 - common methods for, [259](#)
 - mutable variables with, [265](#)
 - Vector class, go to immutable sequence, [266](#)
- immutable objects
 - case classes as immutable records, [138](#)
 - preferring, [644–647](#)
 - using val + mutable and var + immutable, [645](#)
- imperative programming, [15](#), [393](#)
- implicit classes
 - defining, [25](#)
 - defining methods on, [27](#)
 - annotating method return type, [28](#)
 - returning other types, [29](#)
 - putting in package objects, [26](#)
 - putting into objects, [25](#)
- implicit conversions, [2](#)
 - using to add methods to closed String class, [3](#)
 - using to add methods to String class, [25–29](#)
- implied loops, [246](#)
- import statements, using anywhere in Scala, [194](#), [199](#)
- imports
 - hiding a class during import process, [196](#)
 - one or more members into current program scope, [193](#)
 - renaming members on import, [195](#)
 - static, [197](#)
- incrementing and decrementing numbers
 - replacements for ++ and --, [39–41](#)
- indexed sequences, [327](#)
 - ArrayBuffer class, [268](#)
- IndexedSeq, [248](#), [252](#)
 - creating instance of, returning a Vector, [268](#)
- infix notation, [48](#)
- informational methods (collections), [256](#)
- inheritance
 - classes using a trait by inheritance, limiting, [209](#)
- init method, [293](#)
- inner classes, creating, [143](#)
- Inner.breakable, [69](#)
- input, prompting users for, from Scala shell script, [485–489](#)
- InputStream class (Java), [515](#)
- insert method, Casbah MongoClient class, [537](#)
- Insert object, [535](#)
- Int type, [31](#)
- Integer class (Java), parseInt method, [33](#)
- interfaces
 - extending Java interface like a trait, [216](#)
 - traits versus, [203](#)
 - using traits as, [203](#)
- interpreter
 - savecompiled argument for Scala interpreter, [489](#)
 - starting Scala interpreter in REPL, [463](#)
- intersect method, [302](#)
- intersection of two sets, subtracting, [304](#)
- invariant, declaring a type as, [620](#)
- IOException (Java), [515](#)
- isDefinedAt method, [241](#)
- isInstanceOf method, using match expression instead of, [88](#)
- isValid methods, [37](#)
- Iterable trait, [247](#)
- iterators
 - converting to collections, [279](#)
 - creating for tuples, [314](#)
 - keysIterator and valuesIterator methods on maps, [352](#)
 - using in Scala application, [278](#)
 - working with text files, [377](#)

J

- JAD (decompiler), [468](#)
- JAR files
 - adding to REPL classpath, [461](#)
 - creating with Ant, [608](#)
 - deploying single, executable file from SBT, [597–601](#)
 - generated for Scala scripts, [490](#)
 - SBT projects packaged as, specifying main method to run, [591](#)
- Java
 - arrays, [334](#)
 - BigInteger and BigDecimal classes, [44](#)
 - checked exceptions, [166](#)
 - .class, [174](#)
 - class declaration, comparison with Scala, [102](#)

- closure-like techniques, 233
 - conditional (or ternary) operator (`? :`), 71
 - control structures, 53
 - converting numeric types using casting, 36
 - declaration of classes, class constructors, and field visibility control, 99
 - deep copy technique using serialization, 389
 - `FileInputStream` and `FileOutputStream` classes, 382
 - files, opening and reading, 380
 - files, writing with `PrintWriter` or `FileWriter`, 381
 - `HTMLCleaner` library, 581, 605
 - importing code into current scope, 194
 - inner classes, 144
 - interacting with, 549–568
 - adding exception annotations to Scala methods, 554
 - annotating varargs methods, 560
 - Java code requiring `JavaBeans`, 562–565
 - using `@SerialVersionUID` and other annotations, 556
 - using Spring Framework in Scala application, 557–560
 - wrapping traits with implementations, 565–568
 - interfaces, 203
 - extending like a trait, 216
 - iterators, 278
 - Joda Time project, 32
 - List classes, 321, 323
 - method returning multiple values, 159
 - methods, 147
 - OOP approach to writing classes and methods, 639
 - processes running on system, listing, 401
 - static methods, 181
 - switch statement, 651
 - match expressions versus, 54
 - traits with implemented methods, inability to call, 129
 - Java Decompiler Project, 469
 - Java interpreter, running JAR file created by SBT, 599
 - `java.lang.OutOfMemoryError` errors, 332
 - `java.text.NumberFormat` class, 50
 - `java.util.ArrayList` class, 550
 - `java.util.NoSuchElementException`, 350
 - `java.util.TreeSet`, 366
 - `JavaBeans`
 - getters and setters, 117
 - Java code requiring, interacting with, 562–565
 - `JavaConversions` object, 550–553
 - `asScalaBuffer` method, 551
 - `mapAsScalaMap` method, 552
 - Scala to Java one-way conversions, 553
 - two-way conversions provided by, 552
 - `JavaConverters` object, 554
 - `JavaNet1Repository`, 596
 - `javap` command
 - disassembling case class files, 126, 126, 187, 470
 - disassembling class files, 73, 117, 119, 467, 469, 564, 567
 - JDBC, 527
 - connecting to MySQL with, 528
 - `JLineCompletion` class, 455
 - Joda Time project, 32
 - `js.implicit.ly`, 472
 - JSON, 491
 - conversions of JSON strings to and from objects, 522
 - creating JSON string from classes having collections, 495–500
 - creating JSON string from Scala object, 491
 - Gson library solution, 493
 - Lift-JSON library solution, 492
 - creating simple Scala object from JSON string, 500
 - data sent as POST request, handling in Scalatra web service, 510–514
 - parsing JSON data into array of objects, 501
 - POSTing JSON data to Play Framework web service, 524
 - representations of model objects, 521
 - sending JSON data to POST URL, 518
 - `Json.toJson` method, 522
 - `Json4s` library, 513
 - `JValue` object, 501
 - JVM
 - switches, information on, 74
 - type erasure in, 81
- ## K
- key/value pairs, merging two sequential collections into, using `zip`, 304
 - `keys` method, 351, 352

- keySet method, 352
- keyIterator method, 352
 - using with max method, 361
 - using with reduceLeft method, 361
- Kill message, sending to an actor, 431

L

- labeled breaks, 68
- lazy collections, 250, 254
- lazy fields, 122
- libraries
 - finding Scala libraries, 471
 - JSON, 492
 - publishing your Scala library with SBT, 601
- library APIs, sequence traits commonly used in, 252
- libraryDependencies setting in SBT, 579
 - adding dependencies singly or in multiples, 581
 - adding Java HTMLCleaner library, 581
 - managed dependencies, forms of, 582
 - Maven repository library used with SBT, 605
 - specifying version of managed dependency, 584
- Lift Framework, 491, 493
- Lift-JSON DSL, 494, 495–498
- Lift-JSON library, 492
 - adding dependency to SBT project, 512
 - converting JSON string into Scala object, 500
 - generating JSON strings from Scala Map, 493
- limit method, using to limit query results returned, 542
- line separators, 376
- linear sequences, 327
 - ListBuffer class, 268
- LinearSeq, 248, 252
- lines method, 396
- lines_! method, 396, 399
- LinkedHashMap class, 343
 - sorting, 358
 - storing sorted map data in, 359
- LinkedHashSet class, 366
- ListBuffer class, 268, 270, 327
 - creating and converting to a List, 324
 - deleting elements from, 328
 - Scaladoc description of, 325
- listFiles method, File class (Java), 391, 392
- ListMap class, 343
 - sorting by key or value, 357
 - storing sorted map data in, 358
- lists
 - adding elements to existing List, 325
 - converting Range object to List, 48
 - creating a mutable list, 324
 - creating and populating a List, 322
 - deleting elements from a List or ListBuffer, 328
 - differences between Scala List and Java List, 552
 - exploring to* methods of List in REPL, 455
 - flattening list of lists with flatten method, 285
 - for loops iterating over, translation by compiler, 59
 - LinearSeq and, 248
 - merging or concatenating, 330
 - passing anonymous function to List's filter method to create new List, 218
 - performance of a List, 261, 262
 - prepending elements of one list to another, 303
 - Scala List class, 321, 323
 - Scaladoc description of List class, 330
 - Stream class, similarity to List, 331
 - working with a List in a match expression, 89
- :load command (REPL), 460
- Loan Pattern, 377
- locale, setting with getIntegerInstance method, 50
- logging
 - setting SBT log level, 597
- Long type, 31
- looking up actors, 435
- lookupswitch, compiling match expression to, 73, 74
- loops, 2
 - (see also for loops; for/yield loops; foreach method)
 - implied, 246
 - using in methods to handle varargs fields, 164

M

- main method, 176
 - manually implementing in an object, 177

- overriding, 178
- specifying which to run in SBT, 591
- using in a Scala script, 481
- mainClass settings, build.sbt file, 593
- MAMP, 529
- managed dependencies, 582
- Map class, 249
 - foreach method, 272
- map method, xv, xvi
 - calling getName on files, 393
 - combinator, 444
 - combining map and flatten with flatMap, 286–289
 - creating uppercase string from input string, 13
 - for loops and, 57
 - for/yield construct (for comprehension) and, 64
 - for/yield constructs translated to, 60
 - transforming one collection to another with, 282–284
 - understanding how map works, 15
 - using to call trim method on strings, 8
 - using with a Range, 311, 372
 - views on collections and, 308
 - writing for loop to work like, 14
- maps, 322
 - accessing values, 349
 - adding, updating, and removing elements in immutable map, 347
 - adding, updating, or removing elements in mutable map, 345
 - choosing a map class, 343
 - map classes and traits, 344
 - common map choices, mutable and immutable, 253
 - common map classes, 249
 - common methods for, 259
 - creating, 341
 - immutable map, 341
 - mutable map, 341
 - defined, 250
 - filtering, 354–356
 - finding largest key or value in, 360
 - generating JSON string from Scala Map, 493
 - getting keys and values from, 352
 - iterating over elements, 350
 - iterating over with for loop, 275
 - looping over a Map using for loop, 56
 - performance characteristics, 263
 - reversing keys and values, 352
 - sorting existing map by key or value, 357–360
 - _* operator, 359
 - testing for existence of a key or value in, 353
- mapValues method, 351
- match expressions, 650–654
 - accessing value of default case in, 78
 - assigning result to a variable, 77
 - in actors, 653
 - replacement for Java switch statements and if/then statements, 651
 - resemblance to Java switch statement, 54
 - use in catch block of try/catch, 91
 - using as body of function or method, 652
 - using case classes in, 86
 - using in try/catch expressions, 651
 - using instead of instanceof, 88
 - using like switch statement, 72–76
 - @switch annotation, 73
 - handling default match, 75
 - using Option in, 20
 - using pattern matching, 79–86
 - adding variables to patterns, 84–85
 - constant patterns, 79–83
 - constructor patterns, 79–83
 - sequence patterns, 79–83
 - Some and None, 85
 - typed patterns, 79–83
 - using when searching MongoDB database, 541
 - using with foreach method to loop over map elements, 351
 - using with Option/Some/None pattern, 653
 - working with a List, 89
- MatchError, 79
- math, more powerful math classes, 32
- mathematical methods (collections), 256
- Maven, 466, 569
 - Maven2 repository, 596
 - using repository library with SBT, 604
- max method, using on map, 360
- MaxValue, checking for numeric types, 44
- memory allocation, views and, 309
- messages (sent between actors), 419–422
- method chaining, 167–170
- methods, 147–170
 - abstract, defining, 130

- ul style="list-style-type: none;">
- adding exception annotations to work with Java, 554
- assigning class field to results of, 121
- assigning existing method to a function variable, 222
- avoiding returning null from, 656
- calling on a superclass or trait, 152–154
 - selecting the trait, 153
- chaining, 4
- collection (see collections)
- controlling scope, 148–152
 - object-private scope, 148
 - package scope, 150
 - package-level control, additional, 150
 - private scope, 149
 - protected scope, 149
 - public scope, 151
- declaring can throw exceptions, 165
- declaring in traits, 203
- defined in object, not a class, 179
- defining method that takes function as parameter, 223–225
 - more complex functions, 226–229
- differences between Java and Scala, 147
- forcing callers to leave parentheses off accessor methods, 161
- functions versus, 16
- implemented methods of Scala traits, Java and, 565
- match expressions as body of, 652
- passing a class to, 174
- requiring implementation by class attempting to mix in a trait, 213
- returning multiple items (tuples), 159–161
- returning Option from, 659
- setting default values for parameters, 154–157
- static, creating on a class, 180
- symbols implemented as, in Scala, 41
- taking simple generic type, 617
- taking variable-argument fields, 163–165
- throwing exceptions in Scala, 34
- traits providing method implementations, 205
- upper bound definition on type parameter, 626
- using like anonymous functions, 221
- using parameter names when calling, 157
- vararg, annotating, 560
 - without side effects (see pure functions)
- mixins, using traits as, 208
- mkString method, 256
 - converting collection to a String, 318
- mongo command-line client, 536
- MongoCollection class, 536, 537
 - find and findOne methods, 539
 - find* methods, 542
 - findAndModify or update methods, 543
 - findAndRemove method, 545
- MongoCursor, 541
- MongoDB, 491
 - accessing document ID field, 544
 - connecting to and inserting data, 533–537
 - deleting documents from collection, 545
 - inserting documents with insert, save, or +=, 537
 - searching a collection, 539–542
 - updating documents in a collection, 542
- MongoDBObject, 535
 - getting document ID from, 544
- mounting servlets, 508
- multidimensional arrays, creating, 338
- multiline strings, creating, 6
- mutable collections
 - common methods on, 258
 - maps, common methods for, 259
 - sequences, making ArrayBuffer go to sequence, 268
- MySQL, connecting to, using JDBC, 528
- ## N
- named parameters, 114
 - using for constructors, 115
 - using when calling methods, 156, 157
 - namespace collisions or confusion, avoiding, 195
 - NegativeInfinity, 45
 - nested break statements, 68
 - new keyword
 - creating object instances without using, 110, 185–189
 - newline characters, 376
 - replacing with blank spaces in multiline string creation, 7
 - unexpected, in external command output, 398
 - Newman DSL, 517
 - Newton's Method, implementation of, 243–244

- Nil element
 - ending a list, 323
 - last element in List, 89
- Nil\$, 164
- None and Some, using with match expressions, 85
- NoSuchElementException, 350
- nscala-time project, 32
- Null Object Pattern, 657
- null values
 - eliminating from your code, 654–658
 - converting null to Option or something else, 657
 - initializing var fields with Option, not null, 654
 - not returning null from methods, 656
 - initializing fields with Option/Some/None pattern instead of, 135
 - string comparisons and, 6
- NullPointerException (string equality comparisons), 5
- Number type, specifying for collections, 264
- NumberFormat class (Java)
 - getCurrencyInstance method, 50
 - getInstance method, 50
 - getIntegerInstance method, 50
- NumberFormatException, 33
- numbers
 - caution with methods reading numeric values, 486
 - comparing floating-point numbers, 41–43
 - complex numbers and dates, 32
 - converting Strings to Scala numeric types, 32–36
 - creating range, list, or array of, 47
 - formatting numbers or currency, 49
 - generating random numbers, 45–47
 - handling very large numbers, 43–45
 - incrementing and decrementing, replacements for ++ and --, 39–41
 - Scala's built-in numeric types, 31
- Numeric trait, 628
- numeric types, 31
 - casting with asInstanceOf method, 173
 - converting between, 36–37
 - data ranges of Scala's built-in types, 31
 - overriding default type, 37–39
- numeric var fields, setting type on, 124

O

- object instances
 - comparing equality of, 140–143
 - creating, 39
- object keyword, 171
 - creating Singletons with, 178
- object-private fields, 120
- object-private scope, 148
- object-relational mapping (ORM) tools, 527
- objects, 171–190
 - adding traits to object instance, 215
 - array of, parsing JSON data into, 501
 - casting, 172
 - companion object, 113
 - creating JSON string from Scala object, 491
 - creating object instances without using new keyword, 185–189
 - creating Scala object from JSON string, 500
 - creating static members with companion objects, 180–182
 - determining class of, 174
 - generating JSON string from Scala classes
 - containing collections, 495–500
 - launching an application with, 176
 - providing named constructor parameters for, 115
 - putting common code in package objects, 182–185
- onComplete method, futures, 438, 443
- onFailure method, futures, 439, 443
- onSuccess method, futures, 439, 443
- open files, 377
- operators, methods versus, in Scala, 41
- Option
 - converting a null into, 657
 - declaring fields not initially populated, 94
 - declaring object as before using in try/catch/finally block, 93
 - defining uninitialized var fields as, 123
 - get and isDefined methods, 383
 - initializing var fields with, not null, 654
 - returned by findFirstIn method, handling, 19
- Option/Some/None pattern, 658–666
 - eliminating null values from code, 39
 - getting value from an Option, 659
 - not using get methods with Option, 665
 - returning file contents with fromFile and using methods, 379

- returning Option from a method, 659
 - throwing exceptions in string conversion to integer, 35
 - using match expressions with, 653
 - using Option with other frameworks, 661
 - using Option with Scala collections, 660
 - using to initialize val and var fields, 135
 - Oracle Byte Streams tutorial, 382
 - Ordered trait, 315
 - mixing in with type to be sorted, 317
 - mixing with Array, 338
 - Ordering trait, 315
 - Array holding elements with implicit Ordering, sorting, 338
 - sorted method and, 316
 - orElse method, PartialFunction trait, 240
 - ORM (object-relational mapping) tools, 527
 - Outer.breakable, 69
 - OutOfMemoryError errors (Java), 332
 - output, console, writing, 488
 - overriding var and val fields, 206
- ## P
- package objects
 - process package object, 400
 - putting common code in, 182–185
 - putting implicit class in, 26
 - root package object, 184
 - package scope, 150
 - package-level control for methods, 150
 - package.scala file, 182
 - packages
 - creating, 193
 - import statements, 200
 - importing everything from a package, 194
 - unable to use in REPL, 460
 - packaging, 191–193
 - Scala project, using SBT, 574
 - par method, 256, 449
 - parallel collections, 448
 - converting regular collections to, 449
 - parallel collection classes, 450
 - performance, 451
 - situations for use, 450
 - ParseException, 486
 - parseInt method, Integer class, 33
 - partial functions, 238
 - chaining, 240
 - example of use with collections, 242
 - PartialFunction trait, explained, 239
 - partially applied functions, 222
 - using, 234
 - partition method, 293
 - ParVector class, 449
 - :paste command (REPL), 459
 - PATH environment variable, 407
 - pattern matching, 651
 - (see also match expressions)
 - using in match expressions, 79–86
 - patterns, 82
 - adding variables to, 84–85
 - using to match case classes and objects, 86
 - performance
 - collections, 252, 261–264
 - ArrayBuffer methods, 270
 - maps and sets, 263
 - performance characteristic keys, 261
 - sequential collections, 262
 - parallel collections, 451
 - views on collections and, 308
 - piping commands together, 401
 - and writing output to a file, 402
 - methods for external commands, 409
 - Play Framework, 491
 - Anorm database library, using Option, 661
 - creating GET request web service, 521–524
 - web service, POSTing JSON data to, 524
 - play.api.libs.json.Format object, 522
 - PoisonPill message, 427
 - sending, 428
 - stopping an actor by sending (example), 429
 - pop method (Stack), 370
 - PositiveInfinity, 45
 - POST requests
 - accessing request data with Scalatra, 510–514
 - Play Framework web service allowing users to send JSON data, 524
 - sending JSON data to POST URL, 518
 - PostgreSQL, connecting to with JDBC, 530
 - postRestart method, Actor object, 422, 424, 435
 - postStop method, Actor object, 422, 424
 - precision, floating-point number comparisons, 41
 - Predef object, 192
 - classOf method, 174
 - immutable Maps, 341
 - predicate, defined, 246

- preRestart method, Actor object, 422, 424
- preStart method, Actor object, 419, 422, 424
- print method, 485
- print options, scalac, 467
- printAll method, annotated with @varargs, 561
- printf method, 485
- printf style format specifiers
 - common, 12
 - using inside strings, 11
- println method, if expression embedded in, 71
- PrintWriter class, 381
- private access modifier, 101
 - accessing private members, class and companion object, 181
 - adding to val or var field, 106
 - defining object-private fields, 120
 - making primary constructor private, 112
 - object-private scope, 148
 - preventing generation of getter and setter methods, 119
 - private scope, 149
 - private var constructor parameter, 117
- Process object, 395
 - executing external commands, 397
 - setting environment variables when calling apply method, 407
- process package, 394
- ProcessBuilder trait, 403
- processed string literals, 10
- processes
 - building pipeline of external commands, 401
 - executing external commands, 394
 - executing external commands and using STDOUT, 397–399
 - handling STDOUT and STDERR for external commands, 399
 - handling wildcard characters in external commands, 405
 - methods to execute external commands, 408
 - redirecting STDIN and STDOUT of external commands, 402
 - running in a different directory, 406
 - using AND (&&) and OR (||) commands with, 404
 - using Process and ProcessBuilder commands in REPL, 463
- ProcessIO class, 400
- ProcessLogger class, 399
- project directory structure, creating for SBT, 570–574
- Project Lambda (Java 8), 233
- project/Build.scala file, 593
 - RootProject, 593
 - using instead of build.sbt, 602–604
- Props class, 418, 425
- protected scope, 149
- public scope, 151
- pure functions, 636–644
 - and side effects, 162
 - defined, 637
 - examples of, 638
 - Java approach, problems with, 639
 - fixing the problems, 640
 - methods with no side effects, 644
 - referential transparency, 637
 - write 80% of your application as pure functions advice, 639
- push method (Stack), 369
- put method, adding or replacing element on mutable Map, 346

Q

- queues, 367
 - adding elements to mutable Queue, 367
 - creating a mutable Queue, 367
 - removing elements from, 368
- quickSort method, sorting an Array, 338

R

- r method, converting String to a Regex, 18, 542
- Random class, 45
- random numbers, generating, 45–47
- random-length ranges, 48
- range method, 310, 372
- ranges, 371–372
 - as generators in for loops, 61
 - combining with map method, 372
 - converting Range to Array or List, 48
 - creating and populating sequences, 371
 - creating range of numbers, 47
 - random-length ranges, 48
 - creating Range with and without a view, 306
 - populating a collection with a Range, 309
 - collections created directly from Range, 310
 - using range method, 310

- random-length, 46
 - using as for loop generator, 56, 372
 - using in for loops, 274
- raw interpolator, 11
- read* methods (Console.read*), 485
- Read-Eval-Print-Loop (see REPL)
- readLine method, 485, 488
- receive method, Actor object, 414, 416, 422
- recover combinator method, 444
- recoverWith combinator method, 444
- recursive algorithms, 70
- reduce method, 298, 299
- reduceLeft method, 295
 - reduceRight versus, 298
 - showing each step in process, 296
 - using with keysIterator on maps, 361
 - using with valuesIterator, 361
- reduceLeftOption method, 299
- reduceRight method, 295
 - reduceLeft versus, 298
- reduceRightOption method, 299
- referential transparency, 636
- reflection
 - accessing methods of String class, 174
 - Gson library working via, 498
- Regex object
 - converting String to, 542
 - creating, 19
 - creating by invoking .r method on a String, 18
- regular expressions
 - extracting parts of String that match patterns, 22
 - finding patterns in strings, 18
 - replacing patterns in strings, 21
 - using pattern matching in match expressions, 79–86
 - using with split method, String objects, 8
 - writing to match expected user input, 487
- relative complement of two sets, 303
- remove method
 - deleting elements from ArrayBuffer, 336
 - deleting elements from ListBuffer, 329
 - removing items from MongoDB collection, 546
 - using with mutable Set, 364
- renaming clause, importing members, 195
- REPL (Read-Eval-Print-Loop), xxi, 1, 453
 - adding JAR files and classes to classpath, 461
 - command-line options, 457
 - Console.read* methods, 485
 - customizing with startup file, 464
 - deprecation and feature warnings, 457
 - getting started with, 454
 - pasting and loading blocks of code into, 459–461
 - :load command, 460
 - :paste command, 459
 - greedy nature of REPL, 460
 - running shell command from, 462–464
 - tab completion, 455
- replaceAll method, String, 7
- replaceAllIn method, String, 21
- replaceFirst method, String, 21
- replaceFirstIn method, using with a Regex, 21
- repositories, 584
 - telling SBT how to find, 595, 601
 - using Maven repository library with SBT, 604
- request.body method, Scalatra, 511
- resolvers key in build.sbt file, 595, 601
- response.addHeader method, Scalatra, 512
- RESTful URL, 514
- RESTful web services, 521
- retain method
 - filtering elements in mutable maps, 355
 - using with mutable Map, 346
 - using with mutable Set, 364
- return type
 - annotating for implicit methods, 28
 - declaring for function literals, 220
- rich versions of numeric classes, extending Ordered trait, 315
- RichDouble class, 37
- RichInt class, to and until methods, 48
- root package object, 184
- running Scala project in SBT, 575

S

- s string interpolator, 9, 10
- save method
 - Casbah MongoCollection class, 537
- SBT (Simple Build Tool), 454, 569–609
 - compiling, running, and packaging a Scala project, 574–579
 - common SBT commands, 576
 - continuous compiling, 577

- using last to get info on last command, 578
- connecting to database with Spring Framework project, 531
- connecting to MySQL with JDBC, 529
- controlling version of managed dependency to use, 584
- creating project directory structure, 570–574
 - using Giter8, 571
 - using shell script, 570
- creating project with subprojects, 586–588
- deploying single, executable JAR file, 597–601
- generating project API documentation, 590
- generating Scaladoc documentation with, 478
- generating Scaladoc documentation with sbt doc command, 474
- including libraries into a project, 471
- managing dependencies with, 581–584
- publishing your library, 601
- resolving problems by getting stack trace, 596
- running tests with ScalaTest, 579–581
- setting log level, 597
- specifying main class to run, 591
- telling SBT how to find repository, 595
- using Build.scala file instead of build.sbt, 602–604
- using GitHub projects as project dependencies, 593
- using Maven repository library with, 604
- using with Eclipse, 588
 - .classpath file, 589
 - .project file, 589
 - updates to SBT build definition files, 590
- using with Giter8 to create Scalatra project, 504
- sbt-assembly plug-in, 594
 - creating single, executable JAR file, 598
- sbtclipse plug-in, 588
- Scala
 - installing, xx
 - using as scripting language, 480–483
 - versions prior to 2.10
 - defining implicit classes, 27
- Scala Ant Tasks page, 608
- scala command, 457
 - cp or -classpath argument, 461
 - i option, 463
 - running programs with, 465
- Scala Glossary, 212
- Scala IDE for Eclipse project, 588
- Scala language, xiii
 - characteristics of, xiv
- Scala Language Specification
 - for loop translation by compiler, 57, 60
- Scala package object, 184
- Scala Style Guide, 636
 - accessor methods, 161
 - formatting control structures, 62
- Scala Worksheet, 458, 636
- scala.concurrent.duration._ import, 438
- scala.io.Source, 375, 380
 - (see also Source object)
 - fromFile, 387
 - fromString, 387
 - fromURL, 514
 - setting timeout while using, 515
 - simple use of, 514
- scala.math.Ordering, 315
- scala.sys.process package, 394, 463
 - methods for executing external commands, 408
- scala.util.control.Breaks, 65
- scalac
 - print command, 468
 - compiling with, 465
 - print options, using, 467
- Scaladoc
 - generating documentation with, 472–479
 - example documentation, 474
 - hyperlink tags, 478
 - scaladoc command, 474
 - Scaladoc tags, 472, 476
 - using SBT, 478
 - wiki-like tags, 477
- scaladoc command, 472, 474
- ScalaLab, 32
- ScalaTest, running tests with SBT, 579–581
- Scalatra framework, 491
 - accessing POST request data, 510–514
 - accessing web service GET parameters, 509
 - creating web services with, 503–507
 - replacing XML servlet mappings with Scalatra mounts, 507
 - SBT project preconfigured to use, creating with Giter8, 571

- ScalazBuild.scala file, 603
- scanLeft and scanRight methods, 299
- Scanner class (Java), 486
- scripting language, using Scala as, 480–483
- self type, 212
- Seq trait, 252
 - executing external commands, 397
 - executing series of commands in a shell, 401
 - executing system commands, 395
- sequence comprehension, 281
 - (see also for/yield loops)
- sequence patterns, 83
- sequences, 247
 - adapting to use as argument for varargs field, 163
 - adding elements to mutable sequences, 335
 - choosing, 250–252
 - main immutable sequence choices, 251
 - main mutable sequence choices, 251
 - Scala's general-purpose sequential collections, 251
 - defined, 250
 - extracting unique elements from, 300
 - foreach method called on, 271
 - hierarchy in Scala, 248
 - indexed and linear, 248
 - making ArrayBuffer your go to mutable sequence, 268
 - merging sequential collections, 302
 - merging two sequential collections into pairs with zip, 304
 - performance characteristics for sequential collections, 262
 - populating with a Range, 309
 - sorting a sequential collection, 315
 - splitting into subsets, 293
 - using ranges to create and populate, 371
 - Vector class, making your go to immutable sequence, 266
 - walking through elements in, 295–300
- SeqView, 306
- Serializable trait, 389, 556
- serialization, 389
- @SerialVersionUID annotation, 389, 556
- servlet mappings (XML), replacing with Scalatra mounts, 507
- sets, 249, 322
 - adding elements to, 361
 - immutable Set, 363
 - mutable Set, 362
 - common set choices, mutable and immutable, 253
 - converting sequential collections to in order to eliminate duplicates, 300
 - creating a Set from a Range, 310
 - defined, 250
 - deleting elements, 363
 - immutable Set, 365
 - mutable Set, 363
 - performance characteristics, 263
 - relative complement of, getting with diff, 303
 - using sortable sets, 365
- :sh command (REPL), 463
- shell built-in commands versus external commands, 396
- shell scripts, 480–483
 - accessing command-line arguments from, 483
 - first lines of, 480
 - making Scala scripts run faster, 489
 - prompting for input from Scala shell script, 485–489
 - using to create SBT project directory structure, 570
- shells
 - handling wildcard characters in external commands, 405
 - piping capability and, 401
 - running shell command from REPL, 462–464
- Short type, 31
- shutdown method, ActorSystem object, 433
- side effects, 162
 - statements and expressions, 648
 - writing functions or methods without (pure functions), 636–644
- side-effecting operations, 451
- Simple Build Tool (see SBT)
- Singleton object, creating with object keyword, 178
- Singleton pattern, enforcing in Scala, 112
- slice method, 292
- Slick library, 527, 547
 - documentation, 548
- sliding method, 294
- Some and None
 - using flatten method on sequence of, 286
 - using with match expressions, 85

- sortBy method
 - sorting immutable map by key, 357
 - sorting immutable map by value, 357
 - using with Array, 338
- sorted method
 - inability to use on sequence type without implicit Ordering, 316
 - sorting a collection, 315
 - using on Person class after mixing in Ordered trait, 317
 - using with Array, 338
- SortedMap class, 343
- SortedSet class, 365
- Sorting.quickSort method, 338
- sortWith method, 315
 - passing your sorting function into, 316
 - sorting immutable map by value, 357
 - sorting keys of immutable map, 357
 - using anonymous function to sort Person elements, 317
 - using with Array, 338
- Source object, 375, 380
 - fromFile method, 376, 383, 387
 - using with using method, 379
 - variations of, 380
 - fromInputStream method, 515
 - fromString method, 387
 - fromURL method, 514
 - setting timeout while using, 515
 - simple use of, 514
 - getLines method, 376, 384
- span method, 293
- Sphinx-4 project, 172
- Spire project, 32
- split method
 - from Java String class and StringLike class, 9
 - String objects, 8
- splitAt method, 293
- Spring Framework, 527
 - connecting to database with, 530
 - using in Scala application, 557–560
- Squeryl library, 527
- Stack Overflow
 - finding Scala libraries through, 472
 - Rahul Phulore's post on infix notation, 48
- stack trace, getting for SBT, 596
- “Stackable Trait Pattern” article, 209
- stacks, 369
 - creating a mutable Stack, 369
 - emptying mutable Stack with clear, 370
 - immutable Stack, 371
 - inspecting, 370
 - peeking at next element of mutable Stack with top, 370
 - pushing elements onto mutable Stack with push, 369
 - taking elements off mutable Stack with pop, 370
- statements versus expressions, 648
- static imports, 197
- static members, creating with companion objects, 180–182
- STDERR, 399
- STDIN
 - methods for redirecting, 409
 - redirecting for external commands, 402
- STDOUT
 - handling for external commands, 399
 - methods for redirecting, 409
 - redirecting for external commands, 402
 - using in Scala program, 397–399
- stop method, actors, 427
 - stopping actor using ActorSystem stop method, 429
- Stream class, 331
 - caution with calls to non-transformer methods, 332
 - lazy collection, 255
 - transformer methods called on, 331
 - zip method, 273
- strict and lazy collections, 250, 254, 306
- String class
 - java.lang.String
 - equalsIgnoreCase method, 5
 - split method, 9
 - java.lang.String and Scala String, 1
- string interpolation, 9
 - creating your own interpolator, 12
 - f string interpolator, 11
 - functions for, 10
 - raw interpolator, 11
- string literals, using expressions in, 10
- StringEntity object, 513
- StringLike class, 3
 - split method, 9
- StringOps class, 2
 - capitalize method, 3

- supertype and type hierarchy information for, 3
- viewing methods available from, in REPL, 456
- strings, 1–29
 - accessing a character in, 24
 - adding methods to closed String class, 3
 - adding your own methods to String class, 25–29
 - building from collections with mkString, 256
 - case statement matching, 76
 - converting a String to Scala numeric type, 32–36
 - handling base and radix, 33
 - converting collection to a String with mkString, 318
 - extracting parts that match regular expression patterns, 22
 - finding regular expression patterns in, 18
 - complete example, 20
 - flattening list of strings into list of characters, 286
 - foreach method called on, 2, 271
 - implicit conversions, 2, 33
 - iterating over characters, Scala versus Java approach, 15
 - Java String as Scala String, 1
 - multiline, creating, 6
 - operations on, 2
 - pretending that a String is a File, 387
 - processing one character at a time, 13
 - replacing regular expression patterns in, 21
 - sorting list of strings by string length, 316
 - splitting, 8, 385
 - substituting variables into, 9
 - testing equality, 4
 - using methods available on sequential collections, 2
- stripMargin method, 7
- structural types, 214, 618
- subclasses, constructor parameters in, 125
- subprojects in SBT, 586–588
- Success class, 633, 662
- Suereth, Joshua, 378
- sumToMax function, 69
- superclasses
 - calling methods in, 152–154, 208
- super[traitName].methodName syntax, 154
- @switch annotation, 73

- switch statements (Java), 54
 - match expressions as replacement for, 651
 - using match expressions like, 72–76
- SynchronizedMap trait, 253
- SynchronizedSet trait, 254

T

- tab completion (REPL), 455
- tableswitch, compiling match expression to, 73, 74
- tabulate method, 372
- tail method, 293
- tail recursion, 70
- @tailrec annotation, 70
- take method, 292
 - using on immutable Set, 365
 - using with maps, 356
- takeRight method, 292
- takeWhile method, 292
- ternary operator, using if construct as, 71
- testing and debugging
 - running tests with SBT and ScalaTest, 579–581
- text files
 - opening and reading, 375–381
 - processing every character in, 383
 - writing, 381
- this method, as constructor method call, 108
- this reference, returning from fluent style methods, 168
- this.type as return type of fluent style methods, 167
- thread pools, 437
- Thread.sleep method, 377, 437
- threads, 411
- @throws annotation, 34, 92, 165, 554
- timer, creating, 630
- to method
 - Int class, 47
 - RichInt class, 48
- to* methods on numeric types, converting between numeric types, 36
- toArray method, using with findAllIn method call, 18
- toList method
 - collection types converting their data to a List, 323
 - converting array of directory names, 393
- top method, 370

- toString method
 - case class implementation of, 137
 - using on a collection, 319
- traits, 203–216
 - abstract classes versus, 129
 - adding to object instance, 215
 - calling a method on, 152–154
 - selecting the trait, 153
 - defining properties in, 131
 - ensuring trait can only be added to type with
 - method of given signature, 213
 - extending another trait, 204
 - FunctionN, 222
 - inherited by Vector class, 247
 - limiting classes using a trait by inheritance, 209
 - marking for use only by subclasses of certain type, 211
 - sequence, commonly used in library APIs, 252
 - using abstract and concrete fields in, 206
 - using as interface, 203
 - using as simple mixins, 208
 - using like abstract class, 207
 - wrapping with implementations to use from Java, 565–568
- transform method, Map
 - creating new map from existing map, 351
 - using with mutable map, 355
- transformer methods, 254, 255
 - calling on Stream, 331
 - used with lazy collections, 306
 - views on collections and, 308
- transformer, => symbol, 218, 220
- Traversable trait, 247
 - looping over Traversable types with for loop, 272
 - methods common to all collections via, 256
- TraversableOnce trait, 336
- tree command, 570, 587
- TreeSet class (Java), 366
- trimming strings, 8
- Try, Success, and Failure classes, 662
 - Scala 2.10 Try classes, 633
 - writing your own, 631
- try/catch/finally
 - closing files and resources with try/finally, 377
 - declaring a variable before using in, 92–95
 - handling exceptions when opening files, 379
 - matching one or more exceptions with try/catch, 91
 - reading and writing binary files, 382
 - using match expressions in try/catch, 651
- tuple patterns, in match expressions, 80, 83
- Tuple2 class, 313
- Tuple3 class, 456
- tuples, 254, 312–315, 457
 - converting to collections, 314
 - returned by zipWithIndex and zip when used on sequences, 277
 - returning multiple objects from a method, 159
 - returning sequence of Tuple2 elements using map method with a Range, 311
 - treating as collection by creating iterator, 314
 - using Tuple syntax to access key and values fields in maps, 351
 - working with, 160
- Twitter, Effective Scala document, 636
- type ascription, 38
- type class, 627–630
- type erasure, 81
- type patterns, 83
- types, 611–634
 - assigning manually to an Array, 333
 - bounds, 612
 - building functionality with, 630–634
 - creating a timer, 630
 - writing your own Try classes, 631–634
 - creating a collection whose elements are all of some base type, 624–627
 - creating a method that takes generic type, 617
 - creating classes that use generic types, 614–617
 - type parameter symbols, 616
 - declaring a type when creating collections, 264
 - inferred by Scala for most expressions, 271
 - making immutable collections covariant, 622
 - selectively adding new behavior to closed model, 627–630
 - type constraints, 613
 - type examples in other chapters, 613
 - type variance, 611
- Typesafe
 - Akka actor library, 411

- recommendation to think of actors as people, 412
- sbteclipse plugin, 577
- Slick library, 527, 547, 548
- Typesafe repository in build.sbt file, 414, 471, 578
- Zinc, incremental compiler based on SBT, 578
- typing
 - making mutable collections invariant, 620
 - using Duck Typing (structural types), 618

U

- unapply method, case classes, 137
- Unfiltered, 491
- UnhandledMessage, 416, 416
- Uniform Access Principle, 642
- Unit, 224, 225
 - functions or methods returning, calling from foreach method, 271
 - return type for function returning nothing, 227
- Unix command line, 396
- Unix commands, equivalency of Scala commands to, 401
- Unix systems
 - setting Java properties when starting Scala interpreter, 457
 - using Scala as scripting language, 480–483
- unmanaged dependencies, 582
- until method
 - RichInt class, 48
 - using in creating ranges, 48
- unzip method, 294, 305
- updating elements, mutable maps, 346
- upper bound, 624, 626
- URL class (Java), 403, 515
- URLs
 - getting URL headers, 519
 - MySQL JDBC, 529
 - sending JSON data to POST URL, 518
- using method, 378
 - using with Source.fromFile and returning file contents as List, 379
- utility classes, 113

V

- val + mutable and var + immutable, 645

- val fields
 - abstract, in abstract classes (or traits), 131
 - how they work, 132
 - adding private keyword to, 106
 - case class constructor parameters, 107, 137
 - concrete, in abstract classes, 133
 - constructor fields defined as, 106
 - using in traits, 206
- values method, Map
 - getting values from map as an Iterable, 352
 - iterating over values in maps, 351
- valuesIterator method, Map, 352, 354
 - using with max method, 361
 - using with reduceLeft method, 361
- var fields
 - abstract, in abstract classes (or traits), 131
 - how they work, 132
 - adding private keyword to, 106
 - case class constructor parameters, 137
 - concrete, in abstract classes, 133, 134
 - constructor parameters, 105
 - declaring subclass constructor without var declarations, 126
 - initializing with Option, not null, 654
 - private var constructor parameter, 117
 - setting uninitialized var field types in a class, 122
 - using in traits, 206
- @varargs annotation, 560
- varargs fields, 163–165
 - adapting a sequence to use as argument with `_*` operator, 163
 - annotating varargs methods, 560
 - defining in method declaration, 163
 - with or without arguments supplied, 164
- variable patterns, in match expressions, 83
- variable substitution into strings, 9
- variable-binding patterns, 84–85
- variables
 - assigning to result of match expression, 77
 - assigning type to, 38
 - declaring before using in try/catch/finally block, 92–95
 - declaring immutable map as a var, 348
 - function value, 235
 - function variable, assigning existing function or method to, 222
 - mutable and immutable, 266

- mutable, with immutable collections, 265, 340
- preferring immutable variables, 644
 - benefits to using immutable variables, 645
- using functions as, 219–223
- Vector class
 - creating a Vector and converting to parallel collection, 449
 - making your go to immutable sequence, 266
 - traits inherited by, 247
- view bound, 626
- view method, 256
 - calling before zipWithIndex, 273, 277
- views
 - creating lazy view on a collection, 306–309
 - use cases for views, 308

W

- :warning command (REPL), 458
- watch method, context object, 433, 434
- web services, 491–526
 - accessing POST request data with Scalatra, 510–514
 - creating GET request web service with Play Framework, 521–524
 - creating simple GET request client, 514
 - creating with Scalatra, 503–507
 - Play Framework, POSTing JSON data to, 524
 - Scalatra, accessing GET parameters, 509
- web.xml file, 506
 - replacing XML servlet mappings with Scalatra mounts, 507
- while loops, 383
- whitespace around commands and arguments, 395
- wiki-style markup tags, 472, 477
- wildcard characters in external commands, 404
 - handling, 405

- withDefaultValue method, creating map with, 350
- withFilter method
 - for loop with guard, translation by compiler, 59
 - for loop with yield expression and guard, 57
- WriteResult and WriteConcern classes, MongoDB, 538

X

- XML
 - getting classes for objects, 175
 - searching with \ and \\ methods, 503
 - web.xml file, 506
 - replacing servlet mappings with Scalatra mounts, 507
- Xprint:all compiler option, scalac, 58, 60, 468
- Xprint:parse compiler option, scalac, 58, 60, 467

Y

- yield statement, using with for loop, 63
 - (see also for comprehension; for/yield loops)
 - processing requiring multiple lines of code, 274

Z

- zip method
 - merging two sequential collections into pairs with, 304
 - using with a Stream, 273
 - creating a loop counter, 277
- zipWithIndex method
 - using to create loop counter, 55, 273, 276, 386

About the Author

Alvin took the circuitous route to software development. He managed to get a degree in Aerospace Engineering from Texas A&M University, while all he was really trying to do was play baseball. Once he became a practicing engineer, he realized he liked software and programming more than engineering. So, in approximate order, he taught himself Fortran, C, Unix and network administration, sed, awk, Perl, Java, Python, Ruby, JRuby, Groovy, PHP, and Scala. During this process, he started a software consulting firm, grew it to 15 people, sold it, and moved to Alaska for a few years. After returning to the “Lower 48,” he self-published two books (*How I Sold My Business: A Personal Diary* and *Zen and the Art of Consulting*). He also created DevDaily.com, which receives millions of page views every year, started a new software consulting business, [Valley Programming](#), and started a nonprofit organization named [Zen Foundation](#).

Colophon

The animal on the cover of *Scala Cookbook* is a long-beaked echidna (*Zaglossus bruijnii*, *Z. bartoni*, and *Z. attenboroughi*), a genus of three mammal species found only on the island of New Guinea. Weighing up to 35 pounds, long-beaked echidnas are nocturnal insectivores that prefer to live in forests at higher altitudes.

The first specimen was found in 1961 on New Guinea’s Cyclops Mountains, and the entire species was thought to be extinct in that area until evidence of their activity was found in 2007. According to data collected in 1982, only 1.6 echidnas existed per square kilometer of suitable habitat across New Guinea, adding up to a total of 300,000 individuals. Since then, that number has dropped significantly due to habitat loss as large areas are exploited for farming, logging, and mining. Hunting also remains a large problem since the long-beaked echidna is considered a delicacy to locals in Papua New Guinea. The low population numbers and rapid destruction of habitat make the long-beaked echidna an endangered species, while the short-beaked variety fares slightly better in both New Guinea and Australia.

The echidna is classified as a “monotreme,” or a mammal that lays eggs. The mother holds one egg at a time in her body, providing it with nutrients and a place to live after it hatches. The only surviving monotremes are the four species of echidna and the platypus. All of these mammals are native to Australia and New Guinea, although there is evidence that they were once more widespread. With origins in the Jurassic era some 60 million years ago, monotremes offer evidence of mammal evolution away from reptilian forms of reproduction.

Instead of having teeth, echidnas' tongues are covered in spikes that help draw earthworms and ants into the mouth. The entire body is also covered in fur and spikes that are used for protection; much like a hedgehog, echidnas can curl up into a spiny ball when threatened. Although very little echidna behavior has been observed in the wild, they are believed to be solitary creatures; the short-beaked echidna displays little evidence of grooming, aggression, courting, or maternal behavior. In captivity, these creatures can live up to 30 years.

The cover image is from *Cassell's Natural History*. The cover font is Adobe ITC Garmond. The text font is Adobe Minion Pro; the heading font is Adobe Myriad Condensed; and the code font is Dalton Maag's Ubuntu Mono.